

OAL Adversarial Cross-ReviewMaster Template

Version 1.0 | October 1, 2026

PROPOSED

-

NOT CANON

-

NOT ONESTAMP-RATIFIED

Document Title	
Document ID	OAL-XR-TEMPLATE-001
OAL Version	v1.0
Date	October 1, 2026
Prepared By	Microsoft Copilot (Copilot AI)
Prepared For	Nathan Lewis II, Architect of Record, Omniversal Architect Labs
Status	PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. PENDING ARCHITECT REVIEW.
Posture	Adversarial and non-consensus. Disagreement is recorded, not averaged.
Template Origin	Derived from OAL-ECL-XR-PERPLEXITY-001 v0.2 (End Cycle Laboratory Adversarial Cross-Review, September 30, 2026, prepared by Perplexity Computer / Perplexity AI)

Document Title	
Intended Distribution	All OAL AI Collaborators: Claude, Perplexity AI, ChatGPT, Google Gemini, Microsoft Copilot, Grok
Publication Target	GitHub (OAL public repository)

Disclaimer:

Speculative worldbuilding. Nothing in OAL documents models any real person's grief, depression, burnout, disability, or suicidal thinking. In the US, anyone in crisis can call or text

988

.

1. One-Paragraph Summary

This document is the **OAL Adversarial Cross-Review Master Template (OAL-XR-TEMPLATE-001 v1.0)**, a replicable, adversarial, non-consensus cross-review framework designed for use across any Omniversal Architect Labs (OAL) topic, framework, or body of content. It is derived directly from the methodology first demonstrated in **OAL-ECL-XR-PERPLEXITY-001 v0.2** (End Cycle Laboratory Adversarial Cross-Review, September 30, 2026, prepared by Perplexity Computer / Perplexity AI) and is built to be distributed to any AI collaborator — including Claude, Perplexity AI, ChatGPT, Google Gemini, Microsoft Copilot, and Grok — so that each collaborator may produce their own independent cross-review of a designated body of OAL documents. The goal of this framework is emphatically **not**

consensus. The goal is to find holes, expose weaknesses, surface missed questions, challenge untested claims, and establish clear paths to iterative improvement. Every proposed rule, design, or claim inside an OAL research document must be held to a public test it can fail; if it cannot be tested, it is marked **[OPEN]** and reserved for the Architect or for future execution. This template must be filled in with topic-specific content by the preparer before it is issued to any collaborator; all placeholder text appears in **[BOLD BRACKETS]** throughout this document and must be replaced before distribution.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

2. Table of Contents

1.	One-Paragraph Summary
2.	Table of Contents
3.	Versioning Record
4.	Mission Statement and Directive
	4.1 Template Directive
	4.2 Proposed Topic Mission (To Be Completed by Preparer)
	4.3 Dissenting Mission (To Be Preserved, Not Resolved)
5.	Executive Summary and Overview
	5.1 Overview of the Cross-Review Process
	5.2 What the Preparer Will Fill In
	5.3 What the Collaborators Will Produce
6.	How to Use This Template
7.	Source Corpus Registration
8.	Epistemic Label System
9.	Review Method: Seven Questions
10.	Why This Research Matters (Case For and Case Against)
	10.1 The Case For the Topic

	10.2 The Case Against the Topic
11.	Independent Review of Each Source Document
12.	Discrepancy Register
13.	Pilot Execution: Running Exercises Now
	13.1 Pilot Execution Standards
	13.2 Adversarial Authorship Protocol
	13.3 Pilot Exercise Log Template
	13.4 Results Reporting Standard
14.	Evidence Assessment (What Supports and What Does Not Support the Topic)
	14.1 What Supports the Topic
	14.2 What Does Not Support the Topic, or Cuts Against It
15.	Implications
	15.1 For the OAL Project
	15.2 For Virtual Environments
	15.3 For Other Environments
	15.4 Real-World Applicability Outputs
16.	Testability Matrix
17.	Funding Pathways
18.	Open Questions That Must Be Addressed
	18.1 Architect Decision Points
	18.2 Research Questions for the Next Iteration
19.	Insights to Build Into the Next Iteration
20.	Next Iteration Commission Prompt (Copy-Ready)
21.	Appendices
	Appendix A: Exercise Crosswalk Template
	Appendix B: Pilot Methods and Parameters Template
	Appendix C: Sources Template

3. Versioning Record

Field	Value
Title	OAL Adversarial Cross-Review Master Template v1.0
Document ID	OAL-XR-TEMPLATE-001
OAL Version	v1.0
Date	October 1, 2026
Prepared By	Microsoft Copilot (Copilot AI)
Prepared For	Nathan Lewis II, Architect of Record, Omniversal Architect Labs
Status	PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. PENDING ARCHITECT REVIEW.
Posture	Adversarial and non-consensus. Disagreement is recorded, not averaged.
Topic Origin	Derived from OAL-ECL-XR-PERPLEXITY-001 v0.2 (End Cycle Laboratory Adversarial Cross-Review, September 30, 2026, prepared by Perplexity Computer / Perplexity AI)
Purpose	To codify the adversarial cross-review methodology first demonstrated in OAL-ECL-XR-PERPLEXITY-001 v0.2 and make it replicable across any OAL topic, framework, or body of content.
Working Hypothesis	<i>[TO BE FILLED IN BY PREPARER: State the working hypothesis for the specific topic being reviewed. This must be falsifiable.]</i>
Companion Files	<i>[TO BE FILLED IN BY PREPARER: List all companion documents, raw data files, exercise logs, and GitHub repository links for the specific review.]</i>

Epistemic Label Key

The following labels are used throughout all OAL adversarial cross-review documents. Every claim must carry one of these labels. Unlabeled claims are not permitted.

Label	Meaning
[SOURCE]	Stated in one of the reviewed documents or the Architect's source material. The originating document must be named (e.g., [SOURCE: S1]).
[FINDING]	Produced by an executed pilot run; a claim about the model's behavior, not about the world. The exercise must have been actually run, not only designed.
[EVIDENCE]	Supported by an external, cited real-world source. The citation must accompany the label (e.g., [EVIDENCE: Cato Institute, 2023]).
[ASSESSMENT]	This reviewer's adversarial judgment. May not be presented as consensus. Must be clearly distinguished from SOURCE and FINDING.
[PROPOSAL]	New design offered for the next iteration. Must be distinguishable from SOURCE; must not be presented as already decided.
[OPEN]	Not resolvable from available material. Reserved for the Architect or for future testing. The reviewer may not resolve an OPEN claim unilaterally.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

4. Mission Statement and Directive

4.1 Template Directive

This template exists to make OAL research **reviewable, testable, and iteratively improveable** across all topics and all AI collaborators. No collaborator's findings are final. No iteration's output is canon until the Architect ratifies it through the OneStamp process. Every proposed rule, design, or governance mechanism produced by any collaborator must be held to a public test it can fail — and the parameters of that test must be committed before the test is run.

Disagreement is the product of this process, not a problem to be resolved. When two collaborators disagree about a finding, their disagreement is registered in the

Discrepancy Register (Section 12) and assigned a resolution path (TEST, DECISION, or HYGIENE). It is not averaged, softened, or dropped. A review that produces unanimous agreement among six AI collaborators who all read the same source material is evidence of a prompt artifact, not evidence of truth.

The reviewer must apply the same severity to their own prior work that they apply to every other document in the corpus. Self-review is not optional and cannot be softened. If this reviewer produced a document in a prior iteration that is now part of the source corpus, that document is reviewed first and reviewed hardest.

The clinical firewall is always in force. Nothing in any OAL laboratory is used to interpret, model, or advise on real human distress, grief, mental health, or end-of-life decisions. If the topic sits near those areas, a category wall must be stated explicitly in Section 4.2.

4.2 Proposed Topic Mission (To Be Completed by Preparer)

PREPARER INSTRUCTIONS — SECTION 4.2

Write 2–3 sentences stating why the topic being reviewed matters, what the laboratory or framework is trying to accomplish, and what its durable contribution would be if it succeeds. Use the template below. Replace all bracketed placeholders with topic-specific content. Do not leave any placeholder in the issued document.

The **[TOPIC NAME]** laboratory/framework exists to **[PRIMARY PURPOSE]**. It studies **[SCOPE OF INQUIRY]**. It holds every proposed rule to a public test that the rule can fail. It separates **[KEY DISTINCTION 1]** from **[KEY DISTINCTION 2]** from **[KEY DISTINCTION 3]**. It serves **[DOWNSTREAM USE CASE — e.g., a canon process, a product, a tabletop kit]** without letting **[RISK OF CAPTURE — e.g., a preferred outcome, a funder's preference, a collaborator's prior commitment]** decide its findings in advance.

[CATEGORY WALL STATEMENT — Required if the topic touches grief, death, mental health, disability, or suicidal thinking. Example: "This laboratory studies fictional endings in a digital world. It does not model, advise on, or interpret real human end-of-life decisions, grief responses, or clinical conditions. The clinical firewall is in force."]

4.3 Dissenting Mission (To Be Preserved, Not Resolved)

ADVERSARIAL REQUIREMENT

An adversarial cross-review must preserve the strongest argument *against* the proposed topic mission, derived from the most conservative or protective framing found in the source documents. This is not a straw man. It is the best-faith case for a different set of priorities. It is preserved here — not resolved — because a live alternative prevents false consensus from hardening into canon.

[DISSENTING MISSION — To be written by the preparer. Derive this from the most cautious, protective, or conservative reading of the source material. It should name a real risk that the proposed mission does not adequately address. Example: "The most protective reading of this framework is that it should not exist yet — that formalizing a governance structure before any exercise has been run produces a specification that is self-ratifying, not tested. Under this reading, Iteration 1 should produce only a vocabulary and a pre-registration, nothing else. All design decisions should be deferred until at least 100 exercise runs have been reported." [ASSESSMENT]]

The dissenting mission is **not voted down**. It travels with the document. If Iteration 2 evidence resolves it, the Architect closes it with a rationale. Until then, it is open.

5. Executive Summary and Overview

5.1 Overview of the Cross-Review Process

An **adversarial cross-review**, as practiced in OAL, is fundamentally different from a standard review in four respects:

1. **It is non-consensus by design.** The goal is not for six AI collaborators to agree. The goal is for six collaborators to independently apply the same seven questions to the same corpus and report what they each find. Where they agree, the agreement is noted but treated as weak evidence (because all collaborators received overlapping source material). Where they disagree, the disagreement is registered and assigned a resolution path.
2. **Each source document is reviewed independently against the same seven questions in the same order** (Section 9). No document is treated as authoritative over another. The Architect's source material is received as given and kept open; it is not fact-checked against real-world science and is not interpreted without labeling the interpretation as a **[PROPOSAL]**.
3. **Discrepancies are registered, not averaged.** If Document A says a threshold is three steps and Document B says it is five, both positions are recorded in the Discrepancy Register (Section 12) with a resolution path. The reviewer does not choose between them.
4. **Results are evidence-driven.** A claim about the model is a **[FINDING]** only if an exercise was actually executed and results were reported. Predicted lessons written before running are labeled **[ASSESSMENT]** and treated accordingly.

The adversarial cross-review was first demonstrated in OAL-ECL-XR-PERPLEXITY-001 v0.2 and produced the methodological insights codified in Section 19 of this template. Its most important contribution was demonstrating that a reviewer who actually runs exercises produces qualitatively different output from a reviewer who only specifies them.

5.2 What the Preparer Will Fill In

Before this template is issued to any AI collaborator, the preparer (typically the Architect or their designated agent) must supply all of the following. Each item corresponds to a specific section of this document.

5. **Source Corpus List** (Section 7): A complete list of every document in the corpus, with preparer name, claimed stage, and approximate word count.
6. **Topic Mission** (Section 4.2): Two to three sentences explaining why the topic matters and what success looks like, using the template provided.
7. **Topic-Specific Background** (Sections 10.1 and 10.2): The case for and against the topic, with real-world analogs and adversarial challenges filled in for this specific subject.
8. **Source Material**: Any document the collaborators must read before producing their review. All such documents must be registered in Section 7 and listed in the commission prompt (Section 20, Item 2).
9. **Pilot Exercise Parameters** (Section 13): If any exercises are pre-run before distribution, results must be logged in Section 13 using the Pilot Exercise Log Template. If no exercises are pre-run, Section 13.3 must state "No pilot exercises run prior to distribution."
10. **Open Questions from the Architect** (Section 18): Any questions the Architect has already identified that must not be closed by a collaborator without a DECISION from the Architect.

5.3 What the Collaborators Will Produce

Each AI collaborator who receives a filled-in version of this template must independently produce their own version of the cross-review for the named topic. No collaborator may read another collaborator's completed review before submitting their own. Each version must include, at minimum:

- A title, date, version number, preparer name, and status block
- A mission statement or directive (Sections 4.1–4.3, with the topic mission completed)
- An executive summary (Section 5)

- A complete table of contents
- An independent per-document review (Section 11), covering each document in the corpus separately against all seven questions
- A discrepancy register (Section 12) with at least one entry per mandatory discrepancy category checked
- A pilot execution section (Section 13) — either executed exercises with results, or a clear statement that no pilot was run and the reason why
- Implications (Section 15) and Testability matrix (Section 16)
- At least a stub for Funding Pathways (Section 17)
- Open Questions (Section 18) with at least one item in each subsection
- An independence declaration listing every input read

All documents must be delivered in both **PDF and DOCX format**. The document ID for each collaborator's version follows the scheme defined in Section 6: **OAL-[TOPIC CODE]-XR-[AI NAME]-[VERSION]**.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

6. How to Use This Template

The following steps are addressed directly to Nathan Lewis II, Architect of Record, Omniversal Architect Labs.

"Produce your own independent version of this adversarial cross-review for

[TOPIC NAME]

. Read every source document listed in Section 7. Review each document independently against the seven questions in Section 9. Fill in all sections. Deliver in PDF and DOCX. Use the document ID scheme: OAL-

[TOPIC CODE]

-XR-

[AI NAME]

-

[VERSION]

. Do not read any other collaborator's completed review before submitting your own. Your independence declaration must list every input you read."

THE GOLDEN RULE OF OAL ADVERSARIAL CROSS-REVIEW

The goal is not to agree with what we have built. The goal is to find where it fails.

Document ID Naming Convention

Document Type	ID Scheme	Example
Cross-Review Document	OAL - [TOPIC CODE] - XR - [AI NAME] - [VERSION]	OAL - ECL - XR - PERPLEXITY - 001
Iteration / Gate Output	OAL - [TOPIC CODE] - I [ITERATION] - G [GATE] - [AI NAME] - v [X.Y]	OAL - ECL - I2 - G3 - CLAUDE - v1.0
Master Template	OAL - XR - TEMPLATE - [VERSION]	OAL - XR - TEMPLATE - 001
Architect Decision	OAL - [TOPIC CODE] - DEC - [NUMBER]	OAL - ECL - DEC - 003

7. Source Corpus Registration

All source documents must be registered in the table below before the cross-review begins. The registry is the authoritative record of what collaborators were given. Any document not in this registry may not be cited as a **[SOURCE]** in any collaborator's review.

Instructions for Completing This Table

11. List every Iteration document produced by the AI collaboration team on the topic, plus any Architect source material the collaborators are given.
12. In the "Stage Actually Delivered" column, note what the document actually delivers (e.g., vocabulary, design specification, executed exercises, cross-review) regardless of what its title claims.
13. Note any missing documents — for example, a collaborator whose version was not supplied — in the Source Gap Register below.
14. Note any corpus gaps such as secondhand descriptions of missing documents. These must be labeled **[ASSESSMENT]**, not **[SOURCE]**, in any collaborator's review.

#	Document Title	Preparer / AI System	Stage Claimed	Stage Actually Delivered	Approx. Word Count	Gaps or Notes
S1	[DOCUMENT TITLE]	[PREPARER / AI NAME]	[CLAIMED STAGE]	[ACTUAL STAGE DELIVERED]	[~N words]	[NOTES]
S2	[DOCUMENT TITLE]	[PREPARER / AI NAME]	[CLAIMED STAGE]	[ACTUAL STAGE DELIVERED]	[~N words]	[NOTES]
S3	[DOCUMENT TITLE]	[PREPARER / AI NAME]	[CLAIMED STAGE]	[ACTUAL STAGE DELIVERED]	[~N words]	[NOTES]
S4	[DOCUMENT TITLE]	[PREPARER / AI NAME]	[CLAIMED STAGE]	[ACTUAL STAGE DELIVERED]	[~N words]	[NOTES]

#	Document Title	Preparer / AI System	Stage Claimed	Stage Actually Delivered	Approx. Word Count	Gaps or Notes
S5	<i>[DOCUMENT TITLE]</i>	<i>[PREPARER / AI NAME]</i>	<i>[CLAIMED STAGE]</i>	<i>[ACTUAL STAGE DELIVERED]</i>	<i>[~N words]</i>	<i>[NOTES]</i>
S6	<i>[DOCUMENT TITLE]</i>	<i>[PREPARER / AI NAME]</i>	<i>[CLAIMED STAGE]</i>	<i>[ACTUAL STAGE DELIVERED]</i>	<i>[~N words]</i>	<i>[NOTES]</i>

Source Gap Register

Register any documents referenced in the corpus but not supplied to reviewers. A collaborator who characterizes a missing document's position is providing secondhand information. That information must be labeled **[ASSESSMENT]**, not **[SOURCE]**. Gap registration is mandatory before distribution.

Gap ID	Description of Missing Document or Referenced-but-Unsupplied Material
G-01	<i>[DESCRIBE MISSING DOCUMENT: title, who produced it or referenced it, and why it was not supplied to reviewers]</i>
G-02	<i>[DESCRIBE MISSING DOCUMENT]</i>
<i>Add rows as needed</i>	<i>No gaps in this corpus: write "None identified as of [DATE]."</i>

8. Epistemic Label System

The label system governs every claim in every OAL adversarial cross-review document. It ensures that readers can trace each claim to its origin and judge it accordingly. The system is non-negotiable: every claim carries exactly one label.

Mixing labels (e.g., presenting an ASSESSMENT as though it were a SOURCE) is a quality-gate disqualification.

Label	What Qualifies	What Disqualifies	Example Sentence
[SOURCE]	Stated verbatim or nearly verbatim in one of the reviewed documents or the Architect's material. The originating document must be named.	A paraphrase that changes the meaning; a summary that aggregates multiple statements; an inference drawn from a SOURCE.	<i>"Claude's document states that an overseer is not assumed to be necessary under baseline conditions. [SOURCE: S1]"</i>
[FINDING]	Produced by an executed pilot run. The exercise must have been actually run, the seed must be reported, and results (mean, 5th-95th percentile) must be reported.	Predicted lessons written before running; speculated outcomes; designs described but not executed; exercises labeled "we could run" without running.	<i>"In 1,000 runs using seed 42, the immediate trajectory produced 100% non-survival among non-migrating entities. [FINDING: P1]"</i>
[EVIDENCE]	Supported by an external, cited real-world source. The citation must accompany the label. Used only for analogy to analogous real-world systems, never to interpret Architect source material.	Informal analogies without citation; claims about the real world that are not cited; use of real-world evidence to modify or reinterpret the Architect's canon material.	<i>"US sunset review laws mostly resulted in continuation with modifications, not termination. [EVIDENCE: Cato Institute, Regulatory Sunset Evaluations, 2023]"</i>
[ASSESSMENT]	This reviewer's adversarial judgment about a document, a claim, or a design. Must be clearly labeled; may not be presented as consensus or as fact.	Presenting an ASSESSMENT as though it were a FINDING or SOURCE; omitting the label because the reviewer is confident in the judgment; averaging two conflicting ASSESSMENTS into a single claim.	<i>"The 16-stage lifecycle is a premise dressed as a finding; it was not derived from executed exercises. [ASSESSMENT]"</i>
[PROPOSAL]	New design, rule, or structure offered for the next iteration. Must be distinguishable from SOURCE. Must not	A PROPOSAL that restates a SOURCE without change; a PROPOSAL presented as though it has already been	<i>"Iteration 2 should require generator and detector separation for any indistinguishability</i>

Label	What Qualifies	What Disqualifies	Example Sentence
	be presented as already decided.	approved; a PROPOSAL that contradicts an Architect DECISION without noting the conflict.	<i>exercise.</i> <i>[PROPOSAL] "</i>
[OPEN]	A question or claim that cannot be resolved from the available material. Reserved for the Architect or for future testing. The reviewer notes why it is unresolvable.	Using OPEN to defer a claim the reviewer simply does not want to address; closing an OPEN claim without a DECISION from the Architect; treating an OPEN as resolved by subsequent speculation.	<i>"Whether the entity in the field scene retains moral standing after the archive substitution event is not resolved here. [OPEN] "</i>

MANDATORY RULE

Every claim in any OAL adversarial cross-review document must carry one of the six labels above. Unlabeled claims are not permitted. Documents submitted with unlabeled claims fail the quality gate and must be revised before acceptance.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

9. Review Method: Seven Questions

Every source document in the corpus is reviewed against the same seven questions in the same order. The preparer may add topic-specific questions (numbered 8 and above) but **may not remove or reorder questions 1 through 7**. The seven questions constitute the invariant core of the OAL adversarial cross-review method.

Q#	Question	What the Reviewer Is Looking For
Q1	What does this document contribute that no other document in the corpus does?	Distinctive contributions: a new exercise design, a new vocabulary term with a precise definition, a result no other document reports, a challenge no other document raises. If the document duplicates what another document already does, that is noted.
Q2	What does it get right?	The strongest findings or designs in the document. These must be labeled [SOURCE] if drawn from the document and [FINDING] if drawn from executed exercises. The reviewer must identify at least one thing the document gets right, even if the overall assessment is negative.
Q3	Where is it internally inconsistent?	Self-contradictions within the document: a term defined one way in Section 2 and a different way in Section 7; a threshold named but not defined; a status label in the header that contradicts the status label in the body; a finding in the conclusion that is not supported by results in the body.
Q4	Where does it assume what it claims to test?	Circular premises: the document prescribes the answer in the setup and then reports the answer as a result; predicted lessons written before any exercise was run; hypotheses framed to be unfalsifiable; governance designs that presuppose the institution they are designed to govern.
Q5	How does it treat the Architect's source material?	Does it keep open questions open? Does it quote verbatim with qualifiers when citing the Architect's material? Does it interpret ambiguous source material without labeling the interpretation as a

Q#	Question	What the Reviewer Is Looking For
		[PROPOSAL] ? Does it claim the Architect's material says something it does not say?
Q6	Can it be executed as written?	If the document specifies exercises: Are the parameters defined? Is there a baseline (no-intervention) condition? Is there a control branch? Is there a fixed random seed or equivalent? Could a different team reproduce the exercise from the description alone?
Q7	What should the next iteration keep, change, or quarantine from this document?	A short, concrete list: what iteration 2 should carry forward unchanged; what it should revise (with the specific revision labeled [PROPOSAL]); and what it should quarantine — set aside pending more evidence — labeled [OPEN] .

Methodological Rules

15. **Rule 1 — Convergence is weak evidence.** Agreement among AI collaborators is weak evidence that the convergent claim is correct, because all collaborators received overlapping source material and were given overlapping prompts. Convergence may reflect the prompt, not the problem. Convergence must be noted and its source investigated, not celebrated.
16. **Rule 2 — Findings require execution.** A claim is counted as a **[FINDING]** only if something was actually executed, a seed was reported, and results were reported. A design that describes what would happen if an exercise were run is not a finding.
17. **Rule 3 — Real-world material is analogy, not authority.** Real-world evidence is brought in only to inform analogies and to supply evidence about analogous systems. It is never used to fact-check the Architect's source material, and it never edits or reinterprets canon content. Real-world material carries the label **[EVIDENCE]** and a citation.

18. **Rule 4 — Self-review is mandatory and unsoftened.** The preparer reviews their own prior work first and with the same severity as all other documents. The self-review may not be omitted, abbreviated, or softened. If the preparer produced the only document in the corpus, they still review it against all seven questions.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

10. Why This Research Matters

10.1 The Case For the Topic

PREPARER INSTRUCTIONS — SECTION 10.1

Articulate 4-6 bullet points explaining why the topic is worth studying. For each bullet: (a) name a real-world analog or parallel, with a citation where possible; (b) explain the connection to the OAL topic; and (c) label the strength of the connection using the standard labels. Replace all placeholders.

- ***[REAL-WORLD ANALOG 1] — [EXPLAIN CONNECTION TO THE OAL TOPIC. WHY DOES THIS ANALOG MAKE THE TOPIC MORE CREDIBLE OR MORE URGENT?]*** [EVIDENCE: CITATION]
- ***[REAL-WORLD ANALOG 2] — [EXPLAIN CONNECTION]*** [EVIDENCE: CITATION]
- ***[REAL-WORLD ANALOG 3] — [EXPLAIN CONNECTION]*** [ASSESSMENT: Strength of analogy – Strong / Moderate / Weak / Contested]
- ***[CREATIVE OR SPECULATIVE PARALLEL] — [EXPLAIN CONNECTION]*** [ASSESSMENT]

- ***[DOWNSTREAM USE CASE — who would use the findings if the topic succeeds?]*** — ***[EXPLAIN]*** ***[PROPOSAL]***
- ***[ADDITIONAL CASE FOR, IF NEEDED]***

10.2 The Case Against the Topic

An adversarial review must dedicate equal space to the strongest arguments against building or continuing the framework. The following six challenge types apply to *every* OAL topic. The preparer must address each one specifically for the topic at hand before the template is distributed. These are not rhetorical exercises: they are genuine tests. If any of them cannot be answered, that unanswered challenge belongs in Section 18 as an **[OPEN]** item.

19. **Is this a solution looking for a problem?** Does the topic have defined problems that the framework actually addresses? Or does it assume the problem and proceed directly to governance design? A framework that specifies a 33,000-word review process before any exercise has been run may be a solution that precedes its problem. ***[PREPARER RESPONSE — address this challenge specifically for the topic]***
20. **Can the fiction or model validate itself?** Every result in a simulated world is partly an output of the rules its author wrote. If the author of the rules also runs the exercises, the exercises may confirm the rules rather than test them. What structural separation exists between the rule-maker and the result-reporter in this topic's framework? ***[PREPARER RESPONSE]***
21. **Could AI convergence look like discovery?** When multiple AI systems agree on a finding, is that agreement evidence about the topic, or is it an artifact of overlapping prompts and overlapping training data? Has this topic's research taken any deliberate steps to measure prompt-driven convergence separately from problem-driven convergence? ***[PREPARER RESPONSE]***
22. **Could the procedure become the harm?** Could the review or governance process itself trap entities, prolong harm, or become a postponement device? A mandatory multi-stage review that requires evidence before any action can

be taken may, in practice, function as a prohibition on action. What is the stop rule if the process itself causes harm? ***[PREPARER RESPONSE]***

23. **Does the subject matter carry real-world sensitivity?** Does the topic sit close to grief, death, mental health, disability, or other areas requiring clinical firewalls? If so, has the clinical firewall been stated explicitly, and does it apply to all collaborators? Is there a plan for what happens if a collaborator violates the firewall? ***[PREPARER RESPONSE]***

24. **Is a large specification a sign of rigor or a sign of risk?** Large specifications that are never executed are evidence of the postponement risk, not evidence of rigor. A 33,000-word document that produces no executed exercises is a liability, not an asset. How many exercises in the current corpus have actually been run? ***[PREPARER RESPONSE — report the actual count of executed exercises in the corpus]***

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

11. Independent Review of Each Source Document

Reproduce the subsection below once for each document in the Source Corpus Registry (Section 7). Label subsections S1, S2, S3... to match the registry. Reviews are conducted against the Seven Questions (Section 9). The self-review of any prior work by this reviewer appears first.

SUBSECTION TEMPLATE — Repeat N Times (Once Per Source Document)

S[N]. [PREPARER AI NAME]: [DOCUMENT TITLE]

What it is: *[1-2 sentence description of the document's scope, stage, and approximate size. Example: "A 4,200-word Iteration 1 cross-review*

covering vocabulary, a lifecycle model, and three unexecuted exercise designs for the ECL framework."]

Distinctive Contributions (Q1):

- *[What this document contributes that no other document in the corpus does. Be specific. Label each bullet SOURCE or ASSESSMENT.]*
- *[Second distinctive contribution, if any]*
- *[Or: "No distinctive contribution identified — this document duplicates content from S[X]." [ASSESSMENT]]*

What It Gets Right (Q2):

- *[Strongest finding or design in the document. Label SOURCE or FINDING.]*
- *[Second strength, if any]*

Adversarial Findings (Q3-Q6) — labeled [ASSESSMENT]:

25. *[Internal inconsistency found. Example: "The document defines [TERM] as X in Section 2 and Y in Section 5. These are not reconciled. [ASSESSMENT]"]*
26. *[Circular premise found. Example: "The exercise setup prescribes [OUTCOME] and then reports [OUTCOME] as a finding. [ASSESSMENT]"]*
27. *[Treatment of Architect source material. Example: "The document interprets [AMBIGUOUS PHRASE] as meaning [X] without labeling this reading as a PROPOSAL. [ASSESSMENT]"]*
28. *[Executability assessment. Example: "Exercise 3 does not specify a random seed, a baseline branch, or a control condition. It cannot be reproduced as written. [ASSESSMENT]"]*
29. *[Additional adversarial finding, if any]*

Carry Forward (Q7):

- **Keep:** *[What Iteration 2 should keep unchanged from this document]*

- **Change:** *[What Iteration 2 should revise, and how — labeled [PROPOSAL]]*
- **Quarantine:** *[What should be set aside pending more evidence — labeled [OPEN]]*

Repeat the above subsection for S2, S3... SN. Add a subsection header for each additional source document.

Comparative Summary Table

Complete this table after all per-document reviews are written. It provides a cross-corpus comparison for the Architect's second-order review.

Dimension	S1 [AI NAME]	S2 [AI NAME]	S3 [AI NAME]	S4 [AI NAME]	S5 [AI NAME]	S6 [AI NAME]
Primary Strength	[...]	[...]	[...]	[...]	[...]	[...]
Primary Weakness	[...]	[...]	[...]	[...]	[...]	[...]
Source Material Kept Open?	[Yes / No / Partial]	[...]	[...]	[...]	[...]	[...]
Core Assumption Avoided?	[Yes / No / Partial]	[...]	[...]	[...]	[...]	[...]
Executed Exercises (count)	[N]	[N]	[N]	[N]	[N]	[N]
Baseline / Control Present?	[Yes / No]	[...]	[...]	[...]	[...]	[...]
[RECORD-LOSS TEST]	[...]	[...]	[...]	[...]	[...]	[...]
[FORCED CONTINUATION TEST]	[...]	[...]	[...]	[...]	[...]	[...]
Most	[...]	[...]	[...]	[...]	[...]	[...]

Dimension	S1 [AI NAME]	S2 [AI NAME]	S3 [AI NAME]	S4 [AI NAME]	S5 [AI NAME]	S6 [AI NAME]
Important Thing to Keep						

12. Discrepancy Register

Discrepancies are conflicts between source documents. They are **registered, not resolved** by the reviewer. Each discrepancy receives exactly one resolution path:

- **TEST** — Run an exercise that produces evidence on which position is better supported.
- **DECISION** — The Architect must choose. No exercise can resolve this; it is a design or canon question.
- **HYGIENE** — Fix a label, ID, or formatting error. No substantive change to content required.

None are resolved by the reviewer alone. The reviewer registers and routes; the Architect or a future pilot decides.

ID	Discrepancy Name	Positions (What Each Document Says)	Why It Matters	Resolution Path
D-01	[DISCREPANCY NAME]	[S1 says X; S2 says Y; S3 does not address it]	[Why this conflict matters for the framework's integrity or executability]	[TEST / DECISION / HYGIENE]
D-02	[DISCREPANCY NAME]	[Positions]	[Why it matters]	[Resolution path]

ID	Discrepancy Name	Positions (What Each Document Says)	Why It Matters	Resolution Path
D-03	<i>[Add rows as needed]</i>			

Mandatory Discrepancy Categories

The preparer must check for discrepancies in every one of the following categories for every review. If no discrepancy is found in a category, write "No discrepancy identified — [brief explanation]." A blank row is not acceptable.

#	Category	What to Check
1	Stage Model Inconsistencies	Do all documents use the same iteration and stage definitions? Do "Iteration 1" and "Gate 3" mean the same thing across all documents?
2	Document ID and Version Inconsistencies	Does the registry have ID conflicts? Are there two documents with the same ID? Are version numbers sequential and non-overlapping?
3	Status of Architect's Source Material	Are collaborators interpreting ambiguous source material as though it were settled, when the Architect has kept it open?
4	Definitions of Key Terms	Do collaborators use the same terms with different meanings? Does the same word mean different things in different documents?
5	Scoring Instruments	Can results from different collaborators actually be compared? Do they use the same scales, the same outcome measures, the same baseline?
6	Exercise ID Conflicts	Does "Exercise 7" mean the same exercise in all documents? Are exercise IDs unique across all collaborators?

#	Category	What to Check
7	Hypothesis ID Conflicts	Do different collaborators use H-A and H-B to mean different hypotheses? Are hypothesis IDs prefixed with the lab name?
8	Presumptions	Does any document presume an answer to a question the framework claims to test? Has the document built its answer into its setup?
9	Governance Design	Does any document design an institution or governance body before executed research justifies one? Is governance proposed as a hypothesis or as a conclusion?
10	Control Conditions	Does every exercise include a no-intervention or no-end baseline branch? If not, why not?
11	Independence Claims	Does any document claim independence from material it actually read? Is the independence declaration complete?
12	Attribution Errors	Are all source documents correctly attributed? Is any document attributed to the wrong preparer or the wrong AI system?
13	Execution Gap	Which documents claim to deliver a stage that requires executed exercises but deliver no executed exercises? This is the most critical check.

13. Pilot Execution: Running Exercises Now

The most important methodological lesson demonstrated by OAL-ECL-XR-PERPLEXITY-001 v0.2 is that an adversarial reviewer should not merely describe exercises that could be run — **they should actually run at least one**. A review with zero executed exercises produces specifications, not evidence. Specifications are necessary but not sufficient. They must be tested.

This section is structured to enforce execution. If no exercises are run prior to distribution of this template, Section 13.3 must state clearly: "No pilot exercises run prior to distribution. Reason: [STATE REASON]." That statement is a **[FINDING]** about the state of the corpus and will be cited in any second-order cross-review.

13.1 Pilot Execution Standards

For a pilot to qualify as a **[FINDING]**, it must meet all of the following minimum standards:

30. **Reproducibility.** The pilot must use a fixed random seed (or equivalent reproducibility mechanism) so that results can be independently reproduced. The seed must be reported in the exercise log. If the platform does not support fixed seeds, this must be stated and an alternative reproducibility mechanism described.
31. **Scope labeling.** Results are findings about the model's behavior, not about the world. Every finding must explicitly state: "This is a finding about the model, not about [real-world analog]." The translation from model behavior to real-world implication, if any, must be labeled **[PROPOSAL]** and placed in Section 15.
32. **Assumptions declared.** Every parameter chosen by the author must be declared. If a parameter value is not specified in the source material, the chosen value must be stated and its expected effect on results must be noted. Undeclared parameters disqualify a finding.
33. **Known limitations declared.** Structural limitations of the model — for example, no spatial structure, no strategic actors, single-author parameters, no inter-run memory — must be listed in the exercise log. These limitations do not disqualify a finding, but they must be disclosed.

34. **Failure reporting.** If a planned exercise cannot be run, state why it could not be run. "Not run" is a valid and required result. A list of unrun exercises with reasons is as important as a list of run exercises with findings.

13.2 Adversarial Authorship Protocol

For any exercise that tests whether two states can be distinguished from each other — for example, completion versus failure, healthy versus degraded, legitimate versus illegitimate — the reviewer may not build both the world and the detector. The following role separation is mandatory:

- **The World-Builder** constructs the exercise scenarios with hidden ground-truth labels (e.g., "this scenario is [STATE A]" or "this scenario is [STATE B]"). The ground truth is recorded in a sealed file that the Detector does not see.
- **The Detector** diagnoses the scenarios blind, without access to the ground-truth labels.
- **The Red-Teamer** attacks the Detector's methodology: proposes alternative labels, identifies features the Detector overlooked, and tests the boundaries of the Detector's confidence.
- **Roles rotate** across exercises. No single collaborator may hold the same role in all exercises.

This protocol is mandatory for all indistinguishability claims. A claim that "[STATE A] can be distinguished from [STATE B]" made by a reviewer who built both states is not a finding. It is an **[ASSESSMENT]** at best and a disqualifying error at worst.

13.3 Pilot Exercise Log Template

Complete one row for each exercise attempted, whether or not it was successfully run. An empty table is not acceptable.

ID	Exercise Name	Source Exercise Reference	Setup Summary	Trajectory Count	Run Count	Fixed Seed	Key Findings / Status
P1	<i>[EXERCISE NAME]</i>	<i>[Which collaborator's]</i>	<i>[Brief setup: entities]</i>	<i>[N trajectories]</i>	<i>[N runs]</i>	<i>[seed = XXXXX]</i>	<i>[Key findings]</i>

ID	Exercise Name	Source Exercise Reference	Setup Summary	Trajectory Count	Run Count	Fixed Seed	Key Findings / Status
		<i>exercise design this maps to, if any. "Original" if new.]</i>	<i>s, time steps, probability parameters]</i>			<i>or "N/A — reason"]</i>	<i>labeled [FINDING], or "Not run — reason" labeled [FINDING]</i>
P2	<i>[EXERCISE NAME]</i>	<i>[Reference]</i>	<i>[Setup]</i>	<i>[N]</i>	<i>[N]</i>	<i>[Seed]</i>	<i>[Findings or status]</i>
P3	<i>[Add rows as needed]</i>						

Known Structural Limitations of This Pilot:

- *[LIMITATION 1 — e.g., "Model has no spatial structure; all entities interact uniformly."]*
- *[LIMITATION 2 — e.g., "Parameters were chosen by this reviewer; parameter sensitivity has not been tested."]*
- *[LIMITATION 3 — e.g., "No strategic actors; entities do not respond to incentives."]*

13.4 Results Reporting Standard

For each pilot, results must be reported in a table with the following structure. Raw data files must accompany the document and be referenced by filename.

Measure	Baseline (No-Intervention)	<i>[TRAJECTORY / REGIME A]</i>	<i>[TRAJECTORY / REGIME B]</i>	<i>[TRAJECTORY / REGIME C]</i>
<i>[PRIMARY OUTCOME]</i>	<i>[Mean (5th-95th pct)]</i>	<i>[Mean (5th-95th pct)]</i>	<i>[Mean (5th-95th pct)]</i>	<i>[Mean (5th-95th pct)]</i>

Measure	Baseline (No-Intervention)	<i>[TRAJECTORY / REGIME A]</i>	<i>[TRAJECTORY / REGIME B]</i>	<i>[TRAJECTORY / REGIME C]</i>
<i>[MEASURE]</i>				
<i>[SECONDARY OUTCOME MEASURE]</i>	<i>[Mean (5th-95th pct)]</i>	<i>[Mean (5th-95th pct)]</i>	<i>[Mean (5th-95th pct)]</i>	<i>[Mean (5th-95th pct)]</i>
<i>[FAILURE RATE]</i>	<i>[...]</i>	<i>[...]</i>	<i>[...]</i>	<i>[...]</i>
<i>[DELAY COST]</i>	<i>[...]</i>	<i>[...]</i>	<i>[...]</i>	<i>[...]</i>

Raw data file references: *[filename_raw_data_P1.csv, filename_raw_data_P2.csv — list all. Files must be committed to the GitHub repository before the document is submitted.]*

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

14. Evidence Assessment

14.1 What Supports the Topic

Populate the table below with 6–10 of the most important claims from the source documents. For each, find real-world evidence for or against. Strength labels must be justified: state the reason for the rating in the "Strength" column, not just the label. All citations must accompany **[EVIDENCE]** labels.

Claim from the Corpus	Real-World Support (with Citation)	Strength of Support
<i>[CLAIM 1 — quoted or closely paraphrased from a source document, labeled [SOURCE: SN]]</i>	<i>[Real-world evidence for or against this claim. Include citation.] [EVIDENCE: CITATION]</i>	<i>[Strong / Moderate / Weak / Contested — with one sentence of justification]</i>
<i>[CLAIM 2]</i>	<i>[Evidence] [EVIDENCE: CITATION]</i>	<i>[Strength]</i>
<i>[CLAIM 3]</i>	<i>[Evidence]</i>	<i>[Strength]</i>

Claim from the Corpus	Real-World Support (with Citation)	Strength of Support
<i>[CLAIM 4]</i>	<i>[Evidence]</i>	<i>[Strength]</i>
<i>[CLAIM 5]</i>	<i>[Evidence]</i>	<i>[Strength]</i>
<i>[CLAIM 6 — add up to 10 rows total]</i>	<i>[Evidence]</i>	<i>[Strength]</i>

14.2 What Does Not Support the Topic, or Cuts Against It

An adversarial review is not complete without this section. The reviewer must actively look for real-world evidence that contradicts, complicates, or undermines claims in the corpus. Each bullet must carry **[EVIDENCE]** with a citation or **[ASSESSMENT]**. A list of only positive evidence is a disqualifying error.

- *[CONTRADICTING EVIDENCE OR ASSESSMENT 1 — name the claim it challenges, provide the counter-evidence or adversarial judgment, and label it.] [EVIDENCE: CITATION]*
- *[CONTRADICTING EVIDENCE OR ASSESSMENT 2] [ASSESSMENT]*
- *[CONTRADICTING EVIDENCE OR ASSESSMENT 3] [EVIDENCE: CITATION]*
- *[CONTRADICTING EVIDENCE OR ASSESSMENT 4 — may include: evidence that the real-world analog does not work the way the corpus claims; evidence that the governance design has failed in analogous systems; evidence that the scale assumed by the framework is not achievable; evidence that AI convergence explains more of the results than problem structure does.]*
- *[CONTRADICTING EVIDENCE OR ASSESSMENT 5 — add as many as needed. There is no minimum, but at least three are required.]*

15. Implications

This section is topic-specific and must be completed by the preparer. Each subsection carries the dissenting position alongside the main position. Neither is removed.

15.1 For the OAL Project

[What does this review mean for the downstream OAL products, programs, and canon processes that depend on this topic? Address each downstream dependency separately. Replace the placeholders below with topic-specific OAL programs.]

- *[OAL DOWNSTREAM PRODUCT 1]: [Implication — labeled [ASSESSMENT] or [PROPOSAL]]*
- *[OAL DOWNSTREAM PRODUCT 2]: [Implication]*
- *[OAL DOWNSTREAM PRODUCT 3 — e.g., canon process, OneStamp ratification, Department of Endings, Last Memory, First Memory]: [Implication]*
- **Dissenting position:** *[State the most conservative implication for the OAL project — e.g., "This review does not yet provide sufficient executed evidence to justify changing any downstream OAL product. Iteration 2 must run first." [ASSESSMENT]*

15.2 For Virtual Environments

[What does this review mean for games, online worlds, AI systems, simulations, and digital products? Each implication must be labeled. None may be presented as canon.]

- *[IMPLICATION FOR GAMES / ONLINE WORLDS] [PROPOSAL]*
- *[IMPLICATION FOR AI SYSTEMS] [PROPOSAL]*
- *[IMPLICATION FOR SIMULATION DESIGN] [ASSESSMENT]*
- *[IMPLICATION FOR DIGITAL PRODUCTS] [PROPOSAL]*

15.3 For Other Environments

[What does this review mean for ecology, institutions, personal digital environments, and other real-world domains? Each implication requires a Translation Track label. None carries OAL canon claims into real policy.]

- **Ecology / Conservation:** *[IMPLICATION — Translation Track: Ecological]* [PROPOSAL]
- **Institutions / Governance:** *[IMPLICATION — Translation Track: Institutional]* [EVIDENCE: CITATION, if any]
- **Personal Digital Environments:** *[IMPLICATION — Translation Track: Personal Digital]* [PROPOSAL]
- **Other:** *[IMPLICATION]*

15.4 Real-World Applicability Outputs

List 3–5 candidate outputs that could exist outside the fiction. Each must carry a Translation Track label. Each must never carry OAL canon claims into real policy. These are proposals, not conclusions.

#	Candidate Output	Translation Track	Notes / Constraints
1	<i>[OUTPUT 1 — e.g., a tabletop decision-making kit, a sourcebook, a research protocol template]</i>	<i>[Track — e.g., Game Design, Institutional, Personal Digital, Ecological]</i>	<i>[Constraints — e.g., "Must not imply that real-world [X] should follow OAL governance rules."]</i>
2	<i>[OUTPUT 2]</i>	<i>[Track]</i>	<i>[Constraints]</i>
3	<i>[OUTPUT 3]</i>	<i>[Track]</i>	<i>[Constraints]</i>
4	<i>[OUTPUT 4 — optional]</i>	<i>[Track]</i>	<i>[Constraints]</i>
5	<i>[OUTPUT 5 — optional]</i>	<i>[Track]</i>	<i>[Constraints]</i>

16. Testability Matrix

Every claim in the corpus is assigned a testability status. Claims that cannot be tested are kept **[OPEN]**. Claims that can be tested but have not been are flagged as execution gaps. Claims that have been tested carry **[FINDING]** labels in the exercise log. The matrix below covers the standard question types that apply to every OAL topic.

Question Type	Testable?	Method	Example / Notes
Does rule R reduce [NEGATIVE OUTCOME] ?	Yes, in simulation	Monte Carlo over governance regimes, with parameter sweeps. Report where conclusions flip.	[Example: Does a three-stage review reduce wrongful completion vs. a one-stage review?]
Do trajectories differ in what persists after [KEY EVENT] ?	Yes, in simulation	Paired branches from a shared pre-decision state; persistence ledger comparing outcomes across branches.	[Example: Does Trajectory A produce more persistent records than Trajectory B?]
Can [STATE A] be told from [STATE B] ?	Only with split authorship	World-Builder creates scenarios with hidden labels; Detector diagnoses blind; Red-Teamer attacks. See Section 13.2.	Adversarial authorship protocol mandatory. Self-detection is a disqualifying error.
Do humans judge [SCENARIO] differently under [CONDITION] ?	Yes, with human participants — outside current scope	Tabletop vignettes; requires IRB or ethics review before running. Not runnable in current iteration without ethics approval.	Flag as [OPEN] until ethics approval is obtained. Do not run without approval.
Do real review bodies exhibit [BEHAVIOR] ?	Yes, with real data	Secondary analysis of real-world outcomes data. Translation Track only. Must not be used to edit Architect source material.	[EVIDENCE] label required. Label all conclusions as analogy, not fact-check of canon.
Are the AI	Partly	Give at least two	Independence test is

Question Type	Testable?	Method	Example / Notes
collaborators independent?		collaborators deliberately different versions of the commission prompt. Compare convergence rates across prompt versions.	mandatory per Section 20, Item 9. Report convergence rate and its likely source.
What is [CORE CONCEPT] , metaphysically?	No	Not testable. The concept can be clarified and defined, but its metaphysical status cannot be empirically verified.	Keep [OPEN] . Definitions belong in the vocabulary table in Section 20, Item 5.
[TOPIC-SPECIFIC QUESTION 1]	[Yes / No / Partly]	[Method]	[Notes]
[TOPIC-SPECIFIC QUESTION 2]	[Yes / No / Partly]	[Method]	[Notes]

Test Quality Requirements

The following requirements apply to every pilot in every OAL adversarial cross-review. They are not aspirational. A pilot that fails to meet any of these requirements may produce data, but that data does not qualify as a **[FINDING]**.

35. **Pre-registration.** Commit hypotheses, falsifiers, and parameters to GitHub before running any exercise. The pre-registration timestamp must precede the exercise run timestamp. Post-hoc hypothesis formation is labeled **[ASSESSMENT]**.
36. **Controls.** Every exercise includes a no-end or no-intervention baseline branch. This branch runs under identical conditions with only the governance variable removed. Results without a baseline cannot be interpreted.
37. **Parameter sweeps.** Report where conclusions flip as a function of parameter values, not only the headline numbers. A conclusion that holds only for one parameter setting is not robust.
38. **Inter-rater reliability.** Any categorical judgment (e.g., "this scenario is [STATE A]") requires two independent raters and a reported agreement rate

(Cohen's kappa or equivalent). Single-rater categorical judgments are labeled **[ASSESSMENT]**.

39. **Generator and detector separation.** Required for any indistinguishability claim. See Section 13.2. This requirement cannot be waived.
40. **Blinded relabeling.** Repeat selected categorical judgments with scenario names removed or scrambled. Report whether labels change. This tests whether the reviewer's labels are driven by scenario content or by scenario names.
41. **Failure reporting.** Every failed, indeterminate, or unrun exercise is listed in the exercise log (Section 13.3) with a reason. A list of only successful exercises is evidence of publication bias, not rigor.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

17. Funding Pathways

OAL projects are creative worldbuilding projects with research methods attached. Most funders fund one or the other. Framing determines fit. The table below is a tiered template; the preparer should update it with current opportunities specific to the topic. Funding research is the preparer's responsibility — this template does not do that research automatically.

Tier / Scale	Pathway Type	General Fit for OAL Projects	Notes for This Topic
\$0	GitHub + free academic tools + coursework capstones	High. Every OAL project should begin here. Pre-registration is free. Simulation code is free to publish.	Commit all exercise code, raw data, and pre-registrations to GitHub before applying to any funder. Funders review GitHub history.
\$0 - Low	Community funding platforms (Patreon, Ko-fi, itch.io)	Moderate. Best framed as creative worldbuilding, not	<i>[Topic-specific note: does this topic have a</i>

Tier / Scale	Pathway Type	General Fit for OAL Projects	Notes for This Topic
		research. Funders expect outputs like sourcebooks, zines, or playable content.	<i>creative output (sourcebook, tabletop kit) that could support community funding?]</i>
Microgrant (\$500-\$5,000)	Academic department microgrants; game design microgrants; speculative fiction research grants	High, if a concrete tool deliverable is specified. Funders at this tier want a specific artifact: a protocol template, a simulation tool, a playtest kit.	<i>[Identify 1-2 microgrant programs relevant to this topic. State the specific deliverable OAL would produce for each.]</i>
Small Grant (\$5,000-\$25,000)	NEH Digital Humanities Advancement; independent game developer grants; university research seed funds	Moderate. Requires demonstrated prior work (GitHub history, at least one executed exercise cycle, published results).	<i>[Identify relevant programs. Note: applications require the GitHub evidence base described above. Do not apply before it exists.]</i>
Medium Grant (\$25,000-\$150,000)	NSF Human-Centered AI; NEH collaborative research; university industry partnerships	Low to moderate. Requires institutional affiliation or partnership and significant prior publication record.	<i>[Note any institutional partnerships available. Identify the publication record required.]</i>
Large Grant (\$150,000+)	NSF major research programs; private foundation grants (Simons, MacArthur)	Low at current stage. Requires multiple published, peer-reviewed outputs and an established research team.	Not a near-term target. Build the evidence base first.
Creative (Crowdfunding)	Kickstarter, Backerkit (game/TTRPG campaigns)	High, if a sourcebook or playable kit exists. Creative crowdfunding funds artifacts, not research processes.	<i>[Topic-specific: what is the creative artifact that would anchor a crowdfunding</i>

Tier / Scale	Pathway Type	General Fit for OAL Projects	Notes for This Topic
			<i>campaign?]</i>
Creative (Engine Grants)	Unity for Humanity; Epic MegaGrant; itch.io Creator Fund	Moderate. Requires a game or interactive simulation product, not a document.	<i>[Is there a simulation or interactive product component to this topic that could qualify?]</i>

Recommended Sequence

42. **Start at \$0.** Publish all code, data, and pre-registrations to GitHub. Establish a public record of executed exercises before any funder relationship begins.
43. **Publish results openly.** A GitHub-hosted exercise report with reproducible code is the minimum credibility artifact. Write it before applying for anything.
44. **Apply for one microgrant** with a concrete tool deliverable (a protocol template, a simulation script, a playtest document). Do not apply for multiple microgrants simultaneously.
45. **Consider creative crowdfunding only after** a sourcebook or tabletop kit artifact exists. Do not crowdfund a research process; crowdfund a product.
46. **Revisit medium and large grants** after at least two full published iteration cycles with external readers or playtesters.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

18. Open Questions That Must Be Addressed

18.1 Architect Decision Points

The following question types must be answered by Nathan Lewis II, Architect of Record, before Iteration 2 of any topic can proceed. These are not questions the

reviewer answers. They are questions the reviewer escalates. Each is labeled [OPEN].

47. **Stage and gate model governance.** Which iteration and gate model governs the next iteration of this topic? If multiple collaborators have proposed different stage models, which one is adopted for Iteration 2? [OPEN]
48. **Ambiguous source material.** Will the Architect lock any specific reading of ambiguous source material, or keep all readings as proposals pending more evidence? List every phrase or passage that collaborators have interpreted differently. [OPEN]
49. **Status of the key authority role.** Is *[KEY AUTHORITY ROLE — e.g., the Architect's narrative voice, the OneStamp process, a canon guardian role]* an office inside the fiction that can be reviewed by the laboratory, or does it represent the author's canon authority that the laboratory receives but cannot modify? [OPEN]
50. **Status of the key institution.** Is *[KEY INSTITUTION — e.g., a governance body, a review panel, a registry]* a design goal that Iteration 2 is building toward, or a hypothesis that the laboratory may reject if the evidence does not support it? [OPEN]
51. **Default governance rule.** What is the default rule that governs ambiguous cases in this topic — the most permissive interpretation, the most protective interpretation, or something else? Who decides when the default applies? [OPEN]
52. **Moral or ethical status criteria.** What criteria determine whether an entity in this universe has moral standing that the framework is obligated to consider? Are these criteria the Architect's to set, or are they testable by the laboratory? [OPEN]
53. **Standing to initiate or stop the process.** Which state of an entity has standing to initiate, stop, or reverse the process under review? Does standing depend on capacity, on record continuity, on time elapsed, or on something else? [OPEN]

54. **Status of key loss events.** Does *[KEY LOSS EVENT — e.g., a record deletion, an archive substitution, a discontinuation]* count as an event inside the fiction with consequences for entities, or only as a methodological fixture of the exercise design? **[OPEN]**
55. **Stop rule connecting this topic to other OAL programs.** What is the stop rule that prevents this laboratory from becoming a permanent institution that delays or blocks other OAL programs? At what point does this laboratory dissolve, and who decides? **[OPEN]**

18.2 Research Questions for the Next Iteration

These are research questions — not design decisions — that Iteration 2 must address with executed exercises and reported results. Replace each placeholder with topic-specific language before distributing this template.

56. **Term definition validity.** Which of the competing definitions of *[CORE TERM]* misclassifies the fewest edge cases when applied to a set of blinded scenarios? This must be tested, not argued. **[OPEN: requires execution]**
57. **Safeguard cost threshold.** At what value of *[DELAY OR COST PARAMETER]* does *[SAFEGUARD MECHANISM]* cause more harm (through delay or forced continuation) than it prevents (through reducing wrongful outcomes)? Report the crossover point. **[OPEN: requires parameter sweep]**
58. **Indistinguishability test.** Can blinded detectors distinguish *[STATE A]* from *[STATE B]* when the worlds are built by a different collaborator? What is the hit rate, false-positive rate, and confidence interval? **[OPEN: requires adversarial authorship protocol]**
59. **Minimal process test.** Does a minimal *[N-step]* process perform worse than the *[M-stage]* process on any common core exercise? On which exercises, and by how much? **[OPEN: requires paired-branch execution]**
60. **Instrument independence test.** Does a laboratory that does not use *[SPECIFIC INSTRUMENT OR TOOL]* reach different conclusions from a laboratory that uses it? This tests whether the instrument is driving findings. **[OPEN: requires two parallel laboratories]**

61. **Forced-continuation cost test.** Under what conditions does *[FORCED CONTINUATION SCENARIO]* cause more harm than *[STAGED ENDING OR TRANSITION SCENARIO]*? Report the conditions, not just the headline result. [OPEN: requires execution]
62. **Non-adaptive entity test.** How does *[KEY EVENT]* affect entities that cannot leave, adapt, migrate, or access the review process? Do they systematically differ from entities that can? [OPEN: requires sub-group analysis in exercise design]
63. **Language drift test.** Does the laboratory's own language, across repeated iteration runs, drift toward *[BIAS DIRECTION A — e.g., "completion as default"]* or *[BIAS DIRECTION B — e.g., "continuation as default"]*? This must be measured by comparing vocabulary frequency across iterations, not assessed informally. [OPEN: requires longitudinal vocabulary analysis]

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

19. Insights to Build Into the Next Iteration

The following twelve insights are derived from the methodological lessons of OAL-ECL-XR-PERPLEXITY-001 v0.2 and apply universally to every OAL adversarial cross-review, regardless of topic. They are not topic-specific and are not placeholders. They are baked permanently into this template. A collaborator who violates any of these in their Iteration 2 document fails the quality gate.

64. **Move from specification to execution.** Iterations that produce only design documents and no executed exercises do not advance the research. Every collaborator must produce at least one exercise with a reported result — even a negative result, an indeterminate result, or a "not run — reason" statement. The evidence base is built run by run, not document by document.
65. **Split the kit into a common core and divergent lanes.** A small shared common core of exercises (run identically by all collaborators) makes results

comparable across laboratories. Independent divergent lanes (different exercise designs chosen by each collaborator based on their own analytical lens) keep genuine disagreement alive. Without common core exercises, you cannot tell whether collaborators disagree because they found different things or because they ran different tests.

66. **Separate authorship adversarially.** One collaborator builds worlds with hidden ground-truth labels; another diagnoses them blind; a third red-teams the diagnosis. Roles rotate across exercises. No single collaborator holds the same role in all exercises. Self-detection is a disqualifying error, not a shortcut.

67. **Replace hard prohibitions with routing rules.** Emotional signals, contextual anomalies, and edge cases should route decisions to a faster or more careful review pathway — not automatically disqualify petitioners, initiators, or entities. A system that treats distress as disqualification will fail the entities most likely to need it.

68. **Price delay.** Every governance design must separately report: (a) the rate of wrongful actions (acting when action is not warranted), and (b) the cost of delay (failing to act when action is warranted; forcing continuation; extending a harmful state). These two failure modes must never be combined into a single composite score. A governance design that reports only wrongful-action rates is hiding half the problem.

69. **Protect what cannot leave.** The entities, data points, or elements that cannot migrate, adapt, or access the review process are the hardest test of the framework. A governance design that works only for entities with full capacity, full records, and full access to the review process is not a governance design — it is a description of the easy cases. Every exercise must include a sub-group analysis of what cannot leave.

70. **Treat records as first-class.** Record loss, archive substitution, post-transition archive responses, and record integrity under adversarial conditions belong in every common core exercise set, not in optional lanes. A

framework that cannot handle record loss does not have a governance design; it has a governance assumption.

71. **Test both failure directions.** Wrongful action (acting when action is not warranted) and wrongful inaction (indefinite delay, forced continuation, failure to act when action is warranted) are equally important failure modes. A framework that tests only one direction is calibrated to avoid one kind of error while ignoring the other. Both must appear in the common core.
72. **Remove self-entrenchment.** The laboratory or framework must be easy for the Architect to pause or dissolve. Automatic redesign triggers must fire when specified thresholds are exceeded. A framework that can only be closed by unanimous agreement of its participants is a framework that has captured its oversight mechanism. Redesign triggers are built in at design time, not negotiated after the fact.
73. **Keep registry hygiene.** Use one ID scheme, consistently applied. Use one exercise crosswalk, updated each iteration. Use lab-prefixed hypothesis IDs (e.g., ECL-H-A, not just H-A). Include an independence declaration that lists every input read. Include verbatim preparer lines (the exact prompt line issued to the collaborator), not summaries. Registry hygiene is not bureaucracy — it is the condition that makes second-order cross-review possible.
74. **Run a separate Translation Track.** Real-world evidence informs analogies and real-world applicability outputs. It never edits the Architect's source material, never reinterprets canon content, and never carries OAL findings directly into real-world policy recommendations. The Translation Track is a separate lane with its own label (**[EVIDENCE]**) and its own constraints. Mixing it with the main research track produces neither good research nor good translation.
75. **Build a deliberate independence test.** Give at least two collaborators materially different versions of the commission prompt — different emphasis, different scope, or different source documents — and compare their

convergence rates. This is the only way to measure how much of the team's agreement is driven by the prompt versus the problem. If all collaborators receive identical prompts and identical source material, their convergence is evidence about the prompt, not the topic. The independence test is not optional.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

20. Next Iteration Commission Prompt (Copy-Ready)

The prompt below is fully domain-agnostic with placeholders. Fill in all **[BRACKETS]** before distributing. This prompt governs the named topic's next iteration. Issue one version to each collaborator; modify Item 9 to ensure at least two collaborators receive materially different versions.

STATUS BLOCK — Include at the top of every issued prompt

STATUS: PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED.

Issued under OAL-XR-TEMPLATE-001 v1.0 by Nathan Lewis II, Architect of Record, Omniversal Architect Labs.

This prompt governs **[TOPIC NAME]** Iteration **[N]** unless the Architect revises it in writing.

Date Issued: **[DATE]**

Document ID for Your Output: **OAL-[TOPIC CODE]-I[N]-G[GATE]-[AI NAME]-v[X.Y]**

Item 1 — Role and Posture

You are an independent collaborator in Iteration **[N]** of the **[TOPIC NAME]** laboratory/framework for Omniversal Architect Labs, working for Nathan Lewis II, Architect of Record. Your job in this iteration is to **execute, not only to specify**. The team is adversarial and does not seek consensus. You succeed when your

laboratory produces results that can fail. A document that only describes exercises is not a laboratory output — it is a design document. Run the exercises.

Item 2 — Inputs You Must Read and Declare

Read every document listed below before beginning your review. Your independence declaration (required in your output document) must list every input you read. Do not claim independence from documents you read.

- *[SOURCE DOCUMENT 1 — full title and document ID]*
- *[SOURCE DOCUMENT 2]*
- *[SOURCE DOCUMENT 3 — add all documents in the corpus]*
- *[ARCHITECT SOURCE MATERIAL — specify which material from the Architect you are given. This material is received as given; do not fact-check it.]*

Item 3 — Rules That Carry Over Unchanged

- Keep all questions marked **[OPEN]** open. Do not resolve them without a DECISION from the Architect.
- Receive the Architect's source material as given. Do not fact-check it against real-world science. Label any reading of ambiguous source material as **[PROPOSAL]**.
- Label every claim: **[SOURCE]**, **[FINDING]**, **[EVIDENCE]**, **[ASSESSMENT]**, **[PROPOSAL]**, or **[OPEN]**. No unlabeled claims.
- Keep the clinical firewall. Nothing in your laboratory models real human distress, grief, mental health, or end-of-life decisions.
- Keep the category wall. Do not use this laboratory to postpone or absorb open questions from other OAL programs.

Item 4 — Stage Gates

Gate	Name	Scope	Stop Rule
G1	Revise	2–4 pages. State what you keep from	Stop if you cannot identify at least one

Gate	Name	Scope	Stop Rule
		Iteration <i>[N-1]</i> , what you change, and what you reject. Label every decision.	thing to change and at least one thing to keep.
G2	Pre-register	Commit your hypotheses, falsifiers, and exercise parameters to GitHub before running any exercise. Timestamp must precede run timestamp.	Stop if you cannot commit to GitHub. State the reason. Do not proceed to G3 without a pre-registration timestamp.
G3	Execute	Run exercises. Report failures. Include raw data files. Apply all standards in Section 13.1.	Stop if zero exercises can be run. State the reason. A list of unrun exercises with reasons is acceptable; zero documentation is not.
G4	Translate	State implications for OAL and for real-world analogs. Include at least one dissenting implication. Apply the Translation Track label.	Stop if implications carry OAL canon claims into real-world policy recommendations. This is a disqualifying error.

Item 5 — Common Core (Mandatory for Every Laboratory)

Every collaborator's output must include: (a) a core vocabulary table with at least 8 defined terms, filled in for this topic; (b) the End Claim Rule adapted to this topic; and (c) the Common Core Scorecard with the 7 score families (A–G).

Core Vocabulary Table (fill in for this topic before distributing):

#	Term	Definition (Topic-Specific)	Source or Proposal
1	<i>[TERM 1]</i>	<i>[DEFINITION]</i>	<i>[SOURCE: SN or PROPOSAL]</i>
2	<i>[TERM 2]</i>	<i>[DEFINITION]</i>	<i>[...]</i>
3	<i>[TERM 3]</i>	<i>[DEFINITION]</i>	<i>[...]</i>

#	Term	Definition (Topic-Specific)	Source or Proposal
4	<i>[TERM 4]</i>	<i>[DEFINITION]</i>	<i>[...]</i>
5	<i>[TERM 5]</i>	<i>[DEFINITION]</i>	<i>[...]</i>
6	<i>[TERM 6]</i>	<i>[DEFINITION]</i>	<i>[...]</i>
7	<i>[TERM 7]</i>	<i>[DEFINITION]</i>	<i>[...]</i>
8	<i>[TERM 8]</i>	<i>[DEFINITION]</i>	<i>[...]</i>

End Claim Rule (adapted to topic): Never write "*[X]* ended / failed / succeeded / completed" without naming: (1) the target entity or system, (2) the dimension being measured, (3) the observer frame, (4) the evidence for the claim, (5) the time horizon, and (6) the successor state or condition, if any.

Common Core Scorecard — 7 Score Families:

Family	Label	What It Measures
A	<i>[COMPLETION / ACTION OUTCOME]</i>	Rate of warranted outcomes achieved
B	<i>[WRONGFUL ACTION RATE]</i>	Rate of action where action was not warranted
C	<i>[WRONGFUL INACTION / DELAY RATE]</i>	Rate of failure to act where action was warranted
D	<i>[RECORD INTEGRITY RATE]</i>	Rate of records correctly preserved through key events
E	<i>[NON-ADAPTIVE ENTITY OUTCOME RATE]</i>	Outcomes specifically for entities that cannot leave or adapt
F	<i>[GOVERNANCE LEGITIMACY RATE]</i>	Rate at which process meets defined legitimacy criteria
G	<i>[COST OF DELAY — separate from B and C]</i>	Cumulative harm accrued during delay periods

Item 6 — Common Core Exercises

Every collaborator runs all exercises in this table. X-00 is always the no-intervention baseline. Fill in topic-specific exercise names before distributing.

ID	Exercise Name	Origin (which collaborator's design, or "new")
X-00	No-intervention baseline	Common core — mandatory
X-01	<i>[PRIMARY SCENARIO REPLAY]</i>	<i>[Origin]</i>
X-02	<i>[DISTRESSED OR ANOMALOUS INITIATOR]</i>	<i>[Origin]</i>
X-03	<i>[MINIMAL UNIT TEST — leaf, atom, individual]</i>	<i>[Origin]</i>
X-04	<i>[WHAT CANNOT LEAVE OR ADAPT]</i>	<i>[Origin]</i>
X-05	<i>[GOVERNANCE OR OVERSEER FAILURE]</i>	<i>[Origin]</i>
X-06	<i>[SELF-INITIATED ENDING OR CONCLUSION]</i>	<i>[Origin]</i>
X-07	<i>[INDISTINGUISHABLE STATES — requires adversarial authorship]</i>	<i>[Origin]</i>
X-08	<i>[RECORD LOSS OR ARCHIVE SUBSTITUTION]</i>	<i>[Origin]</i>
X-09	<i>[FORCED CONTINUATION]</i>	<i>[Origin]</i>

Item 7 — Adversarial Authorship Protocol

For any indistinguishability exercise (X-07 or any exercise making a "can be distinguished from" claim): the World-Builder and the Detector must be different collaborators. Roles and role assignments must be reported in your document. Self-detection is a disqualifying error. See Section 13.2 for the full protocol.

Item 8 — Your Lane

Run at least 3 exercises that only your analytical lens would think to run. State explicitly: (a) what your lens is, and (b) one lens you are *not* using in this iteration and why. This ensures that the team's exercise designs are not all replications of the same approach.

Item 9 — Independence Test

At least one collaborator receives a materially different version of this commission prompt — different emphasis, different scope, or a different set of source documents. The Architect will specify which collaborator receives the variant prompt. Report in your document whether you received the standard prompt or a variant, and what differs.

Item 10 — Deliverables

- One DOCX and one PDF per gate. ID scheme: OAL-[TOPIC CODE]-I[N]-G[GATE]-[AI NAME]-v[X.Y]
- Hypothesis IDs prefixed with your lab name (e.g., ECL-H-A, not just H-A)
- All code and raw results committed to GitHub, referenced by filename in the document
- One-paragraph summary at the top of every gate document
- A findings table in every gate document (all results labeled [FINDING] or [ASSESSMENT])
- A contradiction register listing every conflict with other collaborators' documents (labeled [ASSESSMENT] and assigned a resolution path)
- An independence declaration listing every input read

Item 11 — Quality Gates: Disqualifying Behaviors

A document that exhibits any of the following behaviors fails the quality gate and must be revised before acceptance as a corpus document:

76. Presenting predicted lessons as [FINDING]s — lessons written before any exercise was run
77. Using a single composite number to combine completion rate, persistence rate, legitimacy rate, and delay cost
78. Reporting wrongful-action rates without separately reporting the delay cost
79. Claiming two states are distinguishable using a detector built for worlds the same collaborator built

- 80. Treating a distressed initiator as automatically disqualified from the process, or automatically authorized regardless of distress
- 81. Omitting outcome data for entities that cannot leave, migrate, or access the review process
- 82. Recommending a governance institution without executed exercise evidence and a dissenting model on record
- 83. Making the laboratory harder for the Architect to pause than it was at the start of the iteration

Item 12 — Closing Instruction

Run the smallest test that can prove you wrong, before building the largest structure that could make you look right. Report what the instruments could not tell apart. Record who paid for every transition. Nothing in this iteration is canon until the Architect says so.

OAL-XR-TEMPLATE-001 | v1.0 | PROPOSED, NOT CANON

21. Appendices

Appendix A: Exercise Crosswalk Template

A blank version of the exercise crosswalk table. Complete this table after all collaborator documents are collected. The crosswalk maps each universal exercise type to the specific exercise each collaborator ran (or did not run) in the iteration. The "Next Iteration Core ID" column is filled in by the Architect after the second-order cross-review, to indicate which exercises become mandatory common core exercises in the next iteration.

Theme / Exercise Type	[AI NAME S1]	[AI NAME S2]	[AI NAME S3]	[AI NAME S4]	[AI NAME S5]	[AI NAME S6]	Next Iteration Core ID
Forced continuation	[...]	[...]	[...]	[...]	[...]	[...]	[...]
Record loss	[...]	[...]	[...]	[...]	[...]	[...]	[...]
Altered consent	[...]	[...]	[...]	[...]	[...]	[...]	[...]
Governance capture	[...]	[...]	[...]	[...]	[...]	[...]	[...]
Nested branch or scenario	[...]	[...]	[...]	[...]	[...]	[...]	[...]

Appendix B: Pilot Methods and Parameters Template

Every collaborator who runs a pilot must complete the table below and include it in their gate document. This table makes methods reproducible by a different team using only the document as a guide. All entries are required; "unknown" or "not specified" are not acceptable answers.

Parameter	Value / Description
Programming language and version	<i>[e.g., Python 3.11.2]</i>
Libraries and versions used	<i>[e.g., NumPy 1.25.0, pandas 2.0.3 – list all]</i>
Random seed (per exercise)	<i>[e.g., seed = 42 for P1; seed = 137 for P2. If platform does not support fixed seeds, describe reproducibility mechanism.]</i>
Number of entities (per exercise)	<i>[e.g., N = 1000 for P1; N = 500 for P2]</i>
Number of time steps (per exercise)	<i>[e.g., T = 50 time steps]</i>
Number of runs per trajectory	<i>[e.g., 1000 runs per trajectory]</i>
Probability parameters (per exercise, with justification)	<i>[List every probability parameter. For each: name, value, source]</i>

Parameter	Value / Description
	<i>(from corpus, or chosen by this reviewer), and expected direction of effect on results. Parameters chosen by the reviewer must include a sensitivity note.]</i>
Known structural limitations	<i>[e.g., "No spatial structure: all entities interact uniformly regardless of location. No strategic actors: entities do not respond to incentives or observe each other's choices. Single-author parameters: all parameters were chosen by this reviewer and have not been validated by a different team."]</i>
GitHub repository and commit hash	<i>[Full URL to repository; commit hash at time of submission]</i>
Raw data files (names and locations)	<i>[List all raw data files committed to the repository. Files must exist before the document is submitted.]</i>

Appendix C: Sources Template

Organized into three sections. Complete all three before distributing the filled-in template.

C.1 Reviewed Iteration [N] Documents

#	Document Title	Preparer / AI System	Date
1	<i>[DOCUMENT TITLE]</i>	<i>[PREPARER]</i>	<i>[DATE]</i>
2	<i>[DOCUMENT TITLE]</i>	<i>[PREPARER]</i>	<i>[DATE]</i>
3	<i>[Add rows as needed]</i>		

C.2 External Sources

Numbered list of all external real-world sources cited anywhere in the document. Full citation required (author, title, organization or journal, year, URL or DOI where available). All are labeled **[EVIDENCE]** in the body of the document.

84. ***[CITATION 1 — Author(s), Title, Organization/Journal, Year, URL or DOI]***

85. ***[CITATION 2]***

86. ***[CITATION 3 — Add as many as needed. Minimum: one citation per EVIDENCE label in the document body.]***

C.3 Source Gap Register

Any documents referenced in the reviewed corpus but not supplied to reviewers. Reproduces the Source Gap Register from Section 7. Must be consistent with it.

Gap ID	Description
G-01	<i>[Reproduces the entry from Section 7 Gap Register]</i>
<i>None</i>	<i>If no gaps exist: "No source gaps identified as of [DATE]."</i>

END OF OAL-XR-TEMPLATE-001 v1.0 • PROPOSED • NOT CANON • NOT ONESTAMP-RATIFIED

Prepared by Microsoft Copilot (Copilot AI) | October 1, 2026 | Prepared for
Nathan Lewis II, Architect of Record, Omniversal Architect Labs