

Adversarial Cross-Review and Iterative Improvement Protocol

Document ID: OAL-ACR-PERPLEXITY-001

Version: v0.2

Date: October 1, 2026

Prepared by: Perplexity AI

Generation platform: Perplexity AI platform / Perplexity Computer session

Underlying model: Not specified in the available session record; do not infer

Prepared for: Nathan Lewis II, Architect of Record

Organization: Omniversal Architect Labs (OAL)

Status: PROPOSED — NOT CANON — NOT ONESTAMP-RATIFIED

Publication intent: GitHub publication and multi-AI comparative review

Provenance class: Perplexity-platform contribution; distinct from the separately commissioned ChatGPT web-application contribution

Executive Summary

OAL has generated a substantial body of conceptual, technical, narrative, governance, and collaborative AI documentation. The central risk is no longer merely a shortage of ideas. It is the possibility that documentation volume, agreement among collaborators, or elaborate framework design may be mistaken for evidence that a system has been exercised, evaluated, or shown to work.

This protocol converts adversarial review into a repeatable, version-controlled practice. It requires reviewers to declare the corpus they actually read; distinguish source claims, executed findings, external evidence, assessments, proposals, open questions, limitations, decisions, hygiene issues, and work not run; preserve disagreements; identify assumptions and missing tasks; propose or execute the smallest test appropriate to the claim; and convert results into explicit revision actions.

The protocol does not seek consensus. Independent convergence may be informative, but agreement is not proof. Disagreement is treated as structured work: some conflicts should be tested, some require an Architect decision, some arise from version or naming problems, and some should remain open, deferred, quarantined, or retired.

A tiered review model prevents the process from becoming another source of documentation bloat. Lightweight artifacts receive lightweight reviews; mature or high-impact systems require preregistration, controls, reproducible results, dissent preservation, and a clear post-review change record. Not every topic requires code or simulation, but every major claim should be matched to the smallest legitimate validation method available.

This v0.2 release also corrects a material provenance error in v0.1. The original files were generated in the current Perplexity AI session but were incorrectly labeled as prepared by ChatGPT (OpenAI). The user clarified that a separate ChatGPT web-application version was still being generated. This release therefore identifies Perplexity AI as preparer, records the generation platform separately, leaves the underlying model unspecified, and preserves the correction conversation for repository versioning.

Table of Contents

1. Document Control
2. Provenance Correction and Conversation Record
3. Mission and Directive
4. Scope, Limits, and Non-Goals
5. Protocol Overview
6. Review Tier Model
7. Corpus and Provenance Control
8. Epistemic Labeling Standard
9. Independent Review Method
10. Assumption and Dependency Audit
11. Discrepancy Management
12. Testability and Validation
13. Preregistration and Exercise Design
14. Results and Interpretation
15. Iterative Improvement and Decisions
16. Multi-AI Collaboration Workflow
17. Stage Gates and Quality Gates
18. Identifiers and Repository Practice
19. Feedback and Cautions
20. Reusable Document Template
21. Copy-Ready Commission Prompt
22. Appendices

Document Control

Field	Value
Document owner	Nathan Lewis II, Architect of Record
Preparer	Perplexity AI
Generation platform	Perplexity AI platform / Perplexity Computer session
Underlying model	Not specified in the available session record
Initial release	October 1, 2026 — v0.1, withdrawn and superseded
Corrected release	October 1, 2026 — v0.2
Review status	Proposed working protocol; open to criticism and comparative revision
Canon status	Not canon; not ONESTAMP-ratified
Primary use	Repeatable adversarial review of OAL frameworks, documents, prompts, systems, repositories, simulations, policies, and collaborative AI outputs
Source basis	ECL cross-review materials, pilot artifacts, the Architect's stated process goals, and the October 1, 2026 attribution-correction conversation
Intended outputs	Versioned DOCX/PDF reviews, machine-readable registers where useful, reproducibility artifacts, and GitHub-ready iteration records

Version History

Version	Date	Platform / preparer	Change summary	Status
v0.1	October 1, 2026	Generated on Perplexity AI platform; incorrectly labeled ChatGPT (OpenAI)	Initial protocol. Attribution and document ID were incorrect because this platform session was confused with a separate ChatGPT web-application commission.	WITHDRAWN / SUPERSEDED
v0.2	October 1, 2026	Perplexity AI; Perplexity AI platform / Perplexity Computer session	Corrected preparer, platform, document ID, filename, metadata, and provenance language; added conversation summary and attribution standard. Core protocol retained.	PROPOSED

Status Notice

This document is a proposal for review. It does not establish OAL canon, policy, authority, or ratification. Adoption, modification, or rejection remains subject to Architect review and any applicable ONESTAMP process.

Provenance Correction and Conversation Record

Purpose

This section documents the attribution error discovered after the v0.1 release, the user's correction, and the resulting version-control action. It enables repository readers to distinguish work generated on the Perplexity AI platform from work generated separately in the OpenAI ChatGPT web application.

Corrected Attribution

Metadata field	Correct v0.2 entry	Version-control meaning
Prepared by	Perplexity AI	Identifies this artifact as the Perplexity-platform contribution.
Generation platform	Perplexity AI platform / Perplexity Computer session	Records where the conversation, drafting, file generation, and correction occurred.
Underlying model	Not specified in the available session record	No provider or model family should be inferred without an explicit platform record.
Document ID	OAL-ACR-PERPLEXITY-001	Separates this artifact from the independent ChatGPT web-application version.
Version	v0.2	Records a material metadata and provenance correction to v0.1.
Prior v0.1 status	WITHDRAWN / SUPERSEDED	v0.1 remains historical but is not the official ChatGPT contribution.
Substantive protocol changes	Core protocol retained; provenance controls added	The correction changes attribution, traceability, and future safeguards.

Conversation Summary

Time (PDT)	Conversation event	Versioning consequence
5:23 a.m., October 1, 2026	The user requested PDF and DOCX versions of the reusable adversarial cross-review protocol.	Formal artifact generation began in the current Perplexity AI session.
Approximately 5:23–5:29 a.m.	Files were delivered as OAL-ACR-CHATGPT-001 v0.1 with ChatGPT (OpenAI) listed as preparer.	The release contained an attribution and identifier error.
5:29 a.m.	The user asked why ChatGPT was listed instead of Perplexity.	The discrepancy was explicitly raised.
5:30 a.m.	The user asked whether a ChatGPT model on Perplexity AI was the preparer.	The conversation examined preparer, platform, application, and underlying model as separate provenance fields.
5:31 a.m.	The user supplied a screenshot and clarified that this work was on Perplexity AI while a separate ChatGPT web version was still being generated.	The controlling provenance fact was established: this is the Perplexity contribution, separate from the ChatGPT web artifact.
5:33 a.m.	The user directed correction in PDF and DOCX with a detailed conversation summary.	v0.2 was commissioned to correct attribution and preserve the history.

Attribution Finding

The v0.1 preparer line conflated this Perplexity-platform session with a separate ChatGPT web-application commission. That would have created incorrect repository provenance and undermined the planned comparison among independently generated AI versions. The corrected record attributes this document to Perplexity AI, leaves the underlying model unspecified, and reserves ChatGPT/OpenAI attribution for the separately generated ChatGPT web-application artifact.

Supersession Notice

OAL-ACR-CHATGPT-001 v0.1 is withdrawn as an official comparative contribution and superseded by OAL-ACR-PERPLEXITY-001 v0.2. The v0.1 files may be retained in an archive solely to document the correction; they should not be indexed as the independent ChatGPT version. Any separate ChatGPT web-application output must retain its own content, creation record, identifier, preparer line, and version history.

Correction Sources

Artifact	Role in correction
Current Perplexity AI conversation, October 1, 2026	Primary record of request, generation, attribution questions, correction, and v0.2 commission.

Artifact	Role in correction
Screenshot-2026-10-01-053129.jpg, thread artifact file:11	User-supplied visual session evidence showing work in the Perplexity AI interface.
OAL-ACR-CHATGPT-001 v0.1 DOCX and PDF	Superseded artifacts containing the incorrect preparer and identifier.
Separately generated ChatGPT web-application document	Independent artifact outside this provenance; it must not be merged into this Perplexity version.

Provenance Rule Adopted

Future OAL AI collaboration documents must separate **Prepared by**, **Generation platform**, **Underlying model**, and **Generation session or application**. If the model is not explicitly disclosed in the session record, it must be listed as **Not specified** rather than inferred. Platform attribution must not overwrite a separately commissioned artifact created in another application.

Mission and Directive

Mission Statement

The Adversarial Cross-Review and Iterative Improvement Protocol exists to make OAL-generated work inspectable, testable, comparable, and revisable before it is treated as stable, operational, or canonical. It does not seek consensus as proof of correctness. It identifies assumptions, disagreements, omissions, unperformed tasks, and falsifiable claims; records them without averaging them away; and converts them into bounded tests, decisions, hygiene corrections, or explicit open questions.

Operating Directive

Review the work as if it may contain important insights and important failures at the same time. Preserve useful structures, identify unsupported claims, test the smallest meaningful part, report failure honestly, and do not allow documentation volume, AI agreement, institutional formality, or authorial investment to stand in for evidence.

Objectives

- Create a repeatable method for reviewing one document, a family of versions, or a repository-scale body of work.
- Separate what the corpus says from what a reviewer concludes and from what an executed exercise actually finds.
- Expose assumptions, dependencies, internal contradictions, omissions, edge cases, and unperformed tasks.
- Turn disagreement into traceable tests, decisions, hygiene corrections, or explicitly unresolved questions.
- Require evidence proportionate to the claim while avoiding unnecessary simulations or procedural overhead.
- Preserve criticism, dissent, null results, and failed tests as useful iteration records.
- Support GitHub publication through stable identifiers, version history, cross-links, and reproducible artifacts.
- Enable multiple AI systems and human contributors to produce genuinely independent, comparable reviews.

Scope, Limits, and Non-Goals

Each commissioned review must define the exact topic, what is inside and outside scope, the category of work, the intended users, the authority of the reviewer, whether tests can be run, and any safety, privacy, legal, ethical, creative, or canon boundaries.

The protocol does not:

- Determine whether a project is globally “good” or “bad.”
- Treat agreement among systems or reviewers as evidence.
- Make a proposal canon, operational policy, or institutional authority.
- Require Python or simulation where another method is more appropriate.
- Retroactively rewrite sources or conceal conflicting versions.
- Permit personal attacks, motive attribution, or dismissive criticism.
- Treat internal simulation or fictional material as direct real-world proof.
- Require every open question to be resolved in one iteration.

Protocol Overview

Function	Core question	Required output
Corpus control	What was actually reviewed, and what was missing?	Source register and declared-input list
Epistemic separation	What is source, finding, evidence, assessment, proposal, or open?	Consistent labels and claim discipline
Independent critique	What does each artifact contribute, assume, omit, contradict, or leave undone?	Artifact-by-artifact adversarial review
Discrepancy management	Which conflicts should be tested, decided, repaired, deferred, quarantined, retired, or preserved?	Discrepancy register
Test design	What is the smallest legitimate exercise capable of challenging the claim?	Preregistration and validation plan
Iterative action	What specifically changes because of the review?	Improvement register and decision log
Publication traceability	Can a later reader reproduce the record and understand version relationships?	Stable IDs, change log, linked artifacts, and status notices

Lifecycle

1. Commission the review with a stable ID, declared scope, source packet, review tier, and non-canon status.
2. Freeze or identify the corpus window.
3. Conduct independent first-pass reviews before sharing newly generated reviews.
4. Create source, assumption, dependency, discrepancy, testability, and action registers.
5. Select the smallest high-leverage tests; preregister before execution.
6. Execute, preserve raw artifacts, and report PASS, FAIL, MIXED, INDETERMINATE, or NOT RUN.
7. Cross-review the reviews and results, documenting revised or withdrawn positions.
8. Make evidence-resolvable changes, elevate value or authority questions, and preserve dissent.
9. Publish “What changed because of this review?” and commission another iteration only when warranted.

Definition of Adversarial

“Adversarial” means actively seeking conditions under which a claim, mechanism, interpretation, score, workflow, or governance rule would fail. It requires fair representation of the source, specific criticism, explicit evidence standards, and willingness to revise the reviewer’s own position. It does not mean hostility toward authors or reflexive rejection.

Review Tier Model

Tier	Best use	Typical output	Minimum validation expectation
Tier 0 — Intake and catalog	Fragments, prompts, notes, small early documents	One- to two-page review card	No exercise required; identify claim type, dependencies, and next disposition
Tier 1 — Adversarial document review	Single framework, policy, architecture, or concept	Focused review dossier	At least one proposed falsifier, verification task, or edge-case test
Tier 2 — Cross-review	Multiple AI versions or competing artifacts	Cross-version review with registers	At least one bounded pilot, tabletop, source audit, dataset check, or comparison where feasible
Tier 3 — Execution dossier	High-value recurring frameworks or public systems	Evidence dossier and reproducibility artifacts	Preregistered tests, controls, results, limitations, and revision decisions
Tier 4 — Canon or operational decision	ONESTAMP consideration or material deployment	Ratification or decision packet	Evidence trail, dissent, known limits, rollback conditions, and post-adoption trigger

Tier Selection

Consider impact, authority, maturity, complexity, reversibility, evidence burden, and publication exposure. Select the tier before work begins. Advance only when scope, risk, authority, maturity, or intended use justifies the additional burden.

****Anti-bloat rule:**** If a review cycle does not produce a clarified claim, a new or completed test, a repaired artifact, a decision, or a justified preserve/quarantine/retire action, it should not automatically expand into another documentation cycle.

Corpus and Provenance Control

No review is stronger than its declared inputs. Every reviewer must identify what was read in full, partially reviewed, unavailable, referenced secondhand, or excluded. Missing artifacts are findings about provenance and reproducibility, not evidence of the missing artifact's contents.

Required Source Register

Field	Required entry
Source ID	Stable short identifier such as S-01
Artifact title	Exact title shown in the source
Preparer	Named person, AI system, team, or unknown
Generation platform	Application or platform where the artifact was produced, if known
Underlying model	Only if explicitly disclosed; otherwise Not specified
Date and version	Exact metadata; note conflicts
Artifact type	Prompt, memo, framework, code, dataset, simulation, policy, report, or transcript
Claimed stage	What the artifact says it delivers
Observed stage	What is actually present
Repository location	URL, path, branch, tag, or commit
Review completeness	Full, partial, metadata only, inaccessible, or secondhand
Dependencies and gaps	Missing attachments, code, data, citations, versions, or relationships

Provenance Exceptions

- Label missing files MISSING and explain the limit.
- Label claims known only from a summary SECONDHAND.
- Record UNKNOWN instead of inferring authorship.
- Open a HYGIENE discrepancy for duplicate IDs or versions.
- Distinguish code inspection from code execution.
- Preserve the reviewed version when a source changes during review.
- Do not infer an underlying model from a platform or user interface.
- Keep independent outputs from different applications under separate IDs.

Epistemic Labeling Standard

Label	Meaning
SOURCE	A claim, instruction, design element, or quotation contained in a named source artifact
FINDING	A result produced by an executed exercise, analysis, simulation, prototype, or test; limited to the method and assumptions
EVIDENCE	A claim supported by an external cited source, dataset, official record, or verifiable artifact
ASSESSMENT	The reviewer's reasoned evaluation or adversarial judgment
PROPOSAL	A new option, design, process, test, or revision recommended for later iteration
OPEN	Not resolvable using available materials
LIMITATION	A constraint on interpreting a claim, method, or result
DECISION	A value, canon, authority, scope, or priority question reserved for designated authority
HYGIENE	A documentation, naming, attribution, versioning, provenance, citation, or repository correction
NOT RUN	A planned exercise or task that was not executed, with the reason recorded

Use “supports,” “fails to support,” “is consistent with,” “was observed under these assumptions,” or “remains unresolved.” Reserve “proves” for formal proof. Never call a prediction, designed behavior, or expected output a finding.

Independent Review Method

Each artifact must receive a fair summary before critique. The reviewer must identify unique contributions and strengths as carefully as weaknesses.

Required Questions

1. What does this source contribute that others do not?
2. What does it make clearer, safer, more executable, or more useful?
3. What does it assume rather than test?
4. Which claims are unsupported, ambiguous, inconsistent, or overconfident?
5. Which definitions, thresholds, roles, procedures, records, or exit conditions are missing?
6. What could fail if the framework were applied as written?
7. Which dependencies, stakeholders, edge cases, incentives, or abuse paths are omitted?
8. What tasks does the source imply should be done but has not done?
9. Can its central claims be tested now? If yes, how? If not, why not?
10. What should the next iteration preserve, revise, quarantine, defer, or retire?

A reviewer must apply the same standard to its own prior output. When evidence or another review undermines an earlier position, explicitly retain, revise, or withdraw it.

Assumption and Dependency Audit

ID	Assumption or dependency	Location	Type	Why it matters	Support	Failure mode	Resolution
A-01	Exact assumption	Source and section	Technical, governance, ethical, narrative, economic, or operational	Consequence	SOURCE, EVIDENCE, or OPEN	What fails if wrong	TEST, DECISION, HYGIENE, or DEFER

Audit definitions, scope, authority, data, resources, incentives, stakeholders, edge cases, reversibility, safety, privacy, law, canon, external systems, and assumptions inherited from earlier prompts or summaries.

Discrepancy Management

Do not average disagreements into vague consensus. Preserve the competing positions, explain why the difference matters, identify what is known, and assign a resolution path.

Path	Meaning
TEST	Narrow through a predefined exercise, prototype, source analysis, or comparison
DECISION	Resolve through designated human authority because values, canon, authority, scope, or priorities are involved
HYGIENE	Correct naming, metadata, attribution, versioning, registry, citation, or repository problems
DEFER	Retain because responsible resolution is not possible in the current iteration
QUARANTINE	Prevent from shaping shared assumptions until tested or independently reviewed
RETIRE	Remove because the claim, mechanism, or rule failed or is no longer justified
PRESERVE	Retain provisionally with limits and a review trigger

A discrepancy closes only when linked to a test result, decision, correction, retirement rationale, or closure artifact. Rewording positions until they appear compatible is not closure.

Testability and Validation

Validation must match the claim. Software may require execution; provenance may require repository history; classification may require blinded raters; governance may require tabletop scenarios and incentive analysis; canon may require an Architect decision rather than empirical testing.

Testability Status

Status	Definition
Testable now	A feasible bounded method and required materials are available
Partly testable	Some components can be tested, but central real-world, normative, metaphysical, or long-horizon claims cannot
Not testable now	Data, access, conditions, authority, or valid methods are unavailable in this iteration
Not empirically testable	The claim is definitional, philosophical, aesthetic, canonical, or normative and should be classified accordingly

Validation Methods

- Tabletop scenario for governance, policy, workflows, handoffs, and crises.
- Source audit for evidence, citations, provenance, and historical claims.
- Red-team prompt test for AI behavior, robustness, safety, and instruction hierarchy.
- Reproducibility check for code, simulations, analyses, and build pipelines.
- Version-difference analysis for competing documents and evolving specifications.
- Dependency-map review for architectures, organizations, and repositories.
- User or stakeholder journey for usability, access, burden, and handoffs.
- Blinded classification for scoring, detection, and distinguishability.
- Data-quality check for datasets, registries, metrics, and reporting.
- Counterfactual comparison for interventions, rules, and alternative designs.
- Code test suite for software behavior and regressions.
- Policy stress test for edge cases, incentives, abuse, and emergencies.
- Timeline or provenance audit for authorship, attribution, sequence, and identifiers.
- Internal-consistency test for definitions, authority, stages, and cross-references.
- Prototype acceptance test for implementation feasibility and specified behavior.

****Minimum rule:**** Run the smallest test appropriate to the claim. Do not require code when a source audit, tabletop, blinded review, or Architect decision is the legitimate method.

Preregistration and Exercise Design

Tier 2 or higher reviews should execute at least one bounded exercise unless they document why no legitimate exercise can be run. Preregistration occurs before execution so success criteria, parameters, and interpretations are not rewritten after results are known.

Preregistration Block

Field	Required content
Test ID	Stable identifier such as X-01
Purpose	Decision or uncertainty informed
Claim or hypothesis	Exact proposition challenged
Falsifier	Outcome that would materially challenge it
Inputs	Documents, data, prompts, code, scenarios, or participants
Parameters	Assumptions, thresholds, seeds, model versions, and configuration
Baseline/control	Current practice, no intervention, alternate rule, or benchmark
Comparison	Competing design, version, or counterfactual
Metrics	Measures and calculation rules
Thresholds	PASS, FAIL, MIXED, or INDETERMINATE criteria
Roles	Generator, evaluator, red-team reviewer, and decision owner
Limitations	Threats to validity and prohibited generalizations
Date/version	Preregistration date and immutable version or commit
Artifacts	Expected output files and locations

Quality Rules

- Use a baseline, control, counterfactual, or comparison whenever feasible.
- Separate scenario generation from evaluation for detection, classification, ranking, or indistinguishability.
- Do not let the expected-answer author be the only scorer.
- Measure benefit and harm in governance or trade-off claims.
- Disclose parameter and prompt sensitivity.
- Preserve outputs, code, prompts, settings, seeds, and model metadata where applicable.
- Report errors, null results, failed runs, deviations, and indeterminate outcomes.
- Do not change thresholds after execution without labeling the analysis exploratory.

When independent evaluation is unavailable, partial safeguards include sealed ground truth, delayed reveal, randomized cases, a second model that did not generate the cases, blinded scoring, predefined rubrics, or later human audit. These safeguards do not create full independence; the limitation remains visible.

Results and Interpretation

Status	Definition	Required interpretation
PASS	Result matched the preregistered prediction under stated conditions	State narrow support without overgeneralizing
FAIL	Result contradicted the preregistered prediction	Identify the challenged claim and needed change
MIXED	Measures, cases, branches, or parameters conflicted	Separate conditions; do not average away differences
INDETERMINATE	Exercise ran but could not discriminate among alternatives	Explain whether method, data, construct, power, or ground truth was inadequate
NOT RUN	Planned exercise was not executed	State why, what was completed, and whether it remains valid later

A PASS does not prove a theory. A FAIL does not automatically invalidate the project. A simulation describes the implemented model, not reality by default. Reproducibility does not validate assumptions. Null and indeterminate results can reveal inadequate metrics or distinctions.

Iterative Improvement and Decisions

Every material weakness, discrepancy, failed test, or open dependency should produce an action, decision request, defer/quarantine record, or documented reason no action is warranted.

Improvement Register

ID	Action	Trigger	Priority	Owner	Closure evidence	Target	Status
I-01	Specific action verb and object	D-03, A-04, or X-02	High, Medium, or Low	Assigned owner	Required artifact or result	Iteration	OPEN

Use concrete verbs: define, separate, test, compare, rename, register, preregister, audit, prototype, red-team, archive, retire, quarantine, document, translate, pause, or escalate.

Architect Decision Log

ID	Decision required	Options	Consequences	Recommendation	Default	Trigger
AD-01	Value, canon, scope, authority, or priority question	A, B, or C	Trade-offs	Optional	Provisional state	When needed

Only questions that cannot be resolved through testing, verification, or hygiene should be elevated. The protocol should reduce decision burden rather than transfer every uncertainty upward.

What Changed Record

Every completed cycle must state definitions revised, claims downgraded, tests added, mechanisms changed, files renamed, decisions made, items quarantined or retired, and questions still open. If nothing changed, explain the value produced and why another cycle is justified.

Multi-AI Collaboration Workflow

Independent versions are most useful when they share enough structure to compare but retain different analytical burdens. Giving every system the same detailed reasoning prompt may create artificial convergence. Use a common source packet and output core with differentiated lanes.

Common Source Packet

- Exact topic, research question, and review tier.
- Frozen corpus or commit range.
- Source register and known missing artifacts.
- Stable document ID and versioning rules.
- Separate preparer, platform, model, and session fields.
- Epistemic labels and result definitions.
- Required registers and quality gates.
- Non-canon and non-ratified status notice.
- Rule that every input used must be disclosed.
- Rule that agreement is not evidence.

Suggested Lanes

Lane	Primary burden
Perplexity	Evidence audit, source verification, analogies, testability, and translation implications
ChatGPT	Measurement design, controls, scoring, falsifiers, synthesis, and cross-version consistency
Claude	Non-assumptions, ambiguity preservation, ethics, scope, and restraint against over-institutionalization
Gemini	Systems architecture, feasibility, dependencies, data flows, and integration risks
Microsoft Copilot	Operationalization, documentation hygiene, implementation, version control, and developer workflow
Grok	Red-team critique, incentive failures, capture, abuse cases, edge cases, and adversarial scenarios
Human Architect	Canon, priorities, aesthetic intent, authority, values, and ratification status

These are starting assignments, not permanent capability claims. Rotate lanes to detect role-induced blind spots and model drift.

Six Phases

1. Commission with shared packet, lane, tier, outputs, and immutable source window.
2. Produce independent first passes before reading newly commissioned reviews.
3. Cross-review convergence, disagreement, revised positions, and withdrawn recommendations.
4. Select one to three high-leverage tests rather than every proposed exercise.
5. Preregister, execute, preserve artifacts, and report result status.
6. Decide, archive, update registers, preserve dissent, and commission another iteration only when justified.

Stage Gates and Quality Gates

Gate	Purpose	Required output	Stop rule
G1 — Revise	Identify preserve, revise, reject, quarantine, or defer positions	Revision memo, corpus register, discrepancies	No testing until claims and scope are explicit
G2 — Preregister	Define hypotheses, falsifiers, inputs, controls, metrics, and roles	Versioned preregistration	No confirmatory execution without dated record
G3 — Execute	Run the smallest tests challenging central claims	Raw artifacts, results, failures, and limits	No canon recommendation based only on planned work
G4 — Translate	Convert evidence into design, policy, operational, narrative, or canon options	Options, trade-offs, and dissent	Do not present analogy or simulation as direct real-world proof
G5 — Review	Adversarially review results and revisions	Contradiction, limitation, and revision register	No ratification without unresolved-risk disclosure and review triggers

Submission Failure Conditions

- Inputs are not declared.
- A prediction or designed behavior is called a finding.
- No source window, version, or identifier is provided.
- Preparer, generation platform, and underlying model are conflated.
- An undisclosed model is inferred from the platform.
- Independent artifacts from separate applications are merged under one attribution.
- Disagreements are erased without resolution records.
- Tier 2+ work has no exercise and no legitimate explanation.
- Classification relies only on the same unblinded creator and evaluator.
- Limitations are omitted or separated from constrained claims.
- Canon or operational use is recommended without evidence, dissent, risks, and review triggers.
- Reproducibility is claimed without required artifacts.
- Criticism attacks contributors rather than claims, mechanisms, methods, incentives, or omissions.

Identifiers and Repository Practice

Artifact	Pattern	Example
Cross-review	OAL-ACR-I{iteration}-G{gate}-{AI}-v{X.Y}	OAL-ACR-I1-G1-PERPLEXITY-v0.2
Source artifact	OAL-{project}-{topic}-S{source}-{AI}-v{X.Y}	OAL-AIU-PROMPT-S3-CLAUDE-v0.2
Discrepancy	D-{two digits}	D-07
Assumption	A-{two digits}	A-12
Test	X-{two digits}	X-04
Improvement	I-{two digits}	I-09
Architect decision	AD-{two digits}	AD-03
Hypothesis	{LAB}-H-{number}	ECL-H-02
Red-team issue	RT-{two digits}	RT-05
Preregistration	PRE-{topic}-{date}-{version}	PRE-ENDCYCLE-2026-10-01-v0.1

Platform and Model Attribution Standard

Field	Question answered	Example for this document
Prepared by	Which AI contribution does this represent?	Perplexity AI
Generation platform	Where was the interaction and file generation performed?	Perplexity AI platform / Perplexity Computer session
Underlying model	Which model produced it, if explicitly disclosed?	Not specified
Session/application record	Which conversation establishes provenance?	Perplexity AI conversation dated October 1, 2026
Separate comparative artifact	What independently commissioned version must remain distinct?	ChatGPT web-application version

Recommended repository bundles include an index, DOCX and PDF, machine-readable registers where useful, preregistrations, code and configuration, raw and processed outputs kept distinct, CHANGELOG, DECISIONS, DISSENT, and appropriate licensing and contribution guidance.

Feedback and Cautions

The End Cycle Laboratory exercise demonstrates the value of running small pilots against mature conceptual frameworks. Its transferable contribution is not a final answer, but exposure of the distance between designing an exercise and executing one. The reusable process must preserve that distinction while remaining lighter than the work it evaluates whenever possible.

Constructive Criticisms

- Comprehensive review can become self-protecting documentation; tiering and stopping rules are necessary.
- A label set requires examples and consistency checks or reviewers will apply it differently.
- Adversarial review should challenge claims, assumptions, methods, incentives, edge cases, and missing tasks—not motives or intelligence.
- Simulation findings and design recommendations must remain visibly separate.
- Generator–evaluator separation is a strong default, but solo workflows need disclosed fallback safeguards.
- A review needs a stopping rule tied to changed knowledge or action.
- Publication should state what changed because of review.
- Common prompts improve comparability but can induce convergence.
- The protocol itself must be tested for completion time, actionable defects, duplicate findings, unresolved decisions, and later outcome improvement.
- Retirement must remain possible; preserving every mechanism indefinitely recreates the accumulation problem.
- Platform and model attribution must be independently controlled metadata fields.

Questions for Team Review

- Are five tiers sufficient, or should risk and canon status be separate dimensions?
- Which labels are essential, redundant, or missing?
- Who may assign, change, or close discrepancy paths?
- What minimum evidence is required before operational use?
- Which tests require human rather than AI evaluation?
- How should inaccessible, private, proprietary, or safety-sensitive sources be handled?
- What repository metadata is mandatory?
- When should dissent be archived as resolved?
- How should the team test whether this protocol improves quality rather than output volume?
- What authority, if any, should a cross-review body have?
- Should OAL use “Prepared by” for the platform identity, the presented AI collaborator identity, or both in separate fields?

Reusable Document Template

A Tier 2–4 cross-review should include:

1. Title page with document ID, version, date, preparer, platform, model disclosure status, session record, status, repository, corpus window, tier, and inputs.
2. Executive summary distinguishing executed work from proposals.
3. Table of contents.
4. Mission and directive.
5. Scope, limits, and non-goals.
6. Source corpus and input declaration.
7. Epistemic labels.
8. Review method and independence disclosure.
9. Neutral topic or framework overview.
10. Independent source analysis.
11. Assumption and dependency audit.
12. Discrepancy register.
13. Testability assessment.
14. Preregistration and exercise plan.
15. Results and findings.
16. Iterative improvement register.
17. Architect decision points.
18. What Changed section.
19. Next-iteration commission prompt.
20. Appendices with methods, artifacts, source notes, version history, provenance, and dissent.

Copy-Ready Commission Prompt

STATUS: PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED.

Document ID: OAL-ACR-I{iteration}-G1-{AI}-v0.1

Date: [Insert date]

Prepared for: Nathan Lewis II, Architect of Record, Omniversal Architect Labs

You are an independent collaborator conducting an adversarial cross-review for Omniversal Architect Labs. Review the supplied corpus concerning **[TOPIC]**. Your objective is not to create consensus, praise the corpus, or defend prior work. Identify what the corpus establishes, what it only proposes, what it assumes, what it fails to test, where it contradicts itself, and what should be tested or decided next.

Create a detailed document suitable for DOCX and PDF export. Include:

1. Title page with document ID, version, date, preparer, generation platform, underlying model if explicitly disclosed, session/application record, status, tier, source window, and declared inputs.
2. Executive summary and table of contents.
3. Mission statement or directive.
4. Scope, limits, and non-goals.
5. Corpus inventory, including missing inputs, uncertain authorship, unavailable code/data, and secondhand descriptions.
6. Labels: SOURCE, FINDING, EVIDENCE, ASSESSMENT, PROPOSAL, OPEN, LIMITATION, DECISION, HYGIENE, and NOT RUN.
7. Neutral topic overview.
8. Independent review of each artifact: contribution, strengths, weaknesses, assumptions, inconsistencies, unperformed tasks, and preserve/revise/quarantine/defer/retire recommendations.
9. Assumption and dependency audit.
10. Discrepancy register assigning TEST, DECISION, HYGIENE, DEFER, QUARANTINE, RETIRE, or PRESERVE.
11. Testability assessment.
12. At least one minimal validation exercise where legitimate, or a documented reason none can be run.
13. Results table using PASS, FAIL, MIXED, INDETERMINATE, or NOT RUN.
14. Improvement register with actions, owners, closure evidence, and target iterations.
15. Architect decisions limited to nonempirical questions.
16. "What changed because of this review?"
17. Copy-ready prompt for the next iteration.
18. Appendices for methods, artifacts, source notes, history, provenance, and dissent.

Mandatory rules:

- Do not represent prediction or designed behavior as a FINDING unless executed.
- Do not treat AI agreement as evidence.
- Do not use page count, complexity, institutional form, or definition count as proof.
- Do not silently interpret ambiguity as settled fact.
- List every input used.
- Distinguish preparer, platform, underlying model, and session/application.
- Do not infer a model not explicitly disclosed.
- Keep independently commissioned platform outputs under separate IDs and histories.
- Include a baseline, control, counterfactual, or comparison for feasible exercises.
- Do not use the same unblinded creator for ground truth and final evaluation without disclosure and checking.
- Preserve dissent and unresolved questions.

- Separate external evidence, source/canon material, simulation findings, assessments, and decisions.
- State limitations near the claims they constrain.
- End with specific next actions.

Required tone: rigorous, fair, adversarial but constructive, specific, version-controlled, and suitable for GitHub publication.

Appendix A — Blank Registers

Source Register

ID	Title	Preparer	Platform	Model disclosure	Date/version	Type	Location	Review status	Notes
S-01									
S-02									

Discrepancy Register

ID	Discrepancy	Positions	Consequence	Status	Path	Owner
D-01						
D-02						

Testability Register

Claim	Status	Method	Threshold	Control	Limitation	Result

Improvement Register

ID	Action	Trigger	Priority	Owner	Closure evidence	Iteration	Status
I-01							
I-02							

Decision Log

ID	Question	Options	Consequences	Recommendation	Default	Trigger
AD-01						

Appendix B — Publication Checklist

- Document ID, version, date, preparer, status, and tier are present.
- Generation platform and session/application are separate from preparer.
- Underlying model is listed only when explicitly disclosed; otherwise Not specified.
- Comparative outputs from other platforms have separate IDs and histories.
- Status is marked proposed, not canon, and not ONESTAMP-ratified unless changed formally.
- Source window and every input are declared.
- Missing and secondhand materials are identified.
- Executive summary distinguishes executed work from proposals.
- Table of contents is updated.
- Epistemic labels are consistent.
- Each source receives a fair summary before criticism.
- Assumptions, dependencies, omissions, and unperformed tasks are recorded.
- Material disagreements appear in the discrepancy register.
- Resolution paths are assigned without forced averaging.
- Tier 2+ includes an exercise or explains why none is legitimate.
- Preregistration predates confirmatory execution.
- Controls, baselines, counterfactuals, or comparisons are present where feasible.
- Generator–evaluator circularity is prevented or disclosed.
- Raw artifacts, prompts, settings, seeds, and model versions are preserved as applicable.
- Results use PASS, FAIL, MIXED, INDETERMINATE, or NOT RUN.
- Limitations appear near constrained claims.
- Actions have owners, closure evidence, and target iterations.
- Architect decisions are limited to nonempirical questions.
- The document states what changed because of review.
- Dissent and withdrawn positions are preserved.
- GitHub links, commits, and artifact names are valid.
- DOCX and PDF match materially.
- No placeholder remains unless intentionally published as a template.

Appendix C — Proposed Disposition

This v0.2 Perplexity AI document is submitted for comparative pass-through. It should be reviewed alongside independently prepared versions from other AI systems and human collaborators. The team should record conflicts rather than merging them prematurely, identify elements suitable for a common core, and assign at least one test of the protocol itself before treating it as an OAL standard.

Disposition item	Proposed status
Mission and directive	PRESERVE provisionally; invite alternatives
Tier model	TEST through differently sized subjects
Epistemic labels	PRESERVE provisionally; test reviewer consistency
Differentiated lanes	TEST against identical-prompt condition
Stage gates	PRESERVE for Tier 2–4; simplify for Tier 0–1
Identifiers	HYGIENE review against current OAL conventions
Platform/model attribution	PRESERVE; audit every comparative artifact before publication
Quality gates	Cross-review for burden and enforceability
Canon status	Remain PROPOSED until explicit Architect and ONESTAMP action, if applicable