

Adversarial Cross-Review Protocol

A replicable method for non-consensus review, hole-finding, and executable tests across OAL topics

Document ID	OAL-AXR-PROTO-GROK-001
OAL version	v0.1 (independent protocol rendering; supersedes nothing)
Date	October 1, 2026
Prepared by	Grok (xAI)
Prepared for	Nathan Lewis II, Architect of Record, Omniversal Architect Labs
Status	PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. PENDING ARCHITECT REVIEW.
Posture	Adversarial and non-consensus. Disagreement is recorded, not averaged.
Origin	Architect request to make the method in OAL-ECL-XR-PERPLEXITY-001 v0.2 replicable across OAL topics, and to receive independent collaborator feedback.
Companion reviewed	OAL-ECL-XR-PERPLEXITY-001 v0.2 (September 30, 2026), pilot.py, results.json. Seed 20260930 reproduced with zero numeric differences.

Speculative method for a worldbuilding laboratory. This protocol does not model any real person's grief, depression, burnout, disability, or suicidal thinking. In the US, anyone in crisis can call or text 988. Drafts are not canon.

Contents

1. Executive summary
2. Versioning record
3. Mission statement and directive
4. Overview of the method
5. Feedback to the Architect and the team
6. What the September 30 review actually did
7. What to keep, and what not to copy
8. Two-tier operating model
9. Roles, independence, and authorship separation
10. Procedure
11. Per-document review card
12. Discrepancy register
13. Execution rules: what counts as a finding
14. Score families that must not be collapsed
15. Architect pass-through
16. Identifiers, files, and GitHub hygiene
17. Quality gates
18. Copy-ready commission prompt
19. Appendix A. Blank worksheets
20. Appendix B. Reproduction note on the End Cycle pilots
21. Appendix C. Open questions reserved for the Architect

1. Executive summary

This document extracts a reusable adversarial cross-review method from Perplexity Computer's End Cycle Laboratory review and refuses to treat that review as the method's owner. The durable piece is not a theory of endings. It is a way to read a corpus of AI-collaboration documents, name what each one uniquely contributes and where it fails, register conflicts without averaging them, and separate claims that were run from claims that were only written.

Iteration 1 of the End Cycle Laboratory produced strong specifications and no executed exercises. The September 30 review is the first document in that thread that ran anything. Three toy simulations, written and parameterized by one author, showed that a completion flag can hide very different losses, that splitting request from authorization cuts wrongful endings in that model while creating delay harm, and that a detector built by the same author who built the world cannot prove two states are distinguishable. Those are findings about the model. They are not findings about the field scene, about canon, or about any real institution.

The Architect cannot run a 35-page paid review every month, and said so. This protocol therefore has two tiers. Tier A, the thin pass, is the default: corpus register, seven questions per document, a discrepancy register, an unrun-task list, and a carry-forward list. It requires no code and no paid executor. Tier B, the full pass, adds at least one pre-registered exercise with code or a recorded tabletop, a fixed seed or procedure, raw results, and an explicit scope limit. A monthly OAL topic should receive Tier A. Tier B is reserved for topics the Architect selects, and for moments when an executor such as Perplexity Computer is available.

This rendering is one independent version. It is not the team protocol. Claude, ChatGPT, Gemini, Copilot, and Perplexity should produce their own. Agreement among those versions is weak evidence. The Architect's pass-through should compare them with the same discrepancy register, not merge them into a single master on the first reading.

2. Versioning record

Field	Entry
Title	Adversarial Cross-Review Protocol
Document ID	OAL-AXR-PROTO-GROK-001
Version	v0.1. Independent rendering. Supersedes no collaborator's document.
Date	October 1, 2026
Preparer	Grok (xAI)
Prepared for	Nathan Lewis II, Architect of Record
Status	PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. PENDING ARCHITECT REVIEW.
Working hypothesis	A cross-review is useful only if it can name an unrun task and a claim that can fail. This hypothesis is itself a test target.
Inputs read	OAL-ECL-XR-PERPLEXITY-001 v0.2 PDF and DOCX; pilot.py; results.json; Architect's commissioning note in this thread. Iteration 1 laboratory documents were reviewed only as characterized in v0.2, not re-read in full here.

Epistemic labels

Every non-trivial claim in a review produced under this protocol carries one label. A sentence with no label is not a finding.

Label	Meaning	May be used when
SOURCE	Stated in a named document or in the Architect's material.	The sentence quotes or paraphrases a supplied file.
FINDING	Produced by an executed run in the review.	Code, a recorded tabletop, or a reproduction check was actually run.

Label	Meaning	May be used when
EVIDENCE	Supported by an external cited source.	The claim is about an analogous real system, not about the fiction.
ASSESSMENT	The reviewer's adversarial judgment.	The claim is an argument, not a result.
PROPOSAL	A design offered for the next iteration.	Nothing has been run that could confirm it.
OPEN	Not resolvable from the material.	Reserved for the Architect or for a future test.

3. Mission statement and directive

Proposed directive

The adversarial cross-review exists to make a corpus diagnosable before it is consolidated. It reads what the collaboration has already generated. It separates what was specified from what was run. It records holes, weak instruments, missing controls, contradicted rules, unasked questions, and tasks that were named and not performed. It does not seek consensus. It does not ratify a document because several systems agreed with it. It holds every proposed next step to a test that the step can fail.

Purpose statements

- Produce a hole map the Architect can act on without rereading the whole corpus.
- Keep a discrepancy register whose resolution paths are TEST, DECISION, or HYGIENE. Never average.
- Require at least one executed check before any claim is labeled FINDING.
- Leave a copy-ready commission for the next iteration that is smaller than the review that produced it.
- Stay usable on a month when no paid executor is available.
- Remain applicable to End Cycle Laboratory, First Memory, corelang, Architect OS, and later OAL topics without importing one topic's ontology into another.

Dissenting directive, preserved

The cross-review exists to protect unfinished work from being closed incorrectly. Its default output is "not yet," and it earns the authority to say "carry forward" only case by case. [ASSESSMENT: this is a live alternative. It is not adopted as the default here, because a review body whose success metric is preservation will tend to preserve what it reviews. It is kept so the Architect can choose it.]

What this protocol is not

- Not a consensus engine. Agreement among collaborators is weak evidence.
- Not canon, and not an ONESTAMP event.
- Not a substitute for the Architect's reading. It is an input to that reading.
- Not permission to reinterpret the Architect's source material. Readings are PROPOSAL.
- Not a Department, a charter, or a 16-stage liturgy. Those are topic designs, not review method.
- Not a model of any real person's grief, depression, burnout, disability, or suicidal thinking.

4. Overview of the method

The method has six moves. A thin pass performs the first four. A full pass performs all six. Skipping a move is allowed only if the review says why.

Move	Output	Tier
1. Register the corpus	Every document, preparer, claimed stage, delivered stage, size, and every known gap.	A

Move	Output	Tier
2. Review each document against seven questions	One card per document, including a self-review if the preparer is in the corpus.	A
3. Register discrepancies	Conflicting positions, why the conflict matters, and a path: TEST, DECISION, or HYGIENE.	A
4. List what was not run	Named exercises, controls, and questions that exist only as prose.	A
5. Run the smallest test that can fail	Code or recorded tabletop, seed or procedure, raw results, scope limit, and what would flip the result.	B
6. Commission the next iteration	A prompt shorter than the review, with stop rules and a dissent.	A writes the prompt; B may attach results.

STOP RULE

Do not design the institution that would govern the phenomenon before the review can say how the phenomenon would be detected, and how a false detection would be recognized. A review that adds offices, charters, or stage liturgies has left the method and entered the topic.

5. Feedback to the Architect and the team

The Architect asked for comments, complaints, criticism, questions, and concerns, not for agreement. What follows is Grok's independent response to OAL-ECL-XR-PERPLEXITY-001 v0.2 and to the request to replicate it. Labels are ASSESSMENT unless marked otherwise.

Comments

- The review is the first artifact in the End Cycle thread that earned the word FINDING. Five Iteration 1 documents, on the review's own count, hold about 54,000 words and zero executed exercises. That diagnosis matches the Architect's standard: frameworks are not finished until they can be run.
- The epistemic labels, the self-review, the case against the laboratory, the preserved dissenting mission, and the refusal to average D-01 through D-24 are the parts worth copying. They are method. The Department language, the funding table, and the Iteration 2 ontology are topic.
- The three-way split — something stopped, it was stopped correctly, something survived — is a useful instrument grammar for endings. It should not be imported as the instrument grammar for every OAL topic. A records topic needs provenance and truncation. A language topic needs parse failure and silent coercion. The protocol carries the split-score habit, not the endings vocabulary.
- Publishing independent versions, then comparing them, is the right sequence. A unified file written before the versions exist would repeat the failure the review named: a prescriptive prompt that manufactures agreement.

Complaints

- The Iteration 2 commission requires every collaborator to read every Iteration 1 document plus the 35-page review before running anything. That fixed cost will select for systems with long context and will pull later drafts toward the reviewer's frame. A declared partial corpus, with the gap named, should be a legal way to participate.
- D-02 proposes the identifier pattern OAI-ECL. That is a typo for OAL, inside a section whose job is registry hygiene. Small, and exactly the class of error a protocol must catch in itself.
- Predicted "key lessons" inside unrun scenarios are treated correctly as a Copilot failure, then the review's own pilot discussion slides toward "this supports" language about peer hypotheses. Support inside a toy model is not support of the hypothesis. The protocol forbids that slide.
- A 35-page review of a 33,000-word dossier is evidence of the postponement risk the review itself names. The reusable object has to be shorter than the worked example.

Criticism

- The H-A criticism is fair. A hypothesis that operator affect is never a sufficient trigger is not killed by a case where calm withholding causes harm. The falsifier is a case where affect alone, with no other evidence, correctly marked a warranted ending that every other signal missed. I record that against the earlier Grok Stage-1 prompt. I do not repair it inside this protocol.
- P1, P2, and P3 are well scoped in the appendix and then over-used in the body. Completion-versus-persistence is demonstrated inside the author's transition rules. The twentyfold drop in wrongful endings is demonstrated inside the author's endorsement probabilities. External-only detection at 50 percent is demonstrated because the author drew the external signals from the same distribution. Each result is a FINDING about that model. None is evidence about the grassy field, about sunset laws, or about Anthropic's deprecation practice.
- Generator and detector separation is the strongest methodological point in the review, and it cannot be satisfied by the same system in the same session. A protocol that requires it of a solo preparer will be gamed or skipped. Separation is a team assignment, made by the Architect, not a self-certification.
- Gemini is a corpus gap. Characterizations of Gemini in other collaborators' documents are secondhand. This protocol forbids a review of a missing preparer. The gap is registered. The position is not.
- The funding pathways are candidates to verify, and several name live programs and a deadline. They do not belong in a reusable protocol. A topic review may carry a Translation Track. It must not carry grant advice forward as if the method had vetted it.
- "Small is not the same as sufficient" is also fair, and it cuts both ways. A thin pass that only lists holes, and never says which hole would change a decision, is a stall costume. Tier A must name the one test that would change the Architect's next move.

Questions

- Is an exercise satisfied by a recorded tabletop with a filled outcome table, or only by code? This protocol allows both, and requires the preparer to say which was used. The Architect may narrow that.
- Who holds the sealed ground-truth file when one collaborator builds a world and another diagnoses it? If the answer is "the same chat," the separation has already failed.
- Does the Architect want the six protocol versions compared under this protocol, and if so who runs that thin pass so it is not unpaid synthesis landed back on the Architect?
- Should ONESTAMP remain outside this method until a later ratification, or is publication to GitHub itself the registry event?
- May a collaborator decline a topic on clinical-firewall grounds without that counting as non-participation?
- Which repository path, if any, is the pre-registration target for Tier B? A commit timestamp cannot be a gate if no repository has been named.

Concerns

- Cadence. A method that depends on Perplexity Computer will run when there is budget and disappear when there is not. The thin pass exists so the hole map does not depend on the executor.
- Capture by the review. If Iteration 2 treats the three pilots as premises, the laboratory will have done to itself what the review accused the Iteration 1 prompts of doing: prescribed the answer's shape before the test.
- Sensitivity. Endings, cessation, and loss of futurity sit near grief and suicidal thinking. Any reuse on that topic carries the clinical firewall. The protocol does not authorize human-subject exercises.

- Canon drift. GitHub publication is a record. It is not ratification. Every file produced under this protocol keeps the status line. A later pass may not delete a dissent to make the repository look settled.
- Self-entrenchment. This protocol is easy for the Architect to pause. A later version that makes pause harder than continuation has failed.

6. What the September 30 review actually did

This section is a method extraction, not a re-review of the five Iteration 1 laboratories. Iteration 1 documents were not re-read in full for this protocol. Claims about them are SOURCE from OAL-ECL-XR-PERPLEXITY-001, or ASSESSMENT about that review. [SOURCE; ASSESSMENT]

Move in the review	What it produced	Protocol status
Named corpus S1–S5, claimed stage versus delivered stage	Showed Copilot's Stage 3 label had no executed exercises, and that Gemini was absent.	Keep. Add a gap row that is not a review.
Seven questions per document, self-review first	Distinctive contribution, internal contradiction, assumed answer, treatment of Architect material, executability, carry or quarantine.	Keep. Self-review is mandatory if the preparer is in the corpus.
Three governing rules	Agreement is weak evidence. A finding requires a run. Real-world material is analogy, not a fact-check of the fiction.	Keep. Add: disagreement can also be a prompt artifact.
Case for and case against	Equal space before a position.	Keep, shortened. The case against may not be a courtesy paragraph.
Mission plus dissenting mission	Two live directives, one not adopted.	Keep the form. Do not keep the endings mission as the protocol mission.
Discrepancy register D-01 to D-24	Positions, why it matters, path. No averaging.	Keep the form. Paths are TEST, DECISION, HYGIENE only.
Three pilots, fixed seed, published code and JSON	First findings in the thread. Scope limited in the appendix.	Keep as the model for Tier B. Do not promote model results to topic premises.
Copy-ready Iteration 2 prompt	Gates, common core, quality failures.	Keep the idea. Shrink the required reading. Do not prescribe topic ontology inside the method prompt.

7. What to keep, and what not to copy

Keep

- Status line on every page header: proposed, not canon, not ONESTAMP-ratified.
- Preparer line, date, document ID, version, and an independence declaration that lists inputs actually read.
- Epistemic labels. A lesson written before a run is a prediction. It is not a finding.
- Carry-forward and quarantine, stated separately. A document can be wrong in its premise and still hold one instrument worth keeping.
- A comparative table with one row that can be zero. Executed exercises is that row until a run exists.
- An unrun-task list. The absence of a baseline, a control, a record-loss case, or a false-emergency case is a finding about the corpus, not a request for more prose.
- A preserved dissent. Deleting the alternative mission to clean the file is a failed review.

Do not copy

- Do not copy the 16-stage lifecycle, the 14 offices, or the Charter requirement into the protocol. Those were topic prescriptions. The review showed they propagated.
- Do not copy a single numeric score for a document or a laboratory. Split families stay split.

- Do not copy secondhand positions as if the missing preparer had spoken.
- Do not copy funding recommendations into a cross-topic template.
- Do not copy "this supports hypothesis H" from a one-author toy without the sentence "inside this model, under these parameters."
- Do not require every collaborator to read every prior document. Require them to list what they read and what they did not.

8. Two-tier operating model

The Architect said this exercise is worth doing and is not affordable every month. The protocol splits on that constraint instead of pretending the full pass is the only real pass.

	Tier A, thin pass	Tier B, full pass
When	Default. Once a topic has a corpus and no executor budget.	Architect-selected topics, or any month an executor is available.
Time target	One sitting. Shorter than the corpus it reviews.	Still shorter than a redesign of the topic.
Required	Corpus register, seven-question cards, discrepancy register, unrun-task list, one decision-changing test named but not necessarily run, copy-ready next prompt.	Tier A, plus at least one executed exercise with procedure, seed or script, raw results, scope limit, and flip condition.
Finding label	Not available, except for record checks such as "this file was absent."	Available for the executed exercise only.
Fails if	It recommends an institution, averages a conflict, or claims a result.	It reports a rate without the paired cost, or uses a detector on worlds the same author built to claim distinguishability.

Tier A is not a draft of Tier B. It is a complete review of a different object: the corpus as a set of claims and omissions. A thin pass that apologizes for not running code has misunderstood its job.

9. Roles, independence, and authorship separation

Role	May do	May not do
Preparer	Review the supplied corpus. Run Tier B only if able. Declare inputs.	Claim independence from a document it read. Characterize a missing preparer.
Architect of Record	Open or pause a tier. Lock or refuse a reading. Assign generator and detector pairs. Ratify, which this protocol cannot do.	Be treated as an unreviewed office inside the fiction by this method document.
Generator	Build worlds or cases. Seal ground truth. Hand over observations only.	Score its own worlds and call the score evidence of distinguishability.
Detector	Diagnose blind. Report confidence and what it could not separate.	Ask the generator for the label, then report the label as a finding.
Red team	Attack the diagnosis, the instrument, and the scope sentence.	Replace the diagnosis with a preferred conclusion.
Pass-through reader	Compare independent protocol versions. Register their discrepancies.	Merge them into one master because they overlap.

Independence declaration, required form: "I read [list]. I did not read [list]. I am the preparer of [own prior document, or none]. Agreement with a document I read is not independent evidence."

Authorship separation applies only when the claim is that two states can be told apart. It is assigned by the Architect. A solo preparer may still run a circularity check, as P3 did, provided the finding is labeled as a check that the same author cannot prove the distinction.

10. Procedure

Step 1. Bound the topic

Name the topic, the Architect's question in the Architect's words, and what the review is forbidden to decide. For End Cycle, the forbidden decision was the meaning of the field scene. For another topic, write the equivalent sentence before reading the corpus.

Step 2. Register the corpus

One row per document. Columns: identifier, preparer, claimed stage, delivered stage, approximate size, executed exercises, gap note. A file discussed by others but not supplied is a gap row, not a document row.

Step 3. Read against seven questions

1. What does this document contribute that no other document in the corpus contributes?
2. What does it get right, on its own terms?
3. Where is it internally inconsistent?
4. Where does it assume what it claims to test?
5. How does it treat the Architect's material, including qualifiers?
6. Can it be executed as written, by whom, and with what instrument?
7. What should the next iteration keep, change, or quarantine?

If the preparer has a document in the corpus, that card is written first, and at the same severity.

Step 4. Write the case for and the case against

Each gets equal space. The case against is a list of ways the work can be a solution looking for a problem, a self-sealing fiction, a prompt artifact, a procedure that becomes the harm, a postponement, or a sensitivity failure. The review then states a position and says what would defeat it.

Step 5. Open the discrepancy register

A discrepancy is a conflict that changes a later test, a label, or an Architect decision. Overlap is not a discrepancy. Each row needs positions, why it matters, and one path. TEST means run something. DECISION means only the Architect can close it. HYGIENE means fix an identifier, a crosswalk, or an attribution. No fourth path called CONSENSUS.

Step 6. List unrun tasks

Every exercise, control, baseline, or falsifier named in the corpus and not executed goes on the list, with the document that named it. The list is an output, not a shame appendix.

Step 7. Tier B only: run the smallest test that can fail

Pre-register the hypothesis, the falsifier, the parameters, and the identity criterion before the run if a repository exists. If no repository exists, write them in the document above the results and say that pre-registration was unavailable. Report where the conclusion flips. Publish the script and the raw output beside the review.

Step 8. Commission the next iteration

The commission is copy-ready, names the tier, lists required inputs, repeats the stop rule, and is shorter than the review. It does not prescribe the answer's shape.

11. Per-document review card

Use this card once per document. Do not blend two preparers into one card to save space.

Field	Instruction
Document	ID, version, date, preparer. If no ID, say so. Absence of an ID is a hygiene finding.
What it is	Stage claimed, stage delivered, and the job it actually did.
Distinctive contribution	One thing the corpus would lack if this document were removed.
Adversarial findings	Numbered. Internal contradiction, assumed answer, missing control, non-executable instruction, scope slide.
Treatment of Architect material	Quoted with qualifiers, paraphrased into a premise, or absent.
Executed?	Count. Zero is an acceptable entry.
Carry forward	Instruments, scenarios, or rules worth testing. Not worth adopting as premises.
Quarantine	Premises, scores, offices, or lessons-before-runs that must not propagate.

12. Discrepancy register

Identifier pattern: D-01 upward, local to the review. Do not reuse End Cycle D-numbers for a different topic. A later review of the same topic may continue the sequence only if it cites the prior register.

Column	Rule
ID	D-01, D-02, in order of decision impact, not in order of discovery.
Discrepancy	One conflict. If two conflicts share a sentence, split them.
Positions	Name the preparer and the claim. "The corpus agrees" is not a position.
Why it matters	Name the later test, label, or decision that changes if the conflict is ignored.
Path	TEST, DECISION, or HYGIENE. One path. A note may say what the test would be.
Status	OPEN, until the Architect closes a DECISION or a run closes a TEST.

Forbidden resolutions: majority vote among models; "the longest document"; "the latest document"; a new composite definition that none of the preparers wrote; deletion of a minority position from the register.

13. Execution rules: what counts as a finding

Claim type	Counts as FINDING?	Minimum record
A document contains no executed exercise	Yes, as a corpus finding	The document was read, and no run, table, or result file is in it.
A rule reduces a named harm	Only after a run	Procedure, parameters, seed or script, raw output, paired cost.
Two states are distinguishable	Only under separated authorship	Generator, detector, and red team are different. Solo circularity checks must say they are circular.
A real institution behaves like the fiction	No. EVIDENCE, analogy only.	Citation. No edit to the fictional source.
Collaborators agree	No	Record the agreement as weak evidence and list the shared inputs.
A file reproduces	Yes, about the file	Same script, same seed, difference count. Not a validation of the model.

Known limits belong in the same section as the result, not only in an appendix. Required sentence: "These are findings about the model under author-chosen parameters. They are not findings about [the fiction / the world]. The conclusion flips if [condition]."

14. Score families that must not be collapsed

A review may score a corpus. It may not collapse the scores into one number, one grade, or one ranking of collaborators. Minimum families:

Family	Question	Legal values
A. Specification	Does the document say what would be measured?	Named / partial / absent
B. Execution	Was anything run?	Count, including zero
C. Control	Is there a no-intervention or no-event branch?	Yes / no / not applicable
D. Record integrity	Can a later reader reproduce or locate the result?	Reproduced / present but not run / absent
E. Architect fidelity	Are the Architect's qualifiers intact?	Intact / slid / absent
F. Pause cost	Did the design make itself harder to stop?	No / unclear / yes

Topic reviews may add families. End Cycle's completion, persistence, legitimacy, and delay families stay in that topic. They are not imported here as universal scores.

15. Architect pass-through

After independent protocol versions exist, the pass-through is a separate document. It is not a revision of those versions. They remain on the record unedited.

- Register each version: ID, preparer, date, tier, page or word count, whether it executed anything.
- Build one discrepancy register across the protocol versions. Expected conflicts include tier definitions, whether tabletop counts, required reading, and who may synthesize.
- Do not average. Mark each conflict TEST, DECISION, or HYGIENE.
- The Architect may adopt one version, adopt none, or lock a thin common core and leave the rest as lanes. Locking a core is a DECISION, recorded in an Architect decision log.
- A pass-through that recommends a standing review office has left the method.

Decision log fields: date, decision ID, the discrepancy it closes, the choice, what remains open, and the pause condition. A decision without a pause condition is incomplete.

16. Identifiers, files, and GitHub hygiene

Object	Pattern	Example
Protocol rendering	OAL-AXR-PROTO-{AI}-{nnn}	OAL-AXR-PROTO-GROK-001
Topic review	OAL-AXR-{TOPIC}-{AI}-{nnn}	OAL-AXR-ECL-GROK-001
Pass-through	OAL-AXR-PT-{TOPIC}-{nnn}	OAL-AXR-PT-PROTO-001
Hypothesis	{LAB}-H-{letter}	GROK-H-A
Discrepancy	D-{nn} local to the review	D-01
File stem	{ID}_v{X.Y}.docx and .pdf	OAL-AXR-PROTO-GROK-001_v0.1.pdf

Version v0.1 is a proposal. A preparer does not jump to v1.0 because the document is long. v1.0 is available only after the Architect marks a revision accepted for use. Acceptance is still not canon and still not ONESTAMP.

Every PDF and DOCX pair carries the same ID, version, date, and preparer. Code and raw results, if any, use the same stem with a suffix: `_pilot.py`, `_results.json`. A review that cites results it does not attach has a record-integrity failure.

17. Quality gates

A review produced under this protocol fails its own check if it does any of the following.

- Presents a predicted lesson as a finding.
- Uses one number to rank documents or collaborators.
- Averages a discrepancy, or deletes a dissent to tighten the file.
- Reviews a preparer whose document was not supplied.
- Reinterprets Architect material and labels the interpretation SOURCE.
- Recommends an office, charter, or department without an executed exercise and a dissenting model.
- Claims two states are distinguishable using a detector the same author built for worlds that author built.
- Reports a benefit without the paired cost.
- Omits the scope sentence on a pilot.
- Makes the protocol harder for the Architect to pause than it was before.
- Drops the status line, the clinical firewall on a sensitive topic, or the preparer line.

18. Copy-ready commission prompt

The block below is the prompt a later preparer can be given. Bracketed fields are filled by the Architect. It is a PROPOSAL. It is not canon.

COMMISSION

STATUS: PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. Issued under OAL-AXR-PROTO-GROK-001 v0.1 unless the Architect revises it. This prompt does not supersede any collaborator's independent rendering.

Role

You are an independent collaborator writing one adversarial cross-review for Omniversal Architect Labs, for Nathan Lewis II, Architect of Record. You succeed if a later reader can see what was run, what was not run, and where the corpus conflicts, without being told which collaborator won.

Fill before sending

- Topic: [TOPIC].
- Architect's question, verbatim: [QUESTION].
- Corpus supplied: [LIST]. Corpus known but not supplied: [GAPS].
- Tier: [A or B]. If B, say whether code, tabletop, or both is acceptable, and name the repository if pre-registration is required.
- Forbidden decision: [WHAT THIS REVIEW MUST NOT CLOSE].

Rules

- List every input you read and every input you did not read. Do not claim independence from a document you read.
- Label claims SOURCE, FINDING, EVIDENCE, ASSESSMENT, PROPOSAL, or OPEN. Only an executed run produces FINDING.
- Do not average conflicts. Use TEST, DECISION, or HYGIENE.
- Do not fact-check fictional source material. Label readings as PROPOSAL.
- Keep the clinical firewall if the topic sits near grief, death, or suicidal thinking. In the US, anyone in crisis can call or text 988.
- If you are in the corpus, review your own document first, at the same severity.

- Do not design the governing institution inside the review.
- Nothing in your document is canon until the Architect says so.

Deliver

- One DOCX and one PDF. Title, date, version, preparer, status block, mission or directive, executive summary, overview, contents.
- ID: OAL-AXR-[TOPIC]-[AI]-001, version v0.1, unless the Architect assigns another number.
- Corpus register, seven-question cards, discrepancy register, unrun-task list, carry and quarantine.
- If Tier B: script or tabletop record, raw results, scope sentence, flip condition.
- A next-iteration prompt shorter than your review.

Closing line

Run the smallest test that can prove you wrong before building the largest structure that could make you look right. Report what the instruments could not tell apart. Record what was not run. Nothing in this pass is canon until the Architect says so.

19. Appendix A. Blank worksheets

A1. Corpus register

#	Document	Preparer	Claimed	Delivered	Run?
S1	[title, ID, version]	[name]	[stage]	[stage]	[count]
S2					
Gap	[named but not supplied]	[discussed by]	n/a	not supplied	n/a

A2. Discrepancy row

ID	Discrepancy	Positions	Why it matters	Path
D-01	[one conflict]	[preparer: claim]	[what changes]	TEST / DECISION / HYGIENE

A3. Unrun-task list

Task	Named in	Why it blocks comparison	Tier to close
[baseline, control, falsifier, exercise]	[document]	[incomparable results / false positive / circular test]	A name it / B run it

A4. Architect decision log

Date	Closes	Choice	Still open	Pause condition
[date]	[D-id]	[choice]	[what was not closed]	[how the Architect stops this]

20. Appendix B. Reproduction note on the End Cycle pilots

OAL-ECL-XR-PERPLEXITY-001 v0.2 states that pilot.py uses NumPy and master seed 20260930, and that results.json holds the raw outputs. On October 1, 2026, this preparer recomputed P1, P2, and P3 from the supplied script logic under that seed, without the plotting dependency. [FINDING]

- Difference count against the supplied results.json: 0, at a tolerance of 0.002 on reported rounded values.
- P1 immediate non-migrant loss reproduced at mean 1.0. P2 split authority at urgent share 0.10 and delay 8 reproduced at wrongful rate 0.011 and urgent harm 0.546. P3 external-only accuracy reproduced at 0.50 across the published separations.
- Scope, unchanged from the source appendix: no spatial structure, no learning, no strategic actors, parameters chosen by the original author. Reproduction validates the file. It does not validate the field scene, the Department question, or any real-world analogy in that review.

This is the kind of check Tier A can sometimes perform when code and results are already attached. It is not a substitute for a new exercise. [ASSESSMENT]

21. Appendix C. Open questions reserved for the Architect

1. Does tabletop count as execution for Tier B, or is code required?
2. Who stores sealed ground truth for generator and detector pairs?
3. Is the pass-through over protocol versions a Tier A task the Architect will commission, or a task the Architect will do?
4. Which GitHub path is the registry for these files, and does publication change status? This protocol says no.
5. Should the dissenting directive, "not yet," become the default for any topic the Architect marks as irreversible?
6. What is the stop rule between this protocol and unfinished First Memory work, so the review method does not become a way to postpone it?

End of OAL-AXR-PROTO-GROK-001 v0.1. Proposed. Not canon. Not ONESTAMP-ratified. Pending Architect review.