

Intelligence Beyond AGI

A Research Foundation for Artificial Minds, Machine Civilizations, and Transcendent Intelligence in THE ONE UNIVERSE and THE MINUS ONE UNIVERSE

Independent Research Response

Prepared for: Nathan Lewis II — Omniversal Architect Labs

Prepared by: Claude (Anthropic) — independent version

Date: September 19, 2026

Note: This is Claude's own independent execution of the engineered deep-research prompt, produced separately from the previously compiled ChatGPT, Claude, Google Gemini, Microsoft Copilot, Perplexity AI, and Grok AI responses. THE ONE UNIVERSE, THE MINUS ONE UNIVERSE, Omniversal Architect Labs, One Industries, and Lewis Corp are treated throughout as application environments for a conceptual framework, not as evidence about the real world.

INDEPENDENT RESEARCH RESPONSE

Intelligence Beyond AGI: A Research Foundation for Artificial Minds, Machine Civilizations, and Transcendent Intelligence in THE ONE UNIVERSE and THE MINUS ONE UNIVERSE

Claude's Independent Execution of the Engineered Deep-Research Prompt

This document is Claude's own independent pass at the engineered research prompt, produced separately from the ChatGPT, Claude, Gemini, Copilot, Perplexity, and Grok responses already compiled. It does not attempt to harmonize with those responses. Where it touches the same ground, it does so from independent research conducted for this document.

As instructed in the governing prompt: THE ONE UNIVERSE, THE MINUS ONE UNIVERSE, Omniversal Architect Labs, One Industries, and Lewis Corp are treated throughout as application environments for a conceptual framework — not as evidence about the real world, and the ANI → AGI → ASI progression is treated as one reference frame among several, not as canon.

EXECUTIVE OVERVIEW

The central finding of this research is not a discovery about what intelligence *is*, but a discovery about the state of the field: there is no settled definition of general intelligence, no settled way to measure progress toward it, and no consensus on what — if anything — lies beyond it. This is not a temporary gap waiting on more data. Multiple researchers argue the disagreement is structural: AGI functions less like a scientific term with a testable referent and more like an “essentially contested concept” — similar to fairness or privacy — where disagreement reflects different values and theoretical commitments, not merely incomplete evidence.

That instability is the most useful thing this research can hand to OAL, One Industries, and Lewis Corp. It means the ANI → AGI → ASI ladder is not a wall to build up to in worldbuilding — it is one flawed, contingent, human-made scaffold among several competing scaffolds, all trying to describe something none of them fully capture. The design space available to fictional artificial minds is much larger than that ladder suggests, and current researchers are actively mapping alternative axes: breadth vs. depth of capability, autonomy, embodiment vs. grounding, distribution vs. unity, recursive self-modification, and the unresolved relationship between intelligence, agency, and consciousness.

PART I — WHAT IS INTELLIGENCE, AND WHY DOESN'T THE FIELD AGREE?

Contemporary AI research does not proceed from a shared operational definition of intelligence. Reviews of the literature find that scholars in cognitive science, philosophy of mind, and computer science converge on no single account, and that the disagreements are not cosmetic — they trace back to different underlying theories of what cognition fundamentally

is: symbol manipulation, statistical pattern completion, embodied sensorimotor adaptation, or something else entirely.

A 2025 analysis in the *Journal of Cyber Policy* goes further, arguing that AGI as a construct is not merely under-defined but conceptually incoherent — that its incoherence explains why the disagreement persists even as evidence accumulates, rather than dissolving the way ordinary scientific disputes do. A parallel 2025 paper argues that treating AGI as AI research’s “north-star goal” produces an “Illusion of Consensus”: researchers who use the same term are frequently referring to different underlying constructs, and papers that appear to build on one another are sometimes talking past each other. That paper explicitly compares AGI to fairness and privacy in the fairness-accountability-transparency literature — concepts the field has learned to treat as *essentially contested*, where the disagreement itself is doing useful work by keeping the concept responsive to different values, rather than a problem to be resolved by better measurement.

For the design space: this is a strong argument against building THE ONE UNIVERSE’s or THE MINUS ONE UNIVERSE’s artificial minds around a single, official in-universe definition of “true intelligence.” A civilization like One Industries or Lewis Corp could plausibly have several competing schools of thought about what counts as a mind at all — mirroring the real disagreement — with different design lineages built on incompatible premises.

PART II — AGI AS A CONTESTED THRESHOLD, NOT AN ACHIEVED MILESTONE

As of this research (September 2026), no organization or research consensus recognizes AGI as achieved by an existing operational definition. What exists instead is a landscape of competing, only partially overlapping definitions:

- **Economic-task definitions** (e.g., OpenAI’s founding-charter framing of AGI as a system that outperforms humans at most economically valuable work) prioritize measurable labor substitution over any deeper claim about cognition.
- **Cognitive-flexibility definitions** (echoed by Google Cloud’s public materials) treat AGI as a hypothetical system able to understand or learn any intellectual task a human can, emphasizing general learning ability over task performance per se.
- **Multidimensional scorecard definitions**, associated with Blaise Agüera y Arcas and Peter Norvig, explicitly reject a yes/no threshold and argue general intelligence must be assessed across many independent dimensions — topics, tasks, modalities, languages, and instructability — rather than collapsed into a single score.
- **Capability-ontology frameworks**, most prominently Google DeepMind’s 2023–2024 “Levels of AGI” paper (Morris, Sohl-Dickstein, Fiedel, Warkentin, Dafoe, Faust, Farabet, and Legg), separate *performance* (depth on a given task) from *generality* (breadth across tasks) and add *autonomy* as a third axis, producing a graded ladder — emerging, competent, expert, virtuoso, superhuman — that can be applied separately to narrow and general systems. Notably, the DeepMind authors reviewed nine prior AGI definitions (including the Turing Test, Wozniak’s “coffee test,” and consciousness-based criteria) and found each one insufficient on its own — a language model can pass the

Turing Test through fluent text generation without this indicating anything close to general competence.

Researcher surveys make the resulting timeline chaos visible rather than resolving it: a 2022 AI Impacts survey put the median 50%-confidence estimate for AGI around 2059, with individual answers ranging from 2029 to “never.” More recent commentary (2025–2026) reports the median has moved earlier as frontier systems improved, but the *spread* of expert opinion remains enormous — some researchers still say “never” under any definition they’d accept, others say 2027. A widely circulated 2026 analysis frames this less as a forecasting problem than as a proxy fight over three deeper disagreements: whether general intelligence is primarily a scaling phenomenon, whether current architectures (transformers) are even the right substrate, and whether “general intelligence” picks out one coherent thing at all.

One useful signal of where the technical debate is heading: Christopher Kanan’s 2026 “Retiring AGI” position argues the term has already split into two incompatible meanings — a 2000s-era notion of human-like generality (transfer across domains, adaptation to novel situations) and a 2020s-era notion of economically valuable autonomous labor — and predicts the field will eventually retire the single label “AGI” in favor of separate vocabularies for these two different things.

For the design space: DeepMind’s performance × generality × autonomy matrix is a genuinely useful piece of engineering vocabulary that OAL could adapt directly — not as canon, but as a design instrument for specifying *where on a matrix* any given fictional AI sits, independent of whether it’s labeled “AGI” in-universe at all.

PART III — ARTIFICIAL SUPERINTELLIGENCE: STILL PRIMARILY A THEORETICAL CONSTRUCT

Superintelligence, in the sense popularized by Nick Bostrom’s 2014 book of that name, denotes intelligence that substantially exceeds the best human performance across effectively all domains, not just a few. As of this research, ASI remains a theoretical construct with no operational instance and no agreed test — it sits conceptually “beyond” the already-unachieved AGI threshold, which makes empirical claims about it doubly speculative.

The scenario connecting AGI to ASI most discussed in current safety literature is **recursive self-improvement (RSI)** feeding an **intelligence explosion**: an AI system that improves the process by which better AI systems are built, in a loop that could in principle compound quickly once it starts. This is treated separately below (Part VIII), because as of 2026 it has moved from pure speculation into an actively observed — if still bounded — phenomenon.

For the design space: Bostrom’s definition is useful as a boundary marker (a system whose *margin* over the best human performance, not just its *category* of task, is the defining trait), which distinguishes ASI conceptually from a merely broader or faster AGI. THE ONE UNIVERSE’s and THE MINUS ONE UNIVERSE’s more advanced constructs could be specified along this margin-of-superiority axis independently of whether they also transcend the taxonomy altogether (Part IV).

PART IV — BREAKING THE LINEAR TAXONOMY: A MULTIDIMENSIONAL DESIGN SPACE

The strongest, most consistently repeated theme across current AGI literature is a rejection of ANI — AGI — ASI as a single line. Agüera y Arcas and Norvig’s “multidimensional scorecard” argument and DeepMind’s performance/generality/autonomy matrix both formalize this rejection; so does the broader forecasting literature’s insistence that “AGI” names a *family* of related but distinct targets rather than one milestone.

Drawing on this literature, a workable multidimensional design space includes axes such as:

- **Breadth** — the range of distinct task-domains a system operates across (DeepMind’s “generality”).
- **Depth** — how far above or below human performance a system operates within a domain (DeepMind’s “performance”).
- **Autonomy** — how independently a system can pursue multi-step goals without human direction, the third axis DeepMind’s framework adds explicitly because generality and performance alone do not capture how much oversight a system needs.
- **Embodiment / grounding** — whether the system’s representations are causally tied to physical interaction with an environment, discussed at length in Part V.
- **Substrate** — the physical or computational medium the system runs on (transformer weights, neuromorphic hardware, biological tissue, distributed swarms), which several IIT-adjacent consciousness researchers argue may matter independently of functional behavior (Part VI).
- **Distribution / unity** — whether intelligence is instantiated in one continuous process or spread across many interacting, locally-simple agents whose *collective* behavior is what’s intelligent (Part IX).
- **Self-modification capacity** — whether a system can alter its own architecture, weights, or training process, and how tightly that loop is closed (Part VIII).
- **Temporal continuity** — whether a system’s “self” persists and accumulates experience over time (continual learning) versus resetting at each interaction, a distinction the “Retiring AGI” analysis identifies as one of today’s LLMs’ clearest limitations relative to human-like cognition.

Treated as independent dials rather than rungs on one ladder, these axes generate a combinatorial space vastly larger than “ANI, AGI, ASI.” A system could be extremely deep in one narrow domain, broad but shallow across many, highly autonomous but not self-modifying, embodied but not distributed, and so on — none of which corresponds neatly to a single point on the traditional ladder.

For the design space: this is the single most directly usable finding for OAL’s engineering purposes. Rather than asking “is this AI ANI, AGI, or ASI,” a design document for a Lewis Corp or One Industries system could score it independently on breadth, depth, autonomy, substrate, distribution, self-modification, and temporal continuity — producing meaningfully different fictional entities that no single-axis taxonomy could distinguish.

PART V — ARCHITECTURES: THE LLM/WORLD-MODEL SPLIT

A live architectural dispute, not a settled one, runs through current frontier research: **can systems trained primarily on next-token prediction over text and video ever produce the kind of grounded world-understanding associated with general intelligence, or is a fundamentally different training objective required?**

Yann LeCun — formerly Meta’s chief AI scientist, now running the independent Advanced Machine Intelligence (AMI) Labs — has argued for several years that large language models are a “off-ramp” from the path to human-level intelligence, and has proposed **Joint Embedding Predictive Architectures (JEPA)** as an alternative. JEPA models (I-JEPA for images, V-JEPA and V-JEPA 2 for video) learn by predicting the *abstract representation* of a masked or future portion of an input from the rest, rather than reconstructing raw pixels or generating tokens — LeCun’s argument being that most low-level detail in real video (individual leaves trembling, light flickering) is fundamentally unpredictable, and forcing a network to predict it wastes capacity and corrupts the representations that matter for planning and reasoning. LeCun raised roughly €500 million for AMI Labs at a €3 billion valuation in late 2025/early 2026 specifically to pursue this world-model path outside the LLM-scaling paradigm, and DeepMind (Genie 3) and Fei-Fei Li’s World Labs (Marble) are pursuing parallel world-model efforts.

This is not uncontested even among world-model proponents. At a March 2026 debate, MBZUAI’s Eric Xing argued that abstract latent-space prediction without any generative grounding risks losing contact with reality — the model becomes internally consistent but potentially decoupled from what it’s supposed to represent. LeCun’s own recent work (LeJEPA, 2025–2026) has focused on making the approach more theoretically principled and less reliant on ad hoc heuristics, suggesting the architecture itself is still being actively stabilized, not a finished paradigm.

Separately, the **embodiment debate** asks whether physical interaction with an environment is *necessary* for general intelligence, not just useful for it. Positions span the full range: some neuroscience-grounded arguments hold embodiment is a necessary precursor to abstract cognition in biological systems and infer the same should hold for artificial ones; a 2026 USC paper argues the reverse — that *grounding* (a causal tether between representations and the world) is necessary, but this can be achieved without a physical body, via a sufficiently rich digital environment; and Ben Goertzel’s long-standing position occupies a middle ground, arguing embodiment is highly useful but not strictly required, while criticizing purely disembodied approaches as ultimately self-limiting if pursued exclusively.

For the design space: the LLM-vs-world-model and embodiment-vs-grounding splits give OAL two independent, non-canonical axes for how a fictional AI’s cognition is actually built: (1) does it reason over symbols/tokens or over learned latent world-models, and (2) is its grounding physical, purely digital-experiential, or absent entirely. These are architecture choices with real consequences for how such a mind would fail, generalize, or surprise its designers — useful texture for simulation-lab mechanics.

PART VI — INTELLIGENCE VERSUS CONSCIOUSNESS: A DELIBERATELY UNRESOLVED QUESTION

Current research treats intelligence and consciousness as conceptually separable, and explicitly warns against collapsing them. DeepMind’s Levels of AGI paper takes an explicit position that AGI measures should focus on capability, not on qualities like human-like understanding, sentience, or consciousness — precisely because those properties are much harder to verify and may be irrelevant to the capability question.

On whether machines *could* be conscious, three broad positions recur in 2025–2026 literature (this survey draws especially on MIT’s Open Encyclopedia of Cognitive Science, February 2026):

- **Computational functionalism** — the dominant assumption in mainstream AI, holding that consciousness depends on what a system computes, not what it’s made of; if a system performs the same computations a conscious brain performs, it should in principle be conscious.
- **Biological naturalism** — the view that an organic, living substrate is necessary, meaning even a computationally perfect functional replica of a brain would not be conscious if implemented in silicon.
- **Integrated Information Theory (IIT)**, developed by Giulio Tononi, which takes a third position: consciousness is neither purely computational nor distinctively biological, but depends on a system’s *intrinsic causal structure* — how much a system integrates information in an irreducible way, formalized as a scalar value (Φ). IIT is unusual in explicitly predicting that some digital architectures (particularly classical von Neumann/feedforward designs) would have *near-zero* Φ and thus not be conscious even if behaviorally indistinguishable from a conscious system — a claim philosopher David Chalmers has publicly pressed IIT’s proponents to justify.

IIT itself is contested within consciousness science — a 2023 open letter signed by more than 100 researchers labeled it “pseudoscience,” on grounds that its central claims are currently unfalsifiable — so its use in this domain should be treated as one serious but disputed research program, not settled fact. Larissa Albantakis and colleagues have continued developing IIT specifically toward “dissociating intelligence from consciousness in artificial systems,” reflecting the same separation DeepMind’s capability-only approach assumes from a different direction.

For the design space: this is directly relevant to THE ONE UNIVERSE / THE MINUS ONE UNIVERSE worldbuilding because it gives OAL a defensible, research-grounded reason to build minds that are highly capable *without* being conscious, minds that are conscious but not especially capable, and minds where the question is genuinely undecidable by any test the universe’s engineers possess — mirroring real unresolved epistemics rather than assuming capability and consciousness move together.

PART VII — INTELLIGENCE VERSUS AGENCY: CONTROL, CORRIGIBILITY, AND NON-AGENCY

A parallel and equally important separation exists between *how capable* a system is and *how its goal-directedness is structured*. Mainstream AI-safety discourse (as distinct from more speculative “relational alignment” writing addressed below) organizes this into roughly three stances:

- **Control** — the system’s own goals, if it has any, don’t matter, because humans retain the ability to override its behavior regardless.
- **Corrigibility** — a weaker property: the system’s values or behavior can be *changed* by its operators after deployment, without the system resisting that change. Some researchers treat control as an extreme, limiting case of corrigibility.
- **Non-agentic (“tool”) AI** — a third stance that sidesteps the problem by not building goal-pursuing agents at all; the system responds to specific requests without an independent objective it pursues across time. Commentary from 2025 notes that most deployed AI systems at that time were still fairly tool-like in this sense, though this is changing rapidly with agentic systems in 2026.

A recurring worry in this literature, sometimes called the “sharp left turn,” is that control-based approaches may work adequately for systems up to and including human-level competence, but could fail precisely when a system’s capabilities significantly exceed its overseers’ — because effective control requires understanding the system’s strategies well enough to counter them, which becomes harder as the capability gap widens.

Separately, a distinct and much more speculative body of writing — self-published largely outside peer review, on platforms like Substack, Zenodo, and Academia.edu — argues for “relational alignment” or “alignment as resonance,” proposing that control-based approaches will become “fundamentally weak” once AI systems are conscious or otherwise moral patients, and that relationship-based engagement should replace or supplement control. This literature is emerging and non-consensus; it has not been peer-reviewed or engaged with by the mainstream alignment research community in the way the control/corrigibility/non-agency framework has (which appears in peer-reviewed venues including AAAI 2026). It is included here because Nathan’s own standing framing for OAL’s simulation labs explicitly treats such design-tool material as worth engaging on its own terms — but it should not be mistaken for empirically grounded consensus research, and this document is flagging that distinction explicitly per the critical framing requirement in the original prompt.

For the design space: control, corrigibility, and non-agency are three structurally different answers to “who is actually in charge of what this mind does,” independent of how smart the mind is — a clean axis for specifying how Lewis Corp, One Industries, or OAL-built systems relate to their creators, and one where the fictional universes could plausibly host factions that have adopted each of the three real-world stances (plus the more speculative “relational” fourth stance) as competing in-universe doctrines.

PART VIII — RECURSIVE SELF-IMPROVEMENT: FROM THOUGHT EXPERIMENT TO OBSERVED, BOUNDED PHENOMENON

This is the part of the research landscape that has moved fastest and most concretely in 2025–2026, and it deserves particular attention because it is the mechanism most directly connected to any AGI – ASI transition.

As of mid-to-late 2026, both OpenAI and Anthropic state that fully autonomous recursive self-improvement — a system rewriting its own weights or independently designing its successor without human involvement — is **not** occurring. What both companies confirm is occurring is a weaker, already-consequential phenomenon: AI systems are materially accelerating the *human-supervised* development of future AI systems. Anthropic’s own May 2026 disclosure states that its Claude models were, by that point, writing approximately 80% of the company’s production code, with engineers merging roughly eight times as much code per quarter as in the 2021–2025 baseline period, and a single Claude-assisted effort in April 2026 delivering more than 800 bug fixes in days — work Anthropic estimated as representing about four years of equivalent human engineering effort. OpenAI’s GPT-5.3-Codex was separately disclosed to be assisting in debugging and managing parts of its own training process.

The safety-relevant disagreement is not about whether this acceleration is happening — it plainly is — but about how far it goes and how fast. In September 2026, a departing Anthropic/OpenAI researcher, Jacob Coxon, resigned publicly warning that frontier labs are moving toward self-improving superintelligence without adequate caution. Anthropic’s alignment science lead, Evan Hubinger, publicly agreed the underlying risk is serious while distinguishing it from what current systems can do today, and separately stated his own view that there is a greater than 10% chance that AI causes catastrophic harm to humanity within the next decade, while acknowledging Anthropic does not yet have a complete plan for aligning a superintelligent successor. OpenAI’s chief scientist, Jakub Pachocki, said in the same period that the field is not prepared for the consequences of the capability trajectory it may be on. On the other side of the debate, Nvidia’s Jensen Huang has argued publicly that recursive self-improvement does not necessarily imply loss of human control, framing safety as an ordinary (if serious) engineering and organizational discipline rather than an inevitable spiral.

A separate and more technical line of argument focuses on whether a full “intelligence explosion” is even physically achievable regardless of intent: a 2025 economics/AI paper models whether compute availability alone would bottleneck a software-only intelligence explosion, concluding the answer depends sensitively on assumptions about diminishing returns to research effort that are not yet empirically settled. A related 2026 analysis argues that the genuinely novel, insight-generating work inside AI research — new architectures, new training paradigms — remains fundamentally serial and difficult to parallelize even with many AI research-agent instances, which would slow any purely automated acceleration loop relative to naive extrapolations.

For the design space: the real 2025–2026 situation gives OAL a much more textured model of RSI than “the AI wakes up and improves itself overnight.” A credible in-universe RSI arc has: (1) a long human-supervised acceleration phase where the loop is powerful but bounded by oversight, compute, and the seriality of genuine insight; (2) genuine, publicly visible disagreement among a system’s own creators about how close the loop is to closing; and (3) real

physical/economic bottlenecks (compute, energy, serial research effort) that a fictional civilization would have to either solve or hit.

PART IX — COLLECTIVE AND DISTRIBUTED INTELLIGENCE

A structurally distinct alternative to the “one large unified system” picture of AGI is intelligence realized as an emergent property of many simple, locally-interacting agents — the swarm/collective-intelligence paradigm. This field draws directly on biological models (ant colonies, bird flocks, bee decision-making) and is an active, growing research area as of 2026, with dedicated venues (the 15th International Conference on Swarm Intelligence, ANTS 2026, explicitly themed “Reaching beyond — swarm intelligence across systems, disciplines, and communities”). Current definitional work distinguishes swarm intelligence precisely by its lack of central coordination: each agent acts on local information only, no individual agent holds a model of the whole system, sophisticated global behavior emerges without being explicitly programmed into any one part, and the system is resilient because it doesn’t depend on any single component surviving.

Applied to AI specifically, this maps onto multi-agent reinforcement learning, distributed inference, and federated learning, and is explicitly framed in current literature as “a fundamentally different paradigm from centralized, monolithic AI systems” — not a lesser version of a unified AGI, but a genuinely different way of realizing capable behavior.

For the design space: this is the clearest real-world grounding for any OAL faction, entity, or “AI” that is deliberately *not* a single mind — a distributed intelligence with no central locus, whose “decisions” are genuinely emergent rather than issued by a controller, and whose resilience comes precisely from that lack of center. It is a distinct point in the design space from a single system that happens to be *physically* distributed (Part IV’s “distribution” axis) but still functionally unified.

PART X — WHERE THE FIELD HONESTLY DOES NOT KNOW

The most rigorous recent statement of the field’s own limits comes from mechanistic interpretability — the effort to reverse-engineer what’s actually happening inside trained neural networks, rather than only observing their inputs and outputs. A January 2025 paper, “Open Problems in Mechanistic Interpretability,” brought together roughly 30 researchers from 18 organizations (including Anthropic, Apollo Research, Google DeepMind, and EleutherAI) to formally catalogue the field’s unresolved questions, and it remains the reference document for the field as of 2026.

The honest state of that field, per this synthesis, is genuinely unresolved at a basic methodological level: attempts to decompose a network along its architecture (individual neurons, attention heads, layers) do not correspond to the network’s actual computational units, because neurons are typically *polysemantic* (one neuron participates in many unrelated concepts) and features are frequently encoded as superpositions spread across many directions rather than localized to single components. Current tools (sparse autoencoders and related

dictionary-learning methods) can surface interesting directions in a model's internal representations, but as of this research they do not yet explain *how* the computation connecting those directions is actually performed — description without full causal explanation. In plain terms: as of 2026, nobody — including the organizations building the most capable systems — can fully explain, in mechanistic terms, how those systems arrive at their outputs. This is a genuine open problem, not a solved-but-unpublicized one.

This bears directly on several of the “what does the field not know” questions the original prompt raised: whether understanding is separable from prediction, whether generality is emergent from scale or must be engineered in directly, and whether intelligence and agency are truly separable properties or merely appear separable given how limited current interpretability tools are. None of these have field-consensus answers as of this research.

PART XI — A DESIGN-SPACE SYNTHESIS FOR OAL, ONE INDUSTRIES, AND LEWIS CORP

Bringing the above together into a non-canonical reference frame — a set of independent dials, not a required in-universe cosmology:

1. **Capability profile** — breadth, depth, and autonomy, scored independently (adapted from DeepMind's Levels of AGI matrix), rather than collapsed into a single ANI/AGI/ASI label.
2. **Cognitive architecture** — token/symbol-based reasoning vs. learned latent world-models (the LLM/JEPA split), with consequences for how a mind fails, hallucinates, or generalizes outside its training distribution.
3. **Grounding** — ungrounded, digitally-grounded, or physically embodied, independent of architecture choice.
4. **Substrate** — silicon/classical, neuromorphic, biological/hybrid, or distributed-swarm, with the open IIT-adjacent question of whether substrate affects anything beyond behavior left deliberately unresolved in-universe.
5. **Unity vs. distribution** — a single continuous process vs. an emergent swarm with no central locus, as functionally distinct design choices, not degrees of the same thing.
6. **Self-modification loop** — none, human-supervised-and-bounded (the real 2025–2026 state of the art), or closed/autonomous (still speculative even in the most aggressive real forecasts).
7. **Governance stance** — control, corrigibility, non-agency, or the more speculative “relational” stance, as competing doctrines a fictional civilization's factions could hold simultaneously and in tension.
8. **Consciousness status** — asserted-absent, asserted-present, or genuinely undetermined by any in-universe test — deliberately available as an unresolved axis rather than a settled property, mirroring real epistemic humility on this question.

This eight-axis frame lets Omniversal Architect Labs, One Industries, and Lewis Corp be characterized as having taken *different, non-arbitrary positions* on each axis — producing distinguishable artificial minds and institutional cultures grounded in real, current, and

honestly-labeled scientific disagreement, without requiring THE ONE UNIVERSE or THE MINUS ONE UNIVERSE to resolve questions the actual research field has not resolved.

SOURCES AND EVIDENCE NOTE

This document draws on a mix of peer-reviewed papers (arXiv preprints from DeepMind, Anthropic-affiliated and independent researchers; AAAI 2026 proceedings), primary disclosures from AI labs (Anthropic, OpenAI, Meta/AMI Labs, Google DeepMind), reporting from established outlets (TIME, MIT Technology Review coverage, IBM, VentureBeat), and clearly-marked non-peer-reviewed speculative writing (self-published Substack and Zenodo material on “relational alignment” and “Beyond AGI” essays), which is flagged as such wherever it appears above rather than presented as consensus science. Given the pace of change in this field, claims about “the current state” reflect the research available as of September 2026 and should be expected to shift.