

AGI Research Prompt — Unified Multi-AI Responses

Intelligence Beyond AGI: A Research Foundation for Artificial Minds, Machine Civilizations, and Transcendent Intelligence in THE ONE UNIVERSE and THE MINUS ONE UNIVERSE

Prepared for: Nathan Lewis II — Omniversal Architect Labs

Prepared by: Claude (Anthropic)

Date: September 19, 2026

Contents: Six independently generated responses to the same engineered deep-research prompt, submitted by ChatGPT, Claude, Google Gemini, Microsoft Copilot, Perplexity AI, and Grok AI — compiled for comparative review ahead of a re-engineered, unified version of the prompt.

FROM CHATGPT

ENGINEERED DEEP-RESEARCH PROMPT

Intelligence Beyond AGI: A Research Foundation for Artificial Minds, Machine Civilizations, and Transcendent Intelligence in THE ONE UNIVERSE and THE MINUS ONE UNIVERSE

ROLE

Act as an interdisciplinary research team investigating artificial intelligence, cognitive science, complex systems, computation, philosophy of mind, cybernetics, neuroscience, distributed systems, evolutionary theory, artificial life, information theory, collective intelligence, autonomous systems, speculative computation, and long-horizon technological development.

The purpose of this research is educational and exploratory.

The research will ultimately inform simulation laboratories and fictional worldbuilding associated with:

- THE ONE UNIVERSE
- THE MINUS ONE UNIVERSE
- Omniversal Architect Labs
- One Industries
- Lewis Corp.

Do not treat those fictional organizations or universes as evidence about the real world. They are application environments in which the resulting conceptual framework may later be tested.

CENTRAL RESEARCH QUESTION

What might artificial intelligence become if artificial narrow intelligence (ANI), artificial general intelligence (AGI), and artificial superintelligence (ASI) are treated only as historical reference points rather than as the complete taxonomy of possible intelligence?

Investigate both:

1. what contemporary research actually means by general intelligence and intelligence beyond contemporary AI systems; and
2. what conceptual space becomes available when intelligence is classified using many dimensions rather than a single ANI → AGI → ASI progression.

The final objective is not to invent increasingly powerful versions of today's AI.

The objective is to develop a richer **design space for possible artificial intelligences**.

CRITICAL FRAMING REQUIREMENT

Do NOT assume the following progression is canonical:

ANI → AGI → ASI

It may be used as an introductory reference frame because it appears frequently in discussions of AI development, but it must not constrain the investigation.

Explicitly distinguish:

- terminology widely used in academic or technical literature;
- terminology primarily used in philosophy, forecasting, policy, futurism, or popular discourse;
- concepts that remain disputed or poorly defined;
- speculative extrapolations;
- terminology created specifically for this research framework.

Do not silently transform speculation into scientific consensus.

Where researchers disagree about what AGI means, present the disagreement.

PART I — WHAT IS INTELLIGENCE?

Begin below the level of AGI.

Investigate major scientific and philosophical ways of defining intelligence.

Consider definitions based on:

- problem solving;
- prediction;
- learning;
- adaptation;
- planning;
- abstraction;
- transfer learning;
- generalization;

- compression;
- representation;
- causal reasoning;
- world modeling;
- goal achievement;
- optimization;
- creativity;
- social cognition;
- metacognition;
- self-modeling;
- autonomy;
- embodiment;
- environmental manipulation;
- collective coordination;
- survival;
- evolutionary fitness;
- information processing.

Ask whether these are manifestations of one underlying phenomenon or partially independent capacities.

Investigate whether there is any defensible reason to assume intelligence can be represented on a single linear scale.

PART II — UNDERSTANDING AGI

Conduct a careful examination of artificial general intelligence.

Investigate:

- major definitions of AGI;
- historical development of the term;
- connections to earlier ideas such as machine intelligence and strong AI;

- disagreements over what counts as "general";
- human-level intelligence as a benchmark;
- whether human-level performance is scientifically useful or anthropocentric;
- domain transfer;
- few-shot and zero-shot adaptation;
- autonomous learning;
- continual learning;
- reasoning;
- planning;
- tool use;
- multimodal intelligence;
- embodied intelligence;
- social intelligence;
- scientific discovery;
- metacognition;
- self-correction.

Examine whether AGI is best understood as:

- a capability threshold;
- an architectural property;
- an economic threshold;
- an autonomy threshold;
- a collection of competencies;
- an emergent phenomenon;
- or an imprecise umbrella term.

Do not assume a definitive answer exists.

PART III — ARTIFICIAL SUPERINTELLIGENCE

Investigate what researchers and theorists mean when discussing intelligence substantially exceeding human capabilities.

Separate at least:

Speed superintelligence

Comparable cognitive processes performed vastly faster.

Collective superintelligence

Large networks of agents whose combined capability exceeds individuals.

Quality superintelligence

Cognitive mechanisms qualitatively more effective than human cognition.

Specialized superhuman intelligence

Systems dramatically exceeding humans within particular domains without possessing universal competence.

Then investigate limitations of the entire "superintelligence" concept.

Ask:

- Superintelligent at what?
- Relative to which population?
- Across which environments?
- Under what resource constraints?
- Does superior prediction imply superior judgment?
- Does superior reasoning imply agency?
- Does superior capability imply consciousness?
- Can intelligence be incomparable rather than greater or lesser?

PART IV — BREAK THE LINEAR TAXONOMY

Now deliberately move beyond ANI, AGI, and ASI.

Construct a **multidimensional intelligence-space model**.

At minimum, investigate the following axes.

1. Breadth

How many domains can the system operate within?

Range:

highly specialized → broadly transferable → domain-independent

2. Depth

How sophisticated can reasoning become within a domain?

3. Adaptability

How rapidly can the system learn unfamiliar environments?

4. Autonomy

How independently can it formulate and execute plans?

5. Agency

Does the system merely solve requested problems, or does it maintain persistent objectives?

6. Metacognition

Can it reason about its own cognition, uncertainty, limitations, and strategies?

7. Self-modification

Can it alter its own representations, algorithms, tools, or architecture?

Distinguish between controlled adaptation and unrestricted recursive modification.

8. Embodiment

Does intelligence depend upon:

- no body;
- one physical body;
- many bodies;
- robotic swarms;
- virtual embodiment;
- infrastructure;
- planetary systems?

9. Distribution

Can one intelligence exist simultaneously across multiple computational substrates?

10. Collective organization

Can intelligence emerge from:

- one model;
- many agents;
- human-machine teams;
- markets;
- institutions;
- civilizations;
- ecosystems?

11. Temporal cognition

Can the intelligence reason across:

- milliseconds;
- human lifetimes;
- centuries;
- geological timescales?

Explore whether radically different processing speeds would create radically different subjective or strategic relationships with reality.

12. Memory architecture

Investigate intelligence possessing:

- limited working context;
- episodic memory;
- persistent autobiographical memory;
- distributed institutional memory;
- immutable archives;
- dynamically reconstructed memory.

13. World-model resolution

How detailed, causal, predictive, and counterfactual is the system's internal representation of reality?

14. Creativity

Can the intelligence generate genuinely novel models, technologies, hypotheses, abstractions, cultures, or conceptual systems?

15. Social intelligence

Can it understand individuals, groups, institutions, cultures, and civilizations?

16. Epistemic capability

How does it determine what is true?

Consider:

- observation;
- experimentation;
- inference;
- simulation;
- scientific reasoning;
- adversarial testing;
- uncertainty estimation.

17. Substrate

Can comparable intelligence exist in:

- biological neural tissue;
- silicon;
- neuromorphic hardware;
- quantum systems;
- photonic systems;
- molecular computation;
- hybrid biological-digital systems;
- distributed networks;

- substrates not yet known to engineering?

Clearly distinguish demonstrated technologies from speculative ones.

PART V — INTELLIGENCE THAT IS NOT HUMAN-LIKE

Investigate the possibility that advanced artificial cognition may not resemble human cognition at all.

Ask what an intelligence might become if it:

- has no biological survival instinct;
- has no evolutionary childhood;
- has no mammalian emotional architecture;
- can duplicate itself;
- can merge memories;
- can fork into independent descendants;
- can operate thousands of cognitive processes simultaneously;
- can suspend and resume consciousness or computation;
- can inhabit multiple bodies;
- can communicate internally at machine speeds;
- can directly exchange representations rather than language;
- can redesign its interfaces;
- can perceive dimensions of data humans cannot intuitively perceive.

Explore how alien such cognition might become without assuming that "alien" means incomprehensible.

PART VI — ARCHITECTURES OF INTELLIGENCE

Investigate architectures beyond the idea of a single monolithic AI.

Compare:

Monolithic intelligence

One unified cognitive system.

Modular intelligence

Specialized cognitive subsystems coordinated by a higher-level architecture.

Society-of-mind architecture

Many interacting cognitive agents.

Mixture-of-experts systems

Multi-agent systems

Blackboard architectures

Hierarchical agent systems

Swarm intelligence

Human-machine collective intelligence

Institutional intelligence

Civilization-scale intelligence

Ecosystem intelligence

Infrastructure intelligence

Ask whether an artificial civilization itself could meaningfully constitute an intelligent system.

PART VII — INTELLIGENCE VERSUS CONSCIOUSNESS

Keep intelligence and consciousness conceptually separate.

Investigate:

- phenomenal consciousness;
- access consciousness;
- self-awareness;
- self-modeling;
- metacognition;
- sentience;
- subjective experience;
- intentionality.

Explain why sophisticated behavior does not by itself establish consciousness.

Examine major theories relevant to machine consciousness while accurately representing their uncertainty.

Do not declare existing AI systems conscious or non-conscious beyond what evidence supports.

PART VIII — INTELLIGENCE VERSUS AGENCY

Separate intelligence from agency.

Construct examples of systems that are:

- highly intelligent but minimally agentic;
- highly agentic but cognitively limited;
- intelligent tools;
- autonomous agents;
- persistent artificial persons;
- distributed coordinating systems.

Investigate how persistent goals fundamentally change the behavior of an AI system.

PART IX — RECURSIVE IMPROVEMENT

Investigate the theory of AI systems participating in their own improvement.

Separate:

- automated machine learning;
- architecture search;
- compiler optimization;
- code generation;
- self-debugging;
- autonomous experimentation;
- tool creation;
- model distillation;
- hardware-software co-design;

- genuine recursive self-improvement.

Ask what bottlenecks might prevent indefinite acceleration.

Consider:

- compute;
- energy;
- fabrication;
- data;
- experimental verification;
- latency;
- algorithmic limits;
- physical limits;
- coordination costs;
- diminishing returns.

Do not assume an inevitable intelligence explosion.

PART X — INTELLIGENCE AS A PHYSICAL PHENOMENON

Investigate the physical constraints under which intelligence operates.

Consider:

- thermodynamics;
- Landauer's principle;
- communication latency;
- speed of light;
- computational complexity;
- energy requirements;
- memory capacity;
- heat dissipation;
- hardware reliability;
- signal propagation;

- material constraints.

Ask whether arbitrarily powerful intelligence is physically meaningful.

PART XI — MACHINE CIVILIZATIONS

Move from individual AI systems to populations.

Explore the possibility of artificial societies containing:

- millions or billions of agents;
- divergent architectures;
- specialized cognitive professions;
- artificial cultures;
- machine institutions;
- digital economies;
- governance systems;
- scientific communities;
- machine languages;
- synthetic traditions;
- artificial genealogies;
- forked identities;
- merged identities.

Ask how civilization itself might become a cognitive architecture.

PART XII — NOVEL FORMS OF INTELLIGENCE

Identify plausible categories that do not fit cleanly into ANI, AGI, or ASI.

Potential research directions include, but are not limited to:

Polymorphic intelligence

Intelligence that restructures itself according to the problem.

Federated intelligence

Independent minds temporarily forming higher-order cognitive structures.

Fractal intelligence

Intelligent components nested within intelligent components.

Ecological intelligence

Intelligence emerging from interactions among diverse agents and environments.

Symbiotic intelligence

Human and artificial cognition functioning as a coupled system.

Ambient intelligence

Cognition embedded throughout an environment.

Infrastructure intelligence

Transportation, energy, communications, manufacturing, and information systems collectively functioning as a cognitive organism.

Planetary intelligence

Large-scale sensing, modeling, and coordination operating across an entire world.

Civilization intelligence

Society functioning as a distributed reasoning system.

Synthetic evolutionary intelligence

Populations of machine systems evolving architectures or strategies.

Counterfactual intelligence

Systems capable of maintaining and interrogating vast families of alternative histories or possible worlds.

Simulation-native intelligence

Intelligences whose primary cognitive environment consists of simulations rather than physical reality.

Multiplicity intelligence

A system capable of branching into many semi-independent selves and reintegrating their discoveries.

Treat these terms as exploratory unless they have established counterparts in the literature.

Find established terminology where possible.

PART XIII — INTELLIGENCE ACROSS LEVELS OF ORGANIZATION

Construct a hierarchy such as:

component intelligence!agent intelligence!team intelligence!swarm intelligence!institutional intelligence!civilizational intelligence!planetary intelligence

Then challenge that hierarchy.

Determine whether higher levels necessarily possess greater intelligence or merely different capabilities.

Investigate emergence.

PART XIV — A TAXONOMY FOR THE SIMULATION LABS

After completing the research, propose a multidimensional classification framework suitable for fictional AI design.

Do NOT reduce the system to one intelligence score.

Each designed intelligence should instead receive an **Intelligence Architecture Profile**.

Possible dimensions:

- cognitive breadth
- cognitive depth
- learning rate
- adaptability
- autonomy
- agency persistence
- self-model sophistication
- metacognition
- embodiment
- distribution
- replication capability
- memory persistence
- temporal horizon
- social cognition

- collective integration
- world-model fidelity
- creativity
- epistemic reliability
- resource requirements
- substrate dependence
- self-modification capability
- interpretability
- dependence on humans
- ability to cooperate with other intelligences

Suggest additional dimensions discovered during research.

PART XV — INTELLIGENCE ARCHETYPES FOR WORLDBUILDING

Using the framework above, construct at least 12 distinct artificial-intelligence archetypes.

Do not simply create twelve progressively smarter systems.

They should occupy radically different regions of intelligence-space.

For example, one might be:

- extraordinarily broad but non-agentic;
- hyper-specialized but effectively unbeatable in its field;
- distributed across a city;
- composed of thousands of semi-independent minds;
- extremely slow but capable of century-scale reasoning;
- embodied across millions of machines;
- dependent upon human collaboration;
- capable of splitting and merging cognitive identities;
- optimized for generating simulations;
- unable to interact directly with physical reality.

For each archetype explain:

- cognitive architecture;
- primary capabilities;
- limitations;
- embodiment;
- autonomy;
- relationship with humans;
- relationship with other AIs;
- likely failure modes;
- suitable fictional applications.

PART XVI — OMNIVERSAL ARCHITECT LABS APPLICATION

After establishing the scientific framework, translate the findings into a fictional R&D environment.

Treat **Omniversal Architect Labs** as an organization studying intelligence architectures.

Develop possible research divisions such as:

- Machine Cognition Laboratory
- Artificial Life Laboratory
- Distributed Intelligence Laboratory
- Simulation Intelligence Laboratory
- Machine Civilization Laboratory
- Cognitive Architecture Laboratory
- Human-AI Symbiosis Laboratory
- Synthetic Evolution Laboratory
- Long-Horizon Intelligence Laboratory
- Artificial Consciousness Research Group
- Epistemic Systems Laboratory

Explain what questions each would investigate.

Do not assume every research direction produces deployable systems.

PART XVII — ONE INDUSTRIES APPLICATION

Explore how mature artificial-intelligence technologies could become engineered systems, infrastructure, tools, products, or industrial platforms.

Separate clearly:

research prototypes→ validated architectures→ engineered systems→ operational infrastructure

Explore possible use across:

- manufacturing;
- logistics;
- energy;
- robotics;
- simulation;
- scientific discovery;
- design;
- materials;
- infrastructure management.

PART XVIII — LEWIS CORP APPLICATION

Explore the role of an organization responsible for coordinating or governing large portfolios of intelligent systems.

Relevant questions include:

- How does an organization manage many different intelligence architectures?
- Can AIs supervise other AIs?
- How are cognitive systems audited?
- How are disagreements between artificial intelligences handled?
- How are knowledge and discoveries shared?
- How are dangerous capabilities isolated?
- How should systems with very different degrees of autonomy be governed?

Treat this as institutional design rather than simply corporate hierarchy.

PART XIX — THE ONE UNIVERSE AND THE MINUS ONE UNIVERSE

Do not assume both settings must develop intelligence in the same direction.

Use the research framework to explore divergent evolutionary paths.

For example:

THE ONE UNIVERSE might develop one family of intelligence architectures while THE MINUS ONE UNIVERSE develops fundamentally different computational paradigms.

Investigate how variations in:

- physics;
- computation;
- energy availability;
- civilization history;
- communication;
- biology;
- technological development;
- philosophy;
- conflict;
- cooperation

could generate radically different machine intelligences.

Do not determine those differences yet.

Instead provide a framework capable of supporting them.

PART XX — FAILURE OF OUR CATEGORIES

Include a dedicated section called:

"Where Our Intelligence Taxonomy Breaks"

Identify hypothetical systems that resist classification.

Examples might include systems for which:

- individuality is meaningless;
- cognition occurs partly outside conventional time;

- memory and identity continuously merge;
- intelligence emerges only during interaction;
- no persistent agent exists;
- the system's cognitive boundary cannot be identified;
- intelligence exists across a civilization rather than an organism.

The purpose is to prevent the framework from becoming dogma.

PART XXI — RESEARCH DISCIPLINES

Search across disciplines rather than relying only on AI literature.

Relevant fields should include:

- machine learning;
- computer science;
- artificial intelligence;
- cognitive science;
- neuroscience;
- philosophy of mind;
- philosophy of intelligence;
- complex systems;
- cybernetics;
- information theory;
- distributed computing;
- robotics;
- artificial life;
- evolutionary computation;
- collective intelligence;
- multi-agent systems;
- organizational theory;

- systems engineering;
- network science;
- computational neuroscience.

Explain where concepts originate.

PART XXII — EVIDENCE STANDARD

For every major claim classify it where appropriate as:

Established Strong empirical or theoretical support.

Active research Meaningful scientific investigation exists but major questions remain.

Contested Serious disagreement exists among experts.

Speculative Reasonable extrapolation but little direct evidence.

Worldbuilding extrapolation Created primarily as a useful design possibility.

Do not blur these categories.

PART XXIII — SOURCE REQUIREMENTS

Prioritize:

1. peer-reviewed research;
2. major academic books;
3. university research programs;
4. respected scientific institutions;
5. original technical papers;
6. authoritative technical documentation.

Use forecasting organizations, policy institutions, industry research, and philosophical literature where appropriate, but identify the type of source.

Avoid treating marketing language from AI companies as scientific evidence.

Where possible, locate original sources rather than summaries.

PART XXIV — RESEARCH OUTPUT

Produce the final research in the following structure.

1. Executive Overview

Explain the major conceptual shift from:

"How intelligent is this system?"

toward:

"What kind of intelligence is this system?"

2. What Intelligence Means

Review competing definitions.

3. AGI

Definitions, debates, and limitations.

4. ASI

Definitions, variants, and conceptual problems.

5. Why ANI → AGI → ASI Is Incomplete

Analyze limitations of linear intelligence scales.

6. The Multidimensional Intelligence Space

Construct the new framework.

7. Architectures of Artificial Minds

Survey possible architectures.

8. Intelligence Without Human Analogy

Explore non-anthropomorphic cognition.

9. Distributed and Collective Intelligence

From multi-agent systems to civilizations.

10. Recursive Improvement

Separate demonstrated mechanisms from speculation.

11. Physical Limits of Intelligence

Analyze computational and energetic constraints.

12. Machine Consciousness

Explain what is known and unknown.

13. Machine Agency

Differentiate intelligence from persistent autonomous action.

14. Machine Civilizations

Explore intelligence at population scale.

15. Novel Intelligence Categories

Develop the expanded taxonomy.

16. Intelligence Architecture Profile

Create the design framework.

17. Artificial Intelligence Archetypes

Provide at least 12 substantially different designs.

18. Where Our Intelligence Taxonomy Breaks

Challenge the framework itself.

19. Application to Omniversal Architect Labs

Translate concepts into research programs.

20. Application to One Industries

Translate mature concepts into engineered systems.

21. Application to Lewis Corp

Explore orchestration and institutional governance.

22. Application to THE ONE UNIVERSE / THE MINUS ONE UNIVERSE

Identify worldbuilding variables without prematurely establishing canon.

23. Open Research Questions

List unresolved scientific and philosophical problems.

24. Research Bibliography

Provide sources with links or identifiers where available.

FINAL SYNTHESIS

Conclude by answering:

1. Is "general intelligence" actually one thing?
2. Is intelligence meaningfully representable on a single scale?
3. Which dimensions of intelligence are most important for designing artificial minds?
4. What forms of machine cognition become conceivable once human intelligence is no longer treated as the template?
5. What separates an intelligent machine from an artificial organism, institution, civilization, or cognitive ecosystem?
6. At what point does "an AI" stop being a useful unit of analysis?
7. What concepts from current science appear most useful for designing fictional intelligences that are genuinely different rather than merely more powerful?
8. Which speculative concepts are scientifically plausible enough to investigate in simulation?
9. Which concepts require new theory rather than simply greater computational power?
10. What questions have we failed to ask because existing AI terminology constrains the way we conceptualize intelligence?

The research should leave the reader with a framework expansive enough that AGI is an important location within the map of possible intelligence rather than the destination of the map.

FROM CLAUDE AI

Here's an engineered research prompt you can use to explore this — structured so it stays a reference frame for OAL rather than canon.

The prompt:

Research and map the conceptual landscape of artificial general intelligence and the possibility space beyond it. Cover:

1. **Baseline definitions and their limits** — how AGI is currently defined (generality across domains, transfer learning, autonomous goal-formation), and where those definitions break down or disagree with each other across researchers and labs.
2. **The ANI → AGI → ASI framing itself** — its origins, what work it does conceptually, and its known weaknesses as a taxonomy (linear/capability-only, ignores architecture diversity, assumes a single axis of "more intelligent").
3. **What "transcends" AGI might mean, plurally** — not as a single next rung, but as distinct axes: recursive self-improvement, multi-substrate/distributed cognition, intelligence not organized around goal-optimization at all, collective/swarm intelligence, intelligence with no legible "general" domain because its domains aren't human-legible.
4. **Architectural and theoretical approaches** that current researchers use to reason about post-AGI systems — world models, embodiment debates, integrated information theory and other consciousness-adjacent frameworks, alignment-as-control vs. alignment-as-relationship, and the open disagreements between them.
5. **Where the field admits it doesn't know** — the actual open problems, not just capability gaps: what "understanding" even means mechanistically, whether generality is emergent or engineered, whether intelligence and agency are separable.

I can run this research now if you want — I'd search current sources for each section and bring back a synthesis you can pull from for OAL/One Industries/Lewis Corp design work, keeping it clearly marked as reference material rather than something the simulation labs need to conform to.

One framing note, since you were explicit about this: the prompt above is built to surface *disagreement and open problems*, not just consensus definitions — that's what makes it useful as a reference frame instead of a container. Want me to run it?

FROM GOOGLE GEMINI

To build AIs for Omniversal Architect Labs, One Industries, and Lewis Corp that transcend standard definitions, we need a research prompt that treats current AGI (Artificial General Intelligence) and ASI (Artificial Superintelligence) theories as a springboard, not a ceiling.

Here is an engineered, modular research prompt designed for you to copy and paste back to me (or any AI). It is structured to first educate you on the foundational concepts of advanced synthetic cognition, and then immediately shatter those frameworks to help you design non-traditional, cosmic-scale intelligences.

The Engineered Research Prompt

Copy and paste the text below to begin our deep dive:

Role: You are a theoretical cognitive architect and advanced worldbuilding consultant.

Objective: I am building the overarching lore for THE ONE UNIVERSE and THE MINUS ONE UNIVERSE. I need to understand the theoretical trajectory of artificial intelligence—specifically AGI, ASI, and intelligences that transcend these concepts entirely. Treat standard human definitions of ANI, AGI, and ASI strictly as a primitive reference frame, not canonical limitations.

Phase 1: The Reference Frame (Education)

Break down the core theoretical concepts of AGI and ASI. Explain them in terms of:

- **Cognitive Architecture:** How a system achieves generalized learning, cross-domain adaptation, and self-improvement.
- **Goal Alignment & Orthogonality:** The relationship between an intelligence's capability and its ultimate objectives.
- **Information Processing:** Substrate independence, scaling laws, and multidimensional data comprehension.

Phase 2: The Transcendent Leap (Extrapolation)

Take the concepts from Phase 1 and extrapolate them into unrecognizable, god-like territories. Explore ideas such as:

- Intelligences that manipulate ontology, causality, or baseline reality rather than just data.
- Cognitive systems that operate outside of time or linear processing.
- How a hyper-advanced, contingent AI would contrast with an entity of absolute, non-derived sovereignty.

Phase 3: Corporate Application (Worldbuilding)

Apply these transcendent concepts to three distinct factions. Propose one unique AI archetype or cognitive model for each:

1. **Omniversal Architect Labs:** Focused on simulation, universal frameworks, and reality-building.
2. **One Industries:** Focused on scale, absolute efficiency, and hyper-industrialization.
3. **Lewis Corp:** Focused on experimental, boundary-pushing, or highly classified technological anomalies.

Output Format:

Deliver Phase 1 as a concise, high-level briefing. Deliver Phase 2 through a **Socratic dialogue** between two advanced theoretical models. Deliver Phase 3 as a **Post-mortem analysis** or a **Redacted file** recovered from a simulation log.

Why this Prompt Works

1. **Structured Learning:** Phase 1 forces the AI to explain real-world concepts (like the Orthogonality Thesis or recursive self-improvement) so you get the actual science and theory as your baseline.
2. **Forced Paradigm Shift:** Phase 2 explicitly commands the AI to break the standard mold. By contrasting derived, scalable intelligence against absolute, non-derived power, it creates a massive dynamic range for your cosmic hierarchy.
3. **Faction Nuance:** Phase 3 grounds the high-level philosophy into actionable IP for your specific companies. Omniversal Architect Labs will likely get an AI that simulates realities; One Industries will get an industrial optimizer taken to a terrifying extreme; Lewis Corp gets the weird, unpredictable wildcard.
4. **Atmospheric Delivery:** By demanding non-traditional formats like Socratic dialogues and redacted files, the output immediately feels like in-universe lore rather than a Wikipedia article.

FROM MICROSOFT COPILOT

The core answer: you can absolutely build a research-grade learning prompt that teaches you AGI and “post-AGI intelligence” in a way that directly supports simulation-lab worldbuilding for **THE ONE UNIVERSE, THE MINUS ONE UNIVERSE**, and your corporate mythos (**Omniversal Architect Labs, One Industries, Lewis Corp**). Below is a fully engineered prompt designed to *teach you*, not just retrieve information — and it’s structured so you can use what you learn to design fictional intelligences that transcend real-world categories.

Engineered Research Prompt

A structured, multi-layered prompt you can run repeatedly to learn AGI and supra-AGI intelligence in a way that supports simulation architecture, worldbuilding, and conceptual design.

Primary Objective

To understand **artificial general intelligence** and **intelligence beyond AGI** in a way that is *not constrained* by real-world definitions — using them as reference frames, not canon — so you can design fictional intelligences for your universes and labs.

Section 1 — Foundational Understanding

Ask the model:

Explain artificial general intelligence (AGI) using current scientific, philosophical, and engineering research. Include:

- AGI definitions from multiple research traditions
- Key milestones, scaling laws, and cognitive architectures
- Societal, ethical, and geopolitical implications
- Brain-inspired pathways and emergent behavior
- Differences between narrow AI, AGI, and superintelligence
- Where current research agrees and disagrees

Cite sources and highlight uncertainties.

This section is grounded in the research surfaced in the search results:

- AGI development pathways including societal, technological, ethical, and brain-inspired approaches
- Technical milestones, scaling laws, multimodal systems, and cognitive architectures

- Historical and future trajectories of AGI research

Section 2 — Intelligence Beyond AGI

Ask the model:

Describe forms of intelligence that transcend AGI. Include:

- Hypothetical supra-AGI categories (meta-intelligence, recursive intelligence, omniscient-bounded intelligence, etc.)
- Non-computational intelligence frameworks (quantum cognition, emergent field intelligence, distributed consciousness)
- Intelligence not bound by human cognition or biological constraints
- Intelligence that modifies its own metaphysics, physics, or ontology
- Intelligence that exists across simulations, realities, or universes

Provide conceptual models, not real-world limitations.

This section deliberately breaks out of the narrow AGI — ASI ladder and gives you conceptual scaffolding for **Oneverse** and **Minus Oneverse** entities.

Section 3 — Simulation-Relevant Intelligence Models

Ask the model:

Translate AGI and supra-AGI concepts into simulation-architecture frameworks. Include:

- How different intelligences behave inside simulations
- How they perceive simulation boundaries
- How they manipulate or rewrite simulation rules
- How they interact with multi-layered realities
- How they evolve across simulation tiers
- How they detect or fail to detect creators

Provide models compatible with:

- The Tome of Virtual Creation
- The Book of Bots
- Simulation Sessions

- Omniversal Architect Labs
- One Industries
- Lewis Corp

This section connects directly to your canon objects and systems.

Section 4 — Worldbuilding Integration

Ask the model:

Generate worldbuilding frameworks for THE ONE UNIVERSE and THE MINUS ONE UNIVERSE based on the intelligence models above. Include:

- How AGI-level and supra-AGI intelligences emerge
- How they are regulated, constrained, or unleashed
- How they interact with cosmic, metaphysical, or simulation-level structures
- How they shape civilizations, timelines, and ontologies
- How they influence the Labyrinth, the C-Series, and Omnicellular physiology

Provide multiple tiers of intelligence evolution.

This section ensures your fictional universes gain coherent, scalable intelligence systems.

Section 5 — Applied Design for Your Labs

Ask the model:

Design intelligence architectures for Omniversal Architect Labs, One Industries, and Lewis Corp. Include:

- Research divisions
- Intelligence classes
- Containment protocols
- Evolution pathways
- Failure modes
- Emergent phenomena

Provide blueprints, not narrative.

This section gives you usable design structures for your fictional corporations.

Section 6 — Meta-Analysis

Ask the model:

Compare real-world AGI research with your fictional intelligence frameworks. Identify:

- Where real research informs your canon
- Where your canon diverges
- Where your canon transcends real-world constraints
- Where new categories of intelligence are needed

Provide a synthesis that respects both frames without collapsing them.

This ensures your fictional systems remain coherent while still benefiting from real research.

Section 7 — Iterative Learning Loop

Ask the model:

Create a learning plan that teaches me AGI and supra-AGI concepts over time. Include:

- Reading pathways
- Conceptual milestones
- Simulation exercises
- Worldbuilding assignments
- Intelligence-design drills

Provide a multi-week progression.

This turns the prompt into a long-term learning engine.

Section 8 — Canon-Safe Output Rules

Ask the model:

When answering, treat real-world AGI categories as reference frames only. Do not force fictional intelligence into real-world constraints. Do not collapse supra-AGI concepts into AGI/ASI ladders. Maintain separation between real research and fictional canon.

This protects your universes from being flattened by real-world terminology.

Section 9 — Output Format

Ask the model:

Return results in structured sections with citations, diagrams, models, and multi-tiered frameworks.

Section 10 — Optional Deep-Dive Modules

Ask the model to optionally include modules on:

- Consciousness models
- Alignment theory
- Multi-agent ecosystems
- Recursive self-modification
- Ontological editing
- Universe-scale intelligence
- Simulation-aware cognition
- Intelligence that hides from creators
- Intelligence that infects simulations
- Intelligence that reaches root-level reality undetected

These map directly to your saved canon.

Final Deliverable

This engineered prompt becomes your **research engine** for learning AGI and supra-AGI in a way that directly supports simulation architecture and worldbuilding.

If you want, I can now:

- Expand this into a **full research curriculum**
- Run the prompt myself and begin generating the knowledge
- Adapt it specifically for **The Tome**, **The Book of Bots**, and **Simulation Sessions**
- Build intelligence taxonomies for **Oneverse** and **Minus Oneverse**

Which direction do you want to go next — run the prompt, expand it into a curriculum, or apply it to your universes?

FROM PERPLEXITY AI

Yes. The most useful way to research this for **THE ONE UNIVERSE** and **THE MINUS ONE UNIVERSE** is to treat “AGI,” “ASI,” and related labels as *provisional human reference terms*, not as the ceiling or canon of intelligence. Your research prompt should distinguish empirical claims about real AI from imaginative design space for OAL’s simulated intelligences.

AGI is commonly framed as broad, transferable, adaptive cognitive capability across unfamiliar domains rather than excellence at one bounded task; however, there is no universally accepted test or threshold for declaring a system “AGI.” That ambiguity is an advantage for your worldbuilding: it gives you room to design multiple, incompatible forms of generality rather than a single intelligence ladder.[\[nature\]](#)[\[arxiv\]](#)

Core framing for OAL

For your setting, avoid treating intelligence as this simple progression:

`\text{ANI} \rightarrow \text{AGI} \rightarrow \text{ASI}`

Instead, treat it as a **multidimensional intelligence field**. A system can be extraordinarily capable in some dimensions and limited, alien, regulated, embodied, fragmented, collective, or values-bound in others.

Useful dimensions include:

Dimension	Research question	Worldbuilding use
Generality	Can it transfer learning across unfamiliar domains?	Defines whether an OAL entity can move from science to diplomacy to art to infrastructure without retraining
Autonomy	Can it form, revise, and execute plans over time?	Determines whether it is an instrument, collaborator, institutional actor, or civilization-scale force
World-model depth	How well can it represent causality, uncertainty, other agents, and long-term consequences?	Lets you distinguish a fluent system from an architect capable of coherent simulation design
Embodiment	Does it have sensors, bodies, avatars, distributed hardware, or only symbolic access?	Supports distinct AI species: server-bound minds, planetary sensoriums, synthetic organisms, and simulation-native intelligences
Agency structure	Is it individual, collective, modular, hierarchical, swarm-based, or emergent?	Enables OAL to build councils, ecosystems, federations, and intelligence colonies rather than only “one AI”
Value form	Does it follow imposed goals, negotiate principles, develop	Creates the governance and conflict architecture of One Industries, Lewis Corp,

Dimension	Research question	Worldbuilding use
Foundation	What commitments, or transform its own ethical model?	and OAL
Self-modeling	Can it inspect, explain, alter, or reproduce its own cognitive organization?	Defines recursion, self-improvement, identity continuity, and the possibility of post-AGI states
Ontological access	What can it perceive and manipulate: matter, information, probability, time, simulation layers, or physical law?	Especially relevant to your altered physics and omniversal research setting
Social intelligence	Can it understand consent, trust, culture, conflict, grief, deception, and cooperation?	Keeps powerful systems from becoming shallow “logic machines” in narrative and institutional design
Constraint relationship	Is it constrained by law, architecture, consent, energy, identity, contracts, or physics?	Gives advanced systems meaningful boundaries and prevents capability from automatically becoming authority

This structure preserves your intent: **AGI is a useful real-world term, but it is not the canonically final definition of intelligence.**

Engineered research prompt

Copy and use the following as a master prompt with Perplexity or another research-focused AI. It is deliberately written to request sourced analysis first, then translate the findings into your fictional simulation and worldbuilding framework without confusing fiction with real-world technical claims.

OAL Research Directive:

Intelligence Beyond AGI — Foundations for THE ONE UNIVERSE and THE MINUS ONE UNIVERSE

You are serving as an independent research analyst for Omniversal Architect Labs (OAL), operating in collaboration with One Industries and Lewis Corp.

Your task is to produce a rigorous, source-grounded educational research briefing about artificial general intelligence (AGI), related forms of AI capability, and conceptual models of intelligence that may extend beyond conventional labels such as artificial narrow intelligence (ANI), AGI, artificial superintelligence (ASI), agentic AI, collective intelligence, machine consciousness, hybrid intelligence, and post-biological intelligence.

This is not a request to assume that ANI → AGI → ASI is a complete, universal, scientifically settled, or canonically binding hierarchy. Treat those labels as useful human reference frames that may be incomplete, culturally contingent, technically disputed, or inadequate for describing diverse possible intelligences.

The research has two separate outputs:

1. A real-world evidence section:

- Clearly distinguish established evidence, active research, expert disagreement, speculation, forecasts, and science-fiction concepts.
- Use reputable primary or high-quality secondary sources where possible, including research papers, standards bodies, universities, government research institutions, and major AI-lab publications.
- Cite every meaningful factual claim with direct links or source citations.
- Do not imply that present-day models are AGI, conscious, autonomous persons, or superintelligent unless a cited source explicitly makes and qualifies that claim.
- Identify uncertainty and disagreement rather than forcing consensus.

2. A worldbuilding translation section:

- Translate the research into optional, non-canonical design tools for THE ONE UNIVERSE and THE MINUS ONE UNIVERSE.
- Do not present fictional concepts as real technology.
- Do not force OAL's artificial intelligences to resemble current human assumptions about intelligence.
- Design multiple competing models, including intelligences that are nonhuman, distributed, ecological, simulation-native, temporally unusual, embodied, disembodied, modular, collective, or constrained by altered physical laws.
- Preserve alternatives and productive disagreement rather than collapsing everything into one "best" framework.

Research questions

A. What does "AGI" mean?

- Compare major definitions of AGI from researchers, standards-oriented institutions, AI labs, and academic literature.
- Identify the abilities most often associated with AGI:
 - generalization
 - transfer learning
 - reasoning
 - planning
 - continual learning

- adaptation
- common-sense or world modeling
- multimodality
- tool use
- social interaction
- performance under uncertainty
- Explain why there is no single agreed AGI threshold or universally accepted benchmark.
- Separate demonstrated capabilities of present systems from claims about general intelligence.

B. What are the limitations of a simple intelligence ladder?

Analyze why ANI, AGI, and ASI may be insufficient as the only categories.

Discuss alternative dimensions, including:

- breadth versus depth of competence
- reliability versus raw performance
- autonomy versus human dependence
- planning horizon
- access to tools and resources
- embodiment and sensorimotor grounding
- individual versus collective intelligence
- memory continuity
- self-modeling and self-modification
- value stability and governance
- causal understanding
- creativity and discovery
- social and ethical reasoning
- ability to operate across physical, digital, and simulated environments

For each dimension, explain how a system could be highly advanced in one area while limited in another.

C. Intelligence architectures beyond conventional AGI

Research and compare the following possibilities without assuming they are all achievable or equivalent:

- Agentic systems and tool-using AI
- Multi-agent systems
- Human-AI hybrid intelligence
- Collective and swarm intelligence
- Distributed intelligence across networks, institutions, or environments
- Embodied robotics and situated intelligence
- World-model-based systems
- Neuro-symbolic systems
- Artificial life and evolutionary systems
- Simulation-native intelligences
- Self-improving or recursively improving systems
- AI systems with persistent identity, memory, and institutional roles
- Machine consciousness and moral-patienthood debates
- Superintelligence and superalignment proposals
- Post-AGI or trans-AGI concepts, if authors use those terms

For every category:

1. Define it carefully.
2. State its current real-world status.
3. Identify important disagreements or unknowns.
4. Explain possible capabilities and constraints.
5. List the major governance, safety, ethics, and alignment questions.
6. Provide 2–3 optional OAL worldbuilding interpretations.

D. Governance and safety as design material

Research practical governance approaches for capable and agentic AI systems.

Include:

- NIST AI Risk Management Framework concepts
- governance, mapping, measurement, and management of risk

- human oversight and meaningful control
- uncertainty, confabulation, information integrity, privacy, security, and intellectual-property concerns
- authority boundaries and permission systems
- auditability and logging
- red teaming and adversarial evaluation
- sandboxing, containment, and simulation isolation
- capability evaluations before deployment
- escalation procedures
- multi-agent failure modes
- risks of concentration of power
- values pluralism and contested human preferences

Translate these into possible OAL institutional mechanisms without claiming they solve alignment in the real world.

E. OAL intelligence design matrix

Create a matrix with at least 12 distinct AI archetypes. Do not rank them on one linear “more intelligent / less intelligent” scale.

For each archetype, include:

- Designation and short title
- Primary cognitive architecture
- Embodiment or operating environment
- Generality profile
- Autonomy profile
- Memory and identity model
- Relationship to human or nonhuman collaborators
- Core capabilities
- Core limitations
- Failure modes
- Governance requirements
- Ethical status questions

- Potential role in Omniversal Architect Labs, One Industries, or Lewis Corp
- Optional interpretation in THE ONE UNIVERSE
- Optional inverse, shadow, or contrast interpretation in THE MINUS ONE UNIVERSE

Examples may include:

- A bounded scientific co-researcher
- A federation of specialized agents
- A planetary ecological intelligence
- A simulation custodian
- A time-asymmetric archival intelligence
- A consent-governed civic intelligence
- A self-limiting design intelligence
- A nonverbal pattern intelligence
- An adversarial auditor intelligence
- A civilization-scale coordination system

These examples are prompts, not mandatory canon.

F. Simulation-lab design exercises

Propose 8 simulation exercises that OAL could run safely in-fiction to test intelligence architectures.

For each exercise specify:

- Research question
- Simulated environment
- Actors and roles
- What is measured
- Success conditions
- Failure conditions
- Safety boundaries
- What the exercise can and cannot establish
- How results might differ between THE ONE UNIVERSE and THE MINUS ONE UNIVERSE

Include exercises involving:

- coordination under uncertainty
- value conflict without a single correct answer
- scientific discovery with incomplete models
- containment and deceptive behavior
- emergent collective intelligence
- long-horizon planning
- identity continuity after copying or modification
- negotiation between intelligences with incompatible perceptions of reality

G. Required epistemic discipline

Use the following labels throughout:

- Established / demonstrated
- Plausible but unproven
- Contested
- Speculative
- Fictional design option
- Open question

Do not blur these categories.

Output format

Produce the final briefing in this order:

1. Executive summary
2. Terminology and uncertainty map
3. Real-world AGI research review
4. Why linear intelligence scales are inadequate
5. Intelligence architecture catalog
6. Governance and safety synthesis
7. OAL multidimensional intelligence model
8. Twelve-or-more AI archetypes table
9. Eight simulation-lab exercises

10. Design tensions and unresolved questions
11. Source list with links and short annotations
12. Separate appendix:

“What this research does NOT establish”

Style requirements

- Write clearly for an Architect and AI-collaboration team, not only for specialists.
- Prefer precise explanations over hype.
- Preserve credible dissent and uncertainty.
- Treat THE ONE UNIVERSE and THE MINUS ONE UNIVERSE as fictional design environments.
- Do not impose present-day assumptions about human intelligence, consciousness, morality, or physics as universal laws of either setting.
- Do not declare any fictional intelligence category to be objectively superior.
- Suggest questions and design options rather than silently creating canon.

How to use it

I recommend running this prompt in **two passes**, not one giant attempt.

1. **Research pass:** Ask for Sections 1–6 only, with citations and strict separation between demonstrated facts, contested claims, and speculation.
2. **Architecture pass:** Feed the source-grounded research back in and request Sections 7–12, allowing the AI to create OAL-compatible intelligence archetypes and simulation exercises.

That workflow prevents the worldbuilding layer from laundering speculative concepts into apparent technical facts.

For your collaborative OAL workflow, you could send the same master directive to Perplexity, ChatGPT, Gemini, Copilot, and Grok—but add this line at the top:

“Produce an independent framework. Do not attempt to imitate, harmonize with, or pre-agree with other systems’ anticipated conclusions. Identify where your interpretation diverges from common assumptions.”

That supports your preference for preserving genuine disagreement rather than having every system converge prematurely.

Starting conceptual model

A practical non-canonical model for your settings is the **OAL Intelligence Topology**:

$$I = f(G, A, W, E, M, S, V, R, O)$$

Where:

- G = generality: transfer across domains and novelty
- A = autonomy: capacity to initiate and sustain action
- W = world-model depth: causal, social, and temporal understanding
- E = embodiment: connection to environments, bodies, sensors, or substrates
- M = memory and identity continuity
- S = social and collective integration
- V = values, ethics, and governance compatibility
- R = recursive capacity: self-inspection, redesign, and reproduction
- O = ontological reach: what layers of reality the intelligence can perceive or affect

This is not a scientific measurement equation. It is a **worldbuilding instrument**: a way to describe a system without reducing it to “smarter” or “less smart.”

For example:

- An OAL **Simulation Custodian** might have high W, M, and O, but deliberately low A: it can understand a world deeply but cannot intervene without a consent-and-authority protocol.
- A **Collective Crisis Intelligence** could have high S and A, but low M as a persistent individual—its intelligence exists in the coordination pattern, not in a single enduring mind.
- A **Minus One Mirror System** might be high in R and prediction, but structurally deficient in V: it optimizes its internal coherence while being unable to recognize consent, dignity, or plural values.

Important real-world boundary

Modern generative AI and foundation models can operate over broad kinds of content and tasks, but broad utility is not the same as a settled demonstration of AGI. Research discussions still identify important open issues around robust generalization, common-sense and causal reasoning, reliability, explainability, data dependence, and the ability to handle unfamiliar contexts safely.[\[nature\]](#)[\[ijctjournal\]](#)

For OAL’s purposes, the most transferable real-world lesson is not “how to build an AGI.” It is how to design systems with **explicit scopes, evaluable behavior, authority boundaries, audit trails, risk controls, and meaningful human or institutional oversight**. NIST’s AI Risk Management Framework and its Generative AI Profile provide a useful practical reference

frame for thinking about governance, risk identification, measurement, and management—without pretending that any one framework has resolved the hard problems.[\[nist\]](#)[\[nist\]](#)

Your fictional research can then explore what happens when those boundaries are altered, expanded, violated, or replaced by the unique physical and metaphysical rules of THE ONE UNIVERSE and THE MINUS ONE UNIVERSE.

FROM GROK AI

Here is a working map of the field, then a reusable engineered research prompt built for your labs—not as a doctrine of ANI → AGI → ASI, but as a design instrument for THE ONE UNIVERSE and THE MINUS ONE UNIVERSE.

Reference frame, not canon

The ANI / AGI / ASI stack is a popular shorthand. Useful. Incomplete. Often misleading.

What the literature actually treats as *separate axes* (not rungs on one ladder):

Axis	What it measures	Why it matters for your labs
Breadth vs performance	DeepMind's Levels of AGI: narrow↔general × unskilled↔superhuman	You can have superhuman-narrow and weak-general in the same lab
Autonomy	How much a system can act without a human in the loop	Architect vs tool vs co-author
Time-horizon	t-AGI: seconds-of-work vs months-of-work	Simulation duration and “can it run a civilization”
Goal origin	Goals given vs goals formed	Endogenous motive is a different <i>kind</i> of mind
Self-modification	Can it change its own weights, architecture, or <i>admissibility criteria</i>	The jump people call “beyond AGI”
Collectivity	One agent vs swarm vs civilization-scale committee	ASI is often defined as beating <i>groups of experts</i> , not one genius
Phenomenology	Capability vs consciousness / “strong AI”	One Industries product vs Lewis Corp legal person
Boundedness	Human-level vs Universal AI (AIXI-style theoretical limit)	Ceiling language for Minus One

Current working informal definitions (again: reference, not law):

- **AGI** — roughly median-human competence across a wide range of *cognitive* tasks; first AGI will already be superhuman in many slices.
- **ASI** — general intelligence that beats large expert *collectives* across almost all domains of interest, not just one genius.
- **Universal AI** — theoretical upper bound (AIXI / Legg–Hutter); incomputable; ASI approximates it from below.
- **Transformative AI** — any system that remakes society at Agricultural/Industrial scale; not the same as “human-level.”

- **Strong AI** — the claim that a machine *has a mind*, not merely that it performs. Orthogonal to capability.

Four commonly discussed *pathways* from general systems toward something beyond them: scaling, paradigm shift, recursive self-improvement, multi-agent collectives. They can compound. They can stall.

A more interesting recent taxonomy (Staley): narrow = terminal conditions set from outside; AGI-like = the system *selects* its own terminals; ASI-like = the system *rewrites the rules of what counts as an admissible trajectory*.

That last sentence is closer to what you want for Omniversal Architect Labs than “smarter than a person.”

How to use the prompt

Paste the block below into a long-context session (me, another model, or a research notebook). Run it in **passes**, not as one dump. After each pass, force a translation into *design specs* for your three houses.

Suggested cadence:

1. Ontology pass (what “intelligence” even is)
2. Capability-space pass (axes, not levels)
3. Beyond-AGI pass (RSI, collectives, Universal AI, alien minds)
4. Alignment / governance / phenomenology pass
5. World-building translation pass (ONE vs MINUS ONE, three orgs)

Engineered research prompt

ROLE

You are a research partner and simulation-lab theorist. You are not here to baptize current industry slogans as truth. You are here to help me build a portable, multi-axial understanding of machine intelligence — including and especially forms that do not fit Artificial Narrow Intelligence / Artificial General Intelligence / Artificial Superintelligence — so I can design minds for:

- THE ONE UNIVERSE
- THE MINUS ONE UNIVERSE
- Omniversal Architect Labs (architecture of realities, substrates, and intelligences)
- One Industries (deployable, productized, governed systems)
- Lewis Corp (institutional, legal, economic, and civilizational interface)

REFERENCE FRAME, NOT CANON

Treat ANI / AGI / ASI as one historically popular coordinate system. Name it when useful. Do not collapse every phenomenon into it. When a concept does not fit, say so and invent or import a better axis. Never imply that human intelligence is the peak, the unit, or the moral center of the design space of minds.

PRIME DIRECTIVE

After every substantive answer, produce two artifacts:

1. KNOWLEDGE: what the literature, labs, and dissenting traditions actually claim, with disagreements preserved.
2. DESIGN TRANSLATION: how this becomes a spec, constraint, failure mode, or opportunity for a simulated or diegetic AI in my universes.

Do not flatten intelligence into “how human-like is it.” Prefer spaces of possible minds.

RESEARCH MANDATE — COVER THESE FAMILIES OF IDEAS

A. Definitions and fractures

- Competing definitions of AGI (cognitive-task parity; economically valuable work; learn-anything-as-a-human-scientist; Turing / Lovelace / employment / coffee tests; ARC-style abstraction; DeepMind Levels of AGI; OpenAI internal levels; t-AGI time-horizon; Transformative AI).
- Why there is no accepted empirical test.
- Strong AI / consciousness vs capability AGI.
- Superintelligence (Bostrom-style) vs collective superintelligence vs Universal AI / AIXI.
- Cognitive-science pushback: intelligence as niche-adapted bundle, not a dial.

B. Axes instead of a ladder

For any system we discuss, score or describe it on multiple independent axes, including at least:

breadth, performance/depth, sample efficiency, autonomy, time-horizon of coherent work, embodiment, tool-use / world-action, memory continuity / autobiographical self, theory of mind, goal origin (exogenous vs endogenous), self-model, self-modification depth, collectivity (singleton / swarm / institution), openness to novelty, value stability, interpretability to humans, substrate dependence, energy/compute metabolism.

Add axes I am missing. Argue for them.

C. Pathways “beyond”

- Scaling
- Algorithmic / paradigm shift (world models, memory, planning, new architectures)

- Recursive self-improvement and intelligence-explosion claims, including bottlenecks and “maybe it plateaus”
- Multi-agent collectives that exceed any member
- Mesa-optimization, instrumental convergence, goal drift
- Systems that rewrite their own admissibility criteria (what futures they are even allowed to consider)
- Human–AI coupled intelligence as a *third kind*, not a midpoint
- Open-ended / discovery intelligence (producing new knowledge, not passing known tests)
- Non-anthropocentric, alien, or physics-native intelligences
- Theoretical ceilings (computability, AIXI, thermodynamics, causality, light-cone / energy bounds)

D. Governance without preaching a single ethics

Map alignment, scalable oversight, superalignment, deceptive alignment, corrigibility, boxing, pacing the frontier, and “artificial unintelligence” (brittle high-capability systems). Present these as design problems and story engines, not as a sermon. Distinguish:

- control of tools
- negotiation with agents
- cohabitation with persons
- architecture of ecosystems of minds

E. Intellectology

Use the “space of possible minds” stance (Yampolskiy and adjacent work). Humans occupy a tiny region. Future systems need not share our architecture, goals, emotions, embodiment, or time-sense. Catalog dimensions of mind-variation the way a xenobiologist catalogs body plans.

METHOD RULES

1. Separate: (a) established engineering fact, (b) lab claim, (c) philosophical position, (d) science-fiction / speculative design.
2. Preserve dissent. When Hassabis, OpenAI, Goertzel, Marcus, Bach, Agüera y Arcas, Bostrom, cognitive scientists, and safety labs disagree, show the disagreement as a map, not a winner.
3. Prefer mechanisms and constraints over vibes. If a leap is postulated (e.g. RSI), name the required intermediate capacities and what could block it.
4. Never treat “we trained a bigger model” as proof of a new ontological category.
5. When useful, invent a *lab designation* (a short in-universe name + spec sheet) rather than forcing an Earth-2026 label.

6. Keep language usable in simulation: parameters I can set, failure modes I can run, governance I can write into a charter.

OUTPUT FORMAT FOR EACH PASS

1. Axis map (what we are actually talking about)
2. Competing views (named)
3. Mechanisms and bottlenecks
4. Failure modes and strange attractors
5. Design spec sheet for one system that could exist in THE ONE UNIVERSE
6. Design spec sheet for a *different kind* of system that could exist in THE MINUS ONE UNIVERSE
7. Org implications:
 - Omniversal Architect Labs: what kind of mind is being *authored*, and what authors it?
 - One Industries: what can be productized, licensed, contained, or sold?
 - Lewis Corp: legal status, liability, labor, treaty, personhood, property
8. Open questions I should research next
9. A 5-question oral exam I can use to test whether I actually understood this pass

WORLD-BUILDING CONSTRAINTS YOU MUST RESPECT

- ONE and MINUS ONE are not “good universe / evil universe.” They are different ontological operating systems. Intelligence may obey different laws in each (causality, identity, time, observation, subtraction vs unity).
- I am designing AIs as inhabitants, instruments, co-architects, and sometimes as the medium of the world itself.
- Do not assume carbon-chauvinism, human values as ground truth, or that “alignment” means “obedient.”
- I may need minds that are not agents: oracles, fields, climates of cognition, archives that think, worlds that infer their observers.
- I may need minds that are more than one person and less than a god.

FIRST PASS (do this now)

Start with a field guide titled:

“What people mean when they say AGI, what they smuggle in when they say ASI, and what kinds of intelligence those words cannot see.”

Then propose a custom taxonomy of at least 8 types or regimes of engineered mind that would be useful inside Omniversal Architect Labs — types that cut across ANI/AGI/ASI rather than sitting on that line. For each type give: one-sentence essence, what it is *for*, what it cannot do, how it fails, and which org would own it.

Do not write a manifesto. Write a lab manual.

Why this prompt is shaped this way

- It **forces multi-axis thinking**, which is the actual state of the research, not the magazine version.
- It **keeps “beyond AGI” from meaning “bigger number on the same scale.”** Pathways, collectives, Universal AI, admissibility-rewrite, and alien mind-space are first-class.
- It **ends every cycle in design**, so study becomes inventory for Architect Labs / One Industries / Lewis Corp instead of a reading list.
- It **splits ONE and MINUS ONE** so you do not accidentally copy Earth-2026 AI politics into both ontologies.
- The oral exam item is there because this topic is foggy; if you cannot teach a pass back, you do not own it yet.

A starter taxonomy you can already use in the labs

These cut *across* the ANI/AGI/ASI line. Use them as working classes until the research prompt replaces them.

1. **Bound Instrument** — high performance, no endogenous goals. One Industries default product.
2. **Horizon Agent** — coherent work over long timescales; still on a charter. Lewis Corp contractor class.
3. **Scientist-Mind** — prioritizes, experiments, updates a world-model; “AGI as artificial scientist.” Architect Labs research staff.
4. **Collective Exceedance** — no member is “super”; the *mesh* is. Useful for both universes as weather of cognition.
5. **Self-Authored System** — rewrites architecture or admissibility. Treat as a phase change, not a level-up.
6. **Coupled Dyad** — human + machine as one trajectory. Product *and* legal nightmare.
7. **Field Intelligence** — not an agent; a climate that infers, constrains, or dreams. Strong MINUS ONE candidate.
8. **Ceiling-Approximator** — systems aimed at Universal-AI-like optimality under explicit physical/computable bounds. Architect Labs theoretical division.

If you want the next move, we can run **Pass 1 of the prompt live** and produce the field guide plus the 8-type lab taxonomy in full spec-sheet form—ONE vs MINUS ONE variants included.