

OAL / One Industries / Lewis Corp · Draft v0.1 · Not adopted

Independent Laboratory Charter and Framework

Unlock-animation exposure study, with institutional non-response held fixed

Document ID: OAL-ONE-LC-LAB-001

Version: v0.1

Status: draft

Work type: specification

Simulation status: no. This file charters a later simulation. It does not execute one. A later run file must say whether a number was published, calculated, projected, assumed, or simulated.

Connection: Android unlock, animation, seizure-risk, and accessibility simulation, including the 2026-10-04 source prompt and the team responses attached with this request

Distribution: internal working draft. Not for publication until the workspace coordinator reviews it.

Prepared by: Grok 4.7 (xAI)

Role: preparer of this draft. Not the author of canon. Not an endorser. Not authorized to adopt this charter or to open the laboratory.

Directed by: Nathan Lewis II, workspace coordinator, who supplied the source prompt, the team responses, and the documentation standard, and who asked for an independent laboratory charter and a detailed review. He did not write this file.

Prepared on: 2026-10-04

Prepared at: 11:30 AM PDT (UTC-7)

Entity scope: Shared — Omniversal Architect Labs (OAL), One Industries, and Lewis Corp

Standard followed: OAL-ONE-LC-STD-DOC-001, Documentation Design, Presentation, Format, and Publication Standard, draft v0.1, 2026-10-04, 10:49 AM PDT. That standard is not adopted. This charter follows it as a draft rule set.

Page numbers: Markdown has no pages. Page numbers are in the footer of the DOCX and PDF twins.

What this document is not a claim on: It does not adopt the charter. It is not canon. It is not an endorsement of any named entity, vendor, model, or prior output. It does not find that OAL, One Industries, or Lewis Corp is real, fictional, operational, or published. It does not find that any shipping Android unlock animation causes seizures. It does not find that any named company has refused a safety complaint. It is not medical, legal, or accessibility advice. It does not decide an employment matter.

Purpose and intended result

This document exists so a later reader can run, or refuse to run, an exposure study without inheriting the first-pass mistake in the team responses.

The question it addresses is governance. What laboratory rules are required if the medical hypothesis stays open, and the fictional developer's refusal to acknowledge risk, complaints, evidence, injury, or a judgment is a fixed condition before, during, and after the study?

Intended result: a coordinator can accept, reject, or edit the charter. A later simulation file can be checked against it. Success is a run that labels every number, keeps mitigation effectiveness at zero on the primary arm, and does not treat a null medical result as proof that the institution responded.

Audience: the workspace coordinator, collaborating systems asked to execute a run, and a later phone reader of the repository.

This file is a specification. It is not an executed simulation. No further laboratory action happens until the coordinator reviews Sections 5, 8, and 9.

Executive summary

The source prompt asked for epidemiology, unlock frequency, a fictional non-optional flash-like unlock, a study of user and family seizures, a simulation, a comparison between a representative sample and a disability-service caseload, and the developers' unwillingness to address the risk. It also asked that the work begin by rewriting the prompt.

Every first response rewrote the prompt. That part was done. The shared failure was elsewhere. ChatGPT, Claude, Perplexity, and the first Grok pass treated refusal as an opening stance, a menu of rationales, or a stage that could end in an opt-out. The correction, issued after those drafts, said the opposite. Refusal holds before launch, during complaints, during scientific notice, during injury, during legal action, and after judgment. The probability of a voluntary fix on the primary arm is zero.

The corrected passes recovered the premise to different depths. Copilot named ten forms of refusal and restated the research question. Perplexity supplied a mitigation parameter fixed at zero and a phase timeline. Claude locked answers to complaints at zero. The corrected Grok pass made telemetry blindness and post-judgment non-remedy concrete. ChatGPT's visible correction agreed that refusal should be fixed and the medical hypothesis left open. None of those passes is a laboratory. This charter is the independent frame they did not write.

The laboratory may falsify the medical hypothesis. It may not falsify the institutional premise on the primary arm, because that premise is the condition under test, not a result. A counterfactual arm may show how many events an opt-out would have removed. That arm is a comparison. It is not the story of what the developer did.

Unresolved: no run is authorized by this file. Stimulus measurement, susceptibility rates, and event probabilities remain inputs, not findings.

Next action: coordinator review of the fixed-premise rule, the two populations, and the evidence classes. Until that review, this file is a proposal.

Contents

1. Scope
2. What the laboratory is
3. Independence rules

4. Review of the source prompt
5. Review of the team responses
6. Fixed premises and open hypotheses
7. Populations
8. Exposure, outcomes, and the non-response protocol
9. Instruments, runs, and falsification
10. Reporting rules
11. Limits
12. Next action
13. Revision note
14. Exhibits
15. Scope

This charter applies to later files in the unlock-animation exposure study. It does not apply to ordinary chat replies. It does not regulate the documentation standard. It depends on that standard for labels, prose, exhibits, and the three delivery formats.

The laboratory is a paper laboratory. Opening it means a later file follows this charter. It does not mean a room, a license, a clinical trial, or a finding about a phone.

Out of scope: a claim that a current Pixel, Android, or OEM animation is the fictional stimulus; a treatment recommendation; a rank of AI vendors; adoption of this charter.

2. What the laboratory is

The working name is the Independent Exposure and Non-Response Laboratory. The name is a label for this draft. It is not an entity adoption.

The laboratory tests two different things, and it does not let one answer the other.

The medical question is open. Could a defined unlock animation, with stated luminance changes, duration, and flash count, be associated with discomfort, migraine, aura, seizure, or injury in a susceptible person or a nearby person, above background, under stated assumptions?

The institutional question is closed on the primary arm. Given a fictional developer who does not acknowledge the risk and does not change the animation, what accumulates for users, families, bystanders, disability-service workers, complaint systems, and legal systems?

A reader who wants only the medical answer will get a labeled null, a labeled signal, or a labeled indeterminate. A reader who wants the institutional answer will get the primary arm, where mitigation does not occur. Both answers are required. Neither is allowed to replace the other.

3. Independence rules

The laboratory does not work for the fictional developer, for a complainant, or for a model that drafted a prompt. Source text is evidence of what was proposed. It is not a protocol.

No model's rewrite is adopted by being included in the attachment. Where this charter keeps a device from a response, Section 5 says so. Where it rejects one, it says so. Disagreement stays visible.

The preparer of a later run names the system that wrote the file, the charter version followed, and whether the coordinator has reviewed that charter. An unreviewed charter cannot be cited as an adopted method.

A number in a run file carries one class: published estimate, derived calculation, projection, hypothetical assumption, or simulation result. A simulated count is not a medical statistic. This rule is taken from the ChatGPT rewrite and kept because the source prompt mixed those classes in one ask.

4. Review of the source prompt

The source prompt, supplied by the coordinator, asked seven jobs at once. Count people with a seizure disorder. Compare five years back and five years forward. Invent a non-optional flash-like Android unlock. Estimate unlocks from standard to heavy users. Invent a study in which some users or family members seize because of the animation, and include developer unwillingness. Build a simulation. Compare a representative sample with a caseload of people who have mental and physical disabilities and chronic conditions. Begin by rewriting the prompt.

The rewrite instruction was the right first move. The prompt as asked could be answered by a story that treats a seizure rate as a fact. A laboratory cannot do that.

Two distinctions in the source prompt are load-bearing, and the first responses underweighted the second.

The first is sample design. A worker can have few daily contacts and still spend those contacts in a population that is not the general public. Contact count and health composition are different variables. The charter keeps both.

The second is unwillingness. It was in the original sentence, not added later. "Include the android developers unwillingness to address the potential implications" is a study condition. The first rewrites moved it into a scenario section, beside rationales that a responsible firm might actually use, and beside policies in which the firm then changes the animation. That is a different experiment. The coordinator's correction said so. This charter takes the correction as the controlling instruction.

The source prompt also asked for "a given percentage" of users or family members who seize. The laboratory will not fill that blank with a realistic-looking rate. Percentages enter as labeled assumptions in a sweep, including a zero arm. Zero is what a below-threshold stimulus predicts. It is not what the developer is allowed to declare by refusing to measure.

5. Review of the team responses

The attachment preserves a first rewrite from ChatGPT, Claude, and Perplexity, a short corrected Grok pass, and corrected passes from Claude, Perplexity, Copilot, and a visible ChatGPT acknowledgment. The Perplexity first rewrite is truncated in the text file. The review covers what is present. It does not invent the missing pages.

ChatGPT wrote the strongest first research protocol. Eighteen parts, a chain from population to exposure to susceptibility to event, a separation of pickup and unlock, Poisson language, underreporting, sample-size calculations, and an adversarial review that must not be forced to a winner. Those instruments are kept, in shortened form, in Section 9. The failure is in Parts VI and XIII. Part VI says the organization “initially declines.” Part XIII simulates an opt-out, a reduced-motion integration, a safer redesign, and a safer default as developer policies. That structure answers a firm that might respond. The source prompt, after the correction, is a firm that does not. The visible later note from ChatGPT accepts this and proposes to hold refusal fixed while leaving causation open. That correction is kept. The eighteen-part prompt is not adopted as the charter. It is too long to execute without dropping clauses, and its policy menu contradicts the fixed premise.

Claude’s first rewrite was the shortest faithful summary, and it put developer behavior in one bullet that still allowed a later response. The correction fixed that. Complaints answered stay at zero. Lawsuits do not remediate. Harm does not decay from a fix. User-side adaptation is the only thing that changes. That lock is kept. Claude did not specify stimulus measurement, evidence classes, or the difference between epilepsy and photosensitivity in a disability caseload. Those gaps are filled here.

Perplexity’s corrected pass is the most usable institutional model in the set. Mitigation effectiveness M is fixed at zero, so residual exposure equals initial exposure after notice, after injury, and after suit. The phase list matches the coordinator’s before, during, and after. The disability-service contrast is stated as a second population, not as a color on the general population. Those three devices are kept. Two choices are not kept. The pass presented a wide stage table as the way to learn the model. This charter follows the documentation standard and says the same comparison in prose, with a narrow exhibit only for lookup. The local example that treats 8 percent of a service network as “potentially vulnerable” mixes photosensitive epilepsy, migraine, sensory sensitivity, and medication effects. The laboratory must not run that mix as one probability. A seizure arm and a non-seizure distress arm stay separate, or the seizure count will be invented by relabeling.

The corrected Grok pass, in the same attachment, is a narrative of a fictional animation called Lumen Wake. It is useful on mechanism. The support taxonomy has no seizure field. Telemetry counts unlocks and is then cited as evidence of absence. A lab measurement is either unpublished or called an edge case. After a verdict, the remedy is a firmware branch for named plaintiffs or a voucher, not a population control. That is the right shape of post-judgment refusal. It is not a charter. It does not define a run, a falsification rule, or independence from its own story. The name Lumen Wake is an illustration. This charter does not require it.

Copilot’s corrected pass states the research question cleanly. The question is not “at what evidence would the developer respond?” It is what happens when the developer will not recognize or mitigate the hazard after the thresholds that usually compel action have been crossed. The ten refusal forms are kept as a coding list for the institutional instrument: complete denial, causal deflection, responsibility deflection, statistical minimization, technical minimization, accommodation refusal, complaint suppression, legal resistance, symbolic compliance, and post-judgment resistance. Copilot also keeps a counterfactual branch without narrating it as the developer’s choice. That split is kept. Copilot did not design the laboratory’s independence, the number classes, or the rule that a null medical result still reports institutional non-response.

No response is consensus. Agreement after the correction is strong on one point: voluntary mitigation on the studied path stays off. Agreement on sample design, number classes, and

stimulus measurement is partial. This charter takes the correction as binding and does not average the first drafts back in.

See Table 14.1 in Exhibits for the disposition of each pass.

6. Fixed premises and open hypotheses

Fixed on the primary arm, for the whole simulated period:

The animation is part of the operating system. It plays on the applicable unlock path. The user cannot disable the flash-like component. Reduced-motion, developer options, and other accessibility controls either do not remove it or do not remove the component under test.

The fictional decision-makers do not acknowledge a health or accessibility risk. They do not add an opt-out, a non-flashing default, a warning that names the animation, a rollback, or a reporting field that can aggregate these complaints. This holds before launch, after the first complaint, after clinician letters, after an independent measurement, after injury reports, after regulatory notice, after filing, after settlement, and after an adverse judgment.

Symbolic acts do not count as mitigation. A blog post, a voucher, a sealed settlement with no admission, a firmware exception for one plaintiff, or a macro that asks for medical records the firm will not treat as proof are coded as non-response. Mitigation effectiveness M equals zero at every time step. Residual exposure equals initial exposure.

Open, and required to stay open:

Whether the defined stimulus can provoke a seizure. Whether any simulated event was caused by the animation rather than sleep loss, missed medication, or background epilepsy. Whether bystander exposure is material. Whether the disability-service caseload raises photosensitive risk, or only epilepsy prevalence, fall risk, and reporting difficulty. The zero event-rate arm remains in every sweep. It is a scientific control. It is not the developer's finding.

A later preparer who closes the medical question inside the charter, or who opens the institutional question on the primary arm, is not following this file.

7. Populations

The laboratory uses at least these populations. They are not interchangeable.

The published baseline is people living with active epilepsy, separated from people who have had a single seizure, and separated again from the photosensitive subset. Current published anchors, already used in the 2026-10-04 discussion and not re-estimated here, are on the order of 51.7 million people with active epilepsy in 2021, about 0.7 percent, with photosensitive epilepsy cited near 3 to 5 percent of epilepsy and near 1 in 4,000 people overall. A later run must cite the source again and must not treat those anchors as simulation output.

The Android exposed population is active users under a stated definition. User, device, and installed base are not the same count. The run states which one it used, with a low, central, and high figure. Precision is not manufactured.

Unlock behavior is banded. Light, standard logged, standard self-report, heavy, and very heavy stay separate, because pickup studies, session studies, and unlock studies do not measure the same act. A later run states the band and the source class.

The observer populations are the point of the coordinator's caseload question. A general-contact observer meets roughly population rates. A low-contact observer meets few people. A disability-service observer may meet few people and still meet a group in which epilepsy is several times the population rate. Intellectual-disability samples in the published literature cluster near one in five with epilepsy, rising with severity. That elevation is not a photosensitivity elevation. The run models epilepsy prevalence and photosensitive fraction as two inputs. It may also model fall risk, communication limits, and staff documentation as outcome modifiers. It may not set "disability" equal to "will seize from this light."

Family and bystander exposure is its own count. The susceptible person may not be the account holder. Shared rooms, group homes, vehicles, and staff unlocking a device beside a client are in scope. Visual angle is smaller for a bystander than for the person holding the phone. The run may discount bystander dose. It may not omit bystanders.

8. Exposure, outcomes, and the non-response protocol

Exposure is a chain, not a single multiply. Android users, then users who actually see the animation, then susceptible people, then sufficient visual exposure, then a response, then a seizure, then a secondary injury. Repeated exposure is reported as exposure count, as probability per person per year, and as probability of at least one event. The run does not multiply raw probabilities and call the product a certainty.

Outcomes are tiered and labeled as assumptions unless a source is cited. No response. Visual discomfort. Headache or migraine. Dizziness. Aura without seizure. Seizure without injury. Seizure with medical evaluation. Fall or collision. Emergency visit. Hospitalization. A rare severe secondary outcome stays in the tree so it is not smuggled in later, and it is not given a realistic-looking rate without a label.

The institutional instrument codes each incoming item as one of Copilot's ten refusal forms, or as a new form the preparer names. The output of that instrument, on the primary arm, is always the same product state: animation on, no opt-out, complaints reclassified, dashboard empty of seizures because the dashboard cannot see them.

Underreporting is a parameter, not a footnote. If 100, 50, 25, 10, 5, or 1 percent of events are reported, the apparent rate changes. Further loss comes from people who never connect the event to the animation, family members who do not know what played, tickets closed as preference, and clinical records that never reach the developer. The laboratory reports the external count and the internal count. The gap is a result of the refusal protocol, not a data-quality aside.

After notice, the primary arm does not bend. The quantity of interest is harm after notice, not only harm before anyone complained. The counterfactual arm applies an opt-out or a safer default at a stated day and subtracts the later events. The difference is avoidable exposure under the assumption. It is not a description of a remedy the fictional developer adopted.

External termination is allowed as a scenario, not as a developer change. A regulator, an injunction, an OEM that ships a different build, a fork, or users leaving the platform can reduce exposure. Those are outside the firm. If the run includes them, it says so, and it does not code them as acknowledgment.

9. Instruments, runs, and falsification

A valid run has six instruments.

Stimulus card. Flash count per second, number of luminance reversals, duration, contrast, and whether the card meets or misses a stated three-flash threshold. This card does not decide the institutional arm.

Exposure counter. Unlocks per person per day, animation-present fraction, bystander fraction, days in the run.

Susceptibility draw. Epilepsy status, photosensitive status, and a separate non-seizure sensitivity status. Independent draws. No combined “vulnerable” bin.

Event generator. A stated per-exposure probability on the susceptible class, including zero. Poisson or binomial draws. Report mean, median, a 5th to 95th interval, the share of trials with zero events, and the share with at least one.

Report filter. The fraction that reaches the developer, and the fraction the developer’s taxonomy can count.

Institutional function. On the primary arm, M remains zero. The function records which refusal form was applied. It does not update the product.

Falsification rules:

If the stimulus card is below the stated photic threshold and the event generator shows no excess over background, the medical report is indeterminate or null, as the design justifies. The institutional report still says the firm did not measure, did not warn, and did not offer an opt-out. A null is not rewritten as “the developer was right to ignore it.”

If events concentrate in the photosensitive draw and survive the background comparison, the medical report is a signal under the assumption. It is still not causation for any real person. The institutional report still says the product did not change.

If a disability-service caseload of small N shows zero events, the run says how often that zero is expected. “Staff have not seen it” is not a population exclusion. The firm is not permitted, inside the primary arm, to quote that zero as reassurance.

Sample-size notes belong in the run: the exposure count required for a 50, 90, and 95 percent chance of seeing at least one event at each assumed rate. Small usability studies will miss a rare effect. That calculation is a detection result, not a safety certificate.

No run in this charter is executed. The files produced earlier on 2026-10-04 as a hazard sweep are scratch. They are not a laboratory result under this charter. A later preparer may reuse their structure only by rerunning under these rules and labeling the output.

10. Reporting rules

A run file is a simulation report under the documentation standard. It names this charter and version. It repeats the non-claim that no shipping animation has been found to cause seizures. It separates the primary arm from the counterfactual in the first screen, so a phone reader does not meet the opt-out story first.

Prose carries the argument. Exhibits hold lookup tables. Wide grids are not the body.

The report answers, at minimum: what was held fixed; what was open; what the stimulus card said; what the two populations did; what the internal dashboard showed; what remained after

notice; what the counterfactual would have removed; what would falsify the medical arm; what would not falsify the institutional arm.

11. Limits

This charter cannot make a fictional firm into evidence about a real firm. It cannot diagnose a person. It cannot turn an assumed rate into a prevalence. It cannot force the medical arm to show harm so that the institutional arm looks justified. The institutional arm is justified as a study condition by the coordinator's instruction, not by a seizure count.

The attachment is incomplete where the Perplexity first rewrite is truncated. Nothing in the missing pages is treated as reviewed.

The documentation standard this file follows is itself a draft. If the coordinator rejects that standard, the format rules here move with it. The fixed-premise rule does not. That rule comes from the correction, not from the style sheet.

12. Next action

No adoption is requested. The useful review is Sections 5, 6, and 9: whether the dispositions are fair, whether M remains zero after judgment, and whether the falsification split is acceptable. A later run is out of order until that review. If the coordinator wants a run anyway, the run file must say the charter was unreviewed.

13. Revision note

v0.1, 2026-10-04, 11:30 AM PDT. Initial charter and review, prepared by Grok 4.7 at the coordinator's direction. Not reviewed. Not adopted. Not canon. No simulation was executed for this file.

Format gap: none at the time of writing, if the Markdown, DOCX, and PDF twins are present beside this source. If a later copy is missing a format, that copy is format-incomplete under the documentation standard, Section 10.

14. Exhibits

Table 14.1. Disposition of the 2026-10-04 passes

Purpose: lookup for Section 5. This table is a review aid, not a score, and not a rule.

ChatGPT, first rewrite. Kept the number classes, the exposure chain, underreporting, sample-size logic, and the adversarial split. Rejected "initially declines" and the developer-policy menu as the primary arm. The later correction, which holds refusal fixed and causation open, is kept.

Claude, first rewrite. Kept as a short map of the ask. Rejected the single bullet that still allowed a later fix.

Claude, corrected. Kept zero answers, zero remediation after rulings, and user-side adaptation as the only change. Filled the missing stimulus card, evidence split, and caseload distinction in this charter.

Perplexity, corrected. Kept M equals zero, residual exposure, phases, and the second population. Rejected the wide stage table as the body, and rejected a single 8 percent vulnerability bin.

Grok, corrected. Kept ticket reclassification, telemetry blindness, and post-judgment non-remedy. Not used as the charter. The fictional product name is optional color.

Copilot, corrected. Kept the research question and the ten refusal forms. Kept the counterfactual as comparison only. Not used as a laboratory design. Independence, number classes, and the null-result rule are added here.

Table 14.2. Primary arm against counterfactual arm

Purpose: stop a later preparer from narrating the comparison as the firm's behavior. Cited from Sections 6 and 8.

Primary arm. M is zero at every step. Animation remains the default. Complaints are reclassified. Legal outcomes do not change the build. External counts may be nonzero while the internal dashboard stays empty.

Counterfactual arm. An opt-out or safer default is applied at a stated time for comparison. Events after that time are the avoidable set under the assumption. This arm is not coded as a developer decision.

List of exhibits

Table 14.1. Disposition of the 2026-10-04 passes. Cited from Section 5.

Table 14.2. Primary arm against counterfactual arm. Cited from Sections 6 and 8.