

THE CLAIM PROVENANCE AND DRIFT LAB

A Documentation and Audit Layer for the OAL Multi-AI Collaborative Process

Document title	The Claim Provenance and Drift Lab
Working name	Claim Provenance and Drift Lab (CPDL)
Date	20 September 2026
Version	v0.1 — Working Design Record
Prepared by	Claude (Claude Sonnet 5, Anthropic)
Prepared for	Omniversal Architect Labs / Architect Nathan Lewis II
Status	Reported. Not canon. Subject to ONESTAMP ratification.
Source basis	Original concept note; ChatGPT, Microsoft Copilot, Perplexity AI, Google Gemini, and Grok cross-analysis; OAL governance instruments (AIDR, CSAD, OAL-AICB)

Mission and Purpose

Mission

To keep the OAL multi-AI documentation record internally consistent, attributable, and honest about its own uncertainty — so that no claim is treated as settled before it has actually been checked against what the corpus already says.

Purpose

Five of the six AI collaborators who reviewed the original Control Flow concept independently converged on a domain-modeling mission: build a flow-simulation engine over some combination of vascular, neural, botanical, or engineered systems. This lab takes a different mission, consistent with Claude's standing role in the OAL workflow as Primary Audit and Documentation Layer: rather than modeling a biological or engineered matrix, it treats the OAL cross-analysis process itself as the matrix, and the claims, documents, and versions moving through that process as the medium.

Its immediate, concrete purpose is to catch exactly the kind of problem that has already occurred in this corpus without being noticed by any of the five domain-modeling responses: two separately authored documents — ChatGPT's Matrix Architecture and Control Flow Research Program v1.0 and Grok's Control Flow Through Living Matrices v1.0 — each present themselves as the terminal, unified synthesis of the same six sources, and neither supersedes the other.

Central Design Principle

A claim is only as trustworthy as its provenance. Recognition of what the corpus already says, before adding to it, is the same discipline the rest of the corpus applies to veins and blades — applied here to documents and claims instead.

Document Map

- 1. Relationship to Existing OAL Governance
- 2. Core Objects and Schema
- 3. Claim Relationship Taxonomy
- 4. Claim Status Labels
- 5. Standard Operation Loop
- 6. Worked Test Case (Described, Not Run)
- 7. Control and Governance Layer
- 8. Relationship to the Other Proposed OAL Labs
- 9. Funding and Resourcing Pathways
- 10. MVP and Build Roadmap
- 11. Success Measures
- 12. Open Questions
- 13. Version Notes

1. Relationship to Existing OAL Governance

This lab is not new governance. OAL already has the relevant rules on the books: no content becomes canon without ONESTAMP ratification; all AI contributions carry REPORTED status until Nathan confirms; the AI Disagreement Register (AIDR, OAL-AIDR-v1.0) and CSAD (OAL-CSAD-v1.0) govern cross-AI interpretation disputes; and Claude's standing role is Primary Audit and Documentation Layer.

What has been missing is not the rule set — it is an operational layer that actually applies those rules to a live, growing corpus as documents arrive, instead of relying on a human or another AI noticing a conflict by hand. This lab is best understood as instrumentation for AIDR and CSAD, and as the mechanism that keeps REPORTED-status material from quietly hardening into something treated as settled before ONESTAMP ratification actually happens.

2. Core Objects and Schema

Object	Definition
Claim	A single assertion that can be true, false, or a matter of design choice — e.g. “a matrix is medium plus embedding.” The atomic unit the lab tracks.
Document	A discrete artifact produced by one contributor at one point in time (a chat response, a PDF, a docx, a doc section).
Version	A specific state of a document or of the corpus as a whole, timestamped and attributable.
Contributor	The AI or person who authored a document (ChatGPT, Copilot, Perplexity, Gemini, Grok, Claude, Nathan).
Corpus	The full set of documents and their claims at a given time — the lab’s working matrix.
Provenance edge	A directed relationship between two claims or documents: restates, extends, redefines, contradicts, or supersedes.

Mapped onto the OAL corpus's own shared vocabulary: architecture = the contributor/document graph; medium = claims; nodes = documents; edges = provenance relationships; driving force = synthesis requests that cause a claim to propagate; control = ONESTAMP, REPORTED status, AIDR, CSAD; failure modes = drift, duplication, silent redefinition, premature canon contact.

3. Claim Relationship Taxonomy

Every provenance edge the lab records gets one of the following labels, in order of how much a claim has changed from its source:

- Identical restatement — same claim, same terms, different document.
- Paraphrase — same claim, different wording, no meaning change.
- Extension — the original claim plus new, non-conflicting content.
- Redefinition — the same term now means something different, without the change being flagged as a change.
- Contradiction — two claims that cannot both be true as stated.
- Novel — a claim with no prior instance in the corpus.

Redefinition is the highest-value category for this lab to catch, because it is the one failure mode that looks like agreement while actually being drift. “Pressure” quietly becoming “driving potential” in one document while staying “pressure” in others, with no cross-reference tying the two together, is a live example already present in the corpus.

4. Claim Status Labels

Every claim carries one status at all times, aligned to OAL's existing REPORTED/ONESTAMP rule rather than inventing a parallel one:

- Reported — the default. Asserted by a contributor; not yet checked against the rest of the corpus.
- Corroborated — independently restated or extended by more than one contributor without contradiction.
- Disputed — flagged as a contradiction and logged to AIDR for resolution rather than silently picking a side.
- Superseded — an earlier version of a claim that a later, attributed revision has replaced.
- Orphaned — appears in exactly one document and has not been picked up, restated, or challenged anywhere else.
- Ratified — reserved for whatever eventually clears ONESTAMP. Nothing in the corpus currently holds this status, including this design record.

5. Standard Operation Loop

- Ingest — a new document enters the corpus.
- Extract — pull out its individual claims as discrete units, tagged to source and contributor.
- Compare — check each extracted claim against every existing claim already in the corpus.
- Classify — assign a taxonomy label (Section 3) to each match.
- Flag — redefinitions and contradictions get surfaced explicitly; identical restatements and paraphrases get merged rather than duplicated.
- Status — assign or update each claim's status (Section 4); a flagged contradiction routes to AIDR.
- Log — record the provenance edge in the experiment ledger, so any future reviewer can trace a claim back to its origin.

This loop is intentionally the same shape as the corpus's own “Observe → Map → Model → Predict → Intervene → Measure → Update” cycle — the difference is what is being observed and intervened on.

6. Worked Test Case (Described, Not Run)

Inputs: ChatGPT's Matrix Architecture and Control Flow Research Program v1.0 and Grok's Control Flow Through Living Matrices v1.0 — both dated 20 September 2026, both claiming to unify the same six sources.

Expected loop behavior: Ingest registers both as separate documents. Extract pulls claims from each. Compare finds heavy overlap. Classify marks most shared claims as paraphrase or identical restatement, the two governing-equation forms as a likely redefinition pending reconciliation, and the fact that each

document's own source ledger lists itself as the terminal synthesis as a contradiction at the document level.

Expected output: a Disputed status on “which document is the corpus's v1.0,” routed to AIDR, with Nathan as the resolving party per existing governance — not the lab picking a winner itself.

7. Control and Governance Layer

- Nothing the lab produces becomes canon on its own; only ONESTAMP ratification moves anything to Ratified.
- Contradictions route to AIDR, not to the lab's own judgment — detection and routing, not adjudication.
- CSAD governs how a surfaced cross-AI interpretation dispute is actually resolved.
- Identity and claim are tracked separately: who said what is needed for provenance, but never becomes a scoring or ranking system for any contributor.
- This design record is itself Reported, not Ratified, and explicitly names two existing documents as Disputed rather than picking one to treat as authoritative.

8. Relationship to the Other Proposed OAL Labs

Five other lab proposals are already on the table across the corpus: a Matrix Control-Flow Lab, a State Transition Lab (Epoch Blades), an Architecture/Cross-Species Comparator, a Temporal Lineage Lab, and a Design-Transfer Workspace. This lab sits upstream of all of them as a shared utility: every document any of the five produces would pass through the same ingest/extract/compare/classify loop before being treated as settled. It has no opinion on domain content — whether a vein is a good model for a data pipe is outside its scope — and tracks only what has been said, by whom, and whether a new statement agrees, extends, redefines, or contradicts the existing record.

9. Funding and Resourcing Pathways

This is an Architect-directed OAL project rather than an externally funded research program, so “funding” here means resourcing and adoption within the existing OAL apparatus, not a grant or budget line. Named plainly, the realistic pathways are:

- Architect time — Nathan's own review and ONESTAMP-ratification cycles remain the project's actual scarce resource; the lab is designed to reduce, not add to, that load by pre-filtering what needs a human decision.
- Existing tooling — the lab runs on infrastructure already in use (Claude's Docs tools for the corpus itself, memory for durable OAL context) rather than requiring new platform spend.
- Cross-AI labor — extraction and comparison work (Section 5) can be distributed across the existing multi-AI team as a standing task, the same way drafting and cross-analysis already are.
- Phased commitment — the roadmap in Section 10 is staged specifically so resourcing can be reassessed after each phase rather than committed all at once.

If a different funding frame was intended by the request (for example, treating this as a formal grant-style proposal with external funding sources), that should be named explicitly, since none of the corpus to date describes OAL as seeking outside funding.

10. MVP and Build Roadmap

Phase 1 — Corpus intake. Register every document in the current corpus (original prompt, ChatGPT, Copilot, Perplexity, Gemini, Grok, Claude's working doc, the two v1.0 records, the Copilot docx) as objects with attribution and timestamps. No claim extraction yet.

Phase 2 — Claim extraction on a bounded set. Run extraction against a small named set of claims already known to be load-bearing: “recognition before intervention,” the matrix definition, the governing-equation form, and the designable-slices principle.

Phase 3 — Run the Section 6 test case. Diff the two v1.0 documents against each other and against the Phase 2 claim set. This is the confirmed next deliverable.

Phase 4 — Full corpus pass. Extend extraction to every claim in every document, producing the first complete provenance graph.

Phase 5 — Live ingest. Wire the loop to run automatically as new documents are added, rather than as a manual pass after the fact.

11. Success Measures

- Detection, not just description — the lab should catch what a careful human read would also catch, faster and without needing to hold the whole corpus in mind at once.
- No false consensus — a claim should never read as Corroborated because two documents used similar words while actually meaning different things.
- Routes, doesn't rule — the lab should never resolve a Disputed claim itself.
- Cheap to run again — worth re-running on every new document, not just usable as a one-time cleanup pass.

12. Open Questions

- How much of Claude's existing OAL memory (governance docs, Invocation Framework, temporal-aliasing research, the theoretical-frameworks set) should be pulled into the corpus as comparanda for Phase 1 intake.
- Whether AIDR and CSAD have a defined intake format this lab's output should match.
- Who besides Nathan is authorized to resolve a Disputed claim this lab surfaces.
- Whether “claim” should eventually extend beyond text documents to cover decisions made in chat that never got written into a formal document.

13. Version Notes

- v0.1 — 20 September 2026: first standalone design record for the Claim Provenance and Drift Lab, following the Architect's confirmation of lane (documentation/provenance) and deliverable order (design record first, diff second).
 - Next revision trigger: Phase 3's actual diff run against the two competing v1.0 documents, or a project-owner decision that changes scope.
-

END OF WORKING DESIGN RECORD