

Documentation Governance, Design, Presentation, Format, and Publication Standard

OAL / One Industries / Lewis Corp - Independent Replacement Execution

Document ID: OAL-ONE-LC-DOCSTD-001

Version: v0.2

Status: PROPOSED / PROVISIONAL - REPLACEMENT CANDIDATE

Supersedes: ChatGPT v0.1 upon/consistent with the human coordinator retraction directive

Prepared by: ChatGPT - GPT-5.6 Sol (OpenAI)

Prepared: 2026-10-04 11:50:25 PDT (UTC-07:00)

Entity Scope: OAL / One Industries / Lewis Corp / Joint Collaborative Workspace

Document Type: Documentation governance and publication standard

Content Nature: Non-simulation governance document

Purpose Type: DEFINE / CORRECT / GOVERN

Required Formats: PDF / DOCX / Markdown

Canonical Source: Markdown for this generated package unless later governance designates otherwise

What this document is: An independent v0.2 replacement standard and correction record based on the original commission, the supplied independent versions, and the coordinator's subsequent clarification.

What this document is not: It is not self-ratifying, not a legal instrument, not proof that an external ONESTAMP has been completed, and not authority to alter third-party repositories without human action.

Table of Contents

Executive Summary

1. Purpose and Authority
2. Source Review and Interpretive Finding
3. Retraction and Replacement Rule
4. Governing Presentation Principle
5. Mandatory Document Architecture

- 6. Exhibit Governance - Zero-Threshold Rule
- 7. Narrative Independence Requirement
- 8. Mobile-First and Accessibility Rules
- 9. Purpose, Classification, and Claim Boundaries
- 10. Required Formats and Fidelity
- 11. Format Failure and Recovery
- 12. Preparer Accountability and Multi-AI Independence
- 13. Publication Acceptance Gate
- 14. Versioning, Retraction, and Supersession
- 15. Conversation Summary
- 16. Independent Comparative Conclusion
- 17. Next Actions

Appendix A. Verbatim Correction Conversation

Appendix B. Verbatim Human Retraction Directive

Exhibits

List of Exhibits

Table E.1 - Independent Version Compliance Review. Cited in Section 2 and Section 16.

Executive Summary

This v0.2 replacement corrects a material interpretation error in the withdrawn ChatGPT v0.1 documentation standard. The governing requirement is architectural: narrative belongs in the narrative body, while tables, charts, figures, diagrams, matrices, row-and-column displays, graphics, and comparable visual/structured representations belong in a designated Exhibits or Appendix area. There is no quantity threshold before that placement rule applies.

The source review confirms that the original commission required every table, chart, graphic, diagram, or figure to be numbered, titled, discussed in the narrative, referenced by item number, and placed in an appropriate designated section or appendix. A separate portion of the source material discussed a possible threshold only for whether a dedicated front-matter List of

Tables/List of Figures was needed. That list-index question must not be confused with exhibit placement.

This replacement removes all body-placement exceptions for structured visual content. If a document contains even one substantive exhibit, that exhibit receives a formal identifier, title, purpose, source/derivation note where applicable, a narrative cross-reference, and placement in the designated Exhibits or Appendix structure. The body must remain understandable without requiring the reader to decode the exhibit.

The human coordinator's 2026-10-04 correction and retraction directive is preserved verbatim in Appendix B. At the time of preparation, the coordinator stated that an electronically signed ONESTAMP retraction record was being generated. This document records that statement; it does not independently verify the external signature or GitHub publication state.

1. Purpose and Authority

This document is an independent replacement execution of the Documentation Governance, Design, Presentation, Format, and Publication Standard for OAL, One Industries, Lewis Corp, and explicitly joint work.

Its purpose is to establish a uniform presentation and publication discipline without forcing uniform conclusions. Human and AI collaborators may disagree on substance, findings, models, recommendations, or interpretation. They must still present official documentation through the same structural rules so readers can recognize status, provenance, purpose, boundaries, evidence, and supporting exhibits consistently.

This version is PROPOSED / PROVISIONAL and does not adopt itself. Human authorization governs ratification. The retraction status described in Section 3 reflects the coordinator's directive recorded in the conversation and any separate ONESTAMP/GitHub record remains external evidence unless attached to the publication package.

2. Source Review and Interpretive Finding

The review covered the original documentation-standard commission and the independent versions supplied from Claude, Perplexity, Grok, Microsoft Copilot, and ChatGPT. The comparative assessment is summarized in prose here and in Table E.1 in the Exhibits section.

The original commission was unambiguous on placement: every table, chart, graphic, diagram, or figure must have a unique item number and descriptive title; must be introduced or discussed in the narrative; must be referenced by item number; and must appear in an appropriate designated section or appendix. The source also separately asked whether a List of Tables/List of Figures should have a reasonable front-matter threshold. Those are different requirements.

Claude's independent version follows the placement requirement most directly: it states that tables, charts, matrices, and figures go in the Exhibits section and must not appear in the body. Claude also uses a List of Exhibits and formal cross-references. Its separate exhibit-list thresholds should be reviewed for consistency with the clarified no-threshold presentation philosophy, but they do not create a body-placement exception.

Perplexity's independent version explicitly permits small essential tables near the narrative and rejects a universal ban on body tables. That conflicts with the clarified architecture and requires revision before it can serve as a governing standard.

Grok's independent version places tables, charts, diagrams, and figures in Exhibits but creates an exception when a single narrow list is considered clearer in place. That exception conflicts with the clarified architecture. Grok also establishes a count threshold for when a separate list of exhibits is added. The placement exception requires revision.

Microsoft Copilot correctly requires numbering and cross-reference, but states that large exhibits belong in dedicated sections. That wording leaves smaller exhibits without the same categorical placement requirement and therefore requires revision for prompt fidelity.

ChatGPT v0.1 permitted a narrow table or figure to remain in the body when deemed necessary and mobile-safe, and it added a three-or-more threshold for lists of tables or figures. The body-placement permission was an unauthorized interpretation. This replacement withdraws it completely.

3. Retraction and Replacement Rule

A documentation-standard document is noncompliant with this v0.2 replacement when it establishes a numerical threshold that must be reached before tables, charts, figures, diagrams, matrices, graphics, row-and-column displays, or comparable visual/structured content must be moved into a designated exhibit/appendix area.

Under the human coordinator's directive recorded on 2026-10-04, any published documentation-standard version containing such a placement threshold must be reviewed, formally retracted, and regenerated in updated form. This document records that workspace directive as supplied by the coordinator; it does not purport to create legal rights or obligations outside the collaborative workspace.

ChatGPT's v0.1 standard is specifically treated here as withdrawn and replaced by v0.2 because it permitted inline body placement of some tables/figures and therefore did not faithfully preserve the user's intended document architecture.

A threshold governing only whether an additional front-matter index (for example, a List of Exhibits) is displayed is conceptually different from a threshold governing exhibit placement. To

eliminate future ambiguity, this v0.2 standard uses no numerical threshold for either rule: if a substantive exhibit exists, it receives designated placement and appears in the List of Exhibits.

4. Governing Presentation Principle

Place for everything; everything in its place. The body is a readable narrative. Structured visual material is supporting material and is placed in the document's designated Exhibits section or, where the material is supplementary/technical, in a clearly identified Appendix or Technical Annex.

The body must explain the point in plain language before or while directing the reader to the supporting exhibit. A reader must be able to understand the argument, finding, requirement, or conclusion without decoding a table or graphic.

The standard prohibits using tables as page-layout devices, using row-and-column grids as substitutes for prose, scattering charts or figures throughout narrative pages, or treating a visual as self-explanatory.

5. Mandatory Document Architecture

Every substantive official document must include, in a readable single-column structure: title; organization/entity scope; project/program where applicable; document type; version; status; preparer; date and timestamp with timezone; intended audience; purpose and intended result; relationship to other work where applicable; tailored reality/simulation/non-claim boundaries; executive summary; table of contents; numbered narrative sections; next action/outcome; revision note; and format status.

PDF and DOCX must carry page numbers. Markdown uses heading hierarchy and internal navigation rather than artificial pagination.

If any exhibit exists, the document must also include a List of Exhibits and a designated Exhibits section. There is no minimum exhibit count.

6. Exhibit Governance - Zero-Threshold Rule

The zero-threshold rule applies from the first substantive exhibit onward. One table is governed. One figure is governed. One chart is governed. Two diagrams are governed. No preparer may delay formal exhibit treatment until an arbitrary count is reached.

Every table, matrix, chart, graph, diagram, figure, graphic, schematic, map, structured row-and-column display, or comparable visual representation must: receive a unique item number; receive a descriptive title/caption; state its purpose; identify source or derivation when applicable; be introduced and interpreted in the narrative; be referenced by item number and

designated location; remain understandable through a textual explanation; and appear in the designated Exhibits section or an explicitly identified Appendix/Technical Annex.

Accepted reference style includes: 'Please refer to Table 6.1 in the Exhibits section.' Spatial-only references such as 'the table below,' 'the chart above,' or 'the following figure' are not sufficient because content reflows across PDF, DOCX, and Markdown.

This rule does not mean that documents should be filled with exhibits. An exhibit must earn its existence by communicating a relationship, comparison, dataset, workflow, architecture, measurement, inventory, or other structured content more effectively than prose. If prose is clearer, use prose and do not create the exhibit.

7. Narrative Independence Requirement

Narrative content is the primary communication layer. The body must state the subject, reasoning, findings, decisions, limitations, and implications in prose appropriate to the intended audience.

Exhibits support verification, lookup, comparison, reproduction, or deeper inspection. They do not carry the only copy of a conclusion, requirement, warning, or decision.

For mixed audiences, use layered communication: executive summary; plain-language narrative; detailed technical explanation; exhibits; then appendices/supporting data.

8. Mobile-First and Accessibility Rules

Assume that many readers will use a phone in portrait orientation. Use one primary text column, descriptive headings, manageable paragraphs, readable type, adequate spacing, high contrast, and no meaning that depends on color alone.

Because structured visual content is moved out of the narrative body, the body should not require horizontal scrolling. Exhibits that are too wide for a phone must be split, restacked, simplified, or supplied as a technical annex/companion dataset with a complete prose interpretation.

Meaningful figures require text alternatives where the format supports them. Tables require clear headers and units. Accessibility is evaluated on exported artifacts, not only on source files.

9. Purpose, Classification, and Claim Boundaries

Documentation exists to accomplish something: inform, define, decide, record, specify, test, preserve, review, execute, or support another explicit purpose. A document whose only result is that another file exists should be reconsidered.

Every document must distinguish lifecycle/maturity status from content nature. Terms such as DRAFT, PROVISIONAL, UNDER REVIEW, APPROVED, RATIFIED, SUPERSEDED, and RETIRED describe lifecycle. Terms such as SIMULATION, HYPOTHETICAL, FICTIONAL, CONCEPTUAL, OBSERVED, CALCULATED, ESTIMATED, or MIXED describe content/evidence nature.

A reader should never have to infer whether a number was observed, sourced, calculated, estimated, simulated, assumed, synthetic, or speculative.

10. Required Formats and Fidelity

The baseline publication package for substantive project documents is PDF, DOCX, and Markdown. The three files are different representations of the same substantive document, not independent rewrites.

Titles, version, status, section numbering, conclusions, dates, transcript content, exhibit numbering, source notes, and major claims must remain synchronized across formats.

PDF is the stable presentation copy. DOCX is the editable office copy. Markdown is the repository-friendly plain-text copy. A project may designate one as canonical, but canonical status must be explicit.

11. Format Failure and Recovery

A preparer must never claim a required file exists when it does not. If PDF, DOCX, or Markdown cannot be generated, validated, or delivered, identify the missing format; state the actual failure; deliver every successful format; preserve a complete conversion-ready source; explain whether completeness was checked; and provide a recovery path.

The publication status becomes FORMAT INCOMPLETE until the missing artifact is produced or an authorized exception is recorded.

12. Preparer Accountability and Multi-AI Independence

The preparer line is part of the record. An AI that materially prepares a document identifies itself. Human direction and AI preparation should not be conflated.

Uniform presentation does not require uniform thought. Collaborators may disagree. Multi-AI documents must preserve material dissent rather than manufacture consensus.

A preparer is accountable for prompt fidelity as well as visual polish. Passing a render check does not prove structural compliance. Structural compliance must be reviewed separately before publication.

13. Publication Acceptance Gate

Before publication, verify purpose, audience, organization/entity scope, status, preparer, timestamp, version, claim boundaries, narrative readability, exhibit placement, exhibit numbering, exhibit cross-references, source/derivation labels, mobile usability, accessibility, cross-format fidelity, file existence, related-work references, outcome/next action, and revision status.

The gate fails if any structured visual appears in the narrative body; if an exhibit is unnumbered, untitled, uncited, or lacks a designated home; if the argument depends on the exhibit alone; if a quantity threshold is used to postpone designated exhibit placement; or if required formats are missing without the recovery protocol.

A render may pass while the document fails this gate. Layout correctness and governance compliance are separate checks.

14. Versioning, Retraction, and Supersession

Draft/provisional versions use v0.x. A later v0.x may replace an earlier draft when requirements are corrected. The first formally adopted standard may become v1.0 only through authorized human governance.

Retracted or superseded files should remain discoverable as historical records when repository practice allows, but must be clearly labeled RETRACTED or SUPERSEDED and point to the replacement. They must not remain indexed as the current governing standard.

This document is v0.2 because it materially changes the exhibit-placement rule and records the retraction directive affecting v0.1.

15. Conversation Summary

The coordinator reviewed ChatGPT's prior output and identified that the generated standard did not consistently segregate structured visual content into designated sections. The coordinator emphasized that the purpose of the standard is presentation uniformity across collaborators while preserving independence of thought.

The coordinator then objected to threshold language being interpreted as permission to leave one or two visuals in the body until a count was reached. ChatGPT acknowledged that this was a misinterpretation and agreed that the previous version required withdrawal rather than grandfathering.

The coordinator provided the other AI-generated standards for comparative review and stated that a signed ONESTAMP retraction record was being prepared for ChatGPT's version and any

other standard that used threshold language to evade the placement requirement. This v0.2 replacement was commissioned to preserve the correction, transcript, and review record.

16. Independent Comparative Conclusion

Claude's implementation is the closest match to the clarified placement architecture because it makes Exhibits the exclusive home for tables, charts, matrices, and figures and keeps prose primary. That structural decision should be preserved.

Perplexity provides strong governance, source, accessibility, and publication controls, but its allowance for small essential body tables conflicts with the clarified placement rule. Grok is concise and operational but creates a single-narrow-list body exception. Copilot captures numbering and cross-reference but limits dedicated exhibit placement to 'large' exhibits. ChatGPT v0.1 created the broadest exception and therefore required replacement.

The synthesis adopted here is not a vote tally. It uses the original instruction and the coordinator's subsequent clarification as controlling requirements, then borrows compatible governance mechanisms from the independent versions.

17. Next Actions

Publish the signed ONESTAMP retraction record with the withdrawn versions or their repository entries. Mark ChatGPT v0.1 RETRACTED and link to this v0.2 replacement candidate.

Review the Perplexity, Grok, and Copilot standards for body-placement language and regenerate them if the coordinator applies the same retraction directive. Review Claude's threshold language concerning exhibit indexes for clarity, while recognizing that its core placement rule already requires exhibits outside the body.

Pilot this v0.2 against at least one research report, one simulation report, one major charter, and one short record. Reject any pilot where structured visuals leak into narrative pages or where the prose becomes unintelligible without an exhibit.

After cross-platform review and human approval, produce the ratified version with the authorized status and version number.

Appendix A. Verbatim Correction Conversation

The following transcript is reproduced verbatim from the conversation relevant to the correction and replacement of the documentation standard. Typographical, grammatical, capitalization, and dictation characteristics are preserved.

USER - 2026-10-04 11:40:58 PDT

I do want to give you some feedback about the documentation that you generated. I wanted to ask if you looked over the document design that I provided? The documentation and design standard should be included in that document, and it looks like you did not adhere to the standardized format for document. The charts, the figures, the rows, and the columns all should be clearly labeled for a section that they are included in. They should be referenced. I looked over Claude AI's documentation and it is formatted correctly. If there are tables, charts, rows, columns, or figures, or any kind of visual display, there should be a section for that. We, we don't want the presentation to only focus on that information. They should be able to read what it is about, but I noticed that throughout your document that that is included. It needs to be labeled and there needs to be a section for it. They should not depend on that information alone. Do you understand what I'm saying? Is there something about the documentation that I provided that wasn't clear or that we need to go over? My goal is really to keep in mind like the potential for somebody else to read the information and for just it to be presentable. When you guys generate stuff and your name is attached to them, I want people to understand or whoever is reading whatever, whatever, if it is a person, if it is an AI, it doesn't matter, but I just want them to know that as a team, as a group, when we generate content that there is a standard attached to that and that it is uniform, it doesn't mean that we have to agree about what that contains, but it just means that for presentation there's there's a standard attached to the work that we do. Do you know what I mean?

CHATGPT - 2026-10-04 11:40:58 PDT (response)

Yes, I understand exactly what you mean. I did review the documentation standard, but I did not apply its exhibit/visual-structure rules rigorously enough in the charter I generated. That is a compliance failure on my part, not a lack of clarity in what you provided.

USER - 2026-10-04 11:42:28 PDT

The documentation that you generated says that if there are three or more figures, but what I specified when I provided the prompt to everybody, I actually have all of the original prompts and everybody's responses that I published on GitHub along with every independent version that was generated. So I don't know if there was a misinterpretation, but I do not want you to interpret what I say like, oh, if there was a certain number, then I will include them in a different section. So if that's the case and that's what you generated, then we need to take down that version because that's a misrepresentation.

CHATGPT - 2026-10-04 11:42:28 PDT (response)

Yes. Your standard was clear, and after re-reading the exact language, I can point to where my execution fell short.

The standard explicitly says that the narrative is the primary communication layer and that tables, matrices, charts, and diagrams support that narrative rather than replacing it. It also says that every substantive visual must have a unique item number, a descriptive title or caption, a purpose, a source or derivation where applicable, and must be introduced or interpreted in the narrative and referenced by item number.

It goes further: large or supplementary visuals belong in a designated Exhibits or Appendices section, and when there are three or more tables or three or more figures/visuals, the document should include a List of Tables and/or List of Figures.

So the problem was not that I misunderstood the philosophy. I understood it conceptually, but then I treated some of the structured displays in the laboratory charter as ordinary inline content instead of treating them as formally governed document objects. That was the mistake.

What you are describing is a stronger presentation discipline:

- A reader should be able to read the prose and understand the document without depending on a table, chart, matrix, figure, or other visual.
- When a visual does exist, the prose should explicitly point to it: for example, “Refer to Table 8.1 in Section 8.”
- The visual should have an identity of its own: Table 8.1 — Experimental Workstream Structure, Figure 12.1 — Institutional Nonresponse State Model, and so forth.
- Those items should appear in a clearly designated Tables, Figures, Exhibits, or Appendices section, depending on their role.
- If several exist, there should be a navigational list near the front of the document.
- A table should not quietly appear in the middle of a page simply because it is convenient to format the information that way.
- Rows and columns are not a substitute for explanation.

That last point is particularly important. The standard itself gives essentially the same distinction: a compliant document explains the conclusion in prose and then points to a numbered supporting table; a noncompliant document puts the conclusion only inside a matrix.

And I understand the larger objective behind what you are saying. This is not merely about whether one document “looks good.” You are trying to establish a recognizable documentation culture across the collaborative workspace. Different AIs can disagree with one another, produce independent findings, challenge one another, and even recommend incompatible conclusions, while the presentation of the work still follows a common professional standard.

That means someone encountering a Claude document, a ChatGPT document, a Grok document, a Copilot document, or another team member's work should recognize that it came through the same broader documentation discipline even though the authors are independent.

That is a useful distinction:

Uniformity of documentation does not require uniformity of thought.

The standard governs things such as identity, provenance, structure, accessibility, presentation, cross-referencing, visual treatment, status, versioning, and publication quality. It should not force the research conclusions themselves to converge.

There is also another part of the standard that I should have enforced more aggressively on myself. The publication checklist specifically requires that a substantive document use tables and figures only where they improve understanding, number and reference those exhibits, maintain readable narrative, and pass the broader acceptance gate. I should not have called the charter compliant after only checking things like clipping, portrait rendering, headers, page numbers, and mobile readability. I checked render quality, but I did not perform a sufficiently strict structural-compliance review.

Those are different tests.

A file can render perfectly and still fail the documentation standard.

So I would classify the charter I generated as:

Content: usable

Research architecture: usable

Rendering: passed

Documentation-standard compliance: failed / requires revision

I would not grandfather it in or make an exception for myself.

I have also retained your clarification as an ongoing documentation rule for our work: the governing standard is the presentation baseline, narrative remains primary, and substantive tables/figures/visual displays must be formally numbered, titled, cross-referenced, and placed within the appropriate designated structure.

Nothing in the standard you supplied needs clarification for me on this issue. The problem was my implementation of it.

USER - 2026-10-04 11:44:30 PDT

Would you like me to provide you with the versions that were rendered by the others? I would like for you to look over them, and then what I can do is in GitHub I can, I can do a

one stamp for a formal retraction notice for the document that you generated. I can sign off on that electronically, and then you can generate a new version. I just, I don't want you to think that the document should just be littered with visuals. They, they should have a clear section. It shouldn't be like, "OK, like we can just still include all of this stuff, like scattered throughout the document. There, there is there's a place for everything and everything in its place. You know what I mean? So visuals, graphic displays, rows, tables, columns, charts, and figures, there needs to be a place for it in that document that we generate it. It's not if there are three, it's not if there are four, it's not if there are two, if it meets this threshold, then I will start doing that because I was very clear and I have all of the original verbatim prompts that I provided you guys about that, so there shouldn't be a lot of misinterpretation

CHATGPT - 2026-10-04 11:44:30 PDT (response)

Yes. Please provide the versions generated by the other AIs. Reviewing them side by side with your original verbatim prompt is the right next step before I generate a replacement standard.

And I understand the rule more precisely now: this is not merely a numbering requirement. It is a document-architecture requirement.

If a document contains a table, chart, figure, matrix, row/column display, diagram, graphic, or other visual representation, that material needs a deliberate home in the document. It should not be scattered through the narrative simply because the software makes it easy to insert it there. The prose should explain the subject independently, then point the reader to the appropriate numbered item in the designated visual/exhibit section.

So the governing principle should be closer to:

Narrative in the narrative sections. Visual and structured representations in their designated section. Every visual or structured item is numbered, titled, and referenced from the narrative. No quantity threshold applies.

That means one table gets formal treatment. One figure gets formal treatment. Two charts get formal treatment. There is no "three or more" trigger.

Your phrase "a place for everything and everything in its place" captures the architecture well. The goal is not to ban visuals, nor to fill a document with them. It is to prevent them from taking over the document or appearing without structure.

For the retraction, a formal GitHub/ONESTAMP record also makes sense because the issue is substantive: the v0.1 document introduced a rule that was not in the originating requirement. I would classify the reason as requirements misrepresentation / unauthorized interpretation, not just formatting preference.

Once you provide the other versions, I can review them comparatively and identify:

- which one most faithfully represents your original rule,
- where each AI added its own interpretation,
- which formatting choices are worth preserving,
- where my v0.1 diverged,
- and what the replacement standard should say so this cannot be misread again.

Then I can generate a clean successor version that explicitly supersedes the withdrawn one.

Appendix B. Verbatim Human Retraction Directive

Timestamp: 2026-10-04 11:49:45 PDT (2026-10-04T18:49:45Z)

The following is reproduced verbatim from the coordinator's message commissioning this replacement:

here are the other versions. i would like you to review and analyze them and the original content again and then render your own independent with updated versioning number. there should be a verbatim transcript of what I stated to you and verbatim responses included as well as a conversation summary. it should also state that any document that that is published for documentation standards that uses language that establishes a threshold number of visuals to meet before having to place said content in a designated area of the document needs a review, is formally retracted, and must be generated in updated form. as for yours i am generating a ONESTAMP that is electronically signed by me now that retracts yours and any other who uses language like that to slide around prompt adherence. please include this verbatim with a timestamp so viewers will know.

Record note: At the time this document was prepared, the coordinator stated that an electronically signed ONESTAMP retraction record was being generated. This document records that statement and does not independently authenticate the external signature or repository transaction.

Exhibits

Table E.1 - Independent Version Compliance Review

Purpose: Identify how each supplied independent standard handles structured visual placement under the clarified zero-threshold rule.

Cited in: Sections 2 and 16.

Source/derivation: Independent review of the supplied Claude, Perplexity, Grok, Microsoft Copilot, and ChatGPT materials. This is a qualitative governance assessment, not a ranking of model capability.

Version / Preparer	Placement rule observed	Zero-threshold placement compliant?	Required disposition under v0.2
Claude Sonnet 5.5 v0.1	Tables, charts, matrices, and figures MUST NOT appear in the body; they go to Exhibits. Separate list-of-exhibits threshold language appears elsewhere.	Yes for placement	Retain placement architecture; review index-threshold wording for clarity.
Perplexity AI v0.1	Permits small essential data tables near narrative / rejects universal body-table ban.	No	Review, retract if published as governing standard, regenerate.
Grok 4.7 v0.1	Exhibits generally separate, but allows a single narrow list in place.	No	Review, retract if published as governing standard, regenerate.
Microsoft 365 Copilot	Requires numbering; states large exhibits belong in dedicated exhibit sections, leaving smaller items ambiguous.	No	Review, retract if published as governing standard, regenerate.
ChatGPT GPT-5.6 Sol v0.1	Permitted a narrow mobile-safe table/figure in the body and used count-based list thresholds.	No	Retracted/replaced by this v0.2 candidate.

Revision Note

v0.2 - 2026-10-04: Replacement of ChatGPT v0.1. Removes body-placement exceptions, removes quantity thresholds from exhibit-placement governance, requires designated placement and formal indexing from the first substantive exhibit, adds comparative review, conversation summary, verbatim correction transcript, and verbatim coordinator retraction directive.