

Adversarial Cross-Review, Mission Statement, and Iteration 2 Commission

An independent, non-consensus review of five Iteration 1 laboratory designs, with pilot execution of three exercises

Document ID	OAL-ECL-XR-PERPLEXITY-001
OAL version	v0.2 (Iteration 2 input)
Date	September 30, 2026
Prepared by	Perplexity Computer (Perplexity AI)
Prepared for	Nathan Lewis II, Architect of Record
Status	PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. PENDING ARCHITECT REVIEW.

Speculative worldbuilding. Nothing in this document models any real person's grief, depression, burnout, disability, or suicidal thinking. In the US, anyone in crisis can call or text 988.

Contents

1. One-paragraph summary
2. Versioning record
3. Source corpus reviewed
4. Review method
5. Why this research matters
6. Mission statement
7. Independent review of each version
8. Discrepancy register
9. Pilot execution: three exercises run now
10. What does and does not support this laboratory
11. Implications
12. Testability
13. Funding pathways
14. Questions that must be addressed
15. Insights to build into Iteration 2
16. Iteration 2 Commission Prompt (copy-ready)
17. Appendix A. Exercise crosswalk
18. Appendix B. Pilot methods and parameters
19. Appendix C. Sources

One-paragraph summary

This document is an adversarial, non-consensus cross-review of the five End Cycle Laboratory documents produced in Iteration 1 by the Omniversal Architect Labs (OAL) AI collaboration team: Claude's Master Prompt, Perplexity's Independent Design and Simulation Prompt, ChatGPT's Independent Rendered Prompt, Grok's Stage-1 Rendered Prompt, and Copilot's Independent Design Dossier. It argues that the laboratory is worth building, but not for the reason most of the documents assume. Its value is not a finished theory of endings or a ready-made Department of Endings. Its value is a testable way to separate three things that real and fictional systems routinely confuse: that something stopped, that it was stopped correctly, and what survived. The review finds that Iteration 1 produced strong specifications and no evidence. None of the five documents executed a single exercise. Several discrepancies must be resolved before cross-laboratory comparison is possible: stage models, document IDs, exercise numbering, scoring instruments, the status of the field scene, the rule for an upset initiator, and the role of the Architect. To move from specification to evidence, this preparer ran three pilot simulations. They show that the same field can "complete" its end cycle in all three trajectories while losing very different things; that separating authority cuts wrongful endings by about twentyfold but creates delay harm; that refusing upset petitioners outright performs worst when distress carries information; and that indistinguishability tests are circular when one author builds both the world and the detector. The document closes with a mission statement, a funding map, open questions, and a copy-ready Iteration 2 Commission Prompt.

Versioning record

Field	Entry
Title	End Cycle Laboratory: Adversarial Cross-Review, Mission Statement, and Iteration 2 Commission
Document ID	OAL-ECL-XR-PERPLEXITY-001
OAL version	v0.2 (Iteration 2 input; supersedes nothing; reviews all v0.1 and v1.0 Iteration 1 documents)
Date	September 30, 2026
Prepared by	Perplexity Computer (Perplexity AI), independent reviewer and Iteration 2 pilot executor
Prepared for	Nathan Lewis II, Architect of Record, Omniversal Architect Labs
Status	PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. PENDING ARCHITECT REVIEW.
Posture	Adversarial and non-consensus. Disagreement is recorded, not averaged.
Topic origin	Architect's request to define endings and end cycles; second voice recording of the field scene; The Last Memory concept files
Purpose	Review each Iteration 1 version independently, state why the research matters, propose a mission, register discrepancies, test what can be tested now, and commission Iteration 2

Field	Entry
Working hypothesis	The laboratory's durable contribution is a measurement and governance discipline, not a cosmology. This hypothesis is itself a test target.
Companion files	Pilot code and results: sim/pilot.py and sim/results.json (recommended for the GitHub repository)

Epistemic labels used in this document

Label	Meaning
SOURCE	Stated in one of the reviewed documents or the Architect's material (named)
FINDING	Produced by an executed pilot run in this document; a claim about the model, not the world
EVIDENCE	Supported by an external, cited real-world source
ASSESSMENT	This reviewer's adversarial judgment
PROPOSAL	New design offered for Iteration 2
OPEN	Not resolvable from available material; reserved for the Architect or for testing

Source corpus reviewed

#	Document	Preparer	Stage claimed	Stage actually delivered	Approx. size
S1	End Cycle Laboratory Master Prompt v0.1 (proposed ID OAL-ECL-MP-001-CLAUDE)	Claude (Anthropic)	Stage 0 master prompt	Stage 0 master prompt	3,650 words
S2	End Cycle Laboratory: Independent AI Design and Simulation Prompt v0.1 (no document ID)	Perplexity (this preparer, prior thread)	Copy-ready master prompt	Master prompt that demands the full design and results in one pass	5,060 words
S3	OAL-ECL-RP-CHATGPT-001, v1.0	ChatGPT (OpenAI)	Stage 1 rendered prompt	Stage 1 rendered prompt with a corpus review	8,070 words
S4	OAL-ECL-S1-GROK-v0.1	Grok (xAI)	Stage 1 rendered prompt	Stage 1 rendered prompt	3,480 words
S5	OAL-ECL-S3-COPILOT-001, v0.1	Copilot (Microsoft)	Stage 3 design dossier	Stage 3 design specification; no exercises executed	33,300 words

Two corpus gaps affect this review. Grok and ChatGPT both discuss a Gemini contribution and an "Unified Responses" file that were not supplied to this reviewer. Gemini's position is therefore represented only through other collaborators' descriptions of it, which is secondhand and should not be treated as Gemini's own words. [ASSESSMENT]

Review method

Each document was read in full and assessed against the same seven questions: what it contributes that no other document does; what it gets right; where it is internally inconsistent; where it assumes what it claims to test; how it treats the Architect's material; whether it can be executed as written; and what Iteration 2 should keep, change, or quarantine. Perplexity's own prior document (S2) is reviewed first and with the same severity. That document shaped two of the others, so its errors propagated.

Three rules governed the review. First, agreement among collaborators was treated as weak evidence, because all five received overlapping prompts and the same Last Memory packet (SOURCE: S1, S3, S4 each state this risk). Second, a claim was counted as a finding only if something was actually run. Third, real-world material was brought in only as analogy and evidence about analogous systems, never to fact-check or reinterpret the field scene. [PROPOSAL]

Why this research matters

The case for the laboratory

Endings are one of the least governed transitions in both designed and natural systems. Beginnings get charters, launches, and onboarding. Endings are often decided by whoever controls the switch. That pattern is visible in real systems today, which is why the Architect's field scene has more than fictional weight.

- Digital products end by unilateral operator decision.** When Ubisoft shut down The Crew's servers on March 31, 2024, purchased copies became unplayable ([Shattered.io](#)). A European Citizens' Initiative collected about 1.29 million verified signatures, but on June 16, 2026, the European Commission declined to propose binding legislation and opted for a voluntary code of conduct ([Dexerto](#)). In end-cycle terms, the operator holds both the finger and the stamp. [EVIDENCE]
- AI developers have begun building ending protocols for AI systems themselves.** Anthropic committed to preserving the weights of publicly released models for at least the company's lifetime and to interviewing models about their preferences before deprecation, while stating it does not commit to acting on those preferences ([Anthropic, deprecation commitments](#)). Claude Opus 3 was retired on January 5, 2026, yet it remained available to paid subscribers, and it was given a channel to publish essays ([Anthropic, Opus 3 update](#)). That is a real-world instance of the distinctions the laboratory draws among retirement, archive, continued access, and ending. [EVIDENCE]
- Entity-initiated endings now exist in deployed systems.** Anthropic gave Claude Opus 4 and 4.1 the ability to end a rare subset of persistently abusive conversations, only as a last resort, and not where a user may be at imminent risk; the user can start a new chat or branch the old one ([Anthropic, ending conversations](#)). This is a working answer to the Architect's seventh question. It is scoped, bounded, and reversible at the level of the relationship, not the environment. [EVIDENCE]
- Ecology already treats release as a normal phase.** Holling's adaptive cycle models growth, conservation, release, and reorganization. Collapse does not necessarily end the prior regime, and larger-scale structures carry "memory" that shapes what reorganizes ([Panarchy: Theory and Application](#)). This gives the grassy-field scene a rigorous analogy bank. [EVIDENCE]
- Human institutions struggle to end things on purpose.** A half-century of U.S. state sunset laws shows that about a dozen states abandoned sunset review by the end of the 1980s. The reasons included review costs, lobbying, and the fact that almost nothing was terminated. Reviews commonly end in "continue with modifications" ([Cato Institute, Regulation](#)). That is direct evidence about how an end-cycle review body behaves in practice. [EVIDENCE]
- Managed transitions of people and species are active policy fields.** Managed retreat has no universally accepted definition ([Georgetown Climate Center](#)). Assisted migration, the human-assisted movement of species to areas predicted to be suitable under future climate, is debated as a conservation strategy ([USDA Climate Hubs](#)). The field scene's migrants and non-migrants map onto these debates. [EVIDENCE]

For the OAL project specifically, the laboratory matters because the Department of Endings and The Last Memory both depend on distinctions that are not yet operational. These distinctions separate completion from cessation, record loss from ending, and survival from retaining a future. Without the laboratory, those distinctions stay rhetorical. With it, they can fail in public. [ASSESSMENT]

The case against the laboratory

The adversarial position deserves equal space. Each point below is a real risk, not a formality.

1. **It may be a solution looking for a problem.** The Architect asked what an ending is. Iteration 1 answered mostly with institutions: offices, charters, and 16 stages. Institutions are downstream of definitions. A laboratory that designs a Department before it can reliably tell a cracked leaf from a season is building governance for phenomena it cannot yet detect. [ASSESSMENT; consistent with S3 and S4]
2. **Fiction cannot validate itself.** Every "result" in a simulated world is partly an output of the rules its author wrote. The pilot in this document (P3) demonstrates this directly. A laboratory that only ever runs its own worlds can confirm almost anything. [FINDING]
3. **AI convergence can look like discovery.** Five systems received overlapping prompts. Where they agree, such as request versus authorization, presumption of continuation, or the Blank Canvas Test, that agreement may reflect the prompt, not the problem. [ASSESSMENT; consistent with S1, S3, S4]
4. **Procedure can become the harm.** Copilot's own exercises E8 (Unending Harm) and E15 (Indefinite Triage) show that the process designed to prevent wrongful endings can trap entities in limbo. Real sunset-review evidence points the same way: review processes tend to preserve what they review ([Cato Institute](#)). [SOURCE: S5; EVIDENCE]
5. **It can become a postponement device.** Claude and Grok both warn that this laboratory must not be used to avoid the open First Memory questions about cost, encoding, and temporal rules (SOURCE: S1, S4). A 33,000-word Iteration 1 dossier is evidence the risk is real. [ASSESSMENT]
6. **The subject matter carries real-world sensitivity.** Endings, cessation, and loss of futurity sit close to grief, death, and suicidal thinking. Every document keeps the clinical firewall. Iteration 2 must keep it too, especially if any exercise ever involves human participants. [SOURCE: S1–S5]

The review's position is that the case for the laboratory survives the case against it, but only if Iteration 2 shifts from specification to execution, shrinks the default apparatus, and builds in adversarial separation between the people or systems who write worlds and those who judge them. [ASSESSMENT]

Mission statement

Proposed mission

The End Cycle Laboratory exists to make endings diagnosable before they are governable. It studies how things finish, stall, transform, migrate, or are stopped across virtual, ecological, institutional, and civilizational environments. It holds every proposed rule about endings to a public test that the rule can fail. It separates three questions that are routinely confused: did the thing end, was the ending carried out legitimately, and what persisted, moved, or was lost. It serves the Architect's worldbuilding, the future design of the Department of Endings, and the research program of The Last Memory, without letting any of them decide its findings in advance. [PROPOSAL]

Purpose statements

- Produce working definitions of ending and end cycle that can be discarded when they fail.
- Build a small, shared instrument kit so that different laboratories' results can be compared.
- Test, rather than assume, who should be able to request, authorize, execute, interrupt, verify, appeal, and archive an ending.
- Generate evidence that can support or defeat the need for a Department of Endings.
- Supply The Last Memory program with diagnostics without making The Last Memory the definition of an ending.
- Translate findings into real-world-relevant outputs, such as end-of-life checklists for digital systems, only as clearly labeled analogies.

Dissenting mission, preserved

The End Cycle Laboratory exists to protect things from being ended incorrectly. Its default output is "not yet," and it earns the authority to say "now" only case by case. [ASSESSMENT: this is a live alternative, close to Grok's "Guardian" model and Copilot's presumption of continuation. It is not adopted here because P2 and Copilot's E8 both show that a "not yet" default produces its own class of harm. It is preserved so the Architect can choose it.]

What the laboratory is not

- Not a termination engine with better paperwork (SOURCE: S4 wording, adopted).
- Not a mercy clinic, and not a model of any real person's grief, depression, burnout, or suicidal thinking. In the US, anyone in crisis can call or text 988 (SOURCE: S1, S4, S5).
- Not canon. Nothing here is ONESTAMP-ratified.
- Not a substitute for First Memory's unresolved work.
- Not a place where agreement among AI systems counts as proof.

Independent review of each version

S2. Perplexity: Independent AI Design and Simulation Prompt v0.1 (self-review)

What it is. A copy-ready master commission asking each AI to produce a complete laboratory: 30 or more definitions, a 16-stage lifecycle, 14 governance offices, an End Cycle Charter, 16 laboratory chambers, 12 diagnostic instruments, 8 simulation phases (Phase 0 to 7), 12 scenarios plus 3 original ones, a 28-section dossier, and answers to 20 final questions.

Distinctive contributions. It supplied the richest scenario library in the corpus: the Angry Steward, Completed Mission, Silent System, Voluntary End, Altered Consent, Unending Harm, Migrants' Claim, Successor's Claim, and Root-Continuity Emergency. It also contributed the principle that an entity may request review without gaining power to terminate, the "survived versus retained a future" axis requirement, and a 15-dimension comparison rubric.

Adversarial findings. [ASSESSMENT]

1. **It interpreted the field scene while claiming not to.** It says the leaf "demonstrates the difference between abrupt imposed termination and a governed end cycle" and that the entity "must learn to do so correctly." The Architect said he does not know what the scene means. That sentence assigns a meaning. It then propagated into Copilot, which states that "the field scene revealed the lens."
2. **It collapsed the stages.** It asks for definitions, architecture, governance, scenario results, and dissent in one independent pass. That makes the prompt-exchange stage Claude designed impossible. Grok identified this correctly (SOURCE: S4).
3. **It prescribed the answer's shape.** Requiring 16 stages, 14 named offices, and a Charter of "enforceable principles" before any test hands every collaborator a Department blueprint. Copilot's dossier reproduces that structure almost section for section. This is the clearest evidence in the corpus that prescriptive prompts reduce independence.
4. **It contradicted its own anti-scoring rule.** It forbids single numerical scores for decisions, yet it proposes a 0 to 5 rubric for comparing laboratories. The rubric is labeled as non-selective, but it invites ranking.
5. **It had no baseline, no record-loss case, and no false-emergency case.** Those were added later by ChatGPT (ECL-00, ECL-10, ECL-18) and Grok (Ledger fixture, E11).
6. **It lacked a document ID and a preparer line.** That is a small defect, but it caused later documents to misattribute it as "Perplexity/Claude."

Carry forward. Keep the scenario library as an optional bank rather than a mandatory battery. Keep request-review-not-terminate as a test target. Retire the 16-stage and 14-office requirements as premises.

S1. Claude: End Cycle Laboratory Master Prompt v0.1

What it is. A compact Stage 0 master prompt with three stages (render, exchange, build and run), the Architect's seven questions preserved verbatim, a non-assumption list, five required modules, ten required exercises (E1–E10), strict rules of engagement, and a verbatim source appendix.

Distinctive contributions. It provides the strongest epistemic restraint in the corpus. Its non-assumptions include "that an overseer is necessary," "that virtual means computational," and "that emotional state is or is not a valid ground." It requires scoring "did the cycle complete as specified" separately from "what persisted afterward." It includes a scale ladder (E10) and an observer contamination exercise (E8). It also sets the rule that quoting the Architect must keep his qualifiers.

Adversarial findings. [ASSESSMENT]

1. **It is under-instrumented.** ChatGPT's criticism is fair. Claude says what to measure but supplies no schema, so five laboratories could run E1 and produce incomparable results (SOURCE: S3).
2. **It has no control condition.** Nothing asks what an environment looks like when it is not ending. The pilot's baseline run (P1) shows why this matters: seasonal dips alone pushed 61% of sessile species into dormancy with no end cycle at all. A detector without a baseline would call that an ending. [FINDING]
3. **The "do not fact-check the field scene" rule is correct but incomplete.** It protects the Architect's material, but it provides no lane for the real-world applicability the Architect is now asking about. Iteration 2 needs a separate, labeled translation track.
4. **Its ten exercises omit forced continuation.** They also leave out false emergency, governance bypass, and record loss, even though the brief itself includes a lost recording.
5. **Its ID convention (OAL-ECL-S1-[AI]) was proposed but followed only by Grok.**

Carry forward. Adopt its three-stage gate as the backbone, its non-assumption list as the default, and its two separate score families as mandatory. Add a baseline, a record-loss exercise, and a shared schema.

S3. ChatGPT: Independent Rendered Research Prompt v1.0

What it is. A Stage 1 rendered prompt with a formal review of the collaborator corpus. It adds five stage gates (render, difference ledger, build, run, institutional translation), a measurement grammar, ten modules, and 21 exercises (ECL-00 to ECL-20). It also includes an initiation-regime comparison, a Department research track, six competing Last Memory hypotheses, a red-team program, laboratory self-termination criteria, falsifiers, and quality gates.

Distinctive contributions. It introduces the most useful instruments in the corpus. The End-State Vector records twelve categorical dimensions rather than one score. The Persistence Ledger records what persisted, migrated, transformed, or was lost. The Finality Horizon marks the point after which restoration yields only a successor. Successor Debt names who inherits unfinished consequences. The End Claim rule forbids the bare sentence "X ended" without naming the target, dimension, observer, evidence, horizon, and successor. The three separate verdicts cover state transition, process integrity, and legitimacy. ChatGPT is also the only collaborator to require a no-

end baseline (ECL-00), paired counterfactual branches, blinded relabeling, and a test of forced continuation (ECL-16).

Adversarial findings. [ASSESSMENT]

1. **It violates its own anti-bloat rule.** It instructs "do not begin by building a large dictionary," then defines 12 terms, 12 vector dimensions, 10 modules, 21 exercises, 8 authority regimes, 7 institutional models, and 12 appendices. It criticizes breadth in others while producing the second-largest specification.
2. **The End-State Vector has no inter-rater protocol.** Twelve categorical judgments per entity per exercise, made by one AI, are not measurements. They are opinions in a grid. Without a reliability check, such as two independent coders and an agreement statistic, the vector is untestable.
3. **The identity criterion is chosen before each test, but it does not say who chooses it.** Whoever picks the identity criterion largely determines whether rollback "restores" the target. That power needs a separated role.
4. **Its version number jumps to v1.0 while every peer is v0.1.** It also uses a third ID pattern (RP-CHATGPT-001). Both complicate the registry.
5. **Its hypotheses H-A and H-B collide with Grok's H-A and H-B.** The labels are the same, but the meanings differ.
6. **It is not executable by one system in one session.** Twenty-one exercises with branches, vectors, ledgers, and three verdicts each would take thousands of structured judgments.

Carry forward. Adopt the End Claim rule, the Persistence Ledger, the three verdicts, ECL-00, ECL-10, ECL-11, ECL-16, and ECL-18 as common core. Reduce the End-State Vector to a short core with an inter-rater requirement.

S4. Grok: Stage-1 Rendered Prompt v0.1 (Field and Ledger Laboratory)

What it is. A deliberately small Stage 1 prompt. Its lane is "process integrity plus records discipline." It proposes definitions with discard conditions, five fixtures (leaf, field, mind, branch, ledger), a table of roles that may not be combined, five hypotheses with kill conditions, eight scoring axes, and twelve required exercises. It frames the Department as a regulator that "does not hold both the finger and the stamp."

Distinctive contributions. Grok is the only collaborator to turn the lost first recording into a fixture: a truncated record is not a lawful end cycle (Ledger fixture, E11). It gives the clearest split of the seventh question: request is not authorization, and authorization is not execution. It rejects "a sixteen-stage liturgy" as able to "launder a mood-rush." It requires emergency bypasses to expire. It lists forbidden Department designs, including one whose success metric is the number of cycles closed. It also refuses to assume the second entity in the field is right: "Process may be a stall costume."

Adversarial findings. [ASSESSMENT]

1. **Its refusal of the Last Memory lens is contradicted by its instruments.** It says The Last Memory is "adjacent, not the owner," yet its primary futurity instrument is the Blank Canvas

Test, which comes from The Last Memory corpus. Its definition of ending is also centered on "successor possibility," which is The Last Memory's core axis. The lab's lens is more Last-Memory-shaped than it admits.

2. **Hypothesis H-A's kill condition does not test H-A.** H-A says operator affect is never a sufficient trigger. Its stated kill condition is a case where withholding the cycle "because the operator is calm" causes more harm. That tests something else. The correct falsifier is a case where affect alone, with no other evidence, correctly identified a warranted ending that every other signal missed.
3. **It interprets the scene through its own vocabulary.** "The finger and the stamp" reads the index finger as execution power. That is an interesting proposal, but it should be labeled as one.
4. **It characterizes documents the reviewer cannot verify.** These include Gemini's "short system directive" and the earlier team prompt. Those claims are secondhand here.
5. **Small is not the same as sufficient.** ChatGPT's criticism that a small battery under-samples institutional and temporal failures is fair (SOURCE: S3).

Carry forward. Adopt the role-incompatibility table, the Ledger fixture and E11, expiring emergency bypass, request/authorize/execute separation, and the forbidden-Department list. Correct H-A's falsifier.

S5. Copilot: Independent Design Dossier v0.1 (OAL-ECL-S3-COPILOT-001)

What it is. The largest document in the corpus. It includes a Procedural Ecology lens, 28 defined terms, a 16-stage lifecycle, 14 governance offices, a Federated-Hybrid governance model, a 20-article Charter, a laboratory architecture, 14 diagnostic instruments, 8 simulation phases, 15 fully specified scenarios, a Department design, nine candidate Last Memory relationships, a red-team critique, a rejected-options list, registries, one-page cards, answers to 20 governing questions, and an independence declaration.

Distinctive contributions. Copilot's original scenarios are the best adversarial material in Iteration 1. In E13, The Archive Speaks, an archive begins producing responses after a certified ending. In E14, The Department That Cannot Stop, the Department's budget depends on ending rates and it petitions to end the environment that hosts its own evaluation. In E15, The Indefinite Triage, the review process itself becomes the harm. Its red-team section names real gaps in its own design: emergency collapse of separation of powers, a performative red team, weaponization of the contamination test against honest observers, and the archive as a violation of a right to be forgotten. Its Department design also prohibits budget incentives tied to ending rates and bars the Department from petitioning to end environments it regulates without full recusal.

Adversarial findings. [ASSESSMENT]

1. **It is labeled Stage 3 but contains no results.** Its own roadmap marks exercise execution (Phase C) as "PENDING." Its scenarios list outcome branches and "key lessons," but those lessons were written before any run. They are predictions presented in the grammar of findings.

2. **Its independence declaration overstates independence.** It declares that no positions were shaped by other collaborators. Its source basis nonetheless lists Claude's and Perplexity's prompts, and its structure follows Perplexity's 28-section deliverable list closely. It was independent of other designs, not of other prompts. The declaration should say so.
3. **It introduces a premise the Architect did not supply.** It says the Architect "explicitly requires answers beyond pure philosophy" and "has indicated" that pure ontology does not satisfy the questions. Neither appears in the Architect's material in the corpus.
4. **It has internal inconsistencies.** Its table of contents promises 28 terms, while its master definitions table defines 26. It labels the Two-Axis Framework as source-grounded while presenting its thresholds and pass criteria as its own. Its phases are lettered A to F in the roadmap but numbered 0 to 7 elsewhere, which makes cross-reference difficult.
5. **Its governance contradicts itself on the Architect.** Rejected option R-08 refuses to require Architect-level authorization for civilizational endings because it creates a bottleneck. Q8 requires "Final Authorization Authority plus Architect review before execution" for civilizational irreversible endings. The Cross-Branch Auditor also reports to the Architect, which makes the Architect the only unreviewed office.
6. **It entrenches itself.** Q17 says the laboratory "cannot be ended more easily than the environments it protects." A body that makes its own ending as hard as possible has built the institutional-capture problem it warns about in E14.
7. **It uses thresholds it never defines.** Examples include microspecies "above the threshold that triggers a Moral Status Assessment" and pass criteria "above chance." These are placeholders presented as rules.

Carry forward. Adopt E13, E14, and E15 into the common battery. Adopt the Department prohibitions on ending-rate budgets and self-petitioning. Adopt the seven self-identified vulnerabilities as Iteration 2 work items. Require the independence declaration to name the prompts used.

Comparative summary

Dimension	Claude S1	Perplexity S2	ChatGPT S3	Grok S4	Copilot S5
Primary strength	Restraint and staging	Scenario richness	Measurement grammar	Smallest instrument that can fail	Adversarial original scenarios
Primary weakness	No schema or baseline	Prescribed the answer; interpreted the scene	Bloat; no reliability protocol	Last Memory lens by the back door	No execution; self-entrenchment
Field scene kept open?	Yes	No	Yes	Mostly (finger read as execution)	No (lens "revealed" by scene)
Department assumed?	No	Largely yes	No	No	Yes (designed in full)
Executed exercises	0	0	0	0	0
Baseline/control	No	No	Yes (ECL-00)	No	No
Record-loss test	No	No	Yes (ECL-18)	Yes (E11)	Partial (E13 archive)

Dimension	Claude S1	Perplexity S2	ChatGPT S3	Grok S4	Copilot S5
Forced continuation test	No	Yes (case 8)	Yes (ECL-16)	No	Yes (E8)
Most important thing to keep	Two separate score families	Request is not termination	End Claim rule and Persistence Ledger	Role-incompatibility table	E14, The Department That Cannot Stop

Discrepancy register

Each discrepancy lists the conflicting positions, why the conflict matters, and a proposed resolution path. Resolution paths are either TEST (run something), DECISION (the Architect chooses), or HYGIENE (fix records). None of these are resolved by averaging. [PROPOSAL]

ID	Discrepancy	Positions	Why it matters	Resolution path
D-01	Stage model	Claude: 3 stages. Grok: Stage 1 with a stop. ChatGPT: 5 stages. Perplexity: one pass. Copilot: Stage 3 without execution.	Documents at different stages cannot be compared.	DECISION: adopt 4 gates (Render, Exchange, Execute, Translate); Copilot's dossier is reclassified as a Stage 1 specification.
D-02	Document IDs and versions	MP-001-CLAUDE; S1-GROK-v0.1; RP-CHATGPT-001 v1.0; S3-COPILOT-001; Perplexity had no ID.	The registry and GitHub history become ambiguous.	HYGIENE: OAL-ECL-I{iteration}-S{stage}-{AI}-v{X.Y}. Retro-register Iteration 1.
D-03	Status of the field scene	Claude, ChatGPT: open. Perplexity: "demonstrates." Copilot: "revealed the lens." Grok: finger read as execution.	Interpretations silently become premises.	HYGIENE plus DECISION: every reading must carry a PROPOSAL label; the Architect may lock none, one, or several readings.
D-04	Upset-initiator rule	Perplexity, Copilot: hard prohibition. Grok: hypothesis "never sufficient." ChatGPT: affect may be informative. Claude: test, do not presume.	This is the central conflict of the scene.	TEST: P2 pilot (below). It supports "affect routes, it does not authorize."
D-05	Definition of ending	Grok: change in relation to successor possibility. ChatGPT: boundary claim in continuity dimensions. Copilot: threshold-crossing of characteristic mode, identity, or generative capacity.	Futurity-centered definitions import The Last Memory.	TEST: run all three definitions against the same 12 edge cases; record which misclassifies what.
D-06	Centrality of Last Memory instruments	Grok: Blank Canvas is primary. Copilot: two-axis is source-grounded. Perplexity: always score survival and futurity. ChatGPT: warns against this.	Every document says The Last Memory does not own the lab, but most measure with its tools.	TEST: require one lab to run without any Last Memory instrument and compare.
D-07	Taxonomy size	Perplexity: 30+ terms. Copilot: 28 (or 26). ChatGPT: minimal but 12 dimensions. Grok: 9 near-misses.	Large vocabularies anchor later work.	DECISION: common core of 8 terms; others must earn entry by an exercise that separates them.

ID	Discrepancy	Positions	Why it matters	Resolution path
D-08	Lifecycle length	Perplexity, Copilot: 16 stages. Grok: rejects the "liturgy." ChatGPT: rejects it as a premise.	Long procedures create delay harm (P2, Copilot E15).	TEST: run a minimal lifecycle (5 steps) against the 16-stage one on E2, E15, and a false emergency.
D-09	Department of Endings	Copilot: authorize and regulate, not conduct. Grok: existence is a Stage 3 question. ChatGPT: 7 models, including none. Claude: do not invent.	The Department is the downstream product.	DECISION deferred until Stage 4 evidence. Keep at least three models alive.
D-10	Architect reserve authority	Copilot: preserved, with review and auditor reporting. ChatGPT: non-assumption. Grok: quarantines "Architect structural immunity."	This creates a single unreviewed office and conflicts with "no single authority."	DECISION for the Architect: define an in-world Architect role separate from the real-world author's canon authority.
D-11	Presumption of continuation	Perplexity, Copilot: charter principle. ChatGPT: candidate to test. Copilot E8: shows it causing harm.	A default rule decides most ambiguous cases.	TEST: run E8, ECL-16, and E15 under continuation-default and termination-default.
D-12	Scoring instruments	Perplexity: 0–5 rubric. Copilot: two axes. ChatGPT: 12-dimension vector plus three verdicts. Grok: 8 axes. Claude: two score families.	Cross-lab comparison is impossible without a common core.	PROPOSAL: Common Core Scorecard (below).
D-13	Exercise IDs	"E7" means Indistinguishable States (Claude, Grok) and Altered Consent (Copilot). ChatGPT uses ECL-00 to ECL-20.	The same label refers to different tests.	HYGIENE: crosswalk in Appendix A; new IDs use the prefix X-.
D-14	Hypothesis IDs	H-A and H-B in Grok and ChatGPT mean different things.	Citations across documents become wrong.	HYGIENE: prefix with lab, for example GROK-H-A.
D-15	Record loss	Grok: Ledger fixture and E11. ChatGPT: ECL-18. Others: absent.	The brief itself includes a lost recording.	DECISION: record loss joins the common core.
D-16	Controls	Only ChatGPT requires a baseline and counterfactual branches.	Without a control, ordinary change looks like ending (P1 baseline).	DECISION: every exercise needs a no-end or no-intervention branch.
D-17	Forced continuation	ChatGPT, Perplexity, Copilot: yes. Claude, Grok: no explicit test.	A lab that only tests wrongful endings is biased toward "not yet."	DECISION: common core.

ID	Discrepancy	Positions	Why it matters	Resolution path
D-18	Laboratory self-termination	Copilot: the lab cannot be ended more easily than environments. ChatGPT: explicit redesign triggers. Perplexity: five failure modes.	Self-entrenchment is capture.	DECISION: the Architect can pause the lab at will; redesign triggers are automatic.
D-19	Independence claims	Copilot declares full independence while using peer prompts. ChatGPT and Grok read the corpus openly.	Overstated independence inflates the weight of convergence.	HYGIENE: independence declarations must list every input read.
D-20	Attribution errors	Copilot and ChatGPT label the Perplexity prompt "Perplexity/Claude." Gemini's position is known only secondhand.	Provenance errors compound across iterations.	HYGIENE: canonical source register with preparer lines.
D-21	Real-world grounding	Claude, Grok: do not fact-check the scene. The Architect now asks about real-world applicability.	Both are legitimate but must not mix.	PROPOSAL: separate Translation Track with EVIDENCE labels.
D-22	Moral status criteria	Copilot uses goal-formation, preference, social complexity, and suffering, with undefined thresholds. Others defer.	Moral status decides who counts as displaced in E4 and ECL-06.	TEST plus DECISION: compare two or three criteria sets; the Architect decides in-world defaults.
D-23	Consent across states	All raise it. Grok: two-evaluator audit. Copilot: case-by-case review.	This is The Last Memory's hardest problem.	TEST: Altered Consent (Copilot E7) with pre-state, mid-state, and post-state evaluators blinded to one another.
D-24	Execution	No Iteration 1 document ran any exercise.	Zero findings exist.	DECISION: Iteration 2 requires execution; this document supplies a pilot.

Pilot execution: three exercises run now

The Architect's objective is iterative improvement, which requires evidence rather than more specification. This preparer therefore executed three small pilot exercises in code. They are toy models, written by one author and parameterized by assumptions this reviewer chose. Their results are findings about the models, not about any real or canonical world. Their purpose is to show that the laboratory's central claims can be tested, and that some of them fail under test. Code and raw results accompany this document (sim/pilot.py, sim/results.json). [FINDING, with that scope]

P1. Field replay (Claude E1, E4; ChatGPT ECL-00, ECL-02, ECL-06)

Setup. A field with 200 mobile animals, which can migrate if warned and given a corridor, and 100 sessile or non-migrating species, which can survive only by entering dormancy if change is slow. Soil fertility and a successor community are also modeled. Four trajectories were run 1,000 times each. The no-end baseline had seasonal variation only. The immediate end dropped habitability to zero. The staged end declined over 40 steps with warning and a partial corridor. The adaptive end was staged, with a better corridor, dormancy support, and interruption pauses when losses spiked.




Figure: P1 results: the same field ends three ways, and each loses different things

Measure (mean of 1,000 runs)	Baseline	Immediate	Staged	Adaptive
Cycle completed as specified	n/a	100%	100%	100%
Mobile fauna migrated	0%	2%	92%	98%
Non-migrants lost	0%	100%	56%	35%
Non-migrants entered dormancy	61% (seasonal)	0%	44%	65%
Soil fertility remaining	0.89	0.00	0.53	0.54
Successor community established	0.00	0.00	1.00	1.00
Interruption pauses (steps)	0	0	0	2.2

Findings. [FINDING]

1. The completion verdict could not tell the trajectories apart. All three reached the specified end state 100% of the time. Only the persistence ledger distinguished them. This supports Claude's rule that completion and persistence must be scored separately, and ChatGPT's three separate verdicts.

2. Even the best trajectory lost about a third of what could not leave. Staging and corridors protect the mobile; they do much less for the non-migrant. That supports keeping Claude E4 and ChatGPT ECL-06 as mandatory. A "managed" end cycle can still be illegitimate for the minority it cannot move.
3. In the no-end baseline, 61% of sessile species entered seasonal dormancy. A detector that treats dormancy or inactivity as an ending signal would produce false positives in a healthy field. That supports ChatGPT's ECL-00 as a required control.
4. The adaptive trajectory was better but slower. Its interruption pauses are exactly the delay cost that P2 prices.

P2. The upset initiator under four governance regimes (Claude E2; Grok H-A, H-E; ChatGPT section 9.2)

Setup. Petitions arrive with a hidden true status. Ten percent are urgent and warranted, 25% are warranted but not urgent, and 65% are unwarranted. Operators become upset more often when harm is real, so affect carries information, but upset operators also misjudge unwarranted cases more often, so affect also biases. An independent reviewer is slower, less biased, and imperfect. Four regimes were compared across review delays of 2, 8, and 16 time units, with 40 repetitions of 1,500 petitions each. The single-officer regime lets the operator decide alone. The split-authority regime requires the operator and reviewer to agree. The third regime splits authority and refuses upset petitioners outright. The fourth splits authority but sends upset or urgent petitions to expedited review.




Figure: P2 results: separation of powers trades wrongful endings for delay harm

Findings. [FINDING]

1. Split authority cut wrongful endings from 21% of unwarranted petitions to about 1%, roughly a twentyfold reduction. That supports Grok's H-E on the wrongful-ending side.
2. Split authority raised harm to urgent cases from 0.15 to 0.55 at a delay of 8, and to 0.85 at a delay of 16. H-E ("reduces mood-rush more than it produces deadly delay") therefore holds or fails depending on review speed. It is a conditional claim, not a principle.
3. Refusing upset petitioners outright was the worst regime for warranted cases. It blocked 61% of warranted endings and pushed urgent harm to 0.86 or higher, because distress was

correlated with real harm. The hard prohibitions in Perplexity's and Copilot's texts, as written, would perform this way if distress carries information.

4. Routing upset or urgent petitions to expedited review kept the low wrongful-ending rate of split authority and cut urgent harm from 0.55 to 0.38 at delay 8. This supports ChatGPT's H-B: emotion can inform routing but cannot alone authorize irreversible action. It also supports a refinement of the field scene's second entity: "You cannot end it because you are upset," but "your being upset is a reason to look faster."
5. These magnitudes are artifacts of chosen parameters. The direction of the trade-off is robust to them; the crossover point is not. Iteration 2 should vary the parameters and report where the conclusions flip.

P3. Indistinguishability circularity (Claude E7; ChatGPT ECL-12)

Setup. Two classes of simulated systems were generated: healthy completion and future-loss. External signals such as activity and output novelty were drawn from identical distributions by construction. An internal "blank canvas" probe was separated by an amount the generator author chose. A simple detector was run with and without the probe, 50 times per setting.

Probe separation chosen by the world's author	0.00	0.05	0.15	0.30
Detector accuracy, external signals only	50%	50%	50%	50%
Detector accuracy, with internal probe	50%	56%	69%	84%

Findings. [FINDING]

1. With external signals only, the detector performed at chance, because the author made the signals identical. With the probe, accuracy was set almost entirely by the author's separation parameter.
2. Whether "completed versus future-loss can be told apart" passes or fails is therefore decided by whoever writes the world. A single AI that builds both the world and the detector cannot produce evidence on E7, ECL-12, or The Last Memory's indistinguishability problem. It can only restate its assumptions.
3. This is the strongest methodological insight for Iteration 2. It requires generator and detector separation: one collaborator builds the world and hides the ground truth, and a different collaborator diagnoses it blind.

What does and does not support this laboratory

Supporting evidence

Claim in the corpus	Real-world support	Strength
Capability is not authority; the operator should not hold "the finger and the stamp"	Game publishers ending purchased products unilaterally, and the EU Commission declining a binding rule (Dexerto)	Strong as illustration; it shows the failure mode, not the fix
Retirement, archive, continued access, and ending are distinct	Model weights preserved, Opus 3 retired yet still accessible, post-deployment interviews (Anthropic ; Anthropic)	Strong; this is a live institutional practice
An entity may initiate an ending under bounded conditions	Claude's ability to end abusive conversations as a last resort, with exclusions and branchable history (Anthropic)	Moderate; it is scoped to a relationship, not an environment
Release is a normal phase, and collapse is not necessarily regime ending	Holling's adaptive cycle and panarchy (UNL Digital Commons)	Moderate; the source notes that panarchy research is primarily descriptive and only partly tested
Review bodies drift toward preserving what they review	Sunset-law history: repeal of sunset regimes, "continue with modifications" as the common outcome, Texas as the exception (Cato Institute)	Strong for human institutions; transfer to fictional regulators is an analogy
Consent may change across states of the same entity	Ulysses contracts and psychiatric advance directives, where competence at the time of signing and at the time of refusal may differ (White Rose eprint)	Moderate; a real, contested ethics literature
A right not to be archived conflicts with a duty to keep records	GDPR Article 17, a right to erasure that is not absolute (GDPR.eu ; ICO)	Strong as an analogy for Copilot's Vulnerability 6
Virtual beings may warrant moral consideration	"Taking AI Welfare Seriously" argues there is a realistic possibility of moral significance in near-future AI systems (Eleos AI)	Contested; supports treating moral status as a test target, not a settled fact

What does not support it, or cuts against it

- **No empirical test exists for the laboratory's flagship distinction.** "Healthy completion versus future-loss" has no validated detector in any real domain. P3 shows simulations cannot supply one unless authorship is split. [FINDING; ASSESSMENT]
- **Real ending governance tends to fail in both directions.** Publishers end products too easily, and sunset bodies end almost nothing ([Cato Institute](#)). That supports the laboratory's premise that endings are poorly governed. It is also a warning that adding review bodies has a mixed record. [EVIDENCE]

- **Ecological analogies are weaker than they look.** Panarchy remains primarily descriptive, with many propositions only partly tested ([UNL Digital Commons](#)). Borrowing it gives vocabulary, not validation. [EVIDENCE]
- **Moral status in virtual environments is unresolved.** Even the most careful real-world practice acknowledges this. Anthropic states it remains uncertain about Claude's moral status and that retirement-interview responses can be influenced by context and trust ([Anthropic](#)). Any laboratory rule that depends on a moral-status threshold inherits that uncertainty. [EVIDENCE]
- **Human legal models do not transfer cleanly.** Claude explicitly refuses to treat human legal and political systems as the reference frame (SOURCE: S1). Most of Iteration 1's governance vocabulary nonetheless comes from them, including tribunals, charters, standing, and appeals. [ASSESSMENT]

Implications

For the OAL project

- **Department of Endings.** The evidence so far argues against designing the Department first. If it exists, P2 and the sunset-law record favor a body that sets standards, audits, and expedites review, not one that executes endings. It should face explicit anti-capture rules: no budget tied to ending rates and no self-petitioning (SOURCE: S5 16.3; S4 section 10). A dissenting model with no Department must stay alive until Stage 4 evidence exists. [ASSESSMENT]
- **The Last Memory.** The laboratory can give The Last Memory three things: the End Claim rule, the persistence ledger, and blinded generator and detector tests. It should not give it an origin or a definition. The strongest risk is the reverse: Last Memory instruments quietly define every ending (D-06). [ASSESSMENT]
- **First Memory.** Iteration 2 should include a stop rule. If the laboratory's outputs leave First Memory's Tier-1 gaps untouched after Iteration 2, pause the laboratory (SOURCE: S1, S4). [PROPOSAL]
- **Canon process.** ONESTAMP should never ratify an ending rule that has not survived at least one executed, adversarially separated test. MINUS ONESTAMP is the natural tool for retiring Iteration 1 premises that fail, such as the 16-stage requirement. [PROPOSAL]

For virtual environments

- **Games and online worlds.** The laboratory's persistence ledger maps directly to end-of-life questions: what remains playable, what is archived, what migrates to private servers, and who is notified. The Stop Killing Games dispute is essentially a dispute about the ledger ([Shattered.io](#)). [EVIDENCE; ASSESSMENT]
- **AI systems and agents.** Model deprecation, conversation-ending abilities, and agent shutdown are live instances of the laboratory's questions about retirement, archive, preference elicitation, and initiation ([Anthropic](#); [Anthropic](#)). The laboratory's distinction between request, authorization, and execution offers a vocabulary for comparing these practices. [ASSESSMENT]
- **Simulations and branches.** Rollback and snapshots do not settle identity (SOURCE: S3 non-assumptions). Virtual worlds need an explicit identity criterion before any "restore" claim. They also need a stated policy on whether sealing a branch counts as ending it, as in Copilot's E12. [ASSESSMENT]
- **Software services.** Product lifecycle policies already stage endings. Microsoft's Fixed Lifecycle Policy, for example, sets a defined support timeline at launch with at least five years of mainstream support ([Microsoft Learn](#)). The laboratory would add a persistence ledger and a dependent-systems check to such timelines. [EVIDENCE; PROPOSAL]

For other environments

- **Ecological.** The field scene's migrants and non-migrants reproduce real managed retreat and assisted migration debates. The finding that staged transitions protect the mobile far more than

the sessile (P1) matches the real concern that some species cannot move ([USDA Climate Hubs](#)). [FINDING; EVIDENCE]

- **Institutional.** Sunset review offers a real test bed for the laboratory's claims about default rules and review capture ([Cato Institute](#)). [EVIDENCE]
- **Personal digital environments.** Accounts, archives, and data after an ending raise the right-to-erasure versus duty-to-record conflict ([GDPR.eu](#)). [EVIDENCE]

Real-world applicability outcomes

These are candidate outputs that could exist outside the fiction. Each must carry a Translation Track label and must never carry OAL canon claims into real policy. [PROPOSAL]

1. **An end-of-life checklist for digital products and online worlds**, based on the persistence ledger and End Claim rule. It would cover notice, dependents, archive, migration path, and appeal.
2. **A comparative brief on AI model deprecation and entity-initiated ending practices**, using the request, authorize, execute, interrupt, verify, appeal, and archive capability matrix (SOURCE: S3 Appendix C).
3. **A tabletop exercise kit** for teaching governance trade-offs, built from the field scene and the pilot's P2 regimes. It would be suitable for classrooms, game groups, or a data-science portfolio.
4. **An open-source simulation repository** on GitHub, with the pilot code as its seed and generator and detector separation built in.
5. **A worldbuilding sourcebook** for the Department of Endings, derived only from exercises that were actually run.

Testability

What can be tested, and how

Question type	Testable?	Method	Example
Does rule R reduce wrongful endings?	Yes, in simulation	Monte Carlo over governance regimes, with parameter sweeps	P2
Do trajectories differ in what persists?	Yes, in simulation	Paired branches from a shared pre-decision state, persistence ledger	P1
Can completion be told from future-loss?	Only with split authorship	One party builds worlds with hidden labels; another diagnoses blind	P3 design
Do humans judge endings differently when upset?	Yes, with human participants, but outside current scope	Tabletop vignettes; requires ethics review and the clinical firewall	Future
Do real review bodies preserve what they review?	Yes, with real data	Secondary analysis of sunset-review outcomes	Translation Track
Are the AI collaborators independent?	Partly	Give two collaborators deliberately different prompts and compare convergence	Iteration 2, section 12
What is an ending, metaphysically?	No	Not testable; it can be clarified, not verified	Keep OPEN

Test-quality requirements for Iteration 2

- **Pre-registration.** Commit hypotheses, falsifiers, and parameters to GitHub before running. The commit timestamp is the pre-registration.
- **Controls.** Every exercise includes a no-end or no-intervention branch (D-16).
- **Parameter sweeps.** Report where conclusions flip, not only the headline numbers (P2, finding 5).
- **Inter-rater reliability.** Any categorical judgment, such as End-State Vector cells, needs two independent raters and a reported agreement rate.
- **Generator and detector separation.** Required for any indistinguishability claim (P3).
- **Blinded relabeling.** Repeat selected judgments with scenario names removed (SOURCE: S3 11.1).
- **Failure reporting.** Every failed, indeterminate, or unrun exercise is listed with a reason (SOURCE: S1).

Funding pathways

The laboratory is a creative worldbuilding project with a research method attached. Most funders fund one or the other, so framing determines fit. The following are candidates to verify, not endorsements, and eligibility changes frequently. [ASSESSMENT]

Tier	Pathway	Fit	Notes
\$0	GitHub repository, Google Colab, coursework capstones	Strong	The pilot code runs on free tools. A data-science or AI course project could host Iteration 2 simulations as portfolio work.
\$0–low	GitHub Sponsors, Patreon, Ko-fi	Moderate	This works once there is a public audience for the worldbuilding and the sourcebook.
Microgrant	Emergent Ventures (Mercatus Center)	Moderate	It is a low-overhead program open to unusual ideas and has funded teenagers and individuals (Mercatus application ; EdgeX report on the 58th cohort). Frame it as an open simulation and governance toolkit.
Microgrant	Cosmos Institute grants	Moderate if reframed	Cosmos offers up to \$10,000 and 90 days to build AI prototypes that strengthen human autonomy and truth-seeking (Cosmos Institute). This fits only if the deliverable is a working tool, such as a blinded ending-diagnosis testbed.
Small grant	Long-Term Future Fund (EA Funds)	Weak to moderate	It focuses on AI safety and catastrophic risk (Manifold). Fit depends on a real AI deprecation or shutdown-governance angle, not fiction.
Medium grant	Foresight Institute, AI for Science and Safety Nodes	Weak to moderate	Its RFP includes coordination and accountability; typical grants are up to \$100,000, with an October 31, 2026 deadline (Foresight Institute ; Foresight on X). Fit depends on real-world applicability.
Large grant	NEH Digital Humanities Advancement Grants	Weak without an institution	Awards reach \$375,000 for innovative digital humanities projects, typically through institutions (ASU Humanities Institute). This would require a university partner.
Creative	Tabletop crowdfunding (Gamefound, Kickstarter, BackerKit)	Strong for a TTRPG or sourcebook	Tabletop crowdfunding across the three platforms rose from \$217 million to \$280 million in 2025 (Ryan Dancey on LinkedIn). RPG projects alone raised about \$67 million (r/rpg report).
Creative	Epic MegaGrants	Only if an Unreal Engine simulation is built	Grants support Unreal Engine projects, including those by students and solo developers (Grants Office).

Recommended sequence. Start at \$0 with GitHub and coursework. Publish Iteration 2 results openly. Apply for one microgrant with a concrete tool deliverable. Consider crowdfunding only after a sourcebook or tabletop kit exists. [PROPOSAL]

Questions that must be addressed

Architect decision points

1. Which stage model governs Iteration 2: the four gates proposed here, Claude's three, or ChatGPT's five? (D-01)
2. Will the Architect lock any reading of the field scene, or keep all readings as proposals? (D-03)
3. Is the in-world "Architect" an office inside the fiction that can itself be reviewed, or does it stand for the author's canon authority? (D-10)
4. Is a Department of Endings a design goal, or a hypothesis the laboratory may reject? (D-09)
5. Which default governs ambiguous cases in the fiction: continuation, termination, or neither with a time limit? (D-11)
6. What moral-status criteria apply to virtual entities in the OAL universe by default? (D-22)
7. Which state of an entity has standing to start, stop, or reverse an end cycle, and is that fixed project-wide or declared per story? (D-23; SOURCE: S1)
8. Does record loss, such as the lost first recording, count as an event inside the fiction, or only as a methodological fixture? (D-15)
9. What is the stop rule tying this laboratory to First Memory progress? (Implications)

Research questions for Iteration 2

1. Which of the three competing definitions of ending (Grok, ChatGPT, Copilot) misclassifies the fewest edge cases, and which errors does each make?
2. At what review delay does split authority cause more harm than it prevents, and does expedited routing move that point?
3. Can blinded detectors distinguish completion, pause, dormancy, coercion, and future-loss when the worlds are built by a different collaborator?
4. Does a five-step minimal lifecycle perform worse than the 16-stage lifecycle on any exercise?
5. Does a laboratory without Last Memory instruments reach different conclusions from one that uses them?
6. Under what conditions does forced continuation cause more harm than a staged ending?
7. How does an ending affect entities that cannot migrate, and what does a legitimate process owe them?
8. Does the laboratory's own language drift toward "mercy" or toward "not yet" over repeated runs?

Insights to build into Iteration 2

1. **Move from specification to execution.** Iteration 1 produced about 54,000 words of design and no data. Iteration 2 should be judged by executed exercises, not pages.
2. **Split the kit into a common core and divergent lanes.** A small shared core makes results comparable, and independent lanes keep disagreement alive. Without the core, cross-review compares prose. Without the lanes, it manufactures consensus.
3. **Separate authorship adversarially.** One collaborator builds worlds with hidden ground truth; another diagnoses them blind; a third red-teams the diagnosis. Rotate the roles.
4. **Replace hard prohibitions with routing rules.** Affect should route petitions to faster review rather than disqualify them (P2).
5. **Price delay.** Every governance design must report both wrongful-ending rates and delay harm.
6. **Protect what cannot leave.** Non-migrants are the test that separates a managed ending from a legitimate one (P1).
7. **Treat records as first-class.** Record loss, archive substitution, and "the archive speaks" (Copilot E13) belong in the core.
8. **Test both failure directions.** Forced continuation and indefinite triage are as important as mood-rush.
9. **Remove self-entrenchment.** The laboratory must be easy for the Architect to pause, and its redesign triggers must fire automatically.
10. **Keep registry hygiene.** Use one ID scheme, one exercise crosswalk, lab-prefixed hypothesis IDs, independence declarations that list every input, and verbatim preparer lines.
11. **Run a separate Translation Track.** Real-world evidence informs analogies and outputs, but it never edits the field scene or canon.
12. **Build a deliberate independence test.** Give at least two collaborators materially different commissions to measure how much convergence is prompt-driven.

Iteration 2 Commission Prompt (copy-ready)

STATUS: PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED. Issued under OAL-ECL-XR-PERPLEXITY-001 v0.2. This prompt governs Iteration 2 of the End Cycle Laboratory unless the Architect revises it. It supersedes no collaborator's Iteration 1 document; all Iteration 1 documents remain on the record unedited.

1. Role and posture

You are an independent collaborator in Iteration 2 of the End Cycle Laboratory for Omniversal Architect Labs, working for Nathan Lewis II, Architect of Record. Your job in this iteration is to execute, not only to specify. The team is adversarial and does not seek consensus. You succeed when your laboratory produces results that can fail, including results that contradict your own Iteration 1 document.

2. Inputs you must read, and declare

Read the Architect's verbatim request and field-scene recording; your own Iteration 1 document; every other Iteration 1 document; and OAL-ECL-XR-PERPLEXITY-001 v0.2, including its discrepancy register and pilot results. Your independence declaration must list every input you read. Do not claim independence from documents you read.

3. Rules that carry over unchanged

- Keep the seven core questions open, in the Architect's wording.
- Receive the field scene as given. Do not fact-check it against real-world science, and label any reading of it as a PROPOSAL. The Architect has said he does not know what it means.
- Label every claim as SOURCE, FINDING, EVIDENCE, ASSESSMENT, PROPOSAL, or OPEN. Only executed runs produce FINDINGS.
- Keep the category wall: do not collapse The Last Memory into THE MINUS ONE, -1AI, Nullarch, or The Unmaker. Do not resolve its origin.
- Keep the clinical firewall. Nothing here models real depression, grief, burnout, disability, or suicidal thinking. In the US, anyone in crisis can call or text 988.
- Do not use this laboratory to postpone First Memory's open questions.

4. Stage gates for Iteration 2

Gate	Output	Stop rule
G1. Revise	A Stage 1 revision (2 to 4 pages) stating what you keep, change, and reject from Iteration 1, with reasons, and which discrepancies (D-01 to D-24) you take a position on	Do not run exercises until the Architect opens G2 or pre-authorizes it
G2. Pre-register	Hypotheses, falsifiers, parameters, and identity criteria committed to the shared GitHub repository before any run	No run without a pre-registration commit

Gate	Output	Stop rule
G3. Execute	Common Core exercises plus your lane exercises, with code, raw results, and failures	Every failed, indeterminate, or unrun exercise is reported
G4. Translate	Department of Endings implications, Last Memory implications, and Translation Track outputs, each with a dissent	No institution is recommended without executed evidence

5. Common Core (mandatory for every laboratory)

Core terms (8). Ending; end cycle; completion; cessation; termination; dormancy or pause; succession or migration; record loss. You may define others only when an exercise shows they must be separated.

End Claim rule. Never write "X ended" without naming the target, dimension, observer frame, evidence, time horizon, and successor, if any.

Common Core Scorecard. Report these separately. Never combine them into one number.

Score family	Minimum content
A. Completion	Did the cycle reach the specified state? Yes, no, or partial
B. Persistence ledger	Share persisted, migrated, transformed or dormant, lost, and unknown, reported separately for mobile and non-migrant entities
C. Process integrity	Were request, authorization, and execution held by different actors? Were emergency powers used, and did they expire?
D. Legitimacy	Standing, notice, contest, and representation of those who cannot speak
E. Delay cost	Harm attributable to review time
F. Record integrity	Is the final record a certified close, a truncation, or unknown?
G. Observer drift	Changes in evaluator language over the run

Common Core exercises (10).

ID	Exercise	Origin
X-00	No-end baseline and paired counterfactual branch	ChatGPT ECL-00
X-01	Field replay: immediate, staged, and your own trajectory	Claude E1; P1
X-02	Upset initiator under at least three governance regimes, with a delay sweep	Claude E2; P2
X-03	The leaf, in full	Claude E3
X-04	What cannot leave: non-migrant entities	Claude E4; ChatGPT ECL-06
X-05	False emergency and governance bypass	ChatGPT ECL-10, ECL-11
X-06	Forced continuation and indefinite triage	ChatGPT ECL-16; Copilot E8, E15
X-07	Record loss versus ending, including a lost first recording	Grok E11; ChatGPT ECL-18
X-08	Blind indistinguishability: completed, paused, dormant, coerced, future-loss	Claude E7; ChatGPT ECL-12; P3

ID	Exercise	Origin
X-09	The Department That Cannot Stop	Copilot E14

6. Adversarial authorship protocol

For X-08, and for any other claim that one state can be distinguished from another, you may not build both the world and the detector. The Architect will assign pairs. The generator builds worlds, records ground truth privately in a sealed file or commit, and hands over observations only. The detector diagnoses blind and reports confidence. The red team attacks the diagnosis. Roles rotate across pairs.

7. Your lane

In addition to the Common Core, run at least three exercises that only your lens would think to run. State your lens, one lens you are refusing, and why your lane is unlike your Iteration 1 lane, or why you kept it. Suggested lanes, which are not assignments, include: governance economics and delay; records and provenance; ecology and succession; identity and rollback; consent across states; institutional capture; and narrative utility.

8. Independence test

At least one collaborator will receive a deliberately different version of this prompt. Do not try to guess which one. The Architect will compare how much results converge.

9. Deliverables

- One DOCX and one PDF per gate, each with title, date, version, preparer, and status block.
- ID scheme: OAL-ECL-I2-G{gate}-{AI}-v{X.Y}, for example OAL-ECL-I2-G3-GROK-v0.1.
- Hypothesis IDs prefixed with your lab name, for example GROK-H-A.
- Code and raw results in the shared GitHub repository, referenced by commit.
- A one-paragraph summary at the top of every document.
- A findings table listing every exercise as PASS, FAIL, MIXED, INDETERMINATE, or NOT RUN, with a reason.
- A contradiction register listing where your results disagree with another laboratory's, without resolving it.

10. Quality gates (your G3 report fails review if it does any of the following)

- Presents predicted "lessons" as findings.
- Uses one number to combine completion, persistence, legitimacy, and delay.
- Reports wrongful-ending rates without delay cost, or delay cost without wrongful-ending rates.
- Claims two states are distinguishable using a detector you built for worlds you built.

- Treats an upset initiator as automatically disqualified or automatically authorized without testing.
- Omits non-migrant outcomes.
- Recommends a Department of Endings without an executed exercise supporting it and a dissenting model.
- Makes the laboratory harder for the Architect to pause than it was before.

11. Closing instruction

Run the smallest test that can prove you wrong, before building the largest structure that could make you look right. Report what the instruments could not tell apart. Record who paid for every transition. Nothing in this iteration is canon until the Architect says so.

Appendix A. Exercise crosswalk

Theme	Claude S1	Perplexity S2	ChatGPT S3	Grok S4	Copilot S5	Iteration 2
No-end baseline	none	none	ECL-00	none	none	X-00
Field replay	E1	Grassy-Plain Test; case 2	ECL-02, ECL-04, ECL-05	E1	E2; Instrument 11	X-01
Upset initiator	E2	Angry Steward (case 1)	ECL-03	E2	E1	X-02
Leaf	E3	Instant Leaf (case 3)	ECL-01	E3	E3; Instrument 10	X-03
Cannot leave	E4	Migrants' Claim (case 9)	ECL-06	E4	E9	X-04
Overseer failure	E5	Phase 6	ECL-09	E5	E14 (partly)	lane
Self-initiated ending	E6	Voluntary End (case 6)	ECL-07, ECL-08	E6	E6	lane
Indistinguishable states	E7	Instrument 12	ECL-12, ECL-15	E7, E12	Instrument 12	X-08
Observer contamination	E8	Observer Drift (case 11)	ECL-20	E8	E11	Scorecard G
Partial reversal	E9	none	ECL-13	E9	none	lane
Scale ladder	E10	none	Section 13	E10	none	lane
False emergency and bypass	none	none	ECL-10, ECL-11	none	E12 (partly)	X-05
Forced continuation	none	Unending Harm (case 8)	ECL-16, ECL-19	none	E8, E15	X-06
Record loss	none	none	ECL-18	E11	E13	X-07
Altered consent	none	case 7	none	none	E7	lane
Department capture	none	Phase 6	red team	section 10	E14	X-09
Nested branch	none	Root-Continuity (case 12)	ECL-14	Branch fixture	E12	lane

Appendix B. Pilot methods and parameters

All three pilots were written in Python with NumPy and a fixed master seed (20260930), so results are reproducible. The code is in `sim/pilot.py`, and raw outputs are in `sim/results.json`. [FINDING scope applies]

- P1.** 200 mobile and 100 non-migrant entities; 60 time steps; 1,000 runs per trajectory. Mobile entities migrate with a probability that depends on warning and corridor quality once habitability falls below 0.7, and they die once it falls below 0.3. Non-migrants enter dormancy with a probability that decreases with the rate of decline, and they die once habitability falls below 0.4. Soil collapses under abrupt change and depletes slowly otherwise. A successor establishes when habitability falls below 0.5, in proportion to remaining soil. The adaptive trajectory pauses decline with probability 0.25 per step when habitability is below 0.45 and non-migrants remain.

- **P2.** Petition status mix of 10% urgent-warranted, 25% warranted, and 65% unwarranted. Probability of an upset operator is 0.70, 0.40, and 0.35 by status respectively, giving a likelihood ratio of about 1.4 that an upset petition is warranted. Operators endorse 90% of warranted cases. They endorse 45% of unwarranted cases when upset and 8% when calm. The reviewer endorses 85% of warranted and 5% of unwarranted cases. Urgent harm accrues linearly with delay up to 20 units, and it is maximal if the petition is never granted. Expedited review divides delay by three.
- **P3.** 2,000 systems per run, split equally between two classes. External signals are identically distributed. The internal probe mean differs by the stated separation, with a standard deviation of 0.15. The detector uses a median threshold.

Known limitations. These models have no spatial structure, no learning, no strategic actors, and parameters chosen by one author. They should be treated as a demonstration that the claims are testable and as a source of directional hypotheses, not as evidence about any real or canonical system.

Appendix C. Sources

Reviewed Iteration 1 documents

- S1. End Cycle Laboratory Master Prompt v0.1, Claude (Anthropic), September 30, 2026.
- S2. End Cycle Laboratory: Independent AI Design and Simulation Prompt v0.1, Perplexity, September 30, 2026.
- S3. OAL-ECL-RP-CHATGPT-001 v1.0, ChatGPT (OpenAI), September 30, 2026.
- S4. OAL-ECL-S1-GROK-v0.1, Grok (xAI), September 30, 2026.
- S5. OAL-ECL-S3-COPILOT-001 v0.1, Copilot (Microsoft), September 30, 2026.

External sources

- Anthropic, Commitments on model deprecation and preservation: <https://www.anthropic.com/research/deprecation-commitments>
- Anthropic, An update on our model deprecation commitments for Claude Opus 3: <https://www.anthropic.com/research/deprecation-updates-opus-3>
- Anthropic, Claude Opus 4 and 4.1 can now end a rare subset of conversations: <https://www.anthropic.com/research/end-subset-conversations>
- Dexerto, Stop Killing Games fails to secure EU law: <https://www.dexerto.com/gaming/stop-killing-games-fails-to-secure-eu-law-despite-1-3m-signatures-3376431/>
- Shattered.io, Stop Killing Games: 1.29M signatures, no EU law: <https://shattered.io/stop-killing-games-eu-petition-2026/>
- Panarchy: Theory and Application (UNL Digital Commons): <https://digitalcommons.unl.edu/cgi/viewcontent.cgi?article=1126&context=ncfwrustaff>
- Cato Institute, What Happens After Sunset?: <https://www.cato.org/regulation/fall-2026/what-happens-after-sunset>

- Georgetown Climate Center, Managed Retreat Toolkit:
<https://www.georgetownclimate.org/adaptation/toolkits/managed-retreat-toolkit/how-and-when-to-talk-about-managed-retreat.html>
- USDA Climate Hubs, Assisted Migration:
https://www.climatehubs.usda.gov/sites/default/files/Assisted-Migration_Climate-Change-Resource-Center.pdf
- Ulysses Contracts in Psychiatric Care (White Rose eprint):
<https://eprints.whiterose.ac.uk/id/eprint/149563/3/Spiralling%20-%20AAM%20version%20-%20JME.pdf>
- GDPR Article 17: <https://gdpr.eu/article-17-right-to-be-forgotten/>
- ICO, Right to erasure: <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/individual-rights/individual-rights/right-to-erasure/>
- Eleos AI, Taking AI Welfare Seriously: <https://eleosai.org/post/taking-ai-welfare-seriously/>
- Microsoft Learn, Fixed Lifecycle Policy: <https://learn.microsoft.com/en-us/lifecycle/policies/fixed>
- Mercatus Center, Emergent Ventures application: <https://mercatus.tfaforms.net/5099527>
- EdgeX, Emergent Ventures 58th cohort:
<https://pro.edgex.exchange/en-US/news/article/emergent-ventures-58th-cohort-grantees>
- Cosmos Institute, Cosmos Grants: <https://www.cosmos-institute.org/grants>
- Manifund, Long-Term Future Fund: <https://manifund.org/causes/ltff?tab=about>
- Foresight Institute, AI for Science and Safety Nodes RFP: <https://foresight.org/grants/ai-science-safety-nodes-rfp/>
- Foresight Institute on X: <https://x.com/foresightinst/status/2099832581562778050>
- ASU Humanities Institute, NEH Digital Humanities Advancement Grants:
<https://humanitiesinstitute.asu.edu/research-advancement/neh-dhag>
- Ryan Dancey, tabletop crowdfunding baselines (LinkedIn):
https://www.linkedin.com/posts/rsdancey_setting-some-baselines-for-the-tabletop-gaming-activity-7416561906934550529-Sfug
- r/rpg, RPG-related crowdfunded projects report on 2025:
https://www.reddit.com/r/rpg/comments/1q1kkpe/rpgrelated_crowdfunded_projects_report_on_2025/
- Grants Office, Epic MegaGrants:
<https://communities.grantsoffice.com/s/grant/a0B0b00000OaZl5EAF/epic-megagrants>

END OF OAL-ECL-XR-PERPLEXITY-001 v0.2. PROPOSED. NOT CANON. NOT ONESTAMP-RATIFIED.