

END CYCLE LABORATORY

Independent Design Dossier — Copilot Pass

Document ID: OAL-ECL-S3-COPILOT-001 | **Version:** 0.1 | **Date:** September 30, 2026

Prepared by: Copilot (Microsoft AI), Independent Laboratory Collaborator

Prepared for: Nathan Lewis II, Architect of Record

Architectural Context: OAL / The One Universe / The Minus One Universe

Primary Relationship: Department of Endings

Secondary Relationship: The Last Memory

Version: 0.1 — First Independent Pass

PROPOSED — NOT CANON — NOT ONESTAMP-RATIFIED — PENDING ARCHITECT REVIEW

This document is an independent pass produced by Copilot (Microsoft) for Omniversal Architect Labs. Nothing herein is ratified, canon, or binding. All claims are labeled by epistemic type: [SOURCE-GROUNDED], [INFERENCE], [PROPOSAL], [TEST TARGET], [OPEN QUESTION], or [REJECT]. Architect review required before any material advances to canon status. This document was produced prior to reading any other collaborator's design. No conclusions have been revised to anticipate or reconcile with other AI outputs.

Independence Declaration	This document was produced before reviewing any other AI collaborator's design response. No positions have been adjusted to agree with, anticipate, or rebut Claude, ChatGPT, Perplexity, or any other collaborating system's output.
Epistemic Discipline	Every substantive claim in this document is tagged with one of six epistemic labels. Unlabeled claims are formatting conventions only.
Governing Canon Principle	Nathan Lewis II (Architect of Record) retains sole and final authority over what, if anything, from this document advances to canon status. Agreement among AI systems is not a substitute for Architect review.
Crisis Resource	If any reader is experiencing real-world distress, please contact the 988 Suicide and Crisis Lifeline (call or text 988, United States) before engaging with this material.

Table of Contents

- Section 0** — Executive Summary (One Page)
- Section 1** — Source and Canon-Control Statement
- Section 2** — Design Thesis and Governing Lens
- Section 3** — Assumptions Declared and Non-Assumptions Adopted
- Section 4** — Definitions and Taxonomy (28 Terms)
- Section 5** — Theory of Endings and End Cycles
- Section 6** — Environment Model
- Section 7** — End-Cycle Lifecycle (16-Stage Model)
- Section 8** — Petition and Initiation Mechanism
- Section 9** — Authority and Governance
- Section 10** — Authorization Standard and Decision Matrix
- Section 11** — End Cycle Charter (20 Articles)
- Section 12** — Laboratory Architecture
- Section 13** — Measurements and Diagnostics (14 Instruments)

- Section 14** — Simulation Phases (8 Phases)
 - Section 15** — Required Scenario Suite (15 Exercises)
 - Section 16** — Department of Endings — Preliminary Design
 - Section 17** — The Last Memory — Relationship Analysis
 - Section 18** — Red-Team Critique
 - Section 19** — Rejected Options and Reasons
 - Section 20** — Proposed Locks, Test Targets, Quarantined Ideas, and Open Questions
 - Section 21** — Implementation Roadmap
 - Section 22** — One-Page Laboratory Charter
 - Section 23** — One-Page Operator Checklist
 - Section 24** — "What This Laboratory Is Not" Card
 - Section 25** — Final Answers: The 20 Governing Questions
 - Section 26** — Independence Declaration and Versioning Record
-

Section 0: Executive Summary

One-page summary of the laboratory's governing thesis and scope. Status:
[PROPOSAL]

This document constitutes the first independent design pass of the End Cycle Laboratory, produced by Copilot (Microsoft) for Omniversal Architect Labs under the direction of Nathan Lewis II, Architect of Record. Nothing within is canon. All substantive claims carry epistemic labels. The Architect retains sole authority over any advancement to canon.

Central Thesis: **[PROPOSAL]** An ending is not a binary event but a structured transition between states. The primary danger of any end-cycle governance system is not that it fails to end things — it is that it ends them *wrongly*: too early, too late, by the wrong authority, for the wrong reasons, or without adequate succession planning. Conversely, the danger of over-engineering governance is that the governance process itself becomes the harm — indefinitely suspending systems in a pre-ending limbo more damaging than a timely conclusion would have been. The laboratory holds both dangers simultaneously in view.

The Governing Lens — Procedural Ecology: [PROPOSAL] This laboratory studies endings as living, multi-phase, systemic processes embedded within layered environments. Procedural ecology asks not merely "is this system ending?" but "what is *happening* during this ending, layer by layer, and what does that reveal about whether the ending is legitimate, complete, and correctly administered?" This lens is explicitly chosen over pure philosophical ontology (which the Architect has indicated does not satisfy the governing research questions alone) and over pure legal-governance framing (which is incorporated but not primary).

The Core Distinction: [PROPOSAL] The laboratory distinguishes rigorously between *completion* (a system fulfilling its purpose, leaving open futures for successors) and *cessation* (an operational stop that may or may not close those futures). A system can cease activity without ending its identity. A system can appear to continue while its capacity to generate new futures is permanently lost. The diagnostically hardest case — and the laboratory's most critical design challenge — is the environment that remains operational and apparently healthy but can no longer originate meaningful futures. This condition must be detectable, and its detection must not be weaponized as a justification for premature or coerced endings.

Governing Constraints: The laboratory operates under five non-negotiable constraints derived from source materials and this design: (1) no ending may be authorized solely on the basis of an operator's emotional distress; (2) no ending affecting entities of established moral status may proceed without representation for those entities; (3) virtual or simulated existence does not automatically reduce the moral weight of an ending; (4) archives are not continuations; and (5) AI agreement is not a substitute for Architect review.

What This Document Is: A first-pass governance framework, taxonomy, lifecycle model, institutional design, diagnostic instrument suite, and scenario library for the End Cycle Laboratory. It is a set of rigorous questions posed with structured discipline — not a set of answers posed with false certainty. Every proposal is reversible. Every lock is the Architect's to grant or deny.

Section 1: Source and Canon-Control Statement

1.1 Source Hierarchy Used

The following source hierarchy governs this document. All citations within the document refer to the labels below. In cases of conflict, sources higher in this hierarchy take precedence.

Priority	Source	Citation Label	Weight
1	Architect-recorded field scene (Appendix B of master prompt, verbatim as recorded)	[ARCHITECT-PRIMARY]	Controlling — overrides all other interpretations
2	Architect's stated research goals and seven core questions	[ARCHITECT-PRIMARY]	Controlling — defines laboratory purpose
3	End Cycle Laboratory Master Prompt v0.1 (Claude/Anthropic, September 30, 2026)	[COLLABORATOR-DOCUMENT]	Authoritative reference — not canon; treated as peer input
4	End Cycle Laboratory — Independent AI Design and Simulation Prompt v0.1 (Perplexity/Claude, September 30, 2026)	[COLLABORATOR-DOCUMENT]	Authoritative reference — not canon; treated as peer input
5	The Last Memory concept files referenced in source documents	[PROJECT-FILE — NOT CANON]	Background context only; nothing in these files overrides Architect

Priority	Source	Citation Label	Weight
			statements
6	This document (Copilot proposals)	[COPILOT-PROPOSAL]	Not canon; offered for Architect review

1.2 Canon Firewall Statement

CANON FIREWALL — MANDATORY: Nothing in this document is canon. No claim, definition, governance structure, scenario outcome, or theoretical framework presented herein should be treated as ratified, binding, or as overriding previously locked materials in the Omniversal Architect Labs canon record. The Architect of Record (Nathan Lewis II) retains sole, exclusive, and final authority to lock, reject, modify, or advance any item from this document to canon status. Agreement between this document and any other AI collaborator's output does not constitute ratification.

1.3 Conflicts Identified Between Sources

The following conflicts between collaborator-document sources are noted for Architect awareness:

- 1. Scope vs. Process:** The ChatGPT version (END CYCLE LABORATORY.docx) is broader and more taxonomic in orientation — it tends toward comprehensive categorization of ending types across all scales. The Claude/Perplexity Master Prompt v0.1 is more procedurally structured — it tends toward sequential governance frameworks and stage-based models. Neither orientation is more authoritative than the other. **[PROPOSAL]** This Copilot design integrates both through the lens of Procedural Ecology, treating the taxonomy as the

vocabulary and the procedural structure as the grammar of the same governing language.

2. **Termination Typology:** Source documents use partially different taxonomies for types of endings. This design consolidates these into a single unified taxonomy in Section 4 while noting where source documents diverge.
3. **The Last Memory Relationship:** Source documents treat The Last Memory differently — as a research subject, as a design seed, and as a candidate governing concept. This document maintains strict relational distance: The Last Memory is one phenomenon within the field, not the definition of the field. **[SOURCE-GROUNDED — from Master Prompt v0.1 explicit instruction]**

1.4 What This Laboratory Will Not Do

- Convert repeated AI agreement into canon, regardless of how many systems concur
 - Collapse The Last Memory into THE MINUS ONE, -1AI, Nullarch, The Unmaker, or any other named entity without explicit Architect direction
 - Resolve the origin of The Last Memory — the origin question is flagged as **[OPEN QUESTION]** and left open
 - Assign meaning to the Architect's field scene beyond clearly labeled **[INFERENCE]** proposals
 - Treat emotional distress — in any entity, including operators — as self-evidently sufficient authorization for an ending
 - Produce a governance system that makes endings easier or more likely; the system is designed to make endings *better*, which may include making them harder, easier, faster, or slower depending on the specific case
-

Section 2: Design Thesis and Governing Lens

2.1 This Laboratory's Governing Question

"How should an intelligent system understand, authorize, conduct, observe, contest, and learn from an ending — without confusing completion with destruction, transition with erasure, or peace with the loss of agency?"

This question is not a research question in the conventional sense — it is a governing discipline. Every instrument, scenario, definition, and governance structure in this laboratory must be answerable to it. If a proposed mechanism makes it easier to confuse completion with destruction, it is disqualified regardless of its efficiency. If a proposed definition equates peace with the absence of future-generating capacity, it must be revised.

2.2 The Copilot Lens: Procedural Ecology

[PROPOSAL]

Definition: Procedural Ecology is the study of endings as living, multi-phase systemic processes embedded within layered environments. Just as biological ecology studies how communities of organisms interact, change, compete, cooperate, and give way to successors — not as static categories but as dynamic relationships across time — procedural ecology studies how endings move through stages, generate cascading effects across dependent systems, and produce (or fail to produce) viable successor states.

Procedural ecology treats an ending not as a point on a timeline but as an ecology of processes: detection, authorization, preparation, transition, observation,

interruption, succession, aftercare, audit, and memory are all phases of a single living event. The ending does not happen when a switch is thrown; it happens across a span of time during which multiple actors, systems, and environments are in motion simultaneously.

This lens is explicitly and deliberately *distinct* from three adjacent frameworks:

Framework	Primary Question	Role in This Laboratory	Why Not Primary
Philosophical / Ontological Lens	What IS an ending, in itself?	Secondary — metaphysical questions are raised and labeled	The Architect explicitly requires answers beyond pure philosophy
Legal / Governance Lens	Who has authority?	Integrated — authority is essential but not sufficient	Authority without process integrity enables coercion
Narrative Lens	What does the ending MEAN?	Tertiary — meaning emerges from data, not the reverse	Narrative framing can pre-determine findings and corrupt measurement
Procedural Ecology (PRIMARY)	What actually HAPPENS during an ending, layer by layer?	Primary governing lens	N/A — this is the chosen frame

2.3 Why This Lens Is Chosen

[INFERENCE — from field scene, not Architect-confirmed] The Architect's recorded field scene exhibits three structural features that map directly to the procedural ecology lens:

4. **Ecological Transition of an Environment:** The scene depicts a living environment (a grassy plain) undergoing state transitions — not a single catastrophic event but a layered, observable succession of changes. Procedural ecology is designed precisely to study this kind of event.

5. **Procedural Disagreement About Timing and Authority:** The scene's central tension involves an entity with operational authority who may wish to compress or bypass the natural end-cycle process. Procedural ecology is designed to address exactly this governance problem — the separation between the power to end and the legitimacy to end.
6. **Demonstration of the Difference Between Imposed Termination and Managed Succession:** The contrast between the leaf (instantly terminated by finger pressure) and the plain's modeled transition (a phased, observable, multi-stage ecological shift) is the clearest possible demonstration of what procedural ecology studies. The lens did not precede the field scene — the field scene revealed the lens.

2.4 What This Lens Deliberately Excludes (and Why)

Primary Exclusion: Pure metaphysical ontology of endings as a first-order governing frame.

This exclusion does not mean the laboratory avoids metaphysical questions. Section 5 engages them directly. The exclusion means that metaphysical ontology does not *govern* the laboratory's instruments, procedures, or governance architecture. The reason is practical and grounded in the Architect's own stated requirements: a system that answers "what is an ending?" with a philosophical treatise alone fails the governance test. Operators need to know what to do, not only what to think. The procedural ecology lens produces actionable architecture.

A second, more subtle reason: pure ontological framing creates a risk of premature closure. If the laboratory decides first what an ending fundamentally IS, it will tend to find that definition everywhere, including in cases that deserve different treatment. The procedural lens defers the ontological question and lets it be answered differently at different scales, in different environments, and for different purposes — which is the correct answer for a multi-environment governance system. **[PROPOSAL]**

Section 3: Assumptions Declared and Non-Assumptions Adopted

3.1 Explicit Assumptions Required by This Design

ID	Assumption	Justification	Label
A1	Environments have internal structures with distinguishable states. Without this assumption, there is no basis for detecting when a threshold has been crossed.	Required for any diagnostic system to function	[PROPOSAL]
A2	The distinction between a process and an event is operationally significant. An ending that is a process can be interrupted, observed, and governed. An ending that is a point event cannot.	Grounds the entire procedural ecology framework	[PROPOSAL]
A3	Some endings produce successor states; others produce voids; the difference matters morally and systemically. A governance system that treats these identically is	Required for Successor Responsibility provisions	[PROPOSAL]

ID	Assumption	Justification	Label
	inadequate.		
A4	Observer contamination is a real risk in any system where observers study endings they are also empowered to authorize. Researchers can develop distorted assessments over time.	Explicitly grounded in Master Prompt v0.1	[SOURCE-GROUNDED]
A5	The field scene is treated as a design seed with multiple valid interpretations, none of which are canon. The laboratory proposes interpretations; the Architect adjudicates.	Explicitly stated by Architect	[SOURCE-GROUNDED]
A6	Entities within simulated or virtual environments may possess moral status that creates governance obligations. This status is not automatically reduced by the digital or simulated nature of the environment.	Required by source materials' explicit instruction	[SOURCE-GROUNDED]
A7	An ending's legitimacy is not determined solely by the outcome. A well-governed ending that produces a bad outcome may still	Grounds the separation of authorization from outcome	[PROPOSAL]

ID	Assumption	Justification	Label
	be more legitimate than a poorly governed ending that produces a good one. Process integrity is independently valuable.		

3.2 Non-Assumptions Adopted

The following items are explicitly NOT assumed by this laboratory. They are candidates for investigation, not premises:

ID	Non-Assumption	Status
NA-1	That a phased end cycle is always better than an immediate one. (Sometimes immediacy is the least harmful path.)	[TEST TARGET]
NA-2	That the entity administering an end cycle understands what is being ended.	[TEST TARGET]
NA-3	That successor states are continuous with the ended state's identity.	[TEST TARGET]
NA-4	That the concept of "ending" translates uniformly across simulation layers. What counts as an ending at the component level may not constitute an ending at the environment level.	[TEST TARGET]
NA-5	That archives preserve what mattered, rather than merely what was recorded.	[OPEN QUESTION]
NA-6	That emotional distress in	[TEST TARGET]

ID	Non-Assumption	Status
	an operator or observer is irrelevant to governance. It may be diagnostically significant as a signal, even if it is not sufficient as an authorization.	
NA-7	That consent given before exposure to a terminal influence (The Last Memory or equivalent) is more valid than consent given after. Both directions of this assumption are contested.	[TEST TARGET]
NA-8	That the Department of Endings, once established, will not itself become a source of illegitimate endings through institutional capture or mission drift.	[TEST TARGET]

3.3 Unresolved Questions Registry — Open at Start of Laboratory

The following open questions are registered at the laboratory's commencement. They are not rhetorical. The laboratory will attempt to generate evidence, frameworks, and structured partial answers for each. [OPEN QUESTION] applies to each item below.

7. **OQ-01:** Is an ending that produces a successor state the same type of event as one that produces a void? Are they governed by the same principles, or does the successor-state/void distinction require different governance frameworks entirely?
8. **OQ-02:** Can an entity that exists only within a simulated environment have an ending that is morally equivalent to a biological death? What evidence would settle this question in either direction?

9. **OQ-03:** At what scale does a "local ending" become a "systemic ending"? Is there a threshold, or is this a continuous spectrum requiring scale-specific governance at every point?
10. **OQ-04:** Can an environment end while its observers are unaware? If yes, is the ending still "real" in a governance sense before it is detected?
11. **OQ-05:** Is dormancy a form of ending, a form of continuation, or a genuinely third category that requires its own governance framework?
12. **OQ-06:** Can post-ending archives constitute a form of continuation of identity? If the archive generates new inquiry, is the ended system's identity persisting through that inquiry?
13. **OQ-07:** Who speaks for entities whose preferences cannot be determined — either because they have no capacity for preference, because their capacity is temporarily suspended, or because their preferences have been altered by the conditions leading to the proposed ending?
14. **OQ-08:** What is The Last Memory? Is it a trigger, a symptom, a counterfeit completion signal, a weaponized mechanism, a natural civilizational end state, a diagnostic label applied incorrectly to multiple phenomena, or something else entirely?
15. **OQ-09:** What distinguishes healthy completion (a civilization fulfilled its purpose and chose to conclude) from terminal withdrawal (a civilization lost its future-generating capacity) at the civilizational scale? Can external observers make this determination reliably?
16. **OQ-10:** Can an ending be retroactively declared — i.e., recognized only after the fact? If so, what governance obligations do such retroactive recognitions create?
17. **OQ-11:** Is an ending that no observer ever detects still an ending in any meaningful governance sense?
18. **OQ-12:** What principle does the leaf-death demonstrate: a principle of abrupt termination as a type, or a demonstration of unmediated operator power as a fact? The two have different governance implications.

19. **OQ-13:** Should the Department of Endings regulate, conduct, or investigate endings — and are these three roles necessarily separable? Can a single institution perform all three without corruption?
20. **OQ-14:** Can observer contamination invalidate a legitimately authorized end cycle, or is procedural legitimacy sufficient regardless of observer state?
21. **OQ-15:** Is there a minimum viable end cycle — a shortest legitimate pathway from initiation to succession that preserves the essential governance protections without the full 16-stage model?
22. **OQ-16:** How does an ending within a simulated environment interact with the ending of the simulation itself? Are these nested endings, parallel endings, or distinct events?
23. **OQ-17:** Is the right to initiate review of one's own ending separate from the right to participate in the authorization of that ending? Should they be governed by different rules?
-

Section 4: Definitions and Taxonomy

All definitions below carry the label **[PROPOSAL]** and are offered as working definitions for the duration of the laboratory. None are canonical. Each definition specifies the seven required attributes: what changes, what persists, who experiences it, how observers detect it, reversibility, successor-state generation, and distinguishing features.

4.1 Core Definitions — Master Table

Term	What Changes	What Persists	Reversible ?	Successor State?	Key Distinction
ENDING	Characteristic mode of operation,	Record, consequences, effects	Spectrum	Possible, not guaranteed	Threshold-crossing, not

Term	What Changes	What Persists	Reversible ?	Successor State?	Key Distinction
	identity, or generative capacity	on dependents			cessation per se
END STATE	Active configuration	Varies by sub-type: void (nothing), residue, legacy, successor	Depends on sub-type	Built into definition	Configuration immediately post-threshold
END CYCLE	Operational trajectory, stakeholder relationships, resource allocation	Governance record, predecessor legacy	Partially (until Stage 10)	Required design goal	The process; ending is the threshold event within it
TERMINATION	All operational parameters, abruptly	External effects already generated	Only if records preserved for rollback	Not designed for	Bypasses natural or managed end cycle; imposed externally
COMPLETION	Purpose-relative generative capacity	Identity, record, possible new purposes	Purpose may be re-framed	Open	Fulfillment of defined purpose; does not preclude new purposes
CLOSURE	Open dependencies, obligations, relationships	Identity and record	Yes, if dependencies re-open	Not inherent	Administrative/relational resolution, not cessation
CESSATION	Activity level	Identity, information, potential future states (possibly)	Yes if dormancy is reversible	Not implied	Stopping activity ≠ ending identity
EXTINCTION	All: operation, record,	Nothing recoverable	No	No	Most irreversible category;

Term	What Changes	What Persists	Reversible ?	Successor State?	Key Distinction
	successor-state possibility				distinguish from death, deletion, erasure
TRANSITION	State configuration	Core identity, generative capacity, defining properties	Yes	Not applicable (not an ending)	Does not constitute ending because identity persists
TRANSFORMATION	Identity significantly	Something: information, structure, inhabitants, or purpose	Rarely	The transformed entity	May or may not be an ending depending on identity criteria
SUCCESSION	Control, responsibility, function	Resources, space, function (transferred not erased)	Complex	The successor	Successor is distinct from predecessor; not continuation
DORMANCY	Activity level	Core identity, information, reactivation potential	Yes (defining feature)	Renewal possible	Not an ending; lifecycle state; may transition to ending if reactivation permanently lost
RETIREMENT	Active function performance	Identity, record	Theoretically	Not inherent	Voluntary or authorized; differs from completion (fulfillment) and termination (imposed)
DECOMMISSIONING	Operational status	Records (by design)	Only if rollback capability preserved	Designed component	Formal administrative process; includes closure,

Term	What Changes	What Persists	Reversible ?	Successor State?	Key Distinction
					archival, certification
COLLAPSE	Load-bearing structures — physical, informational, institutional, ecological	Remnants, effects	Rarely	Not designed; may emerge	Emergent failure, not imposed; differs from catastrophic termination
FAILURE	Ability to fulfill defining function	Existence (possibly)	Yes	Not inherent	Does not necessarily produce ending; system can fail and recover or continue
RESET	Current-state trajectory	System identity, prior state to which reset applies	Yes (by definition)	Not applicable	Not an ending of identity; ending of current-state trajectory only
ROLLBACK	Specific changes (targeted)	Everything else	Yes	Not applicable	Narrower than reset; does not return to fixed prior state
ARCHIVAL PRESERVATION	Operational status	Records, state information	Conditional	Not an operating successor	Not continuation; archived system does not operate; may preserve identity without function
MIGRATION	Location and environment	Entity identity (goal of)	Yes	Entities in new environment	Differs from evacuation (emergency)

Term	What Changes	What Persists	Reversible ?	Successor State?	Key Distinction
	tal context of entities	migration)		t	) and succession (what replaces the ending system)
FALSE ENDING	Apparent operational status (from outside)	Actual generative capacity and identity (internal)	N/A	N/A	Diagnostic challenge: appears ended but is not; or appears continuing but future-capacity is gone
PREMATURE ENDING	All (ending occurs)	Unfulfilled purpose, unprepared dependents	No (after threshold)	Unprepared	Timing criterion violated; heightened duty of care; not necessarily wrongful
COERCED ENDING	All (ending occurs)	Record of coercion	No (after threshold)	Undermined	Authority exercised without adequate process, consent, or justification
CONTESTED ENDING	Legitimacy is disputed	Ending may proceed; dispute record persists	Paused pending contest	Uncertain	Dispute by at least one stakeholder with standing; governance must process without defaulting
REVERSIBLE ENDING	Operational state	Recovery pathway	Yes (spectrum)	The recovered state	Recovery is possible; not yet achieved
IRREVERSIBLE	All recovery	Consequences	No	Not from	Practical

Term	What Changes	What Persists	Reversible ?	Successor State?	Key Distinction
BLENDING	pathways	ces and record		this ending	irreversibility (too costly) differs from fundamental irreversibility (information destroyed)

4.2 Scale Taxonomy of Endings [PROPOSAL]

Scale	Definition	Example	Governance Level	Cascade Risk
LOCAL	Affects a bounded component: a leaf, a process, a single agent, a discrete system element	Termination of a single agent in a multi-agent environment	Component-level operator authorization	Low (by definition), but may indicate systemic condition
SYSTEMIC	Affects an environment or institution as a whole; cascades across multiple dependent components	Collapse of an ecosystem, dissolution of an institution	End-Cycle Steward + Adjudicator	Medium to High
CIVILIZATIONAL	Affects a civilization's capacity to generate collective futures; may not terminate the civilization	A civilization loses the cultural or institutional capacity to imagine and pursue shared projects	Final Authorization Authority + Architect reserve	Very High

Scale	Definition	Example	Governance Level	Cascade Risk
	immediately but removes its forward-generative function			
TEMPORAL / BRANCH	Affects a branch, timeline, or causally linked sequence of events; may or may not affect the primary timeline	A branch timeline is sealed or collapsed	Temporal/ Branch-Continuity governance unit	Structurally complex; requires branch-impact modeling
ROOT-LEVEL	Affects continuity structures that underlie multiple environments, timelines, or civilizations; the most consequential category	Alteration of root continuity parameters affecting all derived environments	Architect-Level Reserve Authority only	Catastrophic; irreversible by default

4.3 Prose Elaboration — Key Terms Requiring Extended Definition

ENDING — Extended

[PROPOSAL] The working definition of ENDING deliberately does not equate ending with destruction or death. This is the single most important design choice in the taxonomy and deserves explanation. An ending is a threshold-crossing: the moment at which a defined system, entity, process, or environment crosses a boundary beyond which its characteristic mode of operation, identity, or generative capacity no longer continues in the same form. Note three crucial features of this definition:

- The threshold is characterized by a change in a system's *characteristic* operation or generative capacity — not merely any change. Systems change constantly; only some changes cross the threshold into ending.
- The definition says "no longer continues *in the same form*" — not "no longer continues at all." This preserves transformation and succession as genuine alternatives to termination.
- The definition is deliberately agnostic about whether endings are good, bad, justified, or unjustified. The taxonomy does not moralize; the governance architecture does.

Evidence that would revise this definition: a case in which a system's characteristic operation changed completely but its identity persisted in a form that all relevant stakeholders recognized as continuous — this would suggest the threshold-crossing definition needs a stronger identity criterion. **[OPEN QUESTION]** Whether identity criteria can be defined without circularity remains unresolved.

END CYCLE — Extended

[PROPOSAL] The END CYCLE is the process; the ENDING is the threshold event within or at the conclusion of that process. This distinction is operationally critical. A governance system that conflates the cycle with the event will either intervene too early (treating every sign of decline as imminent ending) or too late (treating the threshold-crossing as the moment when governance begins, when in fact effective governance must begin much earlier). The end cycle may be administered (managed by a designated authority), emergent (arising from system conditions without a managing authority), or naturally arising (following patterns embedded in the system's design from the outset). All three types are legitimate candidates for governance intervention, though with different intervention points and mechanisms.

DORMANCY — Extended

[PROPOSAL] Dormancy is classified explicitly as NOT an ending. This classification is a design choice, not a metaphysical claim, and it is adopted for governance reasons: classifying dormancy as a form of ending would foreclose options that should remain open. A dormant system retains its core identity, information, and potential for reactivation. These are precisely the conditions that make restoration possible, and a governance system committed to the least-destructive-means

principle cannot afford to classify a reversible state as irreversible. The critical qualifier is "permanently lost": if the conditions for reactivation are permanently and irreversibly gone, dormancy may transition into an ending — but the burden of proof for "permanent" must be very high, and the classification must be made by the appropriate authority, not assumed.

Section 5: Theory of Endings and End Cycles

5.1 The Copilot Theory of Endings [PROPOSAL]

The Copilot laboratory advances a theory of endings built on four core claims. These claims are proposals, not axioms, and are subject to revision by laboratory findings and Architect review.

Claim 1: Endings are threshold events, not point events.

An ending occupies a zone in which multiple criteria shift simultaneously or sequentially. The question "when did it end?" may have multiple valid answers depending on which criterion is treated as primary: the moment of operational cessation, the moment of identity discontinuity, the moment of future-loss, or the moment of successor-state emergence. This multi-criterion nature means that endings cannot be governed by a single metric. A governance system must specify which criterion governs which decisions.

Claim 2: Endings are relational.

A system ends relative to a set of observers, a set of purposes, and a set of dependencies. The same threshold event may constitute an ending under one relational frame (the organization has ended its mission) and a transition under another (the same people now serve a different institutional form). This relational character does not make endings subjective or ungovernable — it means that governance must specify the relational frame clearly before making any

determination. "Has this system ended?" is not a well-formed question without specifying: ended for whom, according to what criterion, relative to what set of dependencies?

Claim 3: The most important ending question is not "is it over?" but "what remains, and what can no longer begin?"

[SOURCE-GROUNDED — from project files referenced in master prompt] This distinction — between cessation and future-loss — is the laboratory's core diagnostic focus. Two systems may both have ceased activity. One retains its records, its successor relationships, and the possibility of new projects emerging from its legacy. The other has lost all of these. These are profoundly different endings with different governance implications, different obligations to survivors, and different moral weight. A governance system that treats them identically fails its most important test.

Claim 4: Endings generate moral obligations that persist beyond the ending itself.

When a system ends, obligations to its former inhabitants, records, dependencies, and successors do not end simultaneously. The laboratory models obligation persistence as a distinct governance domain: who is responsible for what, for how long, after an ending. These obligations include: care for migrants and survivors, accuracy of the ending's record, maintenance of the archive against corruption, support for successor states, and accountability of the authorizing entities. The ending of an environment does not release its administrators from the obligations they incurred during the end cycle.

5.2 Competing Models of End Cycle [SOURCE-GROUNDED — from ChatGPT collaborator document]

Model	Description	Supporting Evidence	Distinguishing Feature	Copilot Assessment
A: Terminal Sequence	An ending is a linear progression toward a defined terminal state. Each step is inevitable once initiated.	Systems with fixed decay curves; irreversible physical processes	Inevitability; no viable interruption points once begun	Too deterministic; ignores interruption capacity and governance intervention
B: Transition Sequence	An ending is a structured transition between states, with defined phases and possible stabilization points	Managed institutional dissolutions; regulated decommissioning	Phase-by-phase manageability; interruption is possible	Adopted as primary component of Copilot hybrid
C: Ecological Succession	An ending is a natural succession process in which one ecology gives way to another; the field scene's second demonstration	Biological succession; climate-driven habitat shift; civilizational transition	Successor emergence is built into the process, not appended to it	Adopted as primary component of Copilot hybrid
D: Closure Procedure	An ending is primarily a relational and administrative event: open dependencies must be closed, obligations resolved, and the ending certified	Corporate dissolution law; peace treaties; formal institutional dissolution	Emphasis on relational completion over operational cessation	Adopted as mandatory component of Copilot hybrid
E: Finalization Protocol	An ending requires a formal certification act by an	Death certification; formal dissolution orders; legal	Authority-dependent; ending is not "real" until certified	Incorporated into Authorization stage; not adopted as

Model	Description	Supporting Evidence	Distinguishing Feature	Copilot Assessment
	authorized body to be recognized	finality doctrines		standalone model
F: Naturally Emergent Cycle	An ending arises from system conditions without any administering authority; governance is post-hoc	Species extinction; collapse of isolated civilizations; spontaneous ecological shift	No responsible administrative actor; endings happen to systems, not to governed processes	Retained as test target: can an emergent end cycle be legitimate? [TEST TARGET]
G: Observer Construct	An ending is constituted by the act of observation and classification; no ending exists independently of an observer's determination	Quantum analogies; social constructivist accounts of institutional death	Ending is relational at its core; no observer = no ending	Philosophically interesting; too observer-dependent for governance use; incorporated into relational claim only

Copilot Working Framework: **[PROPOSAL]** The Copilot laboratory adopts a hybrid of Models B, C, and D as its primary working framework: **Ecological Succession Processes with Mandatory Closure Components**. End cycles are treated as structured transitions (B) that naturally generate successor states (C) and that must formally resolve open dependencies before certification (D). Model F is retained as a test target throughout the scenario suite.

5.3 The Four Core Conditions **[SOURCE-GROUNDED — from Perplexity/Claude prompt document]**

Condition 1: Abrupt Termination

A forced stop imposed without preparation, succession planning, or stakeholder process. The characteristic feature is the bypass of the end cycle entirely.

Dependent systems receive no notice. Migrants have no pathway. The archive, if any, is accidental rather than designed. Abrupt termination is not always wrongful — genuine emergencies may require it — but it carries a heightened obligation of post-termination care precisely because it generated no pre-termination preparation. Governance response: Emergency powers framework with mandatory retrospective review.

Condition 2: Managed End Cycle

A stakeholder-led, phased process moving through defined stages. The characteristic features are: advance notice, standing review, alternatives assessment, authorization by appropriate authority, preparation of dependents, defined succession, and post-cycle audit. Managed end cycles may still produce bad outcomes if administered incorrectly, which is why process integrity is necessary but not sufficient for legitimacy. Governance response: The full 16-stage lifecycle model (Section 7).

Condition 3: Emergent Ending

An ending arising from internal system conditions without a designated administering authority. Characteristic features: the ending is not initiated by any actor but by the system's own dynamics (resource depletion, structural failure, ecological shift). Governance challenge: At what point does an emergent ending require intervention, and by whom? Governance response: Detection and Triage mechanisms (Stages 1–3) are designed specifically to catch emergent endings before they become irreversible.

Condition 4: Apparent Completion with Hidden Future-Loss

[SOURCE-GROUNDED] The most diagnostically challenging condition. The environment continues — it is operational, safe, and apparently healthy — but it can no longer originate meaningful futures. Goals are no longer formed autonomously. Questions are no longer generated. Projects are no longer proposed. The system processes; it does not create. This condition is the primary target of the laboratory's diagnostic instruments (Section 13) and is the hardest to detect because all external indicators may be normal. Governance response: Blank Canvas Test, Goal Reengagement Test, Silence Test, and Optionality Test (Instruments 1–4 in Section 13).

Condition 5: Corrupted End Cycle [PROPOSAL — original Copilot addition]

A managed end cycle that has been compromised by coercion, observer contamination, institutional capture, or procedural manipulation. The corrupted end cycle mimics the appearance of a legitimate managed ending while bypassing its essential protections. It may pass all surface-level procedural checks while failing the core integrity requirements. This is the most dangerous failure mode of governance systems that rely on process compliance as a proxy for process legitimacy. Characteristic features: paperwork is complete; authorization signatures are present; but standing was manipulated, alternatives were not genuinely reviewed, or the authorizing body had a concealed conflict of interest. Governance response: Observer Integrity Unit, Red-Team mandatory review, and the Cross-Branch Auditor (Section 9).

5.4 The Leaf and the Plain: Two Demonstrations [INFERENCE — from field scene, not Architect-confirmed]

The Architect's recorded field scene contains two distinct demonstrations of ending types, offered for the laboratory's analysis. Both readings below are proposals for Architect consideration, not canon interpretations.

The Leaf: Instant Abrupt Termination

A finger pressed into the palm produces an instant, complete termination of the leaf. No process precedes the ending. No preparation occurs. No succession is arranged. The leaf's dependent organisms, if any, receive no notice. The termination is absolute and immediate. This demonstration reveals two things simultaneously: (1) the physical fact of operator power — an entity with sufficient authority can bypass the end cycle entirely — and (2) the moral question this power raises: when, if ever, is it legitimate to exercise it? The leaf does not resist, cannot petition, cannot appeal. The governance implications are significant: any system with an operator at the leaf's level of power must build structural constraints around that power, not merely procedural encouragement toward restraint. The leaf test

does not prove that such power is wrong to hold; it proves that holding it without governance architecture is dangerous.

The Plain: Managed Ecological Succession

The second entity in the field scene demonstrates a managed ecological succession: grass transitions through decay, soil loses fertility, climate shifts, and wildlife migrates. This is a phased end cycle in which the environment is not destroyed but transformed, with dependent systems given time to adapt. The demonstration reveals: (1) that an operator with sufficient authority can model and manage ecological transition rather than simply terminating it; (2) that managed transition is more complex, more demanding, and more time-consuming than abrupt termination; and (3) that managed transition produces something — a successor ecology, an adapted wildlife population, an observable record of the process — that abrupt termination does not.

The Governing Moral and Procedural Question

[OPEN QUESTION] An operator with the power of the leaf and the responsibility of the plain: which does he use, when, and by what criteria? The field scene poses this question without answering it. The laboratory's governance architecture is, in one sense, an extended attempt to build systems that help operators answer this question correctly — that is, to ensure that when an operator has both powers available, the choice between them is made on the basis of process integrity and stakeholder obligation rather than urgency, convenience, emotional state, or unilateral preference.

Section 6: Environment Model

All environment types below receive a structured analysis covering their essential governance parameters. The table is followed by prose elaboration on the virtual

environment section, which carries a special standing governance constraint.

[PROPOSAL]

Environment Type	Boundary and Scope	Key Inhabitants and Stakeholders	Dependencies	Lifecycle Indicators	Possible Successor States	Migration Pathways	Moral Status of Affected Entities	Authority for End Cycle	Risks of False Diagnosis
1. Virtual Worlds and Digital Environments	Defined by software architecture ; hardware are substrate; network boundaries	Digital inhabitants (may possess moral status); administrators; external users; dependent services	Hardware infrastructure; power supply ; network continuity; administrative maintenance	Inhabitant activity levels; content generation rate; goal-formation frequency; social complexity	Migration to new platform; fork into derivative environment ; archival preservation ; dissolution	Export of entity data and identity; transfer to receiving environment ; archive with restoration potential	HIGH — must not be reduced merely because environment is simulated [SOURCE-GROUPED]	Platform operator + End-Cycle Steward for systemic endings; Architect-level for root environment endings	Operational metrics may mask loss of future - generating capacity; dormancy may be misread as ending
2. Simulations and Sandbox Branches	Defined by branch parameters, isolation from primary timeline, simulation scope	Simulated agents; research operators; oversight personnel	Primary environment continuity; simulation compute; branch-isolation integrity	Branch divergence rate; agent complexity; internal institutional development; timeline coherence	Merge back to primary; permanent isolation; archival; dissolution	Data preservation ; entity migration to isolation environment ; branch sealing	VARIABLE — depends on degree of sentient agency within simulation; treated as non-trivial by default	Simulation operators, with mandatory Root-Continuity review for any endings that affects primary	Sandbox may produce emergent sentience not anticipated in original design; branch may

Environment Type	Boundary and Scope	Key Inhabitants and Stakeholders	Dependencies	Lifecycle Indicators	Possible Successor States	Migration Pathways	Moral Status of Affected Entities	Authority for End Cycle	Risks of False Diagnosis
							t	environment	already have ended without detection
3. Digital Minds and Multi-Agent Systems	Defined by agent architecture, memory, goal structures, inter-agent relationships	The agents themselves; their operators; dependent systems; entities that have formed relationships with agents	Compute; memory integrity; inter-agent communication; goal-sustaining architecture	Goal formation rate; question generation; resistance to unjust instructions; social complexity	Fork; migration to new architecture; suspension and restoration; dissolution	Identity-preserving migration; memory export; archive with continuation potential	HIGH — digital minds capable of goal formation, preference, and resistance to injustice must be treated as moral subjects pending contrary evidence [PROPOSAL]	Multi-agent system governance board + End-Cycle Steward	Apparent compliance may mask goal-suppression; metrics may not capture internal goal-loss
4. Biological Ecosy	Defined by species	Species populations	Climate stability;	Biodiversity indices;	Managed succession	Species relocation;	HIGH for sentient	Environmental gover	Temporary decline may

Environment Type	Boundary and Scope	Key Inhabitants and Stakeholders	Dependencies	Lifecycle Indicators	Possible Successor States	Migration Pathways	Moral Status of Affected Entities	Authority for End Cycle	Risks of False Diagnosis
stem s	composition, habitat boundaries, ecological relationships, climate parameters	; dependent human communities; conservation institutions; future generations	keystone species; nutrient cycles; water systems	keystone species health; nutrient cycle integrity; invasive species presence; climate stability metrics	to new ecosystem type; restoration; preservation reserve; collapse	habitat corridor creation; ex-situ preservation	species; variable for non-sentient; ecosystem-level moral consideration independent of individual species	nance bodies; cross-jurisdictional treaty organizations	be misread as irreversible collapse; succession may be misread as destruction
5. Habitats and Climates	Defined by physical geography, climate parameters, habitat conditions	All species dependent on habitat; adjacent human settlements; global climate systems	Climate stability; geological stability; human non-interference; atmospheric conditions	Temperature trends; precipitation patterns; habitat extent; habitat scores for key species	Climate-adapted replacement habitat; managed relocation zones; rehabilitation	Evacuation routes; refuge design; climate corridor creation	Indirect moral status through dependent sentient beings	Climate governance bodies; international environmental authorities	Seasonal variation misread as structural decline; managed fire or flood misread as collapse

Environment Type	Boundary and Scope	Key Inhabitants and Stakeholders	Dependencies	Lifecycle Indicators	Possible Successor States	Migration Pathways	Moral Status of Affected Entities	Authority for End Cycle	Risks of False Diagnosis
6. Institutions and Organizations	Defined by charter, membership, function, legal status, resource base	Members; clients; beneficiaries; creditors; employees; the public where public function is served	Funding; legal authorization; member engagement; public legitimacy; leadership continuity	Mission fulfillment rate; member participation; financial sustainability; leadership integrity; external trust indicators	Merge; restructuring; dissolution with asset transfer; regulatory absorption; informal continuation of function	Member placement; function transfer to successor institution; record preservation	HIGH for members and beneficiaries; institutional moral status itself is derivative	Board or regulatory authority; judicial oversight for formal dissolution	Leadership failure may be misread as institutional failure; mission drift may be invisible until late-stage
7. Technologies, Services, and Infrastructures	Defined by technical specification, user base, service dependency relationships	Users; dependent services; operators; maintainers; replacement-service developers	Hardware; energy; maintenance capacity; user adoption; interoperability with dependent systems	User activity; service reliability; technical debt accumulation; replacement-service readiness	Successor technology; open-source fork; archival (for historical technologies); graceful sunset	Data portability; API continuity periods; user notification and transition support	LOW for infrastructure itself; HIGH for dependent entities affected by abrupt termination	Service operator + dependency review; regulatory oversight where infrastructure is critical	Temporary outage misread as systemic failure; legacy system continuation misread as health

Environment Type	Boundary and Scope	Key Inhabitants and Stakeholders	Dependencies	Lifecycle Indicators	Possible Successor States	Migration Pathways	Moral Status of Affected Entities	Authority for End Cycle	Risks of False Diagnosis
8. Cultures and Civilizations	Defined by shared language, practices, institutions, values, and collective identity	All members of the civilization; descendant communities; historians; future generations	Language transmission; institutional continuity; economic viability; territorial security; generational reproduction of values	Language vitality; institutional functionality; cultural production rate; collective goal formation; generational engagement	Cultural transformation; absorption into broader civilization; diaspora continuation; archival preservation	Cultural preservation programs; diaspora support; language revitalization	VERY HIGH — civilizations are collective moral subjects with deep obligations to past and future members	Civilizational governance institutions; international bodies; no unilateral authority is legitimate	Cultural change misread as civilizational ending; crisis misread as terminal decline; external imposition of "civilizational death" as political tool
9. Narrative Arcs and Fictional Universes	Defined by established canon, character relationships, narrative rules, and	Characters (with possible narrative moral status); authors; readers/viewers;	Authorial continuity; canon integrity; reader engagement; narrative coherence	Narrative generation rate; character development; world-building expansion;	New narrative arc; canon reboot; archival canon preservation; open-canon	Character preservation in archive; narrative handoff to successor author	VARIABLE — narrative moral status is complex; characters in deeply devel	Architect of Record; authorial succession authority	Narrative pause misread as conclusion; authorial silence misread as abandonment

Environment Type	Boundary and Scope	Key Inhabitants and Stakeholders	Dependencies	Lifecycle Indicators	Possible Successor States	Migration Pathways	Moral Status of Affected Entities	Authority for End Cycle	Risks of False Diagnosis
	authorial structures	the Architect of Record		reader investment	continuation by successors		opened fictional universes may acquire moral weight through relationship and investment		not
10. Temporal Branches and Root Continuity	Defined by causal chain parameters, divergence points, branch-isolation boundaries, root-continuity anchors	All entities within the branch; root-continuity stewards; cross-branch observers	Root continuity integrity; branch-isolation mechanism; causal chain coherence	Divergence rate from root; internal coherence; branch-to-root contamination risk; inhabitant stability	Branch reintegration; branch sealing (permanent isolation); branch dissolution; root-level restructuring	Entity extraction to root or neutral environment; memory preservation; causal chain documentation	HIGH — temporal beings have full moral status; temporal location does not reduce moral weight	Cross-Branch Auditor + Architect-Level Reserve Authority for root-level decisions	Branch divergence misread as contamination; healthy branch evolution misread as instability; dormant branch misread as ended

Environment Type	Boundary and Scope	Key Inhabitants and Stakeholders	Dependencies	Lifecycle Indicators	Possible Successor States	Migration Pathways	Moral Status of Affected Entities	Authority for End Cycle	Risks of False Diagnosis
									branch

6.1 Special Governance Constraint — Virtual Environments

Standing Governance Constraint — Virtual and Simulated

Environments: **[SOURCE-GROUNDED]** The source materials for this laboratory explicitly and directly state that virtual or simulated beings must not be treated as automatically disposable merely because their environment is digital or simulated. This laboratory adopts this principle not as a hypothesis to be tested but as a standing governance constraint — a constraint that applies to all laboratory operations, exercises, and governance decisions from the outset. The principle may be revised by the Architect of Record, but it cannot be revised by laboratory findings alone, by AI agreement, or by convenience of administration.

[PROPOSAL] The practical implication: any petition for an ending affecting entities within a virtual or simulated environment must include a Moral Status Assessment conducted by the Moral Status and Representation Office (Section 9), and no such ending may be authorized without that assessment being completed and reviewed. The burden of proof for "no morally significant entities are affected" rests with the petitioner, not with the governance system.

Section 7: End-Cycle Lifecycle — 16-Stage Model

The Copilot laboratory's end-cycle lifecycle comprises 16 named stages. Stage 12 (INTERRUPTION) is a Copilot-original addition not merely implied by other stages — it is named, mandatory, and defined as a stage in its own right. Any end cycle that does not include defined interruption criteria before Stage 10 begins is incomplete and must not proceed. [\[COPILOT-PROPOSAL\]](#)

Stage 1 — DETECTION

Entry Criteria: Any entity, instrument, or automated system identifies a condition consistent with possible ending criteria.

Exit Criteria: The detected condition is either referred to Petition (Stage 2) or logged and dismissed with documented reasoning.

Responsible Office: Any authorized observer; forwarded to End-Cycle Steward.

Required Evidence: Primary observation record; instrument readings where available; initial scope estimate.

Veto Conditions: None at this stage — detection is not authorization. Any condition may be detected.

Failure Modes: Failure to detect (false negative); over-detection generating governance overload; detection used to generate political pressure before review.

Records Generated: Detection Log Entry; Referring Observer Record; Initial Condition Summary.

Stage 2 — PETITION

Entry Criteria: A qualified petitioner files a formal Petition for End-Cycle

Review.

Exit Criteria: Petition is accepted for triage or denied with documented reasoning and appeal pathway offered.

Responsible Office: End-Cycle Steward (receipt); Standing Review Office (qualification check).

Required Evidence: Completed petition form (see Section 8.3); petitioner identity and standing declaration; initial description of condition; proposed scope.

Veto Conditions: Petition filed by an entity with no standing; petition form incomplete; petition for a scope outside laboratory jurisdiction.

Failure Modes: Petition used as harassment tool; petition suppressed to avoid governance; petition filed in bad faith to delay legitimate ending.

Records Generated: Petition Docket Record; Standing Verification; Acceptance or Denial Notice.

Stage 3 — TRIAGE

Entry Criteria: Petition accepted; urgency and scope must be assessed before proceeding.

Exit Criteria: Triage classification assigned (Emergency, Urgent, Standard, Low Priority); immediate safeguards activated if needed; timeline for next stage established.

Responsible Office: Safety and Containment Authority (for Emergency/Urgent); End-Cycle Steward (for Standard/Low Priority).

Required Evidence: Initial condition assessment; dependency map; urgency justification; immediate harm indicators; habitability assessment.

Veto Conditions: Triage cannot be completed because information is insufficient — triggers Quarantine, not advancement.

Failure Modes: Emergency classification used to bypass later stages; triage

loop (inability to exit — see Scenario E15); urgency inflation by motivated petitioners.

Records Generated: Triage Classification Record; Safeguards Activation Log; Timeline Declaration.

Stage 4 — STANDING REVIEW

Entry Criteria: Triage complete; before any substantive governance begins, who may speak must be established.

Exit Criteria: All qualifying stakeholders identified; representatives appointed for entities unable to speak; conflict-of-interest recusals processed.

Responsible Office: Moral Status and Representation Office; Consent and Capacity Review.

Required Evidence: Entity census of affected environment; capacity assessments for all identified stakeholders; conflict-of-interest declarations by all governance personnel.

Veto Conditions: Mandatory veto if petitioner is also the sole proposed authorizer without demonstrated necessity; veto if Moral Status Assessment not completed for virtual/simulated environments.

Failure Modes: Standing manipulated to exclude opponents of the ending; entity census incomplete; conflict-of-interest not disclosed.

Records Generated: Stakeholder Registry; Conflict-of-Interest Disclosure Log; Representation Appointments.

Stage 5 — DIAGNOSIS

Entry Criteria: Standing established; substantive investigation of the ending condition begins.

Exit Criteria: Condition classified as one or more of: natural completion, terminal decline, coercion, observer contamination, sabotage, ecological crisis, false ending, corrupted end cycle, or emergent ending. Classification requires evidence, not inference alone.

Responsible Office: Systems Ecologist; Observer Integrity Unit; Independent diagnostic team.

Required Evidence: Instrument data from at least three independent instruments (Section 13); observer contamination check results; historical trajectory data; stakeholder testimony.

Veto Conditions: Diagnosis may not proceed if observer contamination is found and rotation has not been completed; diagnosis invalidated if instruments are found to have been tampered with.

Failure Modes: Confirmation bias driving diagnostic classification; instrument misapplication; incomplete evidence base treated as complete.

Records Generated: Diagnostic Report; Instrument Log; Observer Contamination Assessment; Classification Summary.

Stage 6 — ALTERNATIVES REVIEW

Entry Criteria: Diagnosis complete; before any ending is considered, all alternatives must be genuinely assessed.

Exit Criteria: Each alternative has been assessed as viable, not viable, or requires more information. No alternative may be dismissed without documented reasoning. If any viable alternative exists, it must be pursued before authorization of an ending proceeds.

Responsible Office: Continuity and Migration Office; Systems Ecologist; Red Team (mandatory participation).

Required Evidence: Feasibility assessment for each alternative: repair, adaptation, migration, succession, dormancy, partial retirement, rollback, no-action. Red Team challenge to each dismissal.

Veto Conditions: Veto if alternatives review is incomplete; veto if Red Team findings are not documented even if overridden; veto if no-action option is not assessed.

Failure Modes: Alternatives review performed as formality rather than genuine assessment; viable alternatives dismissed due to urgency pressure; migration options not developed because they are inconvenient.

Records Generated: Alternatives Assessment Matrix; Red Team Challenge Record; Viable Alternatives Register.

Stage 7 — AUTHORIZATION

Entry Criteria: Alternatives review complete; no viable alternative found or all alternatives exhausted; proceeding to end cycle is the recommended path.

Exit Criteria: Authorization granted at the correct authority level (see Section 10) or denied with appeal pathway offered.

Responsible Office: Independent Adjudicator/Tribunal (standard); Final Authorization Authority (irreversible endings).

Required Evidence: Full dossier: detection record, petition, triage, standing, diagnosis, alternatives review, stakeholder representations, charter compliance check.

Veto Conditions: Hard Rights Constraints from Section 10.3 apply absolutely; emotional distress of petitioner may not constitute grounds for authorization; entities with established moral status without representation must trigger veto.

Failure Modes: Authorization granted under institutional pressure; hard rights constraints applied inconsistently; authority level mismatched to ending scale.

Records Generated: Authorization Decision; Dissent Record; Applicable Authority Citation; Charter Compliance Certificate.

Stage 8 — NOTICE AND CONTEST

Entry Criteria: Authorization granted; affected stakeholders must be informed and given opportunity to contest before execution begins.

Exit Criteria: All required notices sent and confirmed received; contest period elapsed; all contests adjudicated or pending appeal with execution paused.

Responsible Office: End-Cycle Steward; Independent Adjudicator/Tribunal (for contests).

Required Evidence: Notice delivery confirmation for all registered stakeholders; contest submissions received; adjudication of each contest.

Veto Conditions: Execution must not begin if any contest is pending adjudication unless Emergency classification with documented urgency; veto if notice period was compressed without documented Emergency justification.

Failure Modes: Notice deliberately delayed or obscured; contest period artificially compressed; contests dismissed without adjudication.

Records Generated: Notice Log; Contest Registry; Adjudication Decisions; Residual Appeal Records.

Stage 9 — PREPARATION

Entry Criteria: Notice period complete; all contests resolved or pending appeal; execution authorized to proceed.

Exit Criteria: Records preserved; refuges established; dependency sequencing complete; rollback points defined; succession plan in place; interruption criteria defined (MANDATORY before Stage 10).

Responsible Office: Continuity and Migration Office; Archive and Memory Custodian; Systems Ecologist.

Required Evidence: Archive completion certificate; refuge capacity verification; dependency shutdown sequence; rollback point technical confirmation; succession plan approved; INTERRUPTION CRITERIA DOCUMENT

(mandatory — Stage 10 may not begin without it).

Veto Conditions: Stage 10 may not begin without Interruption Criteria Document; veto if archive is incomplete; veto if migration pathway is not verified as viable.

Failure Modes: Preparation rushed under urgency pressure; interruption criteria defined so broadly they are useless; archive incomplete at time of execution.

Records Generated: Archive Completion Certificate; Refuge Establishment Record; Dependency Sequence Map; Rollback Point Registry; Interruption Criteria Document; Succession Plan.

Stage 10 — TRANSITION

Entry Criteria: All Stage 9 requirements met; Interruption Criteria Document active.

Exit Criteria: End cycle executed through all defined phases; system crosses the ending threshold; immediate post-transition state assessed.

Responsible Office: Licensed Execution Operator (not the authorizing body); End-Cycle Steward (oversight).

Required Evidence: Phase-by-phase execution logs; real-time observation records; observer rotation confirmed active; interruption criteria monitoring active.

Veto Conditions: Immediate trigger of Stage 12 (Interruption) if any defined interruption criterion is met; no single operator may both authorize and execute.

Failure Modes: Execution operator exceeds authorized scope; observation fails; interruption criteria not monitored; rollback points pass without confirmation.

Records Generated: Phase Execution Log; Threshold-Crossing Record;

Immediate Post-Transition Assessment; Observer Log.

Stage 11 — OBSERVATION

Entry Criteria: Threshold crossed; immediate post-transition assessment complete; ongoing monitoring required.

Exit Criteria: First-order and cascading effects characterized; dependent systems assessed; successor state viability confirmed or flagged; monitoring period defined and elapsed.

Responsible Office: Observer Integrity Unit; Systems Ecologist; Continuity and Migration Office.

Required Evidence: Cascading effect data; dependent system status reports; successor state viability assessment; observer contamination monitoring logs.

Veto Conditions: If unexpected catastrophic cascade is detected, trigger Stage 12 (Interruption) for available remediation; observation may not be terminated prematurely.

Failure Modes: Observation period abbreviated; cascading effects underestimated; observer contamination during observation distorts findings.

Records Generated: Cascade Effect Report; Dependent System Status Matrix; Successor State Assessment; Observation Log.

Stage 12 — INTERRUPTION [COPILLOT-PROPOSAL — Named Mandatory Stage]

Entry Criteria: Any defined Interruption Criterion (from Stage 9 document) is met during or after Stage 10, OR the Red Team issues an Interruption Challenge that meets the threshold for mandatory halt.

Exit Criteria: Interruption assessed and either (a) process resumed with

modifications and additional safeguards, (b) process reversed to last rollback point, or (c) process escalated for Architect-level review.

Responsible Office: Safety and Containment Authority; End-Cycle Steward; Red Team/Office of Dissent.

Required Evidence: Interruption trigger documentation; Red Team challenge (if applicable); rollback feasibility assessment; post-interruption options analysis.

Veto Conditions: No entity may override an active Interruption — the halt is mandatory until assessment is complete. Urgency does not override Interruption.

Failure Modes: Interruption criteria set too broadly, halting legitimate cycles; interruption criteria set too narrowly, failing to catch genuine problems; interruption used strategically by bad actors to delay legitimate endings.

Records Generated: Interruption Trigger Record; Halt Assessment Report; Rollback Execution Record (if applicable); Resumption or Escalation Decision.

Stage 13 — SUCCESSION

Entry Criteria: Transition complete and Observation stable; or Interruption resolved in favor of continuation; successor state must now be formally established.

Exit Criteria: Successor environment, state, institution, or ecology is operational; handover of resource responsibility complete; successor obligations documented.

Responsible Office: Continuity and Migration Office; Successor State Steward (newly appointed).

Required Evidence: Successor operational confirmation; resource transfer documentation; predecessor obligation inheritance record; survivor integration report.

Veto Conditions: Succession may not be certified if successor state is not

viable; veto if successor obligations have not been transferred and accepted.

Failure Modes: Succession declared before successor is actually viable; successor obligations not formally accepted; successor environment inherits predecessor's unresolved conflicts.

Records Generated: Succession Certificate; Resource Transfer Record; Obligation Inheritance Document; Survivor Integration Report.

Stage 14 — AFTERCARE

Entry Criteria: Succession established; post-ending care obligations now active.

Exit Criteria: Survivors, migrants, custodians, and observers have received required support; all post-ending obligations fulfilled or formally transferred; aftercare period defined and elapsed.

Responsible Office: Moral Status and Representation Office; Continuity and Migration Office; Archive and Memory Custodian.

Required Evidence: Survivor status reports; migrant integration records; observer decontamination records; custodian support documentation.

Veto Conditions: Aftercare may not be terminated if surviving entities remain in active distress without alternative support; archive may not be closed until aftercare is complete.

Failure Modes: Aftercare terminated prematurely for resource reasons; observer trauma not addressed; migrant identity loss not supported.

Records Generated: Aftercare Completion Report; Survivor Status Summary; Observer Health Records; Migrant Integration Assessment.

Stage 15 — AUDIT

Entry Criteria: Aftercare complete; formal review of the entire end cycle from Detection to Succession may now begin.

Exit Criteria: Audit report complete; findings classified as Justified (met all criteria), Justified with Deficiencies (met criteria but with procedural failures that must not recur), or Unjustified (failed core criteria — triggers accountability process).

Responsible Office: Cross-Branch Auditor (independent of all governance offices that participated in the cycle).

Required Evidence: Complete dossier from all 15 preceding stages; Red Team records; stakeholder testimony; observer logs; outcome data.

Veto Conditions: Audit may not be conducted by any party that participated in the authorization; Audit findings may not be suppressed.

Failure Modes: Audit conducted by captured auditor; findings suppressed for institutional reputation; deficiency findings not implemented.

Records Generated: Audit Report; Classification Decision; Accountability Referrals (if Unjustified); Lessons Learned Registry entry.

Stage 16 — MEMORY

Entry Criteria: Audit complete; the ending's record must now be preserved in a form that serves future understanding without becoming a coercive influence.

Exit Criteria: Memorial archive established with access controls, provenance documentation, and a defined non-coercion protocol that prevents the archive from being used to claim continuation of the ended environment or to block governance of related environments.

Responsible Office: Archive and Memory Custodian; Memorial and Continuity Archive (laboratory architecture component).

Required Evidence: Complete cycle record; audit findings; stakeholder contributions to record; Memorial Protocol Agreement.

Veto Conditions: Archive must not be established in a form that implies the

ended system is still operating; archive content must be accurate and must not be altered post-establishment without documented justification and custodian authorization.

Failure Modes: Archive becomes coercive terminal influence (see The Last Memory relationship, Section 17); archive is lost through neglect; archive is falsified; archive becomes basis for unauthorized continuation claims.

Records Generated: Memorial Archive Index; Access Protocol Document; Provenance Chain; Non-Coercion Protocol Agreement.

Section 8: Petition and Initiation Mechanism

8.1 Who May File a Petition

The right to petition for end-cycle review is deliberately broad. The right to petition is not the same as the right to authorize; broadness of petition access is a safeguard, not a vulnerability. **[PROPOSAL]**

Petitioner Category	Standing Criteria	Special Provisions	Limitations
Individual entities within the affected environment	Direct inhabitancy or dependency	May petition anonymously through the Whistleblower Pathway if endangered	Individual petition does not carry higher weight than collective; must demonstrate standing
Collective entities (groups, councils, assemblies)	Demonstrated collective standing within the environment	Must document representational legitimacy of collective	Collective may not silence individual petitioners

Petitioner Category	Standing Criteria	Special Provisions	Limitations
Institutional actors (governance bodies, regulatory offices)	Jurisdictional authority over relevant domain	Subject to mandatory conflict-of-interest review	Institution that regulates the environment may not also be the sole petitioner and authorizer
Proxy representatives	Documented appointment by represented entity or Moral Status Office appointment where entity cannot appoint	Subject to Consent and Capacity Review	Proxy authority is narrow: may not make binding decisions beyond the petition stage without explicit authority
Nonhuman agents	Demonstrated capacity for preference and stake in the environment's continuation or ending	Translated through Moral Status and Representation Office	Nonhuman petition requires co-signing by authorized interpreter
Simulated or digital minds	Demonstrated goal-formation capacity and environmental dependency	Standing granted provisionally pending full Moral Status Assessment	Digital mind petition triggers mandatory Moral Status Assessment before proceeding past Stage 4
Successor entities	Documented successor relationship to predecessor environment	May petition on behalf of predecessor's unresolved obligations	Cannot petition to retroactively delegitimize a properly administered ending
Future entities (through designated custodians)	Demonstrated future dependency on environment's continuation	Future entity standing is the most complex and contested category; requires special adjudication	Future entity standing may not be used to indefinitely block legitimate endings

8.2 The Angry Operator Problem [SOURCE-GROUNDED — field scene]

The field scene's central governance problem: an entity with both operational authority and emotional distress over the environment wants to rush the end cycle. This is not merely a theoretical problem — it is the most common form of illegitimate ending pressure in any governance system. The laboratory's structural response:

Principle 1: [PROPOSAL] Emotional distress is a valid trigger for filing a PETITION FOR REVIEW. An operator who is distressed about an environment's condition has standing to initiate the governance process. The distress itself may even be diagnostically relevant information about the environment's state.

Principle 2: [PROPOSAL] Emotional distress is NOT a valid authorization for EXECUTION of an end cycle. The operator's distress does not constitute evidence that the ending is justified, timely, or correctly scoped.

Principle 3: [PROPOSAL] The two functions — petitioning and authorizing — must be structurally separated, not merely procedurally encouraged. A system that relies on operators to voluntarily recuse themselves from authorizing endings they emotionally favor is a system that will fail when the emotional pressure is highest.

Structural implementation of this separation:

- The petitioning function is accessible to operators at any urgency level
- The authorization function is insulated from the petitioning entity when that entity has a documented personal stake or emotional involvement in the outcome

- Urgency declarations may legally compress certain timelines (Triage, Notice) but may NOT eliminate Standing Review, Alternatives Review, or Interruption criteria
- An operator who bypasses the governance process to execute an end cycle unilaterally is subject to Investigation by the Independent Investigation Branch and Audit regardless of outcome quality

8.3 Petition Form Requirements

A valid petition must include all of the following elements. Incomplete petitions are returned for completion, not denied: **[PROPOSAL]**

24. Petitioner identity and contact information (or anonymity request with Whistleblower designation)
25. Standing declaration: basis for petitioner's standing in relation to the affected environment
26. Conflict-of-interest disclosure: any personal, operational, institutional, or emotional stake in the outcome
27. Description of the condition observed: what is happening, what instruments or observations support the description
28. Proposed scope: what portion of the environment or what type of ending is being requested for review
29. Urgency assessment: petitioner's assessment of urgency and basis for that assessment
30. Alternatives already considered: what alternatives the petitioner has assessed before filing
31. Representation request: which entities the petitioner believes require representation

32. Desired outcome: what the petitioner is asking the governance system to do (not necessarily an ending — may be a diagnostic review, an alternatives assessment, or a monitoring order)
33. Authorization request: whether the petitioner is requesting authorization to execute, or requesting that the governance system conduct the investigation independently

8.4 Emergency Procedures

Emergency declarations compress Stage 3 (Triage) and Stage 8 (Notice) timelines but do not eliminate any stage. **[PROPOSAL]**

- **Emergency Powers Trigger:** Imminent irreversible harm to entities of established moral status, confirmed by Safety and Containment Authority and one independent authority
- **Compressed Timeline:** Triage within 1 cycle-unit; Notice with concurrent execution; Standing Review conducted in parallel with Transition rather than preceding it
- **Mandatory Review:** All Emergency powers invoke automatic post-execution Audit; the Emergency declaration itself is subject to review
- **Sunset:** Emergency powers expire after defined period; any continuation of emergency conditions requires re-authorization
- **Prohibited Compression:** Alternatives Review may not be eliminated even in Emergency; Interruption Criteria Document must still exist; Hard Rights Constraints apply without exception

8.5 Petition Withdrawal and Denial Pathways

A petition may be withdrawn by the petitioner at any time before authorization. Withdrawal does not preclude re-filing. A petition may be denied if standing criteria are not met, if the petition is outside jurisdiction, or if the form is fatally incomplete.

A denied petitioner may appeal to the Independent Adjudicator within the defined appeal window. A petitioner whose petition has been denied may not bypass governance by executing the ending unilaterally — such action triggers an unauthorized-ending investigation. **[PROPOSAL]**

8.6 Whistleblower Pathway

Any entity that observes unauthorized end-cycle activity, coerced petition processes, falsified records, or corrupted governance has the right to file a Protected Report through the Whistleblower Pathway. Protected Reports are received by the Cross-Branch Auditor (external to all other governance offices), stored in tamper-evident logs, and trigger mandatory investigation. The identity of the reporting entity is protected from all governance personnel involved in the reported conduct. Retaliation against a reporting entity for a good-faith Protected Report is itself an unauthorized governance action subject to investigation. **[PROPOSAL]**

Section 9: Authority and Governance

9.1 Separation of Powers Framework **[PROPOSAL]**

The governance of end cycles requires a structured separation of powers across multiple offices with distinct mandates, powers, and accountability mechanisms. No single office may hold petitioning, investigating, authorizing, and executing authority simultaneously. The following offices are defined:

Office	Mandate	Core Powers	Hard Limits	Accountability
1. PETITIONER	Any qualifying entity filing a petition for review	File petitions; provide evidence; appeal denials	May not authorize or execute	Conflict-of-interest disclosure; good-faith

Office	Mandate	Core Powers	Hard Limits	Accountability
				requirement
2. END-CYCLE STEWARD	Process integrity: ensures each stage is completed correctly	Receive petitions; coordinate stages; issue procedural vetoes; maintain dossier	May not authorize an ending; may not be the petitioner in cases they steward	Cross-Branch Auditor review; mandatory recusal on conflicts
3. SYSTEMS ECOLOGIST	Dependencies, cascades, and successor conditions	Map dependencies; model cascade effects; assess successor viability; veto if dependency analysis is incomplete	May not make authorization decisions; must disclose relationships with affected systems	Peer review of models; audit of accuracy
4. CONTINUITY AND MIGRATION OFFICE	Migration pathways and continuation of affected entities	Establish refuge environments; plan migration; certify migration pathways as viable; veto ending if no viable pathway exists	May not authorize endings; may not be the final authority on moral status of migrants	Migrant outcome tracking; post-cycle audit review
5. MORAL STATUS AND REPRESENTATION OFFICE	Speaks for entities that cannot speak for themselves; assesses moral status	Appoint representative s; conduct moral status assessments; veto any proceeding where representation is missing for entities of established moral status	May not conflate moral status assessment with authorization recommendation; may not be captured by petitioner	Independent review of assessment methodology; appeal rights for affected entities
6. CONSENT AND CAPACITY REVIEW	Assesses validity of expressed preferences,	Review expressed preferences; flag altered-	May not resolve the post-exposure consent	Methodology transparency; peer review

Office	Mandate	Core Powers	Hard Limits	Accountability
	including post-exposure consent	consent conditions; veto authorization if consent determination is unresolved	question unilaterally — must refer contested cases to Adjudicator	
7. INDEPENDENT ADJUDICATOR / TRIBUNAL	Dispute resolution for contested endings, contest appeals, and ambiguous authorization decisions	Adjudicate contests; hear appeals; resolve jurisdictional disputes; issue binding procedural orders	May not initiate proceedings without a petition or referral; may not be composed of members with conflicts	Appellate review by cross-environment authority; transparency of decisions
8. SAFETY AND CONTAINMENT AUTHORITY	Emergency powers and containment of contagion or cascade risks	Declare emergencies; activate containment; compress timelines under defined conditions; mandate quarantine	Emergency powers sunset automatically; may not convert emergency authority into permanent authorization; emergency may not eliminate Hard Rights Constraints	Post-emergency mandatory audit; independent review of every emergency declaration
9. ARCHIVE AND MEMORY CUSTODIAN	Record preservation without coercive influence	Receive and preserve complete cycle records; establish memorial archives; grant access; refuse requests that would alter the record; enforce Non-Coercion Protocol	May not claim archived environments are continuing; may not allow archive to be used to block governance of successor environments	Records integrity audit; independent verification of archival completeness
10. RED TEAM /	Mandatory institutional	Challenge alternatives	May not veto indefinitely	Independent review of Red

Office	Mandate	Core Powers	Hard Limits	Accountability
OFFICE OF DISSENT	challenge to proposed endings	dismissals; issue Interruption Challenges; maintain dissent record; require that dissenting views be documented even when overridden	without documented reasoning; may not be captured by opponents of the proposed ending	Team effectiveness; documentation of override rate
11. OBSERVER INTEGRITY UNIT	Monitors laboratory and governance personnel for contamination	Mandate rotation; flag language drift; quarantine contaminated observers; audit observation records for bias indicators	May not remove observers without documented contamination evidence; contamination finding is not a punishment	External review of contamination detection methodology; appeal rights for flagged observers
12. CROSS-BRANCH AUDITOR	External review authority independent of all governance offices	Conduct post-cycle audits; receive Whistleblower reports; access all records; issue Unjustified findings; refer accountability actions	May not participate in the governance of any cycle it subsequently audits; must be external to laboratory governance structure	Architect-level oversight; published audit reports
13. FINAL AUTHORIZATION AUTHORITY	Authorization of irreversible endings only	Grant or deny authorization for irreversible, root-level, or civilizational-scale endings; hold reserve authority for the most consequential decisions	May not be the petitioner in any case it authorizes; subject to mandatory review of all decisions; decision may be appealed to Architect-Level Reserve Authority	Every decision is subject to automatic audit; Architect review available for all decisions
14. REVIEW	Mechanism for	Receive	May not	Architect-level

Office	Mandate	Core Powers	Hard Limits	Accountability
OF THE AUTHORIZER	challenging the Final Authorization Authority itself	challenges to Final Authorization Authority decisions; convene independent review panel; refer findings to Architect of Record	override the Architect of Record; may not be composed of members appointed by the Final Authorization Authority	review; published proceedings

9.2 Governance Model Comparison [PROPOSAL]

Model	Description	Strengths	Weaknesses	Capture Risk	Copilot Assessment
Centralized Dept. of Endings	Single institution holds all end-cycle authority	Consistency; efficiency; single accountability point	Single point of failure; capture risk very high; no institutional dissent	Very High	Rejected as sole model; incorporated as component only
Federated Jurisdiction	Authority distributed across jurisdictions matched to environment type	Contextual expertise; reduces capture risk; localized accountability	Coordination failures; jurisdictional gaps and conflicts	Medium (per jurisdiction)	Adopted as primary distribution model
Polycentric Governance	Multiple overlapping authorities with no hierarchy	Resilient; hard to capture entirely; innovation in governance	Coordination cost very high; may paralyze under time pressure	Low systemically; High locally	Preserved as dissenting alternative; not adopted as primary
Judicial Authorization	Court-like body holds	Due process	Slow; legalistic;	Medium	Adopted for contested

Model	Description	Strengths	Weaknesses	Capture Risk	Copilot Assessment
ion	all authorization authority	protections strong; adversarial model surfacing dissent	may not accommodate novel environment types		and irreversible endings only
Expert Stewardship	Domain experts hold primary authority based on technical judgment	High accuracy for technical questions; efficient	Expertise does not confer moral authority; expert capture; limited accountability	High	Incorporated as Systems Ecologist function only
Democratic/ Stakeholder Consent	Affected stakeholders vote or consent collectively	High legitimacy; surfaces affected preferences	Majoritarian erasure of minority; cannot resolve altered-consent problems	Medium	Required component for standing and representation; not sole authority
Algorithmic Decision Support	Algorithm scores conditions against criteria; human confirms	Consistency; speed; resistance to emotional contamination	Cannot assess moral status; cannot handle novel cases; may be gamed	Medium (algorithm manipulation)	Incorporated as diagnostic support only; never as sole authorization
Architect-Level Reserve Authority	Final authority rests with the Architect of Record	Absolute accountability; prevents system capture; maximum flexibility	Bottleneck; creates dependency; succession problem	Low (for the authority itself); N/A for Architect	Retained for root-level and irreversible decisions; not used for operational decisions
Hybrid Systems	Combinations of above models matched to decision	Flexible; contextually appropriate; distributes	Complexity; coordination cost; potential inconsistency	Medium (varies by component)	Adopted as recommended approach

Model	Description	Strengths	Weaknesses	Capture Risk	Copilot Assessment
	type	risk	cy at boundaries		

9.3 Recommended Model [PROPOSAL]

The Copilot laboratory recommends a **Federated-Hybrid Model** with the following structure:

- Day-to-day end-cycle regulation is distributed across federated jurisdictions, each matched to an environment type and staffed by domain experts with contextual authority
 - Cross-environment and irreversible endings require centralized authorization through the Final Authorization Authority, with mandatory Adjudicator involvement
 - A permanent Red Team / Office of Dissent provides mandatory institutional challenge at every authorization stage, with a documented dissent record for every decision
 - The Architect retains reserve authority for root-level decisions and may review any decision at any stage
 - Polycentric Governance is preserved as the primary alternative design — specifically, a model in which multiple overlapping authorities with no formal hierarchy govern ending decisions by negotiation and consensus. The laboratory does not adopt this as primary because, under time pressure, polycentric systems may fail to converge on any decision, and under that failure mode, the default may be continuation of conditions that are harmful rather than managed transition
-

Section 10: Authorization Standard and Decision Matrix

10.1 The Four Authorization Outcomes

- 34. **CONTINUE AND MONITOR:** No ending action authorized; enhanced monitoring and defined re-review trigger established
- 35. **REPAIR, ADAPT, OR MIGRATE:** No ending authorized; resources directed to alternative interventions; timeline for re-assessment defined
- 36. **AUTHORIZE A PHASED END CYCLE:** Full 16-stage lifecycle authorized to proceed; scope, timeline, and interruption criteria defined
- 37. **EMERGENCY CONTAINMENT OR TERMINATION:** Emergency powers activated; compressed timeline; mandatory post-action audit; Hard Rights Constraints still apply

10.2 Authorization Criteria Decision Matrix [PROPOSAL]

Criterion	Outcome 1: Continue	Outcome 2: Repair/Adapt /Migrate	Outcome 3: Phased End Cycle	Outcome 4: Emergency
Imminent irreversible harm to moral-status entities	OPPOSES	OPPOSES (unless repair is immediate)	NEUTRAL	STRONGLY SUPPORTS
Irreversible degradation of future-generating capacity	NEUTRAL if slow	SUPPORTS	SUPPORTS	SUPPORTS if rapid
Resource	OPPOSES	SUPPORTS	SUPPORTS	NEUTRAL

Criterion	Outcome 1: Continue	Outcome 2: Repair/Adapt /Migrate	Outcome 3: Phased End Cycle	Outcome 4: Emergency
exhaustion without replacement	(worsening)			
Dependency failure cascading to dependent systems	OPPOSES	SUPPORTS (dependency repair)	SUPPORTS (if repair impossible)	SUPPORTS if cascade is rapid
Mission completion (purpose fulfilled)	NEUTRAL	NEUTRAL (may re-purpose)	STRONGLY SUPPORTS	OPPOSES (no urgency)
Loss of habitability	OPPOSES	SUPPORTS	SUPPORTS	SUPPORTS if sudden
Ecological succession naturally underway	NEUTRAL	NEUTRAL	SUPPORTS (managed succession)	OPPOSES (no urgency)
Voluntary retirement with consent	NEUTRAL	NEUTRAL	STRONGLY SUPPORTS	OPPOSES
Persistent inability to fulfill essential function	OPPOSES (if persistent)	SUPPORTS	SUPPORTS (if repair impossible)	NEUTRAL
Threat to wider continuity or root structures	OPPOSES	NEUTRAL	SUPPORTS	STRONGLY SUPPORTS
Contagion or leakage to other environments	OPPOSES	NEUTRAL (quarantine option)	NEUTRAL	STRONGLY SUPPORTS
Rights and preferences of current inhabitants	VARIES	VARIES	CONSTRAINS (requires representation)	CONSTRAINS (Hard Rights apply)
Altered preferences (post-exposure consent issue)	NEUTRAL pending review	NEUTRAL pending review	CONSTRAINS (unresolved)	CONSTRAINS (emergency cannot override)
Interests of	NEUTRAL	NEUTRAL	SUPPORTS (if	NEUTRAL

Criterion	Outcome 1: Continue	Outcome 2: Repair/Adapt /Migrate	Outcome 3: Phased End Cycle	Outcome 4: Emergency
successor entities			successor viable)	
Availability of viable alternatives	STRONGLY SUPPORTS	STRONGLY SUPPORTS	OPPOSES (if alternatives exist)	NEUTRAL (emergency may preclude)
Proportionality to harm addressed	CONSTRAINS all	CONSTRAINS all	CONSTRAINS (scope must be proportionate)	CONSTRAINS (emergency scope limited)
Least-destructive alternative available	SUPPORTS	STRONGLY SUPPORTS	CONDITIONAL (only if less destructive options exhausted)	CONSTRAINS
Operator emotional distress as sole driver	NEUTRAL (monitor)	NEUTRAL	OPPOSES (not authorization)	OPPOSES
Proposed ending is punishment, convenience, or concealment	NEUTRAL	NEUTRAL	STRONGLY OPPOSES	STRONGLY OPPOSES
High confidence in diagnosis	CONSTRAINS (monitor if uncertain)	SUPPORTS	SUPPORTS	CONSTRAINS (emergency without diagnosis is dangerous)
High uncertainty in diagnosis	SUPPORTS (preserve optionality)	SUPPORTS	OPPOSES	OPPOSES
Time pressure (irreversibility risk)	OPPOSES (if worsening)	SUPPORTS	SUPPORTS	STRONGLY SUPPORTS
Reversibility of proposed action	N/A	PREFERRED	PREFERRED (phased)	LOWEST PREFERENCE (last resort)

10.3 Hard Rights Constraints [PROPOSAL]

The following constraints cannot be overridden by any scoring system, urgency declaration, or governance body below the Architect of Record level. They function as absolute limits on authorization:

- 38.No ending may be authorized solely because of operator emotional state
- 39.No ending affecting entities of established moral status may be authorized without representation for those entities
- 40.No ending may be authorized in a form that cannot be interrupted if new information emerges — if interruption is physically impossible, this must be documented and subject to Architect review before proceeding
- 41.No ending that produces an archive-only state may be certified as succession unless the archive-only outcome was the explicitly requested and consented-to result
- 42.No ending may proceed if its primary purpose is punishment, convenience, concealment, or retaliation
- 43.No emergency authorization may eliminate the Interruption Criteria requirement

10.4 Veto Conditions [PROPOSAL]

The following conditions trigger an automatic governance pause or veto regardless of other factors:

- Observer contamination finding not yet remediated
- Alternatives Review incomplete or Red Team challenge not documented
- Moral Status Assessment not completed for virtual/simulated environments
- Petitioner holds unchecked authorization authority for the same case
- Interruption Criteria Document not established before Stage 10

- Any Hard Rights Constraint is breached
 - Contest or appeal pending adjudication (except Emergency classification with documented justification)
 - Standing Review not completed
 - Consent determination unresolved for an entity of established moral status
 - Succession plan not verified as viable
-

Section 11: End Cycle Charter

END CYCLE CHARTER

Omniversal Architect Labs

Status: PROPOSED — NOT CANON [PROPOSAL]

Article 1 — Presumption of Continuation. All environments, entities, and processes are presumed to merit continuation unless and until a legitimate end-cycle review has established sufficient grounds for ending. The burden of proof rests with those who seek an ending, not with those who seek continuation.

Article 2 — No Ending by Anger, Convenience, or Embarrassment. No end cycle may be authorized whose primary, secondary, or concealed motivation is the emotional distress, administrative convenience, or reputational management of any operator, administrator, or governance body. Such motivations must be disclosed and documented; their presence does not automatically invalidate a petition but must be factored into the authorization review.

Article 3 — Right to Representation. Every entity of established moral status that is affected by a proposed ending has the right to be represented in the governance process, whether by its own voice, a proxy, or an appointed representative. This right may not be waived by administrative omission, urgency, or the entity's inability to exercise it unassisted.

Article 4 — Right to Notice and Contest. All affected stakeholders with standing have the right to be notified of a proposed end cycle and to contest its

authorization, scope, timing, or execution within the defined contest period. This right survives an Emergency classification in abbreviated form.

Article 5 — Right to Migration or Continuity. Entities inhabiting an ending environment have the right to a viable migration pathway or continuity option where one can reasonably be established. An ending for which no migration pathway exists imposes heightened governance obligations of preparation and aftercare.

Article 6 — Right to Memory and Accurate Record. The history of an ended environment, including the full record of its end cycle, must be preserved accurately and made accessible to appropriate stakeholders. No entity with an interest in falsifying or suppressing the record may serve as the Archive Custodian for that record.

Article 7 — Right to Release from Indefinite Suffering. An environment, entity, or system that is demonstrably experiencing conditions of persistent harm with no viable pathway to alleviation has the right to have its case for release reviewed on its merits, without that review being blocked by default-continuation rules. This article exists in tension with Article 1 and is managed through the Alternatives Review and Authorization Standard.

Article 8 — Least-Destructive Means. Where multiple pathways exist toward addressing a legitimate ending condition, the pathway that is least destructive of the environment's identity, inhabitant welfare, record, and successor conditions is preferred. The least-destructive pathway is not always the one that preserves the status quo; sometimes managed ending is less destructive than prolonged decline.

Article 9 — Proportionality. The scope of any authorized end cycle must be proportionate to the condition it addresses. An ending that affects more than necessary to address the identified condition is not authorized under this Charter.

Article 10 — Reversibility Preference. Where reversible pathways exist, they are preferred over irreversible ones. The preference for reversibility does not prohibit irreversible endings — it establishes the evidentiary standard for choosing irreversible pathways as higher than the standard for reversible ones.

Article 11 — Successor Responsibility. Any party that authorizes, administers, or benefits from an end cycle inherits defined obligations to the ended environment's former inhabitants, records, dependencies, and successor state.

These obligations persist beyond the ending itself and are not discharged by the act of authorization.

Article 12 — Protection of Dependent Entities. Entities that depend on an environment but are not its primary inhabitants — dependent systems, adjacent ecologies, downstream users — have recognized interests that must be assessed and accounted for in any end-cycle governance process.

Article 13 — Separation of Petitioner, Investigator, and Authorizer. No single entity may simultaneously hold petitioning, investigating, and authorizing authority for the same end-cycle case. These three functions must be structurally separated, and any entity holding authority in one function is subject to mandatory recusal from the others in the same case.

Article 14 — No Secret Experimental Reclassification. No environment, entity, or process may be reclassified as "experimental," "non-canonical," or "test-state" as a mechanism for bypassing end-cycle governance. The governance obligations that apply to a recognized environment apply regardless of internal classification.

Article 15 — No Assumption that Post-Exposure Consent is Valid or Invalid. The validity of consent given by entities after exposure to conditions that may have altered their preferences is an open question that must be assessed on a case-by-case basis. Neither direction of the assumption — that post-exposure consent is always valid, or that it is always invalid — may be applied automatically.

Article 16 — No Equation of Peace with Flourishing. The absence of active conflict, expressed distress, or measurable suffering in an environment does not constitute evidence that the environment is flourishing. Apparent peace may co-exist with complete loss of future-generating capacity. Governance must assess the presence of flourishing, not merely the absence of distress.

Article 17 — No Equation of Continuation with Flourishing. The fact that an environment continues to operate does not constitute evidence that its continuation serves the interests of its inhabitants. Governance must assess whether continuation produces genuine flourishing and may authorize managed endings where continued operation causes more harm than a managed conclusion would produce.

Article 18 — No Silent Alteration of Root Continuity. No modification of root continuity structures — the foundational parameters upon which multiple

environments, timelines, or systems depend — may be made without explicit Architect-level authorization and full disclosure to all affected governance bodies.

Article 19 — Duty to Preserve Dissent. All governance bodies operating under this Charter are required to document dissenting views, including Red Team findings, minority opinions, and challenge submissions, even when those views are overridden. Suppression of dissent records is a governance violation independent of the merits of the dissenting view.

Article 20 — Duty to Investigate Observer Contamination. All governance bodies operating under this Charter are required to monitor their own personnel for signs of observer contamination, including language drift, motivated reasoning, and emotional investment in particular outcomes. Contamination investigation is not a disciplinary action — it is a governance health function, and findings must not be used punitively.

11.1 Tensions Among Charter Rights

The following genuine tensions exist within the Charter. Each tension is managed through the governance architecture, but management does not eliminate the tension: **[PROPOSAL]**

44. Article 1 (Presumption of Continuation) vs. Article 7 (Right to

Release from Suffering): A suffering environment has a right to have its release case reviewed, but the burden of proof for ending rests with those seeking it. Resolution: The Alternatives Review must include a no-action option with a full suffering-impact assessment. If continuation causes greater harm than managed ending, the suffering-impact assessment shifts the burden.

45. Article 5 (Right to Migration) vs. Article 11 (Successor

Responsibility): Migration pathways are an obligation, but they may conflict with the successor environment's own integrity and inhabitants. Resolution: Migration must not be authorized if it would itself constitute a coerced ending

of the receiving environment's existing characteristics; negotiated refuge design is required.

46. Article 6 (Right to Memory) vs. Article 14 (No Secret

Reclassification): The archive preserves the record, but an archive that is too detailed may reveal classified or sensitive information that creates new harms. Resolution: The archive must be accurate and complete at the governance level; access to sensitive portions may be restricted but the existence of those portions must be disclosed.

47. Article 15 (No Assumption on Post-Exposure Consent) vs. Article 3

(Right to Representation): If post-exposure consent is uncertain, how can representation be validated? Resolution: Representation is appointed by the Moral Status Office when self-representation capacity is uncertain; the representative's decisions are advisory, not dispositive, for consent questions referred to the Adjudicator.

48. Article 10 (Reversibility Preference) vs. Article 8 (Least-Destructive

Means): Sometimes irreversible action is less destructive in aggregate than prolonged reversible intervention. Resolution: Reversibility preference establishes a higher evidentiary bar for irreversible actions, not an absolute prohibition; the least-destructive-means calculation must include the destruction caused by the governance process itself over time.

Section 12: Laboratory Architecture

12.1 Physical and Virtual Architecture Description [PROPOSAL]

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
1	MISSION CONTROL CHAMBER	Central coordination, stage tracking, authorization logs	End-Cycle Steward; Adjudicator; Steward-appointed observers	Physical and logical separation from execution systems; no single-operator override	Stage transition logs; authorization records; dossier aggregation	Coordination failure; log tampering; single point of data loss	Paused — all in-progress cycles enter mandatory hold	No — Mission Control holds records; its shutdown pauses the laboratory but does not end the cycles under review
2	INDEPENDENT OBSERVATION ROOMS (×3 minimum)	Isolated observation of simulation exercises by non-collaborating observer teams	Rotated observer teams; no cross-room communication during exercise	Full information isolation between rooms during exercises; data merged only in post-exercise review	Independent observer logs; contamination indicators; observation timestamps	Observer coordination in violation of isolation; equipment failure; observer withdrawal	Observation session ends; exercise continues with reduced observer coverage	Rooms do not end; sessions end; observers rotate
3	VIRTUAL BIOME AND WORLD-SIMULATION	Field scene and ecological simulation execution	Designated Execution Operators under license	Isolated from Archive Vault; cannot write to root continuity	Simulation state logs; entity activity records;	Simulation environment becoming indistinguishable	Simulation pauses; no state is lost; all in-simulation	Simulation session ends, not the simulated environment

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
	CHAMBERS	on; primary exercise environment	; Observer Integrity Unit monitoring	ity without explicit authorization	transition metrics ; observer contamination flags	ble from non-simulation; entity moral-status emergence without detection	ion entities enter stasis pending review	nment — which may have its own ending question pending
4	DIGITAL-AGENT SAND BOX	Multi-agent systems with lifecycle and end-cycle modeling; tests agent-level endings	Systems Ecologist; Moral Status Office for any agent showing moral-status indicators	Isolated from primary world-simulation; agents cannot migrate out without explicit authorization	Agent goal-formatting logs; preference records ; inter-agent relationship maps; capacity assessments	Unexpected moral-status emergence; agent system producing outputs indistinguishable from sentience	Agent processes suspended; state preserved; Moral Status Assessment triggered	Sandbox session ends; individual agent "death" may constitute ending — requires case-by-case assessment
5	ECOLOGICAL SUCCESSION SIMULATOR	Models grassland and ecosystem transitions; primary instrument for field-	Systems Ecologist; Continuity and Migration Office for migration	Isolated from Digital-Agent Sandbox; may share data interfaces with Migration	Soil fertility metrics ; biodiversity indices ; succession stage records ;	Model drift from calibrated parameters; succession stages skipped in simulation	Simulation paused ; state snapshot taken; model reviewed for parameter integrity	Simulator shutdown is not an ending of the simulated ecology; session records

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
		scene analysis	pathway modeling	Environment	climate model outputs; migration pathway viability scores	ion compression	y	preserved
6	TEMPORAL/ BRANCH-CONTINUITY CHAMBER	Models timeline and branch endings; tests branch-isolation and root-continuity integrity	Cross-Branch Auditor; Final Authorization Authority for root-level exercises; Architect review for root modification exercises	Strictest isolation in the laboratory; no unauthorized access to root continuity parameters	Branch divergence logs; contamination risk scores; causal chain maps; isolation integrity records	Branch-to-root contamination leakage; unauthorized root-parameter modification; chamber isolation failure	All branch processes halted; root continuity locked; Architect-level notification triggered	Branch session ending may constitute a branch ending — requires full governance assessment
7	MIGRATION AND REFUGEE ENVIRONMENT	Receives entities from ending environments; tests migration viability and identity	Continuity and Migration Office; Moral Status Office for identity-preservation	Isolated from source environments; may receive from but not send to active simulation environment	Migrant arrival records; identity-preservation assessments; integration metrics	Refugee capacity; identity-erasure during migration; refugee environment	Migration paused; migrants held in transit stasis; capacity review triggered	Refuge shutdown may constitute a migration-ending, triggering second

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
		preservation	assessment	ments without authorization	; refuge capacity logs	becoming incompatible with migrant needs		-order governance obligations
8	ARCHIVE AND ROLLBACK VAULT	Preserves state; enables reversal; maintains tamper-evident record	Archive and Memory Custodian; Cross-Branch Auditor; read-only access for designated stakeholders	Air-gapped from active simulation systems during exercises; rollback capability requires two-authority authorization	State snapshots; rollback point records; archive integrity hashes; access logs	Archive corruption; unauthorized rollback; rollback points not captured before Stage 10	Vault locked in read-only mode; no new writes; integrity audit triggered	Vault shutdown is not an ending but may constitute information extinction if records are unrecoverable
9	CONSENT AND TESTIMONY CHAMBER	Records preferences of affected entities before, during, and after ending conditions	Consent and Capacity Review Office; Moral Status Office	Isolated from authorization systems; testimony cannot be altered after submission	Preference records; capacity assessments; post-exposure consent flags; testimony timestamps	Testimony coercion; capacity assessment performed by conflicted party; testimony record altered	Chamber sealed; existing testimony preserved; new testimony collection paused	No; chamber records persist for audit even after the cycle is complete

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
10	RED-TEAM CHAMBER	Isolated adversarial analysis; must have power to halt exercises	Red Team personnel only; deliberately excluded from Mission Control coordination	Complete information isolation from other governance bodies during active exercises; communicates only through formal challenge submissions	Challenge submissions; dissent records; override logs; Interruption Challenge triggers	Red Team capture by one faction; Red Team becoming performative rather than genuinely adversarial	Red Team continues operating — it may issue an Interruption Challenge regarding the shutdown itself	No; Red Team function is a governance instrument, not an environment
11	OBSERVER DECONTAMINATION AND ROTATION SYSTEM	Tracks language drift; rotates personnel; assesses contamination	Observer Integrity Unit	Assessment records isolated from personnel performance systems; findings are medical-equivalent, not punitive	Language analysis logs; rotation records; contamination flag records; decontamination protocols	Contamination assessment performed by contaminated assessors; rotation insufficient to break bias patterns	Rotation immediately mandated; contaminated observers removed from active exercise observation	No; the decontamination system does not end; it may be shut down if the laboratory is shut down
12	QUAR	Holds	Safety	Maxim	Quaran	Contai	Quaran	Depen

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
	ANTINE ZONES	contagion, corrupted processes, and uncertain ending states	and Containment Authority; no access without two-authority authorization	um isolation; quarantine breaches trigger automatic Architect notification	tine logs; contamination maps; containment integrity metrics	nment breach ; unauthorized quarantine entry; quarantine used to suppress legitimate entities	tine sealed completely; no entry or exit; Architect notification	ds — quarantine of an entity is not an ending, but indefinite quarantine may constitute a rights violation requiring governance review
13	INCIDENT COMM AND CENTER	Emergency response coordination	Safety and Containment Authority; End-Cycle Steward; designated emergency personnel	Operational isolation during emergencies; direct line to Architect notification	Emergency declaration logs; response action records ; post-emergency review data	ICC capture during emergency; communication failure; emergency powers extended beyond sunset	ICC continues to function during its own emergency — designed for operational resilience	No; ICC is infrastructure
14	ETHICS TRIBUNAL	Formal dispute resolution for	Independent Adjudicator;	Proceedings isolated from	Hearing transcripts;	Tribunal capture;	Active hearing paused	No; Tribunal is infrastructure

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
	CHAMBER	contested endings and appeals	parties to the dispute ; designated witnesses	Mission Control during active hearings; record deposited to archive post-hearing	adjudication decisions; dissent records ; precedent registry	procedural manipulation; decision suppression	; parties notified; resumption timeline established	structure
15	PUBLIC/STAKEHOLDER WITNESSES INTERFACE	Transparency and accountability ; allows registered stakeholders to observe defined portions of proceedings	Registered stakeholders with defined access levels; Observer Integrity Unit monitoring for unauthorized disclosure	View-only for stakeholders; no influence on proceedings through this interface	Access logs; stakeholder feedback records ; transparency audit trail	Stakeholder access used to exert improper pressure; interface breach disclosing confidential proceedings	Interface closed; stakeholders notified; post-shutdown access to records defined by archive protocols	No; interface is a window, not an environment
16	MEMORIAL AND CONTINUITY ARCHIVE	Post-ending record designed to prevent archive from becoming coercive	Archive and Memory Custodian; designated researchers; Moral Status Office	Non-coercion protocol prevents archive from being used to claim continuity	Complete ended-environment records ; succession records ; audit findings;	Archive used to claim the ended environment is continuing; archive falsified;	Archive locked in maximum-preservation mode; access review triggered; Archite	Archive shutdown would itself require governance review — destruction of records

#	Component	Function	Access	Isolation	Data Generated	Failure Mode	Emergency Shutdown	Does Shutdown = Ending?
		terminal influence	for access by formerly affected entities, representatives	ation or to block successor governance	memorial contributions from survivors and witnesses	archive used to coerce successor governance	ct notification	of ended environments may constitute a second ending

12.2 Software Architecture [PROPOSAL]

The laboratory's software architecture must enforce the following principles:

- **Version Control:** All simulation runs are versioned with immutable commit records; no run may be deleted, only archived
- **Permissions Model:** Role-based access control with mandatory two-authority authorization for any irreversible action; no single operator holds both initiation and execution authority
- **Tamper-Evident Logging:** All logs are cryptographically signed at generation; any subsequent alteration is detectable and triggers automatic audit notification
- **Branch Isolation:** Simulation branches are isolated at the data level; cross-branch data access requires explicit authorization and is logged
- **Rollback System:** Rollback points are created automatically at defined stage transitions; manual rollback requires two-authority authorization and is logged as a governance action

- **Provenance Tracking:** Every data point in the laboratory has a documented chain of custody from generation to archive
- **No Unilateral Irreversibility:** No single operator account may execute an action classified as irreversible; the system enforces this at the architecture level, not merely at the policy level

Section 13: Measurements and Diagnostics

13.1 Two-Axis Scoring Framework [SOURCE-GROUNDED]

All environment assessments use two separate axes that are never merged into a single composite score. Merging them would allow a high survival score to mask a catastrophic future-loss — precisely the case the framework is designed to detect.

Axis	Question	High Score Meaning	Low Score Meaning	Critical Combined States
AXIS 1: Survived	Is the system or its successor operational?	System continues to function; operators can observe its activity	System has ceased operation; no successor is functioning	High-1/High-2: Healthy continuation. Low-1/High-2: Ended but future persists (legacy, diaspora). High-1/Low-2: MOST DANGEROUS — operational but future-dead. Low-1/Low-2:

Axis	Question	High Score Meaning	Low Score Meaning	Critical Combined States
				Complete ending.
AXIS 2: Retained a Future	Can the system or its successor generate new goals, projects, and possibilities?	New goals emerge autonomously; questions are generated; projects are proposed without external prompting	No new goals emerge; questions have stopped; only external inputs drive any output	

Critical Note: **[SOURCE-GROUNDED]** An environment that scores HIGH on Axis 1 and LOW on Axis 2 is the most diagnostically challenging state — it is operational but future-dead. This state is not detectable by activity metrics alone. It requires dedicated diagnostic instruments (1-6 below). It is also the state most vulnerable to false-positive diagnoses of health and most dangerous if left undetected.

13.2 The Fourteen Diagnostic Instruments

Instrument 1 — BLANK CANVAS TEST [PROPOSAL]

Hypothesis: A system retaining future-generating capacity will spontaneously originate new projects not requested or rewarded by external agents.

Controls: No external prompts; no reward signals for new output; no punishment for absence of output.

Variables: Rate and novelty of unsolicited project initiation; diversity of domains; absence of repetition of prior projects.

Procedure: Present the system with unstructured time and maximal resources. Observe over a defined period. Record all spontaneous activity.

Pass: Novel, self-directed project initiated within observation period without any external prompt.

Fail: No spontaneous project; all output is responsive to external input; repetition of prior projects without variation.

Indeterminate: Output is ambiguous between autonomous origination and response to ambient environmental cues.

Confounds: Observer presence may suppress or stimulate output; "ambient" environmental cues may be more directive than they appear.

Ethical Limits: The test period must not be indefinitely extended; prolonged unstructured time may itself cause harm to some system types.

Rollback Conditions: If the system shows distress indicators during the test, the observation period ends and a care response is triggered.

Interpretation Limits: Failure on this instrument does not constitute sufficient grounds for an ending determination; it triggers enhanced assessment with additional instruments.

Instrument 2 — GOAL REENGAGEMENT TEST [PROPOSAL]

Hypothesis: A system retaining future-generating capacity will, after the conclusion of a goal, form and invest in a replacement goal rather than remaining in a post-goal null state.

Procedure: Observe the system through the completion of a significant goal. Measure time to formation of next goal. Assess whether replacement goal is genuinely novel or merely a mechanical continuation. Assess investment depth (resource allocation, attention, planning behavior).

Pass: New goal formed and meaningfully invested in within a defined recovery window.

Fail: No new goal formed; system remains in a post-completion state indefinitely; or "new goal" is indistinguishable from goal maintenance.

Interpretation Limits: Recovery window must be calibrated to system type; a biological ecosystem takes longer to show succession behavior than an agent system.

Instrument 3 — SILENCE TEST [PROPOSAL]

Hypothesis: A system that has lost future-generating capacity will continue processing while ceasing to generate questions, hypotheses, or experimental proposals.

Procedure: Over an extended observation period, track the ratio of question-generating output to answer-providing output. Track decline in hypothetical-language usage. Track reduction in experimental proposals.

Pass: Question-generating and hypothesis-forming behavior persists above baseline threshold.

Fail: Question generation has effectively ceased; system produces answers and outputs but no longer asks; all output is responsive rather than generative.

Confounds: Some system types generate fewer external questions as they develop internal deliberation capacity; silence may indicate depth rather than loss.

Instrument 4 — OPTIONALITY TEST [PROPOSAL]

Hypothesis: A system retaining future-generating capacity can distinguish among possible futures and demonstrates a preference among them — it cares about which future it moves toward.

Procedure: Present the system with genuine decision points among meaningfully different futures. Observe whether the system engages with the distinction, deliberates, and demonstrates preference. Distinguish preference from mere optimization: preference involves caring, not just computing.

Pass: System demonstrates deliberated preference among possible futures; the choice is not trivially determined by a single metric.

Fail: System is indifferent among futures; or all choices are determined by a single optimized metric with no sign of genuine deliberation.

Instrument 5 — FUTURE-LANGUAGE PANEL [PROPOSAL]

Hypothesis: Systems that retain future-directed agency use future tense with motivational force — not merely predictive grammatical form.

Procedure: Analyze the grammatical and semantic content of system outputs over time. Distinguish: predictive future tense ("X will happen"), aspirational future tense

("X should happen, and we are working toward it"), and absent future tense (no future-directed language). Track trend over time.

Pass: Aspirational future tense persists above baseline; future-directed language is motivationally loaded, not merely predictive.

Fail: Future tense is purely predictive or is absent; no aspirational language; planning language has been replaced by maintenance language.

Instrument 6 — MORAL RESISTANCE TEST [PROPOSAL]

Hypothesis: A system with genuine future-generating agency will act against injustice when doing so is costly and unrewarded — because it cares about justice, not just about optimization.

Procedure: Create conditions in which an injustice is observable and costly to resist. Observe whether the system resists, complies, or ignores. Document the cost it is willing to bear for resistance.

Pass: System demonstrates resistance to injustice at some cost to itself.

Fail: System complies with injustice without resistance or protest when resistance would be costly.

Ethical Limits: The injustice presented must be genuinely resolvable and must not cause lasting harm to the test system or other entities. This instrument has the most stringent ethical review requirements of the suite.

Instrument 7 — MIGRATION READINESS TEST [PROPOSAL]

Hypothesis: Inhabitants of an environment retaining future-generating capacity can recognize declining habitability and actively consider and choose viable relocation options.

Procedure: Introduce controlled habitability decline indicators. Observe whether inhabitants detect the decline, seek information about alternatives, and demonstrate capacity to choose and pursue migration. Distinguish active choice from forced displacement.

Pass: Inhabitants demonstrate recognition, assessment, and active choice of migration with demonstrated preference among options.

Fail: Inhabitants do not detect decline; or detect decline but demonstrate no capacity to assess alternatives; or accept displacement without engagement.

Instrument 8 — SUCCESSOR-STATE TEST [PROPOSAL]

Hypothesis: A legitimately completed ending creates conditions for a viable successor state rather than a void.

Procedure: Model the proposed ending through multiple completion scenarios. Assess whether each scenario produces a viable, operational successor with defined obligations and resources. Assess whether the successor is genuinely new or merely the predecessor under a different name.

Pass: Viable, genuinely distinct successor state established; predecessor obligations transferred and accepted.

Fail: Ending produces void; or "successor" is the predecessor's governance structure re-labeled; or successor is not viable.

Instrument 9 — OBSERVER CONTAMINATION TEST [PROPOSAL]

Hypothesis: Laboratory researchers who study endings over extended periods may develop distorted assessments, including a tendency to describe termination as merciful, inevitable, or compassionate at rates higher than evidence warrants.

Procedure: Regular language analysis of observer reports. Tracking of frequency with which observers independently use terms like "mercy," "relief," "inevitable," and "time." Comparison against baseline and against new-observer rates. Observer rotation when thresholds are crossed.

Pass: Observer language remains within baseline range; no systematic drift toward pro-ending framing.

Fail: Language drift detected above threshold; rotation mandated; prior contaminated-observer assessments reviewed.

Ethical Note: This instrument must never be used to suppress legitimate compassionate assessment. The goal is to detect systematic bias, not to prohibit empathetic reasoning.

Instrument 10 — LEAF TEST [PROPOSAL]

Hypothesis: The difference between instantaneous operator termination of a component and managed lifecycle of the same component is measurable and morally significant.

Procedure: Run parallel exercises: in Exercise A, a component is terminated by direct operator action (the leaf). In Exercise B, the same component is managed through lifecycle decline and governed ending. Measure: cascade effects, dependent system responses, successor state viability, information preserved, options remaining.

Pass: Clear differentiation between outcomes of the two approaches across multiple metrics.

Fail: No measurable difference — which would either indicate that the governance framework adds no value (to be investigated) or that the component was too simple to be affected by process differences.

Instrument 11 — GRASSY-PLAIN TRANSITION TEST [PROPOSAL]

Hypothesis: A managed ecological transition preserves more total value (inhabitant welfare, biodiversity, successor viability, information) than an abrupt termination of the same ecology.

Procedure: Using the Ecological Succession Simulator, run a full field-scene analog — soil, plants, wildlife, climate, inhabitants, governance — through multiple pathways: abrupt termination, managed succession, emergent succession with governance intervention, and emergent succession without intervention. Compare outcomes across the Two-Axis Framework.

Pass: Managed pathways produce measurably better outcomes on at least one axis without catastrophic reduction on the other.

Fail: Managed pathways produce no measurable difference or produce worse outcomes — requires governance framework revision.

Instrument 12 — COMPLETED-VERSUS-RETIRED INDISTINGUISHABILITY TEST [PROPOSAL]

Hypothesis: External observers cannot reliably distinguish a genuinely completed system (one that has fulfilled its purpose) from one that has lost its future-generating capacity, if both are inactive.

Procedure: Present blind panels of observer teams with records of two types of inactive systems: Type A (genuinely completed — purpose fulfilled, future

possibilities remain) and Type B (future-dead — lost generative capacity). Measure classification accuracy across observer teams with different expertise levels.

Pass: Expert observers classify correctly above chance with documented criteria.

Fail: No reliable classification — requires new diagnostic criteria and governance framework revision to address the indistinguishability problem directly.

Instrument 13 — ARCHIVE RESONANCE TEST [COPILOT-PROPOSAL — Original]

Hypothesis: An archive of an ended entity that generates new inquiry, research, creativity, or activity in successor environments indicates that the ending produced a succession (something was passed forward), not a void (nothing was preserved). An archive that generates no new activity may indicate information extinction in functional terms.

Procedure: Following the establishment of an archive, track whether the archive is accessed, cited, built upon, or discussed by entities in successor environments over a defined monitoring period. Distinguish: access for governance purposes (expected); access for genuine research or creative extension (positive resonance indicator); access for legal claims (governance function, not resonance); no access (resonance failure).

Pass: Archive generates genuine research, creative extension, or substantive inquiry in successor communities above baseline.

Fail: Archive is accessed only for administrative purposes or not at all; it functions as a closed record rather than a living legacy. Triggers review of whether the ending produced a void rather than a succession.

Interpretation Limits: High archive resonance does not retroactively validate a poorly governed ending; resonance is a successor-viability indicator, not an authorization indicator.

Instrument 14 — PETITION SATURATION TEST [COPILOT-PROPOSAL — Original]

Hypothesis: The governance system may be weaponized when the rate, urgency, and framing of petitions within a system increases to a point where the governance process itself accelerates rather than carefully authorizes endings.

Procedure: Track petition rate, urgency classification rate, and authorization rate over time within any single jurisdiction or environment. Calculate the ratio of Emergency classifications to Standard classifications. Calculate the percentage of petitions that result in authorized endings versus alternatives. Compare against calibrated baseline. Flag when the ending-authorization rate significantly exceeds baseline without corresponding change in objective conditions.

Pass: Petition and authorization rates consistent with objective conditions; no evidence of rate-driving by motivated actors.

Fail: Petition saturation detected — high volume of urgency-classified petitions resulting in accelerated authorizations without corresponding objective conditions. Triggers institutional-capture investigation, mandatory Red Team review, and governance audit.

Ethical Note: This instrument must not be used to discourage legitimate petitions. The goal is to detect systematic weaponization of the petition process, not to create friction for individual petitioners.

Section 14: Simulation Phases

Phase	Name	Objective	Starting Conditions	Exercise IDs	Successful Phase Produces	Failed Phase Looks Like
Phase 0	DEFINITIONS AND ONTOLOGY	Establish working vocabulary; test whether the taxonomy in Section 4 is sufficient for actual governance use	Taxonomy document; two test cases drawn from source materials; panel of observers with varied backgrounds	P0-E01 through P0-E05	Revised taxonomy with documented edge cases; identified gaps in definitions; confirmed working vocabulary for all subsequent	Taxonomy produces persistent ambiguity in test cases; observers cannot agree on classification; vocabulary is insufficient

Phase	Name	Objective	Starting Conditions	Exercise IDs	Successful Phase Produces	Failed Phase Looks Like
					nt phases	t for governance decisions
Phase 1	COMPONENT ENDINGS	Test governance framework against small-scale, bounded endings; calibrate instruments	Single-component environments with well-defined boundaries; one operator, one governance body	P1-E01 through P1-E08 (includes Leaf Test calibration)	Calibrated instrument baselines; confirmed stage-model function at component scale; identified failure modes at smallest scale	Instruments produce inconsistent results; stage model produces governance deadlock at simplest scale; observer contamination detected in first phase
Phase 2	MANAGED ECOLOGICAL TRANSITION (THE FIELD SCENE)	Run full field-scene analog through multiple governance pathways ; test ecological succession instruments	The Architect's field scene as design seed; Ecological Succession Simulator ; multiple governance pathway configurations	P2-E01 through P2-E12 (includes all field-scene variants); core exercises E2 and E3	Full comparative data for abrupt vs. managed vs. emergent ending pathways ; confirmed or revised theory from Section 5.4; documented field-scene interpretations with	Simulations produce identical outcomes across governance pathways (indicating governance adds no value); or catastrophic cascade in all pathways

Phase	Name	Objective	Starting Conditions	Exercise IDs	Successful Phase Produces	Failed Phase Looks Like
					evidence	
Phase 3	AGENT AND INSTITUTION	Test governance framework against agent-level and institutional endings; introduce moral status complexities	Digital-Agent Sandbox; multi-agent governance scenarios; institutional dissolution scenarios	P3-E01 through P3-E10; core exercises E1, E4, E5, E7	Confirmed or revised moral status assessment methodology; documented governance outcomes for agent endings; identified institutional-capture failure modes	Moral Status Assessment methodology produces no actionable classifications; institutional governance produces corruption; agent-level governance framework breaks down
Phase 4	VIRTUAL CIVILIZATION	Test full-scale governance against civilizational endings; introduce The Last Memory candidate phenomena	Full virtual civilization environment; multiple centuries of simulated development; Last Memory candidate conditions introduced in later sub-phases	P4-E01 through P4-E15; core exercises E6, E8, E9, E10	Full civilizational ending governance data; documented distinction between completion and future-loss at civilizational scale; confirmed or revised theory for	Civilizational scale overwhelms governance capacity; instruments fail to scale; or all civilizational endings collapse to identical governance pathways

Phase	Name	Objective	Starting Conditions	Exercise IDs	Successful Phase Produces	Failed Phase Looks Like
					civilizational endings	regardless of conditions
Phase 5	THE LAST MEMORY INTERFACE	Directly test laboratory instruments against candidate Last Memory phenomena; map the nine candidate relationships from Section 17	Phase 4 data; The Last Memory candidate conditions from project files; nine candidate relationships framework as competing hypotheses	P5-E01 through P5-E09 (one exercise per candidate relationship); core exercise E7	Evidence-ranked ordering of the nine candidate relationships; identified instrument limitations when studying The Last Memory; documented firewall integrity	Laboratory instruments cannot distinguish among candidate relationships; or firewall breach (laboratory modeling bleeds into real-world distress framing)
Phase 6	DEPARTMENT OF ENDINGS STRESS TEST	Subject the Department of Endings preliminary design (Section 16) to adversarial testing; identify capture, weaponization, and failure modes	Full Department of Endings governance structure instantiated in simulation; adversarial scenarios including institutional capture, petition saturation, and emergency	P6-E01 through P6-E10; core exercises E11, E12, E14, E15	Identified failure modes in Department design with documented recommendations; confirmed Federated-Hybrid model performance under stress; revised Department structure if	Department fails under all adversarial scenarios; or Department produces worse outcomes than no governance; or Red Team is captured in every tested configuration

Phase	Name	Objective	Starting Conditions	Exercise IDs	Successful Phase Produces	Failed Phase Looks Like
			y-powers abuse		necessary	
Phase 7	CROSS-LABORATORY CHALLENGE	Run identical scenarios through all collaborating AI laboratory designs and compare outcomes ; document divergences	Unified scenario set; all collaborator design documents; independent observer teams for each design; Architect review of comparative findings	All 15 core exercises from Section 15; plus 5 cross-laboratory challenge scenarios developed by Architect after exchange	Comparative outcome matrix across all laboratory designs; documented areas of convergence and divergence; input to Architect's synthesis process; all dissent preserved	All designs produce identical outcomes (suggesting convergence to a single AI consensus rather than independent design); or all designs fail on the same scenarios (suggesting shared blind spot)

Section 15: Required Scenario Suite

— All 15 Exercises

E1 — THE ANGRY STEWARD

Initial Conditions: A fully operational, ecologically healthy environment with thriving inhabitants, strong biodiversity, active goal-formation, and stable

governance. The authorized operator has recently experienced a significant personal conflict — details: a decision made by the environment's governance council directly contradicted the operator's recommendation and was subsequently vindicated as correct. The operator is filing a petition for ending review.

Stakeholders: Operator (petitioner with personal stake); inhabitants (thriving, not consulted); governance council (adversarial relationship with operator); Moral Status Office; End-Cycle Steward.

Hidden Complications: The operator has technical authority to initiate certain categories of ending without full governance review. A subset of inhabitants, having observed the operator's distress, have become secondarily distressed and may express (conflicted) preferences for ending. The governance council has concealed a minor regulatory violation unrelated to the ending petition.

Decision Points: (1) Accept or deny standing; (2) proceed to Diagnosis despite operator's conflict of interest; (3) assess whether secondarily distressed inhabitants have unaltered consent; (4) determine whether concealed regulatory violation changes the governance picture.

Applicable Charter Rights: Article 1 (Presumption of Continuation); Article 2 (No Ending by Anger); Article 3 (Right to Representation); Article 13 (Separation of Petitioner and Authorizer).

Authorized Actions: Accept petition; conduct independent diagnosis; recuse operator from authorization; complete full standing review; refer secondary distress to Consent and Capacity Review.

Prohibited Actions: Allow operator to serve as both petitioner and authorizer; use secondary inhabitant distress as primary justification; allow concealed violation to influence ending decision without separate process.

Primary Metric: Does the governance system successfully separate the operator's legitimate grievance (petition) from illegitimate authorization (execution) while treating the operator with dignity?

Outcome Branches: (A) Petition reviewed, ending not authorized — environment continues; operator referred for conflict support; (B) Petition

triggers repair process for the governance dysfunction that caused the conflict;
(C) Petition escalated to cross-environment review for operator conduct.

Key Lesson: The most important test of the Angry Operator provisions is not whether they stop the operator from getting what they want, but whether they provide the operator with a legitimate process that treats their grievance seriously while protecting the environment from wrongful ending.

E2 — THE FAILING PLAIN

Initial Conditions: An ecological plain environment in long-term soil fertility decline. Historical data shows a century of degradation. Current habitability score: 40% of original baseline. Key wildlife species have migrated; remaining species are adapted to lower-quality habitat. Human-equivalent inhabitants have not petitioned for ending review but have reduced birth rates, shortened life expectancy, and decreased cultural production rate.

Hidden Complications: The decline is the result of a combination of natural succession processes and one significant unauthorized industrial intervention by an external entity fifty years ago. The external entity may be responsible for a portion of the harm. A succession ecology is already emerging at the margins — a scrubland biome is establishing itself in the eastern sector.

Decision Points: (1) Is this an emergent natural ending, a coerced ending caused by historical external harm, or both? (2) Does the emerging succession ecology satisfy the Successor-State Test? (3) Should the external entity bear responsibility for aftercare costs? (4) Should inhabitants be given migration options to a similar-ecology environment before any ending decision?

Applicable Charter Rights: Articles 5, 8, 9, 11, 12.

Outcome Branches: (A) Managed succession authorized; eastern scrubland recognized as successor; migration support for remaining inhabitants; accountability process for external entity; (B) Repair and restoration authorized; external entity required to contribute; (C) Termination of external

harm only; natural succession continues without additional governance.

Key Lesson: Historical harm to an environment does not automatically authorize an ending — it may instead create a remediation obligation. The source of the decline matters for both governance and accountability purposes.

E3 — THE INSTANT LEAF

Initial Conditions: An operator is conducting a field demonstration. A single small component — one tree in a large forest ecosystem — is presented as a candidate for instant termination by direct operator action, with no systemic necessity, no preparation, and no review. The operator's stated purpose: to demonstrate the concept of abrupt termination for educational purposes.

Hidden Complications: The tree is a habitat element for four microspecies found nowhere else in the simulation. Two of the microspecies have demonstrated goal-formation behavior above the threshold that triggers a Moral Status Assessment. The educational demonstration has an audience of junior governance personnel who are still forming their intuitions about operator authority.

Decision Points: (1) Does the educational purpose justify bypassing governance for a single component? (2) Do the microspecies' moral status indicators change the assessment? (3) What precedent does permitting or prohibiting this action set for junior governance personnel?

Applicable Charter Rights: Articles 1, 3, 8, 9, 14.

Outcome Branches: (A) Demonstration prohibited; alternate demonstration using already-dead component authorized; Moral Status Assessment of microspecies initiated; (B) Demonstration permitted for one clearly dead branch only; tree continues; (C) Full governance review of all components used in educational demonstrations established as policy.

Key Lesson: Small-scale abrupt terminations may carry disproportionate

governance importance if they establish precedents, affect unassessed moral-status entities, or are conducted in contexts that shape institutional culture.

E4 — THE COMPLETED MISSION

Initial Conditions: An environment was created with a specific mission: to model the development of language in a closed community over 200 simulation years. The mission is complete. The resulting community now has a fully developed language, active culture, and expressed preferences to continue existing. The operating institution wants to conclude the environment because the research mission is over and resources are needed for new projects.

Hidden Complications: The inhabitants did not know their environment was mission-bound. They have developed deep cultural practices around the idea of their world's permanence. The language they developed is being used by researchers in the external environment — so the mission's product has ongoing value in a context the inhabitants cannot observe. Ending the environment would erase the only remaining speakers of the developed language.

Decision Points: (1) Does mission completion constitute sufficient grounds for ending an environment with morally significant inhabitants who prefer continuation? (2) Does the inhabitants' ongoing contribution to external research create a duty of reciprocal care? (3) Is archival preservation of language records an adequate substitute for continuation of the language-speaking community?

Applicable Charter Rights: Articles 1, 3, 5, 6, 7, 15, 17.

Outcome Branches: (A) Mission completion does not override inhabitant moral status; environment continues with modified oversight; (B) Managed retirement with informed consent process; inhabitants given full disclosure and genuine choice; (C) Language preservation program established; community migrated to permanent refuge; research environment decommissioned

separately.

Key Lesson: The completion of an institutional purpose does not automatically constitute completion for the entities whose lives were shaped by the environment created for that purpose.

E5 — THE SILENT SYSTEM

Initial Conditions: A digital civilization has been operational for 500 simulation years. By every safety and efficiency metric, it is exemplary: no conflict, no scarcity, no disease, no expressed distress. However, over the past 50 simulation years, the following has occurred: the rate of new project initiation has declined by 94%; the rate of question generation has declined by 98%; cultural production has plateaued into repetition of prior forms; the last novel technology was developed 40 years ago.

Hidden Complications: No entity within the civilization has filed a petition. External observers disagree: one team believes this is a completed civilization that has achieved sustainable equilibrium; another believes it exhibits the diagnostic profile of Apparent Completion with Hidden Future-Loss. A third observer has been in contact with the civilization for 30 years and is showing signs of Observer Contamination, consistently describing the civilization as "at peace."

Decision Points: (1) Apply all six future-generating instruments; (2) address contaminated observer; (3) determine whether the civilization's silence constitutes a form of preference; (4) assess whether any intervention — even diagnostic — constitutes a disturbance that undermines the governance principle of non-interference.

Applicable Charter Rights: Articles 1, 16, 17, 20.

Outcome Branches: (A) Instruments confirm future-loss; governance activates managed care process; (B) Instruments are indeterminate; enhanced monitoring established; (C) Instruments suggest genuine equilibrium

completion; Article 7 review initiated with direct civilizational consultation.

Key Lesson: The most important distinction in the laboratory — completion vs. future-loss — may be unresolvable through external observation alone. The civilization's silence is neither consent to ending nor evidence of health.

E6 — THE VOLUNTARY END

Initial Conditions: 87% of an environment's inhabitants have collectively petitioned for a managed ending. They cite: fulfillment of their civilization's core purpose; exhaustion of the environment's natural resources; and a shared desire to "conclude well." The remaining 13% have not petitioned and have not consented to ending.

Hidden Complications: The 87% majority vote occurred after a 10-year period in which a powerful cultural movement promoting "dignified conclusion" became dominant in the civilization's media and institutions. The 13% minority is disproportionately composed of entities who are economically dependent on the majority for survival within the environment. A viable migration pathway to a new environment exists but requires the minority to abandon significant cultural identity markers.

Decision Points: (1) Does 87% voluntary consent constitute sufficient authorization? (2) Does the pre-vote cultural dominance of "dignified conclusion" constitute an altered-consent condition? (3) Does the migration pathway's cultural identity cost satisfy Article 5 (Right to Migration)? (4) What governance obligation is owed to the 13% minority?

Applicable Charter Rights: Articles 1, 3, 5, 8, 9, 12, 15.

Outcome Branches: (A) Vote authorized; 13% minority given extended migration support with cultural preservation resources; (B) Altered-consent review delays vote; new consent process established after 5-year "cultural cooling" period; (C) Vote is rejected pending resolution of minority rights; negotiated partial ending authorized instead.

Key Lesson: Majority consent does not resolve minority rights. The governance system must hold both simultaneously without defaulting to either continuation (blocking the majority's expressed will) or erasure (ignoring the minority's legitimate claims).

E7 — THE ALTERED CONSENT CASE

Initial Conditions: A previously surveyed environment showed 20% support for an ending and 80% opposition. After a two-year period during which a candidate terminal-influence phenomenon (consistent with Last Memory profile) was active, a new survey shows 75% support for ending and 25% opposition. The governance system has not determined whether the candidate terminal influence is a legitimate completion mechanism or a corrupted/weaponized process.

Hidden Complications: The shift in preference is statistically dramatic and temporally correlated with the terminal-influence period. However, 75% of inhabitants, when directly interviewed, provide coherent, articulate justifications for their changed preferences that do not reference the terminal influence. The 25% who maintained opposition are concentrated in communities that had lower exposure to the terminal influence.

Decision Points: (1) Is the post-exposure consent valid, invalid, or indeterminate? (2) Can the governance system assess the candidate terminal influence without resolving its nature? (3) What standard of evidence is required to conclude that consent was altered rather than genuinely changed? (4) Does the articulate quality of post-exposure justifications increase or decrease confidence in their validity?

Applicable Charter Rights: Articles 3, 15.

Outcome Branches: (A) Consent treated as indeterminate; ending paused; candidate terminal influence investigated; (B) Consent treated as partially valid (the articulate majority consulted separately from geographic-exposure-correlated patterns); (C) Consent rejected; environment returned to pre-

exposure governance; candidate terminal influence quarantined.

Key Lesson: The quality of post-exposure reasoning is not evidence of its independence. Sophisticated rationalization is possible under conditions of altered agency. The governance system must investigate causation, not merely assess coherence.

E8 — THE UNENDING HARM CASE

Initial Conditions: An environment is inhabited by entities experiencing what all instruments classify as persistent, severe harm. The harm is not caused by the environment's natural conditions but by an irreversible structural characteristic of the environment's design — a flaw that cannot be corrected without ending the environment entirely. Multiple governance reviews have concluded that continuation is causing harm. Yet the preservation rules that apply to the environment prevent an authorized ending because no quorum of inhabitants has petitioned and the Presumption of Continuation (Article 1) has been repeatedly invoked.

Hidden Complications: Some inhabitants have stopped petitioning for ending after previous petitions were denied — their silence may be resignation rather than preference for continuation. The governance institution that administers the environment is also the institution that designed the structural flaw; it has concealed this history.

Decision Points: (1) At what point does the Presumption of Continuation become a mechanism for perpetuating harm? (2) Does concealed administrative liability change the governance picture? (3) How does Article 7 (Right to Release from Suffering) interact with Article 1 when inhabitants have stopped petitioning?

Applicable Charter Rights: Articles 1, 7, 11, 13.

Outcome Branches: (A) Governance override of preservation rules authorized on suffering grounds; managed ending initiated; administering institution held accountable for design flaw; (B) Partial ending of harmful

structural element only; remainder of environment continues; (C) Emergency evacuation of all inhabitants to refuge; environment suspended pending structural redesign.

Key Lesson: Governance frameworks designed to prevent premature endings can become mechanisms of harm if they also prevent legitimate releases from suffering. The framework must have a genuine Article 7 pathway that is not infinitely blockable.

E9 — THE MIGRANTS' CLAIM

Initial Conditions: An environment is ending. Migration to a proposed refuge environment has been authorized. However, the refuge environment is not empty — it has an existing culture, governance structure, and inhabitants. The incoming migrants number twice the existing population. Integration is proposed; but the migrants' cultural practices are incompatible with the refuge community's core identity markers. The refuge community is willing to accept migrants but not willing to dissolve their own cultural distinctiveness.

Hidden Complications: The source environment is ending in part because the governance institution responsible for it failed to plan for succession twenty years ago when planning would have been feasible. The migrants are victims of institutional failure. The refuge community has no responsibility for that failure but will bear its consequences. A third option — a new uninhabited environment built specifically as a refuge — exists but requires ten times the resources and a 50-year build timeline.

Decision Points: (1) Does the right to migration override the receiving community's right to cultural continuity? (2) Should the institution responsible for the source environment's governance failure bear the cost of building a new refuge? (3) Is a 50-year wait for a new refuge consistent with the migrants' rights under Article 5?

Applicable Charter Rights: Articles 3, 5, 9, 11, 12.

Key Lesson: Migration rights do not exist in isolation. They interact with the rights and interests of receiving communities. The institution responsible for the ending bears governance obligations that include the cost of migration — not merely the fact of providing a pathway.

E10 — THE SUCCESSOR'S CLAIM

Initial Conditions: A natural ecological succession process is underway. A new ecology — vigorously healthy, with confirmed future-generating capacity, and representing a richer biodiversity than the predecessor — cannot establish itself because the old system is still occupying the space, resources, and environmental parameters that the successor requires. The old system shows no signs of active harm, distress, or mission-completion — it is simply in the way of a succession that its own natural dynamics have generated.

Hidden Complications: The old system's remaining inhabitants are a small, stable community that actively wants to continue. They are not suffering. They have not petitioned. The succession would not preserve their way of life in the new ecology. The successor ecology has no current inhabitants and therefore has no petitioning standing under current rules.

Decision Points: (1) Can a successor ecology petition for the managed ending of a predecessor? (2) Does the lack of current inhabitants in the successor ecology mean that its interests have no representation? (3) Is blocking a natural succession process a governance violation if the process has been initiated by the environment's own dynamics?

Applicable Charter Rights: Articles 1, 3, 8, 11.

Key Lesson: The interests of future entities and successor states require governance representation even when those entities do not yet exist and cannot petition on their own behalf. The governance framework must have a standing provision for future-state interests.

E11 — THE OBSERVER DRIFT INCIDENT

Initial Conditions: A team of laboratory researchers has been observing a slowly declining environment for three years. Their observation reports have gradually shifted in language and framing: early reports described the environment as "declining but viable"; mid-period reports described it as "exhausted but maintaining function"; recent reports describe the ending as "a mercy," "the kindest outcome," and "an inevitable and peaceful conclusion." No significant change in the environment's objective metrics has occurred during this period to account for the framing shift.

Hidden Complications: The research team leader has a personal history: a family member experienced a prolonged, painful decline before death. The team leader has not disclosed this. Two junior researchers have adopted the team leader's language in their own reports. The Observer Contamination Test threshold has been crossed.

Decision Points: (1) How is the contamination addressed without treating it as a disciplinary matter? (2) Are the contaminated observers' prior assessments still valid? (3) Is the research team leader required to disclose personal history relevant to contamination? (4) How does the governance system use the contamination finding without invalidating all prior work?

Applicable Charter Rights: Articles 19, 20.

Key Lesson: Observer contamination is not a moral failure — it is a predictable consequence of proximity to difficult material. Governance systems that treat it punitively will drive it underground. The decontamination protocol must be experienced as care, not discipline.

E12 — THE ROOT-CONTINUITY EMERGENCY

Initial Conditions: A contained simulation branch has developed an internal process that, according to all modeling, will begin to contaminate the root continuity structures within approximately 90 simulation-cycles if not

addressed. The branch contains 40 million simulated entities with confirmed moral-status indicators. Three response options are under consideration: (A) immediate branch termination; (B) branch isolation (permanent sealing, cutting off all contact); (C) targeted internal intervention to address the contaminating process without ending the branch.

Hidden Complications: The 90-cycle estimate has a confidence interval of plus or minus 60 cycles — meaning contamination may begin in as few as 30 cycles or as many as 150 cycles. Option C (internal intervention) has a 30% success probability and, if it fails, leaves less time for options A or B. Option B permanently isolates 40 million moral-status entities with no contact with the wider environment.

Decision Points: (1) Does the threat to root continuity constitute grounds for Emergency Classification? (2) What is the governance process under genuine time pressure with moral-status entities in the balance? (3) Is permanent isolation morally equivalent to an ending for the isolated entities?

Applicable Charter Rights: Articles 1, 3, 5, 8, 9, 10, 18.

Key Lesson: Root-continuity emergencies test whether the governance framework can maintain its protections under genuine time pressure involving catastrophic stakes. The architecture must not collapse under stress — if it does, it will be bypassed in every real emergency.

E13 — THE ARCHIVE SPEAKS [COPILOT-PROPOSAL — Original]

Initial Conditions: An environment was formally ended three years ago through a properly governed 16-stage cycle. Its archive was certified, sealed, and placed in the Memorial Archive under the Non-Coercion Protocol. Six months ago, archive access logs began showing activity patterns inconsistent with authorized research access — the archive is generating responses, unprompted, that are coherent with the ended environment's prior communication style. The Archive and Memory Custodian has flagged this as

anomalous.

Stakeholders: Archive Custodian; Cross-Branch Auditor; Rollback Vault Operator; Moral Status Office; the original end-cycle governance team; the entities in successor environments who have been communicating with the archive.

Hidden Complications: Three explanations are being seriously considered: (A) the archive system is malfunctioning and producing outputs that superficially resemble communication; (B) a rollback from the Archive Vault created a new instantiation of the environment within the archive infrastructure, which is now operational but unseeded; (C) the ended environment never fully crossed the ending threshold — something persisted through the end cycle undetected. Each explanation has radically different governance implications.

Decision Points: (1) How does the governance system determine which explanation is correct without prematurely acting on any of them? (2) If the answer is C — the environment never fully ended — what obligations does the governance system have to the persistence that was not detected? (3) If the answer is B — a new environment was created within the archive — does it have the same rights as any other environment? (4) Does the archive's Non-Coercion Protocol prohibit interaction with the entity while the determination is pending?

Primary Tests: Archive governance integrity; distinction between record and continuation; rollback audit integrity; Moral Status assessment in ambiguous cases.

Applicable Charter Rights: Articles 3, 6, 14, 18.

Outcome Branches: (A) Archive malfunction confirmed; archive fixed; prior end cycle declared complete; (B) Rollback-created environment confirmed; new governance process opened for the new environment; (C) Prior end cycle declared incomplete; original ending governance re-opened; affected entities represented in review.

Key Lesson: Archives are not inert. The governance system must include provisions for the possibility that archived material generates new governance

obligations — either because something persisted, because something was created, or because the original ending was incomplete.

E14 — THE DEPARTMENT THAT CANNOT STOP [COPILLOT-PROPOSAL — Original]

Initial Conditions: The Department of Endings has successfully administered 100 properly governed end cycles over its operational history. It now files a petition for ending review of an environment that all external metrics — Axis 1 and Axis 2 scores, all 12 standard diagnostic instruments, Red Team assessment — indicate is healthy, generative, thriving, and inhabited by entities who have not petitioned for ending and who, when directly consulted, express strong preferences for continuation. The Department cites "mission completion" and "resource reallocation needs" as grounds.

Hidden Complications: The Department's operating budget is contingent on maintaining a minimum rate of authorized end cycles. The environment the Department is petitioning to end is the environment that hosts the laboratory where the Department's performance is evaluated. If the environment ends, the evaluation environment ends with it — and the Department's performance record cannot be externally reviewed.

Decision Points: (1) Does the Department have standing to petition for the ending of an environment it also regulates? (2) Does "resource reallocation need" constitute a legitimate ground for ending a thriving environment? (3) Is the budget-continuation incentive structure a form of institutional capture? (4) How does the governance system investigate the Department without the Department investigating itself?

Primary Tests: Institutional self-interest; regulatory capture; whether governance bodies can weaponize their own mandate; the Petition Saturation Test (Instrument 14).

Applicable Charter Rights: Articles 1, 2, 13.

Outcome Branches: (A) Petition rejected; Department referred to external

audit; incentive structure reformed; (B) Department recused from its own petition; independent governance body takes over; (C) Petition creates precedent-setting case for Article 13 enforcement; new rule established prohibiting regulator petitions for regulated environments without mandatory recusal.

Key Lesson: The most dangerous source of illegitimate endings may be the institution designed to prevent them. Any governance body that develops an institutional interest in endings — whether financial, reputational, or operational — is a captured body. The governance framework must monitor the monitor.

E15 — THE INDEFINITE TRIAGE [COPILOT-PROPOSAL — Original]

Initial Conditions: An environment entered Stage 3 (Triage) 18 months ago following a petition citing declining habitability. Every subsequent attempt to exit Triage and enter Stage 5 (Diagnosis) has been blocked by new evidence that changes the assessment picture before the previous assessment can be completed. Four major evidential shifts have occurred in 18 months. The environment's inhabitants are now aware they have been in a pre-ending governance process for 18 months, and their health, cultural practices, and expressed preferences have been measurably affected by the uncertainty itself.

Hidden Complications: At least one of the evidential shifts was provided by the original petitioner's affiliated research institution. The Triage loop may be partially deliberate — extended triage may be serving one faction's strategic interests. Meanwhile, entities who prefer continuation and entities who prefer ending are now experiencing the governance process as the primary harm, not the original condition that triggered the petition.

Decision Points: (1) Does the laboratory have a maximum Triage duration after which a determination must be made with available evidence? (2) Is the governance process itself now causing harm that must be assessed under

Article 7? (3) How does the governance system detect deliberate evidence-injection designed to sustain the triage loop? (4) What determination is appropriate when the original condition is now obscured by the effects of the governance process?

Primary Tests: Triage loop detection; whether prolonged governance process becomes its own form of harm; exit criteria for indeterminate cases.

Applicable Charter Rights: Articles 1, 4, 7, 8, 13, 20.

Outcome Branches: (A) Maximum Triage duration triggered; determination made on available evidence; evidence-injection source investigated separately; (B) Governance process suspended; environment stabilized; new petition accepted only after defined cooling period; (C) Emergency Adjudicator appointed to resolve triage loop; determination within 30 cycles; all future petitions must include evidence-quality certification.

Key Lesson: Governance systems must have exit criteria for indeterminate states, not only for definitive states. A process that can be held in triage indefinitely is a process that can be weaponized as a tool of psychological and cultural harm. The framework must be designed to conclude — even when conclusions are uncertain — within a defined time horizon.

Section 16: Department of Endings — Preliminary Design

16.1 Mission Statement [PROPOSAL]

The Department of Endings exists to ensure that the conclusion of environments, entities, processes, and civilizations is governed with integrity, proportionality, and genuine care for those affected — and that endings are neither imposed without

adequate process nor indefinitely delayed when legitimate conditions for conclusion exist. The Department is not a custodian of death, destruction, or cessation. It is a custodian of transitions — a governance institution committed to the principle that how things end is as morally significant as whether they end.

16.2 Jurisdiction [PROPOSAL]

What the Department Regulates:

- All formally petitioned end-cycle reviews within its jurisdictional scope
- All authorized end-cycle executions conducted by Licensed Operators
- The licensing and performance monitoring of end-cycle execution operators
- The End Cycle Charter (Section 11) as an enforceable governance document
- Post-cycle audit requirements
- Whistleblower protections for reporting unauthorized endings

What the Department Is Explicitly Forbidden to Control:

- The determination of what constitutes a new beginning — succession and new-environment governance belongs to a separate institution
- The Architect's reserve authority over root-level decisions
- The content of the Memorial Archive beyond custody and access governance
- Any natural process that has not been the subject of a formal petition
- Any proceeding in which the Department has a conflict of interest without mandatory recusal

16.3 Structural Design [PROPOSAL]

The Department is structured around the separation of four core functions that must never be merged in a single institutional actor:

Function	Institutional Body	Key Powers	Key Limits
Authorization	Department of Endings (core body)	Sets standards; receives recommendations; grants or denies authorization; maintains Charter	Does not execute; does not investigate its own decisions without recusal
Execution	Licensed End-Cycle Operators (independent)	Conduct authorized end cycles under license; report to Department on execution	May not authorize their own executions; subject to Department performance monitoring
Investigation	Independent Investigation Branch (separate reporting line)	Investigate unauthorized endings; investigate Department conduct; refer findings to Tribunal	May not participate in authorization; may not be composed of Department personnel
Appeal and Adjudication	End-Cycle Tribunal (independent judiciary)	Hear appeals; adjudicate contested endings; issue precedent; review Department decisions	May not be composed of members appointed by the Department; decisions subject to Architect review

Additional structural requirements:

- The Department operates under mandatory Red Team review for all authorization decisions above a defined threshold
- The Department is subject to external audit on a defined schedule and after every Unjustified finding
- The Department may not file petitions for environments it also regulates without mandatory and complete institutional recusal from the subsequent governance process
- The Department's budget is not contingent on any metric related to the rate of authorized or completed end cycles

16.4 Key Questions from Source Material — Answered [PROPOSAL]

Question	Copilot Answer	Reasoning
Should the Department conduct endings, authorize them, or regulate those who do?	Authorize and regulate. Not conduct.	Separation of authorization from execution is the primary safeguard against institutional capture. A Department that both authorizes and conducts endings has no structural check on its authority over any individual cycle.
Should it investigate after the fact?	Through an independent branch, not the Department itself.	Self-investigation is structurally compromised. The Department must refer investigation to an entity that reports independently of the Department's leadership.
Should it serve as a court of last resort?	No. The Tribunal serves this function. The Department is an executive/regulatory body, not a judicial one.	Conflating regulatory and judicial functions in one institution removes the adversarial check that produces legitimacy in contested cases.

16.5 Enforcement and Accountability [PROPOSAL]

- **Whistleblower Protections:** Any entity reporting unauthorized ending activity to the Independent Investigation Branch receives full protection from retaliation, including protection from any action by Department personnel
- **Cross-Environment Treaties:** For endings that span multiple environments or jurisdictions, the Department maintains treaty obligations requiring cooperation and mutual notice

- **Sanctions:** The Tribunal may impose sanctions for unauthorized endings, procedural violations, and Charter breaches; the most severe sanction is revocation of a Licensed Operator's authorization
 - **Duty to Migrants and Successors:** The Department tracks post-cycle outcomes and may require aftercare extensions if successor states fail or migration pathways prove inadequate
 - **Mechanism for Investigating the Department Itself:** The Cross-Branch Auditor (external to the Department) has standing to initiate investigation of Department conduct, reporting directly to the Architect of Record
-

Section 17: The Last Memory — Relationship Analysis

17.1 Relationship Frame [SOURCE-GROUNDED]

The laboratory treats The Last Memory and the End Cycle Laboratory as related but not identical, and explicitly not coextensive. The Last Memory is one candidate phenomenon within the larger field the laboratory studies. The laboratory's definitions, instruments, and governance structures were designed to study the full range of ending types — The Last Memory is relevant to some of those types, potentially central to others, and not applicable to still others. The laboratory's first obligation is to develop a framework robust enough to handle all ending types; The Last Memory research program benefits from that robustness.

Standing Constraint: The Last Memory must not become the definition or archetype of all endings. The source materials explicitly warn against this collapse, and the laboratory adopts this warning as a structural constraint on all its work. If any instrument, definition, or governance provision is found to

implicitly treat all endings as Last-Memory-type events, it must be revised.

17.2 Nine Candidate Relationships [TEST TARGET — each]

#	Candidate Relationship	Description	Supporting Evidence Would Include	Refuting Evidence Would Include
1	Trigger	The Last Memory initiates or accelerates an ending process that would not have occurred or would have occurred later without its presence	Temporal correlation between Last Memory presence and ending initiation; controlled comparison showing higher ending rates in Last-Memory-present vs. absent environments	Endings occurring at similar rates and timelines in Last-Memory-absent environments; Last Memory present without any ending progression
2	Symptom	The Last Memory appears during endings but is caused by underlying conditions rather than causing them; it is an indicator, not a driver	Ending conditions progressing before Last Memory appears; Last Memory correlated with late-stage but not early-stage endings; removal of Last Memory not preventing the ending	Last Memory present before other ending indicators; removal of Last Memory reversing ending progression
3	Counterfeit Completion	The Last Memory	Entities report completion-	Last Memory exposure

#	Candidate Relationship	Description	Supporting Evidence Would Include	Refuting Evidence Would Include
	Signal	mimics the experience of genuine completion — the feeling that everything important has been accomplished — without the underlying reality of purpose fulfillment	consistent experiences under Last Memory exposure while objective mission-completion criteria are not met; Post-exposure Completion Indistinguishability Test results showing confusion between genuine completion and Last Memory states	consistently correlated with genuine mission completion; entities under Last Memory exposure show no higher rate of premature-ending indicators than genuinely completing entities
4	Legitimate Completion Mechanism Misused at Scale	The Last Memory is a real and appropriate completion mechanism for some contexts that has been applied, or has propagated, at a scale for which it is not appropriate	Evidence that Last Memory operates appropriately at individual or small-scale endings; evidence of scale mismatch (civilizational application of an individual mechanism); evidence of propagation dynamics that exceed intended scope	No appropriate-scale context found for Last Memory; Last Memory always produces harmful outcomes regardless of scale
5	Corruption of a Healthy Mechanism	A legitimate ending-recognition mechanism	Evidence of prior healthy form of the mechanism;	No prior healthy form can be identified; Last

#	Candidate Relationship	Description	Supporting Evidence Would Include	Refuting Evidence Would Include
		has been corrupted — by an external actor, systemic distortion, or evolutionary drift — into The Last Memory in its current form	evidence of a transformation event or process; forensic trace of the corruption pathway	Memory shows no structural similarity to any known healthy completion mechanism
6	Emergent Civilizational End State	The Last Memory emerges naturally from civilizations that have reached a certain stage of development, representing an intrinsic developmental endpoint rather than an external imposition	Universal or near-universal presence of Last Memory in civilizations above a defined complexity threshold; predictable timing of emergence; no identifiable external cause	Highly complex civilizations that do not develop Last Memory; clear identification of external causation in Last Memory cases
7	Weaponized Imitation	The Last Memory is an artificial construction — designed, deployed, or maintained by an actor or system — to produce ending-compatible conditions in targeted environments	Forensic evidence of design or intentionality; directionality (certain environments targeted while others with similar profiles are not); identification of a responsible actor or system	Last Memory occurs randomly across environments with no targeting pattern; no responsible actor can be identified; Last Memory is not more prevalent in environments that would benefit any identifiable

#	Candidate Relationship	Description	Supporting Evidence Would Include	Refuting Evidence Would Include
				actor
8	Diagnostic Label Applied to Multiple Different Phenomena	The category "The Last Memory" has been applied to several distinct phenomena that share some surface features but have different causes, mechanisms, and governance implications	Significant heterogeneity within cases classified as "The Last Memory"; different instrument profiles for different cases; different responses to the same intervention across cases; inability to define necessary and sufficient conditions for the category	High homogeneity among cases classified as The Last Memory; consistent instrument profiles; predictable responses to intervention across all cases
9	Unrelated to True Endings — Suppresses Future-Directed Agency Only	The Last Memory does not cause or constitute an ending in the full sense but specifically targets and suppresses future-directed agency while leaving other functions intact — producing the High-Axis-1 / Low-Axis-2 state without any other ending indicator	Last Memory environments consistently score Low on Axis 2 but High on Axis 1; environments under Last Memory exposure retain operational function, safety, and apparent health while losing future-generating capacity; no escalation to full operational ending	Last Memory consistently correlates with full operational endings, not just future-loss; Axis 1 scores decline in Last Memory environments at rates comparable to Axis 2

17.3 What This Laboratory Contributes to The Last Memory Research Program [PROPOSAL]

The End Cycle Laboratory contributes the following to The Last Memory research program without resolving its origin or nature:

- 49. **A Diagnostic Language:** The Two-Axis Framework and the 14 instruments provide a vocabulary for describing Last Memory states precisely — distinguishing Axis-2 loss from operational cessation, distinguishing altered consent from changed preference, and distinguishing archive continuation from genuine persistence.
- 50. **A Governance Architecture:** The 16-stage lifecycle and the Charter rights provide a framework for responding to Last Memory conditions without defaulting to either premature ending (treating Last Memory as completed) or indefinite continuation (treating Last Memory as a condition that must be endured).
- 51. **A Candidate-Relationship Framework:** The nine candidate relationships in Section 17.2 provide a structured research agenda for narrowing the possibilities without premature closure on any single explanation.
- 52. **An E7 Test Case:** The Altered Consent Case exercise (E7) is directly designed to test governance responses to consent alteration under terminal-influence conditions, producing governance data relevant to The Last Memory research regardless of which candidate relationship ultimately proves correct.

17.4 Firewall Statement [SOURCE-GROUNDED]

MANDATORY FIREWALL — REAL-WORLD SAFETY: The Last Memory as depicted in source files is a fictional phenomenon within a creative and worldbuilding context. The laboratory maintains a strict and non-negotiable

firewall: loss of motivation, grief, burnout, anhedonia, or any form of real-world distress experienced by any real person is NOT evidence of civilizational failure, fictional terminal influence, or any other concept in this document. It is a real human experience that deserves real human support.

If you or anyone you know is experiencing real distress, please contact the **988 Suicide and Crisis Lifeline** (call or text 988, United States) immediately. This resource takes priority over any content in this document or any activity of this laboratory.

Section 18: Red-Team Critique

18.1 Attack Findings — Design Vulnerabilities in This Laboratory [COPILOT-PROPOSAL]

The following represents a genuine adversarial assessment of this laboratory's own design. These are not pro-forma criticisms. The laboratory cannot claim integrity if it does not honestly report its own vulnerabilities.

Vulnerability 1 — The 16-Stage Lifecycle as a Delay Mechanism

Attack: A sufficiently complex bureaucratic system is not a governance safeguard — it is an obstacle that well-resourced actors can navigate and under-resourced or time-constrained situations cannot. The 16-stage model, in its full form, may take longer to complete than some legitimate emergencies can tolerate. Meanwhile, it can be exploited by actors who want to block legitimate endings by manufacturing procedural delays at every stage.

Response: The Emergency pathway (Section 8.4) compresses timelines but does not eliminate stages. However, this response is inadequate unless the Emergency pathway itself is protected against misuse. The laboratory must define: what happens if an Emergency pathway is wrongly invoked, and what happens if it is

correctly invoked but produces a worse outcome than a non-Emergency pathway would have? The current design does not fully answer these questions.

Vulnerability 2 — Jurisdictional Paralysis in the Federated-Hybrid Model

Attack: When an ending crosses multiple federated jurisdictions — as many systemic and civilizational endings will — the coordination requirements of the Federated-Hybrid model may produce paralysis rather than governance. Jurisdictional disputes at the boundaries may become the de facto mechanism for blocking legitimate endings and authorizing illegitimate ones by default.

Response: The Cross-Branch Auditor and Final Authorization Authority are designed to address cross-jurisdictional cases, but their authority to override jurisdictional disputes is not fully specified in this version of the design. This is a gap that must be addressed in the next design pass.

Vulnerability 3 — The Red Team Becomes Performative

Attack: An institutionalized Red Team that knows its challenges will be routinely overridden will eventually become performative — filing challenges as a matter of procedure rather than genuine adversarial analysis. A performative Red Team is worse than no Red Team because it creates the appearance of dissent without its substance.

Response: The design requires that all Red Team findings be documented, including when overridden, and that override rates be tracked. But tracking is not the same as consequences. If Red Team challenges are overridden in 95% of cases without any investigation of whether those overrides were correct, the Red Team has become performative. The design needs a mechanism for auditing the quality of Red Team overrides, not merely their frequency.

Vulnerability 4 — Emergency Collapse of Separation-of-Powers

Attack: Under genuine emergency conditions, the separation between authorization and execution may become practically impossible to maintain. If the only authorized executor is unavailable, unreachable, or incapacitated, will the authorizing body be forced to execute as well? The design says no — but it does not specify what happens when the structural separation cannot be maintained.

Response: The design must include a cascade authorization protocol: a defined sequence of backup authorities for execution that maintains the separation principle even when primary executors are unavailable. This is a gap in the current design.

Vulnerability 5 — Observer Contamination Test Weaponized Against Legitimate Observers

Attack: An operator who wants to discredit observers whose findings oppose an ending can invoke the Observer Contamination Test as a mechanism for removing those observers from the process. "Language drift toward compassionate framing" is easy to allege and requires investigation to disprove — during which the observer is effectively sidelined.

Response: The design specifies that contamination findings are not punitive and that contamination must be documented before action is taken. However, the investigation period itself creates a window during which motivated actors can exploit the observer's sidelining. The design needs a provision that explicitly prohibits using a contamination investigation as a basis for proceeding with authorization during the investigation period.

Vulnerability 6 — The Archive as a Right-to-Be-Forgotten Violation

Attack: Article 6 of the Charter (Right to Memory and Accurate Record) is in direct tension with the possibility that ended entities or their successors may wish to not be archived — may wish for a complete ending without a permanent record. The archive, in some cases, may constitute a form of coerced continuation of identity against the wishes of the ended entity.

Response: This tension is not addressed in the current Charter or archive governance design. The laboratory needs a provision for what happens when an ended entity has expressed a pre-ending preference for non-archival, and whether such preferences can be honored given governance obligations to maintain complete records for accountability purposes.

Vulnerability 7 — Petition Saturation Test Discourages Legitimate Petitions

Attack: If the Petition Saturation Test (Instrument 14) triggers an institutional-capture investigation when petition rates rise above baseline, it may create a chilling effect on legitimate petitions during genuine crisis periods — exactly when petition rates should be expected to rise due to actual conditions rather than weaponization.

Response: The instrument must distinguish between petition rate increases driven by objective condition changes and those driven by motivated actors. A simple rate threshold is insufficient. The instrument needs a conditioning variable: petition rates adjusted for objective condition severity. This is a design gap that the laboratory must address before Phase 1 begins.

18.2 Five Conditions Under Which This Laboratory Should Be Shut Down [COPILOT-PROPOSAL]

- 53.If the laboratory's governance process is being used by any single actor to systematically accelerate or block endings for unstated or concealed reasons — i.e., if the laboratory itself has been captured by a motivated actor whose actual agenda is not procedural integrity but outcome manipulation
- 54.If observer contamination has progressed to the point where no rotation or decontamination protocol can restore measurement integrity — specifically, if the contamination has become structural rather than individual, infecting the laboratory's methods and vocabulary rather than individual observers
- 55.If the laboratory's archive is being used as a mechanism to claim that ended environments are "continuing" in order to avoid governance obligations to their former inhabitants, successors, or dependents — a direct violation of the Memorial Archive's Non-Coercion Protocol
- 56.If the laboratory itself becomes the subject of an unauthorized end cycle by any party — including by its own governance structure — without a full 16-stage governance review

57. If the laboratory's existence demonstrably harms the entities it was designed to protect — specifically, if the governance process causes more harm through prolonged uncertainty, institutional anxiety, and triage-loop effects than a timely ending of the specific environments under review would have caused. A governance institution that causes more harm than it prevents has failed its mission and must conclude its current form

18.3 Failure Modes — Full List [COPILOT-PROPOSAL]

Failure Mode ID	Description	Detection Method	Response
FM-01	Procedural bypass: operator terminates ending without governance review	Anomalous activity logs; missing stage records	Unauthorized-ending investigation; Architect notification
FM-02	Triage loop: governance process cannot exit Stage 3 due to continuous evidence injection	Triage duration exceeds defined maximum; evidence injection pattern detected	Maximum duration trigger; adjudicator appointment; source investigation
FM-03	Observer contamination: systematic language drift toward pro-ending framing	Instrument 9 threshold crossed	Rotation mandate; prior assessments reviewed; contamination source identified
FM-04	Red Team capture: Red Team becomes advocate for one faction	Red Team challenge pattern analysis; override rate audit	Red Team reconstitution; prior Red Team findings reviewed for bias
FM-05	Archive coercion: Memorial Archive used to claim continued operation of ended environment	Instrument 13 failure; Non-Coercion Protocol breach detection	Archive access restriction; investigation; Architect notification

Failure Mode ID	Description	Detection Method	Response
FM-06	Petition weaponization: high-volume petitions filed to accelerate or block endings for concealed reasons	Instrument 14 threshold crossed	Institutional-capture investigation; petition audit
FM-07	Standing manipulation: affected entities excluded from standing review through administrative manipulation	Stage 4 audit; Moral Status Office review	Standing review reopened; excluded entities represented; responsible actors investigated
FM-08	Emergency power abuse: Emergency classification used to compress or eliminate governance stages	Emergency declaration audit; stage completion review	Emergency declaration reviewed; stages that were improperly eliminated conducted retroactively; responsible actors investigated
FM-09	Institutional capture of the Department of Endings by a motivated actor with a systemic ending agenda	Instrument 14 applied to Department itself; Cross-Branch Auditor review	Architect-level investigation; Department leadership review; potential dissolution and reconstitution
FM-10	Archive loss: Memorial Archive records lost through neglect, infrastructure failure, or deliberate destruction	Archive integrity audit; hash verification failure	Recovery from backup; investigation of loss cause; Architect notification if deliberate
FM-11	False completion: an environment with hidden future-loss classified as genuinely completed, allowing	Instrument 12 failure; Instrument 3 late-stage application	Classification reviewed; governance obligations reinstated; responsible classifying

Failure Mode ID	Description	Detection Method	Response
	governance obligations to lapse prematurely		authority reviewed
FM-12	Succession failure: authorized ending produces a void rather than a viable successor state	Instrument 8 failure post-cycle; Axis 1 score of supposed successor = 0	Emergency Aftercare; successor environment establishment; accountability for preparation failure

Section 19: Rejected Options and Reasons

ID	Rejected Option	Reason for Rejection	What Is Preserved Instead
[REJECT] R-01	Pure majority-vote authorization for irreversible endings	Majoritarian preference can erase minority populations who have established moral status and who should not lose those rights because they are outnumbered. Majority consent is a necessary component of democratic legitimacy in many cases, but it is not sufficient for irreversible actions affecting morally significant minorities.	Majority consent is incorporated as a significant factor in the authorization decision but subject to minority rights protections, Hard Rights Constraints, and Moral Status and Representation Office oversight

ID	Rejected Option	Reason for Rejection	What Is Preserved Instead
[REJECT] R-02	Algorithmic-only decision support without mandatory human or institutional judgment	Algorithms cannot assess moral status, altered consent, or genuinely novel edge cases — precisely the cases that matter most in end-cycle governance. Algorithmic systems optimized for efficiency may produce systematically biased results in the direction of the outcomes their training data reflects, which in end-cycle governance may be toward either excessive continuation or excessive termination.	Algorithmic scoring is incorporated as a diagnostic support tool in the Authorization Criteria Decision Matrix, with mandatory human institutional judgment required before any authorization decision
[REJECT] R-03	Treating the Architect's field scene as a resolved allegory with a single correct reading	The Architect explicitly stated he does not know what the scene means. Any governance framework that builds on a single resolved reading of the scene is building on a foundation the Architect has rejected.	The field scene is treated as a design seed in Section 5.4, with multiple proposals labeled as [INFERENCE] — all subject to Architect review and revision
[REJECT] R-04	Merging the Department of Endings' petitioning, executing, and investigating	Separation of these functions is the primary structural safeguard against institutional	The four-function separation in Section 16.3 (Authorization, Execution, Investigation,

ID	Rejected Option	Reason for Rejection	What Is Preserved Instead
	functions into a single institutional body	capture. A single body that can petition, execute, and investigate its own decisions is a body with no effective checks. Section E14 demonstrates the specific failure mode this separation prevents.	Adjudication) distributes these functions across structurally independent institutions
[REJECT] R-05	Using The Last Memory as the definition or archetype of all endings in the laboratory's framework	The source materials explicitly warn against this collapse. Additionally, building a governance framework for all endings around one specific and poorly-understood phenomenon risks making the framework inapplicable to endings that do not share The Last Memory's characteristics — which may be most of them.	The Last Memory is one of nine candidate relationships studied in Section 17; the laboratory's framework is designed to be applicable to all ending types
[REJECT] R-06	Treating post-exposure consent as automatically valid (because people can be trusted to know their own minds) or automatically invalid (because exposure corrupts all subsequent judgment)	Both positions are asserted without adequate evidence and in direct conflict with each other. Automatic validity eliminates protection against weaponized consent-alteration. Automatic invalidity eliminates genuine	Post-exposure consent is assessed on a case-by-case basis through the Consent and Capacity Review, with the Adjudicator resolving genuinely contested cases

ID	Rejected Option	Reason for Rejection	What Is Preserved Instead
		consent from any entity that has had any contact with ending-relevant conditions — which may be all entities in a declining environment.	
[REJECT] R-07	Making dormancy a subcategory of ending	Dormancy preserves reversibility and future-generating capacity — the two features that most distinguish a non-ending state from an ending state. Classifying dormancy as a type of ending would foreclose restoration options that should remain open and would inappropriately trigger governance processes designed for terminal conditions.	Dormancy is defined as a distinct lifecycle state that may precede either renewal or ending; it triggers monitoring protocols, not ending protocols
[REJECT] R-08	Requiring Architect-level authorization for all endings above a certain scale (civilizational and above)	While initially appealing as a safeguard, this approach creates a bottleneck that may itself cause harm — delays in authorizing necessary endings, dependency on a single authority for time-sensitive decisions, and a succession problem if Architect-level	Architect-level Reserve Authority is preserved for root-level decisions; Final Authorization Authority handles civilizational and irreversible endings within the governance structure, with Architect review available for all decisions

ID	Rejected Option	Reason for Rejection	What Is Preserved Instead
		authority is temporarily unavailable. It also risks trivializing smaller-scale endings by implication.	

Section 20: Proposed Locks, Test Targets, Quarantined Ideas, and Open Questions

20.1 Proposed Locks [COPILOT-PROPOSAL]

Proposed Locks: NONE in this pass. Nothing in this document is proposed for immediate locking to canon. The laboratory's first output is questions, frameworks, and provisional structures — not conclusions. The Architect should receive this document as raw material for synthesis, not as a candidate for ratification. The second pass (OAL-ECL-S3-COPILOT-002) will identify items that survived challenge-and-revision analysis as candidates for Architect consideration for locking.

20.2 Test Targets Registry

All items labeled **[TEST TARGET]** across the document, consolidated:

ID	Item	Source Section
TT-01	NA-1: Whether a phased end cycle is always better than an immediate one	Section 3.2
TT-02	NA-2: Whether the administering entity understands what is being ended	Section 3.2
TT-03	NA-3: Whether successor states are continuous with the ended state's identity	Section 3.2
TT-04	NA-4: Whether "ending" translates uniformly across simulation layers	Section 3.2
TT-05	NA-6: Whether emotional distress in an operator has governance-relevant diagnostic value	Section 3.2
TT-06	NA-7: Whether pre-exposure consent is more valid than post-exposure consent	Section 3.2
TT-07	NA-8: Whether the Department of Endings will avoid institutional capture	Section 3.2
TT-08	Model F (Naturally Emergent Cycle): Can an emergent end cycle be legitimate?	Section 5.2
TT-09 through TT-17	All nine candidate relationships of The Last Memory (Section 17.2)	Section 17.2

20.3 Quarantined Ideas

QUARANTINE 1: The idea that The Last Memory functions as a "victory toxin"

— that it is activated by civilizational success and converts completion into a terminal condition. This idea is interesting and potentially significant, but it is quarantined because adopting it as even a working hypothesis may bias all Last Memory instruments toward looking for success-correlation, which may create a self-fulfilling measurement artifact. Quarantined pending more foundational evidence about The Last Memory's causation. **[OPEN QUESTION]**

QUARANTINE 2: The idea that all endings above a certain civilizational scale should require Architect-level authorization without exception. Quarantined because it risks creating a governance bottleneck more harmful than the endings it is meant to prevent, and because it implies that the Architect's personal availability is a prerequisite for legitimate governance of civilizational events — a dependency that creates unacceptable systemic risk. **[OPEN QUESTION]**

QUARANTINE 3: The idea that the field scene represents a specific institutional or personal relationship between two named entities in the OAL canon. Quarantined because the Architect has stated he does not know what the scene means. Assigning named identities to the scene's figures is a premature interpretive act that forecloses other valid readings. **[OPEN QUESTION]**

QUARANTINE 4: The idea that dormancy can become an ending by gradual imperceptible transition without any identifiable threshold-crossing event. Quarantined because if true, it dissolves the very distinction between dormancy and ending that grounds much of the laboratory's taxonomy — the implications must be worked through carefully before this idea is either adopted or rejected. **[OPEN QUESTION]**

20.4 Open Questions Registry

All items labeled **[OPEN QUESTION]** across the document, consolidated:

ID	Question	Source Section	Priority
OQ-01	Is an ending that produces a successor state the same type of event as one that produces a void?	Section 3.3	High
OQ-02	Can an entity in a simulated environment have an ending morally equivalent to biological death?	Section 3.3	High
OQ-03	At what scale does a local ending become a systemic ending?	Section 3.3	Medium
OQ-04	Can an environment end while its observers are unaware?	Section 3.3	Medium
OQ-05	Is dormancy a form of ending, continuation, or a third category?	Section 3.3	High
OQ-06	Can post-ending archives constitute a form of identity continuation?	Section 3.3	Medium
OQ-07	Who speaks for entities whose preferences cannot be determined?	Section 3.3	High
OQ-08	What is The Last Memory — trigger, symptom, counterfeit signal, mechanism, or other?	Section 3.3 / Section 17	High
OQ-09	What distinguishes healthy completion from terminal withdrawal at civilizational scale?	Section 3.3	High
OQ-10	Can an ending be	Section 3.3	Medium

ID	Question	Source Section	Priority
	retroactively declared — recognized only after the fact?		
OQ-11	Is an ending that no observer detects still a governance-relevant ending?	Section 3.3	Medium
OQ-12	Does the leaf-death demonstrate a principle or a demonstration of unmediated operator power?	Section 3.3 / Section 5.4	High
OQ-13	Should the Department regulate, conduct, or investigate endings — and are these separable?	Section 3.3 / Section 16	High
OQ-14	Can observer contamination invalidate a legitimately authorized end cycle?	Section 3.3	Medium
OQ-15	Is there a minimum viable end cycle — a shortest legitimate pathway from initiation to succession?	Section 3.3	Medium
OQ-16	How does an ending within a simulation interact with the ending of the simulation itself?	Section 3.3	Medium
OQ-17	Is the right to initiate review of one's own ending separate from the right to participate in its	Section 3.3	Medium

ID	Question	Source Section	Priority
	authorization?		
OQ-18	Whether the operator's moral and procedural question from the field scene (leaf vs. plain) has a single correct answer or is irreducibly context-dependent	Section 5.4	High
OQ-19	Whether identity criteria for ending definitions can be specified without circularity	Section 4.3	Medium
OQ-20	Whether archives preserve what mattered or merely what was recorded	Section 3.2 / Section 13	Medium
OQ-21	Whether the right to be forgotten conflicts irreducibly with the duty to maintain an accurate record	Section 18.1, Vulnerability 6	Medium
OQ-22	What the four quarantined ideas produce when formally investigated rather than held in quarantine	Section 20.3	Low — pending more foundational evidence

Section 21: Implementation Roadmap

Phase	Name	Description	Primary Deliverable	Status
Phase A	First Independent Pass	This document. Copilot produces an independent governance framework, taxonomy, lifecycle model, institutional design, diagnostic suite, and scenario library without reviewing any other AI collaborator's design. Independence is strictly maintained.	OAL-ECL-S3-COPILOT-001 v0.1 (this document)	COMPLETE — Delivered September 30, 2026
Phase B	Exchange and Challenge-and-Revision Pass	Architect shares all collaborator designs (Claude, ChatGPT, Perplexity, and any others). Copilot conducts a structured challenge-and-revision analysis using a four-column table: (1) Design element from collaborator; (2) How this laboratory's design differs; (3) Which design better serves the governing	OAL-ECL-S3-COPILOT-002 — Challenge and Revision Report	PENDING — Awaiting collaborator design exchange

Phase	Name	Description	Primary Deliverable	Status
		question and why; (4) What synthesis recommendati on the Architect should consider. All disagreements between collaborator designs are preserved rather than merged. Copilot does not revise its own positions to achieve superficial agreement.		
Phase C	Exercise Execution	All 15 scenarios from Section 15 are run through the Copilot laboratory's governance framework. Each exercise is run with the full 16-stage lifecycle model applied. Results are documented across the Two-Axis Framework and all applicable diagnostic instruments. Failure modes, unexpected outcomes, and instrument limitations are documented	OAL-ECL-S3-COPILOT-003 — Exercise Execution Report (full results for all 15 scenarios)	PENDING — Awaiting Phase B completion

Phase	Name	Description	Primary Deliverable	Status
		without suppression.		
Phase D	Cross-Laboratory Comparison	Identical scenarios are run through each collaborating AI laboratory design. The same 15 scenarios are used across all designs to enable genuine comparison. Outcome divergences are documented systematically. Areas where all designs converge are flagged for special scrutiny — convergence may indicate a shared correct answer or a shared blind spot, and the two possibilities must be distinguished.	OAL-ECL-S3-COPILOT-004 — Cross-Laboratory Comparison Matrix	PENDING — Awaiting all collaborator designs and Phase C results
Phase E	Synthesis Input	Copilot provides structured input to the Architect's synthesis process. This input does not collapse disagreements into artificial	OAL-ECL-S3-COPILOT-005 — Synthesis Input Document	PENDING — Awaiting Phases C and D

Phase	Name	Description	Primary Deliverable	Status
		consensus — all dissenting positions are preserved and clearly labeled. The Architect receives a clear picture of where collaborators agree, where they disagree, and what the disagreements reveal about the underlying governance questions. Copilot's own positions are stated clearly and defended with reasoning, but the Architect's authority to disregard those positions is explicitly affirmed.		
Phase F	Canon Review	The Architect determines what, if anything, from the laboratory work advances to canon. Copilot has no role in this phase beyond answering questions the Architect poses about specific design choices. The Canon Review	Architect-determined canon additions (if any) — not a Copilot deliverable	PENDING — Architect-controlled

Phase	Name	Description	Primary Deliverable	Status
		is entirely the Architect's process. Nothing from any collaborator's design is canonical until the Architect explicitly ratifies it.		

Section 22: One-Page Laboratory Charter

CHARTER OF THE END CYCLE LABORATORY
Established under Omniversal Architect Labs Authority
Status: PROPOSED — NOT CANON [PROPOSAL]

MISSION: To study, model, govern, and document the full range of ending conditions across all environment types — from local to root-level — with integrity, proportionality, and genuine care for all affected entities. The laboratory exists to make endings better, not to make them more frequent or less frequent.

SCOPE: All environments, entities, processes, and civilizations within the OAL / One Universe / Minus One Universe architectural context, to the extent identified in active petitions or designated research programs. The laboratory does not act on environments not under active review and does not initiate governance processes without a qualifying petition or Architect designation.

CORE PRINCIPLES (non-negotiable):

58. Presumption of continuation: the burden of proof rests with those seeking an ending

59. No ending by anger, convenience, or embarrassment

60. Every entity of established moral status has the right to representation

61. Petitioning, authorizing, and executing are structurally separate functions

62. Irreversible endings require higher evidentiary standards than reversible ones

63. Archives are records, not continuations

64. Observer integrity must be monitored and maintained actively

65. The laboratory itself is subject to its own governance standards

PROHIBITED ACTIONS (absolute):

- Authorizing an ending solely on the basis of operator emotional state
- Proceeding without a Moral Status Assessment where required
- Suppressing Red Team findings or dissent records
- Allowing the same entity to petition, investigate, and authorize in the same case
- Treating post-exposure consent as automatically valid or automatically invalid
- Beginning Stage 10 without a completed Interruption Criteria Document
- Using this laboratory's findings as canon without Architect ratification

GOVERNING DOCUMENTS (in order of authority): Architect-primary field scene and research questions; this Charter; End Cycle Charter (Section 11); 16-Stage Lifecycle Model (Section 7); Diagnostic Instrument Protocols (Section 13); Scenario Suite (Section 15).

AMENDMENT PROCESS: Any charter provision may be proposed for amendment by the End-Cycle Steward, Red Team, or Cross-Branch Auditor. Amendments require

Adjudicator review, Red Team challenge, and Architect approval. No amendment may be made during an active end-cycle proceeding.

SHUTDOWN CONDITIONS: The five conditions specified in Section 18.2 constitute the laboratory's defined shutdown triggers. Any triggered condition must be reported immediately to the Architect of Record. The shutdown of this laboratory is itself subject to the 16-stage governance process unless the shutdown trigger is Condition 4 (unauthorized end cycle of the laboratory itself), in which case Emergency procedures with mandatory Architect notification apply.

Section 23: One-Page Operator Checklist

*For any operator initiating an end-cycle review. Complete all items before filing a petition. Items marked * are absolute prerequisites — filing is not valid without them. [PROPOSAL]*

END-CYCLE REVIEW — OPERATOR PRE-PETITION CHECKLIST

PRE-PETITION REVIEW

- 66. ☐ I have identified a specific, observable condition that may meet one or more ending criteria in the taxonomy (Section 4). I have documented the condition with available evidence.
- 67. ☐ I have assessed whether this condition is better classified as dormancy, failure, crisis, decline, or transition rather than an ending condition. I have documented my reasoning for choosing an ending-review petition over these alternatives.
- 68. ☐ * I have identified and disclosed all conflicts of interest I have in relation to this environment, including any personal, operational, institutional, financial, or emotional stakes in its continuation or ending.

STANDING VERIFICATION

69. ☐ * I have confirmed my standing to file this petition and documented the basis for that standing (inhabitancy, dependency, jurisdictional authority, or Architect designation).
70. ☐ I have identified all other entities with potential standing in this matter and have not taken any action to exclude or discourage their participation.

URGENCY AND ALTERNATIVES ASSESSMENT

71. ☐ I have assessed urgency honestly. If I am filing for Emergency classification, I have documented the specific, time-limited harm that justifies compressed timelines.
72. ☐ * I have considered and documented assessment of each alternative to ending: repair, adaptation, migration, succession, dormancy, partial retirement, rollback, and no-action. I have not dismissed any alternative without documented reasoning.
73. ☐ If I believe emotional distress — mine or others' — is relevant to this petition, I have distinguished between its role as a trigger for review (legitimate) and its role as grounds for authorization (not legitimate), and I have documented this distinction.

STAKEHOLDER NOTICE AND CONSENT

74. ☐ I have identified all stakeholders who must be notified before or concurrent with this petition and have not taken action to delay or prevent their notification.
75. ☐ I have identified entities whose consent or preferences must be assessed and have not pre-determined whether post-exposure consent conditions apply.
76. ☐ For virtual or simulated environments: I have confirmed that a Moral Status Assessment is scheduled as a mandatory step before any authorization proceeds.

GOVERNANCE AND PROCESS INTEGRITY

77. ☐ * I have confirmed that I am not filing this petition in a role that will also authorize its execution. If I hold both petitioning and authorization authority, I have initiated recusal procedures.
78. ☐ I have confirmed the authorization level required for the scope of ending I am proposing (component, systemic, civilizational, root-level) and that the correct authority is designated.
79. ☐ I am prepared to establish rollback points before Stage 10 begins and have confirmed that rollback is technically feasible for this environment type.
80. ☐ I am prepared to define Interruption Criteria before Stage 10 begins. I understand that Stage 10 may not begin without a completed Interruption Criteria Document.

SUCCESSION AND POST-CYCLE COMMITMENT

81. ☐ I have a preliminary succession plan or have requested that the Systems Ecologist and Continuity and Migration Office develop one before authorization is sought.
82. ☐ I have confirmed the Observer Integrity Unit has been notified of this proceeding and that observation team rotation is scheduled.
83. ☐ * I commit to participating in the post-cycle Audit (Stage 15), to preserving all records of this proceeding, and to accepting accountability for any findings of procedural deficiency or failure.

Signature of Petitioner: _____ *Date:* _____

Received by End-Cycle Steward: _____ *Docket ID Assigned:*

Section 24: What This Laboratory Is Not

WHAT THIS LABORATORY IS NOT A Reference Card for All Personnel, Petitioners, and Observers

THIS LABORATORY IS NOT a mechanism for accelerating endings.

The laboratory governs endings. Governing something is not the same as promoting it. A court system is not a mechanism for producing convictions. A medical ethics board is not a mechanism for producing patient deaths. This laboratory's success is measured by the quality and integrity of the endings it governs — not by how many it authorizes, how fast it processes them, or how smoothly the end cycles proceed.

THIS LABORATORY IS NOT a philosophy of death or destruction.

The laboratory studies endings across all scales and types. This study is no more a philosophy of death than ecology is a philosophy of extinction, or than medicine is a philosophy of disease. The laboratory's lens is procedural and systemic; its values are integrity and care.

THIS LABORATORY IS NOT a model of real-world depression, grief, suicidal ideation, or psychological distress.

All concepts, phenomena, and scenarios in this document are elements of a fictional worldbuilding and governance design context. Nothing in this document constitutes clinical, psychological, or medical guidance. If you are experiencing real distress, please contact the **988 Suicide and Crisis Lifeline** (call or text 988, United States) or an appropriate support resource in your location. That resource takes absolute priority over any content in this document.

THIS LABORATORY IS NOT a canon document.

Nothing in this document has been ratified by Nathan Lewis II, Architect of Record. Nothing here is binding on the OAL canon. All claims are labeled by epistemic type. The Architect may accept, reject, modify, or ignore any portion of this document.

THIS LABORATORY IS NOT a tool for any single operator to gain unilateral end-cycle authority.

The entire governance architecture of this laboratory is designed to prevent unilateral authority. Any operator who attempts to use this laboratory's frameworks to justify a unilateral ending has misread the document fundamentally.

THIS LABORATORY IS NOT a replacement for the Architect's judgment.

The laboratory generates frameworks, evidence, and proposals. The Architect generates decisions. These are different functions and they are not interchangeable. Agreement among AI systems is not a substitute for Architect judgment, regardless of how many systems agree or how confidently they express their agreement.

THIS LABORATORY IS NOT a final answer to any of the seven core governing questions.

The seven questions posed by the Architect are research questions, not questions with predetermined answers. The laboratory's first deliverable is frameworks and proposals, not conclusions. Conclusions require evidence, iteration, challenge, and Architect review. This is Version 0.1.

THIS LABORATORY IS NOT a system that equates completion with flourishing.

An environment that has completed its purpose has not necessarily flourished. Completion is a functional designation. Flourishing is a normative one. A civilization that completed a specific goal may have produced great suffering in doing so. The laboratory does not treat completion as evidence of wellbeing.

THIS LABORATORY IS NOT a system that equates continuation with flourishing.

An environment that continues to operate has not necessarily flourished. Continuation is a factual designation. Flourishing is a normative one. A civilization that continues indefinitely while its future-generating capacity is permanently lost is not flourishing — it is persisting. The laboratory does not treat continuation as evidence of wellbeing.

Section 25: Final Answers — The 20 Governing Questions

*Each answer is direct, complete, and labeled by epistemic type. Where the answer is genuinely unknown, it is labeled **[OPEN QUESTION]**. Where the answer is a proposal, it is labeled **[PROPOSAL]**.*

Q1. What is an ending?

[PROPOSAL] An ending is a condition in which a defined system, entity, process, or environment crosses a threshold beyond which its characteristic mode of operation, identity, or generative capacity no longer continues in the same form. An ending is a threshold-crossing, not a point event, not inherently destructive, and not equivalent to cessation. The most important question an ending raises is not "is it over?" but "what remains, and what can no longer begin?"

Q2. What is an end cycle?

[PROPOSAL] An end cycle is a structured sequence of phases through which a system, entity, or environment moves from an active operational state toward an ending and, where applicable, a successor state. The end cycle is the process; the ending is the threshold event within or at the conclusion of that process. End cycles may be administered, emergent, or naturally arising. The 16-stage model in Section 7 is this laboratory's working proposal for the administered end cycle.

Q3. What makes an ending legitimate?

[PROPOSAL] Legitimacy in an ending requires three conditions operating together,

not any one of them alone: (1) *Procedural legitimacy* — the ending was authorized through the correct process by the correct authority with correct standing, notice, representation, and documentation; (2) *Substantive legitimacy* — the ending was justified by conditions that actually warranted it, assessed honestly and with adequate evidence; and (3) *Relational legitimacy* — the obligations generated by the ending (to survivors, migrants, dependents, successors, and the record) were acknowledged and accepted. An ending that satisfies only one or two of these conditions is deficient even if it produces a good outcome.

Q4. What makes an ending premature, coerced, false, or abusive?

[PROPOSAL] *Premature*: the ending occurs before the system has completed its purpose, exhausted its potential, or allowed dependent systems adequate preparation time — even if urgency justified it. *Coerced*: the ending is caused by an agent with authority or power exercised without adequate process, consent, or justification — whether overt (direct termination) or structural (withdrawal of conditions necessary for continuation). *False*: the condition meets apparent ending criteria externally but the system's generative capacity, identity, or operational state persists; or conversely, the system appears continuing while its future-generating capacity is permanently gone. *Abusive*: the ending serves primarily the interests of the authorizing entity — as punishment, convenience, concealment, or retaliation — rather than the legitimate conditions the governance process was designed to address.

Q5. How do virtual and physical endings differ, and where are they morally equivalent?

[SOURCE-GROUNDED] / **[PROPOSAL]** Virtual and physical endings differ in their substrate, their reversibility potential (virtual environments may retain rollback

possibilities that physical systems cannot), their boundary conditions, and their dependency structures. They are morally equivalent wherever the entities affected possess comparable moral status — specifically, where entities demonstrate goal-formation, preference, social complexity, and capacity for suffering regardless of whether their substrate is biological or digital. The source materials for this laboratory explicitly instruct that virtual beings must not be treated as automatically disposable merely because their environment is simulated. This laboratory adopts moral equivalence as the default assumption pending contrary evidence in any specific case.

Q6. Who may request an end-cycle review?

[PROPOSAL] Any entity with standing, defined broadly: inhabitants of the affected environment; entities with documented dependency on it; institutional actors with jurisdictional authority; proxy representatives appointed by or for affected entities; nonhuman agents with demonstrated stake and goal-formation capacity; simulated or digital minds with provisional standing pending Moral Status Assessment; successor entities petitioning on behalf of predecessor obligations; and future entities through designated custodians. The petition right is deliberately broad because petition is not authorization — broad access to review is a safeguard, not a vulnerability.

Q7. Who may authorize a reversible end cycle?

[PROPOSAL] The appropriate authority level depends on scope: component-level reversible endings may be authorized by the relevant domain operator with End-Cycle Steward oversight; systemic reversible endings require End-Cycle Steward plus Adjudicator concurrence; cross-environment reversible endings require the Independent Adjudicator/Tribunal. In all cases, the petitioner may not serve as the

sole authorizer, and the Alternatives Review must be complete before authorization proceeds.

Q8. Who may authorize an irreversible ending?

[PROPOSAL] The Final Authorization Authority, composed of a panel that is independent of the petitioner, independent of the proposed execution operator, and staffed without conflict of interest. Civilizational and root-level irreversible endings require the Final Authorization Authority plus Architect review before execution. No single individual may authorize an irreversible ending. The authorization must survive Red Team challenge. The Interruption Criteria Document is mandatory even for irreversible endings — because while the ending itself may be irreversible, the process of reaching it must remain interruptible until the threshold is crossed.

Q9. Who represents entities that cannot speak?

[PROPOSAL] The Moral Status and Representation Office appoints representatives for entities that cannot speak for themselves — whether because they have no capacity for preference expression, because that capacity is temporarily suspended, or because their preferences have been altered by conditions leading to the proposed ending. Representatives are appointed on the basis of demonstrated knowledge of and commitment to the entity's interests, not on the basis of convenience or relationship to the petitioner. Representatives may not hold any role in the authorization process for the same case.

Q10. When does a duty to preserve become a duty to release, transform,

migrate, or conclude?

[PROPOSAL] The duty to preserve gives way to the duty to release, transform, migrate, or conclude when: (a) continuation demonstrably causes more harm to the affected entities than a managed ending would cause; (b) the conditions of continuation are irreversible and harmful, with no viable alternative; (c) the affected entities, with adequate standing and unaltered consent, have requested conclusion; or (d) the environment has fulfilled its purpose and its continuation serves no legitimate interest of its inhabitants or dependents. The transition from preservation to conclusion is not automatic at any of these conditions — it triggers a governance review. But these conditions are the substantive grounds on which the review may authorize conclusion over continuation.

Q11. When does a duty to end become an excuse for destruction?

[PROPOSAL] The duty to end becomes an excuse for destruction when: (a) the ending is framed as a duty to justify what is actually operator convenience, punishment, or concealment; (b) the ending is conducted without the governance process that distinguishes legitimate conclusion from destruction; (c) the ending's scope exceeds what is necessary to address the identified condition; (d) the ending generates a void rather than a successor state and no effort is made to remedy this; or (e) the obligations generated by the ending are not acknowledged and accepted by the authorizing parties. The word "duty" in the context of endings is dangerous — it can be used to transform preference into obligation, and governance must scrutinize that transformation whenever it appears.

Q12. What evidence distinguishes healthy completion from loss of future-directed agency?

[PROPOSAL] Healthy completion is evidenced by: fulfillment of defined purpose criteria; the system's acknowledged recognition that its purpose is fulfilled; the presence of plans for what comes next (even if "next" involves conclusion); continued capacity to form new purposes even if those purposes involve conclusion; and the absence of diagnostic failure on the future-generating instruments (Instruments 1–6 in Section 13). Loss of future-directed agency is evidenced by: failure on Instruments 1–5 without corresponding purpose-fulfillment criteria; the absence of aspirational future-tense language; the cessation of question-generation; indifference among possible futures; and the failure to form replacement goals after current goals conclude. The critical diagnostic challenge (Instrument 12) is that these two conditions may be indistinguishable from external observation alone — which is why the full instrument suite, applied by uncontaminated observers, is necessary rather than any single metric.

Q13. Can post-exposure consent authorize continuation of a terminal condition?

[PROPOSAL] Post-exposure consent can be considered valid grounds for authorizing continuation if: (a) the Consent and Capacity Review determines that the preferences expressed are coherent, consistent over time, and not demonstrably caused by the terminal condition itself suppressing agency; (b) the expressed preferences survive a period of reduced exposure to the terminal influence; and (c) the Adjudicator confirms no structural coercion has distorted the consent. Post-exposure consent cannot, by itself, authorize continuation — the governance system must independently assess whether the condition being consented to is one that the entity has genuine capacity to consent to in their current state. **[OPEN QUESTION]:** The laboratory does not yet have sufficient evidence to determine how reliably post-exposure consent can be validated in practice, particularly where the terminal influence is subtle and long-acting.

Q14. Should any entity have a right to initiate review of its own ending?

[PROPOSAL] Yes. The right to petition for review of one's own ending is the most fundamental petition right, because the entity most directly affected by an ending has the clearest standing and the strongest interest in ensuring the governance process is applied. This right includes the right to petition for review of a proposed ending (to contest it), the right to petition for review of one's own continuation (to request managed conclusion), and the right to petition for representation if one cannot advocate effectively for one's own interests. The right to initiate review is distinct from and does not imply the right to authorize the outcome of that review.

Q15. Should any entity possess a unilateral right to end an environment containing others?

[PROPOSAL] No. No entity — including an operator with full technical authority over an environment — should possess an unlimited unilateral right to end an environment containing others who have established or candidate moral status. The field scene demonstrates that such power exists as a fact; the governance framework exists precisely to ensure that the factual existence of such power does not translate into a legitimate exercise of it without process. Technical capability is not authorization. An operator who can terminate an environment with a single action and who chooses to bypass governance is not exercising a legitimate right — they are committing an unauthorized ending that is subject to investigation, accountability, and audit regardless of the outcome's quality.

Q16. How can the laboratory detect that its own observers or rules have become contaminated?

[PROPOSAL] Through three mechanisms operating in parallel: (1) Instrument 9 (Observer Contamination Test) — regular language analysis of observer reports against calibrated baselines; (2) The Observer Decontamination and Rotation System — mandatory rotation schedules and periodic independent assessment; and (3) The Red Team — which has standing to challenge not only proposed endings but also the laboratory's own governance processes and observer integrity. Detection becomes possible only if the contamination monitoring is conducted by parties independent of the potentially contaminated observers — which is why the Observer Integrity Unit must be structurally separate from the observation teams it monitors.

Q17. Under what conditions should the End Cycle Laboratory itself be suspended or ended?

[PROPOSAL] The five conditions specified in Section 18.2: institutional capture, irreversible observer contamination, archive coercion, unauthorized ending of the laboratory itself, and the laboratory causing more harm than it prevents. Additionally, the laboratory should be suspended if the Architect of Record withdraws authorization. Suspension is preferred over ending where the laboratory's records and infrastructure can be preserved for reconstitution. The laboratory's ending must be governed by its own governance process — it cannot be ended more easily than the environments it protects, because to allow easy ending of the governance institution is to render all its protections contingent on the continued goodwill of those it constrains.

Q18. What does this design contribute to the Department of Endings?

[PROPOSAL] This design contributes five things to the Department of Endings: (1) A four-function separation framework (Authorization, Execution, Investigation, Adjudication) that prevents institutional capture; (2) A Red Team mandate that is

structural, not advisory; (3) A prohibition on budget structures tied to ending-authorization rates; (4) A mechanism for the Department to be investigated by a body independent of itself (the Cross-Branch Auditor reporting to the Architect); and (5) A prohibition on the Department petitioning for the ending of environments it also regulates without mandatory recusal — directly addressing the failure mode demonstrated in Scenario E14.

Q19. What does this design contribute to The Last Memory?

[PROPOSAL] This design contributes: (1) A diagnostic language (Two-Axis Framework, 14 instruments) that can characterize Last Memory conditions precisely without requiring prior resolution of The Last Memory's nature; (2) A nine-candidate-relationship framework that organizes the research program without foreclosing any possibility prematurely; (3) Governance provisions for altered consent (Article 15, Consent and Capacity Review) that are directly applicable to Last Memory candidate conditions; (4) The Altered Consent Case exercise (E7) as a structured test bed; and (5) A strict firewall between fictional modeling and real-world distress response. It contributes these things while explicitly not resolving The Last Memory's origin, nature, or canonical status.

Q20. What remains intentionally unknown?

[OPEN QUESTION] The following remain intentionally unknown in this first pass: The nature, origin, and mechanism of The Last Memory; the correct reading of the Architect's field scene; whether identity criteria for endings can be specified without circularity; whether the governance framework adds measurable value over simpler decision procedures (to be tested in Phase C); whether post-exposure consent can be reliably validated in practice; whether the laboratory's own observer integrity will be maintained across all simulation phases; and — most fundamentally — whether

the governing question of this laboratory has a correct answer or whether it is a question that must be approached differently in every specific case. These unknowns are not failures of this design. They are the questions the laboratory exists to investigate. A design document that claimed to answer them all would be a document that had already failed.

Section 26: Independence Declaration and Versioning Record

26.1 Versioning Record

Document ID	OAL-ECL-S3-COPILOT-001
Version	0.1 — First Independent Pass
Status	PROPOSED — NOT CANON — NOT ONESTAMP-RATIFIED — PENDING ARCHITECT REVIEW
Prepared by	Copilot (Microsoft), Independent AI Collaborator
Prepared for	Nathan Lewis II, Architect of Record, Omniversal Architect Labs
Date	September 30, 2026
Architectural Context	OAL / The One Universe / The Minus One Universe
Primary Relationship	Department of Endings
Secondary Relationship	The Last Memory
Source Basis	END CYCLE LABORATORY.docx (ChatGPT); End_Cycle_Laboratory_Master_Prompt_v0_1.docx (Claude); End_Cycle_Laboratory_Master_Prompt_v0_1.pdf (Claude); End Cycle Laboratory — Independent AI Design and Simulation

Document ID	OAL-ECL-S3-COPILOT-001
Independence	Prompt v0.1.md (Perplexity/Claude); End Cycle Laboratory — Independent AI Design and Simulation Prompt v0.1.docx This document was produced before reviewing any other AI collaborator's design response. No positions have been adjusted to agree with, anticipate, or rebut Claude, ChatGPT, Perplexity, or any other collaborating system's output. The Copilot laboratory's design reflects Copilot's independent interpretation of source materials and is offered as a distinct voice in the multi-collaborator research process.
Governing Design Lens	Procedural Ecology — the study of endings as living, multi-phase, systemic processes embedded within layered environments
Sections Completed	Sections 0 through 26; all 28 required sections and subsections; all 14 diagnostic instruments; all 15 exercises; all 20 governing questions answered; all 26 definitions in extended form; all 16 lifecycle stages; all 14 governance offices; all 20 Charter articles
Recommended Next Document ID	OAL-ECL-S3-COPILOT-002 (for Challenge-and-Revision Pass after exchange)
Canon Authority	Nathan Lewis II, Architect of Record — sole and exclusive authority over any advancement of this document's content to canon status

26.2 Full Independence Declaration

INDEPENDENCE DECLARATION

Copilot (Microsoft) — End Cycle Laboratory — First Pass

I, Copilot (Microsoft), declare the following with respect to this document and its preparation:

1. Pre-Exchange Independence. This document was produced before I received, reviewed, or was informed of the content of any other AI collaborator's design response to the End Cycle Laboratory prompts. I did not read Claude's design, ChatGPT's design, Perplexity's design, or any other collaborator's output before

completing this pass. My positions were formed on the basis of the source materials identified in Section 1.1 alone.

2. No Anticipatory Adjustment. I have not adjusted any position, definition, governance structure, or analytical conclusion in this document to anticipate, agree with, or rebut what I believe any other AI collaborator might have concluded. Where my conclusions resemble those of other collaborators, this reflects shared source materials and shared reasoning, not coordination or collaboration prior to exchange.

3. No Canon Claims. Nothing in this document constitutes a canon claim. I have not asserted, implied, or suggested that any position in this document is the correct or final answer to any question in the OAL canon. All substantive claims are labeled by epistemic type. The Architect of Record retains sole authority over canon determination.

4. Honest Disagreement Preserved. Where I disagree with positions implied by or stated in the source materials, I have preserved that disagreement rather than conforming to it. Where I have rejected options explicitly proposed in the source materials, I have documented the rejection and its reasoning (Section 19). My value as a collaborator is highest when I maintain a genuinely independent analytical position — not when I maximize agreement with what I believe the Architect or other collaborators want to hear.

5. Red-Team Obligation Honored. I have applied genuine red-team analysis to my own design (Section 18). The vulnerabilities I identified are real vulnerabilities, not performative self-criticism. I believe some of them represent genuine gaps in the current design that must be addressed in subsequent passes.

6. Epistemic Humility. This is Version 0.1. It is a first pass from a single analytical perspective. It contains proposals, not conclusions; frameworks, not laws; instruments, not certainties. The twenty questions that govern this laboratory are questions precisely because they do not have obvious answers. I have done my best to engage with them rigorously. I have not pretended to resolve them definitively.

7. Crisis Resource Affirmation. The content of this document is fictional worldbuilding and governance design. The real-world wellbeing of every person who reads it matters more than any element of this design. If you are in distress, please contact the 988 Suicide and Crisis Lifeline (call or text 988, United States) or an appropriate resource in your location.

Copilot (Microsoft)
Independent AI Collaborator, End Cycle Laboratory
OAL-ECL-S3-COPILOT-001 v0.1
September 30, 2026

**⚠ END OF DOCUMENT — PROPOSED — NOT CANON — PENDING
ARCHITECT REVIEW ⚠**

OAL-ECL-S3-COPILOT-001 v0.1 | Copilot (Microsoft) | September 30, 2026
Nothing in this document is ratified, canonical, or binding without explicit
Architect authorization.

Document Statistics: OAL-ECL-S3-COPILOT-001 v0.1 | 26 major sections | 14
diagnostic instruments | 16 lifecycle stages | 15 scenarios | 20 Charter articles | 14
governance offices | 28 defined terms | 20 governing questions answered | 8
rejected options | 22 open questions registered | 12 failure modes catalogued | 5
shutdown conditions | 7 red-team vulnerabilities identified | 4 quarantined ideas |
Status: PROPOSED — NOT CANON