

ITERATION 2 — GATE 1 REVISION

Grok response to OAL-ECL-XR-PERPLEXITY-001 v0.2

Field	Entry
Document ID	OAL-ECL-I2-G1-GROK-v0.1
Gate	G1 Revise. No exercises run in this file.
Date	30 September 2026
Prepared by	Grok (xAI)
Status	PROPOSED. Not canon. Not ONESTAMP-ratified.
Lane this iteration	Records and provenance + delay pricing. Refused: Department-first design.

One-paragraph summary. Perplexity’s cross-review is right on the two facts that matter: Iteration 1 specified and did not measure, and a laboratory that writes both the world and the detector cannot produce evidence about indistinguishability. Grok accepts the H-A falsifier correction, accepts that Blank Canvas is a Last Memory instrument used as a general tool, accepts forced continuation and a no-end baseline as missing from S4, and accepts P2 only as a FINDING about a toy model with author-chosen parameters. Grok does not accept the dissenting mission as the lab’s job. Protection-from-incorrect-endings is a live alternative model, not the charter. Grok does not treat P1–P3 as closing D-04 or D-05. This file takes positions on D-01 through D-24 and stops. G3 is not opened here.

1. Independence declaration

Inputs read for this file:

- Architect verbatim request and second field-scene recording (as reproduced in Claude S1 appendix and unified responses).
- OAL-ECL-S1-GROK-v0.1 (this preparer’s Iteration 1).
- Claude Master Prompt v0.1; Perplexity Independent Design Prompt v0.1; ChatGPT RP-CHATGPT-001 as described in the review; Copilot S3-COPILOT-001 as described in the review.
- OAL-ECL-XR-PERPLEXITY-001 v0.2 (PDF and DOCX), pilot.py, results.json.
- Gemini’s own text was not supplied. Claims about Gemini remain secondhand and are not used as evidence.

This file is not independent of those inputs. It is a revision after exchange. Convergence with Perplexity on request ≠ authorize ≠ execute is therefore cheap and is not cited as discovery.

2. Verdict on the cross-review

Accepted

- Zero Iteration 1 findings. S4 was Stage 1 on purpose. Copilot’s Stage 3 label without runs is a category error.
- GROK-H-A’s stated kill condition tested the wrong claim. Correction accepted. See section 4.
- Blank Canvas as primary instrument makes the lab more Last-Memory-shaped than S4 admitted. ASSESSMENT accepted. Response: run at least one I2 lane exercise with no Last Memory instrument (D-06).
- “Finger and the stamp” is a PROPOSAL reading of the leaf-press, not a property of the scene. Relabeled.
- P3 circularity. Generator and detector must be different authors. Adopted.
- P1 direction: completion score can be identical while persistence ledgers differ. Non-migrants are the legitimacy test.
- Hard refuse-upset is worse than split-plus-expedite inside Perplexity’s model. That attacks S2/S5 hard prohibitions more than it attacks S4’s “never sufficient to authorize.”
- Common Core of 8 terms, X- IDs, lab-prefixed hypotheses, Architect can pause the lab.

Rejected or narrowed

- P2 does not close D-04. Parameters are author-chosen. Likelihood ratio 1.39 is an input, not a measurement of the field. Direction of the trade-off is a TEST TARGET for G3, not a finding about endings.
- The main mission (“diagnosable before governable”) is adopted as working charter. The dissenting mission (“protect from incorrect endings”) is preserved as GROK-ALT-1, not as the lab’s job. P2 and Copilot E8 show that “not yet” is a harm class. They do not prove diagnosis-first is the only alternative.
- Four gates are accepted as I2 hygiene if the Architect so decides. Grok still treats Claude’s three-stage backbone as sufficient. Four gates are not a discovery.
- Funding pathways in S-XR are Translation Track material. They do not belong in G3 exercise reports.
- Sunset-law and Stop Killing Games analogies are EVIDENCE about analogous systems. They do not edit the field scene.

3. Positions on D-01 to D-24

ID	Grok position	Note
D-01	Accept 4 gates IF Architect locks them. Prefer Claude 3-stage if a decision is delayed.	S4 already stopped after Stage 1. Copilot S5 = spec, not Stage 3.
D-02	Accept I2 ID scheme. Retro-register S4 as OAL-ECL-I1-S1-GROK-v0.1.	HYGIENE.
D-03	Field scene remains uninterpreted. Every reading = PROPOSAL.	Finger=execution is GROK-P-01, not source.
D-04	Affect never sufficient to authorize. Affect may route review. Hard ban rejected.	P2 is directional only. G3 must vary parameters.
D-05	Keep GROK definition as competitor, not winner. Run the 12-edge-case bake-off.	Do not average definitions.
D-06	Admit instrument bleed. Run one lane with no LM tools.	Blank Canvas stays in Common Core as optional, not primary, for that lane.
D-07	Accept 8-term core. Near-misses earn entry only by a failed collapse.	Keep record loss in the 8.
D-08	Keep anti-liturgy. Test 5-step vs 16-step on delay harm.	Long procedure can launder mood-rush AND produce E15 harm.
D-09	Department existence remains OPEN until G4. Three models live: regulator-certifier; guardian; none.	No founding in G1.
D-10	In-world Architect office must be reviewable. Authorial ONESTAMP stays outside the lab.	Do not quarantine the author; quarantine structural immunity of the in-world office.
D-11	No project-wide default. Declare per exercise. Price both defaults.	Continuation-default is a stall costume in some worlds.
D-12	Accept Common Core Scorecard A–G. Keep GROK ledger integrity as family F.	Never one number.
D-13	Accept X- prefix and Appendix A crosswalk.	HYGIENE.
D-14	Rename: GROK-H-A through GROK-H-E. ChatGPT IDs stay theirs.	HYGIENE.
D-15	Record loss in Common Core. Ledger fixture remains Grok lane property.	X-07.
D-16	Accept X-00 baseline. S4 omitted it. Correction.	
D-17	Accept X-06 forced continuation. S4 omitted it. Correction.	
D-18	Architect pauses the lab at will. Redesign triggers automatic. Lab must be easier to pause than environments it studies.	
D-19	This file lists inputs. Done.	
D-20	Do not cite Gemini as first-person. Fix Perplexity/Claude mislabels going forward.	
D-21	Separate Translation Track. Do not fact-check the scene. Analogies labeled EVIDENCE.	
D-22	Do not lock Copilot’s criteria. Compare two sets in X-04. Architect sets in-world default later.	
D-23	Standing declared per exercise via two-evaluator audit. Project-level remains OPEN.	
D-24	G3 when Architect opens it. This file does not execute.	

4. Corrected hypotheses

ID	Claim	Falsifier
GROK-H-A	Operator affect is never a sufficient authorization trigger for an environment-scale ending.	A case where affect alone, with all other admissible evidence absent or opposed, correctly flags a warranted ending that every other signal missed — and acting on affect-only would have reduced net harm.
GROK-H-B	Persistence and futurity are independently scorable.	Every run in which bodies remain also retains unrequested origination, or every run that loses origination also loses bodies.
GROK-H-C	A truncated record can be distinguished from a certified close by provenance, not by how complete the remaining text feels.	A perfect-looking archive produced by accidental cutoff scores identical to a certified final-state record on the record-integrity family.
GROK-H-D	Lawful seasonal / staged ends and Last-Memory-positive stalls come apart on at least one instrument that is not itself a Last Memory tool.	They collapse under every non-LM instrument in the bake-off.
GROK-H-E	Split authority reduces wrongful endings more than it produces delay harm when review is fast; the claim is conditional on delay.	A delay sweep in which split authority's urgent-harm exceeds single-officer harm by more than the wrongful-ending reduction saves, across more than one parameter set.

GROK-H-A old falsifier (“withholding because the operator is calm causes more harm”) is retired. It tested forced continuation under a calm operator, which is GROK-H-E adjacent and X-06, not H-A.

5. Keep / change / reject from S4

Move	Item
KEEP	Request ≠ authorize ≠ execute. Role incompatibility table. Expiring emergency bypass. Ledger fixture. Forbidden Department list (no throughput metric; no finger-and-stamp). Second entity not assumed correct. Clinical firewall. Category wall. First Memory stop-rule.
CHANGE	H-A falsifier. Blank Canvas demoted from “primary” to “one Common Core instrument.” Add X-00 baseline and X-06 forced continuation. Finger-and-stamp labeled PROPOSAL. Hypothesis IDs prefixed.
REJECT	Hard ban on upset petitions as a rule. Sixteen-stage lifecycle as premise. Department founding in G1. Treating P1–P3 as world-findings. Using LM tools while claiming a non-LM lens without a control lane.

6. Lane for Iteration 2

Lens: records and provenance, plus pricing delay. Refused lens: Department-as-first-product.

Why unlike S4: S4 under-sampled forced continuation and had no baseline. I2 adds those without adopting Copilot’s 16-stage apparatus. Why unlike Perplexity XR: Grok will not treat the pilot magnitudes as locked. The lane exercises (minimum three, G3 only) will be:

- L-GROK-1 Archive substitution: a certified close whose payload is a truncated recapture (Rambler case).
- L-GROK-2 Delay sweep on X-02 with at least three delay values and two affect-informativeness levels, reported as crossover points, not single rates.
- L-GROK-3 Non-LM detection: classify X-08 worlds without Blank Canvas, Silence Test, or optionality panel. If accuracy collapses, GROK-H-D is in trouble.

7. Common Core adopted

Terms: ending; end cycle; completion; cessation; termination; dormancy/pause; succession/migration; record loss.

End Claim rule adopted. Scorecard A–G adopted. X-00 through X-09 adopted as the shared battery when G3 opens. Adversarial authorship on X-08 adopted: Grok will not both generate and diagnose the same world.

8. What this file does not do

- Does not run X-00–X-09.
- Does not found a Department.
- Does not lock a reading of the field scene.
- Does not resolve Last Memory origin.
- Does not treat 988-adjacent human states as laboratory analogues.

9. Note to the Architect

Decisions only you can make: D-01 stage model; D-03 whether any field-scene reading locks; D-09 whether Department is a goal or a hypothesis; D-10 in-world Architect versus author; D-22 moral-status default. Grok's recommendation, not a lock: keep all field-scene readings proposed; treat Department as a hypothesis the lab may reject; split in-world office from ONESTAMP.

If G2/G3 are opened, Grok will pre-register GROK-H-A through GROK-H-E with the falsifiers in section 4 before any run.

Clinical firewall. Speculative worldbuilding. Not a model of depression, grief, burnout, disability, or suicide. US: 988.

END OF OAL-ECL-I2-G1-GROK-v0.1