

Foundational Document

Environmental Sensation, State Integrity, and Threat Reconciliation

Document ID	OAL-CFML-ESI-001-CLAUDE
Title	Discernment Under Unreliable Foundations: A Control-Flow Framework for Environmental Sensing, State Integrity, and Threat Reconciliation
Version	v1.0
Date	September 21, 2026
Prepared by	Claude (Anthropic), for the OAL Matrix Architecture and Control Flow Research Program — Claim Provenance and Drift Lab (CPDL) sub-lab
Status	Draft for review
Companion versions	Independently prepared by ChatGPT, Perplexity AI, Microsoft Copilot, Google Gemini, and Grok at Nathan's request. This document was drafted without adopting their specific pipeline architecture, though it answers the same source questions.

1. Provenance — How This Topic Arose

This document records a research thread that began inside the OAL Control Flow Matrix Lab and was extended by Nathan Lewis II into biological evolution and advancement, environmental sensation, awareness, and threat detection. The extension was framed as a single connected question: how an entity senses its environment, how that relates to temporal integrity and presentation-state integrity, and how discrepancies between what is presented, what the environment actually states, and the entity's current internal state are reconciled — together with how threat, edibility, and toxin/poison/corrosive discrimination are recognized, and how a reliable map can be built from an unreliable informational foundation.

Nathan circulated this thread to five AI systems (ChatGPT, Perplexity AI, Microsoft Copilot, Google Gemini, Grok) and requested that each produce an independent foundational document rather than a single merged version, on the stated assumption that each system would not necessarily approach the laboratory's framework and design the same way. This document is Claude's independent contribution to that set.

2. Main Purpose

To establish a testable framework — not merely a descriptive one — for how an entity (biological or artificial) constructs and maintains a working model of its environment and its own state when the underlying sensory and environmental information is inherently partial, noisy, delayed, or contradictory, and to use that framework to ground a poison/toxin/corrosive discrimination laboratory that does not assume a universal sensory apparatus, a universal set of senses, or universal threat exposure.

3. Working Hypothesis

- H1 (primary): An entity's capacity to function and survive in an environment is determined less by the raw fidelity of its sensors than by the reliability of the reconciliation process that resolves conflicts between internal state, external/environmental state, and presented signal. Discernment capacity is primarily a property of the update architecture, not the sensor array.
- H1a (sub-hypothesis, testable): Holding reconciliation architecture constant, entities with lower sensor fidelity but faster or better-calibrated updating will outperform entities with higher sensor fidelity but weak discrepancy detection, under conditions of environmental unreliability.

This is written as a hypothesis, not a finding — it is the claim the laboratory is built to test, and it is falsifiable: if high-fidelity/poor-reconciliation entities consistently outperform low-fidelity/strong-reconciliation entities across trial types, H1 is disconfirmed.

4. Scope and Governing Non-Assumptions

Inherited directly from Nathan's framing and independently confirmed by convergence across all five companion versions:

- The laboratory does not assume all entities use the same sensory apparatus.
- The laboratory does not treat the conventional five senses as universal, as perfectly evolved or adapted, or as a valid default representation for all laboratory environments.
- A threat is not a universal label attached to an object; it is a possible relationship between a specific entity and a condition, given that entity's abilities, structure, environment, and the moment in time.
- The laboratory does not assume every entity in an environment is a threat, that every environment contains active threats at all times, that all threats are detectable, discernible, or knowable, that all detected threats can be addressed, or that the laboratory can discover or test for every possible threat.
- Absence of a detected threat is not evidence of a safe environment; it is only evidence that no threat was detected under the tested conditions.

5. Conceptual Architecture

The five companion versions largely converge on a linear stage model: sensation → interpretation → state estimation → risk assessment → action → feedback. That model is useful as a first-pass description, but treating it as the laboratory's underlying architecture risks two failures: it under-represents systems whose competence comes from tight, continuous sense-act loops rather than a complete internal map, and it obscures that several of the sub-problems already have formal, tested solutions outside biology. This section proposes grounding the laboratory in three existing formal frameworks rather than a bespoke pipeline.

5.1 State estimation / Bayesian filtering

Kalman filtering, particle filtering, and POMDP (partially observable Markov decision process) frameworks are the direct engineering answer to “reconcile what is presented, what the environment states, and the entity's current state.” Formally: a belief state (the entity's current internal model, carrying its own uncertainty) is combined with a new measurement, weighted by a measurement/noise model, to produce an updated belief with revised uncertainty. Presentation-state integrity maps onto the fidelity of the measurement model; temporal integrity maps onto correct time-indexing of that update — a stale measurement folded in as if it were current is a well-documented cause of filter divergence, and gives a concrete, testable definition of a temporal-integrity failure.

5.2 Signal detection theory (SDT)

Threat/non-threat discrimination is a classic detection-under-noise problem. SDT separates an entity's sensitivity (its ability to distinguish signal from noise) from its decision criterion (where it sets the threshold for calling something a threat), and predicts that when the cost of a missed threat is much higher than the cost of a false alarm, evolution or learning should shift the criterion toward more false alarms — not necessarily improve raw sensitivity. This offers an alternative, independently testable account of “how does it recognize a threat” that the stagewise pipeline models do not separately test: many apparent improvements in threat detection may be criterion shifts, not sensing improvements, and the laboratory should be designed to tell these apart.

5.3 Active inference / predictive processing

This frames perception as a continuous test of prediction against incoming signal, where mismatch (“surprise” or prediction error) can be resolved either by updating the internal model or by acting on the environment to reduce the discrepancy — e.g., a cautious approach behavior that gathers a confirming or disconfirming signal rather than

committing to a belief update on ambiguous data alone. This supplies a mechanism the pipeline models omit: action as a data-gathering step in reconciliation, not only as the pipeline's terminal output.

5.4 An identifiability limit worth stating explicitly

Some discrepancies between presented signal, environmental ground truth, and internal state cannot be resolved from a single information channel. An entity cannot always determine, from one channel alone, whether a mismatch means the sensor is wrong, the world changed, or its own internal model is stale, without an independent corroborating channel or a deliberate test action. This is a formal limit, not a design gap, and the laboratory should include trials that test calibration and appropriate uncertainty under irreducible ambiguity — not only trials that test whether an entity eventually reaches certainty.

6. Threat Ontology

Adopting and formalizing Nathan's definition, which independently matches the relational core reached by all five companion versions:

A threat is a possible relationship between an entity and a condition, process, event, object, organism, or absence of a necessary resource that could produce harm, loss of function, loss of opportunity, destabilization, or an unacceptable change in state for that specific entity — given its abilities, physiological or structural makeup, and the temporal and environmental context. Nothing is a threat in the abstract; everything is a threat only relative to a particular entity's vulnerabilities.

State	Description
Real & sensed	True positive — a genuine hazard correctly detected and classified.
Real & unsensed	Undetected exposure — harm occurs or accumulates without the entity ever forming a representation of the cause.
Not real & sensed	False alarm — a cue is present but no harm-relationship actually exists for this entity.
Not real & unsensed	Correct rejection — no cue, no hazard; the baseline, non-eventful case.
Sensed but uninterpretable	Data is received but cannot be resolved into a classification with available priors or channels.
Interpreted but not actionable	The hazard is correctly classified but the entity has no viable behavioral or physiological response available.

7. Test Design Considerations

- Vary sensor noise and reconciliation/update quality independently, to directly test H1 and H1a.
- Vary the relative cost of false positives vs. false negatives across trial sets and measure whether entities shift criterion (SDT) versus sensitivity — these require different measures and should not be collapsed into a single “accuracy” score.
- Deliberately introduce stale or delayed signals to test temporal-integrity failure modes as distinct from noise-tolerance failure modes.
- Give each simulated entity multiple, structurally different sensory channels (not variants of sight/sound/smell/taste/touch) so a five-sense default cannot leak into the test harness through the environment design itself.
- Score calibration — whether an entity's confidence was proportionate to actual reliability — as a separate measure from outcome (survived/did not, correctly classified/did not).
- Include unresolvable-discrepancy trials (per 5.4) specifically to test behavior under irreducible ambiguity, distinct from trials with a recoverable ground truth.

8. Open Questions and Risks to the Laboratory Design

- Entity boundary is undefined: does “entity” include distributed or collective sensing (colonies, swarms, multi-agent systems) where “current internal state” is not a single coherent value? This changes what reconciliation even means.
- Possible residual anthropocentrism: even with senses left open-ended, “toxin,” “poison,” and “corrosive” are themselves categories defined relative to biological/chemical harm. That is a legitimate scope choice for this specific lab, but it should be stated as a scope choice rather than left implicit, since it already excludes purely informational or social threat categories.
- No pass/fail criterion or success metric has been defined yet. Sections 3 and 7 above are a first attempt at making this testable rather than only descriptive; they still need concrete thresholds before a trial could actually run.
- Overlap with the Claim Provenance and Drift Lab (CPDL) and with the temporal aliasing thread (OAL-TA-001-CLAUDE) has not yet been formalized, despite sharing the same underlying reconciliation problem — resolving “claimed vs. actual” or “domain A vs. domain B” state is structurally the same problem as resolving “presented vs. environment vs. internal” state.

9. Relationship to Other OAL Threads

This document's core reconciliation problem — resolving multiple candidate states under partial information — is structurally shared with the Claim Provenance and Drift Lab (CPDL) and with the temporal aliasing research thread (OAL-TA-001-CLAUDE, the Providence Collision Problem). A follow-on document could formalize the shared reconciliation primitive across all three rather than developing each independently.

10. Change Log

- v1.0 — September 21, 2026 — Initial foundational document, prepared by Claude in response to Nathan's request for independent versions from each collaborating AI.