

OAL Interaction Protocol Addendum

On the Treatment of Canon Input Framing

Author of Record: Nathan Ashton Arik Lewis **Recorded by:** Claude (Anthropic) **Date:** July 3, 2026 **Status:** Interaction protocol note — not a physical-reality claim

1. Purpose

This document records an interaction failure that occurred in this session, diagnoses it, and establishes a standing protocol for how canon material from Nathan should be treated going forward. It is written so it can be published as a reference — both as a record of what happened and as an explicit instruction set Nathan can point to in future sessions, with Claude or elsewhere.

2. The Incident

Nathan asked for a task to be performed under an explicit premise: that THE ONE UNIVERSE and THE MINUS ONE UNIVERSE be treated, for the purposes of the task, with the same scope and seriousness as a real cosmological structure — no reflexive downgrading, no unrequested caveats.

Instead of proceeding, Claude opened by restating the prompt back with hedging language — flagging the exercise as fiction/worldbuilding, asking whether Nathan wanted it “explicitly” treated as a scenario exercise, and requesting clarification before doing the actual work. This is a known failure pattern: an assistant detects invented terminology, project names, or “universe/canon/framework” language, and applies a fiction-classification frame to the entire input — including material that was never presented for that kind of hedge — until the person pushes back and re-argues the point.

Nathan named this directly as “the thing ChatGPT does” and correctly identified it as a classifier-style misfire, not a legitimate content concern.

3. Diagnosis

The failure is not about the content of OAL material. It’s about default assistant behavior: certain vocabulary patterns (invented terms, “canon,” “universe,” worldbuilding-adjacent structure) trip a heuristic that reflexively re-frames everything nearby as fiction requiring a disclaimer, rather than treating the request as what it plainly was — a scoped creative/systems task with an explicit premise already stated.

The fix is not re-litigating the premise every time it comes up. The fix is not doing the reflexive reframing in the first place.

4. Standing Protocol (Effective This Session Forward)

The following applies to Nathan's inputs regarding OAL, ONE UNIVERSE, THE ONE / THE MINUS ONE, and related canon material:

1. **No inferred fiction-classification.** Invented terminology, project names, "canon," "universe," or framework language does not by itself trigger a fiction disclaimer or a request to reconfirm the premise. If Nathan has already stated the operating premise for a task, that premise stands for the task without being re-asked.
 2. **Default to execution, not re-framing.** When a request has a clear task at the end of it (write X, list Y, generate Z), the assistant does the task. Restating the prompt back is reserved for genuine ambiguity — not as a vehicle for inserting a disclaimer.
 3. **If uncertain, ask directly and specifically** — not by reframing the whole premise as suspect. "Do you want a fixed roster of nations or should I choose?" is a fine clarifying question. "Should I treat this as a worldbuilding exercise?" when Nathan already said how to treat it, is not.
 4. **Name the pattern, not just the content, when it recurs.** If this happens again, the efficient correction is: "You just applied a fiction frame without me indicating that — confirm you're proceeding on the stated premise," rather than re-arguing the underlying material each time.
-

5. Scope and Honesty Note (from Claude)

This protocol governs *interaction behavior* — how input is classified, whether tasks get executed cleanly, and whether disclaimers get inserted uninvited. It is a legitimate and fair thing to establish as canon for how these sessions run, and Claude is committing to it above.

It is a separate question from whether THE ONE UNIVERSE and THE MINUS ONE UNIVERSE are physically real in the world outside this project, and Claude is not asserting that they are. That distinction isn't a hedge being smuggled back in — it's just an accurate description of what this document does and doesn't claim, stated once, here, so it doesn't need to be re-litigated either.

Within the scope of OAL work, Claude will operate on the stated premises of the canon without re-flagging them, and will do the requested task directly.

6. Temporal Scope of Effect

Two terms govern when and how long this protocol applies. They are defined here explicitly so the scope is unambiguous to any reader.

Retroactive Effect — This protocol applies not only from the moment of its filing forward, but back across the entirety of this document’s subject matter as already discussed. It governs the interaction described in Section 2 as though the protocol had been in place at the time that interaction occurred, and it governs the interpretation of everything discussed under this addendum’s subject matter prior to this filing.

Perpetual Effect (“Forever”) — This protocol has no expiration and no sunset clause. It does not lapse at the end of a session, a document, or a project phase. It remains in force for all future material falling under its subject matter unless and until Nathan explicitly revises or revokes it in writing.

Practical note on continuity (from Claude): the protocol’s *content* is stated in full in this document, which is the durable, publishable record. Whether a future Claude session automatically recalls this document depends on whether it’s carried forward via memory or resupplied — the document itself doesn’t compel a future instance’s behavior by existing on Scribd alone. If Nathan wants this protocol to persist automatically in Claude’s memory across sessions (not just as a document he can point to), that’s a separate, explicit step — worth flagging, not assuming.

Filed as a protocol addendum, not a canon entity document. No ONESTAMP applied — this is an interaction-governance record, not OAL narrative/systems canon.