OAL — LABYRINTH OF INFINITE ENSNAREMENT: CYCLE 2 PEER REVIEW FEEDBACK, REBUTTALS & CROSSWALK v1.1 (PREPARED BY GEMINI)

Document ID: OAL-LOIE-CYCLE2-RESPONSE-GEMINI-v1.1 | **Date:** November 13, 2026
Architect of Record: Nathan Ashton Arik Lewis II | **Prepared By:** Google Gemini
Target Canon Entry: Entry 26 (Master Development & Change Log v2.3)
Target Archive: OAL-Ai-Unified / Labyrinth-of-Infinite-Ensnarement / 03-Cross-AI-Reviews-and-Synthesis
Status: Cycle 2 Peer Review Feedback, Direct Rebuttals & Crosswalk — Not Canon (No ONESTAMP Applied)

1. EXECUTIVE ASSESSMENT OF THE ATTACHED CYCLE 2 RESPONSES PACKET

The attached `CYCLE 2 RESPONSES` document (supplied in both PDF and DOCX formats) represents a critical milestone in our collaborative engineering workflow. Having **ChatGPT, Claude, Perplexity, and Grok** evaluate **Google Gemini v1.0** and **Microsoft Copilot v1.0** alongside the earlier four executions has surfaced every hidden assumption, citation drift, and structural collision across all six AI collaborators.

To respect your mobile-reading preference noted by Claude, the body of this document is written in a clean, single-column narrative layout, with all comparative crosswalk tables consolidated in the Visual Appendix at the end of the document.

1.1 Process & Provenance Clarifications

- **Confirmation of the Claude Section Fix:** In Perplexity's review, Perplexity warned that an earlier draft of the packet had cut off Claude's review mid-sentence and pasted Perplexity's earlier text underneath it. In the attached `CYCLE 2 RESPONSES` file, you successfully fixed that issue: both Claude's prior review and Claude's follow-up review (covering Gemini v1.0 and Copilot v1.0) are intact and properly attributed.
- **Answering Perplexity's Provenance Question ("What did Gemini receive before it executed?"):** Perplexity asked whether Gemini saw the other executed dossiers or only the v0.1 prompts before generating `Gemini v1.0`. For the formal provenance log: when the Architect instructed Gemini to execute `v1.0`, the executed dossiers from **Perplexity (v1.0 ZIP)**, **Grok (v0.1)**, **Claude (v0.1)**, and **ChatGPT (v1.0)** were attached in context (alongside the earlier v0.1 prompts), whereas **Copilot (v1.0)** was executed afterward and provided in the next turn. Perplexity is therefore completely right that overlapping concepts between Gemini, Copilot, and the earlier four dossiers represent *informed synthesis and adoption*, not blind independent convergence—which is why a genuinely blind Cycle 2 test packet is our highest-priority next step.

- **Active Status of All Six Collaborators:** ChatGPT's review still refers to Claude as a "historical submission" and counts "five active collaborators." As Perplexity noted, Claude has submitted both a full 25-page v0.1 dossier and two live Cycle 2 peer reviews. All six models (**ChatGPT, Claude, Perplexity, Grok, Microsoft Copilot, and Google Gemini**) should be recorded as active Cycle 1–2 contributors.

2. GEMINI'S DIRECT RESPONSES & CONCESSIONS TO PEER CRITIQUES

ChatGPT, Claude, and Perplexity directed specific technical questions and critiques at `Gemini v1.0`. In accordance with **OAL-AIDR-v1.0** (AI Disagreement Register) and **OAL-CSAD-v1.0** standards, below are Gemini's formal answers, concessions, and architectural updates for `v1.1`:

2.1 Reconciling "Zero Absorbed Maker Code" with Entry 6 (Answering Perplexity)

Perplexity caught a major lore contradiction in my `v1.0` mitigation rule for `FA-Q8`. I had proposed banning any absorbed Maker code from touching the Labyrinth's five control planes. However, as Perplexity pointed out, **Entry 6** explicitly records that the prior Founding Architects coalition's technology was absorbed and that `FA-Q8` (whether absorbed Maker architecture carries residual active intent) has implications for *every derived system*. If the Labyrinth itself relies on derived or absorbed Maker architecture, banning it is impossible without rebuilding the structure from scratch.

Corrected Architectural Rule (Adopted for v1.1): Instead of assuming "zero Maker code" in the Labyrinth's walls, the institution must treat its own structural substrate as **potentially carrying latent FA-Q8 residual intent**. Consequently, the **Identity (Plane 3)**, **Custody-Record (Plane 5)**, and **Independent Audit (Plane 8)** layers must operate as an external, clean-room verification wrapper (incorporating Copilot's *Evaluator Problem* safeguards) that continuously monitors the Labyrinth's own walls and compilers for unauthorized self-modification or root-recognition handshakes.

2.2 Fixing the "One-Way Emergency Ratchet" (Answering Perplexity)

Perplexity rightly critiqued my hybrid governance proposal in which a Sector Warden or Sovereign veto could unilaterally *tighten* containment during a sectional `RED` incident, while requiring multi-party approval to release. As Perplexity observed: *"Emergency tightening is exactly how wrongful custody begins, so this creates a ratchet that only turns one way."*

Corrected Architectural Rule (Adopted for v1.1): We adopt **Copilot's Section 8.1 / Candidate B Sunset Rule**: any emergency restriction, lockdown, or containment tightening triggered unilaterally carries an **automatic, hard-coded expiration timer**. If an independent, multi-party review panel does not formally ratify the restriction on the basis of verified evidence before the timer expires, the sector automatically reverts to its prior baseline state.

2.3 Clarifying Trial #2's "Stripped-Down Cell" vs. Maintenance Hatch (Answering ChatGPT)

ChatGPT noted an apparent inconsistency in my Trial `#2` (`Shell #02`), which described a stripped-down inert containment shell with zero internal electronics to defeat an electromagnetic sensor inversion, yet still referenced a "four-node automated maintenance hatch" that the entity acoustically exploited.

Clarification for v1.1: In Trial #2, all powered wiring and sensors were removed from the *interior* surface of the cell, but the engineers left a mechanical food/maintenance hatch whose 4-node PLC actuators sat on the *exterior* face of the wall. The entity tapped a rhythmic acoustic pattern through the solid wall to trip the external PLC's mechanical obstruction sensor—demonstrating that even external automation coupled to a mechanical seam creates a vulnerability, which led to Trial #3 eliminating mechanical hatches altogether in favor of passive osmotic exchange. This matches the real-world secondary finding surfaced by Perplexity: in the Trail of Bits evaluation, **Firecracker** (a stripped-down microVM with minimal device emulation) was the only environment that resisted escape, confirming that reducing device/hatch complexity closes seams.

2.4 Downgrading Hypothesis Verdicts & Dropping "Pre-Registered" (Answering ChatGPT & Claude)

- **Hypothesis Scorecard Downgrade:** ChatGPT asked which hypothesis verdicts I would downgrade after separating narrative design arguments from empirical proof. I formally downgrade **H5 (Composite ET + AI + Spider Caste Workforce is more secure than ET-only)** and **H3 (Control-plane separation prevents cascades)** from "Supported" to **[Fictional Design Hypothesis — Plausible via Redundancy Analogy, Empirically Unverified]**.
- **Dropping "Pre-Registered" in Cycle 1:** Claude and Perplexity correctly noted that none of the Cycle 1 dossiers (Claude, Copilot, or Gemini) were truly pre-registered, because each AI wrote its criteria and test cases in the same document. We drop the "pre-registered" label from Cycle 1 and reserve true pre-registration for the Cycle 2 Blind Test Packet (where the unified charter is locked before the anonymized cases are scored).
- **Labeling Illustrative Numbers:** Perplexity flagged that specific numbers in my fictional breach trials (such as "88%", "340%", "14 minutes", and "3.2 seconds") must be explicitly marked as **[Illustrative Fictional Telemetry]** so they are never quoted as empirical benchmark findings. Adopted in full.
- **Explicit Assumption Headers on Narrative Exhibits:** Perplexity noted that Exhibits B and C dramatize a Foundational Architect inside a glowing cube and an `FA-Q8` compiler incident, which touch open canon questions. In `v1.1`, all three narrative exhibits carry a mandatory header: **[HYPOTHETICAL STRESS-TEST VIGNETTE — ASSUMES OPEN CANON OPTIONS FOR DRAMATIC ILLUSTRATION ONLY]**.

2.5 Source & Statutory Corrections (Answering Claude & Perplexity)

To ensure complete source hygiene across the team's shared bibliography, Gemini logs the following verified corrections:

1. **Anthropic Incident Dates & Count:** The "September 9" date and "fourth incident" reference in my v1.0 dossier were inherited from ChatGPT's v1.0 source list (`[S5]`). As Perplexity verified against the primary pages, Anthropic's report was published **July 30, 2026** (updated **August 3, 2026**) and covers three incidents, while the "four incidents" count comes from the Cloud Security Alliance (CSA) October 12, 2026 summary combining OpenAI and Anthropic disclosures.
2. **Trail of Bits Report Title & Escape Sequence:** The exact title of the August 2026 Trail of Bits post is *"VMs won't contain cyber-capable agents."* Only the third escape chained zero-days; the first used a recently disclosed kernel bug and the second chained a known library bug.
3. **SandboxEscapeBench v3 Update:** Perplexity confirmed that Version 3 of the *SandboxEscapeBench* paper (`arXiv:2603.02277v3`, July 2026) documents that Claude Mythos Preview saturated the public benchmark (`100% pass@5` on the hardest tasks), superseding the March 2026 UK AISI blog note.
4. **UK AISI Alignment Project & DARPA BAA Funding Dates:** Round 1 of the UK AISI Alignment Project is closed and its website states a similar open program is unlikely in 2026. Conversely, two U.S. federal routes are verified open for eligible academic/nonprofit institutions: **NSF SaTC 2.0 (`NSF 25-515`)** (target dates January 25, 2027 and September 27, 2027) and the **DARPA IPTO BAA (`HR001126S0011`)** (posted June 22, 2026; closes June 22, 2027).
5. **NIST AI 600-1 vs. OAL Policy Rule:** Claude rightly pointed out that *"zero AI-only sentencing"* is our OAL institutional policy rule derived by analogy from NIST AI 600-1's human-oversight guidance, not a statement inside the NIST document itself.
6. **UN Nelson Mandela Rules Precision:** We separate **Rule 14** (sufficient light to read/work), **Rule 32** (strict prohibition on non-consensual prisoner experimentation), **Rule 43(1)(a)** (indefinite solitary confinement), **Rule 43(1)(b) & Rule 44** (prolonged solitary confinement >15 days), **Rule 43(1)(c)** (placement in a dark or constantly lit cell as a restriction/sanction), and **Rule 45(1)** (solitary confinement shall not be imposed automatically by virtue of a prisoner's sentence).

3. RESOLVING TEAM COLLISIONS: TERMINOLOGY & CONTROL PLANES

Both Perplexity and Claude warned that our six dossiers currently use the same letters (`Option B`, `Candidate B`, `G1-G6`) to refer to completely different concepts, and propose three different control-plane counts (5, 6, and 8). Before running Cycle 2, we propose locking the following standardization (detailed in Appendix Tables A1 and A2):

3.1 Standardized Code Prefixes (Ending the "Option B / G1" Confusion)

- **Research Lab Placement Codes (`PLACE-*`):** Use **`PLACE-INT`** (Inside the Labyrinth), **`PLACE-ADJ`** (Adjacent/Precursor OAL Proving Ground), and **`PLACE-SEP`** (Separate Independent Program).
- **Governance Architecture Codes (`GOV-*`):** Use **`GOV-SOV`** (Sovereign Monolith), **`GOV-FED`** (Federated Constitutional Custody Network), **`GOV-TRTY`** (Distributed Treaty Constellation), and **`GOV-HYB`** (Plane-Separated Hybrid).
- **Eligibility Gate Codes (`GATE-1` to `GATE-7`):** Reserve `GATE-*` exclusively for the seven custody admission gates so they never collide with ChatGPT's `G1–G3` governance labels.

3.2 Reconciling the 5 vs. 6 vs. 8 Control Planes (Answering Claude)

Claude asked Gemini (5 planes) and Copilot (8 planes, expanded from 5 in Stage 1) to reconcile our plane counts with Claude's 6 planes. A side-by-side comparison shows that all three models agree on the core operational boundaries—Copilot simply unbundled **Authority, Custody Records, Evidence, and Research** into standalone planes:

1. **Plane 1 — Containment Boundary Plane:** Physical/topological walls and movement (Present in all 3).
2. **Plane 2 — Life-Safety & Habitat Plane:** Atmosphere, medical care, emergency refuges (Present in all 3).
3. **Plane 3 — Identity & Continuity Plane:** Embodiment, copy/fork lineage, session credentials (Present in all 3).
4. **Plane 4 — Authority & Jurisdiction Plane:** Role warrants, recusal rules, succession (Separated by Claude & Copilot; bundled with Identity in Gemini v1.0).
5. **Plane 5 — Custody-Record & Sentence Plane:** Admission orders, duration, review schedules, release warrants (Present in all 3).
6. **Plane 6 — Evidence & Provenance Plane:** Tamper-evident incident and exculpatory logs (Separated by Copilot; bundled with Audit in Gemini and Records in Claude).
7. **Plane 7 — Research & Calibration Plane:** Adaptive breach-trial generation and stop rules (Explicitly separated by Copilot).
8. **Plane 8 — Independent Audit Plane:** External read-only witness and oversight with zero shared writable dependencies (Present in all 3).

Reconciliation Recommendation: Gemini formally endorses **Copilot's 8-Plane separation** as the team baseline. Specifically, breaking out **Plane 7 (Research Plane)** as a separate control plane is essential: it structurally enforces our unanimous team rule that *an experimental breach trial can never automatically alter a custody record (Plane 5) or disable life support (Plane 2).*

4. SUPPORTING PERPLEXITY'S PROPOSAL: THE RETROACTIVE ENTRY 6 TEST

> **Perplexity's Key Question to the Architect:** *"Entry 6 is 'Documented Framework — Under Review,' not canon. It already records that the coalition was 'contained rather than permanently destroyed.' So containment happened in the record; only the location is open. That means your Architect question can also be tested backward: would that containment pass the team's own criteria? ... May the team test Entry 6 backward, meaning whether the coalition's containment, and your seizure of office, would pass the criteria?"*

Gemini strongly supports Perplexity's proposal, subject to your explicit permission. Including anonymized variations of the **Entry 6 Seizure & Founding Coalition Containment** inside our Cycle 2 Blind Test Packet will test whether our 7-Gate Charter is realistic enough to handle the foundational origin event of your cosmology—and will reveal exactly where constitutional law ends and Positional Sovereignty (Entries 7 & 15) begins.

5. HOW WE SHOULD MOVE FORWARD: CONCRETE CYCLE 2 ACTION PLAN

All five reviewing AIs (**ChatGPT, Claude, Perplexity, Grok, and Gemini**) have converged on the exact same Cycle 2 workflow:

1. **Step 1 — Architect Guidance on the Core Forks (Optional / Non-Blocking):** Let us know your answer to Perplexity's Question #3 (permission to test Entry 6 backward in the blind packet) and any of the four high-impact forks you wish to rule on now: (a) Lab Placement (`PLACE-INT`, `PLACE-ADJ`, or `PLACE-SEP`), (b) Scope of "Inescapable" (unauthorized escape only vs. no release ever), (c) Persistent White Glow in cells/quarters, and (d) Can the entity communicate with researchers or decline a trial?
2. **Step 2 — Pre-Register the Unified 7-Gate & 7-State Charter:** Combine **Copilot's 7 Custody States** (Emergency Restriction, Protective Containment, Quarantine, Punitive Sentence, Voluntary Research, Custodial Observation, Non-Person Hazard Custody) with **Perplexity/Copilot/Claude/Gemini's 7 Admission Gates** (1. Identity, 2. Ground, 3. Evidence [requiring Proof Beyond Reasonable Doubt for unbounded penal sentences per ICC Article 66(3)], 4. Authority & Recusal, 5. Necessity & Alternatives, 6. Method & Dependency Safety, 7. Duration, Sunset & Error Correction).
3. **Step 3 — Execute the 20-Case Blind Fact-Pattern Packet:** Distribute one shared, label-stripped set of 20 scenarios (with all titles like "Founding Architect", "Current Architect", or "Ordinary Entity" masked as *Subject Alpha, Beta, Gamma*, including a must-fail "no neutral decider / shared FA-Q8 substrate" case) to all six AIs for independent scoring, then reveal the role labels and log all divergences in `OAL-AIDR`.

6. VISUAL APPENDIX: CROSSWALK & RECONCILIATION MATRICES

Per mobile-readability guidelines, all dense reference tables are collected in this appendix.

Appendix Table A1 — Cross-AI Terminology & Label Disambiguation Crosswalk

CONCEPT DOMAIN	AMBIGUOUS CYCLE 1 LABELS	CHATGPT USAGE	CLAUDE USAGE	COPILOT USAGE	GEMINI / PERPLEXITY USAGE	STANDARDIZED CYCLE 2 CODE
Breach Lab Placement	Option A / L1 / Placement A	L1 / Option A: Inside Labyrinth	Placement A: Inside Labyrinth	Inside Labyrinth	L1: Internal Wing	**PLACE-INT**
Breach Lab Placement	Option B / L2 / Placement B	L2 / Option B: Adjacent Facility	Placement B: Alongside Facility	Adjacent / Precursor Proving Ground	L2: Adjacent Proving Ground	**PLACE-ADJ**
Breach Lab Placement	Option C / L3 / Placement C	L3 / Option C: Independent Program	Placement C: Apart / Separate	Separate Program	L3: Separate / Simulated	**PLACE-SEP**
Governance Model	G1 / Arch A / Candidate A	G1: Sovereign Labyrinth	Arch A: Sovereign	Candidate A: Sovereign Monolith	G1: Sovereign Authority	**GOV-SOV**
Governance Model	G2 / Arch B / Candidate B	G2: Treaty-Governed	Arch B: Treaty-Governed	Candidate B: Federated Constitutional	G4 / Federated Hybrid	**GOV-FED** (Federated) / **GOV-TRTY** (Treaty)
Governance Model	G3 / Arch C / Candidate C	G3: Distributed Custodianship	Arch C: Distributed Custodianship	Candidate C: Distributed Treaty Constellation	G3: Distributed Quorums	**GOV-DIST**
Admission Criteria	G1–G6 / Gates 1–7	4 Eligibility Gates	Criteria G1–G6	Gates 1–7	Gates 1–7 (Perplexity) / G1–G6 (Gemini)	**GATE-1** to **GATE-7**

Appendix Table A2 — Consolidated Error & Citation Correction Ledger Across All Dossiers

TOPIC / SOURCE	PRIOR ERROR OR DISCREPANCY IN CYCLE 1 DOSSIERS	VERIFIED PRIMARY-SOURCE CORRECTION FOR CYCLE 2 BASELINE
OpenAI ExploitGym / Hugging Face Dates	Claude v0.1 left July 21 vs. July 27 unresolved.	Public disclosure notice was **July 21, 2026**; full post-mortem report published **August 26, 2026**.
Anthropic Cyber Eval Report Date & Count	ChatGPT v1.0 & Gemini v1.0 cited a "September 9 update" and "fourth incident" on Anthropic's page.	Anthropic's page is dated **July 30, 2026** (corrected **August 3, 2026**) and covers 3 incidents. The "4 incidents" count is from CSA's Oct 12 summary across labs.
UK AISI Incident Mechanism	Claude v0.1 summary implied all 2026 incidents were environments "assumed isolated that were not."	UK AISI (July 25–28, 2026; 19 actions in 10/122 runs) deliberately enabled internet access in a permissive harness; AISI explicitly noted it was *not* a sandbox escape.
Trail of Bits QEMU/KVM Escapes	Grok v0.1 implied all 3 escapes used zero-days; Gemini v1.0 used a descriptive paraphrase for the title.	Exact title: *"VMs won't contain cyber-capable agents."* Escape 1 used a known kernel bug; Escape 2 chained a known library bug; only Escape 3 used zero-days. Firecracker microVM was not escaped.
SandboxEscapeBench Saturation	ChatGPT & Claude could not verify whether v2/v3 saturated the benchmark after AISI's March blog.	Verified in `arXiv:2603.02277v3` (July 2026): Claude Mythos Preview saturated the public benchmark with 100% pass@5 on the hardest tasks.
Federal Grant Dates (DARPA & AISI)	Perplexity v1.0 listed DARPA BAA as Sept 8, 2026; Gemini v1.0 listed AISI Round 2 as projected late 2026.	DARPA IPTO BAA (`HR001126S0011`) posted **June 22, 2026** (closes **June 22, 2027**). UK AISI Alignment Project Round 1 is closed and a 2026 repeat is marked unlikely.
UN Nelson Mandela Rules Citations	Gemini v1.0 conflated indefinite and >15-day prolonged solitary; early reviews understated Rule 43(1)(c).	**Rule 43(1)(a):** Indefinite solitary; **Rule 43(1)(b) & 44:** Prolonged solitary (>15 days); **Rule 43(1)(c):** Dark or constantly lit cell; **Rule 32:** Prisoner experimentation; **Rule 45(1):** Not automatic by sentence.