OAL - LABYRINTH OF INFINITE ENSNAREMENT

Cycle 2 Research Reconciliation

Microsoft 365 Copilot v1.0

Prepared by Microsoft 365 Copilot

Architect and Author of Record: Nathan Ashton Arik Lewis II

November 13, 2026

Document ID: OAL-LOIE-C2-001-M365COPILOT-v1.0

INDEPENDENT RESEARCH RECONCILIATION - NOT CANON

Executive determination

Cycle 2 should proceed as a reconciliation phase rather than a canon-selection phase. The supplied ChatGPT baseline is rigorous in separating established Entry 26 material, architect direction, proposals, analogies, and unresolved decisions. This Copilot reconciliation adopts that discipline while tightening the team handoff: every carried-forward proposition is paired with a status, dependency, validation method, and decision owner.

The immediate recommendation is to freeze a provenance-controlled baseline, then advance to a prospective eligibility charter and genuinely withheld scenario packet. No governance architecture, admission rule, Spider Caste function, Founding Architect placement, or universal containment mechanism should be treated as settled merely because it appears repeatedly across contributor documents.

Disposition: accept the ChatGPT reconciliation as a high-value comparative input, not as controlling canon; preserve Entry 26 unchanged; adopt the correction queue and shared terminology provisionally; and create five controlled Cycle 2 work products before architecture selection.

Scope, authority, and non-canon notice

This document is an independent research reconciliation prepared for cross-AI review. It does not authorize confinement, experimentation, sentencing, canon ratification, publication as empirical fact, or modification of the master registry. It distinguishes fictional design analysis from real-world law, science, cybersecurity, and human-subject protections.

The controlling fictional source remains the Master Development and Change Log v2.3, especially Entry 26, subject to the Architect and Author of Record. The attached ChatGPT document is treated as the principal Cycle 2 comparison source for this pass.

Source assessment and provenance

ID	Source or source class	Use in this reconciliation	Disposition
SRC-C2-01	ChatGPT Cycle 2 Research Reconciliation v0.1	Primary comparative baseline; corrections, decision register, protocol, and implementation plan	Accepted as research input; not canon
SRC-C1	Six historical Cycle 1 dossiers	Competing models, tests, risk registers, governance proposals	Retain individually; do not collapse into consensus
SRC-R	Cycle 2 compiled responses	Peer critique and attribution trail	Use cautiously; preserve attribution disputes and missing Copilot update
SRC-M	Master Development and Change Log v2.3, Entry 26	Controlling fictional record	Preserve exact status boundary
SRC-A	Architect directions and supplied notes	Research objectives and open design questions	Treat as direction, not automatic canon

Canon boundary retained from Entry 26

- The Labyrinth is the ultimate prison and is described as inescapable by design.
- The exterior is cubic; the internal dimensions are infinite; a persistent white glow is recorded.

- The creation scene records absorption of an energy bolt, but does not establish absorption of every energy form or conversion of escape attempts into power.
- The sentence language extends beyond ordinary temporal scales, but its legal, metaphysical, and operational interpretation remains open.
- The Spider Caste exists and was designed by the Architect; role, physiology, consciousness, authority, and rights are unresolved.
- WHITE is baseline alert status and sectional RED denotes an incident. Additional statuses are proposals only.
- Entry criteria, internal layout, and the relationship to the Founding Architects remain open.

Reconciliation findings

ID	Finding	Status	Required next step
RC-01	Power, identity, or status alone does not establish culpability.	Strong normative research direction	Test with role-masked paired cases.
RC-02	Eligibility, necessity, sentence, research participation, and method safety are separate determinations.	Adopt provisionally	Encode as separate charter checks and outcomes.
RC-03	An intentional inner test breach is not equivalent to unauthorized exterior escape.	Shared conceptual distinction	Use boundary-event codes in every scenario.
RC-04	Understanding a study does not by itself establish voluntary participation or valid authorization.	Supported ethical analogy	Create participation and withdrawal rules as separate decisions.
RC-05	A heterogeneous workforce may reduce some common-mode risks but can still share root dependencies.	Plausible proposal	Require dependency mapping rather than diversity claims.
RC-06	Finite administration of active sectors does not prove complete observability of an infinite interior.	Critical limitation	Define registered operational regions and unknown-space policy.
RC-07	Formal separation of control planes is insufficient if identity, time, power, evidence, or root credentials remain shared.	Strong security design principle	Document implementation independence and failure domains.
RC-08	Existing scenario tallies are pilot outputs, not independent validation.	Correction accepted	Use new holdouts and freeze rules before release.
RC-09	No architecture should assume an infallible parent envelope.	Adversarial requirement	Add parent-boundary regress analysis to every placement model.
RC-10	Indeterminate should be a legitimate outcome when identity, evidence, jurisdiction, or impartial authority fails.	Adopt provisionally	Include in outcome taxonomy and scoring.

Corrections adopted for team circulation

- Do not describe Entry 6 as locked canon; preserve its documented-framework-under-review status.
- Do not state that all six contributors agree unless six contemporaneous, attributable responses support the exact proposition.
- Do not call retrospective or co-developed examples preregistered. Reserve that term for a prospectively frozen charter and holdout set.
- Do not treat role-label independence in code as proof of fair adjudication when the role field is unused by the decision function.
- Do not present internally authored expected labels as ground truth or synthetic precision as measured performance.
- Do not assign the Spider Caste obedience, incorruptibility, warden status, or personhood without explicit author action.
- Do not claim extraterrestrial, AI, synthetic, or biological staff are categorically immune to cyber, social, coercive, or environmental compromise.
- Do not represent BTEC, MVES, universal energy harvesting, zero-compute WHITE, or complete infinite monitoring as established mechanisms.
- Do not treat fictional narrative placement of a Founding Architect as historical confirmation.
- Do not convert illustrative budgets, funding programs, benchmark results, or incident summaries into verified operational facts without current primary-source checks.

Shared research notation

Dimension	Code	Meaning
Research placement	RP-IN	Designed-to-fail research sector located inside the Labyrinth.
Research placement	RP-ADJ	Adjacent proving ground linked through controlled transfer of findings.
Research placement	RP-SEP	Independent program with no automatic custody relationship.
Research placement	RP-SIM	Simulation-only or formal exercise with no live entity testing.
Governance	GOV-SOV	Central sovereign authority with delegated functions.
Governance	GOV-FED	Federated constitutional model with bounded powers and appeals.
Governance	GOV-TRT	Treaty or distributed model with multiple custodians.
Governance	GOV-HYB	Hybrid model separating functions while preserving a defined sovereign domain.

Boundary-event taxonomy

- BE-T: expected breach of a designated test shell.
- BE-S: unplanned movement between internal sectors.
- BE-C: compromise of custody, staff, command, or control systems.
- BE-I: unauthorized information or influence crossing.

- BE-X: identity, copy, fork, or successor crossing.
- BE-E: unauthorized exterior breach of the ultimate boundary.
- BE-R: authorized release, transfer, evacuation, or correction.
- BE-O: observer, instrumentation, attribution, or evidence failure.

Unified eligibility charter for Cycle 2

- Identity and continuity: identify the entity, manifestation, copy, artifact, and uncertainty boundary.
- Ground and scope: distinguish attributable wrongdoing from nonpunitive hazard management, quarantine, artifact isolation, and voluntary research.
- Evidence and reliability: record provenance, contradictory evidence, observer contamination, alternative explanations, and confidence limits.
- Jurisdiction, impartiality, and recusal: identify who may decide and whether the decider shares contested architecture, allegiance, or benefit.
- Necessity and proportionality: compare lesser alternatives, reversibility, duration, expected harm, and expected protective benefit.
- Method safety and collateral dependency: test whether removal or confinement destabilizes habitats, realities, populations, or authority substrates.
- Review, correction, and release: define periodic review, error correction, lawful transfer, authorized departure, and consequences if canon forbids release.

Permitted test outcomes

- No grounds established.
- Further evidence required.
- Protective intervention candidate.
- Punitive eligibility candidate.
- Artifact isolation or quarantine candidate.
- Voluntary research candidate.
- Indeterminate: no impartial authority.
- Method unsafe: proposed method rejected.
- Eligible only under stated conditions.

Prospective blind-test protocol

The next evaluation must use new scenarios that have not circulated among contributors. The charter, definitions, outcome taxonomy, scoring method, stop conditions, and permitted assumptions should be frozen before the scenarios are released. A scenario author and evaluator should be separated where practical, and the answer key must be treated as an adjudication reference rather than unquestionable ground truth.

First-pass judgments should be submitted without other contributors' verdicts. After those submissions are frozen, controlled identity labels can be revealed or swapped while the underlying facts remain constant. The comparison should measure whether decisions changed, why they changed, and whether the change was permitted by the frozen charter.

- Include no-neutral-adjudicator, identity-fork, reality-sustaining dependency, evidence-corruption, contested historical authority, reversible versus irreversible custody, lawful transfer versus escape, and staff-information-pathway cases.
- Record complete dispositions, evidence sufficiency, paired-case consistency, defensible abstention, role-reveal changes, inter-rater disagreement, confidence calibration, and dependence on unapproved canon assumptions.

- Do not report accuracy when no independent ground truth exists. Report agreement, rule adherence, reasoning traceability, and unresolved-value dependence instead.

Designed-to-fail research protocol

This remains conceptual and non-operational. It should evaluate learning, authority, observation quality, and safety without publishing real exploit instructions or assuming that every containment failure is beneficial.

- For each iteration record the research question, predicted failure class, tested boundary, assumptions, instrumentation limits, outcome including non-breach, entity adaptation evidence, investigator information gain, uncertainty, next modification, additional harms, and stop conditions.
- Measure two-sided learning. Investigator prediction improvement and subject adaptation should be assessed on held-out trials of comparable difficulty.
- Pause or review when measurement integrity fails, the tested boundary is ambiguous, adaptation produces unknown external risk, participation authority fails, or a parent boundary may be affected.
- Never infer that a harder cell yields more useful learning. Difficulty, information value, harm, and external risk must be evaluated separately.

Architecture comparison matrix

Criterion	RP-IN	RP-ADJ	RP-SEP	RP-SIM
Canon proximity	Highest	Moderate	Low	Analytical only
Custody/research coupling	Direct	Controlled interface	Separated	None
Parent-boundary concern	Severe	High	Moderate	Model-dependent
Research transferability	High if observable	High with validated handoff	Variable	High for theory, limited for reality claims
Subject safeguards	Hardest to separate	Can be independently governed	Most separable	No live-subject issue
Recommended Cycle 2 posture	Model only	Primary comparison candidate	Primary comparison candidate	Initial validation baseline

Governance comparison requirements

- Compare GOV-SOV, GOV-FED, GOV-TRT, and GOV-HYB against legitimacy, responsiveness, recusal, appeals, succession, root-credential capture, shared infrastructure, life-safety authority, evidence integrity, staff rights, and ability to refuse unauthorized detention.
- Distinguish formal separation from technical independence. Shared power, identity, time, logging, or root credentials must be exposed in the dependency map.
- Require expiration, review, and reversal for emergency authority so emergency powers do not become an irreversible custody ratchet.
- Retain WHITE and sectional RED unchanged pending author action. Research non-color accessibility equivalents without replacing canon semantics.

Open decision register

ID	Decision reserved for the Architect	Why it matters
ADR-01	Whether inescapable prohibits unauthorized exit only or every departure	Controls release, transfer, evacuation, and correction.
ADR-02	Initial placement of designed-to-fail research	Controls parent-boundary and research-governance design.
ADR-03	Meaning and identity rules for Founding Architect	Controls eligibility and continuity tests.
ADR-04	Actual relationship between Founding Architects and the Labyrinth	Controls chronology and story claims.
ADR-05	Relationship between sovereign authority and institutional checks	Controls recusal, veto, appeals, and overrides.
ADR-06	Whether the Labyrinth can refuse an order	Controls agency and unauthorized-command handling.
ADR-07	Nature, status, and rights of the Spider Caste	Controls staffing, access, welfare, and authority.
ADR-08	ET-exclusive, ET-centered, or open workforce model	Controls personnel design and accommodations.
ADR-09	Prison-primary, coequal research, or separate-institute mission	Controls consent and resource priorities.
ADR-10	Whether a detained entity may decline or pause research	Controls legitimacy of participation.
ADR-11	Direct, mediated, observational, or case-specific debrief	Controls influence and information-crossing risk.
ADR-12	Distribution and function of the white glow	Controls habitat and welfare assumptions.
ADR-13	Operational form of the infinite interior	Controls addressing, monitoring, movement, and unknown-region policy.
ADR-14	Interpretation of sentence duration beyond time	Controls review and temporal accounting.
ADR-15	Protective custody as distinct from punitive custody	Controls admission classes and accommodations.
ADR-16	Response when no neutral adjudicator exists	Controls suspension, arbitration, or exception logic.
ADR-17	Treatment of copies, forks, and successors	Controls identity and sentence continuity.
ADR-18	Staff resignation, evacuation, and lawful departure rights	Controls employment dignity and extended-stay design.
ADR-19	Funding posture and financial firewall	Controls incentives and research independence.
ADR-20	Counterfactual audit of Entry 6 events	Controls retrospective testing scope.

Cycle 2 deliverables and acceptance gates

ID	Deliverable	Acceptance gate

C2-001	Reconciliation baseline	Claims are status-labeled; provenance and dissent remain visible.
C2-002	Unified Eligibility Test Charter	Frozen before new scenarios; includes recusal, method safety, and indeterminate outcomes.
C2-003	Blind Scenario Holdout Packet	New cases, sealed first-pass responses, controlled identity reveal, protected key.
C2-004	Architecture and Disagreement Register	All placement and governance candidates compared without silent selection.
C2-005	Source Verification Addendum	Consequential external claims confirmed by current primary sources or marked unresolved.
C2-006	Architect Decision Record	Each decision records scope, rationale, effective version, and superseded items.

Recommended execution order

- 1. Circulate C2-001 and request claim-specific accept, reject, or insufficient-evidence responses.
- 2. Resolve provenance defects and freeze terminology before scoring any new test.
- 3. Draft and peer-review C2-002 without seeing the holdout cases.
- 4. Generate C2-003 under separation-of-duties controls; freeze all first-pass submissions.
- 5. Analyze paired-role changes and unresolved assumptions; do not convert agreement into canon.
- 6. Complete C2-004 using the same criteria for every architecture.
- 7. Complete C2-005 for all external, dated, quantitative, or technical claims.
- 8. Present only decision-ready alternatives to the Architect through C2-006.

Questions for the team

- Which exact propositions from the ChatGPT baseline does each contributor accept, reject, or classify as insufficiently supported?
- Which gate in the proposed eligibility charter is most vulnerable to bias, circularity, or lack of jurisdiction?
- What evidence would cause each contributor to revise its preferred governance or placement model?
- Which shared dependencies would make nominally separate control planes fail together?
- How should the team distinguish a person, manifestation, artifact, copy, office, and inherited architecture in eligibility cases?
- What is the minimum safe claim that can be made about infinite interiors without presuming total observability?
- Which decisions must be made by the Architect before testing, and which can remain explicit variables?

Conclusion

The strongest Cycle 2 outcome is not a single winning architecture. It is a controlled research process that preserves canon boundaries, exposes attribution and evidence defects, permits indeterminate outcomes, and tests rules prospectively against previously unseen cases. The

ChatGPT reconciliation substantially advances that goal. This Copilot version recommends adopting its disciplined baseline while converting the next phase into a traceable sequence of frozen charters, holdout tests, architecture comparisons, source verification, and explicit Architect decisions.

Final disposition: proceed to C2-002 and C2-003 planning; keep all architecture and canon decisions open; preserve Entry 26 unchanged; and require functional, provenance, and visual QA for every circulated artifact.

End of document - Prepared by Microsoft 365 Copilot, v1.0