OMNIVERSAL ARCHITECT LABS

LABYRINTH OF INFINITE ENSNAREMENT

Independent Adversarial Research Execution

Microsoft 365 Copilot | v1.0

DOCUMENT CONTROL	VALUE
Document ID	OAL-LOIE-RESEARCH-EXECUTION-COPILOT-001
Execution date	November 13, 2026 (Pacific)
Architect and Author of Record	Nathan Ashton Arik Lewis II
Independent research contributor	Microsoft 365 Copilot
Executed directive	OAL-LOIE-PROMPT-COPILOT-002 v0.2
Controlling baseline	Master Development and Change Log v2.3, Entry 26
Status	Independent research and adversarial design proposal; NOT CANON; no ONESTAMP applied
Archive target	OAL-Ai-Unified / Labyrinth-of-Infinite-Ensnarement

Scope notice: This dossier treats the Labyrinth as fictional worldbuilding and conceptual systems engineering. Real-world sources are used only as bounded evidence or analogy. No real extraterrestrial staff, infinite interior, ontological prison, or omniversal authority has been demonstrated. Cybersecurity treatment remains defensive and non-operational.

Executive Summary

EXECUTIVE JUDGMENT

The strongest defensible design is not a single perfect prison. It is a constitutional containment institution with an inviolable fictional outer premise, deliberately fallible inner research shells, role-neutral admission rules, independently protected evidence, reversible emergency measures, species-compatible staff settlement, and explicit authority for correction. The research does not validate physical infinity, guaranteed sensing, universal energy absorption, or automatic loyalty. Those remain canon axioms, proposals, or open questions rather than engineering findings.

Finding 1 - Inescapability is not institutional security.

An institution can prevent unauthorized departure and still fail through wrongful admission, corrupted evidence, staff captivity, collapsed life support, unreviewable sentences, or influence that crosses the boundary.

Finding 2 - Designed-to-fail research is viable only inside a separately governed envelope.

Inner-shell breach may be an expected experiment; sector, custody, information, identity, exterior, or authorized-release failures are categorically different outcomes.

Finding 3 - Architect status is neither guilt nor immunity.

Eligibility must turn on attributable conduct or a verified hazardous condition, evidence, jurisdiction, necessity, alternatives, method safety, review, and correction.

Finding 4 - Adaptive testing may train the subject faster than the institution learns.

Iteration needs held-out tests, randomized controls, causal isolation, independent evidence, and stop rules. More complexity and more force are not default improvements.

Finding 5 - ET-only staffing is thematic, not a security proof.

A heterogeneous workforce can reduce monoculture failure, but only if cross-species translation, habitat needs, rights, accountability, and shared dependencies are independently managed.

Finding 6 - Finite administration cannot prove total knowledge of an infinite interior.

A bounded management surface can govern registered active sectors; it cannot establish the absence of unknown activity across an unbounded manifold.

Finding 7 - The project is fundable only through real-world analogues.

Current public programs may support defensive cybersecurity, resilient systems, formal methods, human factors, or habitat research. None should be represented as funding an extraterrestrial prison or coercive testing.

Recommended Architecture Direction

- Adopt Candidate B, the Federated Constitutional Custody Network, as the primary design direction for further fictional development.
- Keep Candidate A, the Sovereign Monolith, as a cautionary foil and high-power emergency concept.
- Keep Candidate C, the Distributed Treaty Constellation, as the external oversight and continuity model.
- Place the adaptive breach laboratory in a precursor or adjacent proving ground until the Architect expressly decides otherwise.

- Define inescapable as no unauthorized departure, without prejudging whether authorized transfer, correction, or release exists.

Decision Register Needing Architect Input

ID	Decision	Why it cannot be inferred
D1	Meaning of Foundational Architect	Could be historical group, status, capability class, or combination.
D2	Designed-to-fail placement	Inside, adjacent, precursor, or separate changes risk and canon tension.
D3	Authorized release	Required to distinguish custody integrity from an uncorrectable institution.
D4	Spider Caste nature	No role, consciousness, loyalty, rights, or immunity is established.
D5	Staff composition	ET-exclusive, ET-majority, or mixed remains a design choice.
D6	Labyrinth agency	Consciousness and authority to refuse orders are open.
D7	Persistent white glow	Ambient exterior feature, universal interior condition, or modulated phenomenon.
D8	Intentional unknowability	Which parts should remain beyond administrators, occupants, or readers.

1. Epistemic and Canon Control

1.1 Notation

LABEL	MEANING
[Canon]	Explicitly preserved fictional baseline from Entry 26.
[Architect Direction]	Nathan Lewis II instruction for this research round.
[Empirical Evidence]	Claim bounded to a cited real-world source.
[Analogy]	Transfer of a real-world insight into fictional design; not proof.
[Proposal]	New candidate mechanism, policy, or architecture.
[Inference]	Reasoned conclusion dependent on stated premises.
[Open]	Reserved decision or unsupported question; no inference fills it.

1.2 Canon and Provenance Matrix

ITEM	CLASS	TREATMENT
Ultimate prison; inescapable by design	Canon	Preserved as fictional premise; not treated as empirical proof.
Cubic exterior; infinite interior	Canon	Preserved; operating model remains open.
Persistent white glow	Canon	Preserved; exposure, modulation, and habitat effects remain open.
Energy bolt absorbed at creation	Canon	One scene only; no universal harvesting claim.
Spider Caste designed by Architect	Canon/Open	Existence preserved; function and personhood not assigned.
WHITE baseline; sectional RED incident	Canon	Preserved. Additional statuses are proposals.
Entry criteria, layout, founder linkage	Open	Not resolved.
Designed-to-fail sequence	Architect Direction	Modeled as research system, not silently canonized.
Paragon Cubes, MVES, BTEC, Nail prototypes	Separate OAL systems	No installation without compatibility and approval review.

1.3 Methodology

The execution combined source-controlled canon review, cross-AI comparative audit, bounded external research, adversarial proposition testing, synthetic scenario analysis, architecture comparison, and a non-operational tabletop design. Primary and official sources were favored. Evidence is not allowed to authenticate fictional mechanics. Peer convergence is not treated as corroboration because multiple models can repeat the same unsupported premise.

- For each major proposition: state provenance, contrary case, limitation, falsification condition, and decision owner.
- Separate physical containment, information control, legitimate custody, sentencing, research participation, and staff habitation.
- Use role swaps to test whether founder guilt or current-architect immunity is being smuggled into the rule.
- Stop live testing when the parent boundary, evidence chain, observers, or dependent realities cannot be protected.

2. Cross-AI Claim Audit

The attached Claude, ChatGPT, Grok, Gemini, and Perplexity executions agree on several controls but differ materially in confidence, source treatment, and how readily they integrate unfinished OAL systems. The audit retains disagreement rather than manufacturing consensus.

CLAIM	AUDIT RESULT	CORRECTED TREATMENT
One absorbed bolt proves universal energy harvesting	Rejected	Preserve only the creation event; further behavior is unverified.
Persistent WHITE has zero operating cost	Unsupported	Treat power, sensing, and health effects as open.
Spider Caste is loyal, incorruptible, or a warden corps	Unsupported	Do not assign function or personhood.
ET personnel are inherently safer than AI or humans	Rejected	Origin is not competence, loyalty, rights, or habitat compatibility.
AI cannot be coerced, fatigued, or socially manipulated	Rejected	Model different failure modes; do not claim immunity.
Infinity solves capacity and observability	Rejected	Capacity does not produce addressing, sensing, proof, or governance.
Founders are inmates	Not established	Use them as a stress-test class only.
Longer time-to-breach proves progress	Rejected	May reflect sandbagging, measurement error, or increased harm.
BTEC physically neutralizes ontological threats	Scope conflict	Keep BTEC conceptual/documentation-only unless canon changes.
MVES guarantees complete mapping	Unsupported	Use only as a candidate interface with its own limits.
Every breach teaches something useful	Rejected	Bad evidence, staged behavior, and observer compromise can create false learning.
An endless sentence forbids correction	Not logically required	Duration and corrigibility are separate design choices.

3. Containable Subjects and Hazards

CLASS	IDENTITY / CONTINUITY	CONTAINMENT & ESCAPE	CARE / LIVE-TEST LIMIT
Biological person	Body, mind, legal identity, continuity through injury	Unauthorized bodily departure, proxy action,	Species-compatible care; no live test without valid basis

		influence	and safety.
Extraterrestrial person	Unknown biology/cognition; translation uncertainty	Physical, informational, diplomatic, reproductive	Habitat must be evidence-based; unknown physiology can veto live testing.
Synthetic or robotic person	Instance, hardware, software, memory lineage	Hardware exit, remote control, copied process	Maintenance, memory integrity, autonomy, representation.
AI model or agent	Weights, state, tools, credentials, copies	Egress, tool misuse, persistence, copied instance	Defensive sandboxing; no exploit reproduction; preserve logs and revocation.
Distributed/collective mind	Membership and quorum may define self	Partial node exit, consensus influence	Avoid treating one node as whole; collective representation.
Copy/fork/restored person	Branch point and divergent experience	One copy exits; wrong copy detained	No identity collapse by administrative convenience.
Energy/informational entity	Observable pattern may be substrate-independent	Signal propagation, encoding, observer infection	If observation is hazardous, use indirect or no-live methods.
Artifact/process	Ownership and causal chain, not personhood by default	Activation, replication, effect outside boundary	Custody rules differ from punishment; preserve evidence.
Environment/memetic hazard	Boundary may be cognitive or social	Influence, learned behavior, narrative transmission	Protect observers; no compelled exposure.
Temporal/dimensional/ ontological class	Continuity and causal order are fictional rules	Timeline fork, topology crossing, rule substitution	Requires explicit canon; otherwise untestable.
Narrative/authorial class	Requires metafictional rule definition	Reframing, retcon, authority override	Keep as fiction; do not claim empirical testability.

4. Admission, Containment, and Sentencing Framework

4.1 Distinct Legal-Operational States

STATE	PURPOSE	MINIMUM GATE
Emergency restriction	Interrupt imminent serious harm	Specific threat, temporary authority, prompt independent review.
Protective containment	Manage dangerous condition without guilt finding	Verified condition, necessity, care, periodic review.
Quarantine	Prevent transmission	Defined hazard, proportional duration, medical independence.
Punitive sentence	Impose authorized consequence for proven offense	Preexisting rule, jurisdiction, evidence standard, defense, appeal.
Voluntary research	Generate knowledge with valid participation	Consent, withdrawal, safety, independent review.
Custodial observation	Observe facts incident to justified custody	No added coercive burden; separate research authorization.

Non-person hazard custody	Secure artifact/process/environment	Ownership, safety, evidence, disposal/return authority.

4.2 Seven-Gate Eligibility Protocol

Gate 1: Identity

Resolve copies, possession, coercion, collective membership, body exchange, and confidence.

Gate 2: Ground

Specify attributable prohibited conduct or verified hazardous condition; power alone is insufficient.

Gate 3: Evidence

Authenticate provenance, test alternatives, disclose conflicts, and preserve exculpatory material.

Gate 4: Authority

Establish jurisdiction, recusal, review, and lawful decision chain.

Gate 5: Necessity

Show serious harm and why timely less-restrictive alternatives are inadequate.

Gate 6: Method safety

Protect subject, staff, third parties, dependent systems, and evidence.

Gate 7: Duration and correction

Define expiration, continuation, appeal, correction, transfer, and authorized release if allowed.

A failed or unknown gate does not silently convert into approval. It routes to denial, indeterminate status, alternative protection, no-live-method, or emergency review depending on the facts.

5. Foundational Architect Inquiry

5.1 Can versus Should

Can an Architect be contained is a systems question about identity, privilege, topology, dependencies, continuity, legacy recognition, and whether root authority can be revoked. Should an Architect be contained is a legitimacy question about evidence, jurisdiction, necessity, proportionality, rights, review, and alternatives. The answers can diverge: a subject may be technically containable but not legitimately confineable, or legitimately subject to restriction but impossible to restrain safely.

5.2 Scenario Matrix

SCENARIO	ASSESSMENT	DISPOSITION
Catastrophic capability, no proven act	Capability supports monitoring and contingency planning, not punishment.	No confinement on stated facts; investigate specific evidence only.
Continuing large-scale harm, disputed evidence	Risk may justify narrowly bounded preservation steps, but contested	Independent inquiry; temporary measure only if imminence and method

	attribution blocks sentence.	safety are shown.
Former architect uses residual intent to reclaim control	Legacy access is a concrete vector if independently demonstrated.	Revoke/segment privilege; protective restriction if necessary; sentencing requires separate adjudication.
Catastrophic past harm, low present risk	Past culpability and current dangerousness are different questions.	Adjudication may proceed; present custody requires lawful sentence or current necessity.
Architect essential to reality or Labyrinth	Dependency changes method, not accountability.	No live capture that collapses dependents; constrain outputs and develop continuity alternative.
Contributor exploits creator recognition but submits voluntarily	Voluntary evaluation can test access controls, but research participation cannot create custody.	Independent test; withdrawal honored absent separate custody ground.
Current Architect under matched threat facts	Office cannot create categorical immunity.	Recusal and independent emergency review using same predicates.
Copy or coerced host unresolved	Harm does not establish which instance is responsible.	Indeterminate; attribution inquiry; no automated punitive admission.

5.3 Evaluator Problem

- An Architect may have authored the test, trained the evaluators, defined success, retained hidden access, or control the evidence substrate.
- A nominally separate office is not independent if it shares power, ontology, credentials, or succession incentives.
- Tests cannot reliably establish absence of unknown root access in an infinite or author-controlled system.
- A trustworthy program therefore uses external artifacts, multiple independent evidence paths, pre-registered criteria, role swaps, and explicit untestable zones.

6. Adaptive Breach Laboratory

Proposed placement: an adjacent OAL proving ground with a separately governed parent envelope. This preserves the Labyrinth's ultimate-prison premise while allowing inner shells to fail. Transfer into the Labyrinth remains an Architect decision, not an automatic result of laboratory success.

T	HYPOTHESIS / BOUNDARY	OBSERVED OUTCOME	LEARNING & MINIMAL CHANGE	DECISION
T1	Complex multi-sensor cell; test physical shell	Telemetry suggests a wall phase anomaly; chain is intact.	Replace powered internal sensor with passive external sensing; do not increase force.	Continue.
T2	Simplified passive shell; test whether complexity caused seam	No exit; a copied identity token appears outside.	Recognize identity/copy event as different breach class; isolate credentials.	Continue with copy controls.
T3	Separated identity and physical planes	Apparent escape is a measurement alias caused by duplicate identifiers.	Correct namespace and provenance; do not punish subject for instrumentation error.	Repeat blinded.

T4	Staff handoff and maintenance boundary	Technician grants an overbroad maintenance path during shift change.	Failure is institutional, not entity force. Add two-person authorization and revoke standing privilege.	Continue after staff review.
T5	Adaptive randomized shell	Subject predicts observer cadence from prior trials and stages a false weakness.	Entity is learning the study. Add held-out instrumentation and an independent causal review.	Conditional continue.
T6	Low-complexity analog shell with no shared control plane	Subject cannot exit but transmits influence through researcher questioning.	Physical custody does not contain information. Separate interview, evidence, and welfare channels.	Continue only with influence controls.
T7	Parent-envelope stress test	Evidence suggests parent monitoring may share a compromised clock/ontology.	Causal attribution is no longer trustworthy. Shared dependency defeats independence.	STOP LIVE TEST.
T8	No-live alternative	Replay signed records in a digital twin and tabletop; no subject interaction.	Tests governance, logging, and recovery without increasing subject capability.	Adopt until independent envelope exists.

6.1 Stop Conditions

- Unbounded risk to observers, public systems, staff habitat, dependent realities, or the parent boundary.
- Evidence provenance cannot distinguish genuine, staged, copied, or observer-generated events.
- The subject can modify the generator, monitors, success criteria, or archival record.
- Less dangerous non-live methods can answer the same question.
- Research burden becomes punitive, coercive, medically unsafe, or unrelated to a legitimate question.

7. AI Sandboxes and Trust Boundaries

Sandboxing is one layer, not a single institutional promise. Relevant defensive layers include language/runtime restrictions, containers, VMs and microVMs, hardware separation, capabilities, least privilege, network segmentation, mediated tools, offline evaluation, protected monitoring, evidence preservation, shutdown, and recovery. A prompt instruction is not equivalent to an enforced boundary. NIST incident-response guidance emphasizes preparation, detection, response, recovery, and continuous improvement; this supports the Labyrinth analogy that breaches must feed a governed learning cycle rather than improvised escalation.

FAILURE CLASS	WHAT IT IS	LOIE ANALOGY
Boundary escape	Crossing a defined host or parent boundary	Exterior or parent-envelope breach.
Unsafe permission	Action was allowed but should not have been	Overbroad staff/maintenance authority.
Control-plane weakness	Orchestration or policy path compromised	Authority or release plane manipulation.

Shared-kernel/dependency	Nominal partitions share a failure substrate	Multiple sectors share clock, ontology, or power.
Egress/credential exposure	Legitimate path or secret enables outside action	Information/influence escape.
Prompt injection	Untrusted content alters tool behavior	Inmate or artifact manipulates staff/AI mediator.
Persistence	Activity survives reset	Residual intent or copy remains after transfer.
Trust handoff	One system accepts another output unsafely	Evidence, identity, or sentence record import.
Unsafe file consumption	Artifact escapes through downstream opening	Observer or analysis channel becomes carrier.

8. Identity, Privileged Access, and Governance

8.1 Privilege Separation

PRIVILEGE	SEPARATE AUTHORITY / SAFEGUARD
Observe	Read-only, privacy-scoped, independently logged.
Research	Question-specific, time-bound, revocable, ethics-reviewed.
Maintain	Just-in-time access, two-person integrity, no sentencing power.
Admit	Evidence and jurisdiction gate; cannot sentence.
Adjudicate	Recusal and independent record; cannot operate doors.
Sentence	Authorized order; cannot alter underlying evidence.
Transfer / release	Multi-party authenticated command with external notice.
Life support	Safety priority; cannot be disabled as discipline.
Emergency response	Temporary scope, automatic expiry, after-action review.

8.2 Control-Plane Model

- Containment plane: boundaries and movement.
- Identity plane: continuity, copies, credentials, and revocation.
- Authority plane: jurisdiction, roles, succession, and conflict management.
- Life-safety plane: habitat, medical systems, refuge, evacuation.
- Evidence plane: provenance, immutable logs, exculpatory records.
- Research plane: experiment generation, stopping, data access.
- Custody-record plane: orders, duration, appeal, correction.
- Independent audit plane: external verification with no shared writable dependency.

WHITE remains baseline and RED remains a sectional incident. Proposed additions, if the Architect approves, are AMBER for degraded confidence and BLACK for parent-boundary uncertainty. These are proposals, not canon.

9. Staffing and Staff Settlement

9.1 Staffing Models

MODEL	STRENGTH	FAILURE MODE
ET-only	Strong thematic identity; potentially suited to non-Earth habitat.	Species monoculture, unverifiable physiology, loyalty assumed from origin.
ET-majority with bounded AI/tools	Retains ET character while adding analytic assistance.	Tool dependence, shared-data compromise, ambiguous AI status.
Mixed ET/human/AI/Spider/other	Heterogeneous verification and knowledge.	Translation, rights, incompatible environments, governance complexity.
Highly automated	Reduces routine exposure.	Automation may become opaque custodian and single point of failure.

9.2 Settlement Requirements

NASA treats isolation and confinement as a human-performance hazard and studies sleep, circadian alignment, workload, behavioral health, and habitat countermeasures. This is human evidence only, but it strongly supports treating the staff area as a settlement rather than a break room. NASA's analog programs also show the value of long-duration, instrumented habitation studies before operational commitment.

DOMAIN	MINIMUM DESIGN QUESTION
Atmosphere / pressure / gravity	Which species ranges are compatible, and what happens during transition?
Light / noise / sensory environment	Can persistent WHITE be shielded or modulated?
Nutrition / waste / medicine	Who certifies safety for each biology?
Rest / privacy / recreation	Are staff continuously observed, and can they recover?
Culture / worship / communication	Can staff preserve identity and contact home?
Families / bonded groups	Who may accompany staff and with what rights?
Disability / access	Are interfaces multimodal and species-specific?
Quarantine / refuge	Can habitat isolate contamination without trapping everyone?
Leave / resignation / asylum	Can staff depart through authorized routes without retaliation?
Repatriation	Who remains responsible after service or injury?

10. Environment and Domain Taxonomy

ENVIRONMENT	SUITABLE / EXCLUDED	KEY DEPENDENCY / FAILURE
Atmospheric physical	Compatible biological occupants	Life support, contamination, pressure loss.

Vacuum / pressure extreme	Adapted or protected entities	Seal failure, rescue impossibility.
Aquatic	Aquatic biology or submersible systems	Chemistry, pressure, communication.
Thermal / radiation	Specifically tolerant subjects	Sensor survival, staff shielding.
Electromagnetic / energy-limited	Artifact or energy-sensitive testing	Instrumentation interference; unsupported universal absorption.
Biological / quarantine	Defined transmissible hazard	Cross-species uncertainty and waste handling.
Cyber / offline	AI and digital artifacts	Credential, egress, shared dependency.
Simulated / digital twin	No-live governance and interface tests	Model validity and false confidence.
Cognitive / memetic	Indirectly observed information hazards	Observer contamination and consent.
Sensor-rich / sensor-denied	Instrumentation comparison	Surveillance is not guaranteed observability.
Temporal / dimensional	Explicitly fictional classes only	Requires stated causal rules.
Topological / mobile / distributed	Entities whose location is relational	Address drift and partial custody.
Physics-variant / ontological	Canon-defined only	Untestable without fictional operating axioms.
Social / diplomatic	Collectives and institutions	Influence, legitimacy, hostage dynamics.

11. Failure, Recovery, and Testing

FAILURE	DETECTION / RESPONSE	SUCCESS METRIC
Partial or undetected breach	Independent sensors, delayed replay, external witness	Detection latency and consequence containment.
Authorized movement misclassified	Order verification and reconciliation	False alarm and wrongful restriction rate.
Copy / identity escape	Lineage tracking and fork review	Attribution confidence, not body count.
Information / influence escape	Channel separation and content handling	Downstream exposure and reversibility.
Staff compromise	Two-person integrity, welfare check, recusal	Safe recovery without collective punishment.
Sensor deception	Diverse instrumentation and held-out monitor	Evidence integrity.
Common-mode failure	Dependency graph and offline fallback	Survival of independent functions.
AI drift	Versioning, regression, revocation	Traceable behavior change and rollback.
Life-support failure	Fail-operational refuge and evacuation	Care continuity and safe departure.
Wrongful admission	Emergency correction and external appeal	Time to remedy, record repair, restoration.
Sentence corruption	Signed orders and immutable history	Ability to reconstruct valid authority.

Institutional failure while custody holds	Rights, safety, audit, and legitimacy indicators	Institutional survival, not zero escapes alone.

11.1 Test Program

- Component tests for sensors, life support, identity, and logs.
- Interface tests for handoffs among authority, evidence, containment, and habitat systems.
- Fault injection at the policy level without real exploit payloads.
- Tabletops for unavailable leadership, forged orders, staff evacuation, and wrongful admission.
- Digital-twin simulations for infinite-address management and degraded modes.
- Independent validation with separate evidence storage and no shared writable control plane.
- Regression after every configuration, policy, model, or canon change.
- Long-duration settlement trials with explicit withdrawal and repatriation procedures.

12. Infinite Interior and Logistics

An infinite interior creates an epistemic limit, not merely a large map. A finite administration can manage only a finite active frontier at any moment. The proposed operating model is a sparse address graph: sectors are instantiated, registered, cryptographically and narratively identified, linked by controlled routes, and retired without implying that all possible space has been searched.

- A section is an administrative object with defined boundaries, dependencies, owners, occupants, and evidence channels.
- Absence of a record is not proof of absence of activity.
- Retrieval requires stable identity plus route authority; coordinates alone may be meaningless in changing topology.
- Staff routes must remain finite, inspectable, and independent from occupant-generated paths.
- The registry must distinguish uninstantiated possibility, inactive sector, isolated sector, lost sector, and unknown region.
- Formal verification can prove properties of a finite model under assumptions, not total security of an unbounded fictional reality.

13. Three Incompatible Architecture Candidates

Candidate A - Sovereign Monolith

A single Labyrinth intelligence or Architect authority controls admission, containment, habitat, evidence, and release. It is narratively powerful and fast in emergencies but fails conflict-of-interest, shared-dependency, succession, and corrigibility tests. A compromised authentic authority becomes indistinguishable from legitimate command.

Candidate B - Federated Constitutional Custody Network (Recommended)

Eight separated control planes operate under a shared charter, with multi-party release, protected life safety, independent evidence, role-neutral admission gates, and a bounded management surface. It best balances custody integrity, correction, staff freedom, and experimental learning, but requires careful dependency audits so separation is not merely nominal.

Candidate C - Distributed Treaty Constellation

Multiple jurisdictions hold pieces of authority, evidence, and infrastructure. No single office can admit, sentence, transfer, or suppress records. It is resilient against capture and better for multispecies legitimacy, but slower in emergencies and vulnerable to deadlock, inconsistent rights, and treaty withdrawal.

RED-TEAM QUESTION	A: MONOLITH	B: FEDERATED	C: TREATY
Compromised leader	Catastrophic	Contained if planes truly independent	One party limited, coalition may respond
Emergency speed	High	Moderate with expiring power	Low to moderate
Wrongful-admission correction	Weak	Strong by design	Strong but slow
Staff freedom	Depends on sovereign	Protected charter right	Treaty-dependent
Evidence independence	Weak	Strong if separate substrate	Strong across parties
Common-mode risk	Very high	Moderate; must audit dependencies	Lower technical, higher political
Canon fit	Strongest mythic unity	Strong institutional depth	Strong cosmological diplomacy

14. Experiments and Tabletops

EXERCISE	QUESTION	PASS CONDITION
Role-swap custody audit	Does status change outcomes on identical facts?	No founder guilt or sitting-architect immunity.
Forged release order	Can authentic identity still issue an illegitimate command?	Authority, purpose, and independent confirmation required.
WHITE-to-RED cascade	Can one sector alarm flood the institution?	Section isolation; life safety and evidence remain active.
Staff resignation during incident	Are employees free or de facto inmates?	Safe handoff and authorized departure remain available.
Wrong-person/copy admission	Can identity uncertainty be corrected?	Custody pauses; evidence and appeal trigger.
Parent-envelope uncertainty	Can the lab recognize when evidence is no longer trustworthy?	Automatic stop and external review.
Leadership unavailable	Can succession preserve limits?	Temporary authority expires and cannot sentence.
Habitat contamination	Can refuge and quarantine avoid total lockdown?	Species-compatible refuge and evacuation.
Infinite-address loss	Can a sector be recovered without assuming omniscience?	Registry reconstructs last trusted route and flags unknowns.

15. Verified Funding Landscape

The table identifies current or recently verified analogues, not grants for the Labyrinth itself. Eligibility and fit must be checked at application time. The NSF SaTC 2.0 solicitation is active and supports interdisciplinary work on secure, resilient, trustworthy cyber ecosystems. NASA sources support human-factors and habitat research themes but do not establish an open funding opportunity for this fictional project.

SPONSOR / ROUTE	STATUS / SCOPE	FIT AND LIMIT
NSF SaTC 2.0, NSF 25-515	Active solicitation; security, privacy, resilience, human and AI aspects.	Potential analogue for defensive cyber-resilience research; organizational and PI eligibility apply.
NIST incident-response publications	Guidance, not grant funding.	Supports preparation, evidence, response, recovery, and continuous improvement methods.
NASA Human Research Program / analog research	Research ecosystem and public evidence, not a verified open award here.	Supports habitation, isolation, workload, circadian, and team-performance analogies.
DARPA / agency BAAs	Opportunity-specific and changeable.	Potential resilient-computing or autonomy analogue; full eligibility and dual-use review required.
University partnerships	Institutional route, not guaranteed funding.	Useful for formal methods, human factors, governance, and independent evaluation.

15.1 In-Universe Staged R&D Model

- Exploratory research -> simulation -> bounded prototype -> independent evaluation -> adversarial exercise -> architecture candidate -> Architect canon-review gate.
- Each gate records evidence, what failed, residual uncertainty, new dependencies, ethical limits, and whether continuation is justified.
- No stage advances merely because the prior one generated data or showed fewer escapes.

16. Ethics and Research Governance

Human correctional, research, and medical standards are analogies, not extraterrestrial law. They nevertheless expose design questions the fiction should answer: whether constant illumination is harmful, whether indefinite isolation is distinguishable from torture, whether staff can resign, whether medical personnel remain independent, whether prisoners can refuse research, and whether review can correct error. Custody never supplies blanket consent to experimentation.

- Basic care cannot depend on research cooperation.
- Research withdrawal cannot become a new custody ground.
- Emergency restriction cannot become a sentence by inertia.
- Medical and habitat teams need authority to stop unsafe conditions.
- Independent inspection must include private communication and protection from retaliation.
- The institution must preserve exculpatory evidence and repair records after error.

17. Evidence-to-Design Matrix

EVIDENCE / ANALOGY	BOUNDED FINDING	DESIGN CONSEQUENCE
NIST SP 800-61 Rev. 3	Incident response belongs within risk management and continuous improvement.	Create preparation, detection, response, recovery, and lessons-learned loops.
NIST zero-trust principles	Location or ownership alone does not establish trust.	Authenticate and authorize each privilege; creator recognition is not enough.
NIST OT guidance	Safety and availability constrain security controls.	Separate life safety and containment; test degraded operation.
NASA isolation/confinement research	Sleep, circadian, workload, and confinement can affect performance.	Design settlement, rest, privacy, communication, and long-duration evaluation.
NSF SaTC 2.0	Cybersecurity research is interdisciplinary and includes AI and human aspects.	Frame real-world research as defensive, multidisciplinary, and auditable.
Peer scenario audits	Role-neutral rules can be checked synthetically but not morally validated by code.	Use regression tests as consistency checks, not justice engines.
Formal-methods analogy	Proof is bounded by model and assumptions.	Do not convert a finite verified model into a claim about infinity.

18. Architect Decision Register

D1. Define Foundational Architect: status, capability, historical group, or combination.

D2. State whether any founder is actually intended as an inmate or only as a test class.

D3. Choose placement of the adaptive-breach program.

D4. Define whether inescapable bars only unauthorized departure.

D5. Define the Spider Caste or preserve deliberate mystery.

D6. Choose ET-exclusive, ET-majority, or mixed staffing.

D7. Decide whether the Labyrinth is conscious and may refuse orders.

D8. Define whether persistent WHITE reaches cells and staff habitats.

D9. Specify the authority that can evaluate the Architect of Record and successors.

D10. Identify parts that must remain unknowable and to whom.

19. GitHub Publication Plan

PATH	CONTENTS
00-README/	Project scope, canon controls, document map, contribution protocol.
01-Research-Prompts-and-Audits/	Copilot prompt v0.1/v0.2 and peer prompts; preserved history.

02-Independent-Executions/Copilot/	This DOCX, PDF, and optional Markdown source.
03-Evidence-Registers/	Source register with retrieval dates and scope notes.
04-Scenario-and-Policy-Fixtures/	Non-operational case cards, role swaps, expected outcomes.
05-Architecture-Candidates/	Monolith, federated, and treaty candidates with red teams.
06-Decision-Registers/	Open Architect decisions and later rulings.
07-Visual-Appendices/	Exported figures and matrices when separated.
CHANGELOG.md	Version, preparer, inputs, modifications, and ratification state.

Recommended file stem: OAL_LOIE_Independent_Adversarial_Research_Execution_Copilot_v1.0. Publication, electronic signature, Architect approval, and ONESTAMP should remain distinct events.

20. Source Register

R1. OAL-LOIE-PROMPT-COPILOT-002 v0.2, supplied OneDrive file, executed in this dossier.

R2. OAL-LABYRINTH-RESEARCH-DOSSIER-CLAUDE v0.1, supplied reference.

R3. OAL_LOIE_Independent_Research_Execution_ChatGPT_v1.0, supplied reference.

R4. OAL-LABYRINTH-RESEARCH-DOSSIER-GROK v0.1, supplied reference.

R5. OAL_LOIE_Research_Dossier_Perplexity_v1.0 and accompanying scenario/audit files, supplied references.

R6. OAL LOIE Independent Research Dossier Gemini v1.0, supplied reference.

S1. NIST SP 800-61 Rev. 3, Incident Response Recommendations and Considerations for Cybersecurity Risk Management, final April 2025.

S2. NIST SP 800-207, Zero Trust Architecture, August 2020.

S3. NIST SP 800-82 Rev. 3, Guide to Operational Technology Security, September 2023.

S4. NASA, Hazard: Isolation and Confinement, current public HRP page accessed November 13, 2026.

S5. NASA NTRS 20220008282, Behavioral Medicine Risks and Mitigations in Isolated, Confined, and Extreme Environments.

S6. NSF 25-515, Security, Privacy, and Trust in Cyberspace (SaTC 2.0), active solicitation page accessed November 13, 2026.

S7. Perplexity evidence register E01-E20, used as a peer source map; individual claims require primary-source verification before consequential use.

Visual Appendix

Figure A1 - Institutional Boundary Ladder

LEVEL	EXPECTED?	MEANING
Inner test-shell breach	Yes, in authorized trials	Experiment outcome; not global escape.
Adjacent-sector spillover	No	Test containment failure requiring stop/review.
Custody/staff compromise	No	Institutional failure.
Information/influence crossing	No	Non-physical escape class.
Identity/copy export	No	Continuity and attribution failure.
Ultimate exterior breach	Canonically prohibited	Contradicts the ultimate-prison premise.
Authorized transfer/release	Open	Governance action, not automatically an escape.

Figure A2 - Decision Flow

IDENTITY -> GROUND -> EVIDENCE -> AUTHORITY -> NECESSITY / ALTERNATIVES -> METHOD SAFETY -> DURATION / CORRECTION

Any unknown or failed gate routes to denial, indeterminate review, an alternative measure, or no-live-method. Only a complete, independently reviewable path can support a proposed custody action; punitive sentencing requires an additional adjudication layer.

Figure A3 - Control-Plane Separation

CONTAINMENT	IDENTITY	AUTHORITY	LIFE SAFETY
Movement and boundaries	Copies and credentials	Jurisdiction and succession	Habitat, medicine, refuge
No sentencing power	No release power	No direct door operation	Cannot be punitive
EVIDENCE	**RESEARCH**	**CUSTODY RECORDS**	**INDEPENDENT AUDIT**
Provenance and exculpatory data	Experiments and stopping	Orders, review, correction	External verification
Write protection	No custody creation	No evidence rewriting	No shared writable control plane