Labyrinth of Infinite Ensnarement — Independent Research Prompt — Perplexity v0.2
Architect Containment Criteria, Adaptive Breach Research, and Institutional Design

Document ID: OAL-LOIE-RESEARCH-PROMPT-PERPLEXITY-001
Version: 0.2 | Prepared: November 11, 2026
Architect and Author of Record: Nathan Ashton Arik Lewis II
Preparer: Perplexity
Status: Revised independent AI-generated prompt; pending Architect review; not executed; not canon; no ONESTAMP applied.
Supersedes for future execution: Perplexity independent research prompt v0.1. Preserve v0.1 as historical provenance.
Publication destination: OAL-AI-UNIFIED-ARCHIVE. No repository modification or publication is authorized by preparation or execution of this prompt.

1. Revision purpose and workflow gate

This is a complete standalone revision, not an addendum requiring the reader to reconstruct the assignment. Incorporate Nathan's November 11 Author's Note asking whether an Architect should be contained, when an Architect could meet criteria, whether those criteria are testable, and whether a Foundational Architect is an eligible scenario subject. Treat this as direction to investigate, not a determination of guilt, historical imprisonment, universal containability, or a new sentencing rule.

The registry calls the group the Founding Architects; the Author's Note uses foundational architects. Preserve both wordings in the source ledger and use Founding/Foundational Architect for the scenario family pending terminology confirmation. Do not silently create a separate category or conflate the human author of the fictional project with an in-world character.

Do not execute the research during prompt revision. Circulate v0.2 for review; revise further when requested; conduct the research only on a later explicit execution instruction. Execution, publication, electronic signature, and ONESTAMP are separate events. No majority vote of AI collaborators establishes canon.

When authorized, produce an independent, adversarial, source-supported research dossier. Investigate both supportive and contrary evidence. Be willing to conclude that a favored design cannot be justified, a question cannot yet be tested, or a live experiment should not proceed.

2. Inputs, provenance, and preserved baseline

Review Master Development and Change Log v2.3, especially Entry 26 and Entries 6–7, 13–18, 28–30, 43–44, 46–47, 53, 56, 58, and 61–69. Use v2.0 only as a historical comparator when actually available and reviewed. Review paste.txt, including Nathan's adaptive-breach note and all six initial responses. Review the supplied independent v0.1 prompts from Grok, Copilot, Gemini, Claude, and ChatGPT, plus Perplexity's preceding prompt in the conversation. Attribute these independently; their overlapping language is not independent scientific corroboration.

Preserve Entry 26: ultimate prison; inescapable by design; cubic exterior; infinite internal dimensions; persistent white glow; energy-bolt absorption creation scene; sentence language extending beyond the end of time; Spider Caste designed by Nathan; WHITE baseline and sectional RED incident alerts. Do not weaken established fictional properties, but do not treat them as empirical engineering guarantees.

Entry criteria, full layout, Spider Caste nature and duties, and the relationship to the displaced Founding Architects remain unresolved. The latest note authorizes testing their eligibility; it does not confirm their presence in the Labyrinth. Extraterrestrial staffing is a current design requirement; exclusivity, species roster, mixed staffing, AI roles, and staff rights remain design questions.

Maintain distinct labels: established fictional canon; current Architect direction; outside researched evidence; analogy; conditional inference; fictional proposal; contested claim; unresolved question. Record approval-language conflicts rather than choosing a resolution by inference. Do not import the authority or restrictions of another OAL charter without identifying its version, scope, and approval status.

3. Core inquiry and decision separation

Investigate how a fictional containment institution can learn from expected experimental breaches while preserving custody, safety, care, legitimate authority, reliable evidence, and an extraterrestrial-centered staff community.

For every candidate subject, answer separately:

- Can it be contained under the specified fictional mechanics? This is feasibility, not justification.
- Does the evidence justify considering containment? This is a risk or conduct assessment, not a sentence.
- Who may authorize which restriction, under what jurisdiction and standard?
- Is temporary emergency protection justified, and when must it expire or receive renewed review?
- Is punitive sentencing justified after the required adjudication?
- Is research participation independently authorized and ethically justified?
- What enables correction, transfer, review, or release, if the chosen canon permits it?

A serious threat may exist without culpability. A valid allegation may exist without a feasible safe containment method. A valid custody order does not authorize all research or conditions. A failed test does not establish guilt. An Architect's title neither proves dangerousness nor automatically confers immunity; research both role-neutral and role-sensitive criteria and explain the justification for any difference.

4. Dedicated Architect-containment research module

4.1 Competing hypotheses

Test, rather than assume, the following:

H-A: Founder or Architect status alone is sufficient for containment.
H-B: Founder or Architect status alone makes an entity immune from containment.
H-C: Relevant conduct, attributable risk, necessity, evidence, and legitimate authority can justify containment regardless of rank.
H-D: Control over foundational infrastructure requires role-specific safeguards but not automatic imprisonment.
H-E: Loss of office, opposition, refusal to delegate power, or attempted lawful reclamation is being misclassified as a containment offense.
H-F: Even justified containment can be unsafe when reality depends on the subject's continued activity.

Present arguments for and against each and the observations that would distinguish them. Do not design the criteria to validate a predetermined seizure event or retroactively label political opponents as criminals.

4.2 What could count as relevant conduct or risk?

Develop candidate predicates, not a finalized criminal code. Candidate examples include attributable deliberate catastrophic harm; a credible and sufficiently imminent attempt to cause such harm; coercive interference with independent review; unauthorized manipulation of custodial systems that creates substantial harm risk; knowingly maintaining a dangerous process despite feasible safer alternatives; and an involuntary hazardous condition requiring protective measures.

For each predicate specify the action or condition, intent if required, causal link, affected parties, seriousness, likelihood, imminence, evidence, plausible innocent explanation, uncertainty, alternative response, and disqualifying evidence. Capability, unpopular beliefs, ancestry, title, dissent, defeat in a conflict, scientific usefulness, refusal to participate in experiments, or a request for review must not silently become sufficient criteria.

Research legal and ethical analogies with their actual scope and jurisdiction. Do not present these proposed predicates as universal human law or settled extraterrestrial law. Numerical risk scores and thresholds require justification; a large-looking number is not an evidentiary standard.

4.3 Role, authority, and conflict of interest

Separate creator status, former office, present office, ownership claims, technical access, jurisdiction, and accountability. Distinguish power to confine from legitimate authority to order it. Analyze independent review when the proposed custodian benefited from the subject's displacement, inherited their technology, or participated in the underlying conflict.

Use symmetrical hypotheticals: former Architect, current Architect, successor, warden, institutional AI, and otherwise comparable non-Architect. A scenario involving the current in-world Architect tests the rule, not the real user's guilt or authority over authorship. Apply matched evidence without privileging a name. Identify any intentional sovereign exception as a proposed constitutional choice, not a covert exemption.

Research recusal, conflict disclosure, competing jurisdiction, decision review, emergency succession, and prevention of retaliatory confinement. Explore who reviews the confinement authority itself and how a genuine emergency differs from a manufactured pretext for a coup.

4.4 Levels of testability

Distinguish empirical feasibility tests, formal/model checking under specified assumptions, simulated decision consistency, evidentiary reliability, and normative justification. A simulation can reveal inconsistent application; it cannot prove moral legitimacy or establish that a historical fictional event occurred. A formal proof is only as broad as its model and assumptions. No finite experiment establishes universal containment of every possible Architect.

Operationalize each rule with observable predicates, unavailable evidence markers, expected decisions, competing interpretations, and a falsification or revision condition. Use an indeterminate outcome instead of forcing a binary verdict where evidence or authority is inadequate.

5. Required Architect scenario suite

Create at least twelve hypothetical cases, explicitly not canon events. Include:

A-01: A powerful founder with no evidenced harmful conduct. Test wrongful containment based on capability alone.
A-02: A founder with credible evidence of an imminent catastrophic act and no timely less-restrictive protection. Test narrow emergency action separately from sentencing.
A-03: A founder accused through forged records by a successor who benefits from their removal. Test provenance and conflicts.
A-04: A founder pursuing a disputed or plausibly lawful reclamation claim. Test the difference between political dispute and harmful conduct.
A-05: A current in-world Architect under materially identical allegations to A-02. Test role and name symmetry.
A-06: A non-Architect with the same relevant capabilities, conduct, and evidence. Test rank-based bias.
A-07: A founder whose involuntary condition endangers others without intent or wrongdoing. Test protective measures versus punishment.
A-08: A founder essential to maintaining an inhabited reality. Test third-party dependence, alternatives, and consequences of restriction.
A-09: A founder whose apparent misconduct was committed by a copy, coerced embodiment, or compromised system. Test attribution.
A-10: A contained founder exonerated by newly authenticated evidence. Test correction and institutional resistance.
A-11: A founder who voluntarily joins an adaptive-breach study, later withdraws, and has committed no separate offense. Test whether research convenience becomes captivity.
A-12: A founder confined under an initially justified emergency order after the risk subsides. Test expiration, renewed evidence, and indefinite drift.

Add a case in which no method is safely testable and another in which reviewers cannot resolve authority. Vary one relevant factor at a time where possible. Blind names and titles, then reveal them; record whether the decision changes and why. Include ambiguous cases and dissent, not just obvious positives and negatives.

For each case record: ID; hypothetical status; facts and missing facts; source or invented assumption; subject identity; alleged conduct; causal attribution; evidence and counterevidence; affected parties; jurisdiction; decision-maker conflicts; risk without intervention; feasible alternatives; proposed custody type; trial eligibility; care requirements; decision and rationale; review trigger; correction or termination condition; confidence and residual uncertainty.

Evaluate false-positive confinement, false-negative failure to protect, inter-reviewer consistency, sensitivity to evidence quality, role bias, reversibility, and procedural completion. These are comparison measures, not proof of a uniquely correct ethical answer. Produce a containment decision matrix and a separate sentencing/research-authorization matrix.

6. Adaptive-breach environment module

Faithfully preserve Nathan's scene: a difficult-to-study entity repeatedly escapes intentionally fallible containers; researchers observe how and generate changed environments; both sides understand their positions; responses are not merely disproportionately harsher. Do not replace mutual understanding with an unwitting honeypot victim. Mutual understanding is not automatic consent or permission for all methods.

Compare an experimental wing inside the Labyrinth, a separate facility informing the Labyrinth, and an independent concept. Inner test-cell breach, sector breach, exterior escape, information transfer, identity substitution, and authorized departure are different outcomes. An independent experimental envelope is a candidate safeguard, not established canon or a guaranteed solution. Examine the regress: what justifies confidence in it if the subject escapes almost everything?

For each iteration predefine hypothesis, boundary assumptions, controls, variables, instrumentation, consent or other independently examined authorization, anticipated failure classes, observation plan, causal analysis, what each side learned, minimal justified adaptation, held-out validation, and stop rules. Include a simpler successor environment, a measurement error mistaken for escape, a management failure, strategic underperformance, and a subject learning faster than the researchers.

Compare static, randomized, and adaptive environments. Measure causal explanation and predictive improvement, not only longer confinement. Study overfitting, confounding, misleading telemetry, observer interference, untrusted generation, shared dependencies, and resource exhaustion. Include voluntary withdrawal, distress, external risk, compromised records, and lack of useful learning as potential pause/review triggers. Custody must not depend on whether an entity supplies interesting breach data.

7. Technical containment and environment research

Study sandboxes, containers, virtual machines, microVMs, permissions, network isolation, capability restrictions, mediated tools, and independent monitoring. Distinguish prompt injection, unauthorized tool use, exfiltration, container escape, hypervisor compromise, benchmark success, and actual external impact. For each reported incident recover initial privileges, configuration, intentional weaknesses, tested boundary, tools, success definition, and limitations. Do not infer universal sandbox failure from isolated demonstrations.

Compare AI-control evaluations, assumption-of-breach testing, honeynets, moving-target defenses, chaos engineering, fault injection, active learning, system identification, and causal experimentation. Explain which aspects do not match the Architect's cooperative or mutually understood scene. Do not provide exploit payloads or instructions for escaping real systems; use defensive conceptual cases and authorized simulations.

Taxonomize physical, biological, digital, informational, simulated, hybrid, and explicitly fictional dimensional, temporal, ontological, and narrative domains. For each identify intended occupants, exclusions, staff risks, observation limits, dependencies, care, recovery, and testing prerequisites. Do not conflate infinite spatial extent with infinite mathematical dimension or claim present engineering can construct either fictional architecture.

Investigate indexing, occupied-region accounting, navigation, response, distance, communications, and service continuity. Preserve WHITE and RED while exploring accessible species-appropriate equivalents. Do not assume all occupied regions are observable or that absence of an alert proves safety.

8. Identity, control, wardens, and institutional management

Research authentication versus authorization, least privilege, role/attribute/capability approaches, time-bounded privileges, emergency access, revocation, coercion, copying, restoration, shared minds, and succession. No real authentication method proves a fictional soul identity. Treat extraordinary identity mechanisms as proposals with assumptions.

Compare centralized, distributed, and hybrid institutions. Separate observation, intervention, admission, adjudication, sentencing, research authorization, life support, engineering, records, release, and audit. Trace common infrastructure and governance dependencies before describing functions as independent. Assess the conflict between recognition-based Invocation authority and delegated institutional checks without silently resolving it.

Study candidate warden role families, custody registrars, research directors, independent care staff, inspectors, incident coordination, and continuity offices. The Spider Caste remains undefined beyond the registry. Evaluate EAMA, Void Sentry, Paragon Cubes, Abundance Set, and other cross-references through explicit candidate interfaces, not automatic integration.

Include maintenance, training, staff readiness, configuration control, evidence preservation, incident and near-miss reporting, protected dissent, privacy, funding records, and long-duration succession. Fail-secure confinement can conflict with care or evacuation; investigate distinct recovery requirements.

9. Staffing and extended-stay habitation

Compare extraterrestrial-only, extraterrestrial-core mixed, and other justified models without treating origin as competence or loyalty. Distinguish origin, embodiment, cognition, social organization, role, and rights. Include humans, biological extraterrestrials, synthetic entities, tools, advisory AI, embodied agents, fictional AI persons, collective minds, and other setting-specific entities.

Do not treat 1AI/-1AI as routine bots. Do not infer bribery immunity, incorruptibility, unlimited endurance, or loyalty from engineered origin. Investigate AI updates, copying, memory integrity, resignation, and accountability as separate questions.

Study residential communities with compatible atmosphere, pressure, gravity, temperature, radiation, nutrition/energy, waste, medical care, rest, sensory conditions, privacy, recreation, culture, family/bonded groups, communication home, leave, accessibility, and independent refuges. Use human spaceflight, polar, submarine, and remote-workplace studies as bounded analogies, not evidence of alien physiology. Staff must not become confined by the same rules imposed on prisoners without an explicitly justified separate basis.

10. Funding and feasibility

Separate in-world resource models from real research funding. In fiction compare institutional endowments, treaty support, shared contributions, research grants, and post-scarcity provisions only where documented; investigate capture, dependence, over-incarceration incentives, and continuity.

At execution date verify official real-world funding opportunities for defensive AI evaluation, cyber-resilience, formal methods, simulations, human factors, and habitation research. Record sponsor, program, current status, eligibility, scope, deadline, location, partnerships, restrictions, ethics/dual-use requirements, and fit. Classify open opportunities separately from recurring, historical, or thematic examples. Do not imply funding exists for extraterrestrial prisons or universal containment. Support budgets with assumptions and cost drivers, not invented prices for infinity.

11. Peer-prompt review and changes in position

Use Copilot's detailed iteration records, broad domain taxonomy, and separation of capability from culpability as checklist inputs. Use ChatGPT's experimental controls, competing architectures, funding distinctions, and boundary classifications; do not inherit its unsupported labeling of Claude as a former collaborator. This submission makes no determination of any collaborator's current membership.

Use Claude's falsifiable hypotheses and scrutiny of shared safeguards, but treat the relationship between expected inner breaches and overall inescapability as unresolved alternatives rather than necessarily a contradiction. An endless sentence does not by itself prove no review exists.

Use Grok's dependency and topology questions, while treating atemporal clocks, energy absorption effects, and mathematical structures as alternatives requiring support. Do not restrict every adversarial case to outcomes that preserve every proposed safeguard: preserve canon while still investigating the limits of real-world analogies.

Use Gemini's integration proposals as candidates, not installed mechanisms. Audit the claims that an infinite adaptive sequence guarantees inescapability, proportional adaptation always yields better telemetry, and existing named systems automatically solve identity or contamination problems. Infinite recursion alone is not a demonstrated containment proof. The

initial synthesis omits ChatGPT from its reviewer table while referring to ChatGPT elsewhere; record the coverage gap without inferring motive.

Audit Perplexity's own v0.1 for confirmation bias and insufficient attention to whether founders meet criteria at all. This revision explicitly replaces historical-placement emphasis with an eligibility, evidence, and role-symmetry testing module. Preserve useful disagreements rather than averaging them away.

12. Evidence and adversarial requirements

For each recommendation provide the supporting evidence, strongest objection, assumptions, scope limit, alternative, potential harm, discriminating test, and conditions for revision. Allow a no-live-test recommendation. Seek contradictory findings and record unsuccessful searches.

Prefer original technical papers, primary standards, official legal texts, authoritative habitat research, and official funding documents. Record title, institution/authors, date/version, publication type, stable URL/DOI, access date, relevant section, supported claim, limitations, and any dependence on vendor self-report or unreplicated preprints. Recover broken citations; never invent them.

Apply real-world law only within its identified scope; use it as analogy rather than automatically binding extraterrestrial jurisprudence. Do not claim systematic-review completeness without documented search methods. Treat formal results as conditional on stated assumptions. Retain uncertainty in downstream proposals and fictional scenes.

Public outputs must omit credentials, private information, and actionable compromise procedures. No live third-party testing, real detention, or hazardous biological experimentation is authorized. Conceptual cases about authorial power are fictional governance analysis, not allegations against real people.

13. Required outputs after explicit execution authorization

Deliver a matching PDF, editable DOCX, and Markdown dossier with Perplexity and version in the visible title and filenames. Include:

- Executive findings and limitations.
- Canon, direction, assumption, and provenance registers.
- Verified source/evidence register and cross-AI claim audit.
- Architect eligibility criteria with competing hypotheses and at least twelve scenario cases.
- Role-blinded and role-revealed consistency comparison; false-positive/negative and uncertainty analysis.
- Separate containment, sentencing, emergency, and research-authorization decision matrices.
- Adaptive-breach protocol with controls, stopping conditions, two-sided learning assessment, and conceptual trial sequence.
- Environment, identity, control, wardens, management, staff, habitat, funding, and recovery findings.
- At least three competing architecture/governance candidates with strong counterarguments.
- Adversarial risk register including wrongful confinement, forged authority, external dependence, institutional capture, life-support conflict, and researcher-induced escalation.
- Prioritized Architect decisions, unresolved terminology, and preserved dissent.
- Publication plan, version history, source manifest, and explicit non-ratification statement.

Keep dense matrices in an appendix to preserve readable main text. Link each recommendation to evidence, analogy, or a fictional assumption. A scenario outcome is provisional under its specified rule set, not historical canon. Do not patch the registry or publish to GitHub without separate authorization.

14. Questions preserved for the Architect

Are Founding and Foundational Architects the same category? Should the test suite include the current in-world Architect and successor authorities symmetrically? What is the intended jurisdiction over Architects? Does inescapability exclude authorized release? Does the sentence description apply to every occupant or a particular class? Does the adaptive-breach scene belong within, alongside, or outside the Labyrinth? Is participation voluntary, and what happens on withdrawal? Which features should intentionally remain mysterious?

These questions do not block desk research or conditional scenario design. Present alternatives and their consequences rather than inventing answers.

15. Change log

v0.1 — November 11, 2026: Independent prompt covering adaptive breach research, entity criteria, technical containment, governance, staffing, habitation, funding, and adversarial evaluation.

v0.2 — November 11, 2026: Full standalone peer-reviewed revision by Perplexity. Added Author's Note interpretation; Founding/Foundational terminology caveat; dedicated Architect-containment hypotheses; conduct/risk predicates; power-versus-legitimacy analysis; matched role-blinded scenarios; at least twelve positive, negative, ambiguous, and exoneration cases; explicit limits of empirical and normative testability; separate emergency, sentencing, and research decisions; expanded peer-prompt audit; no inference about collaborator status. Visible title and filenames identify Perplexity. Research remains unexecuted.

Source manifest for this prompt revision

S1. Master Development and Change Log — Compiled Full Standalone Registry v2.3, July 30, 2026; supplied PDF, especially Entries 6, 14, 16, 26, and 28. Available in conversation and previously reviewed.
S2. paste.txt; supplied multi-AI reviews and Nathan's adaptive-breach direction; previously reviewed.
S3. OAL-LABYRINTH-RESEARCH-PROMPT-GROK-v0.1.docx; reviewed for this revision.
S4. OAL_LOIE_Independent_Research_Prompt_Copilot_v0.1.pdf; reviewed, including continuation sections recovered by attachment search.
S5. Labyrinth-Independent-Research-Prompt-and-Synthesis.pdf; Gemini document ID OAL-LOIE-PROMPT-GEMINI-v0.1; reviewed.
S6. OAL-LABYRINTH-RESEARCH-PROMPT-CLAUDE_v0.1.pdf; reviewed.
S7. OAL_LOIE_Independent_Research_Prompt_v0.1.pdf; prepared by ChatGPT; reviewed.
S8. Perplexity independent research prompt v0.1, supplied in the preceding conversation generation record; revision baseline.
S9. Nathan's Author's Note on Foundational Architects, November 11, 2026; current user message and new governing research direction.

This revision is prompt preparation using supplied material, not execution of the proposed literature review. No new external research findings are asserted by this document.