Network Defense Intelligence Exercise

Independent Research and Adversarial Feasibility Prompt

Document ID: NDIE-RESEARCH-PROMPT-PERPLEXITY-001

Version: 1.0 | Date: November 4, 2026

Prepared by: Perplexity | Requested by: Nathan Lewis

Status: Draft for cross-AI review. Research not executed. No canon ratification or deployment authorization implied.

1. Assignment and decision to support

Conduct an independent, evidence-led investigation of the Network Defense Intelligence Exercise described in OAL-DOSSIER-CANON-002, supplied in PDF and DOCX form. Determine whether a bounded real-world defensive research program is worth undertaking, under what conditions, at what cost, and with what safeguards. Do not assume the original architecture, national scale, three-second cycle, or claims of universal capture are desirable, achievable, or necessary.

The central research question is: Can authorized, selectively deployed deception and adaptive isolation measurably reduce the probability or impact of compromise of real assets, relative to credible alternatives, against adversaries who may avoid the deception altogether?

Produce a decision-ready assessment, not an endorsement, fictional scene, marketing proposal, or implementation tutorial. A finding of insufficient evidence, limited usefulness, redesign, or cancellation is an acceptable result. Distinguish the value of investigating a question from the value of purchasing or deploying a system.

2. Source dossier and scope discipline

Read both supplied documents. Determine whether they are equivalent exports or contain substantive differences; do not count them as independent corroboration. Record section and page locations where available. Preserve the originals as source artifacts.

Treat narrative canon as a creative constraint only within the fictional work. In the real-world research track, every technical assertion is contestable. Separate three tracks: existing demonstrated capability; plausible but unvalidated research; and speculative fiction. Keep Time Wars, First Memory, purpose-reading, and temporal containment outside the engineering evidence base unless explicitly discussing narrative translation.

Do not silently revise canon or apply ONESTAMP. Propose changes through a separate change-control register. Do not infer approval from this research request.

The dossier reports agreement among multiple AIs. Treat that as a reported synthesis, not independent experimental validation. If the underlying AI submissions are unavailable, say their independence and exact claims cannot be verified.

3. Initial document audit

Create a traceable audit before external research. Classify every major assertion as source claim, narrative requirement, engineering hypothesis, measurable requirement, assumption, contradiction, or unsupported conclusion.

Investigate these apparent issues rather than smoothing them over:

- The approval notice says nothing is fully ratified until ONESTAMP, while later sections call elements fixed or locked canon.

- The document refers to five reviewing AIs, lists a broader distribution including Claude, and proposes roles for six AI systems. Separate reviewers, recipients, and proposed participants.

- The asset framework is described as five-tier, while examples use both Tier 0 and Tier 5; critical assets also receive inconsistent Tier 0/Tier 1 labels. Request the authoritative taxonomy.

- Section 4.4 refers to Q17 for the transfer package, while Q14 addresses transfer and Q17 addresses failure modes.

- The roadmap refers to five goes-wrong scenarios, while Section 9.4 labels four scenarios, A-D. Distinguish them from the five failure modes in Section 9.3.

- The executive summary presents broad agreement on pre-Time Wars placement, while the placement table assigns different primary preferences. Do not exaggerate consensus.

- Separate a disappearing visualization line, termination of a session, containment of a workload, identification of an endpoint, and identification or apprehension of a person.

- Establish whether real assets are already protected before deception begins. If so, identify exactly what additional benefit the deception contributes.

For each issue, give severity, consequence for research, proposed correction, responsible decision-maker, and whether it blocks further work. Verify PDF/DOCX parity rather than assuming it.

4. Purpose, beneficiaries, and counterfactual

State the operational problem, protected mission, users, assets, and harm to avoid. Distinguish detection, delay, observation, attribution support, containment, recovery, and deterrence; they are not interchangeable objectives.

Explain why this exercise should take place, why now if evidence supports urgency, and why a simpler exercise would not answer the same question. Identify who benefits and who bears risk or cost.

State what happens if the project is not funded. Compare the proposal with improved identity security, patching, segmentation, endpoint monitoring, recovery capability, conventional honeypots, static decoys, and incident-response improvements. Include an equal-budget comparison and opportunity cost. Do not assume deception replaces preventive controls.

Define a minimum useful research outcome and a minimum useful deployable capability. Identify which original ambitions are unnecessary for either.

5. Evidence method and independence

Search for supporting, contradicting, null, failed, abandoned, and inconclusive results. Use explicit searches for detection limitations, bypass, fingerprinting, escape, operational overhead, deployment failure, negative results, and cost-effectiveness.

Prioritize primary technical papers and artifacts, official evaluations and after-action reports, standards, government program records, and documented operational experience. Use vendor claims as claims requiring corroboration. Record preprint versus peer-reviewed status, publication date, study setting, sample size, funding/conflicts where disclosed, and availability of code or data.

Build a source ledger: source ID; title; publisher/author; date; URL or persistent identifier; access date; source type; exact supporting location; supported claim; limitation; and independence from other sources. Do not invent quotations, DOIs, experiments, award amounts, or unavailable results.

Use an evidence matrix linking each material claim to support and counterevidence. Distinguish missing evidence from evidence of no effect. Search disagreement should remain visible. Do not pad the bibliography to meet an arbitrary source quota.

For every major conclusion state confidence, uncertainty, and what new evidence would change it. Record search terms, databases, cutoff date, exclusions, access constraints, and unresolved gaps. If search access is unavailable, identify the task as unexecuted rather than supplying remembered findings as verified research.

Each AI should complete and preserve its initial analysis before reading peer conclusions when feasible. Record which peer prompts or outputs were visible. Agreement from shared training, shared sources, or copied prompts is not independent replication. No AI speaks for another system or claims another system reviewed its work without an actual artifact.

6. Feasibility and architecture decomposition

Evaluate the chain separately: observe; classify; map reachable attack paths; prevent access to real assets; select a response; instantiate decoys; redirect eligible sessions; monitor; isolate; recover; and transfer a deployable artifact.

For each step document required privileges, ownership, instrumentation, latency, scale, trust assumptions, failure behavior, and validation evidence. Explain which steps depend on control of endpoints, gateways, identity systems, routing, certificates, or applications.

Assess stateful session redirection, authenticated and encrypted protocols, existing endpoint compromise, application consistency, identity continuity, and external dependencies. Distinguish routing a new connection from transparently moving an established authenticated session. Do not assume complete national visibility or administrative authority.

Compare full duplication with selective emulation, pre-provisioned decoys, procedural generation, digital twins, and replay. Establish where each is credible and where external validation or physical processes expose a mismatch.

Treat three seconds as a candidate experimental parameter, not an established security optimum. Define exactly what rotates: address, route, identity, key, service state, decoy content, or presentation. Examine session breakage, synchronization, caching, resource cost, observability, and the possibility that periodicity becomes a detectable signature. Compare static, periodic, randomized, and risk-triggered policies.

Ask whether AI is necessary. Compare deterministic orchestration, statistical detection, human authorization, and AI-assisted decisions. Assess false positives, prompt injection through attacker-controlled telemetry, model poisoning, model drift, unreliable generated content, and compromise of orchestration infrastructure.

Translate the black-line display into auditable states: suspected; observed; diverted; isolation verified; visibility lost; resolved; and unknown. Do not allow visual disappearance to imply verified safety.

7. Adversarial analysis: bypass the process entirely

Assume an attacker values the real asset, not winning a contest inside the decoy. Construct defensive attack trees around loss of confidentiality, integrity, availability, and physical safety. Include patient actors, insiders, supply-chain adversaries, opportunists, automated agents, and multiple independent campaigns.

At a threat-model level, examine:

- Unmonitored ingress, unmanaged devices, forgotten services, alternate connectivity, third-party access, and assets outside the administrative boundary.

- Abuse of legitimate accounts, sessions, delegated access, trusted administrators, identity providers, and control planes.

- Compromise already present before activation, upstream software or update compromise, firmware dependencies, backup and recovery systems, and management infrastructure.

- Direct access to real data through legitimate workflows, external copies, SaaS integrations, contractors, or compromised human processes.

- Physical access, removable media, maintenance workflows, and operational technology paths that do not traverse the proposed deception layer.

- Capacity exhaustion, distraction, concurrent attacks, deliberately misleading telemetry, and attacks that achieve damage before useful engagement.

- Compromise of sensors, policy engines, virtualization, orchestration, logging, update packaging, or the artifact transferred to a partner.

- Out-of-band comparison of reality with the decoy, avoidance of engagement, abandonment of one session, and re-entry through a different path.

- Innocent users or compromised intermediaries being misclassified, and attackers blending into legitimate activity.

For each category record violated assumption, assets at risk, broad feasibility conditions, expected observables, defensive controls, safe laboratory test, residual risk, and evidence confidence. Distinguish bypass from decoy recognition, containment escape, and ordinary detector failure.

Do not provide exploit chains, payloads, real-target reconnaissance, credential abuse instructions, or unauthorized testing. Use abstract scenarios, synthetic identities, owned infrastructure, and written rules of engagement. This boundary must not soften the assessment of architectural weaknesses.

8. Evolution over time

Analyze the first encounter, repeated encounters, adversary learning, defender updates, operational drift, and retirement. Test whether value declines after disclosure or recognition. Distinguish attacker deterrence from displacement to weaker assets and silent non-engagement.

Develop a versioned adaptation plan covering refreshed decoys, policy updates, independent validation, staged releases, rollback, contamination controls, and periodic reassessment. Do not permit attacker-derived data to automatically modify trusted production policy without review.

Assess recurring cost, configuration drift, staff turnover, ecosystem changes, cryptographic and protocol changes, AI-agent behavior, vendor dependence, and accumulated technical debt. Define indicators that trigger redesign, reduced scope, or shutdown. Identify what remains beneficial if the attacker knows deception exists.

9. Comparable systems, incidents, and simulations

Investigate relevant cyber deception, honeypots/honeynets, honeytokens, moving-target defense, cyber ranges, digital twins, software-defined containment, automated response, and adversary engagement research. Candidate search leads include MITRE Engage and D3FEND, NIST cyber-resiliency guidance, DARPA deception or cyber-defense programs, and documented public or academic cyber ranges; verify actual relevance rather than treating these leads as proof.

For each selected comparison report the actual objective, scale, environment, attacker model, trust boundary, technology, study design, baseline, results, negative results, operational cost, and limitations. Identify whether it was a tabletop, simulation, emulation, controlled human study, competition, prototype, field deployment, or incident response. Do not conflate a program announcement with a demonstrated outcome.

Answer explicitly: Has this exact architecture been modeled or deployed? Which components have? Which combination remains unvalidated? If no exact precedent is found, state the search scope and uncertainty rather than claiming none exists.

For each lesson explain what transfers to this proposal and what does not. Successful decoy engagement does not alone establish reduction of real-asset harm, seamless national-scale migration, or universal containment. Protect confidential study details and do not manufacture case-study failures when reports do not disclose them.

10. Experimental design and falsifiable outcomes

Design a staged, authorized program: evidence review; tabletop and threat model; isolated proof of concept; controlled cyber-range evaluation; limited opt-in pilot; and independent scale review. Production deployment is not an automatic final step.

Start with one bounded asset class and explicit control of the relevant infrastructure. Define topology, protocols, realistic benign workload, synthetic data, ground truth, attacker objectives, defender knowledge, and exclusions. Include scenarios in which attackers bypass or never encounter a decoy.

Compare at minimum: standard controls without deception; static deception; adaptive deception without rapid cycling; and adaptive cycling with tested schedules. Keep underlying access controls and investigator effort comparable. Evaluate whether simpler configurations produce the same benefit.

Specify randomization where feasible, repeated trials, unseen scenarios, learning effects, blinded adjudication where feasible, stopping rules, statistical uncertainty, and power/sample-size justification. Report unsuccessful and inconclusive runs. Prevent training/evaluation leakage and overfitting to the range.

Measure real-asset compromise and mission impact as primary outcomes. Secondary measures may include time to detection/isolation, engagement coverage, visibility gaps, false diversion of benign traffic, verified isolation failures, decoy-recognition rate, service latency, availability, analyst workload, information value, compute/storage consumption, and cost per useful outcome. Give denominators and observation windows.

Define capture operationally. A session is not a person or a complete adversary campaign. Report simulated ground-truth coverage separately from real-world unknown coverage. Zero observed compromises is not proof of zero residual risk.

Require an independent observer for ground truth and a safety authority able to halt the test. Define egress controls, recovery rehearsal, kill switch, evidence preservation, and stop conditions for privacy breach, loss of isolation, unexpected external traffic, or unacceptable mission disruption.

11. Cost and economic feasibility

Provide bottom-up, assumption-driven estimates only after scoping. Separate the cost of research from procurement, deployment, integration, operation, maintenance, and decommissioning. No unsupported national-scale price tag.

Estimate distinct scenarios: desk research/tabletop; isolated prototype; cyber-range evaluation; bounded organizational pilot; and multi-organization research demonstration. Treat national-scale operation as a separate feasibility problem; if not responsibly estimable, explain why and specify the missing inputs.

For each scenario state duration, staffing roles and person-months, loaded labor assumptions, compute, storage and logging, network/egress, software licenses, facilities, OT hardware if needed, integration, independent evaluation, legal/privacy review,

insurance where relevant, training, travel if needed, contingency, and applicable indirect costs. Explain what loaded rates include to prevent double counting.

Show low/base/high estimates in dated USD with arithmetic, assumptions, quantity and unit-rate sources, exclusions, uncertainty, recurring annual cost, and three-to-five-year total cost where meaningful. Distinguish illustrative planning estimates from quotes or awarded budgets. Provide sensitivity analysis for fidelity, telemetry retention, staffing, number of assets, concurrent activity, and rotation frequency.

Estimate resource cost of the simpler alternatives and explain break-even assumptions without treating hypothetical avoided losses as guaranteed savings. Identify which next-stage funds should be contingent on demonstrated results.

12. Funding pathways and readiness

Investigate realistic pathways matched to the applicant: university research partnerships; eligible small-business programs; federal or state cybersecurity research opportunities; sector-specific critical-infrastructure research; national-laboratory or research-center collaboration; industry-sponsored work; and in-kind cyber-range or cloud support. Candidate agencies or program families are leads, not presumed active opportunities.

For each pathway verify current status at research time, official source, applicant eligibility, entity and partnership requirements, scope fit, award range if published, period of performance, deadlines or rolling status, cost sharing, eligible costs, data/security conditions, intellectual-property terms, publication restrictions, and application burden. Distinguish open solicitations, recurring programs without an open call, closed examples, procurement contracts, grants, and informal partnerships.

Do not describe historical funding as available now. Flag unclear status, unavailable amounts, eligibility constraints, registration lead times, and policy changes. Explain whether an individual concept owner would need a university, company, nonprofit, or other eligible lead organization.

Recommend a staged funding strategy tied to evidence gates. Identify a likely research owner, beneficiary, operator partner, independent evaluator, and funder type without inventing commitments. Draft a short neutral concept-note outline and list the evidence required before requesting funds.

13. Governance, safety, legal, and human factors

Identify relevant jurisdiction-specific issues for qualified counsel: written authorization, scope of monitoring, interception and privacy, employee/customer notice, third-party consent, data minimization and retention, evidence handling, cross-border data, procurement, liability, intellectual property, export restrictions where applicable, and critical-infrastructure obligations. Distinguish binding law from guidance and project policy. Do not provide a legal conclusion beyond the available evidence.

Exclude hack-back or changes to third-party systems without separate lawful authorization. Distinguish controlled redirection within owned infrastructure from reaching into an attacker-controlled device. Define human decision rights, emergency access, auditability, separation of duties, and accountability for automated errors.

Assess whether sole visibility and sole authority in one Architect creates an unacceptable operational dependency. Propose least privilege and independent safety oversight without unnecessarily exposing the deception. Consider user confusion, alert fatigue, unequal impacts, false accusation, and a partner's ability to inspect, update, and safely operate the transferred artifact.

14. Changes required before proceeding

Produce a ranked change register containing original wording/section, issue, evidence, proposed real-world replacement, canon impact, owner, dependency, and disposition. Preserve creative material separately.

Specifically evaluate replacing universal capture with bounded measured outcomes; purpose-reading with observable indicators and uncertainty; complete national replication with scoped selective deception; a fixed three-second rule with a tested scheduling hypothesis; visual disappearance with verified state transitions; and an automatic reboot transfer with signed, inspected, versioned deployment plus rollback and receiving-operator acceptance tests.

Do not assume these replacements are approved or experimentally validated. Explain which are proposed research-design corrections and which require findings from later tests. Retain useful concepts while removing claims that cannot be evaluated.

15. Questions, decisions, and action plan

Produce a prioritized question registry. Include protected mission and asset, authorized boundary, research-versus-fiction intent, institutional owner, data availability, acceptable disruption, required observer, test partner, applicant eligibility, budget ceiling, timeline, intellectual-property goals, repository visibility, and the decision that each answer enables.

Separate questions that block safe testing from questions that only refine cost or scope. Do not block desk research on narrative decisions. If information is missing, use labeled scenario assumptions rather than silently choosing an answer.

Provide owner roles, dependencies, estimated effort, required evidence, and completion criteria for immediate actions and a conditional 30/60/90-day plan. Do not present the schedule as a commitment. Define go/revise/pause/stop gates, including no measurable benefit over baseline, unacceptable benign-user harm, inadequate isolation, unaffordable operations, or lack of legal authority.

Include the strongest argument against funding, strongest argument for a limited study, most consequential unknown, cheapest discriminating experiment, and what would change the recommendation.

16. Required final research report structure

Deliver: (1) executive decision summary; (2) source-document audit; (3) goal and counterfactual; (4) claims/evidence matrix; (5) feasibility decomposition; (6) bypass-focused threat model; (7) comparable-case table; (8) longitudinal adaptation analysis; (9) experimental protocol and metrics; (10) itemized cost workbook and narrative; (11) verified funding-pathway matrix; (12) governance and safety plan; (13) prioritized questions/comments/concerns; (14) change register; (15) conditional action plan and decision gates; (16) limitations, search log, and source ledger.

Use concise executive summaries supported by detailed appendices. Clearly distinguish findings, author assertions, estimates, hypotheses, and recommendations. Cite material factual claims near the relevant text and preserve usable source details in exported documents. Include disagreement and negative evidence in the main assessment, not only appendices.

No predetermined feasibility verdict, monetary estimate, or funding recommendation is required at prompt stage. The executed report must answer every requested question or state why it cannot yet be answered.

17. Deliverables and proposed GitHub organization

Provide a complete PDF and editable DOCX with matching substantive content, date, version, document ID, status, headings, tables, citations, page numbers, and change history. Verify both files open and are non-empty. Add source Markdown, cost CSV, claims CSV, and structured research registers where feasible; label all estimates.

Propose, but do not create remotely without authorization, a dedicated network-defense-intelligence-exercise folder with: README.md; prompts/; source-dossier/; research/; evidence/; threat-model/; experiments/; costs/; funding/; governance/; decisions/; and deliverables/. Preserve each AI's original prompt and independent first-pass report in separate versioned files. Store subsequent synthesis separately, not over the originals.

Use stable IDs for claims, risks, questions, experiments, sources, decisions, and changes. Include reproducibility instructions, dependencies, synthetic-data provenance, sanitized results, and a decision log. Record software versions and seeds where applicable.

Do not publish credentials, real topology, personal data, confidential telemetry, sensitive third-party details, or harmful operational artifacts. Define a public/private split, license and ownership review, contribution process, disclosure contact, and release checks. No external repository writes, uploads, testing, or purchases are authorized by this prompt.

18. Cross-AI prompt review and execution sequence

Stage A: Submit this v1.0 prompt only. Do not execute the research yet.

Stage B: Read the other AI-generated prompts supplied by Nathan. Produce a comparison identifying omissions, stronger methods, disagreements, duplicated work, safety gaps, and potential groupthink. Accept or reject proposed changes with reasons and a visible revision log.

Stage C: If changes are useful, generate v1.1 while preserving v1.0; otherwise state that no revision is warranted. Keep a common core for comparability, with model-specific research emphasis clearly labeled.

Stage D: Execute only when Nathan explicitly requests execution. Disclose available tools and limitations, carry out the evidence search, and produce the research report and associated files. Do not imply that sharing a prompt authorizes deployment or ratifies canon.

Stage E: Compare completed independent reports without majority-vote truth claims. Resolve factual disputes through source verification or additional tests, and preserve unresolved disagreements.

Source basis for this prompt

The user-supplied PDF and DOCX of OAL-DOSSIER-CANON-002, version 1.0.0, dated April 11, 2026, particularly Sections 2, 4, 6, 8, 9, 10, and 12; and Nathan Lewis's development and research request of November 4, 2026. These are design inputs, not external validation of feasibility. External research has intentionally not been executed in this prompt-development stage.