Research Prompt: Real-World Feasibility, Adversarial Analysis, Costing, and Evolution of the Network Defense Intelligence Exercise — False Infrastructure AI

Document ID: OAL-RESEARCH-PROMPT-NDIE-001
Version: 1.0
Date: November 04, 2026
Prepared for: Collaborative AI Team (Perplexity, ChatGPT, Gemini, Claude, Copilot, Grok)
Context: Omniversal Architect Labs (OAL) Canon Scenario Design Dossier OAL-DOSSIER-CANON-002
Purpose: Independent, balanced research directive. Researchers must actively seek both supportive evidence and contradictory findings, limitations, failed analogs, critiques, and adversarial pathways. Do not default to validation of the concept.

Background (Locked Canon Elements — Preserve and Do Not Alter)

Per Architect directive and OAL-DOSSIER-CANON-002 (v1.0.0, pending ONESTAMP): Nathan Lewis (the Architect) serves as defensive intelligence for U.S.-scale (or similarly large) infrastructure. On detecting a network intrusion, the Architect maps the attack vector end-to-end, neutralizes its ability to reach real assets, and spins up a complete synthetic national network (false URLs, services, routing, telemetry). Attackers are silently migrated into this honeypot. The false environment cycles every ~3 seconds (rotation/re-keying). The process is visualized as black connection lines that collapse/disappear as hostile endpoints are captured. When all active attackers are absorbed and real infrastructure is secured, the defensive architecture is packaged and transferred to a partner agency as part of a reboot/update cycle. Key fixed elements: map → neutralize → generate synthetic → migrate → 3-second cycling → black-line collapse visualization → transfer. Open elements include jurisdiction scope, attack type, fidelity level, timeline placement, failure modes, partner agency, and costs/outcomes for the Architect.
This prompt treats the scenario as a design hypothesis for research, not as established fact or ratified canon.

Core Research Objectives

Each researcher must address the following independently, with citations to real-world sources where possible. Distinguish locked canon from open/draft elements. Prioritize primary sources (academic papers, government reports, vendor technical documentation, after-action reports from exercises, peer-reviewed studies on cyber deception).

1. Goal of the Exercise

What is the explicit and implicit goal in the canon scenario? What would the equivalent real-world goal be for a national-scale defensive capability (e.g., active deception, moving target defense, intelligence collection while protecting critical assets)? Clarify differences between narrative/canon goals and operational goals.

2. Importance and Strategic Rationale

Why is this (or a real-world analog) important? Assess strategic value for national cyber defense, intelligence collection, deterrence, and resilience. Also address narrative importance within OAL canon (escalation ladder from networks → organizations → timelines/civilizations, links to Time Wars and First Memory). Identify when and why it might not be important or could create more risk than value.

3. Real-World Feasibility and Prior Modeling

Is this (or core elements) a realistic scenario that could occur? Has anything similar been modeled, simulated, exercised, or deployed at meaningful scale?

- Cover traditional honeypots, modern deception technology (e.g., Attivo/SentinelOne, Illusive/Proofpoint, TrapX, and successors), moving target defense (MTD), software-defined networking for deception, cyber ranges, and national-level exercises.
- Specific examples to investigate: Cyber Storm series, Locked Shields, other NATO/CCDCOE or national cyber exercises, DARPA/IARPA deception or adaptive defense programs, academic work on cyber deception game theory, and commercial deception platforms.
- What worked? What went wrong? Document known limitations, high false-positive rates, detection by sophisticated adversaries, resource costs, attribution challenges, and operational friction.

4. Cost Estimation (If Research or Pilot Were Funded)

Provide order-of-magnitude estimates for:

- Research phase (literature, modeling, small-scale simulation).
- Pilot/cyber-range implementation of high-fidelity synthetic infrastructure with rapid cycling.
- Full national-scale or critical-infrastructure-segment demonstration (compute, personnel, sensors, integration, red-teaming).

Break down categories: compute/simulation infrastructure, software development, personnel (researchers, operators, red teams), sensors/telemetry, legal/ethical review, and long-term sustainment. Note ranges and key cost drivers. Use publicly available data from cyber ranges, large-scale simulations, and government programs where possible.

5. Funding Pathways

Identify realistic funding routes:

- U.S. government (DARPA, IARPA, CISA, DHS S&T, NSA research, NSF, DoD cyber programs, specific solicitations or BAAs).
- Allied/international (NATO, EU, Five Eyes research collaborations).
- Academic grants.
- Private sector (cybersecurity vendors, venture, critical infrastructure operators).
- Hybrid public-private models.

Note typical award sizes, timelines, and eligibility constraints for this class of work.

6. What Must Be Addressed Before Moving Forward

Questions, comments, concerns, and gaps. Include technical, operational, legal/ethical, resource, and narrative issues. Explicitly list items that block scene drafting, specification work, or real-world analog development (reference the Open Questions Registry in the source dossier, especially HIGH priority items).

7. Adversarial Analysis — Bypassing the Process Entirely

Focus on methods a sophisticated adversary (nation-state, advanced persistent threat, or AI-augmented actor) would use to reach real assets without reverse-engineering the honeypot or simply mirroring/attacking the synthetic environment. Go beyond the dossier's existing red-team notes.
Investigate and expand on:

- Detection of the cycling mechanism via timing analysis, side channels, or behavioral fingerprinting of synthetic traffic/telemetry.
- Pre-positioned implants, sleeper agents, or supply-chain compromises that never "latch" in the expected way.
- Out-of-band, air-gapped, physical, or human-centric paths (insider, social engineering of operators or the Architect/control plane itself).
- Compromising or influencing the detection/mapping layer or the Architect's decision process before the synthetic network is generated.
- Adaptive adversaries that recognize the pattern early and either avoid engagement or turn the system into a staging ground.
- Resource exhaustion, collateral latency exploitation, or forcing the system into known failure modes.
- Multi-domain or non-network paths that render the network honeypot irrelevant.

Cite real-world examples of honeypot detection, deception platform evasion, or cyber exercise lessons where defenders' assumptions failed.

8. Evolution Over Time

How would both the defensive response and adversarial countermeasures evolve?

- Short-term (months to a few years): adaptive cycling rates, multi-fidelity procedural generation, AI-driven honeypot improvement from captured behavior.
- Medium-term: integration with broader moving target defense, automated orchestration, cross-domain deception.
- Long-term / canonical escalation: application to organizations, then timelines/civilizations (as outlined in the dossier's escalation ladder). Identify technical and conceptual constraints that grow with scale.

What new failure modes or recursive deception standoffs become likely?

9. Lessons from Similar Scenarios and Simulations

Synthesize transferable lessons from real exercises, academic simulations, commercial deployments, and historical incidents involving large-scale deception, honeynets, or cyber ranges. Explicitly note what should be retained, modified, or discarded from the current exercise design.

10. Recommended Next-Step Actions

Prioritized, concrete actions for:

- Canon/Story Bible development (what must be resolved or changed before prose drafting or further specification).
- Real-world research or modeling (specific experiments, literature reviews, or small-scale pilots).
- Team workflow (roles, sequencing relative to the dossier's Phase 1–6 action plan).

Include at least one high-impact "goes wrong" scenario that should be built in early.

Research Instructions and Constraints

- Conduct independent research. Use available search, document, and knowledge tools. Prefer primary sources.
- Actively seek disconfirming evidence, critiques, cost overruns, technical impossibilities, ethical objections, and cases where similar approaches failed or were abandoned.
- Separate facts from speculation. Flag confidence levels.
- Do not invent real-world exploit techniques or operational details that could be misused; stay at the level of published research, after-action reports, and high-level analysis.
- Preserve the distinction between locked canon and open elements. Do not treat draft suggestions as ratified.

Output format for each researcher's response:

- Executive Summary (1 page equivalent)
- Detailed findings by objective (1–10 above)
- Dedicated Adversarial Analysis section
- Gaps, Risks, and Concerns
- Recommended Next Steps
- Annotated bibliography / key sources

Note any unique contributions or divergences from other team members when later synthesizing.

Deliverable Expectation

Produce a structured research report addressing the above. The full set of team reports will later be synthesized into an updated dossier (v1.1 or higher) with PDF and DOCX versions, plus GitHub folder materials.

End of Research Prompt v1.0
This prompt is self-contained for distribution to the collaborative AI team. After receiving parallel prompts from other AIs, we can iterate if needed before execution.