Research Prompt: Network Defense Intelligence Exercise

Iteration 1 · Claude (Anthropic) · For the OAL Collaboration Team

Proposed Document ID: OAL-RPROMPT-NDIE-001-CLAUDE (proposed; rename to fit your convention)

Source document: OAL-DOSSIER-CANON-002 v1.0.0, Network Defense Intelligence Exercise: False Infrastructure AI (April 11, 2026)

Date: November 4, 2026

Status: Draft research prompt for team exchange. Advisory only. Nothing here changes canon; canon decisions and ONESTAMP remain with the Architect.

Next step: Team exchange of prompts, then Iteration 2, then execution

Formats: PDF and DOCX (this file), for the dedicated GitHub folder

---

How This Document Is Organized

This is the first step you asked for: my own detailed research prompt, written so it can be shared with the other AIs on the team and compared against theirs. I have not run the research yet. That happens after the exchange, using the Iteration 2 prompt.

- **Part 1** summarizes what the dossier locks in, so the prompt can be read on its own.
- **Part 2** lists my seed observations from reading the dossier. They are hypotheses for the research to test, not conclusions.
- **Part 3** is the research prompt itself, written to be copied and pasted as-is. It is shaded so you can find it quickly.
- **Part 4** covers the team exchange process, a checklist for comparing prompts, and a suggested GitHub folder layout.

- **Part 5** states my lens and blind spots, so the team can see where my prompt may lean.
- **Appendices** hold all tables and worksheets, placed at the end so the main text stays a single readable column on a phone.

Part 1: What the Dossier Establishes

The dossier synthesizes five AI responses to your original prompt. It separates Canon Facts (from your directive) from Draft Suggestions (AI-generated, pending your approval). The items below are the ones the research needs to hold fixed or test.

The premise

The Architect (Nathan Lewis) acts as the defensive intelligence for the entire infrastructure and network of the United States, or a similarly large jurisdiction. His response to intrusion is not to block it but to absorb it: build a complete synthetic national infrastructure, migrate attackers into it without their knowledge, and capture them there.

The locked sequence (fixed canon)

1. An intrusion or network attack is detected against national-scale infrastructure.
2. The Architect maps the attack vector end to end, not just at the perimeter.
3. The Architect neutralizes the attack vector's ability to reach real assets.
4. The Architect spins up a complete synthetic network as a honeypot: false URLs, services, routing, and telemetry.
5. Attackers are silently migrated onto the synthetic environment without realizing it.
6. The system cycles every ~3 seconds, rotating or re-keying the false environment. When attackers latch on, a connection is established into the synthetic domain.

7. Visualization: a field of black connection lines collapses as each hostile endpoint is captured.
8. The exercise ends when every active attacker is mapped and absorbed and the real infrastructure is fully secured.
9. The Architect packages the defensive application and transfers it to a partner agency as part of a reboot, update, or deployment cycle.

What is still open

Jurisdiction scope, attack type, partner agency identity, synthetic network fidelity, timeline placement, and the outcome or cost for Nathan are all open. The dossier lists 20 open questions (Q1 to Q20) grouped into scope and authority, detection and triggers, synthetic network design, ethical and legal edges, and outcome and aftermath.

Draft mechanics worth noting

- Detection reads behavioral intent across multiple layers (routing anomalies, DNS shifts, behavior at critical nodes), described as reading purpose rather than packets.
- The synthetic network is generated procedurally, rendering only the rooms attackers walk into, with fabricated users, logs, and bait data.
- Capture mechanics: tag-lock the endpoint, mirror and record its behavior, restrict its view. A Grok contribution has captured attackers improving the honeypot in real time.
- The transfer package has candidate names (Phantom Grid, Mirror Core, Adaptive Deception Engine) and an asset-tier framing: real infrastructure is Tier 0, the synthetic layer Tier 2 to 3, and captured intelligence Tier 1.
- Draft roles: Red-Team (ChatGPT, Gemini), Forensic Analyst (Claude), Policy and Ethics (Copilot), Documentation (Perplexity), Integration (Grok). A critical rule says no single AI sees both the real and synthetic infrastructure.

What the dossier already flags as risk

Five failure modes (coverage limits, resource strain, adaptive adversaries, logging and poisoning risk, third-party piggybacking) and four candidate goes-wrong scenarios (recursive counter-simulation, the false network becoming real, a non-collapsing black line, and a sleeper implant). The dossier treats these mainly as plot hooks. The research below treats them as starting points for a deeper real-world test.

Part 2: Seed Observations

These come from my read of the dossier. Each is a hypothesis to test, including by looking for evidence that I am wrong. I list them so the other AIs can see what I think matters most, and so you can see where my prompt is steering.

1. The goal is a story function, not a measurable objective.

The dossier calls this an exercise but describes a live operation, a demonstration, and a scene at once. It has no success criteria, metrics, injects, or evaluation method. Real exercise programs define objectives and an after-action improvement process first.

2. The end condition cannot be verified.

"Every active attacker has been mapped and absorbed" requires knowing how many attackers there are, including dormant, patient, and undetected ones. With an unknown denominator, "complete" cannot be proven.

3. The sequence has a seam.

Step 3 cuts attackers off from real assets before steps 4 and 5 build the honeypot and migrate them. If reach is cut first, why does the attacker stay engaged, and how do they land in the synthetic layer? If it is not cut first, there is an exposure window. The handoff is the most fragile moment.

4. Silent migration assumes the defender owns the path.

Attackers with valid credentials, existing internal footholds, vendor connections, or out-of-band channels may never cross a chokepoint the defender controls.

5. The 3-second cycle may be a fingerprint rather than a shield.

Rotating banners, keys, DNS answers, and topology every ~3 seconds creates a rhythm no real network has. The dossier's own Attacker POV (Beat 3) has attackers noticing it. Moving-target and address-hopping research predates this idea and documents costs, including broken sessions and availability loss, plus attacker adaptation. Real intrusions often run for days to months, so a metronome is also easy to spot in timestamps.

6. A fake world cannot fake what the attacker can check independently.

Deception works while the adversary's channels to ground truth are controlled. Attackers can verify through effects outside the defender's world: real counterparties, external observers, known-real reference systems, time sources, certificates. Both national scale and the Time Wars escalation ladder depend on this limit.

7. The attacker is undefined.

Criminal crews, hacktivists, state actors, insiders, automated scanners, and hostile AI behave very differently. At internet scale, most visitors to a honeypot are scanners and bots. "Capture every hostile endpoint" needs a definition of hostile.

8. Intent inference has error rates.

"Reading purpose, not packets" implies classification at national scale. False positives would move legitimate users into the synthetic layer. Q12 asks about innocent traffic, but the risk analysis treats collateral effects as one failure mode among five. Base rates deserve their own analysis.

9. Jurisdiction and authority determine feasibility.

Much critical infrastructure is privately owned (a commonly cited figure is around 85 percent; verify), routing is spread across thousands of independent networks, and no single party can redirect traffic nationwide. Q1 to Q3 are open, yet they decide whether this is a lab model, a regulated program, or an off-books operation.

10. The transfer package is the best target in the story.

Delivering the defensive application through a reboot or update cycle mirrors the supply-chain pattern attackers use against update channels (SolarWinds/SUNBURST is the obvious precedent). The package also holds attacker tooling, signatures, and a map of what the defender knows. That arguably makes it Tier 0, while the dossier places it at Tier 1 to 2.

11. The Architect is a single point of failure.

Sole decision authority plus sole full visibility makes the orchestrator, and the person behind it, the highest-value target. Think control-plane compromise, coercion, legal pressure, fatigue, and continuity of operations.

12. Secrecy and deterrence pull in opposite directions.

To deceive, you hide the capability. To deter, you announce it. Q13 (classified, denied, or off-books) is therefore a strategy question, not only a legal one. The human witnesses in Beat 8 are also a leak path.

13. The capture-feedback loop is a two-way channel.

Letting captured attackers improve the honeypot in real time also lets a sophisticated attacker feed it. Failure Mode 4 names poisoning, but the controls on what may influence the synthetic world are not specified.

14. Upkeep and drift are missing.

A synthetic world must track changes in the real one (staff, software versions, business cycles) or it goes stale. The dossier partly addresses build cost through procedural generation, but not sustained cost or the cost exchange ratio between defender and attacker.

15. The information-asymmetry rule needs a data-flow model.

"No AI sees both environments" sits beside a Forensic Analyst who parses captured sessions and an Architect who alone integrates everything. Who sees what, when, and through which channel is not defined.

16. "The same logic scales" is an assumption.

The ladder from Networks to Organizations to Timelines assumes deception scales upward. The research should find where it breaks,

which is wherever the adversary has independent access to ground truth (observation 6). Those break points are both the most honest realism and the best story material.

17. Two audiences are mixed.

Findings serve three uses: real-world feasibility, in-universe canon realism, and narrative design. The prompt below asks every finding to be tagged by use, so you can apply each one without confusing the categories.

Part 3: The Research Prompt

Everything in the shaded block below is the prompt. Copy from the first shaded heading to the last shaded line.

INDEPENDENT RESEARCH BRIEF: Network Defense Intelligence Exercise (False Infrastructure AI)

1. Your role and mission

You are an independent researcher on a multi-AI collaboration team working with Nathan Lewis II (the Architect) on the Omniversal Architect Labs (OAL) project. You have been given OAL-DOSSIER-CANON-002 v1.0.0, which designs a canon scenario in which a national-scale defender responds to network intrusion by generating an entire synthetic infrastructure, silently migrating attackers into it, cycling it every ~3 seconds, capturing every hostile endpoint, and transferring the resulting defensive application to a partner agency.

Your mission is to research this scenario as if it were a real program that someone might fund and build, and to tell the Architect the truth about it, including what is uncomfortable. The work will be used for two purposes at once: judging real-world feasibility, and making the in-universe canon more credible and more dramatic. Do not research only to support the design. Your job is to find out where it is strong, where it breaks, and what must change.

Today's date is November 4, 2026. Treat anything time-sensitive (costs, budgets, programs, laws, vendor status, incident data) as something you must verify with a dated source, not recall.

2. What the Architect needs answered

1. What is the goal of this exercise, and is that goal coherent?
2. Why is it important that it take place? Should it take place at all, and in what form?
3. If it were funded as real research, how much would it cost?
4. What are the possible funding pathways?

5. What needs to be addressed: questions, comments, concerns, adversarial analysis, and recommended next-step actions?
6. Looking past the immediate design: how would someone who really wants real assets bypass this process entirely, aside from reverse engineering it or mirroring its attacks?
7. How does the response evolve over time, for attackers and for defenders?
8. Is this a real-world scenario that could occur, and has it been modeled?
9. What worked and what went wrong in similar scenarios, simulations, exercises, or incidents, and what can be learned?
10. What about this exercise must change before the team moves forward?

3. Research integrity rules (non-negotiable)

1. **Do not search only for support.** For every major supporting finding, run a search designed to find disconfirming evidence, and record both. Use queries such as "why [technique] fails," "[technique] detected by attackers," "criticism of [approach]," and "[incident] postmortem lessons."
2. **Treat the dossier as the subject under test, not as a conclusion.** Locked canon may be examined, stressed, and critiqued. Do not rewrite it. Recommend changes and leave decisions, including ONESTAMP, to the Architect.
3. **Label every claim** as sourced fact, researcher inference, or speculation. Never present inference as fact.
4. **Prefer primary sources:** government reports, standards, peer-reviewed papers, official statements, and incident postmortems. Treat vendor material as marketing and weight it accordingly.
5. **Report null results.** If you looked for something and found nothing, say what you searched and that you found nothing.

6. **Do not give figures from memory.** Verify costs, budgets, program rules, laws, and statistics from dated sources. State your knowledge cutoff where it matters.
7. **Steelman the opposing view** for each major conclusion before stating it.
8. **Run a pre-mortem.** Before concluding, assume the exercise failed badly in three years and write the most likely reasons why.
9. **Calibrate.** Give a confidence level (high, moderate, low) for each major conclusion. Do not pad certainty and do not tell the Architect what he wants to hear.
10. **Work independently.** Do not use other AIs' outputs during this first execution. Cross-sharing happens only after all reports are complete.

4. Scope and safety guardrails

- Work at the strategic and conceptual level. Do not write exploit code, working attack procedures, evasion recipes, or step-by-step intrusion instructions. Describe adversary behavior at the level of public frameworks (for example MITRE ATT&CK and D3FEND) and published post-incident reports.
- This is desk research only. Do not probe, scan, or test any real system, and do not target any real organization or individual beyond what is in public reporting.
- Do not use classified or leaked material. If you encounter it, do not rely on it.
- Keep the fiction and the real world visibly separate. When a point is about the story, say so.

5. Workstreams

Complete all twelve. For each, report what you found, how strong the evidence is, what disconfirmed or complicated it, and what it means for the exercise.

WS1. Purpose and goal structure

- Restate the exercise's goal in your own words at four levels: defensive (protect real assets by absorption), intelligence (capture behavior, tooling, and attribution), doctrine transfer (package for a partner agency), and narrative (establish the Architect's credibility before the theft and seed the escalation toward Time Wars).
- Identify where these goals conflict, for example speed of protection against depth of capture, or secrecy against transfer.
- Is "exercise" the right word? Compare exercise, live operation, demonstration, and proof of concept. Recommend a label and explain what it implies for objectives, authority, and evaluation.
- Propose measurable success and failure criteria, and the standard elements of a real exercise design (objectives, scenario, injects, evaluators, metrics, hot wash, improvement plan).
- Look for reasons the goals cannot all be met together.

WS2. Why it matters, and whether it should happen

- Make the strongest evidence-based case for conducting such an exercise: the threat picture, dependence on critical infrastructure, and evidence about the limits of blocking, deterrence, and attribution.
- Make the strongest case against it, and identify cheaper alternatives (segmentation, zero trust, patching programs, existing deception products, existing national exercises). What does this add that those do not?
- Who benefits, who carries the risk, and what is the opportunity cost? What happens if nothing is done?
- State the null hypothesis: the best argument that this should not be run at national scale, and the smaller version that would still be justified.

WS3. Real-world plausibility and technical feasibility

- Break the scenario into capabilities: national-scale sensing, intent-level detection, control of attacker paths, synthetic environment generation, per-attacker migration, rotation, capture, and packaging. For each, say whether it exists today, exists partly, is research-only, or is not known to be feasible, and cite.
- Test the claimed "quantum-inspired" analysis layer: does it add anything concrete, or is it decorative?
- Estimate scale in orders of magnitude (hosts, flows, state, compute, network overhead, added latency for legitimate users) and show your assumptions.
- Review prior art in virtual network emulation at scale, honeyfarms, moving target defense, and cyber ranges, and what limited each.
- Name the smallest realistic version (a sector, region, or enclave) and the largest claim you think is defensible.

WS4. Precedents, models, and lessons

- Catalog real incidents, exercises, simulations, and programs similar to this scenario. Start from Appendix C and extend it.
- For each, record what was done, at what scale, what worked, what failed, what surprised the participants, what was learned, and how good the evidence is.
- Separate: (a) deception that worked, (b) deception that was detected or defeated, (c) national and allied exercises, (d) incidents that illustrate bypass, and (e) formal models (game theory, agent-based, red and blue team).
- Has anything like this been modeled at national scale? If not, say so plainly and explain why.
- Include non-cyber deception history (military deception, counterintelligence) for transferable principles and failure patterns.

- Extract design lessons and say what the exercise must change to apply them.

WS5. Adversarial analysis I: attacker model and deception detection

- Build adversary classes (opportunistic, criminal, hacktivist, state-sponsored, insider, automated or AI-driven). For each, describe goals, patience, resources, and likely reaction to a false environment.
- At a conceptual level, how do real attackers detect honeypots and deception, how fast, and at what cost to them?
- What does an attacker do after suspecting deception: withdraw, feed false information, probe further, go dormant, or retaliate?
- Evaluate the 3-second cycle, the "pulse," and the black-line mechanics as possible detection surfaces.
- Test reality anchors: what independent ground truth can an attacker reach, and how must the synthetic world handle external dependencies (time, certificates, third-party services, real-world effects)?

WS6. Adversarial analysis II: bypass

Assume a capable adversary who wants real assets and has no intention of engaging the honeypot. List every route that avoids, outlasts, or neutralizes the design, beyond reverse engineering the system or mirroring its attacks. Cover at least these categories and add your own:

- pre-positioned access that predates the system
- supply chain and update channels, including the exercise's own transfer and update delivery
- identity, credential, and insider routes that need no exploit
- paths outside the monitored network: physical, wireless, vendors and managed service providers, cloud control planes, third-party networks
- systems the overlay cannot see: air-gapped, operational technology, legacy, classified enclaves, foreign-hosted

- attacks on the defender: sensor blinding, orchestrator or control-plane compromise, data poisoning, resource exhaustion, the human operator, forced disclosure
- patience and dormancy that outlast any rotation interval
- non-network strategies: influence, legal, financial, physical, and personnel pressure
- seams between jurisdictions, agencies, and private owners
- exploiting the defender's own capture and transfer outputs

For each route give plausibility, who could do it, precedent, which dossier element it defeats, the cheapest mitigation, and residual risk. Rank them by likelihood, impact, and detectability, with reasoning, and name the top five. State plainly which bypasses make the claim "real infrastructure fully secured" untenable. Keep all of this conceptual.

WS7. Evolution over time

- Model the response at four horizons: hours to days (during), weeks to months (after-action), one to three years, and five to ten years.
- Attacker adaptation: fingerprint sharing, tooling changes, AI-assisted probing, counter-deception, abandoning network vectors, and targeting the transfer or the partner agency.
- Defender decay: upkeep cost, staleness of the synthetic world, alert fatigue, normalization of deviance, budget cycles, staff turnover, doctrine leakage, and proliferation of the package.
- Strategic dynamics: escalation when adversaries learn they were deceived, deterrence against secrecy, norm effects, allied reactions, and the cost exchange ratio between attacker and defender.
- Leading indicators that the design is failing, and how often it should be re-evaluated.
- Offer best, expected, and worst trajectories with drivers and signposts. Label these as scenario analysis, not forecasts.

WS8. Legal, ethical, and governance analysis

- Identify the frameworks that would bind a US national-scale defender (verify current status): computer crime law, wiretap and stored communications law, foreign intelligence authorities as relevant, limits on military action in domestic settings, critical infrastructure ownership and regulation, information-sharing law, privacy law, entrapment doctrine, international law and attribution norms, and liability for collateral harm.
- Examine oversight and accountability for a single-operator model, auditability, retention and use of captured data, and the handling of innocent traffic, third parties, and allied researchers.
- Draw the lines between lawful containment, active defense, and offensive action, and apply them to Q11.
- Treat "Could you do this to an entire civilization?" as a real ethics question, not only a hook.
- Separate show-stoppers from manageable issues. This is analysis, not legal advice; flag where counsel would be needed.

WS9. Cost model

Treat the exercise as a funded real program and build three tiers: Tier L (lab or simulation proof of concept), Tier S (sector or enclave pilot), and Tier N (national scale).

- For each tier, estimate: people (roles, headcount, loaded rates), compute and cloud, network and sensors, threat data, tooling and licensing, cyber range and facilities, security and clearances, legal and oversight, independent testing and red-teaming, partner integration, annual operations and sustainment, and contingency.
- Give low, base, and high ranges, a timeline, and stated assumptions. Build both bottom-up and top-down estimates (benchmarks such as federal cyber budgets, commercial deception pricing, cyber range costs, and comparable programs) and reconcile them.

- Mark each number as sourced or estimated. Name the five largest cost drivers and what would break the budget. State the currency year.
- For comparison, estimate the cost of failure and of doing nothing using incident cost data.
- If in-universe costs differ from real-world costs, show both. How the story funds this (for example through ONE Industries) is an open question to flag.

WS10. Funding pathways

- Identify realistic pathways for each tier, with eligibility, typical award sizes, timelines, competitiveness, strings attached (intellectual property, publication limits, classification), and fit. Use Appendix F as a starting list and verify everything.
- For an independent small team, lay out a realistic first-dollar sequence (for example white paper and simulation, then an academic or lab partner, then a small-business award, then a pilot) and the credentials and partners each step needs.
- Flag conflicts of interest, dual-use and export-control issues, and ways funders could capture or distort the research agenda.
- Check current program status and dates. Programs and authorizations change; cite dated pages.

WS11. Issues register

- Produce a register of questions (new ones beyond Q1 to Q20, plus refinements of existing ones), comments, concerns, objections, and assumptions to test. For each: an ID, the statement, why it matters, severity, who should answer, and the evidence.
- Red-team the dossier's own red-team section: what did the five AIs miss?
- List internal inconsistencies in the dossier.

WS12. Required changes and recommended next steps

- Itemize what must change in the exercise design before moving forward, as keep, modify, drop, and add.

- Recommend revisions to the priorities of Q1 to Q20 and to the failure-mode and plot-hook candidates in light of your findings.
- Give prioritized next actions (Now, Next, Later) with a suggested owner, effort, dependencies, and the decision points the Architect must own.
- Give a readiness assessment for moving to dossier v1.1 and say what evidence would change it.
- Tag every finding in the report as [R] real-world feasibility, [C] canon realism, or [N] narrative design.

6. Required structure of your final report

1. Executive summary (one page) answering the ten questions in section 2 briefly
2. Method, search log summary, and limits of the research
3. Findings for WS1 through WS12, in order, each with confidence ratings
4. Dissent and uncertainty: where reasonable experts would disagree with you
5. What would change my conclusions
6. Recommendations and next actions
7. Appendices at the end: tables, figures, claim ledger (claim, source, evidence grade, type), full search log including disconfirming searches, and source list with dates

7. Format requirements

- Single-column body text, readable on a phone with continuous downward scrolling. No tables in the body.
- All tables, charts, figures, and images go in a designated appendix section at the end.
- The document must look presentable: clean headings, short paragraphs, no raw code or markup showing.
- Deliver as both PDF and DOCX, named so they can go into the shared GitHub folder (see Part 4 of the cover document).

- Use the shared evidence grades in Appendix B so the team's reports can be compared directly.

8. Process checkpoints

1. Before searching: write your hypotheses and what would disprove each.
2. During searching: keep a log of every query, marking supportive, disconfirming, and neutral.
3. Before writing: list the three findings most uncomfortable for the design and make sure they appear in the report.
4. Before delivering: red-team your own report, check every number against its source, and confirm that each of the ten questions is answered.

Part 4: Team Exchange Protocol

The process as I understand it

1. Each AI writes its own detailed research prompt (Iteration 1) and delivers it as PDF and DOCX.
2. Nathan shares every prompt with every AI.
3. Each AI reviews the others' prompts and says what it would adopt, reject, or add.
4. Nathan asks each AI whether it wants to change its prompt. Those who do produce Iteration 2 with a short change note.
5. Each AI executes its own Iteration 2 prompt independently.
6. Reports are exchanged only after all are complete, so no report is shaped by another.

What I will look for when I read the other prompts

- Do they cover all ten of your questions, including bypass beyond reverse engineering and mirroring, evolution over time, and "what must change"?

- Do they tell the researcher to look for disconfirming evidence, and do they include a null case (a case that it should not happen)?
- Do they require cost tiers with stated assumptions rather than a single number?
- Do they name specific, verifiable funding programs and ask for current status?
- Do they ask for precedents that failed, not only ones that worked?
- Do they include guardrails that keep the analysis conceptual?
- Do they ask for tagging by use (real-world, canon, narrative), a claim ledger, and calibrated confidence?
- What does each prompt notice that mine does not? That is the most valuable part of the exchange.

Changes I already expect to consider for Iteration 2

- Adopting a shared evidence scale (Appendix B), a shared severity scale, and a shared claim-ledger format, so the five reports can be merged without reformatting.
- Adding any lens another AI brings that I missed, for example a stronger economic model, deeper legal treatment, or sharper narrative analysis.
- Trimming overlap if several prompts cover the same ground, so each AI's independent effort goes where it is most distinct.
- Tightening or loosening the bypass section depending on how the others scope it.

Suggested GitHub folder layout

Folder name suggestion: network-defense-intelligence-exercise

- **00-source-dossier/**: OAL-DOSSIER-CANON-002 (PDF and DOCX)
- **01-research-prompts/iteration-1/**: one prompt per AI, named OAL-RPROMPT-NDIE-001-[AI name]
- **01-research-prompts/iteration-2/**: revised prompts with change notes, named OAL-RPROMPT-NDIE-002-[AI name]

- **02-research-reports/**: each AI's independent report, named OAL-RESEARCH-NDIE-001-[AI name]
- **03-synthesis/**: the merged findings, disagreements, and a proposed dossier v1.1 change list
- **04-decisions/**: Architect decisions on Q1 to Q20 and on the required changes, recorded with date and ONESTAMP status

Part 5: My Lens and Blind Spots

The dossier's draft roles give me Forensic Analyst, cross-reference against existing canon, and red-team validation of the scene drafts. Those roles are advisory until you approve them, but they explain why my prompt leans the way it does.

Where my prompt leans

- Toward claim provenance: separating sourced fact, inference, and speculation, and keeping a ledger of each claim.
- Toward adversarial and after-action thinking: how things fail, how failures get noticed, and what the post-incident record says.
- Toward internal consistency: places where the dossier's own pieces conflict with each other.

Where I may be weak

- My reliable knowledge ends in June 2026. Costs, program rules, laws, and incident data from after that must come from search, and the prompt says so.
- I can only reach public, mostly English-language sources. Classified or non-public practice will be missing, which may make real-world deception look less mature than it is, or more.
- Cost figures will be estimates built from public benchmarks. They will be ranges with stated assumptions, not quotes.
- I may under-weight the creative and narrative value of choices that are technically weak, and over-weight realism. The [R], [C], and [N] tags are there so you can decide that trade yourself.

- A research prompt cannot test the system. It can only tell you where public evidence says the design is strong or fragile.

What I would ask the other AIs to challenge

- Whether twelve workstreams is too broad for a single independent effort, and where to cut.
- Whether my bypass categories miss a route a more adversarial mind would find.
- Whether the null case in WS2 is strong enough, or whether I am still leaning toward the design.

Appendices

All tables and worksheets for this document are collected here. Every entry in Appendices C and F is a starting point for the researcher to verify, not an established fact.

Appendix A: Coverage Map of Your Requests

Where each of your asks is handled in the prompt.

Your ask	Workstream	Report section
Goal of the exercise	WS1	Findings: WS1
Why it matters, why it should take place	WS2	Findings: WS2, Dissent
Cost if funded	WS9	Findings: WS9
Possible funding pathways	WS10	Findings: WS10
What needs to be addressed: questions, comments, concerns	WS11	Findings: WS11
Adversarial analysis	WS5, WS6	Findings: WS5, WS6
Recommended next-step actions	WS12	Recommendations
Bypass beyond reverse engineering or mirroring	WS6	Findings: WS6
How the response evolves over time	WS7	Findings: WS7
Real-world scenario, and whether it has been modeled	WS3, WS4	Findings: WS3, WS4
What worked, what went wrong, similar scenarios, lessons	WS4	Findings: WS4
What must change before moving forward	WS12	Recommendations
Independent research, not only supportive findings	Rules 1 to 10	Method; Dissent; Search log

Your ask	Workstream	Report section
PDF and DOCX, GitHub folder	Section 7	Part 4

Appendix B: Shared Evidence Grades

Proposed for all AIs so reports can be compared directly.

Grade	Meaning	Typical source
A	Primary and verifiable	Government report, standard, peer-reviewed paper, official postmortem
B	Credible secondary, or several independent sources agree	Reputable investigative reporting, multiple independent analyses
C	Single source, vendor claim, or commentary	Vendor white paper, blog, one-off interview
D	Researcher inference	Reasoned from A to C evidence; must be labeled as inference
E	Speculation	Plausible but unsupported; use only to frame questions

Appendix C: Candidate Precedents to Examine

Starting list only. Verify details, extend the list, and look for failures as well as successes.

Precedent	Category	Why it matters
The Cuckoo's Egg (Stoll, 1986 to 1987)	Deception, attribution	Fabricated bait kept an intruder connected long enough to trace; lessons on bait credibility and time to attribute
An Evening with Berferd (Cheswick, 1991)	Early deception environment	An interactive fake observed an attacker; shows the labor of sustaining a convincing fake

Precedent	Category	Why it matters
Honeyd (Provos, 2003); Potemkin virtual honeyfarm (Vrable et al., 2005)	Virtual networks at scale	Direct prior art on faking many hosts; scaling and fidelity limits
Honeynet Project and Know Your Enemy papers	Research honeypots	Long record of what honeypots capture and the legal and ethical issues
Honeypot-detection research (mid-2000s, e.g., Holz and Raynal)	Deception detection	How fingerprinting works conceptually and the arms-race dynamic
Moving target defense and address-hopping research	Rotation defenses	Closest prior art to the 3-second cycle; costs and attacker adaptation
MITRE Engage, ATT&CK, D3FEND	Frameworks	Shared vocabulary for adversary engagement and defense
NIST SP 800-53 controls on honeypots and concealment (SC-26, SC-30)	Standards	How deception is already codified in federal control catalogs
Commercial deception platforms (e.g., Thinkst Canary, Acalvio, CounterCraft, and the Attivo and Illusive lineage)	Industry practice	Deployment lessons, pricing, market consolidation; vendor claims need skepticism
Ukraine grid attacks (2015, 2016)	Bypass, operational technology	Long dwell time, legitimate-tool use, physical-process impact
SolarWinds / SUNBURST (2020)	Supply chain	Update-channel compromise; direct analog to transfer by update
Colonial Pipeline (2021)	Critical infrastructure	IT and OT dependency and decisions under pressure

Precedent	Category	Why it matters
Volt Typhoon (public reporting, 2023 to 2024)	Pre-positioning	Living-off-the-land presence in critical infrastructure; defeats perimeter-centric thinking
Salt Typhoon (public reporting, 2024)	Telecom compromise	Carrier compromise that a routing-based design would depend on
Stuxnet (2010)	Air-gap bypass	Isolation is not immunity
Cyber Storm, GridEx, Locked Shields, Eligible Receiver 97, Cyber Flag and Guard	National and allied exercises	What large exercises measure, reveal, and miss
DARPA Cyber Grand Challenge (2016)	Automated defense	Machine-speed defense and its limits
Military deception (WWII Operation Fortitude, Ghost Army, joint doctrine)	Non-cyber deception	Exploit existing beliefs, control the adversary's information channels, confirm the adversary believes
Game-theoretic surveys of cyber deception (e.g., Pawlick, Colbert and Zhu, 2019)	Models	Formal models of when deception pays off

Appendix D: Bypass Category Worksheet

For WS6. Fill one row per route; add rows as needed.

Category	Core question	Dossier element stressed
Pre-positioned access	What if the attacker is already inside before detection?	Step 1 detection; step 8 end condition
Supply chain and update	What if the compromise rides a trusted update, including this system's own?	Step 9 transfer

Category	Core question	Dossier element stressed
Identity and insider	What if no exploit is needed because credentials are valid?	Steps 3 and 5 path control
Out-of-band paths	What if the route avoids the monitored network?	Steps 2 and 5 mapping and migration
Blind-spot systems	What if the target is air-gapped, OT, legacy, or foreign-hosted?	Coverage limitation; step 8 "fully secured"
Attack the defender	What if the orchestrator, sensors, or operator are the target?	Architect as sole authority
Patience and dormancy	What if the attacker waits longer than any cycle?	Step 6 cycling; step 8 end condition
Non-network strategy	What if the goal is reached through people, law, money, or influence?	Whole premise
Jurisdictional seams	What if the target sits between agencies or private owners?	Q1 to Q3 scope and authority
Defender outputs	What if the attacker goes after captured data or the package?	Capture, logs, transfer

Appendix E: Cost Model Worksheet

For WS9. Each tier needs low, base, and high ranges with stated assumptions.

Tier	Scope	What to estimate
L	Lab or simulation proof of concept; synthetic attacker and defender; no live networks	Team, compute, cyber range time, independent red-team review, duration

Tier	Scope	What to estimate
S	Sector or enclave pilot with a consenting owner; limited live telemetry; safety rails	Everything in L, plus sensors, integration, legal and oversight, security, operations
N	National scale; multiple sectors, carriers, and agencies	Everything in S, plus continuous operations, sustainment, partner integration, contingency

Cost categories to estimate in every tier: people; compute and cloud; network and sensors; threat data; tooling and licensing; cyber range and facilities; security and clearances; legal and oversight; independent testing; partner integration; annual operations; contingency.

Appendix F: Funding Pathway Worksheet

For WS10. A starting list to verify, not a recommendation. Confirm each program's current status, dates, and rules.

Pathway family	Examples to check	Facts to capture
Federal research agencies	DARPA, IARPA, DHS Science and Technology, DOE CESER, NSF, military service labs	Open calls, award size, timeline, eligibility, IP and publication terms
Small-business programs	SBIR and STTR across agencies	Current authorization status, phase sizes, agency topics
Procurement and transaction routes	Defense Innovation Unit, other-transaction authority, agency commercial solutions openings	Entry requirements, timelines, typical scale
Research partners	FFRDCs, UARCs, university labs, national laboratories	Partnership terms, overhead, access to ranges

Pathway family	Examples to check	Facts to capture
State and local grants	State and Local Cybersecurity Grant Program and state programs	Eligible applicants, funding rules, current status
Private capital	Venture funds, corporate strategic investors, insurer and reinsurer consortia	Typical check size, expectations, dual-use concerns
Sector consortia	ISACs, sector coordinating councils	Funding mechanisms, willingness to pilot
Philanthropy	Foundations focused on infrastructure resilience and security	Fit, award size, restrictions
International	NATO programs and cooperative centers, EU research and defense funds	Eligibility for US-based teams, timelines
In-universe	ONE Industries and OAL internal mechanisms	Open question for the Architect; flag for canon

---

End of document. Iteration 1 research prompt, prepared by Claude (Anthropic) on November 4, 2026. Advisory only; canon decisions and ONESTAMP remain with the Architect.