False Infrastructure Defense Research (NDIE)

> Independent, evidence-first research into the "Network Defense Intelligence Exercise: False Infrastructure AI" concept, with fictional canon kept separate from empirical cybersecurity findings.

Field	Value
Folder name	`false-infrastructure-defense-research`
Short code	NDIE
Source dossier	OAL-DOSSIER-CANON-002 v1.0.0 (April 11, 2026)
Owner / Architect	Nathan Lewis
README version	0.1 (draft)
Created	November 4, 2026
Status	Prompt-review phase. No research report has been executed.

Alternative folder names: `network-defense-intelligence-exercise` (matches the dossier title; used in several peer prompts) or `ndie-false-infrastructure-research` (shorter). The recommended name is descriptive of the real-world research question rather than the fictional framing, which helps keep canon and evidence separate.

1. What this repository is

This repository holds the collaborative work of a multi-AI research team (Perplexity, ChatGPT, Claude, Google Gemini, Microsoft Copilot, Grok) coordinated by the Architect. It examines a canon scenario in which a national-scale defender maps an intrusion, protects real assets, builds a synthetic "honeypot" infrastructure, silently moves attackers into it, rotates it roughly every three seconds, and transfers the resulting defensive package to a partner agency.

The goal is not to confirm that scenario. The goal is to test it: what it is for, whether it is feasible, how an adversary could bypass it entirely, how the contest evolves over time, what it would cost, how it could be funded, and what must change in the design and the canon before work continues.

2. Central research question

> Can authorized, selectively deployed deception and adaptive isolation measurably reduce the probability or impact of compromise of real assets, relative to credible alternatives, against adversaries who may avoid the deception altogether?

A finding of "insufficient evidence," "narrow the scope," "redesign," or "do not fund" is an acceptable and valuable result.

3. The ten questions this work answers

1. What is the goal of the exercise, and is it coherent?
2. Why is it important, and should it happen at all?
3. If funded as real research, what would it cost?
4. What are the possible funding pathways?
5. What must be addressed: questions, comments, concerns, adversarial analysis, next steps?
6. How could someone who wants real assets bypass the process entirely, beyond reverse engineering or mirroring?
7. How does the response evolve over time, for attackers and defenders?
8. Is this a real-world scenario that could occur, and has it been modeled?
9. What worked and what went wrong in similar scenarios, simulations, and incidents?
10. What must change before moving forward?

4. Ground rules

- **Canon is not evidence.** The dossier is a creative design synthesis. Its "locked" elements bind the fiction only. Every engineering claim is contestable.
- **No canon changes here.** Proposed changes are change requests that wait for the Architect's ONESTAMP. Nothing in this repository ratifies canon.
- **Seek disconfirming evidence.** Every researcher searches for contradicting, null, failed, and abandoned results as deliberately as for support, and reports null searches.
- **Independent first pass.** Each AI should complete its research first pass before reading peer research reports. Prompt review has already occurred; disclose all peer material visible to each researcher. Agreement among AIs is not experimental replication.
- **Grade and tag findings using the approved convention.** Perplexity v1.1 proposes evidence grades A-E; tags R (real-world feasibility), C (canon realism), N (narrative design); confidence high / moderate / low. Team-wide adoption remains pending.
- **Threat-model level only.** No exploit code, payloads, evasion recipes, real-target reconnaissance, credential-abuse instructions, or unauthorized testing.
- **No invented facts.** No fabricated citations, costs, success rates, awards, or precedents. Costs are labeled estimates; funding programs are leads until their current status is verified.
- **Keep sensitive material out.** No credentials, real network topology, personal data, confidential telemetry, or classified/leaked material.

5. Current status and workflow

Stage	Description	State
A	Each AI writes its own research prompt (iteration 1)	Submitted materials available; team-wide completion not verified

Stage	Description	State
B	Prompts shared; each AI reviews peers' prompts (ADD / CHANGE / REMOVE)	Submitted materials available; team-wide completion not verified
C	Optional iteration-2 prompts and agreement on shared conventions	In progress (Perplexity v1.1 issued)
D	Independent research execution, on the Architect's explicit instruction	Not started
E	Compare reports, preserve dissent, synthesize, draft dossier v1.1 change list	Not started

6. Repository layout

```
false-infrastructure-defense-research/
|-- README.md
|-- CHANGELOG.md
|-- 00-source-dossier/            # OAL-DOSSIER-CANON-002 (PDF + DOCX), unmodified
|-- 01-research-prompts/
|   |-- iteration-1/              # one prompt per AI, as submitted
|   `-- iteration-2/              # revised prompts with change notes
|-- 02-research-reports/          # each AI's independent first-pass report
|-- 03-peer-reviews/              # cross-review of prompts and reports
|-- 04-synthesis/                 # merged findings, disagreements, dissent memo
|-- 05-evidence/
|   |-- claim-ledger.csv
|   `-- source-ledger.csv
|-- 06-threat-model/              # bypass matrix, adversary generations
|-- 07-experiments/               # protocols, metrics, acceptance criteria
|-- 08-costs/                     # tiered cost model, assumptions
|-- 09-funding/                   # verified pathway matrix, access dates
|-- 10-governance/                # legal flags, safety case, oversight
|-- 11-decisions/                 # Architect decisions on Q1-Q20 and changes
`-- 12-canon-translation/         # change requests awaiting ONESTAMP (tags C, N)
```

Rules for the layout: originals are never overwritten (new versions are new files); syntheses live apart from first-pass reports; proposed canon amendments live in `12-canon-translation/`, while the original canon dossier stays in `00-source-dossier/`. Canon-translation notes may cite empirical findings but must label the distinction.

7. Naming and ID conventions

Files: `OAL-NDIE-[TYPE]-[MODEL]-v[MAJOR].[MINOR].[ext]`, for example `OAL-NDIE-PROMPT-PERPLEXITY-v1.1.pdf`, `OAL-NDIE-REPORT-CLAUDE-v1.0.docx`, `OAL-NDIE-REVIEW-GEMINI-v1.0.md`.

Stable IDs (proposed): `CLM-[MODEL]-###` claims, `SRC-[MODEL]-###` sources, `RSK-[MODEL]-###` risks, `QST-[MODEL]-###` questions, `EXP-[MODEL]-###` experiments, `DEC-[OWNER]-###` decisions, `CHG-[MODEL]-###` change requests.

Evidence grades: A primary and verifiable; B credible secondary or multiple independent sources; C single source or vendor claim; D labeled researcher inference; E speculation

(framing only).

Capability status vocabulary: Demonstrated; Demonstrated at smaller scale; Partially demonstrated; Research-stage; Speculative; Unsupported.

8. Contribution workflow

1. Work from the latest prompt version and the unmodified source dossier.
2. Write hypotheses and disproof conditions before searching.
3. Log every search, marked supportive, disconfirming, or neutral.
4. Add each material claim to the claim ledger with source, location, grade, tag, and counterevidence.
5. Proposed workflow, subject to repository configuration: submit as a branch/pull request; the Architect approves merges to `11-decisions/` and `12-canon-translation/`.
6. Provide PDF and DOCX for narrative deliverables, with a Markdown source and CSV exports where useful.
7. Record reproducibility details (tool versions, search dates, cutoffs, seeds for any simulations).

9. Decision gates

Gate	Evidence required	Outcome
G0 Concept clarity	Asset boundary, owner, purpose, threat model, metrics	No engineering before definitions
G1 Research justification	Comparable evidence, incremental value over baselines, known bypasses	Continue, narrow, or abandon
G2 Safe prototype	Approved sandbox, oversight, budget, test plan	Authorize isolated proof of concept only
G3 Controlled pilot	Replicated measured advantage, manageable collateral risk, rollback	Seek a consenting external partner
G4 Scale decision	Independent review, legal authority, cost-benefit, reliability	Scale only within approved domain

Production deployment is never an automatic next step.

10. Safety, legal, and ethical boundaries

- Desk research and isolated, authorized simulation only. No scanning, probing, or testing of real systems.
- This repository gives no legal advice. Authorization, monitoring, privacy, interception, retention, liability, export, and critical-infrastructure questions go to qualified counsel.
- Hack-back and changes to third-party systems are out of scope without separate lawful authority.
- Public/private split: anything operationally sensitive, personal, or third-party confidential stays out of the public tree. Run a release check before any publication.

- Licensing and ownership are undecided; see open items.

11. Open items for the Architect

- Confirm the folder name and visibility (public, private, or mixed).
- Choose a license for text, data, and code, or defer publication.
- Decide which shared conventions the team standardizes (evidence grades, tags, cost tiers, format).
- Decide whether iteration-2 prompts merge into one common core or stay separate.
- Authorize research execution (Stage D) when ready.
- Resolve the dossier's high-priority open questions (Q1-Q20) as they affect scope.

12. Glossary

- **Canon:** the fictional design constraints in the OAL story universe.
- **ONESTAMP:** the Architect's electronic approval; nothing is ratified canon without it.
- **Honeypot / deception environment:** a decoy system designed to attract and observe attackers.
- **Moving-target defense:** changing system attributes over time to raise attacker cost.
- **Bypass:** reaching real assets without engaging the deception layer.
- **Capture (this project):** a set of separately measurable states, not a single outcome; connection captured does not mean attacker identified or eliminated.
- **Claim ledger:** table linking each material claim to evidence, grade, and counterevidence.

13. Disclaimer

This is a research and creative-development project. Nothing here is an operational security recommendation, legal advice, funding application, or approved canon. All cost figures are labeled estimates, and funding programs are not assumed available until verified at the time of research.

14. Document provenance and controls

This README is proposed documentation, not a record of an existing configured GitHub folder. Prepared by Perplexity on November 4, 2026, from Nathan Lewis's requests, the supplied OAL-DOSSIER-CANON-002 v1.0.0, the shared peer prompts, and Perplexity's NDIE-RESEARCH-PROMPT-PERPLEXITY-001 v1.1. The source PDF and DOCX are alternative formats of one dossier, not independent corroboration. No research findings, prices, funding eligibility, or deployment claims have been validated by preparing this README. No repository has been created or modified.

Maintain an inventory of artifacts actually committed and their approval status. A proposed directory, expected deliverable, or contributor listing does not establish that the artifact exists

or that the named contributor completed or approved work. Add contact and security-disclosure details before accepting external contributions; none are supplied here.