

COLLATION OF PARALLEL VERBATIM RECORDS

Same commission. Four independent renderings. No file is declared winner.

1. What this note is

The Architect supplied one original passage and asked for a foundational verbatim source record plus a discrepancy note on later isolation of “a threat can be anything.” Several models produced independent records. This file does not consolidate them into a single official text. It records where they agree, where they silently rewrite the source, and what remains useful in each.

2. Files compared

Label here	File / ID	Stated preparer
Grok record	OAL-CF-EST-VERBATIM-v0.1.0	Grok independent versioning record from the same commission.
Short record	foundational_verbatim_source_record	Unattributed in the file header. Compact source + discrepancy + scope table.
Claude record	OAL-Verbatim-Source-Record-Threat-Statement	Prepared by Claude, at Nathan Lewis II's request.
ChatGPT record	OAL-CF-VERBATIM-v1.0.0	Prepared by ChatGPT (OpenAI). Includes screenshot exhibits and SHA-256 hashes.

3. Shared finding

All four records agree on the discrepancy itself:

- The source paired “A threat can be anything” with immediate laboratory limits.
- The Grok foundational draft and the Perplexity draft isolated the open clause and scored it as unlimited / unbounded.
- Later operational bounds (EIC, protected endpoint, entity class) may still be design work. That is not the same as saying the source left the clause unbounded.

4. Source-fidelity table

Control text is the wording the Architect typed in the commission, including artifacts.

Source token	Grok	Short	Claude	ChatGPT
<i>tobiological</i>	kept	kept	normalized to “to biological”	kept in source block
<i>research questions is</i>	kept	kept	changed to “research question is”	kept
<i>entiry</i>	kept; noted as artifact	kept	changed to “entity”	kept in source block; PDF purpose note lists “entity” as the artifact, which is already a silent swap
<i>doesn't seem to discover</i>	kept; noted	kept	changed to “doesn't seek to discover”	kept in source block; commentary sometimes says “does not seek”
Screenshots of the two drafts	cited, not embedded	cited, not embedded	cited, not embedded	embedded as Exhibit A/B with SHA-256

5. What each record adds that is not source text

Grok record

- Splits source into Blocks A/B/C and writes a reading rule for Block C.
- Quotes the exact Grok §8.1 and Perplexity table language.
- Separates “source already bound this” from “design can still require an EIC.” Binding instructions for later quotes.

Short record

- Closest to a clean archival page: purpose table, italic source block, discrepancy, proposed version-history entry, scope table.
- Does not name a model as preparer. Does not embed exhibits. Least extra interpretation.

Claude record

- Clear structure and an explicit “this is not a correction of those drafts’ design conclusions.”
- Counts “five explicit bounding negations.” The source sentence is one clause-stack; the count is an editorial parse, not a source number.
- Highest silent-correction load of the set. That matters because the commission was a verbatim record.

ChatGPT record

- Best exhibit handling: embeds both screenshots and hashes 870.png / 869.png.
- Section 4 “Canonical contextual reading” is useful as commentary and should stay labeled commentary. It is not the source.
- Baseline rule is sound: do not rewrite v1.0.0; issue new versions for operational definitions.

6. Second-order discrepancy

The first-order problem was later drafts isolating “a threat can be anything.” A second-order problem now exists inside the verbatim records themselves: some files that exist to freeze the source also tidy the source. Claude’s “seek” for “seem,” “question” for “questions,” and “entity” for “entiry” are the same class of move as the original complaint, just applied to smaller tokens. ChatGPT’s purpose note listing “entity” as the preserved artifact does the same thing one layer up.

7. Suggested archive practice

- Keep all four files. Do not collapse them.
- When quoting the Architect, quote the commission text or a file that kept the artifacts. Do not quote Claude’s cleaned paragraph as “verbatim.”
- Use ChatGPT’s exhibits when the screenshots themselves need to travel with the record.
- Use the short record when a one-page source freeze is enough.
- Use the Grok record when the need is to bind later design language to Block C as a whole unit.

End of collation. Parallel drafts remain parallel.