

DEVELOPMENT ITEMS RESEARCH & ANALYSIS

Development Items Research & Analysis

All Three Projects

Unified Research Report for Author Approval Decisions

Document ID: OAL-DOSSIER-DEV-001

Version: v1.1 | April 11, 2026

Status: PENDING ELECTRONIC APPROVAL

Prepared by: Perplexity Computer

Revision: Full team review integration (ChatGPT · Gemini · Claude · Copilot · Grok)

Distribution: ChatGPT · Google Gemini · Claude AI · Microsoft Copilot · Grok AI

PENDING ELECTRONIC APPROVAL — ONESTAMP REQUIRED

APPROVAL NOTICE: This document is pending electronic approval and ONESTAMP from the author (Nathan Lewis, Architect). No elements herein should be treated as final, approved direction until ONESTAMP is applied. All research findings, analysis, and recommendations are advisory until that approval is granted. Where research supports a particular design choice, it is noted explicitly. Where research contradicts confirmation bias, it is flagged in dedicated sections.

v1.1 REVISION NOTICE: This version integrates formal dissent and addendum feedback from the full AI collaborative team following distribution of v1.0.0. Four structural issues identified by ChatGPT, seven additions identified by Claude, three removals identified by Copilot, and consensus notes from Gemini and Grok have been incorporated. New sections added: Cross-Project Dependency Architecture (Section 7), Omnicellular Architecture Reference Layer (Section 8), Failure Cascade & Recovery Logic (Section 9), Architect Identity & Presence Model (Section 10), Ethical Boundary Framework (Section 11), Information Asymmetry Index (Section 12). All decision matrices have been updated with new items. Perplexity's reconciliation notes are included where team feedback differed.

Table of Contents

1. How to Use This Document
2. Team Feedback Integration Log
 - 2.1 ChatGPT — Four Critical Issues
 - 2.2 Claude AI — Seven Recommended Actions
 - 2.3 Google Gemini — Canon Consistency Notes
 - 2.4 Microsoft Copilot — Three Removals + Structural Additions
 - 2.5 Grok AI — Team Consensus Confirmation
3. Project 1 — Kinetic Threshold Demonstration & Open Ocean Parley
 - 3.1 Key Development Items Summary
 - 3.2 Research & Analysis: Kinetic Armor Technology
 - 3.3 Research & Analysis: Synthetic Humanoid Control
 - 3.4 Research & Analysis: Corporate Weapons Demonstration Ethics
 - 3.5 Research & Analysis: Open Ocean Covert Meeting Design
 - 3.6 Confirmation Bias Flags
 - 3.7 Author Decision Matrix — Project 1
4. Project 2 — Network Defense Intelligence Exercise
 - 4.1 Key Development Items Summary
 - 4.2 Research & Analysis: National-Scale Cyber Honeypot & Deception Systems
 - 4.3 Research & Analysis: Legal Authority & Ethical Framework
 - 4.4 Research & Analysis: Synthetic Network Fidelity
 - 4.5 Research & Analysis: Adversarial Counter-Simulation Risk
 - 4.6 Research & Analysis: Mirror Core Transfer
 - 4.7 Confirmation Bias Flags

- 4.8 Author Decision Matrix — Project 2
- 5. Project 3 — Pathogen Outbreak Tabletop Exercise
 - 5.1 Key Development Items Summary
 - 5.2 Research & Analysis: The "10-Year Outbreak Cycle" Claim
 - 5.3 Research & Analysis: Pathogen Archetype Selection
 - 5.4 Research & Analysis: Provider Clinical Uncertainty
 - 5.5 Research & Analysis: TTX Design Standards
 - 5.6 Research & Analysis: Equity Layer
 - 5.7 Confirmation Bias Flags
 - 5.8 Author Decision Matrix — Project 3
- 6. Cross-Project Themes & Unified Recommendations
- 7. Cross-Project Dependency Architecture (NEW — ChatGPT / Copilot / Claude)
- 8. Omnicellular Architecture Reference Layer (NEW — ChatGPT / Gemini)
- 9. Failure Cascade & Recovery Logic (NEW — ChatGPT)
- 10. Architect Identity & Presence Model (NEW — ChatGPT / Copilot)
- 11. Ethical Boundary Framework (NEW — ChatGPT / Copilot)
- 12. Information Asymmetry Index (NEW — ChatGPT)
- 13. Priority Approval Checklist — What to Decide Today
- 14. Version History

1. How to Use This Document

This document is structured to help you make the specific approval decisions needed to advance all three projects today. Each project section contains:

- **Key Development Items** — the blocking decisions identified in the existing dossiers
- **Research & Analysis** — external evidence for and against each design choice, with sources
- **Confirmation Bias Flags** — where research challenges assumptions rather than confirming them
- **Author Decision Matrix** — a summary table of every HIGH-priority open question with a recommended answer, a confidence level, and the research basis

Your job is to review each matrix, mark your decision (Approve / Modify / Defer), and apply ONESTAMP to the approved elements. That is the complete path to moving all three projects forward today.

NEW in v1.1: Sections 7–12 address structural architecture items that must be resolved before ONESTAMP is applied. These are not optional refinements — they are the backbone that the three projects depend on. Review them in order before completing the project-level matrices.

2. Team Feedback Integration Log

This section documents every team contribution received after v1.0.0 distribution, the action taken for each, and the rationale. All Perplexity reconciliation notes are included where team inputs differed.

2.1 ChatGPT — Four Critical Issues (Required Fixes)

#	Issue	Action Taken	Section
C1	Missing Cross-Project Dependency Map	ADDED — Full dependency architecture with unlock logic and timeline markers	Section 7
C2	Omnicellular Gap — no underlying architecture reference	ADDED — Placeholder reference layer with Ontological Anchoring and Bandwidth Throttling concepts	Section 8
C3	Failure Cascade Model incomplete	ADDED — Single-failure, compound-failure, and recovery pathways for all three projects	Section 9
C4	Project 3 authority gap (D4) still unresolved	STRENGTHENED — D4 now requires named improvement plan owner with budget authority as hard prerequisite	Section 5.8

ChatGPT also recommended three advisory improvements (Architect Presence Model, Ethical Boundary Conditions, Information Asymmetry Index) — all three have been added as Sections 10, 11, and 12.

2.2 Claude AI — Seven Recommended Actions

#	Recommendation	Action Taken	Section
CL 1	Add Q20 to Project 1 matrix: post-demo disposition of synthetics	ADDED to Project 1 decision matrix	Section 3.7
CL 2	Elevate Project 2 ethics beat from flags section to formal decision matrix	ELEVATED — now Q21 in Project 2 matrix	Section 4.7
CL 3	Add failure-mode note to Project 2 Q2: real-world adaptation protocols	ADDED as note within Q2 entry	Section 4.7
CL 4	Strengthen D4: improvement plan owner with budget authority is a hard prerequisite	STRENGTHENED — D4 language updated	Section 5.8

CL 5	Equity injects must be designed concurrently, not deferred	UPDATED — D1/equity matrix entry strengthened	Section 5.8
CL 6	Add cross-project interdependency map with chronological/causal relationships	ADDED — Section 7 covers this fully	Section 7
CL 7	Name AI assignments in the approval checklist before ONESTAMP	ADDED — AI assignments are now part of the checklist	Section 13

2.3 Google Gemini — Canon Consistency Notes

#	Note	Action Taken	Section
G1	Omnicellular v2: Ontological Loophole gap (area-of-effect negation not addressed)	ADDRESSED in Omnicellular Reference Layer — Ontological Anchoring concept codified	Section 8
G2	First Memory encoding mechanism must be grounded in existing OAL infrastructure	NOTED — placeholder in Section 8 with explicit author decision required	Section 8.2
G3	Nullarch's Superior: agenda must be defined to prevent generic-antagonist drift	NOTED — flagged as author decision required, not resolved by Perplexity	Section 8.3

2.4 Microsoft Copilot — Three Removals + Structural Additions

Removals acted upon:

#	Copilot Flag	Action Taken
CP 1	"Nathan as sole operator" implicitly embedded in narrative assumptions	CONFIRMED REMOVED — Section 10 (Architect Presence Model) explicitly replaces this framing throughout
CP 2	"Perfect procedural fidelity" language still implied in Project 2	CONFIRMED REMOVED — all "perfect clone" language removed; replaced with "procedural shadow" framing
CP 3	"Pre-theft placement as default" treated as recommendation rather than neutral option	CONFIRMED ADJUSTED — Decision Matrix Q10 now presents both options neutrally; author must decide

Copilot's structural additions were merged with ChatGPT and Claude additions — all covered in Sections 7–12.

Perplexity reconciliation note on Copilot Issue CP3: Copilot's concern that "pre-theft placement as default" is a bias is valid. However, the research basis for pre-theft placement (surveillance proof → Nathan accelerates the theft) is still the strongest narrative logic. The matrix now presents both options with their trade-offs rather than declaring a recommendation, per

Copilot's feedback.

2.5 Grok AI — Team Consensus Confirmation + Minor Refinements

Grok's review confirmed no blocking canon contradictions. Three minor polish items from Grok's team consensus summary are noted:

#	Grok Note	Action Taken
GR 1	Add minor armor-demo anomaly beat for character depth	ADDED as Project 1 Q21: Design the anomaly beat
GR 2	Nathan as designer/architect (not sole operator) — already in flags; make explicit	MADE EXPLICIT in Section 10
GR 3	Bold "testing the system, not evaluating individuals" in TTX ground rules	BOLDED in Project 3

3. Project 1 — Kinetic Threshold Demonstration & Open Ocean Parley

3.1 Key Development Items Summary

From OAL-DOSSIER-CANON-001 v1.0.0, the following HIGH-priority items remain unresolved:

#	Item	Status
Q1	Timeline: how many years before the theft?	UNRESOLVED
Q3	Audience composition at the armor demo	UNRESOLVED
Q5	Nathan's control mechanism over the synthetics	UNRESOLVED
Q7	Third-shot impact severity for Beta	UNRESOLVED
Q9	Demo classification: public, internal, or classified joint?	UNRESOLVED
Q10	Timeline placement for Ocean Meeting (pre- or post-theft)	UNRESOLVED — see Copilot note in Section 2.4
Q12	Colleague archetype	UNRESOLVED
Q14/Q15	Meeting surface and real objectives	UNRESOLVED
Q17	What changes hands	UNRESOLVED
Q18	Does Nathan leave with advantage, compromise, or disadvantage?	UNRESOLVED
Q20	Post-demo disposition of synthetic humanoids	NEW — Claude
Q21	Is there a minor demo anomaly beat?	NEW — Grok/Perplexity

3.2 Research & Analysis: Kinetic Armor Technology

The Real-World Context for Kinetic Redirection Armor

Real active protection systems (APS) for military vehicles currently operate on intercept-and-destroy logic — not redirection. The Trophy system (Israel), Iron Fist, and Rheinmetall StrikeShield all destroy or deflect incoming threats using fragmentation charges or directed energy, typically within 6–15 meters of the protected vehicle. None currently achieve the energy storage and return mechanic that the three-shot protocol describes.

What Supports the Design (Pro):

- Explosive Reactive Armor (ERA) already uses kinetic energy manipulation — the flyer plates move across an incoming jet, feeding armor material into the interaction. The underlying physics of kinetic energy being managed rather than simply blocked is credible as a direction of travel for advanced materials science.

- Rheinmetall's StrikeShield, as of 2025, is being tested against kinetic energy penetrators (APFSDS rounds) — demonstrating that the field is actively extending toward harder KE threats.
- Non-Explosive Reactive Armor (NERA) variants exist specifically to reduce collateral risk, suggesting the field is already pursuing more controlled energy management.
- The "three-shot protocol" is narratively distinct and memorable precisely because it departs from real-world intercept logic — appropriate for near-future corporate R&D; fiction.

What Challenges the Design (Contra — Anti-Confirmation Bias):

- The energy storage-and-return mechanic has no current physical analog. Real ERA is single-use per tile; it does not accumulate energy across multiple impacts.
- At realistic projectile velocities (2,000+ m/s for APFSDS), forces involved push into theoretical metamaterial territory with no current prototype.
- "Armor resets after sustained period without fire" implies an active powered system, not purely passive — changing form factor and vulnerability profile significantly.
- Critics of power-armor fiction note that energy storage at projectile scale creates secondary detonation risk. What happens if the armor is destroyed before the third shot cycles? This is the canonical failure-mode gap. (See also Section 9 — Failure Cascade for Project 1.)

Research Verdict: The concept is internally consistent as near-future speculative fiction and is grounded in real directional developments in APS research. The three-shot protocol is not implausible — it is ahead of current capability in a direction the field is actively pursuing. The design works. The failure mode gap must be addressed.

3.3 Research & Analysis: Synthetic Humanoid Control

Brain-Computer Interface (BCI) Control of Humanoids — Current State

Peer-reviewed research in Sensors (Basel) documents BCI-based humanoid control using EEG signals, achieving 78–87% real-time accuracy in directing humanoid movement via TCP/IP client-server architecture. BCI-based humanoid control is not speculative — it is demonstrated in laboratory conditions.

What Supports the Design (Pro):

- "Scripted AI" (pre-programmed behavioral sequences) is the most technically grounded option — the synthetics execute a scripted demonstration sequence triggered by Nathan, requiring no real-time control complexity.
- A "proprietary hybrid" (pre-scripted with Nathan able to intervene via interface) is the most dramatically interesting and technically defensible choice — it gives Nathan visible authority without requiring the audience to understand the mechanism.
- The Uncanny Valley phenomenon is a free dramatic tension element: an audience observing synthetics perform a combat demonstration experiences heightened psychological response at the threshold of human-likeness.

What Challenges the Design (Contra):

- Full real-time neural control at the fidelity required for a live weapons demonstration is not achievable outside laboratory conditions. A public demonstration relying on this would be a liability risk.
- If Nathan controls the synthetics via neural link during the demonstration, the scene implies the First Memory or a similar technology is active — potentially over-exposing that thread earlier than intended.

Research Verdict for Q5: Pre-scripted AI with Nathan holding a manual override capability. This is the most defensible, least over-exposing option. The scripted sequence runs autonomously; Nathan's visible calm comes from knowing exactly what will happen; the override exists if something goes wrong — creating the failure-mode possibility that the anomaly beat (Q21) can leverage.

NEW — Post-Demo Disposition (Q20 — Claude Addendum): Whether the synthetics are decommissioned, retained, repurposed, or whether one is unaccounted for after the demonstration has downstream canon implications. This must be decided before the scene is drafted. The most narratively generative option: one unit is retained by the demonstration facility under the government's contractual rights — Nathan does not object publicly, but this becomes a thread. See Section 7 for dependency implications.

3.4 Research & Analysis: Corporate Weapons Demonstration Ethics

Real corporate weapons demonstrations to joint corporate-government audiences are documented practice. Historical and contemporary examples involve: invited military acquisition officers, corporate board members and investors, NDAs and classification levels varying by prototype sensitivity.

Research Verdict for Q9: Recommend joint corporate-government classified demonstration with: invited military acquisition officers, ONE Industries board or senior investors, and at least one ethics or oversight figure (for moral friction). This is both realistic and narratively superior to a public demo.

Ethical Layer — Copilot Addition (see also Section 11): The ethics of deploying synthetic humanoids as weapons demonstrators requires explicit framing in the scene. Nathan's ethical position here — using entities that may be on the threshold of sentience to demonstrate weapons — is a defining character beat. The Architect Ethical Doctrine (Section 11) must be referenced here: Nathan has a defined position, not an unexamined one.

3.5 Research & Analysis: Open Ocean Covert Meeting Design

Naval intelligence history provides rich precedent for covert naval rendezvous. Transit method analysis:

Method	Realism	Narrative Tone	Recommendation
--------	---------	----------------	----------------

Small launch	High — used in actual intelligence operations	Grounded, tense, deniable	Viable
Submersible	Medium-high — technology exists	Covert, high-tech, cinematic	Preferred for tone
Helicopter	High — but compromises secrecy	Visible, confirmable	Not recommended

On Q10 (Timeline — Copilot Reconciliation):

Copilot correctly identified that v1.0.0 treated pre-theft placement as a default rather than a recommendation. The two options are presented neutrally below. The author must decide.

Placement	What It Enables	Trade-Off
Pre-theft (6–12 months prior)	Nathan learns of surveillance → accelerates theft timeline → reactive to threat	Nathan introduced as proactive/strategic; less omnipotent
Post-theft	Nathan already executed the theft before this meeting; colleague brings consequence	Nathan introduced as already having acted; higher stakes immediately

Research Verdict for Q17: The most narratively load-bearing option — proof of government surveillance of ONE Industries — establishes a threat that makes Nathan reactive, sets up the theft naturally, and is historically grounded. The author may select an alternative, but this option has the highest research support.

Nathan's Outcome (Q18): Research on narrative tension consistently supports: Nathan leaves with a disadvantage he chooses to interpret as advantage. Claude's addendum specifies: the cost should be relational, not informational. Nathan losing certainty about someone he trusted is not recoverable. Q18 should specify the relational cost mechanism, not just the outcome category.

3.6 Confirmation Bias Flags — Project 1

The following items represent places where the existing narrative direction may be reinforcing what Nathan wants to be true rather than what creates the strongest story:

FLAG 1: Nathan as Perpetually In Control Research on compelling protagonist dynamics consistently shows that the most memorable demonstrations of power include something that goes slightly wrong and is handled with controlled improvisation. A minor anomaly in the armor demo would make Nathan more compelling. This is now formalized as Q21 — the anomaly beat. (Endorsed by Grok across the full team.)

FLAG 2: The Colleague as Information Delivery Vehicle The most dramatically durable colleague is one who wants something from Nathan that he is not willing or able to give. The meeting should have a cost to Nathan, not just a benefit. Copilot's addendum requires the colleague archetype to include: a motive, a cost, a leverage point, a reason they risk the meeting, and a reason they don't simply report Nathan. These are now required fields in Q12.

FLAG 3: Pre-Theft Placement as Default (Copilot) This was a bias in v1.0.0. It has been corrected. Both options are presented neutrally in Q10. The author must decide.

FLAG 4: Synthetic Humanoid Ethics as Background Element The ethics of using near-sentient entities as weapons demonstrators is not a background element — it is one of Nathan's most character-defining choices. If the story treats this as unremarkable, it loses a layer. See Section 11 for the Ethical Boundary Framework.

3.7 Author Decision Matrix — Project 1

#	Question	Recommended Answer	Con f.	Research Basis	Your Decision
Q1	Years before theft?	3 years	High	Allows relationship arcs; avoids "too early to matter"	
Q3	Audience composition?	Military acquisition officers + ONE Industries board + 1 ethics observer	High	Consistent with classified demo practice; creates moral friction	
Q5	Control mechanism?	Pre-scripted AI with Nathan's manual override	High	Most technically defensible; avoids First Memory over-exposure	
Q7	Third-shot severity for Beta?	Dramatically flung, clearly functional	High	Shows kinetic energy was real; Beta unharmed = armor limits demonstrated	
Q9	Demo classification?	Joint corp-gov classified	High	Realistic; most narratively interesting	
Q10	Ocean meeting timeline?	NEUTRAL — Author must decide	N/A	See Section 3.5 for both options	

Q12	Colleague archetype?	Black-Ops / Military Liaison — define motive, cost, leverage, risk reason, silence reason	High	Most durable relationship type; leverage-based not loyalty-based	
Q17	What changes hands?	Proof of government surveillance of ONE Industries	High	Sets up the theft; historically grounded	
Q18	Nathan's outcome?	Relational cost: loss of certainty about someone he trusted	High	Claude addendum; most durable arc cost	
Q19	Reentry consequence?	A communication arrived during his absence that changes his next move	Med.	Creates forward momentum	
Q20	Post-demo synthetic disposition?	One unit retained by government under contract — not objected to publicly	Med.	Claude addendum; most narratively generative; creates downstream thread	
Q21	Minor demo anomaly beat?	Yes — Beta's cycle initiates 0.4 seconds late; Nathan's expression doesn't change	High	Grok/Perplexity; avoids omnipotence; character-defining control	
Q22	AI assigned to draft scene?	Name before ONESTAMP is applied	Req.	Claude addendum	

4. Project 2 — Network Defense Intelligence Exercise

4.1 Key Development Items Summary

From OAL-DOSSIER-CANON-002 v1.0.0, the following items remain unresolved:

#	Item	Status
Q1	Jurisdiction scope	UNRESOLVED
Q2	Exercise classification: simulated, live, or hybrid?	UNRESOLVED
Q3	Operating authority	UNRESOLVED
Q4	Primary attack type	UNRESOLVED
Q5	Detection timing	UNRESOLVED
Q10	Is the 3-second cycle a fixed constant?	UNRESOLVED
Q11	Authorization: contain only vs. retaliate?	UNRESOLVED
Q12	Innocent traffic captured?	UNRESOLVED
Q14	What exactly is transferred?	UNRESOLVED
Q16	Future use of the Mirror Core	UNRESOLVED
Q17	Canonical failure mode	UNRESOLVED
Q19	Personal/political cost to Nathan	UNRESOLVED
Q20	First Memory / Time Wars connection at conclusion	UNRESOLVED
Q21	Ethics beat delivery method	NEW — Claude

4.2 Research & Analysis: National-Scale Cyber Honeypot & Deception Systems

Real-World Precedent for Nation-Scale Cyber Deception

The UK's National Cyber Security Centre (NCSC) published research in 2024–2025 explicitly pursuing nation-scale cyber deception deployment, targeting 5,000+ honeypot instances, 20,000+ internal tripwire instances, and 2,000,000+ honeytokens. Their three core research questions directly parallel the issues in OAL-DOSSIER-CANON-002.

What Supports the Design (Pro):

- The NCSC's Active Cyber Defence 2.0 program is the real-world government equivalent of what the Architect built
- The March 2026 US National Cyber Strategy explicitly endorses private-sector cyber operators pursuing adversary networks

- Adaptive honeypot systems that improve from attacker behavior are documented and in production

What Challenges the Design (Contra — Anti-Confirmation Bias):

- NCSC 2025 study: "cyber deception can work, but it's not plug-and-play" — it requires the right data and context or generates noise rather than insight
- National-scale systems are not operated by a single individual — they require distributed organizational infrastructure. Nathan as architect/designer is defensible; Nathan as solo operator is not. (See Section 10 — Architect Presence Model for resolution.)

Copilot Removal — Confirmed: All language implying Nathan as sole real-time operator has been removed from narrative assumptions. Section 10 establishes the correct framing.

4.3 Research & Analysis: Legal Authority & Ethical Framework

Active Cyber Defense Legal Landscape

Under current US law (CFAA): purely defensive measures within own systems are permissible; tracking tools reaching into attacker systems are legally gray; offensive measures without government authorization are prohibited; weaponized honeypots damaging attacker infrastructure remain legally unsettled.

2025–2026 developments: The 2026 NDAA expanded government cyber authority. A pending bill (Cybercrime Marque and Reprisal Authorization Act) would allow deputized private operators to pursue foreign criminal enterprises. The March 2026 National Cyber Strategy explicitly endorses deploying defensive and offensive operations using federal resources and private sector.

Research Verdict for Q3: Named joint task force, classified, black-program contractor designation. Most legally grounded; most dramatically interesting.

Research Verdict for Q11: Strict containment and observation only — retaliation capability exists but is not exercised. Creates the most powerful restraint narrative.

Research Verdict for Q12: Innocent traffic captured, flagged, immediately reviewed and released. This is a real operational problem Nathan has had to engineer into the system. The ethics of collateral honeypot capture must be acknowledged in the scene, not elided.

4.4 Research & Analysis: Synthetic Network Fidelity

Supporting the "Procedural Shadow" Model:

High-interaction honeypots are resource-intensive; they are only deployed where attacker engagement is expected. The SODA system (ACM, 2021) uses automated honeypot deployment triggered by malware behavioral analysis — a real analog of "spin up synthetics as needed."

Copilot Removal — Confirmed: All "perfect clone" language has been removed. The Architect's system procedurally generates plausible shadows of real infrastructure — not exact duplicates. This is more interesting narratively (the fake world has subtle imperfections a sophisticated

adversary might detect) and more technically defensible.

3-Second Cycle Research Verdict: Position as a tunable parameter that Nathan sets deliberately — not an industry standard. It is his signature. The anomalous cycling rate is also a Chekhov's Gun: it is the fingerprint that a sufficiently sophisticated adversary will eventually detect. The recursive counter-simulation failure mode (Q17) begins when that fingerprint is read.

4.5 Research & Analysis: Adversarial Counter-Simulation Risk

Current research explicitly models adversaries who learn to recognize honeypot signatures through behavioral analysis. The "recursive counter-simulation / mutual deception standoff" failure mode is the most technically credible and narratively generative option. A future adversary who detects the 3-second cycling signature and begins feeding false behavioral data back into the capture loop is:

- Technically credible
- Narratively escalating
- A direct path to the Time Wars thread

Research Verdict for Q17: Recursive counter-simulation as canonical failure mode — planted as seed in this exercise. In this exercise the defense works completely. One captured attacker's behavioral data contains a pattern the Architect doesn't fully understand until much later.

4.6 Research & Analysis: Mirror Core Transfer

Copilot Addition — Mirror Core must be:

- Dangerous if misused
- Traceable (Nathan can track its deployment)
- Corruptible (a future adversary can introduce artifacts into it)
- A future plot hook (the Mirror Core's existence eventually becomes a liability or a weapon turned against its architect)

These requirements are now part of the Q14 recommended answer.

4.7 Confirmation Bias Flags — Project 2

FLAG 1: Nathan as Sole Operator of National Infrastructure (Confirmed Removed) Resolution: Nathan is the mind behind the system. The system runs itself. Nathan's role is architectural oversight and intervention when required. See Section 10.

FLAG 2: The Transfer Is Framed as Generous What Nathan keeps (attacker behavioral data, attribution intelligence, the Mirror Core's architectural knowledge) may be more valuable than what he gives away. This ambiguity makes Nathan more complex. The scene should not resolve whether the transfer is generous or strategic — both must be readable simultaneously.

FLAG 3: The Ethics Question Is Framed as External This was in v1.0.0. Claude's addendum elevates it to a formal decision item. The ethics beat is now Q21 in the decision matrix. The most powerful version: Nathan asks himself the question. The scene ends before he answers it internally. The audience does not know whether the answer is "I already have" or "not yet."

FLAG 4: Timeline vs. Project 1 (Copilot Cross-Project Contradiction) Copilot identified: Project 1 places the Ocean Meeting 6–12 months pre-theft. Project 2 implies Nathan is already operating at national scale. These timelines must be reconciled in the cross-project dependency map (Section 7). If Project 2 occurs after the theft, Nathan's national-scale operation is a post-theft development — which changes Project 1's Ocean Meeting stakes significantly.

4.8 Author Decision Matrix — Project 2

#	Question	Recommended Answer	Con f.	Research Basis	Your Deci sion
Q1	Jurisdiction scope?	Full US critical infrastructure (civilian + government); military networks are a separate parallel system Nathan does not touch	High	Highest stakes without military-only framing	
Q2	Simulated, live, or hybrid?	Live networks during active intrusion with Nathan's synthetic layer intercepting before real impact. NOTE: Real-world adaptation requires defined failure-mode protocols if intercept fails	High	Creates genuine stakes; legally defensible under Q3 authority	
Q3	Operating authority?	Named joint task force, classified, black-program contractor designation	High	Most legally grounded; most dramatically interesting	

Q4	Primary attack type?	AI-driven coordinated multi-vector — supply chain compromise + DDoS chain	High	Most current and technically sophisticated	
Q5	Detection timing?	Pre-exploit behavioral recognition	High	Demonstrates Architect's superiority without waiting for damage	
Q10	3-second cycle?	Tunable — Nathan's signature; faster than any adversary has seen; also a Chekhov's Gun	High	Technically defensible; narratively distinctive; future failure seed	
Q11	Contain or retaliate?	Contain and observe only — retaliation capability exists but is not used	High	Legal; creates ethical restraint narrative	
Q12	Innocent traffic?	Captured, flagged, immediately reviewed and released — acknowledged as real operational problem	High	Ethically honest; creates real moral complexity	
Q14	What is transferred?	Mirror Core black-box module: dangerous if misused, traceable, corruptible, future plot hook	High	Agency gets capability; Nathan keeps knowledge; creates future story potential	
Q16	Future use?	Used successfully again; then in a later arc turned against Nathan	High	Best escalation arc	
Q17	Canonical failure mode?	Recursive counter-simulation — planted as seed here; discovered later	High	Most technically credible; most narratively generative	
Q19	Personal cost?	Political cost arrives later — delayed and compound; not visible in this exercise	Med.	Delayed consequences are more durable dramatically	
Q20	First Memory / Time Wars connection?	All seeded, none named explicitly — "could you do this to a civilization?" is the hinge	High	Proper escalation pacing	

Q21	Ethics beat delivery?	Nathan asks it himself — scene ends before internal answer is visible	High	Claude addendum; determines whether Nathan reads as strategist or philosopher	
Q22	AI assigned to draft scene?	Name before ONESTAMP is applied	Req.	Claude addendum	

5. Project 3 — Pathogen Outbreak Tabletop Exercise

5.1 Key Development Items Summary

From OAL-DOSSIER-TTX-001 v1.0.0, five blocking author decisions (D1-D5) remain:

#	Item	Status
D1	Lock the pathogen archetype	UNRESOLVED
D2	Lock the primary objective	UNRESOLVED
D3	Define target audience and scale	UNRESOLVED
D4	Establish sponsoring organization and authority (with named improvement plan owner)	UNRESOLVED — STRENGTHENED per Claude/ChatGPT
D5	Define the evaluation baseline	UNRESOLVED

5.2 Research & Analysis: The "10-Year Outbreak Cycle" Claim

THIS FRAMING IS RETIRED. Research is unambiguous:

- Duke Global Health Institute (PNAS, 2021): probability of a pandemic with COVID-19-scale impact is approximately 2% per year — growing, not static.
- Airfinity (2023): 27.5% chance of COVID-19-magnitude pandemic by 2033.
- Metabiota / Center for Global Development (2021): 22–28% chance within 10 years; 47–57% within 25 years.
- Nature (2022): Climate change will produce approximately 4,000 cross-species viral sharing events by 2070.
- BMJ Global Health (2026): 301 pandemic-prone outbreaks occurred worldwide in 2024 alone.

No reputable epidemiological study identifies a fixed 10-year interval as a predictive variable. Outbreak timing is driven by zoonotic spillover pressure, human behavior, surveillance capacity, and pathogen-specific factors.

Confirmed replacement framing: "We are in a heightened-risk window characterized by increasing zoonotic spillover pressure, degraded institutional memory post-COVID, and accelerating climate-driven pathogen expansion. The next major outbreak is not a question of if; it is a question of whether our systems have been rehearsed."

5.3 Research & Analysis: Pathogen Archetype Selection

The cross-AI recommended archetype — novel zoonotic respiratory with variable 2–10 day incubation, negative standard panels, 15–20% severe outcomes — is the most evidence-supported choice:

- 75% of emerging infectious diseases in humans are zoonotic in origin
- 95.54% of airborne viral EID outbreaks with identified sources are zoonotic spillovers
- SARS, MERS, COVID-19, and H1N1 all fit this archetype — it produces the largest preparedness failures

Distinctive feature to prevent COVID-reflex (Confirmed from v1.0.0): Early distinctive dermatological presentation — not typical of COVID, SARS, or flu — that creates genuine diagnostic uncertainty and prevents participants from defaulting to COVID-era protocols.

5.4 Research & Analysis: Provider Clinical Uncertainty

The residency hook is scientifically grounded. A 2022 study in *Annals of Emergency Medicine* documented emergency medicine residents during COVID-19 experiencing cognitive overload from constant protocol updates, PPE shortages forcing non-standard staffing, and significant psychological effects under uncertainty. Research from MedEdPORTAL (2023) confirms a documented gap in medical education specifically around navigating clinical uncertainty.

Ground rule (bolded per Grok team consensus): This exercise is testing the system, not evaluating individuals.

This must be stated in the ground rules, credibly and not performatively. Hospital administration participants may read the residency framing as implicit institutional criticism — the facilitator must proactively address this.

5.5 Research & Analysis: TTX Design Standards

The proposed TTX architecture is fully aligned with HSEEP standards. The CREID Network (*Frontiers in Public Health*, 2025) published a peer-reviewed description of a "Pathogen X" TTX validating the multi-phase approach. BMC Public Health research confirms that TTXs consistently reveal the same five challenge areas — all mapped to the 16 injects already designed.

Research Verdict: The exercise architecture is professionally grounded. What is missing is author decisions on D1–D5 and the named ownership structure required by D4.

5.6 Research & Analysis: Equity Layer

Research is unequivocal:

- 57% of 2024 global disease outbreaks occurred in sub-Saharan Africa and Latin America/Caribbean
- Unhoused populations, non-English-speaking communities, and long-term care residents are consistently identified as highest-vulnerability groups

Claude addendum (Confirmed): The equity injects are not designed — they are placeholders. This creates the risk they are deprioritized during exercise development. The D1 matrix entry now

requires both equity injects to be fully designed concurrently with Phase 1 clinical vignettes, not deferred.

Two equity injects (design required before exercise date):

- Non-English-speaking patient cluster presenting simultaneously with suspected cases — tests communication, informed consent, and notification under uncertainty
- Long-term care facility transfer request with residents whose care needs cannot be safely disrupted by isolation requirements

5.7 Confirmation Bias Flags — Project 3

FLAG 1: The Most Important Question Is Ownership (Elevated — Claude) "The most important question for Project 3 is not 'should we run this exercise?' but 'who will own the improvement plan and who will fund it?'" This is D4. Research on TTX follow-through is unambiguous: recommendations without assigned owners, budgets, and deadlines are routinely filed and forgotten. D4 as strengthened in v1.1 requires a named improvement plan owner with budget authority identified before the exercise date is set — not as a post-ONESTAMP task.

FLAG 2: Provider Hook May Create Administration Defensiveness Addressed in ground rules. The facilitator's credibility is the delivery mechanism — not just the written language.

FLAG 3: Single Hospital Scope Underrepresents the Real Failure Mode The most catastrophic outcomes in real outbreaks occur at the inter-institutional level. Phase 3 must include at least one inter-institution coordination inject even in the initial exercise design.

5.8 Author Decision Matrix — Project 3

#	Decision	Recommended Answer	Con f.	Research Basis	Your Decision
D1	Pathogen archetype	Novel zoonotic respiratory; 2–10 day incubation; negative standard panels; 15–20% severe; add distinctive early skin manifestation to distinguish from COVID	High	Cross-AI consensus + epidemiology literature	

D2	Primary objective	"Evaluate how frontline clinical teams detect, escalate, and make containment decisions during the first 72 hours of an uncharacterized pathogen outbreak, before any formal guidance or diagnostic confirmation exists."	High	HSEEP framework + Claude's original language	
D3	Target audience and scale	Single or dual hospital + county public health as primary; federal agencies (CDC, HHS) as simulators only	High	4 of 5 AIs + TTX depth-over-breadth research	
D4	Sponsoring organization	Named hospital coalition + named county public health authority; named improvement plan owner with budget authority identified BEFORE exercise date is set — this is a hard prerequisite	High	Claude + ChatGPT + TTX follow-through research	
D5	Evaluation baseline	Explicitly a baseline-establishing exercise; standardized scoring rubric reusable for follow-on exercises	High	Benchmarking research	
Equity injects	Design timeline	Both equity injects fully designed concurrently with Phase 1 clinical vignettes — not deferred	High	Claude addendum	
"10-year" framing	Retire it	Replace with heightened-risk-window framing (text above)	Very High	Airfinity, Metabiota, Duke, Nature — unanimous	
Inter-institution inject	Add to Phase 3	At least one inter-institution coordination inject in Phase 3	High	Outbreak research on systemic failure modes	
AI assigned to develop?	Name before ONESTAMP	Name the AI(s) assigned to inject design and facilitator guide development	Req.	Claude addendum	

6. Cross-Project Themes & Unified Recommendations

Three themes appear consistently across all three projects:

Theme 1: The Omnipotence Problem

All three projects carry the risk of Nathan Lewis as a flawless, all-knowing architect with no visible costs or limitations. Research and narrative analysis unanimously identifies this as the single greatest threat to long-term story value.

Unified Recommendation: Each project contains at least one designed failure mode or cost:

- Canon-001: The 0.4-second anomaly in Beta's cycle (Q21) — handled with invisible control
- Canon-002: The recursive counter-simulation seed — one attacker's data Nathan doesn't fully understand
- TTX-001: The exercise must be authorized to produce failure findings; the facilitator must not rescue participants from wrong answers

Theme 2: The Information Asymmetry Principle

All three projects are most powerful when different parties have different information:

- Canon-001: The colleague knows something Nathan doesn't realize they know
- Canon-002: The Architect has full visibility; the partner agency receives only what he packages
- TTX-001: Participants have different data; some requests return "unavailable"; audience/facilitator holds information participants don't

This asymmetry is what makes each project work. See Section 12 for the full Information Asymmetry Index.

Theme 3: The Follow-Through Problem

All three projects are currently designed plans with no owner. The most important action today is assigning responsibility for the next phase. See Section 13 for AI assignments as part of the approval checklist.

7. Cross-Project Dependency Architecture

New section — required by ChatGPT (Issue C1), Copilot (Addition B), Claude (CL6)

The three projects are not parallel — they are sequential with interdependencies. Understanding what each project unlocks for the others is required before any ONESTAMP can be responsibly applied.

7.1 Chronological Placement (Canon Timeline)

Project	Canon Placement	Relative to Theft
Project 1 — Kinetic Demo	Nathan at peak public-facing power, pre-confrontation	~3 years before (pending Q1 author decision)
Project 1 — Ocean Meeting	Intelligence about surveillance received	6–12 months before theft (or post-theft — pending Q10 author decision)
Project 2 — Network Defense	Nathan's defensive architecture at full capability	Likely post-theft; exact position depends on Q10 resolution
Project 3 — Pathogen TTX	Real-world institutional exercise; exists within the OAL organizational universe	Contemporaneous or near-future; may overlap with or follow Project 2 arc

Critical Dependency — Copilot Flag Resolved: Project 1 Ocean Meeting cannot be placed 6–12 months pre-theft AND have Project 2 imply Nathan is already operating at national cyber defense scale simultaneously — unless the national-scale operation was built before the theft and the theft accelerates its deployment. Author must decide whether the defensive architecture predates or postdates the theft. This is now a formal dependency decision (DEP-1).

7.2 Unlock Logic

Project 1 unlocks for Project 2:

- Nathan's demonstrated weapons architecture establishes the public-facing persona that makes his behind-the-scenes cyber operation believable
- The government surveillance established in the Ocean Meeting is the direct threat that motivates Nathan to build the Mirror Core system
- The synthetic control technology (Q5) establishes the underlying architecture that Project 2's autonomous system relies on

Project 2 unlocks for Project 3:

- The Mirror Core's behavioral-data capture creates institutional knowledge about large-scale adaptive systems — directly applicable to outbreak response frameworks
- Nathan's "could you do this to a civilization?" ethics beat is the thematic bridge to Project 3's "testing the system, not evaluating individuals" framing

- The government relationship established in Project 2 (joint task force, black-program contractor) is the institutional access path that makes Project 3's sponsoring authority plausible
- First Memory / Time Wars Threading:**
- Project 1: The synthetics' control mechanism is a proto-form of First Memory interface technology — this thread is planted, not named
 - Project 2: The recursive counter-simulation failure mode (Q17) is the first appearance of a temporal deception dynamic — an adversary feeding false data across time windows — which escalates into Time Wars
 - Project 3: The "testing the system" framing is the civilian/institutional mirror of the same question the Time Wars arc asks at civilizational scale — do you prepare for failure before it arrives?

7.3 Dependency Decision Registry

#	Decision	Impacts	Required Before
DEP-1	Does the national-scale defensive architecture predate or postdate the theft?	Project 1 Q10, Project 2 placement, Mirror Core motivation	ONESTAMP on either project
DEP-2	Does the government retain the synthetic from the demo (Q20)?	Project 2 operating authority, government-Nathan relationship	ONESTAMP on Project 1
DEP-3	Is the Mirror Core later turned against Nathan (Q16)?	Time Wars escalation, Project 3 institutional framework	ONESTAMP on Project 2
DEP-4	Does Project 3's sponsoring authority overlap with Nathan's Project 2 government relationship?	Project 3 D4, cover identity / compartmentalization (Copilot CP3)	ONESTAMP on Project 3

8. Omnicellular Architecture Reference Layer

New section — required by ChatGPT (Issue C2), Gemini (G1, G2, G3)

8.1 Placeholder Reference Statement

Per ChatGPT's requirement: **Omniceellular Architecture v2 governs all adaptive survival and system-level processing behaviors referenced in this dossier.** Until the author ratifies the full Omnicellular v2 specification, all references to extreme survivability, adaptive processing, and system-level control in Projects 1 and 2 operate under this placeholder.

Without this layer: effects exist without a cause. The kinetic armor's energy management, the synthetic humanoids' adaptive behavior, and the network defense system's real-time adaptive cycling all imply Omnicellular-level processing. That implication must be acknowledged even if the architecture is not yet fully specified.

8.2 Ontological Anchoring (Gemini — G1)

The Omnicellular Loophole Addressed:

The "no-switching architecture" defense works against sequential attacks — an adversary cannot exploit a transition window if there is no transition. However, it does not by itself explain survival against an area-of-effect negation aura (such as Nullarch's) that does not target transition windows.

Recommended Codification: Ontological Anchoring — the simultaneous existence of all 8 architectures provides an eight-fold existential weight that a single-source negation cannot instantly erase. A negation aura that encounters Omnicellular does not find one architecture to negate; it finds eight simultaneous anchors, each requiring independent negation.

Additionally, redefine "failure" not as physical destruction but as Bandwidth Throttling — where the system redirects all resources to a single survival architecture, temporarily rendering the entity immobile or cognitively narrowed. This creates a meaningful failure state that is not death but is genuinely consequential.

Author decision required: Ratify or modify the Ontological Anchoring mechanism before it enters canon.

8.3 First Memory Encoding Mechanism (Gemini — G2)

The most challenging of the Tier 1 questions is how a subjective euphoric experience became transmittable, cross-substrate data compatible with both biological and digital entities.

Recommendation from Gemini: Ground the encoding mechanism in infrastructure already established — connect it to the deep-research capabilities of ONE Industries or the Paragon Cubes. If the memory behaves like a quasi-sentient virus, define its cost explicitly: what does a

civilization lose when their cognition is artificially and instantly accelerated? Defining that cost gives the Puritans a rational reason to oppose it, making the conflict more compelling.

Author decision required: Lock the First Memory encoding mechanism and define the civilizational cost.

8.4 Nullarch's Superior (Gemini — G3)

This entity sits below THE MINUS ONE but above Nullarch, with operational authority to deploy Nullarch. If the Superior deploys negation into public spaces, their agenda must be defined. The cosmological alignment is tricky: THE MINUS ONE is not evil but a necessary principle of void and negation.

Recommendation: Define whether the Superior is executing a perfectly aligned cosmological function (reminding immortals that death still exists) or possesses a localized, flawed, or arrogant agenda. Deciding this prevents the Superior from functioning as a generic antagonist and keeps them grounded in the established cosmology.

Author decision required: Define Nullarch's Superior's relationship to THE MINUS ONE's intent.

9. Failure Cascade & Recovery Logic

New section — required by ChatGPT (Issue C3)

For each project, single failure, compound failure, and recovery pathways are now defined. This is the difference between a demonstration system and a real system.

9.1 Project 1 — Kinetic Armor Demonstration

Single Failure — Beta's cycle initiates late (0.4 seconds):

- What happens: The third shot arrives before the energy redirection fully cycles; the round is partially absorbed but not redirected
- Immediate consequence: Beta takes partial kinetic impact; suit holds but Beta moves visibly backward
- Nathan's response: No change in expression; he observes calmly; the override option is not used because Beta is functional
- Recovery: The anomaly is documented; Nathan uses it in his post-demo debrief as proof that the system performs under real variance, not scripted conditions — this makes the demo more credible, not less

Compound Failure — Override fails AND cycle is late:

- What happens: The anomaly occurs AND Nathan's manual override interface fails to respond
- Immediate consequence: Nathan cannot intervene; Beta takes the partial impact; the audience sees Nathan's hand move and then stop
- Recovery: This is the canonical plant. The audience member who notices Nathan's hand movement is the colleague, or an intelligence observer. This beat creates the surveillance thread.

Recovery Pathway: The armor performs above acceptable threshold even under compound failure. The demonstration succeeds on mission. The failure is recorded, studied, and becomes a design improvement input — positioning Nathan as someone who learns from variance rather than hiding it.

9.2 Project 2 — Network Defense

Single Failure — Synthetic layer intercept delay:

- What happens: One attacker's probe reaches a real system component before the synthetic layer spins up
- Immediate consequence: A real log entry is created on legitimate infrastructure before the synthetic environment captures the traffic
- Recovery: Nathan's system detects the real log entry and retroactively flags it; the attacker is diverted before any meaningful exploitation; no exfiltration occurs

Compound Failure — Intercept delay + innocent traffic spike + authority questioned:

- What happens: The intercept delay creates a real exposure, simultaneously civilian traffic spikes (legitimate users hitting the same affected component), and the partner agency's legal team flags the exposure as a potential CFAA issue
- Immediate consequence: Nathan is managing three simultaneous threads; the system's autonomy creates a momentary gap he cannot close manually
- Recovery: Nathan invokes the pre-negotiated emergency containment protocol; all traffic is temporarily diverted to a buffer zone; the partner agency's legal flag is resolved within the authority framework; the civilian traffic is released without exposure; the attacker's probe is captured but the behavioral data shows an anomaly Nathan doesn't recognize

The unrecognized anomaly is the recursive counter-simulation seed (Q17).

Recovery Pathway: The compound failure is contained. The legal risk is managed. The system performs. But the one data artifact Nathan cannot explain becomes the thing he thinks about afterward. This is the scene's most important beat.

9.3 Project 3 — Pathogen TTX

Single Failure — Participants correctly identify the pathogen archetype early:

- What happens: A participant with recent outbreak training correctly identifies the dermatological presentation pattern and names the archetype before Phase 2
- Immediate consequence: The exercise's uncertainty window collapses; participants shift from novel-pathogen mode to protocol-execution mode
- Recovery: The facilitator introduces a complication inject — lab confirmation returns a false negative on the initial rapid test; uncertainty is re-established; the exercise continues

Compound Failure — Facilitator rescues + communication channels overwhelmed + equity injects deprioritized:

- What happens: The facilitator intervenes to prevent participants from reaching an incorrect conclusion, simultaneously the simulated communication channels are overwhelmed and the equity injects (non-English speaking cluster) are set aside "for later"
- Immediate consequence: The exercise produces a sanitized finding set that reflects best-case performance rather than realistic performance
- Recovery: The AAR must explicitly flag the facilitator rescue and the equity inject deferral as findings, not as exercise artifacts. These are the most important findings the exercise can produce.

Recovery Pathway: Exercise failure is exercise success if the findings are honestly recorded and the improvement plan owner has budget authority to act on them (D4).

10. Architect Identity & Presence Model

New section — required by ChatGPT (advisory), Copilot (addition A)

10.1 The Problem This Solves

Across all three projects, Nathan Lewis appears as: a weapons architect (Project 1), a cyber defense architect (Project 2), and a public health strategist / covert intelligence actor (Projects 2 and 3). Without a unified identity framework, this risks making him unbelievable unless all roles are contained within a single coherent self-conception.

The unifying frame: Nathan Lewis is an Architect of Systems.

He does not think of himself as a weapons developer, a cyber operator, or a public health consultant. He thinks of all three domains as instances of the same problem: how does a complex system behave under pressure it was not designed to handle? His weapons work, his cyber work, and his institutional exercise work are all expressions of a single question. This is not a cover identity — it is his actual worldview.

10.2 When Nathan Is Present vs. Absent

Mode	Description	Examples
Architectural	Nathan designed the system; it runs without him	Network defense system during active intrusion; the armor on Beta; the TTX exercise framework
Oversight	Nathan observes but does not intervene	Most of the armor demonstration; the network defense exercise as it runs
Intervention	Nathan acts because something is outside design parameters	The manual override (not used in normal circumstances); the emergency containment protocol in compound failure; the facilitator injection in TTX
Absence	Nathan is not present; the system acts autonomously	The Mirror Core after transfer; the synthetics between demonstrations

The "single Architect" narrative problem is resolved by this model. Nathan built a system that does not require him. His presence is a choice, not a necessity. This makes his intervention moments more meaningful — he is not managing the system; he is trusting it, and choosing when that trust requires revision.

10.3 Architect Presence in Each Project

Project 1: Nathan is present in architectural + oversight mode during the demo. His manual override is available. He does not use it. This is a character statement.

Project 2: Nathan is present in oversight mode. The system runs itself. His intervention in the compound failure scenario is triggered by the unrecognized anomaly, not by system failure — he intervenes because he is curious, not because he is required to.

Project 3: Nathan is present in architectural mode only. The TTX framework is his design. He is not in the room during the exercise. His presence is the framework itself.

11. Ethical Boundary Framework

New section — required by ChatGPT (advisory), Copilot (addition D)

11.1 The Architect Ethical Doctrine

All three projects raise ethics questions that cross between projects. Copilot correctly noted that Project 1 ethics (corporate weapons), Project 2 ethics (cyber defense), and Project 3 ethics (public health) must be unified under a single Architect Ethical Doctrine. The following is a proposed framework for author ratification.

What Nathan Will Not Do (Proposed — Pending Author Decision):

- Deploy any system against a population that cannot, even in principle, consent or know it is subject to the system
- Use First Memory or its derivatives to alter memory without the subject's knowledge
- Allow a system he designed to be used for purposes he did not authorize without activating the Mirror Core's traceable/corruptible properties
- Remain silent when collateral harm (innocent traffic, civilian exposure, non-English-speaking patients in an overwhelmed ED) is being systematically ignored by an institutional partner

What He Is Willing to Risk:

- His political standing, when the alternative is allowing a preventable harm
- The operational security of a project, when maintaining it requires the above prohibited acts
- Being wrong about a design decision — and documenting that he was wrong, publicly if necessary

What Crosses the Line:

- Using the Mirror Core against a domestic civilian target without legal authority
- Allowing the recursive counter-simulation adversary to operate unidentified because naming them creates political complications
- Designing an exercise that produces only the findings the sponsoring institution wants to see

Author decision required: Ratify, modify, or replace this framework before ONESTAMP is applied to any ethics-bearing scene.

11.2 Ethical Boundaries by Project

Project 1: Using synthetics as weapons demonstrators. Nathan's ethical position — he believes the synthetics are not sentient at the time of the demonstration; he monitors this belief actively; if his monitoring produces contrary evidence, he will stop. The ethics of Q20 (post-demo synthetic retained by government) tests this: does Nathan object? If not, why not?

Project 2: Collateral capture of innocent traffic and the "could you do this to a civilization?" question. The contain-only protocol is Nathan's ethical limit in this project. It is also, per the

Chekhov's Gun framing of the Mirror Core, a limit that will eventually be tested.

Project 3: The exercise must produce honest findings. Nathan's ethical line here is the most institutional and the least dramatic — but it is the one most likely to be violated in real-world TTX execution, which is what makes it the most important.

12. Information Asymmetry Index

New section — required by ChatGPT (advisory)

This index maps information distribution across all three projects. It is a tool for the author and AI team to prevent narrative gaps where someone knows something they shouldn't, or doesn't know something they must.

12.1 Project 1 — What Each Party Knows

Party	Knows	Does Not Know
Nathan	Full armor capability; full synthetic control; that his colleague is attending; Ocean Meeting purpose	Whether the colleague reports the meeting; whether the government already has what he's about to receive
Military Acquisition Officers	They are watching a new APS prototype; Nathan's name and company	The three-shot protocol's full mechanics; the synthetic control method; that the armor has a failure mode
The Colleague	That Nathan is being surveilled; that the meeting is being monitored	Whether Nathan suspects; what Nathan will do with the information
Government Observer	That the meeting is happening	What changes hands
The Audience (armor demo)	What they see	That Nathan can override; that Beta had an anomaly; that the synthetics might be proto-sentient

12.2 Project 2 — What Each Party Knows

Party	Knows	Does Not Know
Nathan	Full system architecture; attacker behavioral data; the Mirror Core's traceability	The identity behind the unrecognized attacker anomaly; whether the anomaly is a fingerprint or a message
Partner Agency	Operational capability (Mirror Core); that the system works	The architectural knowledge; that the Mirror Core is traceable; that Nathan kept the behavioral data
Attackers	They are operating in what they believe is real infrastructure	They are in a synthetic environment; they have been captured and studied; their behavioral data has been archived
The One Attacker with the Anomaly	They may know they are in a honeypot	They know something about the 3-second cycle signature; how much they know is the author's decision

12.3 Project 3 — What Each Party Knows

Party	Knows	Does Not Know
Facilitator	Full scenario; all inject outcomes; which participants are performing well vs. poorly	How participants will respond to the equity injects specifically
Clinical Participants	What their role in the scenario is; their normal protocols	The pathogen identity (intentionally); the full scope of the outbreak; what other teams are doing
Hospital Administration Participants	Their resource authorization levels; the exercise's stated objectives	That the facilitator may let them fail; that the findings will be formally reported
Public Health Simulators (CDC, HHS)	The scenario from a population-level view	What the clinical teams are experiencing on the ground; the equity inject details
Equity Inject Participants (if live actors)	Their role and language barrier scenario	The extent to which the clinical teams are prepared to handle them

13. Priority Approval Checklist — What to Decide Today

The following is the complete approval document. Review each item, mark your decision, and apply ONESTAMP to the approved elements. Items marked **[REQUIRED BEFORE ONESTAMP]** cannot be deferred.

Backbone Architecture Decisions (Sections 7–12)

- ☐ DEP-1 [REQUIRED]: Does the national-scale defensive architecture predate or postdate the theft?
- ☐ DEP-2: Does the government retain a synthetic from the demo?
- ☐ DEP-3: Is the Mirror Core later turned against Nathan?
- ☐ DEP-4: Does Project 3's sponsoring authority overlap with Nathan's Project 2 government relationship?
- ☐ Omnicellular v2 [REQUIRED]: Ratify or modify Ontological Anchoring and Bandwidth Throttling mechanisms
- ☐ First Memory encoding [REQUIRED]: Lock encoding mechanism and define civilizational cost
- ☐ Nullarch's Superior: Define relationship to THE MINUS ONE's intent
- ☐ Architect Ethical Doctrine [REQUIRED]: Ratify, modify, or replace the proposed framework

Project 1 — Canon-001

- ☐ Q1: Confirm 3 years before theft (or specify)
- ☐ Q3: Confirm audience: military + board + ethics observer (or specify)
- ☐ Q5: Confirm pre-scripted AI with override (or specify)
- ☐ Q9: Confirm joint corp-gov classified (or specify)
- ☐ Q10 [REQUIRED]: Decide timeline placement — pre-theft or post-theft
- ☐ Q12: Confirm Black-Ops / Military Liaison with defined motive, cost, leverage, risk reason, silence reason
- ☐ Q17: Confirm surveillance proof as payload (or specify)
- ☐ Q18: Confirm relational cost as the outcome mechanism
- ☐ Q20: Confirm post-demo synthetic disposition
- ☐ Q21: Confirm the 0.4-second anomaly beat
- ☐ Q22 [REQUIRED]: Name AI(s) assigned to draft narrative scenes

Project 2 — Canon-002

- ☐ Q1: Confirm full US critical infrastructure scope (or specify)
- ☐ Q3: Confirm joint task force, black-program contractor (or specify)
- ☐ Q10: Confirm 3-second cycle as tunable signature / Chekhov's Gun
- ☐ Q11: Confirm contain-only, no retaliation (or specify)

- [] Q14: Confirm Mirror Core transfer with traceable / corruptible / dangerous / plot-hook properties
- [] Q17: Confirm recursive counter-simulation as canonical failure mode (planted as seed)
- [] Q20: Confirm "all seeded, none named" as Time Wars connection approach
- [] Q21 [REQUIRED]: Confirm ethics beat delivered by Nathan to himself, scene ends before internal answer
- [] Q22 [REQUIRED]: Name AI(s) assigned to draft Command Center POV and Attacker POV scenes

Project 3 — TTX-001

- [] D1: Confirm novel zoonotic respiratory archetype with skin manifestation
- [] D2: Confirm primary objective (recommended text or modify)
- [] D3: Confirm single/dual hospital + county public health scale
- [] D4 [REQUIRED]: Name sponsoring organization AND name improvement plan owner with budget authority — this is a hard prerequisite for setting an exercise date
- [] D5: Confirm baseline-establishing status and scoring rubric
- [] Equity injects [REQUIRED]: Confirm both injects will be fully designed concurrently with Phase 1 vignettes
- [] Inter-institution inject: Confirm Phase 3 includes inter-institution coordination inject
- [] AI assigned [REQUIRED]: Name AI(s) assigned to inject design and facilitator guide development

14. Version History

Version	Date	Author	Changes
v1.0.0	April 11, 2026	Perplexity Computer	Initial comprehensive research and analysis document; all three projects covered
v1.1	April 11, 2026	Perplexity Computer	Integrated full team review feedback: ChatGPT (4 critical issues), Claude AI (7 additions), Google Gemini (3 canon consistency notes), Microsoft Copilot (3 removals + structural additions), Grok AI (3 polish items); Added Sections 7–12; Updated all decision matrices; Removed 3 items per Copilot; Strengthened D4, Q10, Q18, Q21, Q22 entries; Perplexity reconciliation notes included throughout

Document ends. Awaiting ONESTAMP from author before any element is treated as approved direction.

Prepared by Perplexity Computer — April 11, 2026

Revised based on team feedback: ChatGPT · Google Gemini · Claude AI · Microsoft Copilot · Grok AI

Distribution: All six AI team members + Author