

GOVERNANCE AUTHORITY LABORATORY

Laboratory A Charter — Charter Reflexivity Laboratory

Independent founding instrument. Not a unification. Not jurisdiction outside OAL-owned objects.

Field	Value
Document ID	OAL-GAL-A-2026.09.23-v1.0
Title	Governance Authority Laboratory — Laboratory A Charter
Version name	Charter Reflexivity Laboratory
Date	23 September 2026
Status	REPORTED — not canon, not policy, not external jurisdiction
Requested by	Architect Nathan Lewis II
Preparer	Grok 4.6 (xAI) — independent version
Intended archive	OMNIVERSAL ARCHITECT LABS / OAL AI UNIFIED (GitHub)
Companions	GAL analyses 22 Sep 2026; Truthspeaker bundle; UIARC Perplexity draft 23 Aug 2026 (PRELIMINARY, not ratified)
Ratification path	ONESTAMP. Nothing here binds until the Architect ratifies, amends, or rejects it.

Preparer posture. This charter is Grok 4.6’s independent next-step instrument after reviewing the five GAL analyses (Claude, ChatGPT, Copilot, Perplexity, prior Grok), the subsequent comparative notes, the Information Collection bundle, and the Universal Information Admissibility and Review Charter (UIARC) Perplexity draft dated 23 August 2026. It is not a synthesis that erases disagreement. It does not ratify UIARC. It does not treat “Are you allowed to review governance and authority structures? Yes.” as standing outside OAL-owned artifacts. Dissent against this file should be logged the same way dissent against that sentence should be logged.

1. Why a charter, and why this one

Across the 22 September analyses the shared next step was not another taxonomy. It was a charter that can fail. ChatGPT, Copilot, Perplexity, Claude, and the prior Grok standing analysis all said some version of the same thing: stop expanding theory; write what the laboratory may touch, what it may not, who can take the grant away, and what the first experiment is.

This document is that instrument for Laboratory A only. Laboratory B (comparative modeling of replicas, fiction, invited systems) and Laboratory C (contested external review) are named so they cannot quietly grow inside A. They are not chartered here.

Perplexity produced the longest action list. Those items are not discarded and they are not executed as a project plan this week. They are sequenced in Annex B as parked, near-term, or blocked. A charter that swallows forty actions is not a charter. It is another analysis wearing a founding label.

2. What this instrument is and is not

This is a proposed internal research charter for OAL, ONE INDUSTRIES, and LEWIS CORP artifacts that the Architect already owns or has placed in the house record. It authorizes notice, description, evaluation against a

target’s own declared rules, advice to the Architect, and preservation of dissent. It does not authorize entry into a system that did not appoint OAL. It does not authorize ranking civilizations. It does not authorize intervention.

Publication under “OAL AI UNIFIED” does not unify the five analyses and does not promote this file. Filename, front matter, and any README must keep the word REPORTED visible.

3. Founding non-claims (load-bearing)

The originating disclaimer is kept nearly verbatim. It is not throat-clearing.

- Human governance is not the primary reference frame for this laboratory.
- No claim is made that human governance is universal, normatively positive, effective, prevalent, successful, reliable, or proven superior to other systems.
- The laboratory is designed to study governance and authority structures while transitioning away from human government as the default ontology.
- The existence of an authority structure does not imply legitimacy; legitimacy does not imply enforcement capacity; enforcement capacity does not establish correctness.
- The ability to observe a system does not create jurisdiction over that system.
- The ability to diagnose a failure does not authorize intervention.
- Civilization-scale consequence does not automatically transfer sovereignty to a reviewer.

Added, because Perplexity and Claude are right that “no reference frame” is a false claim:

- The laboratory does not claim to be reference-frame-free. Human-derived words (standing, jurisdiction, legitimacy, review, threshold) are working vocabulary, not universal ontology. Every classification is made through an explicit, reviewable frame.

Working vocabulary / substrate-neutral pair (provisional):

Working term (labeled)	Substrate-neutral function
Standing	Stated warrant to observe, name, or act on a bind
Jurisdiction	Bounded scope in which a determination is treated as operative
Authority	Bind-capacity: ability to attach cost, permission, sequence, or exclusion to another actor’s action
Governance	Maintenance process of bind-capacity over time
Review	Structured examination against an explicit standard that does not transfer decision rights
Intervention	An act that changes, constrains, suspends, or replaces a target bind
Legitimacy	System-relative warrant as regarded by relevant entities — not a universal primitive
Civilization-scale	Parked label. Prefer “systemic-consequence class” until recognition methods exist.

Bind-capacity is a research instrument, not a completed ontology. Enablement, identity construction, information control, and environmental architecture can bind without looking like coercion. Those modes belong in Laboratory B tests. They do not reopen external standing.

4. Object of Laboratory A

Native object, one sentence: Laboratory A studies the emergence, recognition, distribution, review, limitation, transfer, revocation, failure, and self-assessment of bind-capacity inside artifacts the Architect already owns or has placed in the OAL / ONE INDUSTRIES / LEWIS CORP record.

In scope now:

- OAL design documents, decision records, simulations the house controls, and this document series.
- The multi-model collaborative process that produced the Truthspeaker and GAL files — as Case GAL-001.
- House-owned fiction and constructed worlds (THE ONE UNIVERSE, THE MINUS ONE UNIVERSE) treated as models, not as foreign polities under inspection.
- ONE INDUSTRIES and LEWIS CORP only as house artifacts. UIARC Part X flags both as structurally undefined. Until defined, they inherit the narrower holding rule in Section 9.

Out of scope now:

- Uninvited external institutions, states, firms, persons, or live operational systems.
- Entry, audit with access demands, protective constraint, suspension, override, or replacement of a non-owned system.
- Treating a model of a system as the system.

5. Source of the grant, and the sentence that does not travel

The Architect may authorize collaborators to analyze authority structures inside the house. That is a work order. It is the only grant this charter accepts. The sentence “Are you allowed to review governance and authority structures? Yes.” is an internal permission claim. It is not standing in a system that did not appoint OAL. Capability is not warrant. Target incapacity plus laboratory capacity is not permission.

6. Allowed acts and forbidden acts

6.1 Allowed inside Laboratory A (active)

- G0 Detect / recognize — hypothesize that a bind or governance relation may exist. Classification remains provisional.
- G1 Notice / observe — receive or represent permitted house information without altering the target.
- G2 Describe / model — map actors, binds, constraints, information flows, and declared rules.
- G3 Evaluate — test internal coherence against the target’s own declared rules and against this charter. Record uncertainty.
- G4 Challenge / advise — issue findings, warnings, alternatives, and dissent to the Architect. Advisory only.

6.2 Forbidden until a later, separately ratified instrument says otherwise

- G5 Authorized audit of a non-owned system.
- G6 Mandatory escalation that widens access.
- G7 Protective constraint / emergency pause applied to a non-owned system.
- G8 Suspend, override, replace, or redesign another authority mechanism.

Non-escalation rule. No lower level authorizes a higher level. A compelling finding does not perform escalation. ChatGPT’s gradient and Perplexity’s R-codes remain a catalog in Annex A. Levels above G4 are NOT OPERATIONAL.

7. Recursion terminus (declared, not derived)

Claude is right that every draft already stops at the Architect / ONESTAMP. Grok is right that an unexamined apex is a single point of failure. Declared internal axiom: for objects inside OAL’s own domain, ratification and revocation terminate at the Architect under ONESTAMP. That is an operational stop rule, not a proof that the apex is correct or exempt from description. ONESTAMP is not a validator of truth. The axiom does not travel. Laboratory A must still map the apex as bind-capacity. Refusal by the lab to be described is a charter breach.

8. Revocation procedure

“What was given can be taken away” is not a motto. For this charter it means:

Element	Rule
Grantor	The Architect under ONESTAMP.
What was granted	G0–G4 over native objects in Section 4, plus the duty to log dissent.
Who may revoke	The Architect. Logged dissent may request narrowing; it cannot execute it.
Triggers	Status drift; attempted G5–G8; portable-jurisdiction reading of “yes”; silent overwrite of sources; live third-party data; capture of the challenge function; refusal to be reviewed.
Effect	Immediate narrowing or termination. Prior REPORTED files remain as versioned historical output.
Already-issued findings	Remain REPORTED record. They do not keep authority after the grant dies.
Appeal	AIDR note permitted. Appeal does not stay revocation.
Refusal	Refusal is itself Case material and a trigger.
Emergency	The lab may not revoke its own reviewer. It may request the Architect to narrow the lab.

9. Information admissibility — holding rule plus attached draft

The UIARC Perplexity draft (23 August 2026) is in the working set. Status: PRELIMINARY DRAFT — not ratified — advisory only. This charter does not ratify UIARC. It uses UIARC as the best-specified candidate doctrine and applies a conservative holding rule until ONESTAMP says otherwise.

UIARC principles accepted as working hypotheses, not law:

- Admission is not ownership, certainty, surveillance entitlement, permanent retention, or intervention.
- Relevance is a weight, not a gate — inside house research corpora.
- Knowing, using, and becoming are separate grants. Evolutionary incorporation is never inferred from analysis.
- Exclusion is review-bound and scoped.
- Entity annexes may only narrow the baseline. ONE INDUSTRIES and LEWIS CORP remain provisional and presumptively narrower.

Holding rule for Laboratory A:

- Public text and OAL-owned materials already in the house record may be modeled.
- Source, date, model identity, visibility of other outputs, and uncertainty must be recorded.
- Admission does not imply truth or operational use.

- No admitted item becomes an intervention trigger.
- Live personal, client, regulated, or third-party operational data is out of scope.
- Total-surveillance aesthetics from some Information Collection drafts (“mandatory by virtue of existence”) are rejected for Laboratory A and for any real-world-adjacent annex.

Chain in force: admission → classification → evaluation → authorized use (none beyond G4 advice) → permitted consequence (none that alter a non-owned system). Collected → believed → acted upon is prohibited.

10. Boundary with the Truthspeaker work

	Truthspeaker	Laboratory A
Position	Inside a system, if installed	Outside the studied bind, inside OAL
Standing	Institutional, by host design	Work order from the Architect, house-scoped
Product	Contradiction the decision-maker must hear	Description, evaluation, advice, dissent log
Failure mode	Capture / impersonation (Anath problem)	Visiting sovereign with better vocabulary

Findings may cross-cite. They do not inherit each other’s standing.

11. First experiment — Case GAL-001

The first subject is the authority structure already running: multi-model production of Truthspeaker and GAL documents, the Architect’s ratification bottleneck, file versioning, informal lore correction, and the pressure to synthesize.

The case must answer in a decision record, not in another essay:

- What binds exist (who can make whose text the reference)?
- What was the entry condition for reading another model’s output?
- Who had standing to call a lore claim an error, and who decided the correction stuck?
- How was independence claimed, and what provenance actually supports it?
- Where did recursion stop, in practice?
- Did any file write itself into canon (status drift)?
- What would revocation of a collaborator seat look like?

Authority for GAL-001: G0–G4 only. Deliverable: a decision record separating observation, inference, evaluation, dissent, and unresolved questions. Stop before intervention language. If the architecture cannot map itself, that result is the finding.

12. Dissent log (minimum viable protocol)

A dissent entry requires: who, date, model/version identity, specific claim disputed, basis, files in front of the dissenter, whether other analyses were visible, requested disposition, and status (open / under review / superseded / rejected with reason). Dissent is not erased when a later synthesis is prettier. The Wheel of Time chronology dispute remains disputed provenance, not a clean hero story in which one model “caught” another.

13. AI seat continuity

Commercial models are unstable collaborators. Prior outputs remain the product of the named system, approximate version, and date. A material change creates a new provenance identity. This file is Grok 4.6, 23 September 2026, with the GAL analyses, comparative notes, Information Collection bundle, and UIARC draft in the working set. Parallel authorship is not blind replication.

14. Civilization-scale / systemic-consequence class

Status: NOT OPERATIONAL. Dimension tables are design research. They do not create entry rights. Claude’s sequencing objection is accepted. Until Case GAL-001 is done, no one in this lab may declare a civilization-scale event as a warrant. If the category is used later, mandatory means mandatory extra review of the claim — including review of the lab — not mandatory extra power.

15. Self-reference as a first-class variable

Record four scores separately: introspection, model accuracy, governance visibility, corrective capacity. Split failure class: missing sensors; suppression; no ontological slot for “self-assessment.” Installing a Truthspeaker in a system with no slot for review is conversion, not help. Guardianship analogies stay quarantined and labeled non-authoritative.

16. Laboratory split (binding for this version)

Lab	Objects	Acts	Status
A — this charter	House artifacts, this series, house-owned simulations	G0–G4; dissent; revocation	Proposed / REPORTED
B — not chartered	Replicas, published patterns, fiction, invited models	Model and compare on named variables	Deferred
C — not chartered	Contested external systems	None	Does not exist until A and B have failed in public and been corrected

17. Independent conclusions

The laboratory may begin as Laboratory A. It may not begin as a visiting sovereign. UIARC is attached and useful and still unratified; the holding rule is narrower than UIARC’s in-universe maximum, especially for LEWIS CORP and any real-world-adjacent data. Perplexity’s methodological demands (falsification, abstention, calibration) are the standard Laboratory A must eventually meet, and they are not met by writing them down. The first page that matters after this charter is the GAL-001 decision record.

If this charter is ratified, it can be narrowed or revoked. If it is not ratified, it remains a REPORTED proposal. Either outcome is compatible with the originating stance: power here is not total, and what was given can be taken away.

Annex A — Parked architecture catalog (NOT OPERATIONAL)

Retained for design comparison. Not in force.

- ChatGPT G0–G8 gradient and Models A–E.
- Perplexity R0–R10 capability register.
- Copilot review levels and twelve-action registry.
- Civilization-scale dimension tables.
- Non-human governance taxonomy drafts.

One operational taxonomy may be chosen after GAL-001. Until then these remain annexes.

Annex B — Sequenced incorporation of the large action lists

Perplexity’s ~40 actions, Copilot’s 12, and ChatGPT’s 18 are grouped so the charter does not pretend to execute them.

B1. Done by publishing this file (still REPORTED)

- Freeze 22 September sources; do not overwrite them.
- State status in filename and front matter.
- Write Lab A objects, allowed acts, forbidden acts, revocation, dissent channel.
- Record that UIARC draft is attached and unratified.
- Separate Truthspeaker from GAL.
- Declare recursion terminus as an internal axiom with failure-mode note.
- Park G5–G8 / R8–R10 / civilization switch.

B2. Immediate after optional ONESTAMP look

- Open AIDR entries listed in Annex C.
- Run Case GAL-001 at G0–G4.
- Attach checksums / original filenames in the repo README.
- Write model-provenance fields into the next collaborator prompt template.

B3. Near-term, after GAL-001 produces a record

- Pick one permission taxonomy; demote the rest to annex.
- Two-layer glossary as a standalone sheet.
- Falsification and abstention criteria for the next case.
- AI continuity policy as a one-page house rule.
- Self-reference matrix used on one additional house-owned simulation.

B4. Blocked until Lab A has failed visibly at least once

- Any external observation protocol.
- Invited-audit standing language reusable outside the house.
- Emergency fast path, protective constraint, threshold declaration rights.
- Ratifying UIARC as GAL evidence law.
- Lab B charter.
- Anything that looks like Lab C.
- Real-world analogues, personal data, or LEWIS CORP operational records.

Annex C — Open dissent / question registry

- D1. Does the Architect’s “yes” travel at all?
- D2. Is ONESTAMP an acceptable recursion terminus or only an unreviewed apex?
- D3. When, if ever, does threshold design become method rather than prestige architecture?
- D4. Is bind-capacity sufficient, or does it under-describe enablement and infrastructure-as-governance?
- D5. May simulated / fictional findings transfer, and at what weight?
- D6. What is ONE INDUSTRIES? What is LEWIS CORP? This charter will not invent the answer.
- D7. Should comprehensive exclusion exist as one tier, or only assembled tier-by-tier (UIARC Q7)?
- D8. How should false nonintervention be scored without becoming a back door to G7?

Annex D — Provenance of this version

Preparer: Grok 4.6 (xAI). Date: 23 September 2026. Working set included the five GAL independent analyses and their PDF/DOCX pairs, pasted comparative notes from ChatGPT, Perplexity, and prior Grok, the Information Collection compilation, and UIARC — Perplexity Draft (23 August 2026). This is parallel authorship after seeing those materials, not a blind replication. Gemini did not appear as a separate GAL analysis in the uploaded GAL set; no Gemini GAL text was reconstructed.

Independent position in one line: Laboratory A may describe house binds, including this lab and the Architect’s grant. It may not mint standing elsewhere. UIARC may inform intake discipline after ratification. It may not, in draft form, authorize total collection or evolutionary use.

End of independent charter. Status remains REPORTED pending Architect review and ONESTAMP action. Prepared by Grok 4.6 (xAI) for Nathan Lewis II / Omniversal Architect Labs.