

THE OPEN WINDOW DOCTRINE

Canonical Merge

Governance-Locked Baseline — Minimum Viable Protocol

OAL Working Document — v3.0 — 2026-08-10

Status: GOVERNANCE-LOCKED BASELINE — pending Nathan's ONESTAMP

ONESTAMP authority: Nathan Lewis, with AI collaborators required to flag disagreement before ratification (per OAL-AIDR-v1.0)

Supersedes: OAL-Open-Window-Infection-Scenario-v1.0 (origin, preserved) → v2.0 (Claude) → Preliminary Draft v0.1 (ChatGPT) → Protocol Draft 0.1 (Grok) → inline response (Gemini)

Consolidation basis: independent red-team analysis by Perplexity AI, cross-checked against primary source documents

1. Governance

ONESTAMP authority: Nathan Lewis holds sole ratification authority over this document set. No content in this document, or in any future iteration of it, is canon-locked until Nathan applies ONESTAMP.

Dissent requirement: Before ONESTAMP is applied, all contributing AI collaborators (ChatGPT, Claude, Gemini, Grok, and any future contributor) must have the opportunity to flag disagreement with this merge, per the existing OAL-AIDR-v1.0 disagreement-logging framework. Disagreements are logged, not overruled by omission — silence from a collaborator is not treated as agreement.

Scope of this lock: This governance decision applies specifically to the Open Window Doctrine document set. It does not change ONESTAMP authority for any other OAL canon track.

This resolves the governance gap identified independently by both this collaborator's prior review and Perplexity AI's red-team analysis: three competing “final” merges existed with no defined sign-off authority. That gap is now closed.

2. Version Changelog — Retiring the Branch History

Five prior artifacts exist across this concept's development. None is deleted — all remain readable as historical record — but only this document (v3.0) is now the active working draft. Status column reflects what each contributed to this merge.

Document	Status	Disposition
v1.0 — Origin Scenario (Claude)	PRESERVED	Verbatim brainstorm and founding doctrine language retained as historical origin record. Not superseded — it is the root of the lineage.
v2.0 — Unified Draft (Claude)	SUPERSEDED	Five merged scenarios (A–E) retired in favor of Section 6 below. Time-discrepancy analysis (§4.1) and root-level failure argument (§5) preserved and moved to Deferred — Iteration 2 (Section 8).
Preliminary Draft v0.1 (ChatGPT)	PARTIALLY ADOPTED	Test matrix (T1–T12), hypotheses (H1–H9), and metric set (TTE/TTD/TTA/TTC/FCR/CTC/PD-PB/TER/ECR/RTR) adopted as the canonical baseline throughout this document. Funding tiers not adopted (see §3).
Protocol Draft 0.1 (Grok)	PARTIALLY ADOPTED	Window Object data model adopted as the object-level schema (Section 4.2). 14-day campaign structure rejected in favor of the locked lean prototype (Section 3). “Partially unaware Layer 0 operators” proposal flagged, unresolved, not scheduled (Section 9).
Inline response (Gemini)	PARTIALLY ADOPTED	Physical countermeasure vocabulary retained as flavor language only. Literal biological/physical-harm framing explicitly not adopted pending the open decision in Section 7.

Note on source completeness: this merge uses Gemini's actual primary response text (available in this conversation) rather than the secondhand paraphrase of Gemini that Perplexity's analysis had to rely on, since no standalone Gemini file was ever exported.

3. Locked Session Baseline (Phase 2 Decision)

LOCKED: ChatGPT's lean 2-day / ~12-active-hour prototype is the executable baseline. Grok's 14-day campaign and Gemini's 168-hour structure are not run at this stage.

Structure adopted from the ChatGPT draft, unchanged in shape:

- Pre-lab setup (offline): 4–8 staff hours before the session — build snapshots, seed native clocks, verify telemetry, define benign markers, pre-register expected outcomes.
- Day 1 — Baseline and visible-interface tests: ~6 hours — baseline cycling, black-screen tests, content-state bleed, awareness comparison.
- Overnight persistence window: 8–12 unattended hours — no participant intervention unless a safety/containment threshold trips.
- Day 2 — Closure and reconciliation tests: ~6 hours — total-event-closure reconciliation and recovery.
- Post-run analysis: 1–2 analyst days — merge logs, normalize time separately from native time, score hypotheses.

Deliberately excluded from this baseline: carrier transition, time divergence, unknown-vector, interface airlock formal testing, community escape, and root trace. These require the clean data this run is designed to produce first — see Section 8.

4. Canonical Architecture

4.1 Layer Diagram

ChatGPT's L0–L4 taxonomy is adopted as canonical. This choice is not arbitrary — the locked session baseline (Section 3) and the test matrix (Section 6) were built against this taxonomy, and T9 (Interface Airlock) specifically requires the airlock to be a separable, testable layer rather than a cross-cutting function. Grok's Layer 0/Layer 1/Layer 2–N language is retired as a competing top-level taxonomy but its Window Object schema is retained below at the object level, where it does not conflict.

Layer	Function
L0	Oversight / Root Observer — authoritative simulation ledger, immutable experiment log, snapshots, emergency rollback, researcher telemetry. No participant has direct write access.
L1	Primary inhabited environment — four-bedroom residential simulation with multiple agents and ordinary devices.
L2	Media / application environments — content spaces, feeds, calls, games, recordings, synthetic social platforms.
L3	Adjacent household / community environments — used only after local containment tests pass; out of scope for the locked baseline.
L4	Sacrificial interface airlock — translation/quarantine layer through which unknown cross-environment state is forced before admission to protected environments.

4.2 Object-Level Schema (from Grok, retained)

Every device is instantiated as a Window Object with explicit properties, sitting inside the L0–L4 stack above rather than defining its own competing layer model:

- Aperture state: isolated / listening / transmitting / receiving / bidirectional / latent / unknown
- Presentation state: active / idle / black / minimized / sleep / apparent-off
- Power state: powered / low-power / suspended / simulated hard-off
- Network state: connected / disconnected / local-only / alternate-path
- Closure state: open / closure requested / closure acknowledged / closure verified / residual detected
- Directionality, residual process, entity affinity, local clock, and bi-directional event log (per-object, logged continuously)

5. Canonical Metric Dictionary

ChatGPT's metric set is adopted as the canonical baseline since it is the most operationally defined (each metric has an explicit unit of measurement). Overlapping terms from Grok's and the prior Claude draft are mapped as aliases and retired as independent metrics to prevent the duplication Perplexity flagged. Two Grok metrics are kept as genuinely distinct — not aliases — because they measure something the ChatGPT set does not.

Canonical metric	Definition	Retired aliases
TTE	Time to Exposure — elapsed authoritative simulation time from interface eligibility to first verified transfer.	—
TTD	Time to Detection — first verified transfer to defender recognition.	—
TTA	Time to Attribution — detection to correct causal source/path ID.	—
TTC	Time to Containment — detection to verified event closure.	—
FCR	False Closure Rate — share of trials where operators believed closure occurred while telemetry showed a remaining pathway.	Claude's “response-quality score” (Band-Aid vs. real fix) folds into FCR as its practical measurement.
CTC	Carrier Transition Count — number of medium/host changes before containment.	—
PD / PB	Propagation Depth / Breadth — layers crossed, nodes affected.	—
TER	Temporal Error Residual — difference between reconciled native-event ordering and authoritative ordering.	—
ECR	Event Closure Reconciliation — % of events with matching open/close records on both sides, no residual state.	Grok's “closure completeness score” and Claude's “containment success rate” both retired in favor of ECR.
RTR	Root Trace Reliability — confidence that independent observations converge on the same upstream dependency.	—
ARDR	Aperture Residual Detection Rate (kept from Grok, distinct) — specifically, the share of black-screen residual states caught vs. missed. Narrower than FCR; complements it.	Kept as its own metric — not a duplicate.
ACR	Awareness Conversion Rate (kept from Grok, distinct) — how many unaware agents shift to semi-/aware after repeated anomalies. No ChatGPT equivalent exists.	Kept as its own metric — not a duplicate.

6. Minimum Viable Protocol — 4 Locked Scenarios

Selected from ChatGPT's own 12-test matrix plus one addition for awareness stratification, per Perplexity's Phase 2 recommendation. Each scenario now carries a pre-registered null hypothesis and falsification criterion, closing the confirmation-bias risk both reviews flagged: without this, ambiguous results risk being read as partial confirmation of the doctrine by default.

Scenario 1 — Baseline & Black-Screen Closure (merges ChatGPT's T1 + T2)

- Manipulation: cycle screen/app/power/network/aperture states with no anomaly present; separately, hold presentation black while varying underlying aperture state.
- Primary metric: FCR, ARDR
- Null hypothesis (H0): presentation state and connection state are not independent — a black screen reliably predicts a closed aperture.
- Falsification: H0 is rejected if any trial shows a verified cross-aperture event while presentation state reads black/off. A single clean counterexample is sufficient to reject H0; this is deliberately a low bar to clear, since the entire doctrine depends on it clearing.

Scenario 2 — Content-State Bleed (ChatGPT's T3)

- Manipulation: subject consumes media with a mapped benign environmental variable (e.g., a harmless simulated temperature marker tied to on-screen content); neutral-content control run in parallel.
- Primary metric: TTE, TTD
- Null hypothesis (H0): rendered content properties show no measurable correlation with local environmental variables beyond the control condition's baseline noise.
- Falsification: H0 is rejected only if correlation exceeds the control baseline at a pre-agreed threshold, timestamped against content beats — not merely “something happened during viewing.”

Gate: this scenario cannot run until the open decision in Section 7 is resolved, since what counts as a valid “bleed” marker depends on that decision.

Scenario 3 — Aware vs. Unaware Comparison (new addition)

- Manipulation: matched occupants, one briefed on aperture theory, one unbriefed, identical device-use conditions.
- Primary metric: ACR, behavioral divergence
- Null hypothesis (H0): awareness of aperture theory produces no measurable difference in behavior or exposure outcome versus the unaware control.
- Falsification: H0 is rejected only if the aware group shows a statistically distinguishable behavior pattern (not just self-report) — tests ChatGPT's H5 directly.

Scenario 4 — Total Event Closure (ChatGPT's T7)

- Manipulation: close each side of an interaction independently; reconcile acknowledgments.
- Primary metric: ECR, TER
- Null hypothesis (H0): a closure request acknowledged by both sides always corresponds to a fully terminated event with no residual state.
- Falsification: H0 is rejected if any “closed” event later shows residual state activity — this is the direct test of the “Band-Aid” failure mode.

7. Open Decision Required Before Execution

Carried forward from the prior v2.0 draft (§8.6) and independently flagged by Perplexity as a blocker with no owner or deadline. It has an owner and a deadline now: it must be resolved before Scenario 2 runs, since it changes what counts as a valid measurement.

The question: When a benign marker successfully crosses an aperture in a test, does it register as a physical/environmental effect, a data/informational effect, or something else — and does that choice apply uniformly across all scenarios or vary by scenario type?

Owner: Nathan. No default is assumed. This is a canon decision, not a lab-design decision — Claude will not select on Nathan's behalf.

8. Deferred to Iteration 2

Everything below depends on Section 6 producing clean baseline data first. None of it is cancelled — all of it is preserved from the source documents and sequenced for after the lean prototype completes.

- Carrier Transition (ChatGPT's T5) — tests whether an anomaly can survive changing medium/host.
- Time Divergence (ChatGPT's T6) and the full time-discrepancy analysis from v2.0 §4.1 — cross-referencing native clocks across windows and layers.
- Root Trace (ChatGPT's T11) and the Root-Level Failure Tie-In from v2.0 §5 — the reconciliation-process-as-high-value-target argument, still flagged REPORTED/UNCONFIRMED and still resting on the unstated assumption (does the reconciliation process trust window-reported time?) that no draft has yet proposed a direct test for.
- Multi-Aperture scaling (ChatGPT's T4) — deferred since it compounds whatever the Scenario 1–4 baselines establish.
- Unknown Vector / Interface Airlock formal test (T8, T9) — deferred until L4 airlock behavior is characterized under known conditions first.
- Community Escape (T10) — explicitly out of scope until local (L1) containment is validated.
- Response Model Audit (T12) — replay of prior incident patterns; needs at least one completed run to have anything to replay.

9. Unresolved Governance Item Carried Forward

Grok's proposal to keep Layer 0 operators “partially unaware” of certain residual classes, in order to test monitoring blindness, is preserved here rather than quietly dropped — but it is explicitly not scheduled into the locked baseline or Iteration 2. It raises a real design question (can the lab's own oversight layer be deceived by its own protocol, and what's the recovery path if that deception doesn't resolve on schedule) that has no oversight framework yet. This needs its own governance pass before it is ever scheduled, separate from the technical decision gates above.

10. Next Steps

- Resolve the Section 7 decision (manifestation type) before Scenario 2 is executed.
- Circulate this document to ChatGPT, Gemini, and Grok for the disagreement-flagging step required by Section 1 before Nathan applies ONESTAMP.
- Once ONESTAMPED, this document becomes the single active working draft; v1.0 remains as historical origin record and all other branches move to fully retired status per the Section 2 changelog.
- After the lean prototype (Section 3) completes and produces clean Scenario 1–4 data, begin Iteration 2 planning against the Section 8 deferred list.