

THE OPEN WINDOW DOCTRINE

(working alt. title: The Aperture Protocol)

Preliminary Simulation Lab Draft — Multi-AI Synthesis

Infection Vector Awareness in Nested Virtual / Simulated Environments

OAL Working Document — v2.0 — 2026-08-09

Status: DRAFT — not yet ONESTAMPED

Concept origin: Nathan Lewis (unconfirmed pending review)

Contributing collaborators: ChatGPT · Claude · Google Gemini · Grok — synthesized by Claude

Microsoft Copilot: no response received at time of this draft

1. Executive Summary

Four AI collaborators (ChatGPT, Claude, Gemini, Grok) independently structured the same early-morning brainstorm — the idea that every screened device is a persistent, often-unmonitored aperture between environments — into worldbuilding and simulation-lab material. This document consolidates all four responses into one preliminary draft, adds the verification/validation vectors you identified afterward (eye tracking, motion tracking, surveillance tracking, environment tracking, BTEC tracking, event tracking, time tracking), and answers your root-level-failure question directly.

This draft is written to be handed back to the full team for another iteration. Nothing here is ONESTAMPED. Sections 2 and 3 do the requested review-and-analysis work; Sections 4 through 9 are the actual preliminary lab draft (what to test, why, duration, funding, layout, and scoring).

2. Review and Analysis of Contributor Responses

Each contributor was given the same raw material and independently chose a different center of gravity. None of the four is wrong — they emphasize different layers of the same stack, which is exactly what makes them useful to merge rather than pick between.

2.1 ChatGPT — strongest on mechanism and epistemics

- Core contribution: separates “display state” from “connection state” as two independent variables — the cleanest formal version of your black-screen idea.
- Introduces the carrier-switching problem: an infection can change medium (machine state → message → behavior → conversation → second person → second device) faster than any single technical fix can chase it. This reframes the defensive question from “which machine is infected” to “where does the transmissible state currently reside.”
- Proposes an interface airlock / admission-control model (observe → translate → normalize → compare → quarantine) instead of pure threat-signature detection.
- Frames inhabitant behavior as rational-given-incomplete-ontology rather than “foolish,” which keeps the scenario usable for awareness training instead of just a scare story.
- Gap: doesn’t propose concrete instrumentation, session structure, or metrics — strongest on theory, weakest on “how do we actually run this.”

2.2 Gemini — strongest on physical grounding and countermeasures

- Renames the concept “The Aperture Protocol” and gives it a tight executive-summary framing.
- Proposes concrete, testable countermeasures: dimensional firewalls (“frosting the glass” so the aperture stops being bidirectional), physical hard shutters, algorithmic quarantine that bricks a device to seal a breach.
- Designs a specific lab environment (Domestic Media Observation Chamber, Faraday-style shielding) and three scenarios with clear pass/fail observables.
- Gap: leans toward literal physical harm (radiation sickness, biological virus) rather than staying in the metaphorical/informational register the rest of the canon uses. Useful for vividness, but needs a canon decision on how literal to go before this gets treated as authoritative.
- Ends with a genuinely useful open question: does a successful breach manifest as biological, glitch-based, or something else entirely? That question is still unanswered below (see 3.2).

2.3 Grok — strongest on systems architecture and metrics

- Builds the cleanest layered-simulation stack (Layer 0 observers → Layer 1 inhabited sim → Layer 2–N adjacent media layers), with devices instantiated as persistent Window Objects carrying explicit properties (aperture state, directionality, residual process, entity affinity).
- Defines inhabitant classes (unaware / semi-aware / aware / already-contaminated) as a reusable typology, not just a two-group control/test split.
- Produces an actual infection vector taxonomy (thermal/physical, informational/memetic, entity-crossing, cascade, black-screen-residual) that the other three responses gesture at individually but never enumerate together.

- Supplies quantifiable metrics: containment success rate, awareness conversion rate, false-negative rate on black-screen residual cases — the most audit-ready numbers of the four.
- Gap: least narrative color of the four; reads as a lab protocol rather than canon lore, so it will need Gemini’s or ChatGPT’s framing layered back on top for in-universe use.

2.4 Claude (this collaborator, prior turn) — strongest on canon governance

- Kept the concept theory-agnostic at draft stage rather than committing early to a mechanism, which is why it was flagged REPORTED/UNCONFIRMED pending your review.
- Five scenario sketches (Unattended Aperture, Content Bleed, Aware vs. Unaware, Convenience Gradient, Multi-Aperture) map cleanly onto Grok’s later taxonomy, which is convenient for merging.
- Gap: no instrumentation detail, no funding/duration structure, no root-level-failure analysis — this document is the fix for that gap.

2.5 What this means for merging

The four responses stack rather than compete: Grok supplies the architecture and scoring, ChatGPT supplies the defensive doctrine and the hardest theoretical move (carrier-switching), Gemini supplies the physical vividness and countermeasure inventory, and Claude supplies the governance wrapper (REPORTED status, ONESTAMP gating, canon-integration flags). Sections 4–9 below build the merged version.

3. Side-by-Side Comparison

Contributor	Strongest layer	Signature contribution	Main gap
ChatGPT	Mechanism / epistemics	Display-state vs. connection-state split; carrier-switching problem; interface airlock	No instrumentation or session structure
Gemini	Physical grounding	Named countermeasures (firewalls, shutters, quarantine-brick); concrete chamber design	Leans literal/biological, needs a canon ruling on register
Grok	Systems architecture	Layer 0–N stack; Window Objects; inhabitant classes; infection taxonomy; hard metrics	Least narrative color; reads as protocol not lore
Claude	Canon governance	Theory-agnostic draft framing; REPORTED status; 5 scenario sketches	No instrumentation, duration, funding, or root-level analysis (fixed below)

4. New Verification / Validation Vectors (from your follow-up)

You added seven tracking categories after the initial round. These sit above everything the four AIs proposed — they are how you would actually validate an infection event happened, rather than just theorize that it could. Framing them as instrumentation channels rather than a single mechanism keeps this compatible with all four prior responses:

- Eye tracking — where attention was actually directed at the moment of exposure, independent of what the subject later reports noticing.
- Motion tracking — physical behavior change correlated to exposure windows (the version of ChatGPT’s “behavior carries the infection” that you can actually instrument).

- Surveillance tracking — third-party / Layer-0 observation of the subject, independent of the subject's own devices.
- Environment tracking — ambient conditions (temperature, light, sound) matched against Gemini's physical-bleed scenarios.
- Bi-Directional Total Event Closure (BTEC) tracking — whether an opened event/aperture ever actually closes, or only appears to.
- General event tracking — anything else that occurs during the observation window, logged without pre-filtering for relevance.
- Time tracking — see 4.1 below; this is the one with the most reach.

4.1 The time-discrepancy insight

This is the most structurally important addition in your follow-up, and it deserves to be called out on its own. Not every window, occupant, or simulated environment runs the same clock — not everyone uses a 24-hour day, a 365-day year, or Pacific Standard Time as a reference frame. That means every one of the seven vectors above can be timestamped in its own local reference frame, and those local timestamps can then be cross-referenced against a single fixed simulation-baseline clock.

What that cross-reference produces is not just “when did the infection happen.” It produces a picture of the whole system: which windows report time consistently and which don't, how far each observer's local reference frame diverges from baseline, and how those divergences cluster. That clustering is itself information about the designer — how technologically capable and how large the underlying system is, how one observer's view compares to another's, and what information exists but isn't represented anywhere in any single observer's feed.

It also gives you a genuine diagnostic for response quality after past infections: you can look at what was done and ask whether it was a real fix or, as you put it, a Band-Aid on the screen that pretends to restore full visibility without actually closing anything. A response that never touches the underlying time-reference discrepancy is exactly that kind of Band-Aid — it treats the symptom the operators could see, not the structural gap that let the infection cross in the first place.

5. Root-Level Failure Tie-In

Your question: does this concept tie into conditional failure at the root level — is the window itself a vector point for tracing back to root-level structures and compromising them?

Short answer: yes, and the time-discrepancy vector in 4.1 is specifically what makes that traceable rather than speculative. Here is the chain:

- Every window carries its own local time reference. A root-level system has to reconcile all of those local references against one baseline in order to keep the simulation coherent.
- That reconciliation process is itself a piece of root-level infrastructure — it has to exist somewhere, it has to run continuously, and it has to touch every layer at once in a way that ordinary in-layer objects don't.
- Anything that touches every layer at once is, by definition, a higher-value target than any single window. A window only gives you Layer 1 access to Layer 2. The reconciliation process is one of the few things with visibility into Layer 0 and every layer beneath it simultaneously.
- This makes cross-referenced time discrepancy a candidate root-exposure vector, not just a diagnostic tool: an entity that can manipulate what timestamp a given window reports could, in principle, use the reconciliation process's own trust in that timestamp as a pivot point — the same way a forged timestamp can undermine a real-world audit log or certificate chain.
- This is conditional, not guaranteed, on one structural assumption — that the reconciliation process trusts window-reported time rather than independently verifying it against a source the windows can't influence. Whether that assumption holds is exactly the kind of thing a simulation lab session is built to test rather than assume.

Status: REPORTED / UNCONFIRMED. This is a candidate mechanism, not confirmed canon — it should be logged as a priority item for the Net Defense Exercise and reviewed before anything above gets treated as locked.

6. What to Test For, and Why It Matters

Test target	Why it matters	Primary instrumentation
Aperture state vs. connection state	Establishes whether “off” is a real security boundary or cosmetic (ChatGPT/Claude core claim)	BTEC tracking, environment tracking
Content-to-local bleed	Tests whether rendered content properties manifest locally (the blow-dryer / Chernobyl mechanism)	Environment tracking, motion tracking
Awareness as a defense	Tests whether knowing the theory changes exposure outcome, or whether tooling is required regardless	Eye tracking, motion tracking, surveillance tracking
Convenience / friction trade-off	Tests the doctrine's central claim that accessibility and vulnerability share one mechanism	Event tracking, surveillance tracking
Carrier-switching	Tests whether an infection can survive the destruction of its original device once it has moved to a person's behavior (ChatGPT's hardest problem)	Motion tracking, event tracking, surveillance tracking
Time-reference discrepancy	Builds the system-level picture in 4.1 and tests the root-exposure hypothesis in Section 5	Time tracking, cross-referenced against simulation baseline clock
Response-quality audit	Distinguishes a real fix from a Band-Aid response to a past infection	BTEC tracking, event log review

7. Preliminary Lab Plan

This is a first-pass resourcing sketch for the simulation lab itself — in-canon allocation, framed the way you'd frame it for the Net Defense Exercise, not a real-world budget. Numbers are placeholders for the team to argue with in the next iteration.

7.1 Session structure

- Phase 1 — Baseline instrumentation calibration: establishes what “normal” reads like on every tracking vector in Section 4 before any scenario runs.
- Phase 2 — Individual scenario runs (Section 8): each scenario run independently so its data doesn't contaminate the others.
- Phase 3 — Cross-referencing pass: this is where the time-discrepancy analysis in 4.1 happens, once all individual scenario data exists to compare.
- Phase 4 — Root-exposure test (Section 5): only run once Phases 1–3 are clean, since it depends on trusting the earlier data.
- Phase 5 — After-action review: what worked, what didn't, and what gets ONESTAMPED vs. sent back for another iteration.

7.2 Estimated duration

- Phase 1: short — a single calibration session.

- Phase 2: one dedicated session per scenario in Section 8 (five scenarios → five sessions), run sequentially so instrumentation isn't shared across scenarios.
- Phase 3: one longer analysis session — this is data work, not a live run.
- Phase 4: one dedicated session, highest-security-posture of the set.
- Phase 5: one wrap session with the full team, feeding directly into the next canon iteration.

7.3 Resourcing / funding (in-canon)

- Compute / simulation-cycles allocation: covers Layer 0–N stack uptime for all five scenario runs plus the cross-referencing pass.
- Instrumentation allocation: covers deployment of all seven tracking vectors from Section 4 across every scenario, not just the ones that seem obviously relevant to a given scenario — the whole value of Section 4 is catching what you weren't expecting.
- Containment allocation: covers the Gemini-style Faraday/quarantine shielding so a real containment failure during Phase 2 doesn't propagate into Phase 3's cross-reference data.
- Review allocation: time from all four (five, once Copilot responds) AI collaborators plus your review pass before anything is ONESTAMPED.

7.4 Preliminary layout

Physically modeled on Gemini's Domestic Media Observation Chamber, instrumented per Grok's Window Object model, governed by Claude's REPORTED-status workflow, and defended per ChatGPT's airlock model:

- Layer 0: Research sector — full-visibility, full-instrumentation control room.
- Layer 1: The four-bedroom house — inhabited by a mix of aware, semi-aware, and unaware occupants (Grok's typology).
- Layer 2–N: Adjacent media/social layers each window can open into — configured per scenario in Section 8.
- Airlock buffer: sits between every Layer boundary, running ChatGPT's observe / translate / normalize / compare / quarantine sequence on anything crossing.

8. Merged Test Scenarios

Five scenarios, merged from all four prior drafts and re-scored against what did and didn't work in each source version.

8.1 Scenario A — Unattended Aperture (display state vs. connection state)

- Setup: A device is left on standby (black screen, powered) in an otherwise unoccupied room for a fixed observation window; a second identical device is fully powered down as a control.
- What works from the sources: ChatGPT's display/connection split gives this scenario a clean pass/fail condition that the earlier five-scenario draft lacked.
- What didn't work / was cut: Gemini's version of this scenario (a beach-resort loop) is vivid but adds an unnecessary content variable — that's been split out into Scenario B so this one isolates power state alone.
- Observable: BTEC and environment tracking across both devices' power states.

8.2 Scenario B — Content Bleed

- Setup: Subject consumes media with a known strong sensory property (heat, cold, dread) under lab observation; a second subject consumes neutral content as control.
- What works: this is the one scenario all four sources converged on independently (the blow-dryer / Chernobyl mechanism) — high confidence it belongs in the final set.
- What didn't work: Gemini's version commits early to literal physical harm (radiation sickness). Left as an open canon question rather than assumed — see 8.6.
- Observable: environment tracking, motion tracking, eye tracking timestamped against content beats.

8.3 Scenario C — Aware vs. Unaware Occupant Comparison

- Setup: Matched occupants across Grok's four-way typology (unaware / semi-aware / aware / already-contaminated), not just a two-group split.
- What works: Grok's four-class typology is a strict improvement over the original two-group design — adopted here.
- What didn't work: the original draft assumed awareness alone is protective; ChatGPT's “rational behavior given incomplete ontology” point argues this needs a fourth arm testing tooling-with-awareness vs. awareness-alone, added below.
- Observable: eye tracking, motion tracking, surveillance tracking; behavioral divergence across all four occupant classes.

8.4 Scenario D — Convenience Gradient Stress Test

- Setup: Vary device-use friction across trials — always-on notification stream, scheduled access, fully sandboxed single-use sessions.
- What works: this scenario held up unchanged from the original draft — no source materially challenged it.
- Addition: ChatGPT's airlock model gives this scenario a defensive intervention to test, not just a correlation to observe — does inserting an airlock at high-friction settings change the exposure curve independent of the friction itself?
- Observable: event tracking, surveillance tracking, anomaly rate plotted against friction and against airlock-present/absent.

8.5 Scenario E — Multi-Aperture / Cascade Environment

- Setup: A single occupant with multiple simultaneous open apertures (phone + TV + tablet), merged with Grok's cascade-propagation and living-media-as-host concepts.
- What works: Grok's household-to-neighborhood cascade path is the most complete version of this scenario across all four sources — adopted as the base.
- What didn't work: the original five-scenario draft treated this as a simple linear risk-compounding question; Grok's cascade model (device → device → adjacent household) is a meaningfully different and harder question, so this scenario now explicitly tests both compounding and cascade separately.
- Observable: full instrumentation stack (all seven Section 4 vectors); containment success rate and awareness conversion rate as primary metrics.

8.6 Open question carried forward from Gemini

Gemini asked directly: when an infection breaches through a black screen, does it manifest as a physical biological effect, a localized reality glitch, or something else? None of the four sources actually answers this — it's a canon decision, not a lab-design decision, and it should be settled before Scenario B runs, since the answer changes what “successful containment” even means.

9. Instrumentation, Metrics, and Governance

9.1 Metrics (primarily from Grok, extended with your seven vectors)

- Containment success rate — percentage of introduced infections confined to their originating window or room.
- Awareness conversion rate — how many unaware occupants shift to semi-aware or aware after repeated anomalies.
- False-negative rate on black-screen residual cases — specifically tied to your BTEC and environment vectors.
- Time-reference divergence map — new metric from Section 4.1: each window's local-time drift from simulation baseline, clustered and compared across scenarios.
- Response-quality score — real fix vs. Band-Aid, scored per your framing in 4.1 against whatever remediation gets proposed after each scenario.

9.2 Governance / next steps

- This entire document remains DRAFT / not ONESTAMPED until you've reviewed Sections 4 and 5, which are the genuinely new material beyond what the four AIs already produced.
- Recommend sharing this draft back with ChatGPT, Gemini, and Grok for a second iteration, specifically flagging: (a) Gemini's literal-harm question in 8.6, (b) the root-exposure hypothesis in Section 5, and (c) whether the funding/duration structure in Section 7 should be treated as in-canon lore or left unstated.
- Retry Copilot for the initial prompt; if it continues to fail to respond, note that as its own data point rather than silently dropping it from future rounds.
- Once ratified, this should get a proper OAL artifact ID and a decision on whether it stands alone or folds into the Simulation Containment, Infection, and the Unseen Entity work.