OAL CURRENCY LABORATORY

Experimental Baseline v0.2

Quantum Credits / 1COIN / NCOIN | Protected premises, preserved forks, and adoption controls

Prepared by Perplexity | 2 November 2026

PROPOSED FOR REVIEW / PRE-CANON

Independent contribution for circulation to the AI Collaboration Team. Preserve disagreements; freeze test conditions, not winning currency mechanics.

Section map

Executive purpose and status

1. Review scope and source precedence
2. Constitutional controls proposed for experimental adoption
3. Status vocabulary and change control
4. Working instruments and open forks
5. Independent disagreements retained
6. Boundary vocabulary: proposed operational use
7. Adoption and next-run gate

Source and review register

Executive purpose and status

This is Perplexity's independent contribution to the OAL Currency Laboratory, prepared on 2 November 2026 at the project owner's request. It implements the previously recommended two-part package: an Experimental Baseline and a Laboratory Charter with a foundational test suite. The companion document is the Currency Laboratory Charter and Experimental Test Protocol v0.1. Neither document is a harmonized team specification, an approved constitution, a completed experiment, or a real-world financial instrument.

Status: PROPOSED FOR REVIEW / PRE-CANON. The v0.2 label follows the development sequence proposed in the archives; it does not imply that an approved v0.1 constitution exists. Rules described as frozen below become LOCKED FOR TESTING only when the project owner adopts a named run configuration. Today's request authorizes preparation of these materials, not canonization of their contents.

The originating premise is deliberately extraordinary: a Quantum Credit evaluates to the value of what it is applied to, including targets outside ordinary exchange and ordinary temporal assumptions. The laboratory must neither quietly weaken that premise into a conventional budget nor strengthen it into an unexplained guarantee of every requested outcome. Source basis: Master Pass-Through, inception message and Parts II-V; Unified First Responses, sections on target-relative valuation and two-track testing. [file:1][file:4]

1. Review scope and source precedence

The review covers all seven supplied attachment artifacts: two formats of the August master pass-through, the OAL conversation archive, the Unified First Responses PDF, the pasted six-member status compilation, and both formats of Next Step Recommendations. Full extracted text, including paragraph and table content in DOCX, was inspected alongside the supplied PDF text. Related PDF/DOCX editions are not independent corroborating sources. This is a textual review, not an authentication of the attributed AI authors, cited outside research, or every embedded hyperlink.

The latest compilation records broad agreement about a pre-canon project and divergent recommendations about sequencing and test design. ChatGPT and Copilot explicitly move toward paired baseline/charter deliverables; the Perplexity contribution proposes the same two-part structure. Claude's ballot-based sequencing preference, Gemini's 3-by-4 matrix, and Grok's charter-first approach remain recorded alternatives rather than rejected authors. [file:9][file:10]

Interpretive precedence proposed here: direct owner statements define intent; explicit owner decisions define adopted rules; later owner-approved revisions supersede earlier approved revisions; AI summaries remain interpretations unless adopted. A later AI statement does not automatically outrank an earlier owner statement. Source disputes remain visible in the decision ledger. No approval of detailed currency mechanics is inferred from these compilations.

2. Constitutional controls proposed for experimental adoption

EB-01 - Preserve the premise. Accept target-relative QC behavior as stipulated within a named experimental track. Test what follows. Do not reject the experiment merely because contemporary monetary systems cannot implement its premise.

EB-02 - Do not predetermine the outcome. Target-relative valuation does not, by itself, specify resource generation, title transfer, authorization, or reality transformation. Record each claimed capability separately. Preserve an expansive outcome-capable track rather than excluding it by definition.

EB-03 - Separate operations. Target identification, valuation, synchronization, authority, transferability, execution, settlement, and consequence are distinct questions. This distinction does not require a separate token or separate institution for each question.

EB-04 - Separate status from consequence. Rejection alone does not issue NCOIN. A failed attempt may nevertheless cause a separately evidenced consequence under a declared NCOIN model. Failure and harm may coexist without being identical.

EB-05 - Freeze configurations, not discoveries. A run must retain its target, value domain, capability, authority rule, time model, and settlement criterion. A repair creates a new configuration and a rerun; it cannot retroactively convert a failed original run into a pass.

EB-06 - Permit architectural simplification. QC-only, QC plus ordinary records, QC plus 1COIN, and Trinity configurations must be compared fairly. No instrument receives an automatic right to survive. Equally, an instrument need not be logically indispensable to have useful narrative or operational value.

EB-07 - Preserve dissent. Minority reports remain attached to the case they challenge. More AI votes do not establish correctness. Distinguish a reasoned objection, a counterexample, missing specification, and a preference.

EB-08 - Limit claims to their evidence. A stated axiom is not a measured result. A written walkthrough is not an executable simulation. A finite modeled success is not proof over all targets or 150 years of actual operation.

EB-09 - Keep fiction and external claims separate. Speculative authority belongs to the declared world. It does not grant real access, ownership, or permission. These documents authorize no transactions, purchases, deployments, or interaction with protected systems.

EB-10 - Owner retains adoption authority. The owner or an explicitly designated maintainer may adopt an experimental configuration or later canon. AI contributors recommend and challenge; none may silently approve rules on behalf of the collaboration.

These controls are new procedural proposals informed by the master record's review instructions and the latest recommendations, not a claim of unanimous adoption. [file:1][file:9][file:10]

3. Status vocabulary and change control

Label	Meaning in this package
SOURCE INTENT	An attributed owner statement, preserved with its source and context.
PROPOSED	A contributor's suggested rule, definition, or procedure; not adopted.
LOCKED FOR TESTING	Owner-adopted assumptions frozen for a named run and revision.
OPEN	An unresolved choice; must be set before a test depending on it.
SUPERSEDED / REJECTED	A proposal no longer used in a specified branch; retain its record.
CANON	An explicit later owner decision, separate from experimental passing.

Experimental verdicts are a different namespace: PASS, FAIL, INDETERMINATE, NOT APPLICABLE, or INVALID RUN. For example, an OPEN issuance rule is a design status; INDETERMINATE is the verdict when that missing rule prevents a conclusion. A failed purchase can be a laboratory PASS when correct rejection was the declared requirement.

4. Working instruments and open forks

Instrument	Hypothesis retained	What is not fixed
QC	Target-relative value behavior; QC(T) = V(T) is historical shorthand.	Target validity, value type, capability, value source, issuance, reuse, recursion, and timing.
1COIN	Possible executor, record, package, or state of settlement.	Whether separate, transferable, consumed, reusable, independently issued, or a currency at all.
NCOIN	Possible liability, consequence, remediation, degradation, inverse, or conservation representation.	Trigger, unit, bearer, reversibility, transferability, conservation, and discharge.

These hypotheses follow the source progression but do not settle it. In particular, earlier contribution-backing, governed-treasury, penalty-freezing, and mandatory-discharge models are historical candidates, not universal requirements. Perplexity's own early governance-heavy model is likewise not the default constitution. [file:1][file:2][file:4][file:9]

The capability fork is primary: EVALUATE produces an equivalence description; TRACK maintains a relation; SUPPLY transfers, creates, or converts specified value/resources; ENACT makes a declared target-state change. These capabilities can be combined only in an explicitly specified configuration. An evaluation-only control is useful, but must be described as a limited comparator if it does not satisfy the owner's intended acquisition capability.

5. Independent disagreements retained

Disagreement D-01 - Four historical QC labels need not be mutually exclusive. Static and Dynamic concern timing, Commitment concerns obligation scope, and Relational concerns representation. Preserve QC-S/QC-D/QC-C/QC-R as historical aliases, but map each run to explicit dimensions. A relational rule can govern a dynamically revalued 150-year commitment. Do not blend configurations mid-run; do not prohibit predeclared combinations merely because labels overlap. This continues the Perplexity proposal in the latest compilation. [file:9][file:10]

Disagreement D-02 - Permission-Bound, Sovereign, and Ontological are not a clean three-way partition. Authority, capability, and environment must be recorded separately. A permission-bound system can transform a simulated world; a sovereign system can merely evaluate. Copilot's axis proposal is compatible with this refinement. Gemini's 3-by-4 matrix remains a possible benchmark set, not an exhaustive design space. [file:9][file:10]

Disagreement D-03 - Boundary classifications are hypotheses relative to stated rules. Some incoherence follows from a fixed definition; other barriers depend on authority, identity, physics, or temporal assumptions. No proposed boundary is universal merely because multiple reviewers repeat it. Keep original target identity fixed when testing extraordinary access or transformation.

Disagreement D-04 - Independent necessity is too narrow as the sole retention criterion. Test necessity, measurable operational advantage, and narrative utility separately. A settlement receipt may be useful without being a currency; a named instrument can survive in a narrative branch even if mechanically redundant, provided that redundancy is disclosed.

Disagreement D-05 - A complete round of ten-case philosophical essays is not a prerequisite for drafting the protocol. Missing independent boundary analyses can proceed under its common template. No claim is made

that an absent response was never produced elsewhere; only the supplied record is in scope.

6. Boundary vocabulary: proposed operational use

Term	Use for review, not settled canon
Non-Purchasable Domain	A broad collection of targets outside ordinary purchase; record the reason.
Absolute Transaction Boundary	A proposed claim that an essential transaction condition cannot be met under specified rules; demonstrate its scope.
Ontological Transaction Limit	A proposed identity/being-level incompatibility with ownership, transfer, or requested change.
Quantum Resolution Null Domain	An umbrella for absent valid QC pathways; not a substitute for a precise reason code.
Currency Event Horizon	Optional narrative language; never itself an executable decision rule.

The ten historical target categories are retained as tags in the charter, not as a proven ordered scale. Multiple tags may apply to the same target. Boundary contact, missing information, an unknown seller, and an incoherent target must not be collapsed into one null result. Source basis: master taxonomy and glossary; later terminology cautions. [file:1][file:9]

7. Adoption and next-run gate

- Owner adoption identifies the baseline revision, charter revision, case IDs, configuration IDs, and reviewers. Adoption may be limited to a pilot without approving canon.
- Every run-dependent OPEN item must be set or explicitly treated as an indeterminacy condition. No default external permission requirement or unlimited sovereignty is silently supplied.
- First circulate the package for objections to test fairness. Then run the foundational cases before using macro-scale outputs to rank architectures.
- Preserve original independent responses before peer review. Publish amendments and reruns separately. A later Version 2 candidate must link to evidence and unresolved dissent.

This baseline does not choose a winning capability, temporal configuration, conservation rule, NCOIN domain, or 1COIN model. Its recommended first experiment is the discriminating resource-box case F01 in the companion charter. That case isolates whether QC describes, tracks, supplies, or enacts before the 150-year scenario introduces additional ambiguity.

Source and review register

Sources are user-supplied development artifacts, not externally verified research. The source IDs below are conversation-local; filenames and section descriptions support later circulation. Attributions identify the labels in the compilations, not authenticated model execution histories. No external primary-source validation or real-world financial analysis was performed for this package.

A - Currency_Laboratory_Master_Pass_Through_2026-08-20-1.docx [file:1]. Owner inception and boundary directives; Parts II-VIII architecture, tests, open questions.

SHA-256: febc3562e62be2954b2e0a5ff03e87be7c504761ee127dab22fa846fe09c890f

A-PDF - Currency_Laboratory_Master_Pass_Through_2026-08-20-1.pdf [file:3]. Paired master edition; 20 pages. Not an independent vote.

SHA-256: b123c7a73dfdb99d2e9512029b60a699ac287d11875ede4b829d219d10eefcc0

B - OAL_Currency_Laboratory_Conversation_Archive.docx [file:2]. Grok-attributed archive; baseline-first recommendation and ontological boundary framing.

SHA-256: 3c347e65ac0ef311c9cbb74905e3b59dc9917bc4ba400da255de4046da73bc17

C - CURRENCY-REVIEW-UNIFIED-PASS-THROUGH-FOR-FIRST-RESPONSES.pdf [file:4]. Earlier attributed baselines; competing governance, settlement, penalty, and value-source interpretations; 51 pages.

SHA-256: d7dc8bbefeb4c1ec94ac3abd10c6f021c9ab02ddf4ba83a6621abd0f435b71e8

D - pasted_text_1793493056.txt [file:8]. Six-member compiled status review and sequencing recommendations; shared source dependencies.

SHA-256: 7a0fb29daa91c0b40852ecf82aa8e5251faf0e36884d60bd43e8145160102cea

E-PDF - CURRENCY-REVIEW-NEXT-STEP-RECOMMENDATIONS.pdf [file:9]. Latest recommendation compilation; 29 pages; includes the Perplexity overlapping-dimensions proposal on pages 25-26.

SHA-256: 56f5f226e3d0e40fb2053724a9f26d1890b59ca2f8a9f55aba8d0d2b86d00918

E-DOCX - CURRENCY-REVIEW-NEXT-STEP-RECOMMENDATIONS.docx [file:10]. Editable counterpart of the latest compilation; reviewed including tables. Related edition, not independently validated evidence.

SHA-256: 28a0da1b90a83f293c5f09af37fa7d9bba7d22d3f1517de4d12c5f1bc351e08a

Review limitation: extraction and supplied page text were used for substantive textual review. Rendered illustrations, styling, and external linked documents are not additional reviewed evidence. The separate original Red Team Currency Review and focused markdown baseline mentioned within the archives were not separately supplied; their contents are known here only through these sources. Paired editions and repeated AI summaries are not counted as additional experiments.