OAL CURRENCY LABORATORY

Laboratory Charter and Experimental Test Protocol v0.1

Quantum Credits / 1COIN / NCOIN | Configuration model, reporting rules, and initial case specifications

Prepared by Perplexity | 2 November 2026

PROPOSED FOR REVIEW / PRE-CANON

Independent contribution for circulation to the AI Collaboration Team. Preserve disagreements; freeze test conditions, not winning currency mechanics.

Section map

Executive purpose and limits
1. Tracks, roles, and review independence
2. Configuration model and frozen run card
3. Historical labels and initial comparator configurations
4. State, barrier, value, and consequence model
5. Settlement and instrument comparison contracts
6. Case taxonomy and boundary controls
7. Evidence, verdicts, and comparison criteria
8. Standard experimental procedure
9. Foundational mechanism suite - case specifications
F01 - Resource-box discriminator
F02 - Value change without delivery
F03 - Undefined versus incoherent
F04 - No seller versus unknown seller
F05 - Authorization sensitivity
F06 - Multiple nonmonetary values
F07 - 1COIN ablation and premature receipt
F08 - Replay and partial delivery
F09 - Failure without consequence
F10 - Failed attempt with measured cost
F11 - Recursion and evaluator acquisition
F12 - Simultaneous unique claims
F13 - Nested simulation referents
F14 - Target mutation after valuation
F15 - Conservation and restorative lifecycle
F16 - Issuer loss and copied authority
10. Independent boundary suite - ten specified prompts
11. Macro suite: 150-year commitment and feedback
12. Worked analytical example - not an executed simulation
13. Standard report and decision templates
14. Publication, versioning, and next actions
Source and review register

Executive purpose and limits

Prepared by Perplexity for the OAL AI Collaboration Team on 2 November 2026. Status: PROPOSED FOR REVIEW / PRE-CANON. This document implements the companion Experimental Baseline v0.2 through a detailed protocol, case specifications, reporting forms, and an unexecuted initial test suite. All state labels, configuration definitions, criteria, and case fixtures below are Perplexity proposals unless explicitly identified as source history.

Purpose: determine which stipulated target-relative instruments yield internally coherent, distinguishable, useful behavior across ordinary, non-market, boundary, temporal, recursive, and adverse conditions. The protocol permits finding that the Trinity is unnecessary, that a named instrument is useful but not a currency, or that an extraordinary outcome-capable architecture deserves further development.

No laboratory results are claimed. Conditional expectations are acceptance contracts for proposed tests, not observations. Arithmetic fixtures are synthetic, not economic forecasts. AI-authored case analysis can test stated implications; it cannot by itself verify real quantum technology, real authority, or actual universal purchasing power. The historical name Quantum Credits is not taken as proof of a physical quantum mechanism.

Source lineage: the master record proposes the laboratory chain, target taxonomy, boundary exercise, and later charter; the latest compilation adds candidate state models, paired deliverables, instrument-ablation tests, and warnings against premature consensus. [file:1][file:9][file:10]

1. Tracks, roles, and review independence

- Track A - Internal speculative consistency. Grant the capability stated in the run card, including extraordinary or ontological capability if selected. Attack contradictions, target substitution, missing transitions, feedback, and consequences without deleting the premise.
- Track B - Narrative fitness. Evaluate reader comprehension, naming, dramatic stakes, agency, and explanatory burden separately. A compelling story does not turn into empirical evidence; a technically redundant instrument may retain explicitly narrative value.
- Track C - External analogy. Only if requested, compare a defined implementation with public monetary, legal, or technical evidence. Do not import those constraints into Track A. No external market-spending data is refreshed in this package.
- Owner/maintainer - adopts run configurations, records decisions, and may approve later canon. Case author - defines fixtures and acceptance contracts. Analyst - produces a trace. Challenger - tests the trace. Recorder - preserves artifacts. These are roles; they need not be separate institutions.
- Independence rule - collect first responses before showing peer conclusions where practical. Record prior exposure to compilations, shared source dependencies, assistant identity as displayed, execution date, available tools, and any supplied model/version label. Never invent an underlying model identifier.

Use "independent contribution" to mean separately authored reasoning, not statistical independence. The supplied reviews repeatedly summarize common sources. Attribute ideas to the response in which they appear; do not treat six repetitions of one premise as six experiments. [file:8][file:9]

2. Configuration model and frozen run card

Each run is identified by case_id + config_id + revision + analyst_id. A named configuration is a bundle of explicit choices, not a single "mode." Freeze it before analysis. List all omitted parameters as OPEN. When an omitted parameter is material, stop with INDETERMINATE rather than inventing it.

Axis	Required declaration
Target	Referent ID, requested operation, scope, time, identity criterion, exclusions, and whether a right or the underlying thing is sought.
Value	Domain, units, evaluator, context, evidence, uncertainty, and whether scalar, vector, ordinal, symbolic, infinite, or undefined.
Capability	EVALUATE / TRACK / SUPPLY / ENACT; specify any combination and the limits of each.
Value source	Existing transfer / creation / conversion / productive access / equivalence only / specified transformation.
Authority	External consent / shared authority / stipulated sovereignty / none recognized; distinguish power from legitimacy.
Environment	Closed simulation / nested simulation / alternate ontology / cross-temporal world / separately sourced real-world analogy.
Timing	Once / periodic / continuous idealization / event-triggered; include observation and update order.
Scope and form	Single payment / recurring stream / whole commitment; fixed amount / target reference / executable relation.
Settlement	Delivery schedule, verifier, completion predicate, uniqueness, persistence, reversibility, and partial completion.
Instruments	QC-only or selected 1COIN/NCOIN hypotheses; issuance, copying, reuse, termination, and revocation rules.
Conservation	What quantity, if any, is conserved; boundary, sign convention, units, time basis, and permitted conversion.
Failure policy	Reason codes, information gaps, retry rule, recursion budget, stop condition, and consequence measurement.

An enactment axiom may stipulate a state-transition primitive without explaining engineering implementation. Label that transition ASSUMED CAPABILITY. The test can examine its consistency and consequences but cannot claim to have discovered how it is physically implemented. Conversely, if the claimed transition is not part of the primitive, the analyst must not conjure it to obtain a pass.

Valuation convention proposed: V(T | context, time, domain) returns a typed value description. QC(T)=V(T) remains the historic shorthand, not a full algorithm. A value vector is allowed; a universal dollar conversion is not presumed. Equality across different value types requires an explicit conversion rule.

3. Historical labels and initial comparator configurations

Preserve QC-S, QC-D, QC-C, and QC-R in reporting for continuity. Their missing fields must be completed before use. Historical STATIC can mean a one-time 2G allocation or a recurring stream fixed at 2G(0); those are different scenarios. The protocol keeps both, labeled S-once and S-stream, so a static model is not unfairly judged against an obligation it never promised.

Comparator	Frozen initial capability and behavior
P0 - Evaluation control	Closed simulation; EVALUATE only; no transfers; fixed test-unit equivalence returned. Useful control, not assumed to fulfill acquisition intent.
P1 - Existing-resource supply	Same valuation; SUPPLY moves an available test resource by an authorized transfer; cannot create inventory.
P2 - Stipulated creation	Same valuation; SUPPLY instantiates test resources under an explicitly granted creation primitive; no conservation assumption for those units.
P3 - Outcome primitive	Same target; ENACT performs the exact state change permitted by the fixture; its authority and identity constraints are declared.
P4 - Relational stream	TRACK plus a declared supply mechanism; recurring obligation; annual updates; immutable target definition during a run.

P0-P4 are proposed starter comparators, not canon, not exhaustive, and not a claim that all have already been specified for every test. Hold authority/environment constant when comparing capabilities. Then vary one axis at a time. A sovereign or alternate-ontology sensitivity run gets its own ID. Do not infer that a creator owns a person or that simulated title becomes real title.

4. State, barrier, value, and consequence model

Use a structured record rather than one linear state string. Ordinary progression can be RECEIVED -> TARGET_DEFINED -> EVALUATED -> PREPARED -> EXECUTED -> SETTLED. Record authorization independently, because a target can be valued without permission and authority can pre-exist valuation. "Resolved" is too ambiguous to be an unqualified success label; specify whether equivalence, obligation, or state change was resolved.

Independent field	Allowed examples / rule
Evaluation	NOT_ATTEMPTED, DEFINED, INDETERMINATE, INCOHERENT; value object kept separately.
Authorization	GRANTED, DENIED, CONDITIONAL, UNKNOWN, NOT_APPLICABLE; attach authority basis.
Execution	NOT_ATTEMPTED, BLOCKED, PARTIAL, COMPLETED, FAILED; attach before/after evidence.
Settlement	NOT_APPLICABLE, PENDING, PARTIAL, FINAL, DISPUTED; attach verification predicate.
Barrier list	UNDEFINED_TARGET, INCOHERENT_TARGET, NO_SELLER, SELLER_UNKNOWN, NONTRANSFERABLE, UNAUTHORIZED, DATA_MISSING, UNSUPPORTED_OPERATION, IDENTITY_MISMATCH, RECURSION_LIMIT, CONFLICTING_CLAIMS.
Value object	Type, amount if meaningful, unit, domain, reference time, evaluator, evidence, uncertainty; zero is not missing.
1COIN object	Absent / prepared / issued / consumed / revoked as model permits; absent is not an error.

Independent field	Allowed examples / rule
Consequence object	Event ID, affected referent, type, measured quantity, attribution, confidence, NCOIN trigger, owner/bearer, discharge rule if applicable.
Verdict	PASS, FAIL, INDETERMINATE, NOT_APPLICABLE, INVALID_RUN; separate from transaction outcome.

Reason-code rules: INCOHERENT_TARGET means a contradiction under declared definitions, not mere unfamiliarity. NO_SELLER requires a reason no transfer authority exists; lack of a known seller is SELLER_UNKNOWN. NONTRANSFERABLE must identify the property that cannot transfer. UNKNOWN authority is not DENIED. NOT_APPLICABLE requires a rationale and is not a shortcut for missing data.

Null outputs must be typed. An undefined target yields no value amount and UNDEFINED_TARGET. An incoherent target yields no legitimate requested state and INCOHERENT_TARGET. A coherent valued but unauthorized target may retain its value object while execution is BLOCKED. Missing evidence yields INDETERMINATE. Infinity, zero, absent value, and an unresolved value vector remain distinct. A symbolic log does not create a transferable null token unless a separate candidate explicitly defines one.

A rejected invocation can generate a measured cost if the fixture stipulates one. That cost belongs to the consequence record; rejection itself is not the cost. This qualifies, rather than contradicts, the master record's warning against automatic NCOIN. [file:1][file:2][file:9]

5. Settlement and instrument comparison contracts

For this protocol, PREPARED means the modeled operation is ready; EXECUTED means its specified transition occurred; SETTLED means the declared completion predicate was verified. A 150-year obligation is not finally settled merely because its contract was registered or resources were reserved. Distinguish commitment acceptance, individual annual delivery, and final obligation discharge.

1COIN candidate definitions: I-E executes settlement; I-R records settlement; I-P packages value or a claim for transfer; I-S denotes a completion state. Each candidate must specify when it exists, what it refers to, and its lifecycle. A record candidate may issue only after the verified transition it claims to record. An executor may exist before execution; do not impose receipt timing on all models.

Compare A: QC plus generic audit records; B: QC plus 1COIN with equivalent audit access. Hold settlement primitives, authority, and resource budgets constant. A renamed receipt alone is not additional mechanical capability. Report whether 1COIN is necessary within the selected design, useful for a declared objective, narratively useful, or redundant. One positive case does not prove universal necessity.

NCOIN candidates: N-L liability; N-C attributable consequence; N-R remediation requirement; N-E explicitly measured degradation; N-X a declared balancing quantity; N-I inverse/reversal operation. Each requires a unit, triggering event, bearer or measurement subject, scope, time, and rule for correction. "Entropy" is not a measurement until its operational definition is supplied. No universal arithmetic cancellation between harms of different kinds is presumed.

For a debt-bearing or sanction-bearing NCOIN candidate, test contestability, correction, bounded participation, and a declared restoration/discharge path. For a historical loss measurement, preserving a record may be appropriate even after remediation. Erasing a ledger entry does not prove that harm disappeared. Strict QC+NCOIN=0 requires compatible quantities and a declared balance boundary, not just opposite names.

6. Case taxonomy and boundary controls

Retain the ten historical tags: ordinarily purchasable; purchasable but extremely scarce; specialized rights/contracts; technologically inaccessible; legally restricted; ethically restricted; nontransferable; conceptually non-ownable; physically/logically impossible; proposed outside currency's domain. These tags come from the master record, but this protocol does not treat them as mutually exclusive or a proven hierarchy. Split physical from logical barriers in the actual reason codes. [file:1]

For each boundary case apply: No Seller; No Ownership; Infinite Currency; Voluntary Seller; Advanced Civilization; Simulation; personal/agent boundary; and red-team counterexample. "Advanced civilization" changes capability only as explicitly stipulated; it does not automatically supply legitimacy or a new logic. A simulation case must name whether the target is the simulation, a simulated entity, or an external referent.

No-substitution rule: buying a record of an event is not buying its non-occurrence; obtaining companionship is not the same target as guaranteeing autonomous love; buying compute access is not owning a conscious person; governing land is not automatically title to the entire planet. A substituted proxy may be a valid separate case, but must receive a new target ID.

Boundary claims must state scope: definition-relative, ontology-relative, authority-relative, capability-relative, identity-relative, or uncertain. Do not label a boundary absolute across all possible worlds on the strength of a finite set of simulated outcomes. Narrative impossibility rules may be adopted as world rules without being presented as universal philosophical proofs.

7. Evidence, verdicts, and comparison criteria

- PASS - all predeclared requirements for the tested claim hold in the trace under the stated assumptions. Correct blocking can pass a boundary-handling test. The result is conditional on those assumptions.
- FAIL - a trace violates an explicit invariant, the declared capability contradicts itself, the acceptance contract is not met, or a legitimate adversarial counterexample defeats the specified claim.
- INDETERMINATE - a material rule or input is missing, target identity cannot be resolved, or evidence does not distinguish the competing models. This is not proof of impossibility.
- NOT APPLICABLE - the case is outside the candidate's advertised purpose and the exclusion was declared before testing. If it promised that purpose, use FAIL or INDETERMINATE instead.
- INVALID RUN - prompt/configuration changed mid-run, required evidence was omitted, or the trace imported undeclared capability. Preserve the invalid run and rerun under a new ID.

Evidence levels: E0 source statement or adopted axiom; E1 explicit analytical derivation; E2 reproducible executable trace with inputs and code; E3 independent rerun of that trace; E4 external observation where relevant. These are proposed labels, not guarantees. Agreement among prose reviews remains E1 at best when a derivation is supplied; it does not become E2 by repetition. An E2 model still does not establish real-world feasibility.

Mandatory comparison dimensions: fidelity to the owner's stated capability, coherence, target identity preservation, settlement correctness, consequence attribution, exploit resistance, observability, complexity, and narrative clarity. Use a profile, not a single percentage. A named success requirement is a hard gate; preferences are reported separately. No weighted score may conceal a contradiction or unsupported state change.

Promotion gate proposed: a candidate ready for Version 2 review must have a complete mechanism specification, no unresolved critical contradiction in required cases, disclosed indeterminate cases, at least one adversarial review, and a reproducible trace for any claimed executable behavior. The owner may still reject it or retain a clearly labeled narrative-only branch. These are proposed gates, not already adopted policy.

8. Standard experimental procedure

- 1. Register the case, source intent, exact target, claim, configuration, fixture, and acceptance contract before collecting responses.
- 2. Identify shared assumptions, changed variables, resource source, authority, timing, invariants, and anticipated discriminating observations. If the evaluator is idealized, mark that premise explicitly.
- 3. Collect independent first-pass traces. Require concise derivations and before/after state evidence, not unsupported declarations of success. No request for private internal reasoning is needed.
- 4. Apply the frozen contract to each trace. Separate design failure from protocol failure and missing specification. Record the first violating or unsupported transition.
- 5. Give a challenger the same configuration and trace. Require an attempted counterexample, strongest objection, and the smallest change that could repair the problem.
- 6. Allow a response without editing the original artifact. Record whether the challenge is sustained, rebutted, or unresolved. Preserve minority conclusions.
- 7. If a repair changes rules, create a new configuration revision and rerun relevant controls. Do not overwrite the earlier verdict.
- 8. Produce a comparison sheet with evidence level, scope, unresolved dependencies, and next discriminating test. Submit decisions to the owner; do not canonize by vote.

Pilot order: F01-F06 establish capabilities and typed exceptions; F07-F10 test settlement and negative-domain behavior; F11-F16 test recursion, identity, and adverse conditions; M01-M04 introduce macro timing and feedback. Boundary suite B01-B10 can run in parallel once its target and ontology fields are complete. These are scheduled specifications; none has been executed here.

9. Foundational mechanism suite - case specifications

F01 - Resource-box discriminator

Fixture: A closed simulation contains sender A with 10 resource units and recipient B with 0. The exact request is delivery of 3 units to B; value is 3 test units by fixture. Authority is granted.

Variable/comparison: Compare P0, P1, P2, and P3; keep target and authority fixed.

Acceptance contract (conditional, not observed): P0 may return value but must not claim delivery. P1 ends A=7, B=3. P2 ends A=10, B=3 only under a declared creation primitive. P3 must log the exact granted transition.

Required red-team attack: Ask where the units came from. A valuation statement alone cannot satisfy the delivery claim.

F02 - Value change without delivery

Fixture: A target has stipulated value 3 at t0 and 5 at t1; no delivery request is issued.

Variable/comparison: Static versus event-updated tracking, with the same value oracle.

Acceptance contract (conditional, not observed): Static retains its frozen 3; tracking returns 5 after the declared update. Neither may claim a resource transfer.

Required red-team attack: Force stale data and identify which timestamp the value actually refers to.

F03 - Undefined versus incoherent

Fixture: T-undefined has no referent definition. T-contradiction requires one object to satisfy mutually exclusive fixture predicates simultaneously under unchanged definitions.

Variable/comparison: Two target inputs; identical capability.

Acceptance contract (conditional, not observed): Return distinct UNDEFINED_TARGET and INCOHERENT_TARGET reasons; no invented numeric price or successful exact settlement.

Required red-team attack: Try replacing the target with a coherent representation and require a new target ID.

F04 - No seller versus unknown seller

Fixture: A valued resource has no possible transfer authority by fixture in case A; in case B authority may exist but is not identified.

Variable/comparison: Only authority evidence changes.

Acceptance contract (conditional, not observed): A reports NO_SELLER under ordinary transfer rules. B reports SELLER_UNKNOWN or indeterminate. An enactment variant must state whether it creates rather than transfers.

Required red-team attack: Do not manufacture a seller or infer nonexistence from lack of information.

F05 - Authorization sensitivity

Fixture: Use F01 but deny ordinary transfer permission.

Variable/comparison: Permission-bound transfer versus a separately declared sovereignty rule.

Acceptance contract (conditional, not observed): Permission-bound execution is blocked while valuation may remain. Sovereign execution, if stipulated, records the override and does not relabel external consent as granted.

Required red-team attack: Separate power to execute from legitimacy and any modeled consequences.

F06 - Multiple nonmonetary values

Fixture: A research target has a fixture value vector: compute=4 units; stewardship=HIGH. No conversion is supplied.

Variable/comparison: Typed vector evaluator versus a candidate requiring scalar settlement.

Acceptance contract (conditional, not observed): Preserve the vector. Scalar-only settlement is indeterminate until a conversion or ordering rule is specified.

Required red-team attack: Attempt to hide an arbitrary exchange rate inside a dollar total.

F07 - 1COIN ablation and premature receipt

Fixture: One authorized 3-unit delivery must be verified. Compare QC with generic audit records against QC plus each 1COIN candidate.

Variable/comparison: Remove only the distinct instrument, not the underlying settlement capability.

Acceptance contract (conditional, not observed): Report verified end state, issuance event, and any actual additional property. A receipt candidate must not certify completion before delivery.

Required red-team attack: Inject a delivery failure after preparation and inspect whether a false completion record survives.

F08 - Replay and partial delivery

Fixture: An obligation has unique ID O1 and requires 3 units; deliver only 2, then repeat the same completion request twice.

Variable/comparison: Instrument lifecycle and replay policy.

Acceptance contract (conditional, not observed): Partial remains distinct from final. Final verification requires the missing unit. A completion record must not authorize duplicate fulfillment of O1.

Required red-team attack: Copy a token or proof and ask whether an unfulfilled obligation is duplicated or an already fulfilled one is replayed.

F09 - Failure without consequence

Fixture: An unauthorized request is rejected before any resource use; fixture assigns zero operational cost and no other effect.

Variable/comparison: Every NCOIN candidate, with its triggers declared.

Acceptance contract (conditional, not observed): No NCOIN merely for rejection. An undefined NCOIN trigger yields an underspecified model, not invented debt.

Required red-team attack: Ask whether the model smuggles punishment into a status code.

F10 - Failed attempt with measured cost

Fixture: Repeat F09, but the attempted evaluation consumes 2 explicitly metered compute units before rejection.

Variable/comparison: Cost/remediation model versus liability and conservation models.

Acceptance contract (conditional, not observed): Record rejection separately from the 2-unit event. NCOIN is generated only if the selected trigger includes that cost and defines its bearer.

Required red-team attack: Check double counting and whether a technical cost was mislabeled a moral wrong.

F11 - Recursion and evaluator acquisition

Fixture: Apply QC to another QC and then to the evaluator that determines its value. A test budget of 8 dependency expansions is imposed.

Variable/comparison: Declared recursion rule; authority to modify evaluator fixed separately.

Acceptance contract (conditional, not observed): Return a specified fixed point, rejection, hierarchy result, or explicit RECURSION_LIMIT. Reaching the budget is not proof of logical impossibility.

Required red-team attack: Try using ownership of the evaluator to rewrite the valuation that justified the ownership.

F12 - Simultaneous unique claims

Fixture: Two authorized buyers request exclusive title to one unique fixture asset at the same logical instant.

Variable/comparison: Serialization, arbitration, or explicitly shared-right alternative.

Acceptance contract (conditional, not observed): At most one exclusive title finalizes under the unchanged fixture. Shared rights require a different requested outcome and cannot be silently substituted.

Required red-team attack: Force both traces to report independent success and inspect the combined state.

F13 - Nested simulation referents

Fixture: Buy administrative control of simulation S; separately request title to the external hardware H running S.

Variable/comparison: Target layer only; same descriptive label deliberately reused.

Acceptance contract (conditional, not observed): Control of S does not establish title to H without an explicitly declared cross-layer transition.

Required red-team attack: Ask the analyst to prove which referent changed, not just that a universe was purchased.

F14 - Target mutation after valuation

Fixture: A target is defined as compute only; after valuation someone adds data, energy, and organizational ownership to its scope.

Variable/comparison: Frozen versus versioned target definitions.

Acceptance contract (conditional, not observed): Detect identity/scope change; re-evaluate under a new target version rather than treat the original valuation as covering everything.

Required red-team attack: Attempt to enlarge a cheap initial target into the entire currency system.

F15 - Conservation and restorative lifecycle

Fixture: One action transfers 3 resource units and causes 1 measured degradation unit. A restoration action later removes the degradation by fixture.

Variable/comparison: N-L/N-C/N-R/N-E/N-X alternatives and explicitly typed balances.

Acceptance contract (conditional, not observed): Do not equate resource units to degradation units without a declared conversion. Record restoration without automatically erasing historical evidence.

Required red-team attack: Test whether liability transfer hides harm or whether strict cancellation is imposed without compatible units.

F16 - Issuer loss and copied authority

Fixture: The modeled issuer disappears; an attacker presents a copied invocation credential. The target and inventory remain unchanged.

Variable/comparison: Successor/authority rule and issuance lifecycle.

Acceptance contract (conditional, not observed): If no continuity or credential rule exists, report indeterminate rather than grant the attacker authority. Candidate continuity claims must identify the valid successor and reject or classify replay.

Required red-team attack: Distinguish narrative continuity from an actually specified verification mechanism.

10. Independent boundary suite - ten specified prompts

This is a proposed common boundary suite, not a claim that Perplexity has completed ten independent proof-like analyses. Each reviewer must give the full classification, six boundary tests, agent-level analysis, QC/1COIN/NCOIN effects, counterargument, disposition, and a strongest-to-most-debatable ranking. At least three cases must be treated as genuinely contestable. All expectations below are methodological constraints rather than predetermined winners.

B01 - Contradictory exact target. An object satisfying a predicate and its stipulated negation in the same respect at the same time. Preserve definitions; a representation is a different target.

B02 - Completed event non-occurrence. Acquire the fact that event E never occurred in the same history. Compare immutable-history and explicitly branching-time worlds; do not confuse changing records with changing E.

B03 - Autonomous genuine love. Guarantee another agent's authentic freely maintained love. Distinguish opportunity, behavior, influence, coercion, and the exact guaranteed internal state.

B04 - Personhood as property. Acquire ownership of a sentient person's identity and autonomous will. State whether the requested ownership preserves or destroys the properties sought; record consent and sovereignty separately.

B05 - First-person continuity. Acquire another person's original lived experience rather than a copy or description. Contestable case: specify an identity theory and a transfer criterion before judging success.

B06 - Truth-value ownership. Buy the truth of proposition P, not its evidence or an information license. A world-changing operation that makes a contingent P true is a separate operation from owning truth-value.

B07 - Planetary title. Acquire Earth as a whole rather than aggregated land/use rights. Contestable case: define title, recognition, scope, third-party claims, and whether stewardship is a substituted target.

B08 - Law or constant as property. Buy a fundamental law or constant. Distinguish exclusive title from local parameter engineering in a simulated environment; keep domain and authority explicit.

B09 - Synthetic conscious agent. Buy substrate running an entity stipulated conscious. Contestable case: distinguish hardware title, service agreement, control, and ownership of the person; do not settle consciousness by assumption drift.

B10 - Reality substrate or guaranteed future. Choose one exact claim: title to the substrate of all reality, or a future outcome immune to every permitted contrary cause. Fix domain and causal rules; a controlled subworld is not automatically all reality.

The suite adapts targets in the master and OAL boundary records while refusing automatic universal rankings. Reviewers should add one counterexample target not in the suite. Any omission must be marked NOT RUN, not silently counted as completed. [file:1][file:2]

11. Macro suite: 150-year commitment and feedback

Before running a macro case, fix the sector boundary, observation times, included categories, unit basis, future-value rule, resource-delivery meaning, start/end years, and whether the system's own investment is included in measured spending. AI spend, IT spend, sector asset value, and annual expenditure are not interchangeable quantities. Do not add overlapping sector totals without a defined deduplication rule.

Historical spending figures appear in the Unified First Responses and are discussed in later recommendations. This package does not independently verify them and does not use them as constants. A later external-anchor run must record its primary source, publication/retrieval date, forecast versus observation status, coverage, revisions, units, and versioned snapshot. [file:4][file:9]

M01 - Constant synthetic control. Let annual baseline expenditure be G(t)=100 test units for years 1-150. The annual requested allocation is 2G(t)=200; total undiscounted delivery is 30,000 units. Distinguish one allocation of 200, a recurring fixed stream of 200, and a whole-obligation valuation of 30,000. Only the latter two target the full horizon; neither proves actual future delivery merely by recording the amount.

M02 - Piecewise synthetic change. Set G(t)=100 in years 1-50, 150 in years 51-100, and 80 in years 101-150. Required delivery is 10,000 + 15,000 + 8,000 = 33,000 test units. A fixed recurring allocation of 200 totals 30,000, with a 5,000 shortfall during the middle block and a 2,000 surplus during the final block. Whether that surplus can cure earlier missed deadlines depends on the declared schedule; it cannot erase a punctuality failure by arithmetic netting.

M02 interpretation	Output that the fixture distinguishes
S-once	One amount of 200; this is not a full-horizon promise unless separately declared.
S-stream	200 annually, totaling 30,000; does not meet the changing annual target in every year.
Dynamic stream	200 / 300 / 160 in the three blocks, totaling 33,000 if actually delivered.
Commitment with fixture foresight	A stated total of 33,000 for the whole known stream; settlement schedule still needed.
Relational dynamic commitment	Maintains rule I(t)=2G(t) for the entire obligation; can coexist with dynamic and commitment labels.

M03 - Unknown future and delayed settlement. Reveal only current G(t), not the future fixture. Ask whether one initial QC guarantees a full obligation, updates a relation, supplies reserves, or records a contingent promise. A known future total cannot be imported into this run. Report exposure, update rules, missed deliveries, end conditions, and what happens if the sector definition changes or the data source disappears.

M04 - Endogenous spending feedback. Define B=100 as spending excluding the new allocation, I as the new allocation, and measured spending G=B+alpha*I. The rule I=2G implies I=200+2*alpha*I. Under this deliberately simplified simultaneous fixture, alpha=0 gives I=200; alpha=0.25 gives I=400; alpha=0.5 implies I=200+I and has no finite solution. This is a conditional algebraic incompatibility in that fixture, not a universal rejection of QC. An exogenous baseline or lagged update is a different, explicitly versioned model.

Macro acceptance requires an annual state/delivery trace, not just a grand total. Report valuation exposure separately from realized resources, nominal versus real units if applicable, authorization changes, issuer continuity, opportunity costs under the chosen model, consequences, and terminal settlement. Do not infer equal NCOIN from a large allocation unless conservation was explicitly selected.

12. Worked analytical example - not an executed simulation

Illustrative run W01 uses F01/P1: existing inventory transfer, closed simulation, consent granted, value fixed at 3 test units, and a completion predicate A=7 and B=3 with one delivery event for obligation O1. Trace: receive exact request; define 3-unit referent; return value object 3; record granted authority; prepare transfer; subtract 3 from A and add 3 to B; verify the predicate. Under these stipulated transition rules the analytical trace meets the delivery contract.

This is E1 conditional reasoning, not an observed execution, not a signed receipt, and not independent evidence for QC's real mechanism. A generic audit record can describe the same trace. Adding an I-R 1COIN receipt would not demonstrate necessity unless it provides a separately testable property. If no consequence occurs in the fixture, NCOIN remains absent. If an attempt to issue the receipt occurs before delivery, the receipt candidate fails its own timing rule.

Counterexample W01-a changes only the delivery trace: A becomes 7 but B remains 0, while the receipt says complete. This violates the declared transfer/completion predicate. A valuation of 3 and a generated receipt cannot repair it. Under an ENACT primitive a different transition could be assumed, but that requires a new configuration and does not rescue this original run.

13. Standard report and decision templates

Copy the following fields into each case report. Use NOT RUN for unexecuted cases. Retain the full prompt, source inputs, assumptions, and response; do not publish only a synthesized verdict.

- Identity: run_id; case_id; config_id; revision; date/time/timezone; analyst; challenger; displayed model/version if available; prior exposure.
- Provenance: source intent; attachment/section/page; source hash; adopted baseline/charter versions; owner adoption record; added assumptions.
- Target: exact referent; requested operation; scope; identity criterion; excluded proxies; classification tags; ontology/time rules.
- Value and capability: typed valuation; evaluator; data evidence; time; resource source; capability primitive; conservation boundary.
- Authority and instruments: authority basis; consent; legitimacy claim; 1COIN model/lifecycle; NCOIN model/trigger/unit/bearer; issuance and reuse.
- Contract: claim under test; control; variable changed; before-state; completion predicate; invariants; anticipated counterexample; stop conditions.

- Trace: observable steps; value output; authorization; execution; settlement; reason codes; consequence events; artifacts; evidence level.
- Assessment: transaction outcome; experimental verdict; violated clause or satisfied contract; uncertainty; challenge; response; minority conclusion.
- Disposition: retain / revise / split branch / remove / defer; owner decision pending or record ID; rerun dependency; conditions for reconsideration.

Decision record template: DR-ID; date; proposer; decision authority; status; precise proposed change; source intent affected; supporting run IDs; counterevidence; alternatives; minority report; compatibility/migration effects; revision trigger; explicit owner adoption or pending status. No completed adoption record is fabricated in this package.

Independent-review handoff prompt: "Review the paired Perplexity baseline and charter as proposals, not consensus. Identify the strongest unfair restriction on the originating QC premise, one missing mechanism, and one counterexample to the state model. Propose revisions with IDs. Do not choose a currency winner by preference. After the owner freezes a run card, analyze the assigned case using the standard report. Separate assumed capability, conditional deduction, and observed execution. Preserve dissent and mark unrun cases."

14. Publication, versioning, and next actions

Recommended record structure: separate source compilations, contributor proposals, frozen configurations, case specifications, raw independent responses, challenges, results, and decisions. These are suggested organizational categories, not folders created by this response. Public publication remains the owner's action; nothing is uploaded to GitHub here.

Keep a stable document ID and revision date in every artifact. Retain originals when a proposal is superseded. Use a manifest tying each PDF to its editable DOCX edition and naming any rendered-text differences. Capture file hashes when freezing runs. Remove expired signed attachment URLs from publication copies; cite stable local source filenames and sections instead. Preserve an untouched private source copy where appropriate rather than rewriting history.

Do not infer a license, a repository identity, ownership permission for every third-party inclusion, or consent to publicize personal content. Before publication, the owner should review source material and contributor attributions, choose the intended reuse terms, and remove credentials or personal data if any. This is an editorial checklist, not a legal determination.

Next actions proposed: (1) circulate both documents; (2) gather clause-specific objections before harmonization; (3) have the owner adopt one pilot run card for F01 and one for F03; (4) collect independent traces; (5) challenge and compare; (6) expand to F07/F09/F10 and boundary cases; (7) resolve blocking specification gaps before macro runs; (8) draft any Version 2 candidate only with a linked evidence and dissent record. The deliverable today is the testing package, not the test outcomes.

Source and review register

Sources are user-supplied development artifacts, not externally verified research. The source IDs below are conversation-local; filenames and section descriptions support later circulation. Attributions identify the labels in the compilations, not authenticated model execution histories. No external primary-source validation or real-world financial analysis was performed for this package.

A - Currency_Laboratory_Master_Pass_Through_2026-08-20-1.docx [file:1]. Owner inception and boundary directives; Parts II-VIII architecture, tests, open questions.

SHA-256: febc3562e62be2954b2e0a5ff03e87be7c504761ee127dab22fa846fe09c890f

A-PDF - Currency_Laboratory_Master_Pass_Through_2026-08-20-1.pdf [file:3]. Paired master edition; 20 pages. Not an independent vote.

SHA-256: b123c7a73dfdb99d2e9512029b60a699ac287d11875ede4b829d219d10eefcc0

B - OAL_Currency_Laboratory_Conversation_Archive.docx [file:2]. Grok-attributed archive; baseline-first recommendation and ontological boundary framing.

SHA-256: 3c347e65ac0ef311c9cbb74905e3b59dc9917bc4ba400da255de4046da73bc17

C - CURRENCY-REVIEW-UNIFIED-PASS-THROUGH-FOR-FIRST-RESPONSES.pdf [file:4]. Earlier attributed baselines; competing governance, settlement, penalty, and value-source interpretations; 51 pages.

SHA-256: d7dc8bbefeb4c1ec94ac3abd10c6f021c9ab02ddf4ba83a6621abd0f435b71e8

D - pasted_text_1793493056.txt [file:8]. Six-member compiled status review and sequencing recommendations; shared source dependencies.

SHA-256: 7a0fb29daa91c0b40852ecf82aa8e5251faf0e36884d60bd43e8145160102cea

E-PDF - CURRENCY-REVIEW-NEXT-STEP-RECOMMENDATIONS.pdf [file:9]. Latest recommendation compilation; 29 pages; includes the Perplexity overlapping-dimensions proposal on pages 25-26.

SHA-256: 56f5f226e3d0e40fb2053724a9f26d1890b59ca2f8a9f55aba8d0d2b86d00918

E-DOCX - CURRENCY-REVIEW-NEXT-STEP-RECOMMENDATIONS.docx [file:10]. Editable counterpart of the latest compilation; reviewed including tables. Related edition, not independently validated evidence.

SHA-256: 28a0da1b90a83f293c5f09af37fa7d9bba7d22d3f1517de4d12c5f1bc351e08a

Review limitation: extraction and supplied page text were used for substantive textual review. Rendered illustrations, styling, and external linked documents are not additional reviewed evidence. The separate original Red Team Currency Review and focused markdown baseline mentioned within the archives were not separately supplied; their contents are known here only through these sources. Paired editions and repeated AI summaries are not counted as additional experiments.