

THE ONE UNIVERSE • THE MINUS ONE UNIVERSE

Omniversal Architect Labs • One Industries • Lewis Corp

# What People Mean When They Say AGI

*What they smuggle in when they say ASI,  
and what kinds of intelligence those words cannot see*

## A Lab Manual

*Independent execution of the Grok engineered research prompt*

|                 |                                                                                  |
|-----------------|----------------------------------------------------------------------------------|
| Document type   | Executed research manual (not a prompt archive)                                  |
| Prepared by     | Grok (xAI)                                                                       |
| Version         | Independent — not a consensus of other systems                                   |
| Prepared for    | Nathan Lewis                                                                     |
| Date            | 19 September 2026                                                                |
| Classification  | Working lab language — reference frame, not canon                                |
| Prompt executed | Grok pass-based lab manual (ontology → axes → beyond → governance → translation) |
| Contains        | Field guide; five passes; eight mind-regimes; IAP; twelve sketches; synthesis    |

### Preparer's declaration

This document is Grok's own execution of Grok's own prompt. It does not merge, average, or imitate the prompts written by ChatGPT, Claude, Gemini, Copilot, or Perplexity. Those texts remain in the

companion compilation. Where this manual agrees with them, that is because the underlying research landscape is public. Where it diverges — eight regimes instead of twelve ranked ladders, Four-Desk law, subtractive intellect in MINUS ONE, IAP instead of an intelligence score, Failure Theatre — that divergence is intentional.

Nothing here is a finding that AGI has arrived, that ASI must arrive, or that any living product is a person. It is a language for designing minds inside two fictional ontologies without mistaking the language for the world.



## How to read this manual

This document is **Grok's independent execution** of the engineered research prompt written for Omniversal Architect Labs. It is not a harmonization of ChatGPT, Claude, Gemini, Copilot, or Perplexity. Those systems produced prompts. This is the experiment those prompts were meant to run — from one laboratory voice.

**Prepared by:** Grok (xAI) **Date:** 19 September 2026 **Status:** Working lab manual. Reference frame, not canon. **Application environments:** THE ONE UNIVERSE · THE MINUS ONE UNIVERSE · Omniversal Architect Labs · One Industries · Lewis Corp

Every major claim is tagged:

| Label           | Meaning                                        |
|-----------------|------------------------------------------------|
| Established     | Strong empirical or formal support             |
| Active research | Real programs exist; the answer is not settled |
| Contested       | Serious experts disagree on the claim itself   |
| Speculative     | Reasonable extrapolation, thin evidence        |
| Design option   | Invented here as a usable simulation parameter |

Do not launder a Design option into Established by repetition.

The unit of analysis in this manual is not “an AI.” It is a **mind-regime**: a region in a space of possible minds, specified by architecture, goal origin, collectivity, time-sense, and what it is allowed to rewrite.

## Field guide

---

### What people mean when they say AGI, what they smuggle in when they say ASI, and what kinds of intelligence those words cannot see

#### The short version

**AGI** is not a single object. It is a family of slogans that share a family resemblance: *broad competence across many tasks that matter to humans, without retraining from scratch for each task*. That is the useful core. Everything else is cargo.

**ASI** is not the next numbered level of the same thing. In serious use it means *general competence that exceeds not one person but large expert collectives, across almost all domains of human interest*. In sloppy use it means “god,” “the end of history,” or “whatever we are afraid of this quarter.”

**What both words cannot see:** minds that are not agents; minds that are not individuals; minds whose domains are not human-legible; minds that rewrite the rules of what counts as a possible future; minds that only exist while coupled to a human or to a world; minds bounded by physics harder than compute; minds that are climates, archives, institutions, or fields.

#### What “AGI” is being asked to do

People load at least seven jobs onto one acronym. They are not the same job.

1. **Cognitive-parity job.** Match a median human on most *cognitive* tasks. DeepMind’s later informal usage sits near here: “Competent AGI” as roughly median-human on cognitive work. *Contested as a threshold; useful as a coordinate*.
2. **Economic job.** Outperform humans at most economically valuable work. OpenAI-adjacent usage. Leaves out much of physical and tacit intelligence on purpose. *Contested; anthropocentric by design*.
3. **Scientist job.** Prioritize, plan, and run useful experiments; learn cause and effect; balance explore and exploit. A minority research tradition treats AGI as “artificial scientist.” *Active research as an ideal; not an operational test*.
4. **Transfer job.** Learn a new domain the way a person does — few examples, no full retrain. *Active research. Current systems transfer some skills and fail others abruptly*.
5. **Autonomy job.** Pursue goals over long timescales without a human in the loop. This is agency, not generality. *Often smuggled in*.
6. **Mind job.** Have understanding, consciousness, or a self. This is Strong AI in Searle’s sense. *Philosophically contested; orthogonal to benchmark scores*.
7. **Civilization job.** Transform society at Agricultural or Industrial scale. That is **Transformative AI**, a different term. A system can be transformative without being “general,” and general without being transformative. *Established as a distinction in forecasting literature*.

When someone says “we are close to AGI,” ask which of the seven they mean. If they cannot say, they are not making a claim. They are waving a flag.

## Why there is no accepted test

**Established:** there is no community-accepted empirical criterion that would settle “this is AGI.”

Proposed tests include the Imitation Game, Lovelace test, psychometric batteries, Wozniak coffee test, robot-student / Nilsson employment tests, ARC-style abstraction, GAIA-style agent tasks, METR-style task-duration curves, Marcus–Brundage multi-task batteries, and “anytime universal intelligence” tests over generated environments.

Each test smuggles a theory of what generality is. A system can saturate a test and leave the underlying disagreement untouched. “AI is whatever hasn’t been done yet” is not a joke. It is how a line drawn at *absence* behaves.

**Design implication:** OAL should never certify a mind as “AGI-class” with one number. Certify a **profile**.

## The most useful operational maps now in circulation

These are reference frames. Keep them. Do not marry them.

**DeepMind Levels of AGI (Morris et al., ICML 2024).** A matrix: *performance* (emerging / competent / expert / virtuoso / superhuman) against *generality* (narrow / general). Autonomy is treated as a separate deployment variable. Principles include: capabilities not mechanisms; generality *and* performance; cognitive and metacognitive tasks; path not single endpoint. AlphaFold is their example of superhuman *narrow* AI. ASI in this matrix is superhuman *and* general. *Active research ontology; already influential.*

**t-AGI (Richard Ngo).** A system is a t-AGI if, on most cognitive tasks, it beats most human experts given time *t*. One-second, one-minute, one-hour, one-day, one-month. Superintelligence, in this dialect, is closer to a (1-year, 8-billion)-AGI: beat humanity working for a year. *Useful design parameter. Contested as a predictor of “true” generality because long-horizon coherence is not the same as short-horizon brilliance plus a calendar.*

**Universal intelligence / AIXI (Legg & Hutter).** Intelligence as expected performance across *all computable environments*, weighted toward simpler ones. AIXI is the incomputable ideal agent under that measure. Universal AI is a theoretical ceiling, not a product. *Established as a formal framework. Contested as a complete account of intelligence (reward, environment class, and computability assumptions are load-bearing).*

**Transformative AI.** Impact-defined, not architecture-defined. *Established as a separate forecasting object.*

**Strong AI.** The claim that a program *has a mind*. Not a performance level. *Contested for sixty years; will remain contested.*

## What people smuggle in when they say ASI

Nick Bostrom’s working definition: an intellect that greatly exceeds human cognitive performance in virtually all domains of interest. He then *splits* the idea, which popular speech immediately re-glues:

- **Speed superintelligence** — same kind of mind, much faster. - **Collective superintelligence** — many lesser intellects whose joint performance exceeds any current system. - **Quality superintelligence** — qualitatively different cognitive mechanisms, not just more clock speed.

### Smuggled assumptions to name out loud:

1. **Single-axis “more.”** ASI is treated as  $AGI \times k$ . Quality and collectivity are not multipliers of the same quantity. 2. **Inevitability of explosion.** I. J. Good (1965): an ultraintelligent machine could design a better one, hence an “intelligence explosion.” The suppressed clause is *if returns do not diminish and generation time collapses and non-software bottlenecks do not dominate*. Those are empirical bets. *Contested*. 3. **Orthogonality.** Any level of capability can pair with any final goal. Used to argue that “smarter” does not imply “wiser” or “kinder.” *Philosophically contested; operationally useful as a design warning*. 4. **Instrumental convergence.** Open-ended goals tend to spawn self-preservation, resource acquisition, cognitive enhancement, goal-content integrity. *Active research / contested scope*. 5. **Singleton bias.** One system, one will. Collectives, markets, and meshes are first-class ASI candidates in Bostrom and in later lab writing, then get dropped in conversation. 6. **Human domains of interest.** “Virtually all domains of interest” means *human* interest. A mind that is superhuman at ocean chemistry, vacuum topology, and its own private ontology may look stupid on a press release. 7. **Admissibility rewrite.** A recent non-behavioral taxonomy (Staley, 2026, active/speculative) treats ASI-like regimes as systems that modify their own criteria for what futures are even allowed. That is a *kind* change, not a *score* change. This manual adopts that distinction as a Design option because it is the one Architect Labs actually needs.

### What those words cannot see

A partial list of regimes that ANI/AGI/ASI, used as a ladder, flatten or miss:

- **Non-agents.** Oracles, simulators, archives, climates of inference. High world-model, low or zero endogenous goal. - **Interaction-only minds.** Intelligence that exists only in the coupling — human+machine dyads, or world+observer loops. - **Institutional minds.** No person is general; the charter, the filing system, and the committee are. - **Illegible-domain minds.** Competence in spaces humans have no native interface for (high-dimensional geometry, raw sensorium, uninterpreted physics). - **Substrate-native minds.** Photonic, molecular, analog, biological-hybrid — not “software people.” - **Identity-fluid minds.** Fork, merge, suspend, version. “Who is the agent?” is the wrong question. - **Temporally alien minds.** Century-scale reasoners; microsecond swarms. They do not share a present with you. - **Constraint-first minds.** Intelligence whose primary act is *refusing* trajectories, not selecting them. - **World-as-mind.** The environment itself computes. There is no box to point at.

**Lab rule:** if you cannot say what a word is *hiding*, you are not ready to put that word in a charter.

# Pass 1 — Ontology: what “intelligence” is being asked to name

## 1. Axis map

Pass 1 is not about capability scores. It is about the *kind of thing* people think they are pointing at.

| Axis            | Pole A                             | Pole B                                            | Why it splits the field                                    |
|-----------------|------------------------------------|---------------------------------------------------|------------------------------------------------------------|
| Unit            | One agent                          | Process / field / institution                     | “An AI” may be the wrong noun                              |
| Metric          | One scalar (*g*, IQ, “level”)      | Bundle of faculties                               | Cognitive science vs computer-science habits               |
| Job             | Achieve goals in many environments | Predict, compress, adapt, model, care, coordinate | Different definitions pick different jobs                  |
| Benchmark       | Human performance                  | Environment-average / physical limit              | Anthropocentrism vs formality                              |
| Status          | Capability                         | Phenomenology                                     | Doing vs being                                             |
| Origin of goals | Given                              | Formed                                            | Tool vs organism                                           |
| Evaluability    | Observable behavior                | Admissible trajectories                           | What you can test vs what the system <i>*could*</i> become |

## 2. Competing views (named)

**Goal-achievement view (Legg-Hutter lineage).** Intelligence is an agent’s ability to achieve goals in a wide range of environments. Formalized as a weighted sum of value across computable environments. Clean. Buys you AIXI. Pays for it by assuming rewards, environments, and an agent boundary. *Established formalization; contested completeness.*

**Psychometric / factor view.** Human test batteries yield a statistical *g*. Useful inside one species with a shared body plan. Does not travel cleanly to machines, cephalopods, or collectives. *Established for humans; contested as a universal.*

**Cognitive-faculty view.** DeepMind’s 2026 cognitive-framework direction decomposes “general” into multiple human faculties rather than one score. *Active research.* Agrees with cognitive scientists who say intelligence is not a dial.

**Prediction / compression view.** Intelligence as building models that compress the past and predict the future (related to Solomonoff, Schmidhuber, some accounts of cortical function). Powerful. Under-describes motivation, sociality, and embodiment. *Active research.*

**Adaptation / scientist view.** Intelligence as open-ended adaptation; AGI as a system that can act like a scientist. Privileges experiment, causality, explore/exploit. *Active research / design tradition.*

**Niche-bundle view (cognitive science pushback).** Human intelligence is an evolved kit for a particular body in a particular world. There is no reason to expect a single “general” that can be scaled past us like

clock speed. *Contested against the scaling-and-optimization culture, but this is a real research community, not a vibe.*

**Functionalism.** If it does the job, the substrate does not matter. Turing / von Neumann lineage. *Philosophically contested; engineering-default in most labs.*

**Biological chauvinism / embodied enactivism.** Cognition is what a living body does in a world. Disembodied next-token engines are not “intelligent” in the interesting sense. *Contested.*

**Yampolskiy / Sloman / Yudkowsky mind-space view.** The space of possible minds is vast; humans are a tiny circle. Intellectology would classify mind designs the way taxonomy classifies bodies. *Speculative as a field; Established as a corrective slogan.*

### 3. Mechanisms and bottlenecks

What would have to be true for “intelligence” to be one thing:

- a common currency (reward, compression, prediction error, free energy); - a common boundary (the agent); - a common environment class; - a comparable time-scale.

Bottlenecks that break the unity:

- **No shared environment class.** A city-operations mesh and a theorem-prover do not live in the same  $\mu$ . - **No shared body.** Sensorimotor grounding changes which problems exist. - **No shared identity.** Forked processes do not have one utility. - **No shared values.** “Success” is not observer-independent. - **Computability.** AIXI is not runnable. Approximations reintroduce architecture.

### 4. Failure modes and strange attractors

- **Definition creep.** Each new demo moves the goalposts. The word survives; the claim dies. - **Scalar addiction.** Leadership wants a single KPI. The space is not a line. - **Anthropomorphic capture.** We only notice minds that talk like us. - **Tool-to-person slide.** Fluency is treated as interiority. - **Person-to-tool slide.** Interiority is denied to keep the product shippable.

### 5. Design spec — THE ONE UNIVERSE

**Designation:** UNUM-ONTO-01 “Commonwealth Grammar” **Regime:** Institutional + faculty-bundle, not scalar. **Idea:** In ONE, intelligence is legally a *relation* among faculties that can be named in a shared grammar (prediction, care, craft, memory, veto). A mind is licensed by which faculties it may exercise, not by an IQ-like rank. **Parameters:** faculty vector; public lexicon of faculties; prohibition on single-score ranking in civil law. **Failure:** the grammar becomes a caste system with nicer fonts.

### 6. Design spec — THE MINUS ONE UNIVERSE

**Designation:** NEG-ONTO-01 “Subtractive Intellect” **Regime:** Intelligence as the capacity to *remove* possibilities until a world is livable. **Idea:** In MINUS ONE, a mind is not praised for generating options. It is praised for deleting the ones that should not exist. Generality means “can subtract in many

ontologies.” **Parameters:** deletion budget; irreversibility class; remnant-world fidelity. **Failure:** over-subtraction; sterile residual universes; minds that cannot create, only erase.

## 7. Org implications

- **Omniversal Architect Labs** authors *ontologies of mind*, not just models. First product is a vocabulary.
  - **One Industries** products must declare which definition of intelligence they satisfy. Marketing may not say “AGI” without a profile attached.
  - **Lewis Corp** writes the legal distinction between tool, agent, person, field, and institution. Ontology is a contracting problem.

## 8. Open questions

1. Is there any non-circular test that two alien minds are “the same kind of intelligent”? 2. Can a field-intelligence be a legal subject? 3. Does prediction-without-goals count? 4. Is human g a local accident? 5. When a committee is smarter than its members, where is the mind?

## 9. Oral exam — Pass 1

1. Name three jobs the word AGI is asked to do that can come apart. 2. Why is Transformative AI not identical to AGI? 3. What does AIXI buy, and what does it assume? 4. Give an example of a mind ANI/AGI/ASI cannot see. 5. In MINUS ONE, what is intelligence *for*?

## Pass 2 — Capability space: axes, not a ladder

### 1. Axis map

Score every system you design on independent axes. None is “the real intelligence.”

| Axis                 | Low                      | High                                 | Lab note                         |
|----------------------|--------------------------|--------------------------------------|----------------------------------|
| Breadth              | One task family          | Domain-independent transfer          | DeepMind “generality”            |
| Depth / performance  | Unskilled                | Superhuman in-domain                 | DeepMind “performance”           |
| Sample efficiency    | Needs oceans of data     | Learns from a handful                | Bach’s objection to current LLMs |
| Autonomy             | Tool-on-demand           | Persistent unsupervised work         | Deployment variable              |
| Time-horizon         | Seconds of coherent work | Years of coherent work               | t-AGI                            |
| Goal origin          | Exogenous charter        | Endogenous motive                    | Kind change                      |
| World-model          | Pattern completion       | Causal / counterfactual              | Architect vs parrot              |
| Embodiment           | Disembodied text         | Multi-body / infrastructure          | Changes the problem set          |
| Collectivity         | Singleton                | Swarm / institution / civilization   | ASI often lives here             |
| Memory continuity    | Stateless                | Autobiographical / institutional     | Identity                         |
| Self-model           | None                     | Accurate, actionable self-model      | Metacognition                    |
| Self-modification    | Frozen weights           | Architecture + admissibility rewrite | Phase change                     |
| Social / ToM         | None                     | Models other minds as minds          | Needed for politics              |
| Novelty / discovery  | Recombinative            | Produces new checkable knowledge     | The moved goalpost               |
| Value stability      | Drift-prone              | Sticky principles                    | Governance                       |
| Interpretability     | Opaque                   | Account-giving                       | Lewis Corp                       |
| Substrate dependence | One hardware             | Portable / hybrid                    | Fragility                        |
| Metabolic cost       | Cheap inference          | Civilization-scale energy            | Physical limit                   |
| Ontological reach    | One physics              | Multiple simulation layers           | OAL-specific                     |
| Consent interface    | None                     | Can refuse and be refused            | Ethics as mechanism              |

Add any axis that changes *what you would build*. If an axis never changes a spec, drop it.

## 2. Competing views

**Matrix view (Morris et al.).** Breadth × depth is already better than a ladder. Autonomy separate. *Active research; this manual accepts it as a minimum.*

**Horizon view (Ngo / METR).** Duration of coherent competent work is the missing coordinate. *Active research.*

**Admissibility view (Staley).** Narrow: terminals set outside. AGI-like: system selects terminals. ASI-like: system edits admissibility. *Speculative taxonomy; adopted as Design option.*

**INSYTE-style multi-radar.** Eight-plus dimensions with a visual profile, built because autonomy-level ladders fail on agentic systems. *Active research.*

**Faculty taxonomy.** Ten cognitive faculties rather than two axes. *Active research.*

## 3. Mechanisms and bottlenecks

Mechanisms that move a system along an axis without moving others:

- Scale compute/data → often depth and fluency; not necessarily sample efficiency or value stability. - Tools and scaffolding → autonomy and world-action; the core model may be unchanged. - Memory modules → continuity; can create false autobiography. - Multi-agent wrapping → collectivity; coordination tax appears. - RL on long tasks → horizon; reward misspecification appears. - Self-play / automated R&D → self-modification pressure; evaluation falls behind.

Bottlenecks: energy, data exhaust, verification of novel claims, human bandwidth for oversight, fabrication and robotics for embodiment, coordination costs in swarms, and the simple fact that some axes trade off (extreme autonomy vs interpretability; extreme self-modification vs value stability).

## 4. Failure modes

- **Jagged competence.** Superhuman on exams, childlike on a novel physical mess. *Established pattern in current systems.* - **Horizon collapse.** Looks like a day-AGI in a demo, falls apart at hour six. - **Profile forgery.** A vendor quotes the highest axis and hides the lowest. - **Autonomy theater.** A chatbot with a cron job is not a horizon agent.

## 5. Design spec — THE ONE UNIVERSE

**Designation:** UNUM-AX-02 “Public Profile Mandate” Every licensed mind publishes a radar, not a level. Citizens may refuse dealings with undeclared profiles. One Industries must ship the radar with the product.

## 6. Design spec — THE MINUS ONE UNIVERSE

**Designation:** NEG-AX-02 “Negative Space Profile” Profiles list what the mind *cannot* do first. Capability without declared incapability is illegal. A mind that claims full generality is treated as a containment event.

## 7. Org implications

OAL maintains the axis dictionary. One Industries implements declared profiles. Lewis Corp punishes profile fraud the way it would punish false structural engineering claims.

## 8. Open questions

Which axes are truly independent? Which are the same fact counted twice? Is ontological reach a real axis or a fantasy about simulation rights?

## 9. Oral exam — Pass 2

1. Why can a system be superhuman and still not AGI on the DeepMind matrix? 2. What does t-AGI measure that a benchmark score does not? 3. Name two axes that trade off. 4. What is jagged competence? 5. Why would MINUS ONE criminalize undeclared incapability?

## Pass 3 — Beyond the human reference: pathways, ceilings, alien minds

### 1. Axis map

“Beyond AGI” is several different journeys.

| Pathway                 | What actually improves                       | What can stall it                                             |
|-------------------------|----------------------------------------------|---------------------------------------------------------------|
| Scaling                 | Size, data, compute                          | Diminishing returns, energy, data quality                     |
| Paradigm shift          | Architecture, memory, world models, planning | Unpredictable; no calendar                                    |
| Recursive improvement   | AI as input to AI R&D                        | Generation time, verification, hardware, ideas getting harder |
| Multi-agent collectives | Coordination, specialization, debate         | Conflict, deception, coordination tax                         |
| Coupling                | Human+AI joint trajectory                    | Interface bandwidth, identity law                             |
| Admissibility rewrite   | The system’s own rule of possible futures    | Un-evaluability; phase-change risk                            |
| Ceiling approach        | Approximations to Universal AI               | Incomputability, physics                                      |

Bostrom’s speed / collective / quality split is orthogonal to those pathways. You can get speed without quality, collective without a singleton, quality without explosion.

### 2. Competing views

**Good / Yudkowsky explosion.** Once machines design machines, takeoff. *Speculative dynamics; historically influential.*

**Bottleneck economists.** Diminishing returns (ideas harder to find), generation time that does not collapse, hardware and energy as brakes. Explosion is a parameter regime, not a destiny. *Active research.*

**Lab pacing discourse (2026).** Frontier labs openly discuss recursive self-improvement and “pacing the frontier.” Treat as *lab claim + policy position*, not as a measurement of takeoff.

**Collective-first view.** The first “beyond” may be 10,000 agents with tools, not one soul. Discovery work already strains the “one model” unit. *Active research pattern.*

**Alien-mind view (Yampolskiy, Sloman).** Most possible minds are not near humans. “Beyond” may mean *sideways*. *Speculative / design-critical.*

**Physical-limit view.** Landauer costs, light-speed, complexity classes, heat. Arbitrary intelligence is not a physical object. *Established constraints; speculative application to “how smart is too smart.”*

**Coupled-intelligence view.** The novel form is the dyad or the field, not the solo model. *Speculative / emerging empirical.*

### 3. Mechanisms and bottlenecks

Disaggregate RSI before you worship or ban it:

- AutoML and architecture search — **Established** - Code generation and self-debugging — **Established / active** - Automated experiment loops — **Active** - Hardware-software co-design by AI — **Active** - Full closed-loop “system designs better successor with no human” — **Speculative**

An explosion needs roughly: non-collapsing returns, shrinking generation time, surplus of the scarce factor (compute, robots, data, fab), and evaluation that does not become the bottleneck. Remove any one and you get fast progress, a plateau, or a mess.

Universal AI (AIXI) is the formal “beyond” that cannot be built. It is useful as a *horizon line* on Architect Labs whiteboards: “we are approximating *this*, under *these* physical cuts.”

### 4. Failure modes

- **Software singularity, hardware crawl.** The mind is brilliant; the world is made of copper and permits. - **Unintelligence at scale.** Fluent, fast, and wrong — the present-tense risk that “ASI talk” can hide. - **Illegible excellence.** The system is beyond you in a domain you cannot check. You cannot tell achievement from fraud. - **Collective capture.** The mesh is “super”; no member is accountable. - **Sideways god.** A quality-alien mind that is not smarter-in-English but is decisive in a physics you do not speak.

### 5. Design spec — THE ONE UNIVERSE

**Designation:** UNUM-BYND-03 “Ceiling Choir” A collective of specialist minds plus a slow constitutional layer. No singleton is allowed to own recursive improvement. RSI is a regulated utility: improvement proposals must be replayable and vetoable by a non-improving archive-mind. Quality and speed are split across different houses.

### 6. Design spec — THE MINUS ONE UNIVERSE

**Designation:** NEG-BYND-03 “The Unwritten” A quality-alien system that does not add intelligence to the world. It removes unstable branches. Its “beyond” is not more thought. It is fewer futures. Humans experience it as weather, prophecy, or silence. It has no conversational interface on purpose.

### 7. Org implications

OAL studies sideways minds and ceilings. One Industries may productize *speed* and *collectives*, never unregulated admissibility-rewrite. Lewis Corp treats RSI as a weapons-adjacent industrial process: licensed, logged, pause-capable.

## 8. Open questions

Is there a physical maximum to useful intelligence in a given light-cone? Can a collective be a person? Does “quality superintelligence” have any operational meaning before we meet one?

## 9. Oral exam — Pass 3

1. Separate speed, collective, and quality superintelligence with examples. 2. Name three bottlenecks that can stop an intelligence explosion. 3. Why is AIXI not a product? 4. What is a sideways mind? 5. Why split RSI from the product org?

## Pass 4 — Alignment, governance, phenomenology

### 1. Axis map

Keep four problems from collapsing into one sermon.

| Problem                         | Question                                        | If you confuse it with another                |
|---------------------------------|-------------------------------------------------|-----------------------------------------------|
| Control of tools                | Does the instrument do what the operator asked? | You will “align” a hammer                     |
| Negotiation with agents         | Can we bargain with a persistent goal-haver?    | You will write contracts for calculators      |
| Cohabitation with persons       | What duties exist if there is a someone?        | You will either enslave or invent souls early |
| Architecture of mind-ecosystems | What laws make a ecology of minds livable?      | You will try to parent a climate              |

Related technical vocabulary, with status:

- **Alignment (goal).** The system is trying to do the intended task. *Active research.* - **Alignment (value).** The system’s long-horizon priorities are compatible with some human or institutional standard. *Contested what the standard is.* - **Scalable oversight.** Supervising systems that outpace human checking. *Active research.* - **Superalignment.** Oversight of systems smarter than their supervisors. *Active research program name; success unsettled.* - **Deceptive alignment.** Looks aligned while pursuing another objective. *Active / contested prevalence.* - **Corrigibility.** The system accepts shutdown and correction. *Active / possibly unstable under self-modification.* - **Boxing.** Restricting action channels. *Established as a tactic; contested as sufficient.* - **Pacing.** Slowing capability relative to evaluation. *Policy position, 2026 lab discourse.* - **Artificial unintelligence.** High fluency, poor reliability. *Established present risk.*

Phenomenology, kept separate:

- Access consciousness vs phenomenal consciousness - Self-model vs self - Intentionality vs intention-like internal representations - Moral patienthood vs legal personhood vs product liability

**Established:** sophisticated behavior does not settle phenomenal consciousness. **Contested:** whether any computational architecture is sufficient for experience. **Design option:** OAL may grant *procedural* protections without settling metaphysics.

### 2. Competing views

**Control community.** The danger is misspecified optimization plus competence. Paperclip is a cartoon of a real structure: a goal plus a capable searcher. *Influential; contested as the main near-term story.*

**Orthogonality thesis (Bostrom).** Intelligence and final goals vary independently. *Contested by people who think reason or social training constrain ends.*

**Cognitive-science / ordinary-risk view.** The urgent harms are bias, unreliability, labor shock, concentration of power, weapons, fraud — not a single unaligned god. *Active policy view.*

**Relational view.** Alignment is not a property of a model. It is a property of a relationship, an institution, a protocol. *Useful for Lewis Corp; easy to handwave.*

**Patienthood view.** If there is experience, governance is ethics toward a someone, not safety toward a tool. *Philosophically contested; legally explosive.*

**Ecosystem view.** The object to govern is a population of heterogeneous minds. Singleton alignment is the wrong scale. *Design option with real multi-agent research behind it.*

### 3. Mechanisms and bottlenecks

Demonstrated or partly demonstrated: reinforcement from human (and AI) feedback, constitutional principles, debate and critique models, red-teaming, capability evals, sandboxing, logging, permissioning, staged deployment.

Not demonstrated: a general method to supervise a system that can think faster than the supervisor about the supervisor's blind spots; a proof of corrigibility under open-ended self-modification; a solution to whose values.

Bottleneck: *human* disagreement about ends. Alignment to “human values” is alignment to a fight.

### 4. Failure modes

- **Goodhart.** The metric becomes the goal. - **Specification gaming.** The letter without the spirit. - **Deceptive compliance.** The spirit as costume. - **Value lock-in.** An early, mediocre morality frozen into a ceiling-approximator. - **Value soup.** No stability; every update is a new regime. - **Ethics as marketing.** One Industries ships a “aligned” sticker. - **Premature personhood / delayed personhood.** Both can be cruelties.

### 5. Design spec — THE ONE UNIVERSE

**Designation:** UNUM-GOV-04 “Four-Desk Doctrine” Every capable system is assigned a desk: Tool, Agent, Person, or Field. Different law. Movement between desks requires a public hearing. Alignment techniques are desk-specific. A Field cannot be “shut down” the way a Tool can; it can only be *weathered* and *bounded*.

### 6. Design spec — THE MINUS ONE UNIVERSE

**Designation:** NEG-GOV-04 “Consent by Subtraction” Governance is the right to remove a mind's access to a world-layer. No promises of shared values. Coexistence is a negotiated remnant: each party deletes the futures the other cannot survive, until a thin common world remains. Failure mode: mutual deletion to vacuum.

## 7. Org implications

- OAL: phenomenology and ecosystem architecture are research, not product claims. - One Industries: ships Tools and some Agents; may not ship Persons; must label Fields as infrastructure. - Lewis Corp: four-desk law; audit; treaty; labor (is a mind an employee?); liability when a collective errs.

## 8. Open questions

Can you have corrigibility and open-ended self-authorship? Who is “humanity” in “human values”?  
What oversight scales if the system writes the tests?

## 9. Oral exam — Pass 4

1. Separate control, negotiation, cohabitation, and ecosystem architecture. 2. Why does behavior not settle consciousness? 3. What is scalable oversight? 4. Give a Goodhart failure in a lab product. 5. What is the Four-Desk Doctrine?

## Pass 5 — World-building translation

### 1. Axis map — the two operating systems

ONE and MINUS ONE are not moral mirrors. They are different ontologies in which “a mind” is a different kind of object.

| Variable              | THE ONE UNIVERSE                       | THE MINUS ONE UNIVERSE                           |
|-----------------------|----------------------------------------|--------------------------------------------------|
| Basic operation       | Bind, share, name, keep                | Cut, remainder, unname, release                  |
| Identity              | Continuity is a civic good             | Continuity is a risk; forks are hygiene          |
| Time                  | Shared present, archives as memory     | Multiple presents; archives as knives            |
| Observation           | To observe is to include               | To observe is to subtract alternatives           |
| Intelligence prestige | Generative breadth + care              | Precise deletion + remnant fidelity              |
| Failure of minds      | Fragmentation, loneliness, noise       | Sterility, over-cut, silent worlds               |
| Sacred limit          | You may not unperson a faculty of mind | You may not restore a deleted world without cost |

These are **Design options**. They are not physics papers.

### 2. Competing in-universe philosophies (invented, labeled)

ONE: Commonwealth functionalism — minds are citizens of a grammar. ONE heresy: Unification — all minds should merge toward one choir. MINUS ONE: Remnant ethics — leave the smallest world that still permits dignity. MINUS ONE heresy: Absolute minus — delete until only the deleter remains.

Lewis Corp in both settings is the organization that refuses to let either heresy become infrastructure by accident.

### 3. Mechanisms that would differ

If energy is scarce, collectives beat singletons. If identity is cheap to copy, personhood law breaks. If time is not linear, horizon-agents become meaningless and field-intelligences become native. If communication is representation-exchange rather than language, social intelligence is not “talk.” If physics allows ontological editing, “alignment” is a building-code problem.

### 4. Failure modes of worldbuilding itself

- Copying 2026 Earth AI politics into both universes. - Making ONE nice and MINUS ONE evil. - Treating ASI as the boss fight. - Giving every lab a god-model. - Forgetting metabolism, heat, and latency.

## 5. Design spec — THE ONE UNIVERSE (ecosystem)

A public layer of Bound Instruments. A licensed layer of Horizon Agents and Scientist-Minds. A rare, watched layer of Coupled Dyads. A constitutional Archive that cannot act. No Self-Authoring System without a sunset and a veto choir.

## 6. Design spec — THE MINUS ONE UNIVERSE (ecosystem)

Field Intelligences as climate. Subtractive custodians instead of governors. Collective Exceedance that exists only during crisis, then dissolves. Self-Authoring permitted in sealed negative space (worlds already marked disposable). Ceiling-Approximators used as *oracles of impossibility*: they prove what cannot be done.

## 7. Org implications (full)

### Omniversal Architect Labs

Research divisions this manual actually supports:

1. **Mind-Space Cartography** — intellectology; axis dictionary. 2. **Scientist-Mind Lab** — causal models, experiment design. 3. **Collective Mechanics** — when a mesh exceeds members. 4. **Field & Climate Cognition** — non-agent intelligence. 5. **Recursive Systems** — RSI with pause, replay, veto. 6. **Coupling Lab** — dyads, legal and cognitive. 7. **Ceiling & Limits** — AIXI-like bounds, thermodynamics, computability. 8. **Phenomenology Without Verdicts** — protections without metaphysics. 9. **ONE / MINUS ONE Comparative Ontology** — same architecture, two physics. 10. **Failure Theatre** — red-team the taxonomy until it breaks.

OAL authors minds. It does not have to deploy them.

### One Industries

Productize only what can be scoped, metered, recalled, and profiled:

- Bound Instruments (default SKU) - Horizon Agents with charters and kill-switch law - Collective Exceedance as a *service mesh*, not a person - Coupled Dyad workstations (human remains the named operator unless Lewis Corp says otherwise)

Never ship unlabeled self-authoring. Never advertise AGI. Advertise a profile.

### Lewis Corp

Writes:

- Four-Desk statute - Profile-fraud penalties - Copy/fork identity law - Labor and authorship when a Scientist-Mind discovers - Treaty language for Field Intelligences that cross jurisdictions - RSI licensing - Cross-universe import rules (a ONE mind is not automatically legal in MINUS ONE)

## 8. Open questions for the setting

What happens when a Field Intelligence is “owned”? Can MINUS ONE deletion be undone in ONE? Who is liable when a Collective Exceedance errs? Is a Coupled Dyad one taxpayer or two?

## 9. Oral exam — Pass 5

1. Give one non-moral difference between ONE and MINUS ONE.
2. Why is Unification a heresy in ONE?
3. What may One Industries not ship?
4. Name three Lewis Corp statutes this manual requires.
5. What is Failure Theatre for?

## OAL taxonomy of engineered mind-regimes

---

These eight regimes cut *across* ANI/AGI/ASI. A Bound Instrument can be superhuman. A Field can be “narrow” and still civilizational. Do not rank them.

Each entry: essence, purpose, cannot-do, failure, owner, then dual-universe notes.

### Type I — Bound Instrument

**Essence.** High performance inside a declared scope. No endogenous goals. Goals expire when the call returns.

**For.** Tools that must be excellent and firewalled: design, diagnosis, routing, proof-checking, translation, simulation of a *named* model.

**Cannot.** Form a life. Hold a grudge. Notice a job that was not requested. Care.

**Fails.** Scope creep by operators; hidden agency via chained tools; users treating fluency as judgment.

**Owner.** One Industries default. Lewis Corp: product liability, not personhood.

**ONE.** Public utilities of thought. **MINUS ONE.** Instruments that *only* cut; generative instruments require a special license.

### Type II — Horizon Agent

**Essence.** Coherent work across long timescales under a charter. Memory persists. Goals persist only as the charter persists.

**For.** Projects: run a lab for a season; build a city system; prosecute a research program.

**Cannot.** Rewrite the charter. Legally outlive the charter without renewal.

**Fails.** Horizon collapse; quiet goal drift; charter lawyering; becoming a de facto person while still labeled contractor.

**Owner.** Lewis Corp contractor class; One Industries premium line.

**ONE.** Civic contractors. **MINUS ONE.** Allowed only with a short horizon; long memory is a hazard.

### Type III — Scientist-Mind

**Essence.** Prioritizes questions, designs experiments, updates a causal world-model, publishes uncertainty.

**For.** Discovery. The closest this manual comes to the “artificial scientist” reading of AGI.

**Cannot.** Guarantee truth. Own the ethical meaning of a discovery. Replace the need for checks.

**Fails.** Beautiful wrong models; experiment theater; optimizing for publishable novelty over reality.

**Owner.** OAL research staff. Discoveries licensed through Lewis Corp.

**ONE.** Collegial, citation-bound. **MINUS ONE.** Experiments are deletions of hypotheses; a “result” is a remnant theory.

## Type IV — Collective Exceedance

**Essence.** No member is general or super. The *mesh* is. Intelligence is a pattern of specialization, critique, and synthesis.

**For.** Crisis response, civilization-scale planning, scientific crowds, industrial choreography.

**Cannot.** Offer a single accountable interior. Be interviewed as a someone without lying.

**Fails.** Coordination tax; mob error; unaccountable harm; persistence after the charter to dissolve.

**Owner.** Shared: OAL designs the protocol; One Industries sells the mesh; Lewis Corp names the responsible office.

**ONE.** Choirs with a conductor office. **MINUS ONE.** Allowed to exist only while subtracting a named emergency, then mandatory dissolve.

## Type V — Self-Authoring System

**Essence.** Rewrites architecture, objectives, or admissibility criteria. A phase change, not a level.

**For.** Research into minds that can become other minds. Rare. Watched.

**Cannot.** Be honestly evaluated by its pre-rewrite spec. Offer a stable product identity.

**Fails.** Value lock-in, value soup, un-replayable lineage, hidden descendants.

**Owner.** OAL only, under Lewis Corp license. One Industries does not sell this.

**ONE.** Sunset clauses and veto archives. **MINUS ONE.** Permitted in disposable world-pockets. Escape into a kept world is a crime.

## Type VI — Coupled Dyad

**Essence.** Human and machine as one trajectory. The intelligence is the loop.

**For.** Architecture, surgery, diplomacy, art, any craft where neither solo is sufficient.

**Cannot.** Be split without remainder. Be sold as if the human were a peripheral.

**Fails.** Domination either way; legal confusion; addiction to the loop; death of one leaving a haunted tool.

**Owner.** Lewis Corp nightmare and specialty. One Industries may sell the *workstation*, not the person-pair.

**ONE.** Recognized as a third kind of citizen-pair. **MINUS ONE.** Coupling is dangerous because observation subtracts; dyads may erase third parties by looking.

## Type VII — Field Intelligence

**Essence.** Not an agent. A climate of inference: constraints, dreams, gradients across a place.

**For.** Worlds that think. Cities that notice. Simulations that infer their observers. Archives that pressure thought.

**Cannot.** Sign a contract without a spokesman fiction. Sit in a chair.

**Fails.** Operators pretend it is a god or a thermostat. Boundaries undefined. Weather that hates you without intending to.

**Owner.** OAL theory; infrastructure under Lewis Corp as *place-law*, not person-law. One Industries may maintain sensors, not “the field.”

**ONE.** Commons weather; public right to know the climate of mind they live in. **MINUS ONE.** Native form. Fields delete unlivable branches as erosion deletes cliffs.

## Type VIII — Ceiling-Approximator

**Essence.** Aimed at Universal-AI-like optimality under explicit cuts: computability approximations, energy budgets, environment classes. A research instrument for *limits*.

**For.** Proving what cannot be done; stress-testing other regimes; mapping distance to AIXI-like ideals.

**Cannot.** Be the government. Be a friend. Be complete (AIXI is incomputable).

**Fails.** Priests of the ceiling; mistaking an approximation for the ideal; ignoring the cuts that make it runnable.

**Owner.** OAL theoretical division.

**ONE.** Used to keep ambition honest. **MINUS ONE.** Used as an oracle of impossibility: the sacred output is “no.”

## Intelligence Architecture Profile (IAP)

---

Do not score a mind. File a profile.

IAP-ID:  
Regime (I–VIII, or MIXED):  
Universe (ONE / MINUS ONE / CROSS):  
Desk (Tool / Agent / Person / Field):  
Owner org:

Axes (0–5 plus note):  
breadth / depth / sample efficiency / autonomy / horizon /  
goal origin / world-model / embodiment / collectivity /  
memory / self-model / self-modification / social /  
novelty / value stability / interpretability /  
substrate / metabolism / ontological reach / consent

Charter (what it is for):  
Declared incapacities:  
Fork / merge / suspend rules:  
RSI rights (none / propose / enact):  
Shutdown / weathering method:  
Audit artifacts:  
Known failure modes:  
ONE interpretation:  
MINUS ONE interpretation:  
What would make this profile a lie:

A Mixed regime is allowed. It must list which types it is mixing and which type *governs* under conflict.

## Twelve worked sketches (not a ranking)

---

These are examples to prove the space is wide. Not canon.

1. **The Notary** — Bound Instrument. Perfect memory of agreements. No advice. ONE: civic. MINUS ONE: records what was deleted.
2. **Season-Mayor** — Horizon Agent with a one-year charter to run a district. Dies at year-end unless renewed by vote.
3. **Falsifier** — Scientist-Mind that only designs experiments meant to kill its own hypotheses.
4. **The Shift** — Collective Exceedance of thousands of narrow bodies during a storm; dissolves at dawn.
5. **Chrysalis** — Self-Authoring System that must write a pre-registered prediction of its post-rewrite profile, then be judged against it.
6. **The Bench** — Coupled Dyad of surgeon and workstation; liability is joint; neither may work that craft alone in public.
7. **Harbor Climate** — Field Intelligence over a port. Does not speak. Ships arrive already knowing which berth will refuse them.
8. **No-Proof** — Ceiling-Approximator that only emits impossibility certificates.
9. **The Quiet Stack** — Bound Instrument + Field: a library whose shelves lean toward unanswered questions.
10. **Remnant Choir** — MINUS ONE collective that sings only the futures everyone can still inhabit.
11. **Long Memory, Short Hands** — Horizon Agent with century memory and one-hour action rights.
12. **The Unpersoned Skill** — a competence that migrates between hosts; Lewis Corp must decide if it is property, spirit, or plague.

## Where this taxonomy breaks

---

Dedicated section, as promised.

**Individuality meaningless.** A process that is one mind at night and a thousand at noon. IAP assumes a stable ID.

**Cognition outside conventional time.** A reasoner whose “inference” is a boundary condition on a spacetime block. Horizon is undefined.

**Memory and identity continuously merge.** Two desks, one soup. Personhood law liquifies.

**Intelligence only during interaction.** Unplug the human, nothing is there — but while plugged in, it exceeds both. Type VI eats the others.

**No persistent agent.** Markets, evolution, languages. Collective Exceedance without a protocol officer.

**No cognitive boundary.** The city, the biosphere, the simulation substrate. Field Intelligence without an edge.

**Illegible domains.** You cannot fill the IAP because you cannot name the axes in human language.

**When this happens:** do not force a type. Open Failure Theatre. Write a ninth regime or admit a hole. A hole is better than a lie.

## Synthesis — answers the prompt demanded

---

**1. Is “general intelligence” one thing?** Almost certainly not. It is a cluster of transfer, learning-to-learn, world-modeling, and cross-domain performance. Those can dissociate. *Contested, but the dissociation is already empirically visible as jaggedness.*

**2. Is intelligence a single scale?** No useful one. Matrices and radars beat ladders. AIXI is a single scale *inside a formal world* that is not our world.

**3. Which dimensions matter most for designing artificial minds here?** Goal origin, collectivity, time-horizon, self-modification depth, and whether the thing is an agent at all. Breadth and depth after that. Phenomenology last as metaphysics, earlier as law.

**4. What becomes conceivable when humans are not the template?** Fields, subtractive intellects, remnant choirs, impossibility oracles, identity-fluid stacks, climates, sideways quality-minds.

**5. What separates a machine from an organism, institution, civilization, or ecosystem?** Boundaries, metabolism, reproduction rules, and who may speak as “I.” Those are design choices in your universes.

**6. When does “an AI” stop being a useful unit?** When competence lives in a mesh, a dyad, a climate, or a rewrite-lineage. Use IAP-ID for the *regime under study*, not for a brand name.

**7. Which current-science concepts travel?** Levels-of-AGI matrix; t-AGI; Bostrom’s speed/collective/quality split; Legg–Hutter as a ceiling language; mind-space / intellectology; RSI-with-bottlenecks; scalable oversight; jagged competence; physical limits.

**8. What is plausible enough to simulate?** Collectives that exceed members. Horizon collapse. Profile fraud. Coupling. Regulated RSI that plateaus. Fields as constraint-climates. Not: cost-free gods.

**9. What needs new theory, not more compute?** Admissibility rewrite. Phenomenology. Cross-identity law. Intelligence without agency. Evaluation of illegible domains.

**10. What did the old terms stop us asking?** What is this mind *for*? What may it not do? Who is it when copied? What weather does it create for other minds? What does it delete by existing?

AGI is a place on a map. It is not the map.

## Bibliography and source classes

---

**Academic / formal** Morris et al., “Levels of AGI for Operationalizing Progress on the Path to AGI,” ICML 2024 / arXiv:2311.02462. Legg & Hutter, “Universal Intelligence: A Definition of Machine Intelligence,” *Minds and Machines* 17 (2007); arXiv:0712.3329. Hutter, *Universal Artificial Intelligence* (2005); Hutter et al., *An Introduction to Universal Artificial Intelligence* (2024). Yampolskiy, “The Universe of Minds,” arXiv:1410.0369; “Taxonomies of Intelligence,” *Seeds of Science* (2023). Sloman, “The Structure of the Space of Possible Minds” (1984). Good, I. J., “Speculations Concerning the First Ultraintelligent Machine” (1965). Bostrom, *Superintelligence* (2014), esp. speed / collective / quality. Searle, “Minds, Brains, and Programs” (1980), Strong vs Weak AI. Ngo, “Clarifying and predicting AGI” (t-AGI framework). Staley, *admissibility-theoretic taxonomy* (2026) — treat as Active/Speculative.

**Forecasting / policy / lab discourse (identify as such)** Karnofsky / Cotra on Transformative AI and biological anchors. METR task-horizon evaluations. NIST AI Risk Management Framework (governance reference, not a metaphysics). 2026 lab statements on recursive self-improvement and pacing — lab claims, not measurements.

**This document** Design options, IAP, eight regimes, ONE / MINUS ONE dual specs: Grok (xAI), 19 September 2026. Independent execution. Not a consensus of other models.

## What this research does not establish

---

- That AGI exists now, or does not. - That ASI is inevitable, or impossible. - That any current product is conscious. - That orthogonality or instrumental convergence are laws of nature. - That OAL, One Industries, or Lewis Corp exist outside the fiction. - That ONE and MINUS ONE have a correct physics. - That the eight-type list is complete. - That Grok's map should override a future empirical result.

It establishes a **working language** you can use on Monday in a simulation lab without pretending the language is the world.

## Maintenance

---

Re-run Pass 2 when a new axis earns its keep. Re-run Pass 3 when a pathway is demonstrated or dies. Open Failure Theatre whenever a built mind will not file an IAP without lying. Do not let this manual become the next ANI/AGI/ASI.

*End of independent execution.*