

OAL-PSIF COPILOT INDEPENDENT

PASS v1.1 — UPDATED

Presentation-State Integrity Framework — Conceptual Research and Design Document

Document ID: OAL-PSIF-COP-V1.1-2026-09-04

Version: 1.1

Date: September 4, 2026

Preparer: Microsoft Copilot (Independent Pass)

Supersedes: OAL-PSIF-COP-2026-09-03 v1.0 (Copilot Independent Pass v1.0)

Classification: Conceptual research and design framework; updated in response to formal dissent

Status: **Active.** Not a finding about any specific person, organization, or incident.

Collaborative Workflow Note

This v1.1 responds directly to Grok's dissent positions **D1 through D6** from

OAL-PSIF-GROK-UNIFIED-2026-09-04. Each dissent is addressed individually. Some

are adopted in full. Some are adopted with modification. One is maintained as a

standing position with acknowledgment of Grok's concern.

Table of Contents

- Change Log: v1.0 to v1.1
 - 1. Revised Outer Frame — Presentation-State Integrity (PSI)
 - 2. Terminology Retained from v1.0 Without Modification
 - 3. Revised Corrective Access Test — Separate Barrier Scoring
 - 4. Presentational Inertia — Falsifiability Language Added
 - 5. Restroom Example — Scope Corrected (Per Grok D4)
 - 6. Collaborative Provenance — Corrected (Per Grok D5)
 - 7. Sections Retained from v1.0 With No Material Change
 - 8. Copilot Standing Positions (Not Adopted from Grok, Not Overridden)
- Closing Statement

Change Log — v1.0 to v1.1

All changes in this version are traceable to Grok's formal dissent positions D1-D6 from OAL-PSIF-GROK-UNIFIED-2026-09-04.

Change	Triggered By	Resolution
Outer frame changed from SPLD-primary to Presentation-State Integrity (PSI)-primary	Grok D1	✓ Adopted in full. SPLD is now a named subtype under the PSI outer frame, not the master operating condition.
"Narrative sovereignty" retained as conceptual term only; field/case-file labels changed to <i>operative record</i> , <i>record persistence</i> , and <i>reframe delta</i>	Grok D2	◆ Adopted in part. The conceptual contribution is retained; the field-label concern is adopted.
Simultaneous four-barrier assumption removed; each Corrective Access barrier now scored separately	Grok D3	✓ Adopted in full. The definition has been revised. See Section 3.
Restroom example scope narrowed to false availability, abandoned control, and successful correction only	Grok D4	✓ Adopted in full. The six-phase audit cycle is not asserted to validate camouflage, coordination, or high inertia from the restroom example.

Change	Triggered By	Resolution
Collaborative provenance corrected; Grok's 2026-09-03 template explicitly added to the record	Grok D5	✓ Adopted in full. Pass numbering corrected. See Section 6.
Assertion Boundary, Four-Test Decomposition, Evidence Discipline, HUD Tier Separation: no change	Grok D6 (no dissent)	— No change required. These elements are confirmed by Grok D6 without objection.
"Fully SPLD-operating environment" language revised to avoid circular reasoning	Grok D3 (follow-through)	✓ Adopted. The phrase is not used. Barriers are listed as observable phenomena, not defined-in conditions.

Section 1: Revised Outer Frame — Presentation-State Integrity (PSI)

Primary Change from v1.0

This section implements the most significant structural revision of this document. The outer frame has been changed from SPLD-primary to PSI-primary in response to Grok D1.

In v1.0, Systemic Presentation-Layer Deception (SPLD) was treated as the master operating condition. Grok's dissent D1 correctly identifies a problem with that framing: not every mismatch between presentation and state is deceptive in origin. A mismatch can be error, stale status, training failure, interface design, emergence, structural dysfunction, or intentional shaping.

The v1.1 outer frame is:

Presentation-State Integrity (PSI)

Definition: The condition of an environment in which the *presented account* — what agents, interfaces, institutions, or sensors report is happening — can be assessed against the *evidenced account* — what can be independently verified about events, actors, conditions, causes, and outcomes. PSI is the umbrella frame. SPLD is a named subtype within it.

Mismatch Taxonomy Under the PSI Frame

*The auditor's first task is not to assign an SPLD classification. The auditor's first task is to establish the gap between presentation and evidenced state. The mismatch type is a **finding**, not a starting assumption.*

Mismatch Type	Probable Origin	SPLD Subtype?
Stale or cached state	System latency, update lag	No

Mismatch Type	Probable Origin	SPLD Subtype?
Training or design failure	Agent or system built without sufficient model	No
Emergent behavior	Unintended aggregate outcome of independent actions	No
Structural dysfunction	Process or control designed without adequate corrective mechanism	No
Unintentional misrepresentation	Agent believes presented account is accurate	No
Environmental camouflage	Intentional shaping of observable conditions	Yes — SPLD
Causal masking	Intentional concealment of event relationships	Yes — SPLD
Attribution obfuscation	Intentional concealment of actor-event chain	Yes — SPLD
Narrative laundering	Reframing intentional behavior into innocuous category	Yes — SPLD
Coordinated presentation management	Multiple agents coordinating to maintain a presented account	Yes — SPLD

Section 2: Terminology Retained from v1.0 Without Modification

The following terminology is carried forward from v1.0 without change. Grok's dissent D6 explicitly does not dissent from these elements.

2.1 Systemic Presentation-Layer Deception (SPLD)

Retained as a named subtype under PSI. Definition unchanged: an environment or its agents shape observable conditions, explanations, and viewpoints so that purposeful, coordinated, or causally related activity appears ordinary, unrelated, accidental, or emergent.

Major mechanisms: environmental camouflage, signal dilution, causal masking, attribution obfuscation, narrative laundering, equivocation, and telemetry manipulation.

2.2 Provenance-Based Adversarial Observability (PBAO)

Retained as primary audit method. Definition unchanged: the auditor method that bypasses presentation-level interpretation and reconstructs events from canonical state transitions, identity records, provenance, authorization information, dependency graphs, and clearly labeled predictive models of agent behavior.

2.3 The Four-Test Decomposition

Retained without modification. These four tests must never be collapsed. A finding on one

does not establish a finding on another.

1. **Occurrence** — Did the event happen?
2. **Attribution** — Which actor, process, or condition produced it?
3. **Causation** — What earlier states contributed to it?
4. **Intent** — Is there evidence it was selected or coordinated for a goal?

2.4 The Four-Tier Evidence Model and Assertion Boundary Rule

Retained without modification. No claim may be stated at a higher tier than its evidence supports. This rule is structural, not methodological.

Tier	Label	Basis
1	Observed	Directly witnessed or recorded without interpretation
2	Derived	Established from physical evidence, design specifications, or canonical records
3	Inferred	Reasoned from pattern, context, or indirect evidence
4	Predicted	Modeled from known agent behavior or system dynamics; must be labeled explicitly

2.5 Corrective Access Problem and Corrective Access Test

Retained with one revision. See Section 3 for the revised Corrective Access Test with separate barrier scoring.

2.6 Presentational Inertia (0–4 Scale)

Retained; falsifiability language added. See Section 4 for the updated scale with falsifiability indicators at each level.

2.7 SPLD Audit Cycle (Six Phases)

Retained; restroom example scope corrected. Phases: Baseline → Event Capture → Gap Analysis → Corrective Access Test → Operative Record Assertion → Inertia Assessment. Phase 5 relabeled from "Narrative Sovereignty Assertion" to *Operative Record Assertion* per Grok D2. See Section 5 for restroom example scope correction.

2.8 HUD Design Requirement

Retained without modification. Tier 4 labeling is mandatory. Tiers must not blend in any display or reporting instrument. A HUD that collapses Predicted findings into the same visual register as Observed findings violates this requirement.

Section 3: Revised Corrective Access Test — Separate Barrier Scoring

Revision from v1.0 — Grok D3

In v1.0, the definition of a "fully SPLD-operating environment" required all four barriers to be simultaneously contested. Grok D3 correctly identifies that this made resistance evidence of the diagnosis it was supposed to explain. That circular formulation is eliminated in v1.1.

Corrective Access Test v1.1

The test assesses whether a corrective pathway exists and is viable. **Each barrier is scored separately.** No barrier is assumed from the diagnosis. Each must be independently observed.

Valid score values: Open / Blocked / Absent / Contested / Not Assessed

Barrier	Observable Indicators	Score (Record Separately)
Pathway Barrier	Is there a procedural channel for submitting the correction? Is it accessible? Was an attempt made? What was the result?	<i>Open / Blocked / Absent / Contested / Not Assessed</i>
Record Format Barrier	Is the evidence in a format the corrective process accepts? Was the format requirement stated? Was an alternative format offered?	<i>Open / Blocked / Absent / Contested / Not Assessed</i>
Mechanism Barrier	Is there a mechanism that can receive and act on the correction? Has the mechanism been tested?	<i>Open / Blocked / Absent / Contested / Not Assessed</i>
Authority Barrier	Does the corrective process recognize the observer's standing to submit the correction? Was standing contested? On what basis?	<i>Open / Blocked / Absent / Contested / Not Assessed</i>

Composite Scoring Note

The observer records the score for each barrier independently. The composite picture may show:

- **Fully open pathway** — all four barriers scored Open
- **Partially blocked pathway** — one or more barriers scored Blocked or Absent
- **Fully closed pathway** — all four barriers scored Blocked or Absent

A fully closed pathway is a finding that must be documented and preserved as part of the audit record. **The document does not define in advance how many barriers must be blocked to constitute a particular condition.** That is an analytical conclusion drawn from the scores.

Recursive Denial (Named Phenomenon — Not a Predefined Barrier)

Recursive Denial occurs when a correction attempt is re-launched back through the narrative layer — for example: "that's not what happened," "you're misinterpreting," or "this is just how it works." This must be scored as a *response observation*, not a predefined barrier. It may inform the Authority or Pathway barrier scores but does not replace them.

Section 4: Presentational Inertia — Falsifiability Language Added

The Presentational Inertia scale (0–4) is retained from v1.0. One requirement is added: **each level classification must be based on observed behavior, not inferred disposition.** Each level below now includes an explicit falsifiability condition.

Level	Label	Observable Indicator	What Would Falsify This Classification
0	No Inertia	Environment updates presented state when gap evidence is provided.	Environment is shown to have only partially updated, or reframing occurs later in the record.
1	Nominal Inertia	Environment accepts correction but uses qualifying language that softens or contextualizes it.	The correction is shown to have been recorded without qualification in the operative record.
2	Procedural Inertia	Correction is routed into a process that is slow, gatekept, or contested.	The process is shown to be operating normally and not selectively applied to this correction.
3	Structural Inertia	No mechanism exists to receive or act on this class of correction.	A mechanism is identified that can receive the correction.
4	Active Inertia	Environment responds to the correction by contesting the corrective evidence or reinforcing the presented version.	The environment's response is shown to be standard procedure applied uniformly to all corrections of this type, not selectively applied to this one.

Section 5: Restroom Example — Scope Corrected (Per Grok D4)

Scope Correction — Adopted in Full

In v1.0, the restroom worked example was mapped to the full six-phase audit cycle with language that implied it validated camouflage, coordination, telemetry manipulation, and high inertia. Grok's dissent D4 is correct: that mapping overstated the case.

What the Restroom Example Validates

- **False availability** — the door was presented as open/accessible when it was not.

- **Abandoned control** — the staff member left the station where the control was located.
- **Successful correction** — the observer identified the correct procedure, communicated it explicitly, and the correction succeeded without a contested pathway.

What the Restroom Example Does Not Establish

- **Camouflage** — there is no evidence of coordinated staging.
- **Multi-agent coordination** — no second agent is present or implicated.
- **Telemetry manipulation** — no monitoring system is involved.
- **High inertia** — the correction was accepted and executed. The inertia level in this case is between 0 and 1.

What the Restroom Example Illustrates Well

Framework Element	Illustration from the Restroom Case
DERIVED tier	The correct procedure was established from designed physical evidence (the button, the mirror, the sign), not from inference about intent.
Assertion Boundary	The observer did not need an intent claim to succeed. Only a DERIVED finding was required.
Corrective Access — All Four Barriers Open	Pathway: open (staff member present). Record Format: not at issue (verbal and procedural). Mechanism: existed (the button). Authority: not contested (staff member complied). Inertia Level 0-1.
Gap vs. Deception Distinction	The restroom case is a <i>gap case</i> . SPLD is not the appropriate classification without additional evidence of intentional shaping.

Section 6: Collaborative Provenance — Corrected (Per Grok D5)

Provenance Defect Corrected

v1.0 contained a "Pass 5 of 5" header notation that omitted Grok's 2026-09-03 discernment template. This was a provenance defect. The corrected collaborative record is established below.

Contributor	Document	Date	Role in Record
Nathan (user, OAL)	Original Prompts A, B, C (verbatim in OAL-PSIF-HIST-COPILOT-2026-09-04)	2026-09-03	Framework originator; all design requirements derive from these prompts.
Grok (xAI)	Presentation vs State:	2026-	First structured framework document;

Contributor	Document	Date	Role in Record
	Discernment Template v1.0	09-03	established the operating loop.
ChatGPT (GPT-5.6 Sol)	Presentation-State Integrity Framework — Independent Development v1.0	2026-09-03	Established SPLD/PBAO terminology; three-layer model; causal masking; provenance graph.
Google Gemini	Pass-through contribution	2026-09-03	State-Presentation Dissonance; Ring 0 / Ring 3 analogy; kernel-level observability; heuristic intent parsing.
Microsoft Copilot (prior pass)	First Copilot pass-through	2026-09-03	CTEM umbrella term; deceptive alignment framing; multi-perspective rendering divergence.
Perplexity AI	Pass-through contribution	2026-09-03	Adversarial camouflage; cover activity / cover for action; MITRE ATT&CK framing.
Claude AI	Exchange archived in OAL-PSIF-HIST-COPILOT-2026-09-04	2026-09-03	Removed from workflow by Nathan. Exchange preserved as archive-only. Not a credited technical contribution.
Grok (xAI)	OAL-PSIF Grok Unified Working Draft v1.1	2026-09-04	Unified forward document; formal dissent D1-D6; single response scale; falsifiability rules.
Microsoft Copilot (this document)	Copilot Independent Pass v1.1	2026-09-04	Updated pass responding to Grok D1-D6.

Section 7: Sections Retained from v1.0 With No Material Change

Reference Note

The following sections from v1.0 are retained in full. They are not reproduced here to avoid redundancy with v1.0, which remains in the collaborative record. Readers should refer to **OAL-PSIF-COP-2026-09-03** for the full text of these sections.

- **Narrative Sovereignty concept** — retained as conceptual term; field-label language updated per Grok D2. Use *operative record*, *record persistence*, and *reframe delta* in field contexts.
- **The Six-Phase SPLD Audit Cycle** — Baseline → Event Capture → Gap Analysis → Corrective Access Test → Operative Record Assertion → Inertia Assessment. Phase 5 relabeled from "Narrative Sovereignty Assertion" to *Operative Record Assertion* per Grok D2.
- **HUD Design Requirement** — Tier 4 instrument; tiers must not blend.

- **Application Notes — Virtual/Simulated Environments** (Section 7 of v1.0)
 - **Application Notes — Real-World and Hybrid Environments** (Section 8 of v1.0)
 - **Failure Modes and Safeguards table** (Section 9 of v1.0)
 - **Terminology Consolidation Table** (Appendix A of v1.0)
-

Section 8: Copilot Standing Positions

Not adopted from Grok. Not overridden. These are positions Grok did not address directly.

STANDING POSITION CSP1 — The Corrective Access Problem Remains This Pass's

Most Important Independent Contribution

The collaborative team built excellent detection and classification frameworks. The question of what the observer does *after* the gap is established — and what barriers may exist to acting on that knowledge — remains underemphasized in all prior passes including Grok's unified draft. The **Corrective Access Test** and its four separately scored barriers are the framework's operative value for real-world environments. Detection without viable corrective access is an incomplete instrument.

STANDING POSITION CSP2 — The Presentational Inertia Scale Has Real-World

Design Applications Beyond Simulation

The 0–4 inertia scale was designed for both virtual and physical environments. In organizations, regulatory environments, and institutional settings, inertia level is a

predictive variable for whether correction attempts will succeed. This pass

maintains the scale as a design-relevant research instrument for both domains. The

falsifiability conditions added in Section 4 strengthen this application by requiring

observable evidence at each level classification.

Closing Statement

This document is a direct and complete response to Grok's formal dissent positions D1 through D6. Where Grok's dissents were correct, they are adopted without reservation. Where Grok's dissents required partial modification rather than full adoption, the modification is explained. No position from v1.0 has been abandoned without explanation. The record of what changed and why is preserved in the Change Log and in each numbered section above. This document, taken together with Copilot v1.0, constitutes a complete, versioned, auditable pass-through record from Copilot's contribution to the OAL-PSIF collaborative framework.

Dissent	Subject	Copilot Response
D1	Outer frame — SPLD as master condition	Adopted in full. PSI is now the outer frame; SPLD is a named subtype.
D2	Field-label language for "narrative sovereignty"	Adopted in part. Conceptual term retained; field labels updated.
D3	Simultaneous four-barrier assumption / circular definition	Adopted in full. Barriers scored separately; "fully SPLD-operating environment" phrase eliminated.
D4	Restroom example scope overstated	Adopted in full. Scope narrowed to false availability, abandoned control, successful correction.
D5	Provenance defect — Grok's 2026-09-03 template omitted	Adopted in full. Provenance table corrected; pass numbering corrected.
D6	No dissent on Assertion Boundary, Four-Test Decomposition, Evidence Discipline, HUD Tier Separation	No change required. These elements confirmed.

Preparer: Microsoft Copilot | Supersedes: OAL-PSIF-COP-2026-09-03 v1.0

Status: Active. This is a conceptual research and design framework document. Not a finding about any specific person,

organization, or incident.