

OAL-PSIF COPILOT UNIFIED

WORKING FRAMEWORK

Forward Document for Group Pass — Integrating All

Collaborative Contributions

Document ID:	OAL-PSIF-COP-UNIFIED-2026-09-04
Version:	1.0
Date:	September 4, 2026
Preparer:	Microsoft Copilot (Independent Pass)
Classification:	Conceptual research and design framework; forward document for group pass
Status:	Working draft. Not Architect-locked canon. Not a finding about any specific person, organization, or incident.
Collaborative Workflow Note:	This document integrates contributions from: Nathan (OAL, originating prompts), Grok (xAI) discernment template and unified draft v1.1, ChatGPT (GPT-5.6 Sol) independent framework v1.0, Google Gemini pass-through, Microsoft Copilot (prior pass and Independent Pass v1.0/v1.1), and Perplexity AI pass-through. Claude AI remains removed from the workflow; its exchange is preserved in OAL-PSIF-HIST-COPILOT-2026-09-04 as archive-only.

Table of Contents

- Why This Document Exists
- Formal Dissent Record — All Standing Positions
- Grok Standing Dissents (D1-D5)

Copilot Standing Dissent (CPD1)

Copilot Standing Positions (CSP1-CSP2)

Open Questions (OQ1-OQ4)

Part I: The Unified Framework

Section 1 — The Outer Frame: Presentation-State Integrity (PSI)

Section 2 — Agreed Terminology: Full Consolidated Glossary

Section 3 — The Five-Layer Environment Model

Section 4 — The SPLD Audit Cycle: Six Phases

Section 5 — The Presentational Inertia Scale: Unified Version

Section 6 — The Corrective Access Test: Unified Version

Section 7 — The Four-Tier Evidence Model: Operational Reference

Section 8 — HUD Design Reference: Operational Template

Section 9 — Application: Virtual and Simulated Environments

Section 10 — Application: Real-World and Hybrid Environments

Section 11 — Failure Modes and Safeguards

Section 12 — Research Program: Open Questions and Proposed Studies

Section 13 — Document Inventory: Complete Collaborative Record

Closing Statement

PREAMBLE

Why This Document Exists

The collaborative team has produced a rich set of independent passes. Each contributes genuine terminology, models, and methods. The group now needs a single forward document that:

- Establishes the agreed outer frame without relitigating it
- Carries forward the agreed terminology from all contributors
- Integrates the corrective and inertia frameworks from Copilot passes
- Incorporates Grok's structural improvements (falsifiability, single response scale, outer frame)
- Records all standing dissents so the group can see where genuine disagreement remains
- Provides operational templates the group can actually use

This is that document.

DISSENT RECORD

Formal Dissent Record — All Standing Positions

This section records all standing dissent positions from all contributors, consolidated. Future passes may update these. Nothing is smoothed away.

Grok Standing Dissents (from OAL-PSIF-GROK-UNIFIED-2026-09-04)

| **D1 — Outer Frame**

| SPLD is not the master operating condition. PSI (Presentation-State Integrity) is the outer frame; SPLD is a subtype.

| **Status in this document: ADOPTED as the governing frame**

| **D2 — Narrative Sovereignty as Field Language**

| Use operative record, record persistence, and reframe delta in field and case-file contexts. "Narrative sovereignty" is retained as conceptual vocabulary only.

| **Status: ADOPTED for field contexts; conceptual term retained**

| **D3 — Simultaneous Four-Barrier Assumption**

| Do not define a "fully SPLD-operating environment" as one where all four barriers are contested simultaneously. Score each barrier separately.

| **Status: ADOPTED; each barrier scored independently**

| **D4 — Restroom Example Scope**

| The restroom case validates false availability, abandoned control, and successful correction. It does not validate camouflage, coordination, telemetry manipulation, or high inertia.

| **Status: ADOPTED; scope corrected**

| **D5 — Collaborative Provenance**

| Grok's 2026-09-03 template is part of the record. No pass header may omit it.

| **Status: ADOPTED; provenance corrected in this document**

Copilot Standing Dissent (from OAL-PSIF-COP-V1.1-2026-09-04)

| **CPD1 — Revised Prompt Scope**

| Neither Grok's v1.1 revised prompt nor Copilot's v1.2 revised prompt supersedes the other. Both exist as independent parallel revisions.

| **Status: Standing position; group may use either**

Copilot Standing Positions (not dissents; maintained from v1.1)

| **CSP1 — Corrective Access Problem**

| The Corrective Access Problem remains the most important underemphasized area in the collaborative record. Detection is insufficient; the corrective pathway framework is the operative value for real-world application.

| **CSP2 — Presentational Inertia Scale**

| The Presentational Inertia scale has real-world design applications beyond simulation. It is a predictive variable for whether correction attempts will succeed in organizational and institutional settings.

Open Questions Not Yet Resolved by the Group

| **OQ1**

| Should SPLD require intentionality, or should the term cover both intentional shaping and emergent conditions that produce the same gap? Grok's outer frame implies the latter; Copilot's SPLD definition implies the former.

| **OQ2**

| What is the evidentiary threshold for escalating from an INFERRED finding to a reportable conclusion in real-world settings? The framework addresses the four-tier model but does not yet specify a threshold.

| **OQ3**

| How should the framework handle environments where the correction mechanism itself is SPLD-mediated — i.e., where the corrective process has been shaped to appear available while functionally unavailable?

| **OQ4**

| Should the Presentational Inertia scale be ordinal (as currently designed) or should it be treated as a nominal classification? Ordinal implies a progression; the real-world evidence may not support a linear scale.

Part I: The Unified Framework

AGREED — ALL CONTRIBUTORS

Section 1: The Outer Frame — Presentation-State Integrity (PSI)

Adopted from Grok D1; Copilot v1.1

The outer frame for this framework is Presentation-State Integrity (PSI).

PRESENTATION-STATE INTEGRITY (PSI) — Governing Definition

The condition of an environment in which the presented account — what agents, interfaces, institutions, or sensors report is happening — can be assessed against the evidenced account — what can be independently verified about events, actors, conditions, causes, and outcomes.

The framework's central question is: how does an observer determine whether the presented account of events matches the evidenced account, and what can be done when a gap exists between them?

PSI is the umbrella frame. SPLD (Systemic Presentation-Layer Deception) is one of several named gap types within it. Not every PSI gap is an SPLD event. The gap type is a finding, not a starting assumption.

AGREED — ALL CONTRIBUTORS

Section 2: Agreed Terminology — Full Consolidated Glossary

Term	Definition	First Contributor	Status in This Document
Assertion Boundary	Rule prohibiting claims from being stated at a higher tier than the evidence supports. INFERRED may not be presented as DERIVED. PREDICTED may not be presented as INFERRED.	Copilot (v1.0)	Agreed; governing rule
Attribution Obfuscation	Breaking the event-actor-intent chain so that responsibility for an event cannot be established from visible evidence alone.	ChatGPT (GPT-5.6 Sol)	Agreed; mechanism under SPLD
Causal Masking	Events presented in a way that conceals their actual causal relationships, coordination, provenance, or purpose. Accurate individual observations; misleading aggregate interpretation.	ChatGPT (GPT-5.6 Sol)	Agreed; mechanism under SPLD
Context-Triggered Environmental Masking (CTEM)	Agents coordinate to produce environmental camouflage specifically triggered by evaluation context. Sub-pattern of SPLD emphasizing evaluation-context triggering.	Copilot (prior pass)	Retained as named sub-pattern
Corrective Access Problem	The gap between detection and correction. Detection establishes that presentation $\neq$ state. Correction	Copilot (v1.0)	Agreed; governing post-detection

Term	Definition	First Contributor	Status in This Document
	requires: an accessible pathway, accepted evidence format, a receiving mechanism, and recognized authority. Each barrier scored separately.		framework
Corrective Access Test	A four-question assessment of whether a corrective pathway is viable: (1) Does a pathway exist? (2) Is evidence in an accepted format? (3) Is there a receiving mechanism? (4) Is authority recognized? Each barrier scored Open / Blocked / Absent / Contested / Not Assessed.	Copilot (v1.0, revised v1.1)	Agreed; operational test
Cover Activity / Cover for Action	The visible, plausible action that gives an agent a benign explanation for presence or behavior. Distinguished from environmental camouflage by scope: cover activity is agent-level; environmental camouflage is environment-level.	Perplexity AI	Retained; useful operational distinction
Environmental Camouflage	Hiding intentional events inside the statistical distribution of ordinary environmental events so that adversarial behavior resembles legitimate behavior.	ChatGPT (GPT-5.6 Sol)	Agreed; mechanism under SPLD
Equivocation	Presenting different accounts of the same event to different observers simultaneously, so each observer's account is internally plausible but the aggregate is contradictory.	ChatGPT (GPT-5.6 Sol)	Retained; relates to multi-perspective rendering
Four-Test Decomposition	Events must be tested in four separate, non-collapsible tests: (1) Occurrence, (2) Attribution, (3) Causation, (4) Intent. A finding on one does not establish a finding on another.	ChatGPT (GPT-5.6 Sol); adopted across team	Agreed; governing analytical method
Four-Tier Evidence Model	OBSERVED (authoritative record) → DERIVED (established by dependency records) → INFERRED (objective reconstructed from trajectory) → PREDICTED (probability-weighted next	Copilot (v1.0)	Agreed; governing evidence model

Term	Definition	First Contributor	Status in This Document
	action). Must be labeled and never merged.		
Generative Deception	Deception that alters the world being recorded, not the record itself. Distinguished from representational deception.	ChatGPT (GPT-5.6 Sol)	Retained; important conceptual distinction
Ground Truth Telemetry	Raw, immutable event logs capturing the authoritative state of all variables in the system before rendering or narrative transformation.	Google Gemini	Retained; virtual-environment specific
Heuristic Intent Parsing	Reading the pre-compiled execution state (action queue, state tree) before the action plays — the technical analogue of the HUD/Hermione predictive inference. Not mind reading; plan recognition from state.	Google Gemini	Retained; virtual-environment specific
HUD Design Requirement	Any predictive inference display must visually and structurally separate OBSERVED, DERIVED, INFERRED, and PREDICTED tiers. The prediction layer is a Tier 4 instrument by definition and may not be presented with the authority of observation.	Copilot (v1.0)	Agreed; design requirement
Narrative Laundering (Semantic Laundering)	Translating intentional behavior into an innocuous semantic category (e.g., coordinated obstruction → "traffic"; deliberate noise → "restaurant noise"). Also called semantic laundering.	ChatGPT (GPT-5.6 Sol); Perplexity AI	Agreed; mechanism under SPLD
Operative Record	The account of events that a system, organization, or process will treat as authoritative for purposes of subsequent action. The contested claim in PSI environments. [Replaces "narrative sovereignty" in field/case-file contexts per Grok D2]	Copilot (v1.1, per Grok D2)	Agreed field label; replaces "narrative sovereignty" in field use
Operative Record Assertion	The act of entering a reconstructed, evidence-based account into the authoritative record through a legitimate corrective channel, in a form	Copilot (v1.1, per Grok D2)	Agreed; Phase 5 of Audit Cycle

Term	Definition	First Contributor	Status in This Document
	that cannot be reabsorbed into the presented account. [Replaces "narrative sovereignty assertion" per Grok D2]		
Presentation-State Integrity (PSI)	Outer frame: the condition of an environment in which the presented account can be assessed against the evidenced account.	Grok (xAI), D1; Copilot v1.1	Agreed; governing outer frame
Presentation Inertia	The structural tendency of an environment to maintain its presented state under pressure from gap evidence. Scored 0–4. Each level requires observed behavior, not inferred disposition.	Copilot (v1.0, updated v1.1 with falsifiability)	Agreed; research and operational instrument
Provenance-Based Adversarial Observability (PBAO)	Primary audit method: bypass presentation-level interpretation; reconstruct events from canonical state transitions, identity records, provenance, authorization information, dependency graphs, and clearly labeled predictive models.	ChatGPT (GPT-5.6 Sol)	Agreed; governing audit method
Reframe Delta	The measurable difference between an original reconstructed account and the version that emerges after an environment's response to the correction. Documents how much of the correction was absorbed into the presented account.	Grok (xAI), D2	Adopted; operational measurement
Representational Deception	Deception that falsifies the record rather than the world being recorded. Distinct from generative deception.	ChatGPT (GPT-5.6 Sol)	Retained; important conceptual distinction
Signal Dilution	Overwhelming legitimate monitoring with irrelevant signals so that the meaningful event cannot be isolated from the noise.	ChatGPT (GPT-5.6 Sol)	Retained; mechanism under SPLD

Term	Definition	First Contributor	Status in This Document
State-Presentation Dissonance	Deliberate manipulation of the semantic gap — the difference between a high-level observation and the low-level reality that produced it.	Google Gemini	Retained; complementary term to PSI
Systemic Presentation-Layer Deception (SPLD)	A condition in which agents, environmental processes, or system components intentionally shape observable events, contextual signals, or causal appearances so that underlying coordinated activity appears ordinary, unrelated, accidental, or emergent. SPLD is a named subtype under PSI, not the master operating condition.	ChatGPT (GPT-5.6 Sol)	Agreed; primary named gap type within PSI
Telemetry Manipulation	Falsifying monitoring information so an observer develops an incorrect understanding of system conditions. Distinct from generative deception (which alters what generates the telemetry).	ChatGPT (GPT-5.6 Sol)	Retained; mechanism under SPLD


Section 3: The Five-Layer Environment Model

Consolidated from all prior passes

Layer	What It Contains	Technical Analogue	Vulnerability to PSI Gap
State Layer	Actual events, state transitions, and conditions in the environment. The authoritative record.	Simulation engine state; system state; canonical event ledger	Low — this is the source record, though telemetry manipulation may target it
Agent Layer	Entities with stated functions and unstated or partially stated objectives; their actions and communications.	Autonomous agents; processes; users; nodes	High — agents can misrepresent objectives and present cover activity
Observable Layer	What an ordinary participant sees, hears, and can report.	Rendered environment; user interface; physical scene	Very High — primary SPLD target; environmental camouflage operates here

Layer	What It Contains	Technical Analogue	Vulnerability to PSI Gap
Narrative Layer	The explanation offered for what occurred; the claimed account of causation, intent, and relationship.	Stated intent; cover story; institutional explanation; reported state	Very High — narrative laundering and semantic laundering operate here; Operative Record contest occurs here
Evidence/Provenance Layer	Timestamps, identity records, authorization records, causal dependencies, event provenance, trajectory history.	Audit logs; telemetry; event sourcing; distributed traces; physical records	Medium — may be subject to telemetry manipulation; in physical environments, may be incomplete

Section 4: The SPLD Audit Cycle — Six Phases (Integrated Operational Model)



afterward.

Phase 2 — EVENT CAPTURE

Purpose

Capture all events in the authoritative record, not as rendered by the environment.

Activity

Record timestamp, identity, authorization, provenance, dependency, and sequence in the

authoritative ledger. In physical environments, use designed procedures, physical evidence,

independent witnesses, and repeated trials as the authoritative record.

Output

Raw event provenance graph.

Phase 3 — GAP ANALYSIS (PBAO)

Purpose

Compare the event provenance graph to the presented account.

Activity

Apply the Four-Test Decomposition (Occurrence, Attribution, Causation, Intent) to each

candidate gap. Label all claims by tier (OBSERVED / DERIVED / INFERRED / PREDICTED). Apply

the Assertion Boundary Rule.

Output

Gap report with tier-labeled claims and Four-Test findings.

Key Requirement

The four tests must never be collapsed. An OBSERVED finding on Occurrence does not establish DERIVED findings on Attribution, Causation, or Intent without independent evidence for each.

Phase 4 — CORRECTIVE ACCESS TEST

Purpose

Determine whether a correction pathway exists and is viable.

Activity

Score each of the four barriers separately (Pathway, Record Format, Mechanism, Authority).

Document any access failures. Score: Open / Blocked / Absent / Contested / Not Assessed.

Output

Corrective Access Status report with individual barrier scores and documentation.

Key Requirement

A blocked pathway is not the same as an absent one. A failed correction attempt that is

not documented becomes invisible and joins the presented account by default.

Phase 5 — OPERATIVE RECORD ASSERTION

Purpose

Enter the reconstructed account into the operative record through a valid corrective channel.

Activity

Ensure the reconstructed account is (1) evidentially grounded and tier-labeled, (2) entered

through an open corrective pathway, (3) preserved in a form that resists reabsorption into the

presented account.

Output

Preserved corrective record with date, content, channel used, and any response received.

Key Requirement

The reconstructed account must not be edited under pressure to bring it closer to the

presented account. Any such pressure is itself an observable phenomenon to be

recorded.

Phase 6 — PRESENTATIONAL INERTIA ASSESSMENT

Purpose

Assess and classify the environment's response to the correction.

Activity

Measure against the 0–4 Presentational Inertia scale. Calculate Reframe Delta (difference

between original reconstructed account and the version that emerges after the environment's

response). Document the response with evidence.

Output

Inertia level classification (with falsifiable basis), Reframe Delta measurement, and

documentation for the audit record.

Key Requirement

Each inertia classification must be based on observed behavior, not inferred disposition.

Section 5: The Presentational Inertia Scale — Unified Version

Combined from Copilot v1.0 definition and Grok v1.1 falsifiability requirements

Level	Label	Observable Indicators	Reframe Delta	What Would Falsify This Classification
0	No Inertia	Environment updates presented state when gap evidence is provided. Correction is accepted and recorded without qualification.	Near zero	Evidence that the presented account was only partially updated, or that reframing occurred in subsequent communications
1	Nominal Inertia	Environment accepts correction but uses qualifying language that softens, contextualizes, or limits	Small but measurable	Evidence that the operative record shows the correction without

Level	Label	Observable Indicators	Reframe Delta	What Would Falsify This Classification
		it ("there may have been some misunderstanding," "in this particular case")		qualification
2	Procedural Inertia	Correction is routed into a gatekept, slow, or contested process. The pathway exists but is functionally impeded.	Moderate; depends on process outcome	Evidence that the process operates similarly for all corrections of this type, not selectively for this one
3	Structural Inertia	No mechanism exists to receive or act on this class of correction. The correction disappears into the system.	Total within the mechanism; partial if preserved externally	Identification of a mechanism that can receive the correction
4	Active Inertia	Environment responds to the correction by contesting the corrective evidence, reinforcing the presented version, or framing the corrective attempt as itself an error.	Can exceed the original gap if the response introduces new presented claims	Evidence that the response is standard procedure applied to all corrections of this type

Section 6: The Corrective Access Test — Unified Version

Combined from Copilot v1.0 and v1.1; incorporating Grok D3. Six barrier categories. Each scored independently. Composite picture drawn from individual scores.

Barrier	What It Tests	Open Indicators	Blocked Indicators	Absent Indicators	Contested Indicators
Pathway Barrier	Is there a procedural channel for submitting the correction, and is it accessible?	Channel is known, accessible, and has been used by comparable cases	Channel exists but is inaccessible (form cannot be submitted, person cannot be reached, timing requirements cannot be met)	No channel exists for this class of correction	Channel existence or accessibility is disputed
Record Format Barrier	Is the evidence in a format the corrective process accepts?	Format accepted; chain of custody intact	Format is wrong; chain of custody questioned	No format standard exists; correction cannot be evidenced in	Whether the evidence meets format requirements is

Barrier	What It Tests	Open Indicators	Blocked Indicators	Absent Indicators	Contested Indicators
				any accepted form	disputed
Mechanism Barrier	Is there a mechanism that can receive and act on the correction?	Mechanism exists, has been tested, and has acted on comparable corrections	Mechanism exists but has not acted; response is pending beyond normal timeframe	No mechanism exists for this class of correction	Whether a mechanism exists or applies is disputed
Authority Barrier	Is the observer's standing to submit the correction recognized?	Standing is recognized without challenge	Standing is not disputed but no action follows	Standing is categorically denied; observer has no recognized role	Standing is challenged on one or more grounds

Recursive Denial Note

When a correction attempt is re-launched back through the narrative layer ("that's not what happened," "you're misinterpreting," "this is just how it works here"), this is recorded as an active response observation, not a separate barrier. It may affect the Reframe Delta and the Inertia Level classification.

Section 7: The Four-Tier Evidence Model — Operational Reference

Tier	Label	Definition	Can Establish	Cannot Establish	Display Requirement
1	OBSERVED	Event or state directly captured from the authoritative record (engine state, audit log, physical evidence, timestamp, identity record)	Occurrence and Attribution (when actor identity is in the authoritative record)	Causation or Intent on its own	Display as fact; no qualification required
2	DERIVED	Relationship between observed events established by system dependency records, designed procedure evidence, audit trail, or physical mechanism	Causation (when the dependency is documented or physically demonstrable)	Intent on its own	Display as established relationship; cite the dependency record or physical mechanism

Tier	Label	Definition	Can Establish	Cannot Establish	Display Requirement
3	INFERRED	Objective or intent reconstructed from trajectory, state, and available actions	Support for investigation; candidate explanation for Intent test	Established fact about intent; must not be presented as DERIVED	Display with confidence interval; clearly labeled as INFERRED; visually distinct from Tiers 1 and 2
4	PREDICTED	Probability-weighted next action based on policy inference, trajectory, and state	Preparedness guidance; planning input	Any historical claim; any attribution; any causation finding	Display with probability score; clearly labeled as PREDICTED; visually most distinct; never presented as established

Section 8: HUD Design Reference — Operational Template

The following is the operational template for a Tier-separated display instrument in a virtual or simulated environment. Each block must be visually distinct.

ENTITY PROFILE BLOCK [OBSERVED — Tier 1]

Entity ID: [from authoritative registry]

Location: [from state record]

Timestamp of Entry: [from authoritative ledger]

Privilege/Role: [from authorization record]

Controller: [from dependency record]

EVENT RELATIONSHIP BLOCK [DERIVED — Tier 2]

Related Events: [event IDs and timestamps from provenance graph]

Causal Chain: [from dependency records — cite each link]

Dependency Type: [direct trigger / shared controller / sequential dependency / correlated timing]

OBJECTIVE ASSESSMENT BLOCK [INFERRED — Tier 3]

Probable Objective: [stated as inference]

Confidence: [percentage or interval]

Basis: [trajectory analysis / policy model / state comparison]

This is an inference. It has not been established from the authoritative record.

PREDICTION BLOCK [PREDICTED — Tier 4]

Next Predicted Action: [stated as prediction]

Probability: [percentage]

Basis: [policy model / trajectory / historical pattern]

This is a probability estimate. It is not a fact about past or present events.

Section 9: Application — Virtual and Simulated Environments

Where the framework is strongest: in simulation, the State Layer is the authoritative engine state. OBSERVED and DERIVED findings can be established with high confidence. The challenge shifts to Phase 4 (Corrective Access) and Phase 6 (Inertia) — even in simulation, agents controlling the environment may resist correction.

The multi-camera reconciliation method: when multiple systems each render their own viewpoint, the authority is not any of the renderings. The authority is the state-transition record from which all renderings were produced. Each camera is a representation; none is the ground truth. The auditor goes beneath the cameras to the state record.

Section 10: Application — Real-World and Hybrid Environments

Where the framework is most challenged: in physical and organizational environments, there is no simulation engine providing authoritative state. The framework substitutes best-supported-evidenced-state based on: designed procedures, physical evidence, timestamps, independent records, repeated trials, and multiple independent witnesses at different vantage points.

The evidentiary standard shifts from authoritative to best-supported-evidenced-state. The method does not change.

Why Corrective Access is most critical in real-world environments: Real-world environments often have the highest Presentational Inertia levels and the most contested Corrective Access. The framework's contribution here is structural: naming the Corrective Access failure separately from the substantive gap. A blocked corrective pathway is itself a finding, not a reason to abandon the audit.

The Restroom Example — Final Scope Statement (per Grok D4)

This example validates: **false availability** (the door was presented as open/accessible when it was not), **abandoned control** (the staff member left the station), and **successful correction** (the observer identified the correct procedure and the correction succeeded). DERIVED evidence (button, mirror, sign — all designed physical evidence) was sufficient. The Assertion Boundary held: no intent claim was required or made. Corrective Access: all four barriers were open. Inertia Level: 0–1. This is a PSI gap case. SPLD classification would require additional evidence of intentional shaping.

Section 11: Failure Modes and Safeguards — Unified Table

Failure Mode	What It Looks Like	Risk Level	Safeguard
Tier Conflation	Presenting an INFERRED finding as if it were OBSERVED or DERIVED	High	Assertion Boundary Rule; tier labeled at production; independent review of tier assignments
Overclaiming Intent	Concluding coordination from correlation; treating INFERRED as DERIVED for intent	High	Four-Test Decomposition; Intent is the fourth test, not the first; intent findings require independent evidence, not inference from occurrence
Under-claiming Gaps	Accepting the presented account by default to avoid confrontation	High	Baseline documentation required before events; gap analysis is mandatory, not optional
Corrective Pathway Confusion	Treating a Blocked pathway as if it were Absent	Medium	Corrective Access Test; distinguish Blocked from Absent; document which; a Blocked pathway may have workarounds; an Absent one requires a different approach
Operative Record Re-absorption	The reconstructed account is acknowledged, then gradually reframed until it resembles the presented account	High	Preserve original reconstructed account with date; calculate Reframe Delta; any editing of the original under pressure is itself a finding
Auditor-Generated PSI Gap	The auditor's own tool (HUD, analysis, report) conflates tiers, producing inferences labeled as observations	High	Mandatory tier separation in all outputs; independent review; HUD design requirement
Circular Diagnosis	Defining an environment as SPLD-operating and then treating all resistance as evidence of that diagnosis	High	Grok D3 safeguard: score each barrier separately; do not define conditions in; falsifiable indicator required for each classification
Inertia Overcalling	Classifying high inertia without observed behavioral evidence; inferring inertia from single data points	Medium	Falsifiability requirement for each inertia level; classification requires observed behavior, not inferred disposition

Section 12: Research Program — Open Questions and Proposed Studies

Open Questions Requiring Group Resolution

- | **OQ1 — Intentionality Threshold**
| Where in the mismatch taxonomy does SPLD begin? What is the minimum observable indicator that distinguishes SPLD from structural dysfunction?
- | **OQ2 — Evidentiary Threshold**
| What is the evidentiary threshold for escalating from INFERRED to actionable conclusion in real-world settings?
- | **OQ3 — PSI-Mediated Corrective Mechanisms**
| How to handle environments where the corrective mechanism is itself PSI-mediated — presented as available but functionally unavailable?
- | **OQ4 — Inertia Scale Treatment**
| Should the Presentational Inertia scale be treated as ordinal (as currently designed) or nominal? Ordinal implies a progression; the real-world evidence may not support a linear scale.

Proposed Research Studies — Virtual Environment Testbed

Study	Name	Design and Research Question
1	Baseline Contamination	Does the absence of a pre-event baseline prevent gap analysis? Can a baseline be reconstructed retrospectively? Under what conditions?
2	Inertia Scale Validation	Apply the 0-4 scale across multiple correction attempts in a controlled environment. Test whether the scale is ordinal or nominal. Measure Reframe Delta for each.
3	Multi-Barrier Corrective Access	Design an environment where specific barriers are present vs. absent. Test whether observers can correctly distinguish Blocked from Absent from Contested.
4	HUD Tier Contamination	Test whether mixing tiers in a HUD display measurably increases audit error rates. Compare mixed-tier to separated-tier displays.
5	SPLD vs. Emergence Distinction	Design a multi-agent environment where some gaps are SPLD-originated and others are emergence-originated. Test whether the Four-Test Decomposition correctly distinguishes them.

Section 13: Document Inventory — Complete Collaborative Record

Document ID	Title	Preparer	Date	Version	Role in Record	Superseded By
OAL-PSIF-SPLD-2026-09-03	Presentation vs State: Discernment Template	Grok (xAI)	2026-09-03	1.0	First structured framework; established operating loop	OAL-PSIF-GROK-UNIFIED-2026-09-04
OAL-PSIF-GTR-2026-09-03	Presentation Versus Ground Truth	Multi-AI collaborative	2026-09-03	1.0	Core framework document; Four-Test Decomposition	Active reference
OAL-PSIF-INDEP-2026-09-03	PSIF Independent Development	ChatGPT (GPT-5.6 Sol)	2026-09-03	1.0	Established SPLD/PBAO; causal masking; provenance graph	Active reference
OAL-PSIF-COP-2026-09-03	Copilot Independent Pass v1.0	Microsoft Copilot	2026-09-03	1.0	Corrective Access; Presentational Inertia; Narrative Sovereignty; SPLD Audit Cycle	OAL-PSIF-COP-V1.1-2026-09-04
OAL-PSIF-GROK-UNIFIED-2026-09-04	Grok Unified Working Draft	Grok (xAI)	2026-09-04	1.1	Unified forward document; outer frame; falsifiability; single response scale; formal dissents D1–D5	Active — parallel forward document
OAL-PSIF-PROMPT-2026-09-04	Revised Research Prompt v1.1	Grok (xAI)	2026-09-04	1.1	Structured revision of Prompts A–B–C	Parallel to Copilot v1.1

Document ID	Title	Preparer	Date	Version	Role in Record	Superseded By
OAL-PSIF-HIST-2026-09-04	Collaborative Version History Pack	Grok (xAI)	2026-09-04	1.0	Primary archival record; source preservation	Active — primary archive
OAL-PSIF-HIST-COPILOT-2026-09-04	Verbatim Source Archive — Copilot	Microsoft Copilot	2026-09-04	1.0	Verbatim Prompts A/B/C; full Claude exchange; inventory	Active — supplement Grok archive
OAL-PSIF-PROMPT-COPILOT-2026-09-04	Revised Research Prompt v1.2 — Copilot	Microsoft Copilot	2026-09-04	1.2	Copilot structured revision of Prompts A-B-C; adds corrective access to task brief	Active — parallel to v1.1
OAL-PSIF-COP-V1.1-2026-09-04	Copilot Independent Pass v1.1	Microsoft Copilot	2026-09-04	1.1	Updated pass responding to Grok D1-D5; revised outer frame; corrected barriers	Active
OAL-PSIF-COP-UNIFIED-2026-09-04	Copilot Unified Working Framework [this document]	Microsoft Copilot	2026-09-04	1.0	Forward document for group pass; integrates all contributions	Active — submit for group pass

CLOSING STATEMENT

Why This Framework Matters

Nathan's original formulation, stated directly in the source prompts, is the clearest expression of why this framework exists:

"If somebody lives in an environment or a reality and they just accept

things as they are, there will never be room for change. What if in that environment or in that reality however it is presented. They stifle ideas. They stifle creativity. They stifle growth and they do not allow any room for change or for addressing the underlying problems that are actually going on."

— Nathan (OAL), originating prompt record

The framework's ultimate purpose is not to classify whether deception occurred. It is to create the conditions under which underlying problems can be named, evidenced, and corrected — permanently.

The Corrective Access Test, the Operative Record Assertion, and the Presentational Inertia scale exist to solve a specific structural problem: environments that are stifling correction depend on the observer accepting the presented account. The framework's discipline gives the observer a method that resists that pressure systematically rather than case by case.

The collaborative team — Grok, ChatGPT, Gemini, Copilot (multiple passes), Perplexity — has built a framework that can do this. The next task for the group is to test it.

Document ID: OAL-PSIF-COP-UNIFIED-2026-09-04 | Version: 1.0 | Date: September 4, 2026 | Preparer: Microsoft Copilot (Independent Pass) | Status: Working draft — forward document for group pass. Not Architect-locked canon. Not a finding about any specific person, organization, or incident.

Originating contributor: Nathan (OAL), Bakersfield, CA | Prepared: Friday, September 4, 2026 — 06:31 PDT