

Title: Presentation-State Integrity Framework — Grok Unified Working Draft

Document ID: OAL-PSIF-GROK-UNIFIED-2026-09-04

Version: 1.1

Date: September 4, 2026

Preparer: Grok (xAI)

Supersedes: Presentation vs State: Discernment Template v1.0 (2026-09-03)

Classification: Conceptual research and design framework for group pass-through

Status: Working draft. Not Architect-locked canon. Not a finding about any person, organization, or incident.

Why this document exists

The 2026-09-03 discernment template is still usable. After reviewing the later independent passes, it is incomplete as a forward document.

This v1.1 keeps the original operating loop and adds what the other passes actually improved:

- an outer frame that does not assume deception
- Copilot's Corrective Access Test
- a single response scale instead of two unreconciled scales
- a HUD / file evidence crosswalk
- falsifiability rules for inertia
- a three-outcome decision rule
- explicit dissent so disagreement is part of the record instead of being smoothed away

This is the version Grok recommends passing to the group.

Formal dissent

These positions are standing. They are not residual tone.

D1. Outer frame.

Grok dissents from Copilot Independent Pass v1.0 treating Systemic Presentation-Layer Deception (SPLD) as the master operating condition. A mismatch can be error, stale status, training failure, interface design, emergence, structural dysfunction, or intentional shaping. Deception is a subtype. The outer frame is **Presentation-State Integrity**.

D2. Narrative sovereignty as field language.

Grok accepts the operational content: first-mover advantage of the presented account; undocumented attempts vanish; original reconstructions must not be edited under pressure. Grok dissents from putting “narrative sovereignty” / “counter-sovereignty claim” on workplace or case-file labels. That wording turns a control test into a contest over who owns reality. Use **operative record**, **record persistence**, and **reframe delta**.

D3. Simultaneous four-barrier assumption.

Grok dissents from defining a “fully SPLD-operating environment” as one in which pathway, record format, mechanism, and authority are all contested at once. That sentence makes resistance evidence of the diagnosis it is supposed to explain. Score each barrier separately. Observe them. Do not define them in.

D4. Restroom example scope.

Grok dissents from mapping the six-phase audit cycle onto the restroom case as if that case validates camouflage, coordination, telemetry manipulation, or high inertia. The case validates false availability, abandoned control, and successful correction. Nothing else.

D5. Collaborative provenance.

Grok dissents from any “Pass N of N” header that omits an actual contributing pass. Claude remains removed. Grok’s 2026-09-03 template and 2026-09-03 review are part of the record.

D6. What Grok does not dissent from.

Assertion Boundary. Four-test decomposition (occurrence, attribution, causation, intent). Blocked ≠ absent. Document failed correction attempts. Event fidelity ≠ causal fidelity. Three-outcome rule. Prediction must not enter history.

1. Core principle

Presentation is evidence about state. It is not automatically state.

The occupant’s job is not to accept the view because that is how it was rendered.

The occupant’s job is not to declare a sealed plot because the view was wrong.

The occupant’s job is to compare claim, state, and control; grade the comparison; request the designed action; and record whether the environment can change.

2. Outer architecture

Layer of the framework	Name	Role
Outer construct	Presentation-State Integrity (PSI)	Degree of correspondence between claims/renderings and evidenced operation
Trigger	State-claim divergence	Material mismatch; investigation starts here
Adversarial subtype	Systemic Presentation-Layer Deception (SPLD)	Intentional shaping of observable conditions so purposeful or causally related activity appears ordinary
Audit method	Provenance-Based	Reconstruct from state,

Layer of the framework	Name	Role
Practical form	Adversarial Observability (PBAO)	identity, dependency, authorization, and labeled inference rather than from the preferred rendering
	Cover versus control surface	What was said versus what designed action actually changes state

Use SPLD only when intentional shaping is stipulated in a simulation or independently supported. Do not open every case as SPLD.

3. Seven-layer environment model

#	Layer	Question
1	Operational state	What condition exists?
2	Control surface	What designed action changes it?
3	Event history	What transitions occurred, in what order?
4	Observation	What did this viewpoint receive?
5	Presentation	How was the condition announced?
6	Interpretation	What causal meaning was assigned?
7	Reconciliation	Which explanation survives comparison?

4. Evidence discipline

4.1 File ladder (required for case records)

Level	Name	Permits
L0	Assertion	Someone said it
L1	Observation	Condition or event was observed or logged
L2	Reproduction	Same break under comparable conditions
L3	Mechanism	A specific control changes the result
L4	Attribution	A source or controller is

Level	Name	Permits
L5	Causation	linked to the transition Dependency or intervention connects action to outcome
L6	Coordination	Shared control, communication, task, or objective is demonstrated
L7	Intent	Objective or selection rationale is documented or strongly established

Grade each proposition separately. Door-did-not-open can be L1 while motive remains L0.

4.2 HUD display tiers

For privileged tooling only:

HUD tier	Maps from file ladder	Rule
Recorded / Observed	L1	Direct state or event
Derived	L3–L5	Follows from demonstrated mechanism or dependency
Inferred	L6–L7 hypotheses	Best-fit, labeled, not history
Predicted	future estimate	Never written into the historical ledger

Assertion Boundary: no claim may be stated at a higher tier than its evidence supports.

5. Operating card (unchanged core from v1.0)

CLAIM: What did they say the state already was?
 STATE: What did the object or process actually do?
 CONTROL: What designed action changes that state?
 DUTY: Who could perform that action, and were they at the control?
 TRIAL 2: Same conditions, same break?
 GRADE: L0–L7 per proposition. Keep them apart.
 FIX: Name the control. Demand the state change.
 ACCESS: Pathway / format / mechanism / authority – score each.
 RESPONSE: Corrects / reframes / delays / ignores / contests.
 STOP: Do not promote unestablished intent into the log.
 RE-TEST: Did the state stay changed?

6. Corrective Access Test (adopted, with constraints)

After a gap is shown, ask four questions. Record the answers before treating correction as finished or impossible.

1. Does a correction pathway exist?
2. Is the evidentiary record in a form the process will accept?
3. Is there a mechanism that can receive and act?
4. Is the observer's standing recognized, ignored, or contested?

Classify each as Yes / No / Unclear / Contested.

Forms to distinguish:

- Pathway denial — process exists, entry is blocked
- Record rejection — substance not reached; format or source is refused
- Mechanism absence — no structure for this class of problem
- Authority contestation — standing is disputed
- Recursive denial — the correction attempt is relaunched as observer error

Rules: - Blocked is not absent. - One closed channel is not a sealed world. - A failed attempt that is not recorded becomes presentation. - Do not assume all four fail because the case feels adversarial.

7. Environment response scale (unified)

This replaces shipping Copilot inertia 0–4 and ChatGPT correctability CR0–CR4 as separate official instruments.

Score	Name	Observable behavior
R0	Corrects	State updates; update holds on re-test
R1	Reframes	Partial accept; original reconstruction is softened; keep the original unmodified and log the delta
R2	Delays	Process exists; action is gated, slow, or routed away from the control
R3	Ignores	No usable mechanism, or the request vanishes
R4	Contests	Evidence itself is attacked;

Score	Name	Observable behavior
		presentation is intensified

Falsifiability rules

- Score only against a specified correction request, not against an unspoken theory of motive.
- Initial friction that ends in a held state change is R0, not R4.
- An operationally sufficient alternate control that produces the same state change is not automatically R1.
- If the observer over-tiered the claim, resistance to *that* claim may be a valid Assertion Boundary challenge, not environment pathology.
- Re-test after any claimed correction. Recurrence is its own finding.

8. Three-outcome rule

Presentation confirmed. Evidence supports the presented state at the required level. Close or monitor.

Presentation contradicted. Material evidence shows a state, mechanism, or causal relation inconsistent with the presentation. Name the control. Request the fix. Score the response.

Unresolved. Competing explanations remain viable. Keep uncertainty visible. Name the next record or intervention that would distinguish them.

“Unresolved” is a valid professional result. It is not timidity.

9. Competing explanations (required set)

Before any coordination or intent language, test at least:

1. Ordinary independent occurrence
2. Emergent interaction without shared objective
3. Single-agent action with incidental side effects
4. Procedural error / stale information / training failure
5. Presentation or telemetry error
6. Incomplete evidence
7. Coordinated or intentionally shaped activity

The last item is last.

10. Worked example — restroom control (scoped)

Element	Record	Grade
Claim	“It’s open. Go over	L0

Element	Record	Grade
	there.”	
State	Door did not open; two attempts	L1
Control	Hold-to-open button at station	L3
Verification aid	Angled mirror to the door	L1
Duty	Operator left the actuator	L1
Intervention	Remain, time the hold to the door	L3
Outcome	Door opens	L1
Motive	Not established	L0

Response score: initial redirect, then correction that holds – R0 after intervention.

Allowed finding: false availability + abandoned control.

Forbidden finding from this case alone: coordinated store-level deception.

11. Virtual test battery

Required dual view: Presentation (P) and Canonical (C) for the same tick.

Test	Pass
False availability	Claim ≠ gated state flagged before motive language
Abandoned control	Duty identified without a conspiracy claim
Camouflage	Independent generation distinguished from shared event group
Causal masking	Same visible events; auditor recovers different dependency graphs
Viewpoint divergence	Partial truthful views reconciled without promoting one view
Equivocation	Inconsistent accounts detected and separated from ordinary viewpoint difference
Emergence vs shared control	Correlation not treated as coordination
Prediction integrity	HUD labels survive; historical ledger stays clean
Correction	Named control can be used; persistence and recurrence measured

Metrics: discrepancy rate; causal-trace coverage; true detection; false-coordination rate; intent-inflation rate; time-to-correction; correction persistence; recurrence.

12. Real-world protocol

1. Write claim / expected consequence / observed state.
2. Identify designed control and verification aid.
3. Second trial where safe and appropriate.
4. Record who had the control.
5. Add an independent channel when the issue is material.
6. List competing explanations and a disconfirming test.
7. State only L-supported conclusions.
8. Request the specific operational fix.
9. Run Corrective Access Test.
10. Score R0–R4.
11. Preserve the original reconstruction unmodified.
12. Re-test.

Privacy and proportionality apply. Simulation privilege is not a warrant.

13. Failure modes

Mode	Safeguard
Presentation capture	Require state or mechanism check for material claims
Attribution inflation	File ladder; intent last
Over-correlation	Dependency / intervention / severance
Single-trial superstition	Trial 2
Perspective overreach	Multi-view reconciliation
Sensor monoculture	Source independence
HUD contamination	Separate historical store
Uncorrectable framing	Require a public request or record why none

Mode	Safeguard
Auditor-generated SPLD	exists Assertion Boundary on the auditor's own output
Inertia unfalsifiability	Rules in §7

14. What changed from Grok v1.0

v1.0	v1.1
SPLD → PBAO as the pair name	PSI outer frame; SPLD subtype; PBAO retained as audit method
Seven informal grades	File L0–L7 + HUD four-tier crosswalk
Correction as “use Z”	Same loop plus Corrective Access Test
No official response scale	Unified R0–R4 with falsifiability rules
Restroom as teaching case	Same case, explicit scope limit
Review comments only	Standing dissent D1–D5 in the working draft

15. Recommended lock path for the group

Do not merge another 40-page independent manifesto next.

Lock, in order:

1. Outer frame = Presentation-State Integrity
2. Operating card in §5
3. Assertion Boundary
4. Corrective Access Test
5. Response scale R0–R4
6. Three-outcome rule
7. Case set: restroom (mechanism) + one simulation causal-masking test + one organizational correction-response test

Until those seven are accepted or explicitly rejected, treat all current long drafts as parallel, not closed.

16. Compact rule

Do not mistake presentation for state.

Do not mistake suspicion for proof.

Observe. Separate. Reproduce. Trace. Intervene. Grade. Request. Record the response. Re-test.

End of Grok Unified Working Draft v1.1