
OAL-PSIF Revised Research Prompt v1.2

Copilot Independent Pass — Structured Design Brief

Document ID	OAL-PSIF-PROMPT-COPILOT-2026-09-04
Version	1.2
Date	September 4, 2026
Preparer	Microsoft Copilot (Independent Pass)
Supersedes	OAL-PSIF-PROMPT-2026-09-04 v1.1 (Grok, xAI) — this is an independent parallel revision, not a replacement
Classification	Derivative prompt revision for group pass-through
Status	Working revision. Not a substitute for verbatim source. Not a finding about any person, organization, or incident.
Verbatim Source	OAL-PSIF-HIST-COPILOT-2026-09-04 (Copilot Verbatim Source Archive) and OAL-PSIF-HIST-2026-09-04 (Grok Version History Pack)

Authorship Notice — Read Before Use

This document is a Copilot-issued structured revision derived from Nathan's original Prompts A, B, and C. It is **not** a verbatim copy of those prompts. Verbatim text is preserved in the separate archival document (OAL-PSIF-HIST-COPILOT-2026-09-04). This document must not be cited as Nathan's original wording. Every load-bearing requirement from the original prompts is carried forward without omission or softening. Copilot additions are identified explicitly in the versioning differences section below.

VERSIONING Version History and Differences

This document is version 1.2 in the prompt revision series. Version 1.1 was issued by Grok (xAI) as OAL-PSIF-PROMPT-2026-09-04. That version is acknowledged and adopted as the structural baseline for this revision. Version 1.2 maintains all of Grok's structural decisions and adds Copilot's independent revisions in the areas identified below.

Differences: v1.1 (Grok) vs. v1.2 (Copilot)

Change ID	Section Affected	Description
C1	Section 1 — Task	Corrective Access requirement made explicit in the task brief, not only in the method section.
C2	Section 4 — Method Outputs	Narrative Record Persistence requirement added as §4.3.
C3	Section 5 — HUD/Predictive Access	Brief stipulation added on HUD design requirement: Tier 4 (Predicted) labeling must be visually distinct from Tiers 1–3 in all display outputs.
C4	Formal Dissent	Copilot dissent position CPD1 added (scope of revised prompt series). All Grok v1.1 dissent positions preserved without modification.

DISSENT Formal Dissent Positions

The following dissent positions are recorded in full and carry forward into all derivative work from this document. Dissent positions are not negotiating stances — they are structural constraints on how findings, language, and method outputs may be expressed.

Grok Dissent Positions — Adopted Without Modification

- D1 — Outer Frame:** The outer frame is *Presentation-State Integrity*, not SPLD. Deception is a subtype within that frame, not the frame itself. Copilot adopts this position fully and without qualification.
- D2 — Narrative Sovereignty Language:** Use *operative record*, *record persistence*, and *reframe delta* in field and case-file contexts. The term *narrative sovereignty* is retained as a conceptual label only and may not be used as a field label in any method output, classification, or finding.
- D3 — Simultaneous Barrier Assumption:** Score each Corrective Access barrier separately. The four barriers (pathway existence, format acceptance, receiving mechanism, observer authority) must not be defined or scored as a set. A failure on one does not substitute for assessment of the others.
- D4 — Restroom Example Scope:** The restroom example validates: false availability, abandoned control, and successful correction. It does not validate: camouflage, coordination, telemetry manipulation, or high inertia. The example may not be extended beyond the phenomena it was constructed to illustrate.

D5 — Collaborative Provenance: Grok's 2026-09-03 template is part of the record and must be credited as such. Claude remains removed from this workflow. No Claude output is part of this revision series.

Copilot Dissent Position — Added in v1.2

CPD1 — Revised Prompt Scope: Copilot dissents from any future pass that treats Grok's v1.1 revised prompt as the only valid structured brief for this research task. This document exists as an independent parallel revision. Both versions (v1.1 and v1.2) may be passed to the group. Neither version supersedes the other. The group may use whichever version better fits their task, or may use both in parallel. Convergence on a single version must be a deliberate group decision, not an assumed default.

SECTION 1 Task

Design terminology, a layered model, and a testable method for distinguishing:

- how an occupied environment is **presented**
- how that environment actually **operates**
- which underlying problems can be **named and corrected** because they are evidenced and not merely felt

v1.2 Addition — Corrective Pathway Requirement:

The method must also address what the observer does *after* the gap between presentation and state is established. Detection is a necessary but not sufficient outcome. The task includes designing a corrective pathway framework applicable to both virtual and real-world environments.

Core Constraints

- Do not treat the presented view, spoken cover, local rendering, or institutional explanation as complete by default.
- Do not treat every mismatch as coordinated deception, persecution, or readable private motive.
- The goal is **discernment and correction** — not denial and not inflation.
- Do not filter or soften the phenomenon in order to make it more palatable. Name it as it is.
- Do not re-open the debate over whether real-world examples (McDonald's, traffic, conference calls) belong in a simulated or fictional framework. That debate is closed. Analyze the control surface.

SECTION 2 Stipulated Setting

The setting may be: a simulation, a hybrid testbed, a training environment, an organization, or ordinary physical space. Simulated environments may contain branded spaces, real procedures, metrics, partnership objects, and real-world testing protocols. The presence of real-world elements in a simulated environment does not require a debate about categorical boundaries.

The Environment Contains

- Agents and processes with stated functions and unstated objectives
- Environmental events that may be independent, emergent, or coordinated
- A presented account of what those events mean
- A privileged or disciplined observer who can compare presentation to state

The Privileged Observer's Access May Include

Setting Type	Available Access Points
Simulated / Virtual Environment	Engine state, entity identity, event trigger, controller, event group, authorization records, queued action
Physical / Organizational Environment	Designed procedure, physical evidence, timestamps, repeated trials, independent witnesses, control surface identity, object behavior record, statement-vs-requirement comparison

On Predictive Inference: Predictive inference (the HUD / Hermione analogue) is plan recognition from state and affordances — it is not mind reading. It must be labeled as inference and may not be presented with the authority of observation. See Section 5 for HUD design requirements.

SECTION 3 Phenomenon to Name

The phenomenon has the following required properties. **Any framework that omits one of these properties is incomplete.**

Property 1 — The Presented Account Claims Normalcy

Agents and environmental processes produce ordinary-looking events and ordinary-

sounding explanations. The intended reading is: this is chance, this is baseline, this

is how the place always works, nothing here needs to be addressed.

Property 2 — The Underlying State Differs from the Presented Account

There is a difference between what the environment presents and what the

authoritative event record shows actually occurred. This difference may be in

occurrence, attribution, causation, or intent — but these four must be tested

separately and may not be collapsed into a single conclusion.

Property 3 — The Presented Account Bids for Permanence

The environment is not merely presenting a false account once. It is structured to

maintain that account as the operative record indefinitely. The assumption is that

the presented account will go unaddressed in perpetuity.

Property 4 — The Observer's Job Is to Address What Is Actually Happening

Not to accept the presented account. Not to assume every mismatch is coordinated.

To reconstruct what actually occurred from the evidence available, name it, and

create conditions for correction.

Property 5 — The Environment May Resist Correction

The environment may behave in ways that make correction difficult after the gap is

established. This resistance is a phenomenon to be named and documented — not a

reason to abandon the audit.

SECTION 4 Required Method Outputs

The method must produce all four of the following outputs. An output that is missing renders the method incomplete for the purposes of this task.

4.1 — Gap Finding

A structured comparison between the presented account and the evidenced account. Each claim must be labeled by tier:

Tier	Label	Meaning
1	OBSERVED	Drawn directly from authoritative state or physical evidence — entity identity, timestamp, location, authorization
2	DERIVED	Causal relationship established by dependency record, audit trail, or designed procedure
3	INFERRED	Probable relationship or objective reconstructed from trajectory and state — must carry confidence interval
4	PREDICTED	Next action or state, probability-weighted — must be visually and structurally distinct from Tiers 1–3

Claims may not be stated at a higher tier than their evidence supports.

4.2 — Corrective Access Assessment

After a gap is established, the method must assess whether a corrective pathway exists. Each of the following barriers must be scored separately:

- **Pathway Existence:** Does a correction pathway exist? (Yes / No / Unclear)
- **Format Acceptance:** Is the evidence in a format the corrective process accepts?
- **Receiving Mechanism:** Is there a mechanism that can receive and act on the correction?
- **Observer Authority:** Is the observer's authority to correct recognized?

A blocked pathway is not the same as an absent one. A failed correction attempt that is not documented becomes invisible. Both must be distinguished and recorded.

4.3 — Record Persistence Standard *[v1.2 Addition]*

The reconstructed account must be preserved in a form that cannot be reabsorbed into the presented account.

Setting	Required Persistence Form
Virtual / Simulated Environment	Write-protected event ledger entry
Physical / Organizational Environment	Dated, independent, persistently stored record

Prohibited outcome: The reconstructed account is acknowledged and then gradually reframed until it resembles the presented account. This outcome must be treated as a record persistence failure, not as a resolution.

4.4 — Inertia Classification

The environment's response to the corrective attempt must be classified using the following scale. **Each level must be falsifiable. A Level 4 classification requires observed behavior, not inferred intent.**

Level	Label	Description
Level 0	Immediate Update	Environment immediately updates presented state to reflect the corrected account.
Level 1	Acceptance with Reframe	Environment accepts correction but reframes it — softens, contextualizes, or qualifies the corrected account.
Level 2	Gatekept Routing	Environment routes the correction into a gatekept or slow process without updating the current presented account.
Level 3	No Receiving Mechanism	Environment has no mechanism to receive or act on the correction.
Level 4	Active Reinforcement	Environment responds to correction by reinforcing the presented version and contesting the corrective evidence. Requires observed behavior — not inferred intent.

SECTION 5 Privileged Access and HUD Design

The observer's privileged access may include predictive inference — the HUD or Hermione analogue. The following requirements apply to any HUD or equivalent display instrument.

5.1 — Tier Separation Is Mandatory

The HUD or equivalent display must visually and structurally separate all four tiers:

- **OBSERVED** — entity identity, timestamp, location, authorization — drawn from authoritative state
- **DERIVED** — causal relationships established by dependency records or audit trail
- **INFERRED** — probable objective, reconstructed from trajectory and state — labeled with confidence interval
- **PREDICTED** — next action, probability-weighted — labeled as probabilistic, visually distinct from Tiers 1-3

5.2 — The Prediction Layer Is a Tier 4 Instrument *[v1.2 Stipulation]*

Predictive inference is always labeled as inference. It may guide preparedness. It may not be presented as established fact.

Critical constraint: Mixing tiers in the HUD output produces a system that generates its own presentation-layer error — which would be the auditor's own tool performing the same failure mode it is designed to detect. Tier mixing is not a formatting error. It is a method failure.

5.3 — Reconciling Multiple Viewpoints

When multiple cameras, agents, or systems each render their own viewpoint, the reconciliation method is **not** to average those renderings. The authority is the state-transition record from which all renderings were produced. Each camera is a representation, not the authority. Disagreement among renderings is data — it is not a reason to select the majority view.

SECTION 6 Testability Requirements

The framework must be testable in both virtual and real-world settings. The following testability requirements are binding on any derivative framework or method output.

6.1 — Falsifiability

Each claim produced by the method must be stated in a way that could be proven wrong. Inertia classifications must be based on observed behavior, not inferred disposition.

6.2 — Repeatability

The method must produce consistent results when applied by independent observers to the same evidence set.

6.3 — Separation of Tests

The four event tests must be run separately. A finding on one does not automatically establish a finding on another:

- **Occurrence:** Did the event happen?
- **Attribution:** Who or what was responsible?
- **Causation:** What produced the outcome?
- **Intent:** What objective was being pursued?

6.4 — Real-World Applicability

In physical and organizational environments, the framework substitutes *best-supported evidenced state* for *authoritative engine state*. The evidentiary standard shifts but the method does not. Physical evidence, designed procedures, timestamps, repeated trials, and multiple independent witnesses are the physical-world analogues of engine state.

SECTION 7 What This Prompt Does Not Require

For clarity of scope, the following are explicitly outside the requirements of this prompt:

- This prompt does **not** require the observer to prove that events were deliberately coordinated in order to name the gap between presentation and state.
- This prompt does **not** require a debate about whether fictional and real-world settings should be mixed. That debate is closed. Analyze the control surface.
- This prompt does **not** require the observer to have certainty about intent before addressing what is evidenced.
- This prompt does **not** require the framework to apply only in simulation. It applies wherever a gap between presentation and state can be evidenced and a corrective pathway can be identified.

Closing Note — Derived Document Status

This is a derived document. Nathan's original prompts are preserved verbatim in **OAL-PSIF-HIST-COPILOT-**

2026-09-04. If a sentence in this document does not appear in that verbatim archive, it is a Copilot

reconstruction or addition. All Copilot additions are identified as such in the versioning differences section

(Section: Versioning) of this document.

September 4, 2026 | Status: Working Revision