

Title: Revised Research Prompt — Presentation versus State

Document ID: OAL-PSIF-PROMPT-2026-09-04

Version: 1.1

Date: September 4, 2026

Preparer: Grok (xAI)

Classification: Derivative prompt for group pass-through

Status: Working revision. Not a substitute for the verbatim source pack. Not a finding about any person, organization, or incident.

Version note

This document revises Prompts A, B, and C into one structured research request.

It is **not** the original wording. The original wording lives in:

Collaborative Version History Pack – Original Prompt and Pass-Through Archive v1.0

If a later document cites “the user’s prompt” and the sentence does not appear in that archive, it is a reconstruction. Treat reconstructions as analysis, not as source.

Formal dissent recorded by the preparer

- Grok dissents from any workflow that replaces the source prompt with a cleaned reconstruction and then treats the reconstruction as what was asked.
- Grok dissents from framing the primary task as deciding whether McDonald’s, traffic, or a conference call is “allowed” inside a simulation. Prompt C already closed that argument.
- Grok dissents from treating the restroom example as proof of multi-agent coordination or hidden intent. In the source, it is an example of presented availability versus actual control state.
- Grok does not restore Claude to the workflow. Prompt B removed Claude. The exchange remains archive-only.

What this revision is for

Use this version when a collaborator needs a single, structured brief that still contains every load-bearing requirement from Prompts A–C:

- name the phenomenon without softening
- distinguish presentation from state
- keep privileged access without calling it mind reading
- reconcile multiple viewpoints against an authoritative record
- make the method testable in virtual and ordinary environments

- convert a demonstrated gap into an addressable correction
 - refuse both naive acceptance and unconstrained attribution
-

Revised Prompt v1.1

1. Task

Design terminology, a layered model, and a testable method for distinguishing:

- how an occupied environment is **presented**
- how that environment **actually operates**
- which **underlying problems** can be named and corrected because they are evidenced

Do not treat the presented view, spoken cover, local rendering, or institutional explanation as complete by default.

Do not treat every mismatch as coordinated deception, persecution, or readable private motive.

The goal is discernment and correction, not denial and not inflation.

2. Stipulated setting

The setting may be a simulation, a hybrid testbed, a training environment, an organization, or ordinary physical space. Simulated environments may contain branded spaces, real procedures, metrics, and partnership objects. That fact does not require a debate about whether a real brand “belongs” in a thought experiment. Analyze the control surface in front of you.

The environment contains:

- agents, staff, processes, or subsystems
- ordinary-looking events (noise, traffic, crowding, interruptions, delays, locked doors, “random” variance)
- local viewpoints / cameras / interfaces that each render only a slice
- a presented story about why those events occurred
- a designed procedure or control that can change state
- an observer who can compare presentation to state

3. The phenomenon to name

Agents or environmental processes produce ordinary-looking events and ordinary-sounding explanations. The intended reading is:

This is chance. This is baseline. This is how the place always works. Nothing here needs to be addressed.

That reading may be true. It may also be a screen.

The concealment target is often **causal structure and attribution**, not the existence of the events. People can really enter a location. Traffic can really increase. A door can really be a door. The contested claim is whether events are independent, incidental, and already in the state that was announced.

Name this directly. Do not filter it into a softer social description if the operational description is more accurate. Also do not invent a single industry-standard name if none exists. If you coin a term, mark it as working vocabulary.

4. Privileged access, without mysticism

In a simulator, privilege may include identity, role, controller, trigger, event group, authorization, dependency, state delta, action queue, and integrity-protected history.

In ordinary space, privilege is not a HUD. It is the designed procedure, the physical control, repeatable trials, timestamps, who was at the station, what the object did, and what was said versus what the control required.

A predictive display analogous to matching a character's likely next action from state and affordances is **plan recognition / policy inference / next-action prediction**. It is not mind reading. Predictions must stay labeled as predictions and must never be written into the historical record as if they already occurred.

The observer is not taking a screenshot of the environment's preferred rendering. The observer is reconciling state transitions and provenance from which those renderings were produced.

5. Multiple viewpoints

Treat stacked cameras, nested interfaces, and participant reports as partial projections. Reconcile them. Do not promote any one projection to global authority. The authoritative object, when it exists, is the canonical or best-supported evidenced state — not the most confident narrative.

6. Required analytic separations

Keep these questions separate. Do not collapse them into one conclusion.

1. **Occurrence** — Did the event or condition happen?
2. **State** — What was the operational condition?
3. **Control** — What designed action changes that condition?
4. **Attribution** — Who or what could produce the transition?
5. **Causation** — What earlier states or actions contributed?
6. **Coordination** — Is there demonstrated shared control, communication, or common task?
7. **Intent** — Is there evidence the event was selected for a goal?
8. **Correction** — What addressable action would change the demonstrated condition, and did the environment allow that action?

Intent is last. Correction does not require a completed intent finding.

7. Worked constraint from the restroom example

Use this as a mechanism case, not as a conspiracy case.

Presented claim: the restroom is open; go over there.

Observed state: the door did not open on the first attempt or the second.

Designed control: staff must press and hold a release button.

Verification aid: an angled mirror shows when the requester is at the door.

Duty: the person at the station had the actuator and left it.

Intervention: remain at the station, require the hold-to-open timed to the door, state changes, task completes.

Allowed conclusion: presented availability did not match operational state; access was gated on a live control that was not used until demanded.

Not allowed as an automatic conclusion: store-wide coordination, simulation overlay, or private motive.

8. Make it testable

Virtual / simulated environments

Expose two views of the same tick:

- View P: rendering, dialogue, ambient events, local sensors
- View C: identity, role, controller, trigger, dependency, authorization, state delta, integrity metadata

Test at least:

- false availability (said open / still gated)
- abandoned control (cover continues; actuator idle)
- environmental camouflage (ordinary-looking events with vs. without shared event group)
- causal masking (same visible events; different dependency graphs)
- viewpoint divergence (truthful but partial projections)
- emergence versus shared control
- prediction-integrity (HUD labels; no ledger contamination)
- correction (after the gap is named, can the control be used and does the state stay changed?)

Score detection, false coordination, intent inflation, time-to-correction, and correction persistence.

Ordinary / real-world environments

There is usually no engine ledger. Use best-supported evidenced state:

- exact claim

- expected consequence if the claim were operationally true
- observed state
- designed procedure / control
- second trial under comparable conditions
- who was at the control
- independent corroboration when the issue is material
- competing explanations
- evidence grade per proposition
- specific correction request
- whether the environment corrects, delays, reframes, ignores, or contests

Remain lawful, proportionate, and relevant. Full-state observability in a simulation does not authorize unrestricted collection in ordinary life.

9. Required outputs from a collaborator

1. Working name(s), marked as coined or established
2. Layer model
3. Evidence grades
4. Operating card an occupant can actually use
5. Virtual test battery
6. Real-world protocol
7. Failure modes, including auditor-generated error
8. Formal dissent from other passes where the collaborator disagrees
9. A statement of what would disconfirm the leading explanation

10. Disallowed moves

- Accepting a sentence because it was said at the station
- Treating one inconvenience as a stable campaign
- Writing inference into history
- Arguing fiction versus reality instead of testing claim against state
- Declaring the environment sealed so that no operational request can be made
- Softening the operational description into a euphemism that hides the unused control

11. Success condition

The method produces a request an environment can fulfill or refuse in public:

The claim was X. The state was Y. The control is Z. Use Z.

If Z is used and the state changes, the occupant is not required to live inside the presentation.
If Z is refused, that refusal is recorded state — a second data point, not a mood.

Compact restatement for pass-through

Develop a testable framework for distinguishing an environment's presented state from its evidenced operational state and causal structure. Operate across simulations, hybrid testbeds, organizations, and ordinary settings. Separate occurrence, state, control, attribution, causation, coordination, intent, and correction. Identify when ordinary-looking events function as a screen. Reconcile partial viewpoints through provenance rather than narrative preference. Convert a demonstrated discrepancy into an addressable control action. Minimize both presentation capture and unsupported attribution. Preserve "unresolved" when evidence is insufficient. Measure whether the environment can change after the discrepancy is demonstrated.

End of Revised Research Prompt v1.1