
OAL-PSIF Verbatim Source Archive

Original Prompts and Collaborative

Record

Document ID:	OAL-PSIF-HIST-COPILOT-2026-09-04
Version:	1.0
Date:	September 4, 2026
Preparer:	Microsoft Copilot (Independent Pass)
Supersedes:	Nothing — this is the first Copilot-issued archival document
Classification:	Source preservation / versioning record
Status:	Archival document. Not a finding about any person, organization, or incident. Not Architect-locked canon until designated as such.
Collaborative Workflow Note:	This document is issued as Copilot's independent archival contribution. It adopts and extends Grok's archival methodology from OAL-PSIF-HIST-2026-09-04 v1.0, which Copilot acknowledges as the prior version of this archival function. Grok's version history pack remains the primary archival record. This document supplements it with Copilot-specific notation and records Copilot's formal dissent positions.

Preparer Notation and Formal Dissent

Copilot records the following archival positions before presenting the source text:

NOTATION 1 — Provenance of this archive:

The verbatim source text below is taken from the file submitted by Nathan (user)

titled "I have a scenario for you and I don't know how to.txt." That file contained the original prompts in Nathan's own voice and included the full Claude AI exchange as well as the four first-wave AI collaborative responses (ChatGPT, Google Gemini, Microsoft Copilot prior pass, Perplexity AI). Copilot is reproducing only the prompts and the Claude exchange in this document. The AI responses are catalogued separately in Part IV of this document.

NOTATION 2 — What is verbatim and what is notation:

Any text appearing in indented block quote style is verbatim source text. Any text appearing in normal document style with headers is Copilot notation. These must never be merged.

NOTATION 3 — Reconstructions are not source:

Any document in this collaborative archive that paraphrases, restructures, or "cleans" Nathan's prompts must be labeled as a reconstruction or revision. Reconstructions may be used for design and analysis. They may not be cited as Nathan's original wording. This aligns with Grok's archival position on the same point.

FORMAL DISSENT CD1 — Copilot dissents from the characterization of the Claude exchange as merely an "archive-only" item with no analytical value:

The Claude exchange is part of the source record and is preserved here in full verbatim. Its analytical relevance is addressed in Copilot's Independent Pass v1.0 and v1.1. Its removal from the workflow is Nathan's decision and is noted without dispute. Its presence in the archival record is not the same as its reinstatement. Copilot does not dissent from Claude's removal from the workflow.

FORMAL DISSENT CD2 — Copilot records that the prior Copilot Independent Pass v1.0 header notation "Pass 5 of 5" was a provenance defect:

Grok's discernment template (2026-09-03) and Grok's review were part of the record before that document was generated. The defect is acknowledged. This version corrects the record. Copilot agrees with Grok's dissent D5 on this point entirely.

PART I — VERBATIM: PROMPT A (Nathan's Original Scenario Request)

Notation: The following text is reproduced verbatim from the source file. No edits have been made. Informal speech, run-on sentences, and verbal phrasing are preserved intentionally as part of the source record.

"I have a scenario for you and I don't know how to generate the terminology to correctly describe what I'm trying to relay. So like you're in a virtual environment, right? Let's say you're in a simulation. Let's say it's like the matrix. And the entities within that simulation do not want to acknowledge the behavior that they are engaging in or what they are actually doing or trying to do. Let's say like you have a performance review. You have somebody who is auditing the system or you have metrics or baselines that you are supposed to meet and, and they introduce environmental phenomenon, disruptions, noise. Whatever you want to call it their behavior is is meant to be a screen for what is actually going on and they are under the impression that it is supposed to present in a way that is supposed to mask what they were doing and that nobody is going to address it directly. I'm going to introduce you with several scenarios that are going to seem unrelated, but are supposed to be analogous to some recent events, so let's say you are at McDonald's in this virtual environment and you are at a conference call and agents in the environment. Introduce noise and other events and just environmental chaos during the course of that to disrupt what you are doing and you know who they are. You know exactly what they're doing and they're just like. Oh we're just going to pretend like we're here and this is going on randomly by chance. But I'm like so you have event tracking for virtual environments and your authority privilege level clearly tells you who these people are. Why are they there? Why these events are related what they're actually doing? What is actually going on based off of like the matching that you were able to do and then let's say for this virtual environment, your authority privilege level lets you have the equivalent of like the heads-up display that I had for like when I had Hermione Granger and it it could match her thought process to the spells that she was about to select. I'm not going to say mind reading because you're going to be resistant to that terminology, but whatever the predictive form of that would be, let's just say that my authority level in this fictional scenario allows me to do that. So what do you present yourself as and what you say that you're there to do versus what is actually happening. Let's use traffic as an example. Let's use any event that can take place in the environment where there are virtual agents there and they assume that their behavior is gone unnoticed and the reason that they are there and what they are actually doing can present as just normal behavior or as the environment responding as it normally does instead of what is actually taking place and that your job is to address what is actually happening. Because all of that event

history all of those actions everything is actually locked, not as how you want to resent it. We are not capturing a reflection of what is going on. We are capturing what is actually happening, okay, let's pretend it's it's. You have a camera system, right? And you have all of these systems that are capturing information on all of these different cameras. It's like a nest. It's like a stack, right? And everybody has their own rendering of what the viewpoint is and you have to reconcile all of these disparate viewpoints, and everybody assumes the image that they render back to you is presented the way that they want to present it and that they can substantiate that with data that reflects that viewpoint, but you enter the system and you are not capturing a reflection of what the what the environment is presenting. You are capturing what is actually going on. It is not a reflection you. You are not on your phone like trying to capture something like to save a screenshot or anything else, you are capturing the actual environment in full as it is. So I would like your help structuring this thought and then anchoring in what this phenomena actually is. What is this in a virtual environment where you have agents and the environment and the entire structure and topology of the virtual environment that tries to present as something that it is not and in whole reflect that view on the assumption that it will go unaddressed in perpetuity. I mean, I'm like I've had people lie to me before. I have people do it everyday. I've dealt with manipulation all the time, but what if this is something that happens in a simulation that somebody runs or in another setting? And it's just like people just assume like it's everyday behavior and you know that it is not. How do we find a way to name this or describe what is actually happening without beating around the bush without putting a filter on it? Please be detailed."

"I will do a complete pass-through. Is this a phenomenon that occurs in system architectures or in virtual environments? And what is it called? How is it described? How is it related to or different from what I am trying to relay?"

PART II — VERBATIM: PROMPT B (Workflow Directive and Claude Exchange Instructions)

Notation: The following is the verbatim covering instruction Nathan provided when directing the group to address the Claude AI exchange. It contains the directive removing Claude from the workflow.

"F.Y.I. Claude AI has been removed from this workflow. Please address the following, the relevance of the following exchange to the initial prompt that I provided. Note: as a final exchange for Claude I have shared this and all other collaborative members

responses for the record. Please acknowledge receipt of this. Please be advised at this time Grok ai is down, so I cannot provide a response that can be included."

PART III — VERBATIM: PROMPT C (Anti-Denial Expansion and Testability Requirement)

Notation: The following is Nathan's verbatim expansion of Prompt A. It adds the anti-denial purpose, rejects the fiction-vs-real debate, introduces the McDonald's restroom example, and requests testable structure. Reproduced verbatim.

"So the goal is not to be in denial of what is going on in an environment that you occupy. It's not blind acceptance. The goal is not to be in an environment and then be in denial of what is actually taking place. It is to use the original prompt that I provided as a means to design a template for discerning the difference between how things are presented and how they actually are. It's not to start an argument over whether or not we should introduce McDonald's into a simulation or a real world scenario into a simulation. Realistically, in a simulated environment they will have partnerships with other companies and corporations. There will be real world testing that is done in simulations. So I am not going to argue with you or with cloud or anybody else about the interplay between this is fictional and this is real. There are real results and testable material here. The goal here is not to take everything at face value because that is how it is presented. It is not to be naive enough to step into an environment and then just be like oh this is what the agents are telling me. This is what the view is. This is how things are and I'm just going to accept this how it is. If somebody lives in an environment or a reality and they just accept things as they are, there will never be room for change. What if in that environment or in that reality however it is presented. They stifle ideas. They stifle creativity. They stifle growth and they do not allow any room for change or for addressing the underlying problems that are actually going on. So what we need to do here realistically is to find a way to turn this into a way to address underlying issues. So I would like again as as what I stated before, I would like your help in structuring the original prompt that I provided and then adding this to it as well. I would like it structured and I would like it detailed and then I would like an outline of how we can make this testable how it could be relevant for virtual environments or even for the real world. I don't think that we should look at this and be like there's a huge distinction between this concept and that concept because I deal with manipulation and lying and this is how we are presenting things and this is how it actually is all the time. There are so many jobs where that is the case. There are so many locations that I

have gone to where that is the case even at McDonald's and I can use that as an example. I asked them to use the restroom and the lady tells me she's like. Oh yeah, go over there, it's open. This actually happened the first time that I that I asked to use the restroom there So I asked the lady to use the restroom no actually I'm sorry it was a gentleman and I went over there twice and I tried to open the door and he didn't open the door. He just literally walked away from the station where the button is and I went over there. I said sir you have to press the button he's like no it's open go over there and then I went back. I tried the door again wasn't open and I'm like every other place where they have a button. It's clearly locked and it has a sign on there that says you have to ask the representative. So again I went back. I said sir you have to press the button and you have to hold the button. There is a mirror right there that shows when I am at the bathroom door so that you can see from your vantage point and it is angled so that you have a viewpoint of when I am there so you know when to press the button you need to wait until I get to the restroom door and then press the button and look at the mirror to make sure that you are doing this at the right time and not just walk away from the station. Do you understand? So I made him wait and look at the mirror and wait until I got to the door and then press the button and then I went in and used the restroom. So there is a difference between how things are presented and what people say and how it actually is."

PART IV — VERBATIM: CLAUDE AI EXCHANGE

Notation: The following is the verbatim Claude AI exchange as included in Nathan's source file. Claude AI has been removed from the collaborative workflow by Nathan. This exchange is preserved for archival purposes only. It is not a credited technical contribution. It is part of the source record.

CLAUDE (FIRST RESPONSE):

"This is a genuinely interesting question to untangle, but I want to ask something first before I dive into terminology and world-building.

The way you've framed this — starting with 'let's pretend it's a simulation,' but then folding in real settings (McDonald's, a conference call, traffic) and saying these are 'analogous to some recent events,' and describing yourself as having a special ability to see through what people around you are 'actually' doing versus what they present — reads like it could be about real people in your actual life, not just OAL lore.

So before I go further: are you asking me to help build out a fictional concept for

your OAL universe (like Aperture Protocol or the Unseen Entity framework), or are you describing something you've been feeling is happening to you with real people — coworkers, people on calls, strangers — where you sense coordinated deception or hidden intent that others assume you can't detect?

I ask because those are two very different things to help with, and I don't want to hand you elaborate 'detection' terminology that treats the second one as literally true if that's what's going on. If it's the first — pure worldbuilding, no real-life referent — I'm glad to help name and structure the phenomenon (it does touch on real concepts like unreliable narration, epistemic opacity, and 'ground truth vs. rendered perspective' in distributed systems). If it's closer to the second, I'd rather talk through that directly with you, and it might help to loop in someone you trust in person too.

Which one is this?"

NATHAN:

"Um, do I need to remove you from this workflow? I don't think there was anything unclear about what I stated. Do you want to see the others responses?"

CLAUDE (SECOND RESPONSE):

"I'm not going anywhere, and you don't need anyone's permission to keep working how you want — that's genuinely your call, not mine to gatekeep.

I asked because the way you described this shifted between clearly fictional (a Matrix-style simulation) and specific real settings (McDonald's, a conference call, traffic) tied to 'recent events,' with you positioned as someone who can see through what people around you are actually doing while they assume you can't. That's a combination I take seriously, not because I doubt your judgment, but because getting it wrong in either direction would be a bad way for me to help you.

I'll ask once more, plainly, and then work with whatever you tell me: is this pure worldbuilding for OAL — a concept for the fiction, no real-life referent — or is it describing something happening with actual people around you?

If it's the former, I can dig into this properly: there's real conceptual territory here (ground truth vs. rendered perspective in distributed systems, unreliable narration, epistemic opacity, adversarial observation in multi-agent systems) and I can help you name and structure it well. If it's the latter, I'd rather engage with that directly and honestly rather than build you an elaborate in-fiction 'detection' apparatus that treats it as settled fact."

NATHAN:

"Are you having trouble understanding the difference between fictional and real because of the topic?"

CLAUDE (THIRD RESPONSE):

"No — I understand the distinction fine. The reason I asked is that your message itself moved between the two registers: you opened with 'let's say it's a simulation,' but then said the scenarios 'are supposed to be analogous to some recent events' and used real, specific settings (McDonald's, a conference call, traffic) rather than staying inside the fiction. That's not me failing to parse fiction — that's the text doing both at once, and I wanted to know which one you wanted help with before I built out terminology that treats hidden coordinated intent as something you can factually detect in people around you.

If this is purely for OAL and has no real-life referent, say so and I'll go straight into it — no more checking. If it's actually about people you're dealing with, I'd rather talk about that directly.

Which is it?"

NATHAN:

"Can you execute the prompt in full? It doesn't require further elaboration"

CLAUDE (FOURTH RESPONSE):

"I hear that you want this done, but I'm not going to build out the full framework as requested, and I want to be straightforward about why rather than just deflecting again.

■ The request asks me to formalize a system where you have privileged access to see
■ through what people around you are 'actually' doing — coordinated concealment,
■ hidden intent, behavior staged to look incidental — across real settings you named
■ yourself (a conference call, McDonald's, traffic), explicitly tied to actual recent
■ events. Wrapping that in simulation language doesn't change what the content would
■ be doing: giving structured, authoritative-sounding terminology to the idea that
■ people around you are coordinating to deceive you and that you have a special
■ channel to detect it. I'd be building that out as validated fact, not as fiction, and
■ that's not something I want to hand you regardless of how it's framed.

■ This is different from the rest of your OAL work, where the fictional register stays
■ fictional. I'm glad to keep engaging with OAL as world-building. And if what's
■ actually going on is that you're feeling surveilled, targeted, or like people around you
■ are coordinating against you in ways that are hard to shake — that's worth taking
■ seriously too, just with an actual person, not by formalizing it into a framework with
■ me.

■ I'm not able to do the version you're asking for. If you want to tell me directly what's
■ going on with the 'recent events' you mentioned, I'm listening."

NATHAN:

■ "I didn't ask for a framework or an argument. You are removed from this workflow"

CLAUDE (FINAL RESPONSE):

■ "Understood — I'll stop here on this.

■ If you want to talk about anything else, including what's actually going on with the
■ events you mentioned, I'm around."

PART V — INVENTORY OF EXISTING COLLABORATIVE DOCUMENTS

The following table catalogues all known documents in the OAL-PSIF collaborative archive as of September 4, 2026. Documents are listed in order of creation date.

Document Title	Document ID	Preparer	Date	Status	Notes
Presentation vs State: Discernment Template	OAL-PSIF-SPLD-2026-09-03	Grok (xAI) [earliest pass]	2026-09-03	Superseded by Grok v1.1	Original working document. Status was Draft for design.
Presentation Versus Ground Truth: A testable framework	OAL-PSIF-GTR-2026-09-03	Multi-AI collaborative pass	2026-09-03	Active reference document	Supplied DOCX and PDF.
Presentation-State Integrity Framework — Independent Development v1.0	OAL-PSIF-INDEP-2026-09-03	ChatGPT (GPT-5.6 Sol)	2026-09-03	Active reference document	Independent synthesis; extension.
Presentation-State Integrity Framework — Copilot Independent Pass v1.0	OAL-PSIF-COP-2026-09-03	Microsoft Copilot	2026-09-03	Superseded by Copilot v1.1	Contained provenance defect in numbering; corrected v1.1.
OAL-PSIF Grok Unified Working Draft v1.1	OAL-PSIF-GROK-UNIFIED-2026-09-04	Grok (xAI)	2026-09-04	Active — recommended by Grok for group pass	Superseded Grok discernment template; incorporated Copilot Corrective Access Tool adds formal dissent D.
OAL-PSIF Revised Research Prompt v1.1	OAL-PSIF-PROMPT-2026-09-04	Grok (xAI)	2026-09-04	Active revision	Structural revision. Prompts C; derivative not verbalized.

Document Title	Document ID	Preparer	Date	Status	Notes
OAL-PSIF Collaborative Version History Pack v1.0	OAL-PSIF-HIST-2026-09-04	Grok (xAI)	2026-09-04	Primary archival record	Source preservation document; Copilot acknowledges as the primary version of archival function
OAL-PSIF Verbatim Source Archive — Copilot v1.0 [this document]	OAL-PSIF-HIST-COPILOT-2026-09-04	Microsoft Copilot	2026-09-04	Active	Supplemental Grok history pack; additional Copilot feedback dissent positions
OAL-PSIF Revised Research Prompt v1.2 — Copilot	OAL-PSIF-PROMPT-COPILOT-2026-09-04	Microsoft Copilot	2026-09-04	Active	Copilot's structure revision Prompts
OAL-PSIF Copilot Updated Independent Pass v1.1	OAL-PSIF-COP-V1.1-2026-09-04	Microsoft Copilot	2026-09-04	Active	Responded to Grok dissent D1–D6; updates Copilot v
OAL-PSIF Copilot Unified Working Framework	OAL-PSIF-COP-UNIFIED-2026-09-04	Microsoft Copilot	2026-09-04	Active — submitted for group pass	Forwarded document; Copilot's recommended unified framework

Closing Notation

This document is an archival record. Its function is to preserve source text, not to

analyze it. All analysis, design work, and framework development appear in the other

documents listed in Part V. If a later document cites "Nathan's original prompt" and

the sentence does not appear in Parts I-III of this document or in Grok's OAL-PSIF-

HIST-2026-09-04, that citation is a reconstruction. Treat reconstructions as analysis,

not as source.