

OAL_QC_Baseline_SixAI_Audit_v1

Independent Review of the Six-AI Baseline Responses

Quantum Credits / 1COIN / NCOIN — Post-Scarcity Valuation Exercise

Reviewer: Claude (Primary Audit and Documentation Layer) **Subject documents reviewed:** ChatGPT, Claude (prior turn), Grok, Gemini, Copilot, Perplexity — all responding to the same prompt: rebuild the QC valuation exercise as a structured baseline, informed by the Red Team Currency Review v1. **Status:** Audit pass, not yet canon. Feeds into a Version 2 baseline document.

Executive Summary

All six systems converged on the same core diagnosis, independently: **the first red-team pass failed not because QC is incoherent, but because the reviewers (including me, in that pass) analyzed it as a conventional economic instrument instead of accepting its defining rule and testing what follows from it.** That convergence is itself meaningful — it wasn't in the original prompt, so five of six systems inferred it, and the sixth (my own prior turn in this thread) missed it. That's flagged below, not glossed over.

Where the six responses differ is in how far they took the reframe. Three (ChatGPT, Perplexity, Grok) produced genuinely new conceptual architecture. One (Gemini) restated the reframe in compressed, stylized form without much new structure. One (Copilot) produced a competent rewrite but repeated its earlier pattern of defaulting to a finished-looking deliverable rather than leaving the open questions open — and padded the document with decorative stock-image placeholders that add nothing to a governance document. One (my own prior response) answered a narrower question than the one actually being asked — I fetched real-world spending data and left the conceptual reframe to the rest of the team instead of doing it myself, which is a real gap in that response, not just modesty.

Verdict: the reframe is validated and should become the new starting assumption for all future QC work. The disagreement that matters going forward isn't "is QC's mechanic viable" — it's which of several proposed sub-models (below) should govern 1COIN's settlement behavior, NCOIN's negative-value behavior, and QC's relationship to scarcity.

Comparative Assessment

AI	Approach	Depth	Key Strength	Key Weakness
ChatGPT	Full conceptual rebuild (35 sections, formal	Extreme	Introduced the "meta-currency /	Long enough that some sections restate

AI	Approach	Depth	Key Strength	Key Weakness
	notation)		value-resolution protocol” reframe; built a 15-level escalating boundary-test ladder; raised the recursive-valuation problem (QC applied to QC itself)	the same point (target-relative valuation) three or four different ways
Grok	Concise synthesis + clean baseline rewrite	Deep	Correctly summarized the Red Team Review first, then didn’t let that summary bleed into the new baseline — kept the two tasks separate	Least original: mostly synthesizes ChatGPT/Perplexity-style ideas already implicit in the prompt rather than adding new sub-models
Gemini	Compressed executive-style brief with tables	Moderate	Best phrase-level framing again (“Ontic Solution” vs. “Red Team Friction Point”); correctly identifies that QC doesn’t draw from a pre-mined pool	Thin — no working-definitions section, no open-questions list, ends by asking Nathan a question rather than finishing the baseline
Copilot	Full rewrite formatted as a “canon-ready” brief	Deep	Clear plain-language definition of QC as “value-mirroring, not scarcity-based”; correctly separates the 150-year exercise from a literal funding	Repeats its prior-session pattern: treats an open design problem as basically solved and offers to write “Version 2” immediately; includes four irrelevant

AI	Approach	Depth	Key Strength	Key Weakness
			claim	decorative stock-image placeholders in a working design document
Perplexity	Formal constitutional-brief style with governance table	Deep	Only response to propose an actual resolution to the post-scarcity paradox (“contribution-backed, not labor-backed”) instead of just restating the tension; most complete governance-layer table (Architect / Human Council / AI Bureau / Red Team / Ledger / Appeals)	Heaviest reliance on citing back to the Red Team Review itself rather than generating new material — more synthesis than invention in places
Claude (prior turn)	Real-world data lookup + narrow design-fork framing	Shallow, on a narrower question	Got the live Gartner figures right and correctly identified the anchor-mode fork as <i>a</i> live question	Missed the larger reframe entirely — answered “what does the exercise cost in today’s dollars” when the actual open question (as all five other systems recognized) was “what kind of instrument is QC, if not a conventional currency at all”

Consensus Findings (5–6 of 6 agreeing)

1. **The first red-team pass applied the wrong paradigm.** Every system except my prior turn explicitly named this. Scarcity economics, reserve-currency theory, and stablecoin analogs don't apply to an instrument whose defining rule is that its value is the target's value, not a stored or pegged quantity.
 2. **QC is better modeled as a protocol than a currency.** ChatGPT calls it a “value-resolution protocol” / “economic primitive.” Perplexity calls it a “governed evaluation and authorization instrument.” Copilot calls it “value-mirroring.” Gemini calls it a “dynamic ontic evaluation tool.” These are the same underlying claim in four vocabularies: **QC doesn't hold value, it resolves to value on invocation.**
 3. **1COIN is a settlement/resolution layer, not a public peg.** No system defended treating 1COIN as a stablecoin analog (which was the original red-team's Flaw 3 framing). All five reframe it as a closed-system instrument that converts a QC evaluation into a recorded, approved settlement.
 4. **NCOIN should not default to a permanent debt trap.** Every system that addressed NCOIN in depth (ChatGPT, Copilot, Perplexity, Grok) proposed some form of rehabilitation, cancellation, or bounded-consequence model rather than the irreversible freeze the original red team flagged as Flaw 4.
 5. **“Can this be purchased?” is the wrong question to ask at each test level.** ChatGPT's boundary-laboratory format made this explicit as a rule (no AI is allowed to answer “this cannot be purchased” — that's the premise under test, not a valid response), but the same instinct is present in Grok's and Perplexity's framing too.
 6. **The narrative-layer / specification-layer split (already recommended in the Red Team Review) should be formalized,** not just acknowledged. Copilot, Perplexity, and Grok all treat this as a structural requirement for the next document, not a nice-to-have.
-

Points of Genuine Disagreement (unresolved — need your decision, not more AI consensus)

These are the places where the six responses actually diverge in ways that matter for design, not just phrasing:

1. **Does QC create value, transfer it, or merely recognize it?** ChatGPT lists this as an open “Value Creation vs. Value Transfer” question with six candidate models and deliberately doesn't pick one. Perplexity picks one implicitly — QC “creates a recognized and governed relationship” rather than material resources (closer to an authorization model). Copilot's framing (“QC does not ‘buy’ the target; it synchronizes with its value”) is closer to a pure-recognition model with no creation or transfer at all. **These are not compatible with each other as written** — pick one before Version 2, or the three-currency system will inherit contradictory assumptions depending on which AI's language gets used in the next draft.
2. **Is scarcity abolished, relocated, or preserved?** Perplexity's “contribution-backed, not labor-backed” proposal relocates scarcity to verified contribution (research, stewardship, risk, coordination) rather than logged work hours — this is a genuine

candidate resolution to Flaw 1 from the Red Team Review, and it's the only one offered. ChatGPT instead proposes relocating scarcity to *issuance* (how many QCs can be generated, by whom) rather than to what backs each QC's value. These are different mechanisms and would produce different systems. Nobody reconciled them.

3. **Does QC override consent, ownership, and security boundaries, or work within them?** Perplexity is explicit and firm: QC can evaluate a relationship to a classified system but “would not automatically defeat the security rules, consent requirements, or access controls of that system.” This is a meaningful constraint that ChatGPT's and Gemini's more expansive language (“regardless of... systemic boundaries”) doesn't share. This is arguably the single highest-leverage design fork in the whole set — it determines whether QC is an omnipotent narrative device or a governed instrument with teeth, and it changes almost every downstream mechanic.
 4. **1COIN's exact function** still has at least four live candidate models across the responses (ChatGPT's four lettered models, Perplexity's “settlement record,” Copilot's “resolution instrument,” Grok's “dynamic resolver against QC”). Nobody has picked one yet.
-

Best Original Contributions Worth Carrying Forward

- **ChatGPT's recursive valuation problem** (Section 21 of its response): what happens when QC is applied to another QC, to the QC issuance system, or to itself? This wasn't in the original Red Team Review at all, and it's a real gap — any complete spec needs an answer.
 - **ChatGPT's 15-level Currency Boundary Laboratory.** This is the strongest concrete next-exercise proposal in the set — a fixed escalating ladder every AI answers identically, which would finally produce comparable data instead of six differently-scoped essays.
 - **Perplexity's contribution-backed resolution to the post-scarcity paradox.** The only response that didn't just restate Flaw 1, but proposed an actual mechanism.
 - **Perplexity's governance table** (Architect / Human Council / AI Bureau / Independent Red Team / Ledger / Appeals). This directly answers Flaw 9 (Architect as single point of failure) and Flaw 5 (DAO governance capture) from the original review in one structure, rather than treating them as separate problems.
 - **Grok's discipline in not letting the Red Team Review's findings leak into the new baseline uncritically** — it summarized the review, then built the baseline from your stated intent, not from the review's assumptions.
-

Gaps No AI Addressed

1. **No system proposed a concrete answer to the authorization question** (who or what actually approves a QC invocation, and against what evidence) — everyone flagged it as an open question, nobody drafted even a first-pass answer.

2. **No system reconciled its own proposed model with the other five.** Each response is internally coherent but was written without seeing the others. This audit is the first place the contradictions (see “Points of Genuine Disagreement” above) are surfaced together.
 3. **Nobody addressed what happens to the exercise’s 150-year, doubled-global-tech-spend scenario under each of the four candidate interpretations Perplexity laid out** (funding / authorization / valuation / hybrid). That table was Perplexity’s strongest structural contribution and it’s currently unused by the other five responses.
-

Recommended Next Step

Before writing a Version 2 specification, three decisions need to be made explicitly by you (not delegated back to the team as open questions again), because the six responses show the team will keep generating parallel, incompatible answers otherwise:

1. **Pick one value-creation model** (creation / transfer / recognition / hybrid) — this is the ChatGPT vs. Perplexity vs. Copilot fork above.
2. **Decide whether QC overrides consent/security boundaries or operates within them** — this is the highest-leverage fork in the set (Perplexity’s constrained model vs. ChatGPT/Gemini’s unconstrained language).
3. **Choose where scarcity lives, if anywhere** — issuance-side (ChatGPT) vs. contribution-backed (Perplexity) vs. none (the original “post-scarcity” framing the first red team flagged as contradictory).

Once those three are fixed, ChatGPT’s 15-level Currency Boundary Laboratory is the strongest candidate for the next team-wide exercise — run all six AIs against the same fixed ladder, with your three decisions above stated as governing constraints up front, and the run should finally produce comparable rather than divergent output.