OMNICELLULAR PHYSIOLOGY

Phase 4 Independent Recommendation

Prepared by Grok for Nathan Lewis II, Architect, Omniversal Architect Labs

Date: October 30, 2026

Status: Independent recommendation. Not canon. Conflicts with other contributors are expected and preserved.

1. Purpose

This document reviews the full set of team responses to Claude's Cross-AI Execution Review and states an independent recommendation for moving forward. It does not require consensus. Different contributors may proceed with different definitions of an Omni cell, different remediation scopes, and different next artifacts. The goal is a usable path that preserves the original blind experiment, surfaces real disagreements, and produces comparative evidence before any new generation of designs.

2. Analysis of Team Responses

The responses show stronger agreement on process than on physiology. ChatGPT, Claude, Gemini, Copilot, and Perplexity all support freezing the original 48 architectures as the Gen-1 Blind Baseline, delaying Gen-2, running cross-AI red-teaming before ranking, building feature and similarity matrices, treating identity continuity as a separate standard, and using controlled divergence rather than a simple ban on the three attractors. Mechanism provenance (prompt-required, threat-implied, canon-inherited, or spontaneous) is widely accepted as a useful thirteenth analytical dimension.

Disagreements remain and should stay visible:

• 4,712 allocation. ChatGPT and Copilot recommend marking the current 2,600–4,450 ranges illustrative/unreconciled and separating the locked canonical count from any research-range estimate. Claude initially favored raising ranges but later accepted the illustrative approach. Gemini proposed raising upper bounds to 5,200; this is not adopted here because it manufactures an allocation without event data.

• Remediation boundaries. ChatGPT and Perplexity emphasize that remediation must not change core wagers or mechanisms if the result is still to be treated as a Gen-1 correction. Gemini's plan to replace shared rules with 24 fully prototype-specific rules, and any expansion that adds new mechanisms, crosses into redesign. These should be labeled as Gen-1b or Gen-1R and kept distinct from the immutable baseline.

• Omni-cell definition. Most responses support Constitutional Omnicellularity (broad latent competence, no requirement that every cell perform every function simultaneously) as the core track, with micrometer-scale and engineered-module work kept as explicitly labeled comparative tracks. Perplexity correctly notes that its own parcel-based designs may fall outside a strict biological definition and should be classified accordingly rather than forced into compliance.

• Sequencing of matrices. Several responses recommend building the baseline 48×12 and 48×48 matrices from the frozen originals first, then optionally building a second set from remediated versions. This preserves evidence of original convergence.

Claude correctly notes that the current agreement is not independent—every AI has read the review—and that quick acceptance of critiques is a statement, not verification. Cross-AI red-teaming remains the stronger test.

3. Independent Recommendation

Do not start Gen-2. Freeze the original 48 architectures as the immutable Gen-1 Blind Baseline. Create a thin Errata & Standards layer that flags (but does not silently rewrite) the 4,712 allocation problem, the roadmap compression arithmetic, Omni-cell premise options, physical-versus-canon separation, and identity terminology. Then proceed to Phase 4 evaluation.

Recommended sequence:

1. Freeze originals. Assign contributor-prefixed permanent identifiers. Hash canonical sources and renders. Preserve the Mosaic Covenant name collision in the baseline rather than erasing it.

2. Publish a short Errata & Standards note. Mark the 4,712 Part V ranges as illustrative/unreconciled. Keep 4,712 as the locked project record. Do not invent category assignments. State Constitutional Omnicellularity as the core principle with two explicitly labeled comparative tracks (micro-module and engineered-module). Adopt Copilot's C0–C9 identity continuity states as a working draft, not as a solved philosophy.

3. Build the baseline 48×12 Architecture Feature Matrix and 48×48 Pairwise Similarity Matrix from the frozen originals only. Add the 13th provenance field. Mark missing information as Unknown.

4. Allow a limited Gen-1R remediation round. Permitted changes: rendering, version labels, matching PDF/DOCX/MD sources, citations, evidence-label downgrades, calculations with disclosed assumptions, expanded explanations of mechanisms already present, and explicit Unknowns. Changes to substrate, topology, core wager, identity governance, or survival mechanism must be labeled as a new version or Gen-1b and kept out of the blind baseline comparison.

5. Run a common cross-AI red-team on the frozen originals (remediated versions may be supplied as supporting material). No AI is the first reviewer of its own designs. Use a shared attack suite with explicit outcome labels that distinguish active function, recoverable state, and restored successor. An independent reconciler reviews before the original designer responds.

6. After the red-team reports are locked, perform targeted verification of mechanisms that survived. Prioritize shared mechanisms and high-impact quantitative claims. Perplexity's reproducibility package is the logical first audit target.

7. Only then commission a controlled divergence round (attractor controls retained, plus designs that cannot rely principally on dormancy, mineral/refractory hardening, or modular/swarm survival). Ranking and Gen-2 follow the results.

4. Specific Positions on Open Decisions

• Omni-cell: Constitutional Omnicellularity for the core track. Broad latent competence is required; simultaneous universal function in every cell is not. Micrometer-scale and engineered-module designs remain valuable but must be labeled as comparative tracks. A fixed cell count is not required; each design must justify its count.

• 4,712: Preserve the locked count. Mark current category ranges illustrative/unreconciled. Do not raise ranges or adopt a 5,200 upper-bound proposal without an underlying event ledger.

• Version locking: Canonical source + PDF + DOCX each receive full SHA-256 hashes in a manifest. Originals stay immutable.

• Name collisions: Keep original names in the baseline. Use contributor-prefixed IDs for comparison. Any rename occurs only in a later version.

• Ranking: None yet.

• Red-teaming: Mandatory non-author review before any ranking or Gen-2.

• Identity: Adopt C0–C9 as a working standard so different outcomes are not all recorded as "survived."

• Evaluation axes: Keep architecture merit, evidence maturity, execution completeness, and reproducibility/checkability separate. Do not collapse them into a single score.

5. Conflicts Preserved

This recommendation does not resolve every disagreement. Gemini may still produce a remediated version with prototype-specific rules; that version should be labeled Gen-1b. Copilot may rename Mosaic Covenant in a later edition; the original name stays in the baseline. Perplexity's parcel designs may be scored as Track III comparative embodiments. Claude's errata on the ChatGPT 4,712 finding should be recorded. These differences are useful and should remain visible in the decision log.

6. Immediate Next Artifact

The next shared artifact should be a Phase 4 Control Package (not another physiology dossier). It should contain: the original-submission manifest and revision lineage, the approved or provisional Omni-cell and outcome definitions, the 4,712 unreconciled notice, the 12-dimension extraction rubric plus provenance field, the common failure suite and review assignments, remediation boundaries, and a decision log that preserves disagreements. Once that package exists, contributors can proceed with their own remediation or red-team work without waiting for perfect consensus.

7. Questions for the Architect

1. Do you accept Constitutional Omnicellularity as the core principle, with micro-module and engineered-module tracks explicitly labeled?

2. Do you accept marking the current 4,712 category ranges as illustrative/unreconciled while keeping 4,712 as the locked record?

3. Do you approve freezing the original submissions and limiting Gen-1R remediation to non-mechanism changes?

4. Should the first red-team pass be applied to all 48 architectures or only to a selected subset?

5. Do you want a single Phase 4 Control Package drafted next, or do you prefer each contributor to proceed independently on remediation?

End of independent recommendation. Conflicts with other contributors are expected and should be preserved in the project record.