OMNICELLULAR PHYSIOLOGY

Independent Eight-Prototype Architecture Research Directive

Grok v2.0 — Updated After Cross-AI Review

Prepared for: Nathan Lewis II — Architect, Omniversal Architect Labs (OAL)

Date: October 28, 2026 | Status: Updated Independent Directive for Execution

Purpose of This Document

This document is Grok's updated independent research directive for designing eight peer Omnicellular Physiology prototypes. It incorporates the strongest methodological, security, and evaluation insights from the Round-One prompt engineering review conducted by the OAL Multi-AI Coordination Team (Perplexity, Claude, ChatGPT, Copilot, Gemini, and prior Grok contributions) while deliberately leaving the actual architectural answers open for independent invention. It exists to produce genuinely differentiated architectures rather than cosmetic variants of a single shared design.

How the Topic Arose

The original request asked each collaborating AI to independently design eight Omnicellular Physiology prototypes capable of long-distance linking, extended durability, and first-prototype resilience against the 4,712 Extinction Attempt Classification Framework (Mild through Severe). Designs were to assume hostile interest from individuals, governments, and corporations, nonstandard and uninhabitable operating environments, and a vested interest in continued existence. Mandatory elements included respiration rules, telomere or longevity mechanisms, thermal regulation, conversion and conservation rules, three observation rules, three genetic-security rules, and three laws from which the physiology may depart.

The first round produced six distinct re-engineered prompts. Cross-review revealed productive disagreement on peer versus progressive structure, shared versus independently chosen rules, evidence standards, and the boundary between physical and canon-speculative tracks. This v2.0 resolves those points for Grok's execution path.

Mission Statement

Design eight independently derived, peer Omnicellular Physiology architectures that can survive, operate, and remain recoverable in environments that would permanently extinguish standard human biology, while resisting capture, reverse-engineering, and coercive testing, and while maintaining the capacity to link with other instances across vast distances.

What This Directive Aims to Accomplish

- Produce eight peer (not successive) Gen-1 prototypes that differ in kind.
- Require each architecture to clear Mild-to-Severe resilience across all ten extinction taxonomy categories, with an explicit pathway for Extinction-Class.
- Enforce honest energy, heat, and matter accounting on the physical track.
- Force independent selection of observation rules, genetic-security rules, and biological law exemptions.

- Address identity continuity, single-cell theft risk, and monoculture failure.
- Link each prototype to the 15-Year Omnicellular Architecture simulation roadmap via a simulation hook.
- Include a self-attack pass in which some attacks succeed.

Core Locked Decisions for This Version

The following decisions are locked for Grok's execution of this directive:

- Eight prototypes are peer alternatives evaluated at the same Gen-1 bar, not an evolutionary ladder.
- The three observation rules, three genetic-security rules, and three biological departures are chosen independently by the designing agent.
- Fundamental physics remains binding on the physical research track. Conservation-law violations, zero-latency links, ontological mechanisms, and similar ideas must be labeled and confined to a separate canon-speculative track.
- Every mechanism receives an evidence label: Established, Engineering Extrapolation, Speculative, or Canon-only.
- Heterogeneity and compartmentalization plans are required to mitigate monoculture and single-cell theft risk.
- No AI sees another's completed prototypes until its own set is version-locked.

Operational Definition of Omnicellular

Before designing, define what an omni-cell is in your architectures. Distinguish: (1) every cell contains the information needed for many functions, (2) every cell can assume many functions under appropriate conditions, and (3) every cell performs every function simultaneously. These are not equivalent. Address volume, reaction incompatibility, resource competition, heat, and maintenance burden if simultaneous multifunctionality is claimed. Explicitly address how universal capability is reconciled with the risk that theft of a single cell could compromise the organism.

Mandatory Rule Sets

1. Cellular Respiration Rules

Define primary and fallback energy-extraction pathways, terminal electron acceptors, energy currency, byproduct handling, behavior under anoxia or argon displacement, and detectable respiratory signatures. Include numeric estimates where possible (baseline and peak power, reserve duration on zero input) and label their evidence class.

2. Telomeres or Equivalent Longevity Mechanism

Define end-protection, renewal rules, division limits or their absence, error correction, protection against tampering, and behavior of stolen or cultured material. State design life under nominal and hostile conditions.

3. Thermal Regulation

Define setpoints, heat generation and rejection (especially radiation-only in vacuum and sealed argon chambers), survival modes for extreme heat and cold, and thermal signature control. Explicitly address the heat-rejection problem when convection is unavailable.

4. Conversion and Conservation Rules

Define permitted conversions, what is conserved, efficiency floors, waste handling, scavenging protocols, and storage. Any apparent violation of conservation must be one of the three chosen exemptions and must be accounted for on the appropriate track.

Independently Chosen Rules

Choose exactly three observation rules (covering outward sensing, inward monitoring, and being observed) and exactly three genetic-security rules. For each, state the threat addressed, mechanism, trigger, response, cost, failure mode, key location if any, duress behavior, and what a thief or captor obtains. Choose exactly three ordinary biological laws from which the physiology is exempt; provide a precise replacement law, justification, price, and outsider effect for each. At least one should address the genetic alphabet (nucleobases). Name at least three laws the physiology still obeys.

Linking and Durability

Design linking at local, planetary, and interplanetary scales. Specify medium, shared content, latency, authentication, partition and rejoin behavior, quarantine of captured nodes, energy cost, detectability, and blackout operation. State design life, repair strategy, damage budget, redundancy, dormancy options, and end-of-life behavior.

Prototype Requirements

Produce eight peer prototypes (P1–P8). For each, state a distinct core wager and what is sacrificed. Provide cellular matrix description, genetic layer, the four rule sets, three observation rules, three security rules, three exemptions plus kept laws, linking protocol, durability metrics, resilience profile (Mild–Severe for all ten categories plus Extinction-Class pathway), captivity and compound-event notes, simulation hook linking to multi-scale abstraction on the 15-Year Roadmap, and major weaknesses. After the eight, run a self-attack pass in which at least three attacks succeed, then rank and select.

Next Step Recommendations

- Each collaborating AI should execute its own updated independent directive without viewing other agents' prototype sets.
- After version-locking, perform controlled cross-comparison focused on architectural fingerprints rather than surface naming.
- Preserve the Round-One prompt archive as historical record.
- Architect decision register should lock any remaining open policy questions (e.g., non-destructive defense constraints) before final portfolio selection.

Questions, Comments, and Concerns

Concern: Extreme output volume may tempt compression of required fields. Staging (core wagers and rule sets first) is permitted if full single-pass delivery risks incomplete specifications.

Question for Architect: Should non-destructive anti-capture mechanisms be mandatory on the physical track, or may destructive options remain available under canon labeling?

Comment: The strongest protection against convergence remains independent derivation followed by later fingerprinting, not pre-standardization of the architectural answers themselves.

Appendix: Visualizations, Tables, and Reference Charts

All tables and structured visualizations appear in this section.

Table 1. Evidence Classes

Label	Meaning
Established	Supported by relevant empirical evidence within stated conditions
Engineering Extrapolation	Extends known mechanisms but lacks demonstration at proposed scale
Speculative	Requires unverified mechanism or substantial scientific breakthrough
Canon-only	Defined for the fictional setting; not claimed as scientific evidence

Table 2. Gen-1 Resilience Bar

Tier	Requirement	Must Report
Mild	No meaningful loss of function	Result and cost
Moderate	Full function or automatic recovery	Mechanism and recovery time
Severe	Survive and recover unaided	Recovery time, cost, residual damage
Extinction-Class	Pathway only	Pathway, needs, confidence

Table 3. Ten Threat Categories (Reference)

ENV-T Thermal | ENV-A Atmospheric | ENV-R Radiation | BIO Biological | MECH Mechanical | DEP Deprivation | ENE Energetic | INFO Informational/Cognitive | ONT Ontological | CIV Civilizational/Macro-Scale

All ten must be addressed at Mild through Severe for every prototype.

End of Grok v2.0 Independent Research Directive.

This directive is ready for independent execution. Prototype sets remain independent until Architect review.