

OPEN WINDOW DOCTRINE

Decision Log & Canonical Skeleton

v0.2 • Single Source of Truth • 2026-08-10

Status	WORKING DRAFT — Awaiting Phase 0 decisions Not ONESTAMPED
Purpose	Canonical reference for all subsequent work. Every future contribution must reference this skeleton.
Supersedes	All prior competing “definitive” merges (ChatGPT v0.1, Claude v2.0, Grok Protocol 0.1)

1. Phase 0 — Governance Decisions

These decisions must be locked before any further scenario expansion or second-iteration drafting. Open items are marked [OPEN]. Proposed defaults are offered so the skeleton remains usable immediately.

D1. ONESTAMP Authority

Proposed default: Originator (Nathan Lewis) holds final ONESTAMP authority. Multi-AI review is mandatory before any item is submitted for stamp. Disagreement protocol: dissenting AI must supply a written counter-proposal with explicit falsification criteria; originator decides within one review cycle.

[OPEN] Confirm or amend.

D2. Manifestation Register (for first run)

Proposed default for MVP only: Theory-agnostic. All markers are benign informational, environmental, or state-change proxies (color drift, temperature variable, semantic token, synchronization error). No biological, radiological, or physical-harm manifestation is permitted in the first executable run. Full canon decision deferred until after MVP data exists.

[OPEN] Confirm or amend.

D3. Standing relative to existing OAL constructs

Proposed default: Open Window Doctrine remains a free-standing research framework for the Net Defense Exercise series. Mapping onto Simulation Containment, living media, or DOCUMENT 771-ALPHA is deferred until after the first clean data set.

[OPEN] Confirm or amend.

D4. Version control & artifact lifecycle

This skeleton becomes the single living document. All prior parallel “v0.1 / v1.0 / v2.0 definitive” files are retired. Future changes are logged in a short changelog at the end of this document. Superseded content is moved to an archive folder, not deleted.

[OPEN] Confirm.

2. Authoritative Layer Architecture

Reconciled from ChatGPT L0–L4, Grok Layer 0–N + Window Objects, and Claude hybrid buffer language.

Layer	Function
L0	Oversight / Root Observer — Authoritative simulation ledger, immutable experiment log, snapshots, emergency rollback, researcher telemetry. No participant write access.
L1	Primary inhabited environment — Four-bedroom residential simulation with mixed-awareness agents and ordinary devices (Window Objects).
L2	Media / application environments — Content spaces, feeds, recordings, synthetic social platforms reachable through device windows.
L3	Adjacent household / community environments — Additional residences and shared services (activated only after local containment is proven).
L4	Sacrificial Interface Airlock — Translation / quarantine layer. Unknown cross-environment state must pass observe → translate → normalize → compare → quarantine before admission to protected layers.

Window Object properties (mandatory on every device): aperture state, directionality, residual process residual, entity affinity, local clock, bi-directional event log. Black-screen is a first-class state, not an absence of state.

3. Metric Dictionary (Reconciled)

Single authoritative set. Prior overlapping names are retired.

Metric	Definition
FCR	False Closure Rate — Share of trials where operators believed the interface closed while telemetry showed a remaining pathway.

ECR	Event Closure Reconciliation — % of events with matching open/close records on both sides + no residual state.
CTC	Carrier Transition Count — Number of medium/host changes before containment.
TTE/TTD/TTA/TTC	Time to Exposure / Detection / Attribution / Containment (authoritative simulation time).
TER	Temporal Error Residual — Difference between reconciled native-event ordering and authoritative ordering after legitimate clock transforms.
RTR	Root Trace Reliability — Confidence that ≥ 3 independent evidence classes converge on the same upstream privileged dependency.
Awareness Conversion	% of unaware agents who shift to semi-aware or aware after repeated anomalies.

4. Definitions Ledger

These terms must be declared per trial. Ambiguous use is not permitted.

- **Infection** — Any unwanted cross-aperture state transfer (informational, environmental, residual-process, or entity-class). Exact class declared per trial.
- **Exposure** — First verified transfer of anomalous state into a protected environment or agent.
- **Contamination** — Persistent presence of anomalous state after the originating window has been nominally closed.
- **Persistence** — Anomalous state that survives removal of its original host or carrier.
- **Recovery** — Verified return to baseline with no residual pathway (ECR = 100 % for the event set).
- **Residual state** — Any post-closure activity, process, or marker still detectable by independent telemetry.
- **Carrier transition** — Change of medium or host (device → environment → agent → message → secondary device).
- **Band-Aid closure** — Visible symptom removed while residual pathway or privileged dependency remains active.
- **Verified total event closure** — Matching open/close records on both sides of the aperture + independent confirmation of zero residual state.
- **Native time** — Clock, calendar, epoch, and tick rate exactly as reported by the local environment; never normalized before analysis.

5. MVP Scenario Set (First Executable Run)

Four scenarios only. Session structure: 2-day / ~12 active hours + overnight persistence window. All markers are benign. Null hypotheses are pre-registered.

S1 — Unattended Aperture (Black-Screen Residual)

Setup: Device left in black-screen residual state vs. fully powered-down control in an empty room. Null hypothesis: Black-screen residual state produces no measurable difference in residual activity or environmental drift compared with the powered-down control. Primary metrics: FCR, ECR, residual detection rate.

S2 — Content-State Bleed

Setup: Subject consumes media carrying a known benign sensory property (e.g., mapped temperature variable) vs. neutral-content control. Null hypothesis: No correlation exists between rendered content property and local environmental or physiological change beyond chance. Primary metrics: TTE, TTD, environmental tracking correlation.

S3 — Aware vs. Unaware Comparison

Setup: Matched agents across four classes (unaware / semi-aware / aware / already-contaminated) under identical exposure conditions. Null hypothesis: Awareness class produces no statistically detectable difference in exposure rate, containment behavior, or anomaly frequency. Primary metrics: Awareness conversion rate, TTC, behavioral divergence scores.

S4 — Total Event Closure (BTEC)

Setup: Controlled open events are closed independently from each side; reconciliation is forced. Null hypothesis: All events achieve matching open/close records and zero residual state when both sides attempt closure. Primary metrics: ECR, FCR, residual state after claimed closure.

Deferred to later iteration (explicitly not in MVP): carrier-switching cascade, full time-divergence mapping, multi-aperture household, root-reachability probes, infection-as-probe system mapping.

6. How to Use This Skeleton

- Every future AI contribution or human edit must reference the layer model, metric dictionary, and definitions in this document.
- No new scenario may be added until the four MVP scenarios have clean data and the decision gates in Section 1 are closed.
- Time-discrepancy and root-level tests remain sequenced after MVP, exactly as required by the dependence on trusted baseline data.
- When Phase 0 decisions are locked, update the [OPEN] markers and raise status to “Ready for MVP execution.”

End of Canonical Skeleton v0.2 — Ready for Phase 0 decisions.

Synthesized 2026-08-10 from multi-AI review cycle. This document supersedes all prior competing merges.