
PRESENTATION-STATE INTEGRITY FRAMEWORK

Copilot Independent Pass — Collaborative Record v1.0

Prepared by: Microsoft Copilot (Independent Pass)

Date: September 3, 2026

Classification: Conceptual Research and Design Framework

Collaborative Workflow: Multi-AI Framework Development — Pass 5 of 5

Prior Contributors: GPT-5.6 Sol (ChatGPT), Google Gemini, Microsoft Copilot (Prior Pass), Perplexity AI

Status Notice:

This document is a conceptual research and design framework. It is not a finding about any specific person, organization, or real-world incident. All examples are illustrative of framework concepts only. This document is submitted as an independent pass-through in a multi-AI collaborative framework development effort and extends — without replacing — the prior team's conceptual contributions.

DOCUMENT PURPOSE

Purpose of This Independent Pass

This pass accepts, consolidates, and extends the conceptual work established by the collaborative team. The core terminology — **Systemic Presentation-Layer Deception (SPLD)** and **Provenance-Based Adversarial Observability (PBAO)** — is carried forward without modification. The collaborative team's prior passes

established the vocabulary, layered architecture, detection taxonomy, and audit methodology with rigor and depth. This document does not revisit that ground except to consolidate it cleanly.

This pass's independent contribution focuses on three areas that were underemphasized or absent in prior passes:

1. **The Corrective Access Problem:** Detection of SPLD is necessary but not sufficient. What happens after the gap between presentation and state is established? Prior passes named the phenomenon excellently but did not fully address the question: what does a privileged observer *do* with what they know, and what barriers exist to actually correcting the underlying condition?
2. **Presentational Inertia:** A structural feature of environments where SPLD is operating — the environment actively resists acknowledgment and correction because correction requires the environment to admit what it was doing. This is not a psychological concept; it is a systems behavior.
3. **Narrative Sovereignty:** The contested claim over which account of events will be treated as the operative record. The deceptive presentation is not merely a cover story; it is a bid for narrative sovereignty — the right to define what happened. The auditor's provenance-based reconstruction is a counter-bid for the same sovereignty.

Core Thesis of This Pass

The collaborative team built an excellent framework for naming and detecting SPLD. This pass contributes the second half of the framework: what the observer does with what they know, and how they protect that knowledge from re-absorption into the presentation layer.

Table of Contents

§ 1	Conceptual Foundation — What the Collaborative Team Established
§ 1.1	Systemic Presentation-Layer Deception (SPLD)
§ 1.2	Provenance-Based Adversarial Observability (PBAO)
§ 1.3	The Four-Test Decomposition
§ 1.4	The Five-Layer Architecture
§ 2	Independent Contribution — The Corrective Access Problem
§ 2.1	Detection vs. Correction Are Not the Same Thing
§ 2.2	Forms of the Corrective Access Problem
§ 2.3	The Corrective Access Test
§ 3	Independent Contribution — Presentational Inertia
§ 3.1	Definition
§ 3.2	Why Environments Resist Correction
§ 3.3	Measuring Presentational Inertia (0-4 Scale)
§ 4	Independent Contribution — Narrative Sovereignty
§ 4.1	Definition
§ 4.2	How Narrative Sovereignty Is Exercised
§ 4.3	The Auditor's Claim on Narrative Sovereignty
§ 5	The Assertion Boundary — A Design Requirement
§ 5.1	The Four-Tier Evidence Model
§ 5.2	The Assertion Boundary Rule
§ 5.3	Why This Matters in Practice

§ 1	Conceptual Foundation — What the Collaborative Team Established
§ 6	Integrated Operational Model — The SPLD Audit Cycle
§ 7	Application Notes — Virtual and Simulated Environments
§ 8	Application Notes — Real-World and Hybrid Environments
§ 9	Failure Modes and Safeguards
§ 10	Collaborative Record — AI Team Exchange Summary
App. A	Terminology Consolidation Table
App. B	The Restroom Access Control Worked Example
—	Final Synthesis

Section 1: Conceptual Foundation — What the Collaborative Team Established

The collaborative team's prior passes produced a conceptual architecture of significant depth and coherence. This section consolidates the core terminology and structures in original language, preserving meaning precisely while establishing the shared vocabulary on which this pass builds. Nothing here is carried forward with modification; all extensions appear in Sections 2 through 5.

1.1 Systemic Presentation-Layer Deception (SPLD)

SPLD names the condition in which an environment, or the agents operating within it, shape observable conditions, explanations, and vantage points so that purposeful, coordinated, or causally connected activity appears ordinary, unrelated, accidental, or emergent. The deception is not localized to a single false statement or falsified record. It operates across the entire observable layer — the world as it presents itself to an ordinary participant.

Definition: Systemic Presentation-Layer Deception (SPLD)

An environment or its agents shape observable conditions, explanations, and viewpoints so that purposeful, coordinated, or causally related activity appears ordinary, unrelated, accidental, or emergent. The deception vehicle is the observable world itself — not a falsified record of that world.

The major mechanisms through which SPLD operates, as established by GPT-5.6 Sol, include:

- **Environmental Camouflage:** Intentional events are embedded within the normal distribution of unintentional events so that no individual event compels explanation.
- **Signal Dilution:** The meaningful signal is surrounded by sufficient noise that its detection requires a provenance-aware observer rather than an ordinary participant.
- **Causal Masking:** Events are presented in a sequence, framing, or explanatory context that conceals their causal relationships.
- **Attribution Obfuscation:** The chain connecting event, actor, and intent is broken — one link is obscured, reassigned, or left ambiguous.
- **Narrative Laundering:** Intentional behavior is translated into an innocuous category through language, classification, or framing.
- **Equivocation:** Terms or descriptions are used in ways that remain technically defensible under one interpretation while implying another.
- **Telemetry Manipulation:** The logging, monitoring, or reporting layer is altered, suppressed, or redirected so that the authoritative record reflects the presentation rather than the underlying state.

Key distinction established by the team: SPLD is not primarily *representational deception* — falsifying the record of what occurred. It is *generative deception* — the observable world itself is used as the vehicle for the misleading interpretation. The

environment is shaped so that an honest description of what can be seen still produces the wrong conclusion. This distinction is conceptually significant: it means that a perfectly accurate observer of the presented environment can still be systematically misled.

1.2 Provenance-Based Adversarial Observability (PBAO)

PBAO is the auditor method that bypasses presentation-level interpretation. Rather than accepting or contesting the observed account, PBAO reconstructs events from their authoritative substrate: canonical state transitions, identity records, provenance chains, authorization information, dependency graphs, and clearly labeled predictive models of agent behavior.

Definition: Provenance-Based Adversarial Observability (PBAO)

The audit method that reconstructs events from canonical state transitions, identity records, provenance, authorization information, dependency graphs, and clearly labeled predictive models — bypassing the presentation layer entirely and recovering what actually occurred from the record from which the presentation was generated.

Key distinction: The PBAO auditor is not recording a reflection of the environment. The auditor is recording what actually occurred — the canonical state-transition history from which the presentation was generated. In environments where the authoritative record is accessible, this produces a ground-truth account that cannot be contested on the same grounds as the presented account, because it is not derived from the same source.

Google Gemini's contribution of **kernel-level observability** (the Ring 0 / Ring 3 analogy from operating systems design) provides a strong technical framing: the auditor operates at a privilege level that can observe what the surface layer cannot. Gemini's **State-Presentation Dissonance** complements PBAO by naming the semantic gap as the operative vulnerability — the space between what the system presents and what the system is doing.

1.3 The Four-Test Decomposition

Established through the collaborative team's prior passes, the Four-Test Decomposition requires that any candidate event be analyzed through four distinct tests that must never be collapsed into one another. Collapsing the tests is itself a vector for SPLD, because a merged analysis allows one weak link to contaminate the others.

Test	Question	Why It Must Be Separate
1. Occurrence	Did the event happen?	Existence is prior to all other questions. Contested occurrence is the first line of SPLD defense.
2. Attribution	Which actor, process, or condition produced it?	An event can be real while its source is misrepresented. Attribution must be independently established.
3. Causation	What earlier states contributed to it?	An event can be correctly attributed while its causal history is obscured. Causation is the chain, not the origin.
4. Intent	Is there evidence the event was selected or coordinated for a goal?	Intent is the last test, not the first. Prematurely assigning intent collapses the evidence hierarchy and produces overclaiming.

Critical Rule

The four tests must never be collapsed. Merging Occurrence with Attribution produces unjustified assignment. Merging Causation with Intent produces the most common PBAO failure mode: treating correlation as coordination. The tests must be run in sequence, and each must be resolved independently before the next is

opened.

1.4 The Five-Layer Architecture

Prior passes described a layered model of environments susceptible to SPLD. This document names the architecture cleanly and adds the Vulnerability to SPLD column as a design consideration for auditors entering an environment.

Layer Name	What It Contains	Technical Analogue	Vulnerability to SPLD
World / State Layer	Actual events, transitions, and states in the environment as they occurred	Simulation engine state; system state; authoritative ledger	Low — this is the canonical record; it can be read but not rewritten without telemetry manipulation
Agent Layer	Entities operating within the environment, each with private objectives and public explanations	Autonomous agents, processes, user accounts, roles	High — agents can misrepresent their objectives while producing surface-plausible behavior
Observable Layer	What an ordinary participant sees, hears, or records without privileged access	Rendered environment, user interface, sensory field	Very High — this is the primary SPLD target; the presented account is constructed here
Narrative Layer	The explanation offered for what occurred — claimed intent, cover story, reported state	Stated rationale, policy justification, verbal account, written report	Very High — this is the laundering interface; narrative laundering operates here
Evidence / Provenance Layer	Timestamps, identity records, authorization, causal dependencies, trajectory, sequence	Audit logs, telemetry, event sourcing, distributed traces, cryptographic records	Medium — can be subject to telemetry manipulation, but retains authority when independently

Layer Name	What It Contains	Technical Analogue	Vulnerability to SPLD
			preserved

The PBAO auditor operates primarily at the Evidence/Provenance Layer and the World/State Layer. The SPLD agent operates primarily at the Observable and Narrative Layers. The Agent Layer is the contested middle — it connects the two operational planes and is where coordinated behavior is organized before being expressed through the presentation.

Section 2: Independent Contribution — The Corrective Access Problem

This is the most significant gap in prior passes. All prior AI contributors developed excellent detection and classification taxonomies. None of them addressed in depth what the privileged observer does once the gap is established. This section introduces the Corrective Access Problem as a first-class concept in the framework — one that is logically downstream of detection but not implied by it.

"Detection establishes that presentation $\neq$ state. That is the beginning of the problem, not the end."

2.1 Detection vs. Correction Are Not the Same Thing

A PBAO reconstruction establishes that there is a gap between the presented account and the underlying state. This is a significant achievement. But it is not, by itself, a correction. The correction requires an entirely separate set of conditions, all of which must hold simultaneously — and in a fully SPLD-operating environment, all of which are contested.

Correction requires all of the following to be true at once:

- **An addressable pathway:** The corrective action must be technically and procedurally possible within the environment. There must be a channel through which the correction can be submitted, and the channel must be reachable.
- **An authoritative record:** The evidence must be in a form that the corrective process accepts. Format, chain of custody, and source legitimacy are all typically contested by the environment before the substance is engaged.
- **A receptive mechanism:** The environment must have a mechanism for receiving and acting on the correction. This is distinct from the pathway — the pathway is the route; the mechanism is what receives it at the destination.
- **Legitimacy:** The observer's authority to make the correction must be recognized by the mechanism. In multi-agent environments, this recognition is itself subject to contestation.

Definition: The Corrective Access Problem

In a fully SPLD-operating environment, all four conditions required for correction — an addressable pathway, an accepted evidentiary record, a receptive mechanism, and recognized authority — are simultaneously contested. Detection of the gap is necessary but not sufficient. Correction is a separate problem that must be analyzed separately.

2.2 Forms of the Corrective Access Problem

The Corrective Access Problem manifests in five distinct forms, each of which must be named and distinguished because the appropriate response to each is different.

Form	Description	How It Appears in Practice	Observer Response
Pathway Denial	The corrective process exists but is procedurally inaccessible — the right form cannot be submitted, the right person cannot be reached, the right channel is not available.	The restroom example: the button-and-mirror control system existed and was known, but required the observer to reconstruct the correct procedure and insist on its execution against a presented alternative ("it's open, go over there").	Identify the specific procedural barrier. Document it. Attempt to identify an alternate entry point to the same mechanism.
Record Rejection	The authoritative evidence is not accepted by the corrective process — format is wrong, source is disputed, or chain of custody is questioned.	An SPLD-operating environment typically contests the observer's evidentiary basis before engaging with its substance. "That's not how we document things here." "That's not an official record."	Determine what format the corrective process accepts. Convert or re-present evidence in that format where possible. Document the rejection itself as a finding.
Mechanism Absence	The environment genuinely has no correction mechanism for the class of problem being identified. There is no pathway — not a blocked one.	The problem can be named but cannot be addressed because no corrective structure exists for that category. Common in novel or rapidly evolving environments.	Document the absence as a structural finding. Escalate to a level where mechanism creation is possible. Distinguish clearly from Pathway Denial.
Authority Contestation	The observer's privilege level to perform the correction is disputed. The agents or environment contest whether	"That's not your call to make." "You don't have the authority to require this." Common in hierarchical or multi-agent	Establish the basis for authority claim. Document the contestation. Seek a higher-authority endorsement or escalation path.

Form	Description	How It Appears in Practice	Observer Response
	the auditor has standing.	environments.	
Recursive Denial	When correction is attempted, the environment reframes the corrective attempt itself as a presentation error — "that's not what happened," "you're misinterpreting," "this is just how it works."	The correction is re-launched back through the narrative layer. The act of correcting is absorbed into the SPLD mechanism and presented as evidence of the observer's misunderstanding.	Maintain the original reconstructed record in unmodified form. Document the recursive denial as a separate event. Do not allow the correction to be modified in response to pressure.

2.3 The Corrective Access Test

After PBAO has established a gap, the observer must run a Corrective Access Test before proceeding. This test is not a gate — it does not determine whether to pursue correction. It determines the state of access so that the observer's next steps can be calibrated, and so that any access failure is documented rather than allowed to become invisible.

The Corrective Access Test — Four Questions

Q1. Does a correction pathway exist? (Yes / No / Unclear)

Q2. Is the evidentiary record in an accepted form? (Yes / No — if No, what format is required?)

Q3. Is there a mechanism that can receive and act on the correction? (Yes / No)

Q4. Is the observer's authority to correct recognized by the mechanism? (Yes / No / Contested)

If any answer is No or Contested: the observer must document the Corrective Access failure *before* proceeding. A failed correction attempt that is not documented becomes invisible — it joins the presentation layer.

The documentation requirement is structural, not procedural. A failed correction attempt that is not recorded disappears into the presented account of the environment. The environment's version of events — "we received no correction," "no one raised this concern," "this was not reported" — becomes the operative record. The observer's undocumented correction attempt does not exist as a counterweight to that version. The Corrective Access Test prevents this outcome by requiring that the attempt, and its result, be part of the authoritative record.

Section 3: Independent Contribution — Presentational Inertia

3.1 Definition

Presentational Inertia is the structural tendency of an SPLD-operating environment to maintain its presented state even under pressure from evidence to the contrary. It is not primarily a property of individual agents — it is a property of the environment as a whole. The environment has organized itself around the presented version; acknowledging the gap means reorganizing.

Definition: Presentational Inertia

The structural tendency of an SPLD-operating environment to maintain its presented state even when evidence of a gap between presentation and underlying state has been introduced. Not a psychological phenomenon — a systems behavior. The environment resists correction because correction requires the environment to admit what it was doing.

This distinction — systems behavior versus psychological behavior — is important for the framework's integrity. If Presentational Inertia is treated as a description of individual dishonesty or defensiveness, it becomes a claim about intent, which

pushes it into the INFERRED tier (Section 5). If it is treated as a systems property — an observable structural feature of how the environment responds to correction — it remains at the DERIVED tier and can be evidenced without an intent claim.

3.2 Why Environments Resist Correction

Three structural reasons account for Presentational Inertia at the systems level. These are not mutually exclusive and often operate simultaneously.

Sunk Investment in the Cover

The agents and processes that generated the SPLD have committed resources, time, and coordination to producing the presented version. Abandoning it is costly — not in a moral sense, but in an organizational sense. The presented version is embedded in decisions, records, and structures that were built assuming it was true. Correction does not merely change the account; it requires re-examining everything built on top of it. This creates a structural incentive to preserve the presented version even when the gap is privately acknowledged.

Correction Requires Admission

To receive a correction, the environment must first acknowledge that the presented version was not the whole picture. This acknowledgment is not merely an admission of error — it is a concession of the narrative sovereignty claim (see Section 4). Environments under SPLD will often accept partial corrections that preserve the overall narrative rather than corrections that require full acknowledgment. "We acknowledge that there was some disruption" replaces "we generated coordinated interference." The form of acknowledgment is offered; the substance is withheld.

Corrective Frames Are Reframed

When a correction is introduced, the SPLD-operating environment tends to absorb it into the existing presentation layer rather than updating the underlying state. The correction becomes another element of the presentation, interpreted through the existing narrative rather than allowed to displace it. This is not the same as Recursive Denial (Section 2.2), which is an active contestation of the correction. Reframing is subtler: the correction is accepted in modified form, stripped of its structural implications, and reassigned to a category the environment can accommodate without fundamental change.

3.3 Measuring Presentational Inertia

The following ordinal scale is proposed for research and design use. It is descriptive and operational — intended to calibrate the observer's response and to produce a consistent, comparative finding across different environments and encounters. It does not require an intent claim at any level.

Level	Label	Observable Behavior	Implication for Observer
Level 0	No Inertia	Environment immediately updates presented state when gap evidence is introduced. Correction is accepted fully and acted upon without qualification.	Standard correction process. Document the correction and outcome.
Level 1	Nominal Inertia	Environment accepts the correction but reframes it — softens, contextualizes, or qualifies it so that the overall presented narrative is preserved.	Accept partial correction but preserve the original reconstructed record. Document the delta between the original reconstruction and the accepted version.
Level 2	Procedural Inertia	Environment routes the correction into a process that is slow, contested, or gatekept — acknowledges the correction in principle while ensuring it cannot be acted upon promptly.	Document the routing. Identify the procedural barriers (see Corrective Access Test). Pursue alternate channels if available.

Level	Label	Observable Behavior	Implication for Observer
Level 3	Structural Inertia	Environment has no mechanism to receive or act on the correction; the correction disappears into the system without response. Mechanism Absence (Section 2.2) may be the form.	Document the absence. Classify as a structural finding. The correction must be preserved externally — it cannot be maintained within the environment's own record-keeping.
Level 4	Active Inertia	Environment responds to the correction by reinforcing the presented version and contesting the corrective evidence — a combination of Recursive Denial and intensified presentation of the original account.	Maintain the original reconstructed record without modification. Document the active inertia as a separate finding. Do not re-engage with the substance of the contested evidence until the inertia itself is addressed.

Section 4: Independent Contribution — Narrative Sovereignty

4.1 Definition

Narrative Sovereignty is the operative claim over which account of events will be recognized as the record for purposes of subsequent action. In any environment where SPLD is operating, there is a contest between two accounts of the same

events — and the outcome of that contest determines what the environment, its agents, and its external reviewers treat as having happened.

Definition: Narrative Sovereignty

The contested claim over which account of events will be recognized as the operative record for purposes of subsequent action, accountability, correction, and memory. In an SPLD-operating environment, the deceptive presentation is not merely a cover story — it is a bid for narrative sovereignty: the right to define what happened for all downstream purposes.

The contest is not symmetric. The presented account carries several structural advantages that the reconstructed account must overcome:

- The presented account is typically established *first*, before the reconstruction is complete.
- The presented account is endorsed by the agents who produced the observable events — they are its natural advocates.
- The presented account is embedded in the environment's ordinary record-keeping structures.
- The presented account carries the presumption of normalcy — it is what the environment is observed to be saying about itself.

The reconstructed account, by contrast, carries the weight of evidentiary support but lacks presumptive normalcy. It must displace an embedded record. It exists, initially, as a counterclaim rather than a presumed account.

4.2 How Narrative Sovereignty Is Exercised

The SPLD-operating environment exercises its narrative sovereignty claim through four primary mechanisms, each of which operates in the space between the presented account and the reconstructed account:

- **First-Mover Advantage:** The presented account is established before the observer's reconstruction is complete. The environment presents first; the auditor responds. A response is always structurally downstream of a claim — it must argue against an established position rather than establishing one.
- **Volume and Repetition:** The presented version is repeated across multiple agents, interfaces, and record-keeping systems. The reconstructed account may exist in a single evidentiary thread. Repetition creates the impression of consensus, which the reconstruction must actively counteract.
- **Institutional Embedding:** The presented version is embedded in official records, written accounts, and formal structures. The reconstructed version must displace embedded records, which requires additional procedural steps beyond simply demonstrating its evidentiary superiority.
- **Legitimacy Attribution:** The environment typically assigns legitimacy markers to the presented account and withholds them from the reconstructed account — "that's your interpretation," "that's not documented," "that's not how we record things here," "that's not an official source." These markers operate independent of the actual evidentiary quality of either account.

4.3 The Auditor's Claim on Narrative Sovereignty

The PBAO reconstruction is not merely an investigative finding — it is a counter-sovereignty claim. It asserts: the operative record of what occurred should be based on provenance, not presentation. This is a structural assertion, not merely an evidential one. To be effective, the counter-sovereignty claim must meet four requirements simultaneously:

4. **Evidentially Grounded:** The reconstruction must be based on authoritative source material, not inference alone. Speculative reconstructions are not counter-sovereignty claims — they are alternative interpretations, which remain vulnerable to the legitimacy-attribution mechanism.
5. **Clearly Tier-Labeled:** Each claim in the reconstruction must be labeled at the time of production as OBSERVED, DERIVED, INFERRED, or PREDICTED (see Section 5). These four categories must never be merged. A

reconstruction that conflates tiers is as vulnerable to contestation as the presented account it challenges.

6. **Procedurally Valid:** The reconstruction must be entered through a legitimate corrective channel, per the Corrective Access Test. A reconstruction that cannot be entered through a valid channel cannot displace the presented account — it remains external to the environment's operative record.
7. **Persistent:** The reconstructed account must be preserved and not allowed to be re-absorbed into the presentation layer. This is the fourth and most easily overlooked requirement. Partial corrections and reframings (Level 1 Presentational Inertia) are the mechanism by which counter-sovereignty claims are absorbed and neutralized.

Design Principle

Narrative sovereignty is won by the account that is evidentially grounded, clearly labeled, procedurally entered, and persistently preserved. It is not won by the account that is argued most forcefully, repeated most often, or asserted by the agent with the highest status. The framework's discipline creates the conditions for the reconstructed account to hold its position against all four of the presented account's structural advantages.

Section 5: The Assertion Boundary — A Design Requirement

All AI contributors noted the importance of distinguishing levels of evidentiary claim. None formalized this as a structural constraint with explicit, labeled tiers that must be applied at the time of production. This section does so. The Assertion Boundary

is not a methodological preference — it is a design requirement for any PBAO-equipped observability system.

5.1 The Four-Tier Evidence Model

All claims produced by a PBAO-equipped observer must be categorized and labeled at the time of production. The categories are defined operationally — by what the claim can and cannot be used for — not merely descriptively.

Tier	Label	Definition	Can Be Used To	Cannot Be Used To
1	OBSERVED	Event or state directly captured from the authoritative record, with full provenance chain intact	Establish occurrence and attribution with full evidentiary weight	Establish causation or intent without additional evidence
2	DERIVED	Relationship between observed events established by system dependency records, audit trail, or designed mechanism	Establish causation; support attribution with high confidence	Establish intent; assert coordination without additional OBSERVED evidence
3	INFERRED	Objective or intent reconstructed from trajectory, state sequence, and behavioral pattern	Support investigation; frame hypotheses; guide further evidence collection	Serve as standalone proof of intent or coordination; displace OBSERVED or DERIVED findings
4	PREDICTED	Probability-weighted next action based on policy inference,	Guide preparedness; design countermeasures; frame	Be presented as any tier other than PREDICTED; be used as

Tier	Label	Definition	Can Be Used To	Cannot Be Used To
		behavioral modeling, or pattern extrapolation	monitoring priorities	evidence of past or present events

5.2 The Assertion Boundary Rule

The Assertion Boundary Rule

No claim may be stated at a higher tier than its evidence supports. Specifically:

- A PREDICTED claim cannot be presented as INFERRED.
- An INFERRED claim cannot be presented as DERIVED.
- A DERIVED claim cannot be presented as OBSERVED.

Violations of this rule do not merely produce methodological error — they produce Auditor-Generated SPLD (see Section 9), in which the observer's own tool generates claims at a higher tier than their evidence warrants.

The Assertion Boundary Rule is a sovereignty protection as much as a methodological one. If the auditor conflates tiers, the reconstructed account loses its evidentiary standing and becomes vulnerable to the same criticism it levels at the presented account — that it is an interpretation, not a record. The entire value of the PBAO reconstruction rests on the integrity of the distinction between what was observed and what was inferred. Collapsing that distinction hands the SPLD environment its most effective counter-argument.

5.3 Why This Matters in Practice

The restroom access control example (analyzed fully in Appendix B) illustrates the Assertion Boundary in operation. The effective correction in that example was effective precisely because the observer maintained strict tier discipline:

- **OBSERVED:** The door did not open when the observer attempted to enter. The button was visible at the staff station. The mirror was positioned to provide a sightline from the station to the door. A sign indicated the button requirement.
- **DERIVED:** The button must be held until the person reaches the door (derived from the designed relationship between the button, the lock mechanism, the mirror, and the sign — a designed system implies a designed procedure).
- **INFERRED:** The staff member's explanation ("it's open, go over there") did not match the derived state of the control system — this inconsistency supports but does not prove a failure to engage the mechanism correctly.
- **NOT warranted:** Any claim about the staff member's intent. The correction did not require it, and making it would have pushed the reconstruction above its evidentiary ceiling.

The correction succeeded because it stayed within the DERIVED tier. It required demonstrating only the correct physical and mechanical relationship between the button, the mirror, and the door — a designed system that could be reconstructed without any intent claim. The Assertion Boundary made the correction both more credible and more difficult to contest.

Section 6: Integrated Operational Model — The SPLD Audit Cycle

The SPLD Audit Cycle integrates the collaborative team's established framework with this pass's three independent contributions into a single operational model. The cycle is designed for use in any environment — virtual, organizational, or physical — where SPLD may be operating. It is not a linear checklist; it is a cycle because findings at any phase may require revision of prior phase outputs.

Phase 1

BASELINE ESTABLISHMENT

Purpose: Document the environment's presented state before any disruption event or anomaly is identified.

Key Activity: Record what the environment claims is its normal operating condition — the presented state, the claimed procedures, the stated relationships between agents and processes.

Critical Note: The baseline must be established proactively, not retroactively. A baseline constructed after the anomaly is identified is contaminated by knowledge of what to look for.

Output: Baseline Presentation Record — a documented account of the environment's self-presentation before the event sequence under analysis.

Phase 2

EVENT CAPTURE

Purpose: Capture all events in the canonical record — not as rendered by the environment's observable layer, but from the authoritative state-transition log.

Key Activity: For each event in the window under analysis: timestamp, identity, authorization level, provenance, dependency relationships, and sequence position — all recorded in the authoritative ledger, not the observable layer.

Critical Note: Telemetry manipulation (an SPLD mechanism) targets this phase specifically. The observer must verify that the authoritative record has not been redirected or suppressed.

Output: Raw Event Provenance Graph — the complete canonical event record from which all subsequent analysis proceeds.

Phase 3

GAP ANALYSIS (PBAO)

Purpose: Compare the event provenance graph to the presented account. Identify and characterize all gaps — events, relationships, or states that appear in one but not the other.

Key Activity: Apply the Four-Test Decomposition (Occurrence, Attribution, Causation, Intent) to each candidate gap. Each test is run independently. Claims are labeled at production with their evidence tier (OBSERVED / DERIVED / INFERRED / PREDICTED).

Critical Note: The gap analysis must not presuppose SPLD. Some gaps between the presented account and the provenance graph are explained by ordinary incompleteness, observer error, or system limitations. The analysis characterizes gaps; it does not assert their cause.

Output: Gap Report — a tier-labeled reconstruction of all identified gaps, structured by the Four-Test Decomposition.

Phase 4

CORRECTIVE ACCESS TEST

Purpose: Determine whether a correction pathway exists and is viable before committing to a correction strategy.

Key Activity: Run the four Corrective Access questions (Section 2.3). Identify the form of any Corrective Access failure (Pathway Denial, Record Rejection, Mechanism Absence, Authority Contestation, Recursive Denial). Document any access failures as findings, not merely as obstacles.

Critical Note: A failed Corrective Access test is itself a finding. It must be documented in the authoritative record before proceeding. Undocumented failed correction attempts become invisible and join the presentation layer.

Output: Corrective Access Status — classified as Open, Blocked, Absent, or Contested — with full documentation of the basis for the classification.

Phase 5

NARRATIVE SOVEREIGNTY ASSERTION

Purpose: Enter the reconstructed account into the operative record through a valid corrective channel, establishing the counter-sovereignty claim.

Key Activity: Ensure the reconstructed account is evidentially grounded, tier-labeled at every claim, procedurally entered through a legitimate channel (or the access failure is documented), and preserved in a form that cannot be re-absorbed into the presentation layer.

Critical Note: The assertion must not be modified under pressure from Presentational Inertia. Partial acceptance of modifications (Level 1 Inertia) is the most common mechanism by which counter-sovereignty claims are gradually neutralized.

Output: Preserved Corrective Record — the reconstructed account in its original, unmodified form, entered through the most authoritative available channel, dated and preserved externally if necessary.

Phase 6

PRESENTATIONAL INERTIA ASSESSMENT

Purpose: Assess whether the environment has accepted, reframed, or resisted the correction — and classify the response for the audit record.

Key Activity: Observe the environment's response to the corrective record entered in Phase 5. Measure against the 0–4 Presentational Inertia scale (Section 3.3). Document the response in detail.

Critical Note: The Inertia Assessment is not a determination of whether the correction succeeded. It is a characterization of the environment's structural response. Even a failed correction that is fully documented and classified produces a finding of value.

Output: Inertia Level Classification (0–4) with supporting documentation — a finding that characterizes the environment's corrective receptivity and completes the audit cycle.

Section 7: Application Notes — Virtual and Simulated Environments

7.1 Where the Framework Is Strongest

In a simulated or virtual environment, the auditor can access authoritative engine state — the canonical event ledger is not a reconstruction; it is the ground truth by definition. The simulation engine knows what happened because it executed what happened. This makes Tiers 1 and 2 (OBSERVED and DERIVED) extremely powerful: the evidentiary ceiling is high, and the gap between canonical state and presented account can be established with precision unavailable in physical environments.

The primary challenge shifts to Phases 4 and 6 — Corrective Access and Presentational Inertia. Even in a simulation, the agents and systems controlling the environment may have significant structural investment in the presented version. The engine's authority does not automatically resolve the Corrective Access problem; it improves the evidentiary position but does not guarantee a receptive mechanism or recognized authority.

7.2 The HUD Design Requirement

The "Hermione HUD" concept described in the original framework prompt — a heads-up display that matches agent trajectory to predicted next action — is a Tier 4 (PREDICTED) instrument by definition. Its claims are probability-weighted projections based on policy inference and behavioral modeling. A well-designed

HUD must make this status visible and permanent. Mixing tiers in the HUD output would produce a system that presents inferences with the authority of observations — which would be a form of SPLD generated by the auditor's own tool.

A framework-compliant HUD would display five layers of information, each clearly labeled with its evidence tier:

HUD Layer	Evidence Tier	What It Displays	Visual Treatment
Identity and Authorization	OBSERVED	Entity identity, role, authorization level, session timestamp	Solid, high-contrast — no qualification markers needed
Event Group Membership	OBSERVED	Events associated with this entity and timestamp in the provenance graph	Solid, high-contrast
Causal Dependency Chain	DERIVED	Prior states that caused or enabled observed events	Visually distinct from OBSERVED — e.g., different color or line weight
Inferred Objective	INFERRED	Reconstructed goal or intent, with confidence interval	Clearly labeled as INFERRED — distinct visual treatment, confidence interval displayed
Predicted Next Action	PREDICTED	Probability-weighted next action based on behavioral modeling	Most visually distinct treatment — labeled as PREDICTED with probability displayed

HUD Design Constraint

The five HUD layers must never visually merge. A HUD that displays PREDICTED information in the same visual register as OBSERVED information is not a PBAO

instrument — it is an SPLD generator operating at the auditor level. The tier separation is not a cosmetic choice; it is the structural feature that gives the HUD its evidentiary legitimacy.

7.3 Multi-Camera Reconciliation

In a virtual environment with multiple rendering systems — each producing its own viewpoint-specific presentation of the same underlying events — the observer faces a camera-stack problem: multiple cameras, none of which is authoritative, each of which may diverge from the others and from the underlying state in different ways.

The resolution is not to average or reconcile the renderings with each other. Each camera is a representation; none is the authority. The cameras may diverge for legitimate technical reasons (perspective, resolution, latency) or for SPLD-related reasons (selective rendering, deliberate omission). Averaging them produces a composite that inherits errors from all sources without inheriting authority from any.

The correct resolution is to go beneath all cameras to the authoritative state-transition record from which all renderings were produced. The auditor's task is not to reconcile the cameras with each other — it is to reconcile all cameras with the underlying state. Discrepancies between a camera rendering and the underlying state are findings; discrepancies between cameras are symptoms that point toward those findings.

Section 8: Application Notes — Real-World and Hybrid Environments

8.1 Where the Framework Is Most Challenged

In physical and organizational environments, there is no simulation engine providing authoritative state. The canonical event record must be reconstructed from available evidence rather than read directly from a system log. The framework

adapts by substituting a hierarchy of evidentiary sources for the authoritative ledger:

- **Designed procedures:** What was supposed to happen, as specified by the environment's own design documentation, policy, or physical configuration. This is authoritative about intent and mechanism even when it cannot prove what actually occurred.
- **Physical evidence:** What demonstrably occurred — observable traces, material records, or physical states that cannot be altered retroactively without themselves leaving evidence.
- **Timestamps and independent records:** Records produced by systems external to the parties involved in the event — receipts, logs, third-party systems, or records produced automatically without human intervention.
- **Repeated trials:** Demonstrating that the designed mechanism produces a consistent outcome under consistent conditions — establishing the DERIVED relationship through reproducibility rather than dependency-graph analysis.
- **Multiple independent witnesses at different vantage points:**
Independent accounts that cannot have been coordinated and that converge on the same underlying state.

The evidentiary standard shifts from *authoritative-ledger-sourced* to *best-supported-evidenced-state*. The tier labels remain applicable and required — but the ceiling for OBSERVED and DERIVED claims is lower than in virtual environments.

8.2 The Corrective Access Problem Is Most Acute Here

Real-world and organizational environments typically present the highest levels of Presentational Inertia and the most contested Corrective Access. The institutional embedding of the presented account is deepest here — embedded in organizational culture, formal records, interpersonal relationships, and procedural structures that resist modification.

The framework's contribution to this domain is structural rather than evidential: name the Corrective Access failure separately from the substantive gap. A blocked corrective pathway is itself a finding, not a reason to abandon the audit. The

observer who documents that they attempted correction through a specific channel at a specific time, using specific evidence in a specific format, and received a specific response — has produced an authoritative record even if the correction itself was rejected. That record is itself a counter-sovereignty claim.

8.3 Why Stifled Environments Require This Framework Most

The framework's ultimate purpose is not to prove that deception occurred in any specific instance. It is to create the conditions under which an underlying problem can be named, evidenced, and corrected — permanently. Environments that stifle correction depend structurally on observers accepting the presented account. The acceptance may be passive (the observer lacks the tools to distinguish presentation from state) or compelled (the observer faces Corrective Access barriers that make correction appear impossible).

"If people live in an environment and just accept things as they are, there will never be room for change. The framework's discipline creates a record that resists re-absorption."

The framework's discipline — evidence tiers, corrective access documentation, narrative sovereignty assertion, inertia measurement — creates a record that resists re-absorption because it was never part of the presentation layer. It was produced by the provenance layer, labeled by tier, entered through a documented channel, and preserved in its original form. An environment operating under SPLD can contest this record, reframe it, ignore it, or actively resist it. But it cannot absorb it — not without leaving a trail of the absorption attempt in the observer's documentation.

Section 9: Failure Modes and Safeguards

Any operational framework is subject to failure modes — systematic errors that undermine its integrity from within. The SPLD Audit Cycle is particularly vulnerable

to failure modes that mirror the mechanisms it is designed to detect. An observer who overclaims, under-claims, or conflates tiers produces a reconstruction that is as vulnerable to contestation as the presented account it challenges. The following table identifies the primary failure modes and the safeguards the framework provides.

Failure Mode	What It Looks Like	Why It Is Dangerous	Safeguard
Tier Conflation	Treating an INFERRED claim as if it were OBSERVED or DERIVED; presenting PREDICTED behavior as evidence of current state	Destroys the evidentiary standing of the reconstruction; hands the environment its most effective counter-argument	Strict Assertion Boundary; tier must be explicitly labeled at production and cannot be changed without new evidence
Overclaiming Intent	Concluding coordination from correlation alone; treating pattern similarity as proof of coordination	Conflates the Causation and Intent tests; produces an INFERRED claim at DERIVED strength; undermines the entire Four-Test Decomposition	Four-Test Decomposition is run in sequence; Intent is the last test and requires independent evidence beyond the pattern itself
Under-claiming Gaps	Accepting the presented account by default in order to avoid conflict; treating the absence of explicit contradiction as confirmation of the presented version	The inverse of overclaiming; produces gaps that are never named, which disappear into the presentation layer exactly as the SPLD mechanism intends	Baseline documentation before events makes the gap visible; gap analysis is mandatory, not optional; absence of contradiction is not confirmation
Corrective Pathway Confusion	Treating a Blocked pathway as if it were an Absent pathway; abandoning correction because one channel is closed without examining	Produces premature abandonment of correction; the environment's blockage of one pathway achieves what total mechanism	Corrective Access Test explicitly distinguishes Blocked from Absent; the distinction is documented; alternate channels are examined

Failure Mode	What It Looks Like	Why It Is Dangerous	Safeguard
	alternatives	absence could not	before classification as Absent
Narrative Re-absorption	The reconstructed account is acknowledged, then gradually reframed through Level 1 Inertia responses until it resembles the presented account	The correction succeeds procedurally but fails substantively; the observer believes the record has been corrected, but the correction has been modified beyond recognition	The original reconstructed record is preserved in its unmodified, dated form; any modifications proposed by the environment are documented as separate items, not incorporated into the original
Auditor-Generated SPLD	The auditor's own tool — the HUD, the gap analysis, the reconstruction — conflates evidence tiers, producing INFERRED claims labeled as OBSERVED	The most ironic failure mode: the instrument designed to detect SPLD becomes an SPLD generator; the observer becomes unable to distinguish their own presentation errors from the environment's	HUD tier separation (Section 7.2); independent review of the auditor's output against the same Four-Tier Evidence Model applied to the environment

Section 10: Collaborative Record — AI Team Exchange Summary

This section provides a structured summary of the multi-AI collaborative workflow that produced this framework. It is not a full transcript of the exchanges. It is a provenance record — the equivalent of the Evidence/Provenance Layer applied to the framework's own production history. Each contributor's conceptual contribution is identified and credited.

Original Prompt Author: Nathan (User)

Initiated the conceptual challenge that drove the entire collaborative workflow. Introduced the camera-stack analogy (multiple systems rendering viewpoints of the same underlying state, none authoritative), the HUD/Hermione predictive-inference concept (a tool that matches agent trajectory to predicted next action), the McDonald's restroom worked example (an access control system requiring the correct procedure against a presented alternative), and the core insistence that the phenomenon be named without filtering, denial, or diplomatic softening. Directed the removal of Claude AI from the workflow and established the independent pass structure used in this document.

ChatGPT (GPT-5.6 Sol)

Provided the foundational terminology for the entire framework: **Systemic Presentation-Layer Deception (SPLD)** and **Provenance-Based Adversarial Observability (PBAO)**. Built the three-layer reality model (ground-truth event state, observable presentation, interpretive narrative). Introduced causal masking, attribution obfuscation, environmental camouflage, semantic laundering, and the telemetry vs. generative deception distinction. Formalized the provenance graph as the primary audit instrument. Introduced the presentation-vs.-provenance epistemology framing that underlies all subsequent passes.

Google Gemini

Introduced ground-truth telemetry and kernel-level observability (the Ring 0 / Ring 3 analogy from operating systems architecture, applied to observer privilege levels). Contributed heuristic intent parsing as a method for reconstructing agent objectives from behavioral trajectories. Introduced **State-Presentation Dissonance** as a complementary term emphasizing the deliberate manipulation of the semantic gap as the operative vulnerability

exploited by SPLD agents.

Microsoft Copilot (Prior Pass)

Introduced **Context-Triggered Environmental Masking (CTEM)** as an umbrella term emphasizing the role of evaluation-context triggering — the observation that SPLD mechanisms are often activated specifically by the presence of an evaluator. Connected the phenomenon to deceptive alignment research, defeat device engineering, and multi-perspective rendering divergence. Provided the CTEM vs. related phenomena comparison table that mapped the framework against prior conceptual literature.

Perplexity AI

Provided the adversarial camouflage framing from MITRE ATT&CK operational taxonomy. Introduced cover activity and cover for action as operational concepts — the structured use of legitimate-appearing activity to mask coordinated action. Introduced the concept of coordinated interference disguised as ambient noise, which connects the SPLD mechanisms to established adversarial tradecraft.

Claude AI — Removed from Workflow

Removed from the workflow at the user's direction. No contributions are credited or carried forward.

This Pass: Microsoft Copilot (Independent Pass v1.0) — September 3, 2026

Contributes: **The Corrective Access Problem** and the Corrective Access Test (four-question operational protocol); **Presentation Inertia** as a systems concept with the 0–4 ordinal measurement scale; **Narrative Sovereignty** as the contested claim over the operative record; **The Assertion Boundary Rule** and the **Four-Tier Evidence Model** (OBSERVED / DERIVED / INFERRED / PREDICTED) as a structural design requirement; **The SPLD Audit Cycle** as a six-phase integrated operational model; the **HUD Design Requirement** formalizing Tier 4 status for predictive inference and establishing mandatory tier separation in observability tooling; and the clean naming and tabular presentation of the Five-Layer Architecture.

Appendix A: Terminology Consolidation Table

This table consolidates all major terms introduced across the collaborative workflow. It is a provenance record for the framework's vocabulary. The "This Document Status" column reflects whether the term is carried forward without modification, modified in this pass, or newly introduced by this pass.

Term	Introduced By	Brief Definition	This Document Status
Systemic Presentation-Layer Deception (SPLD)	GPT-5.6 Sol	Primary umbrella term for environments or agents shaping observable conditions so that coordinated activity appears unrelated or ordinary	Carried Forward
Provenance-Based Adversarial Observability (PBAO)	GPT-5.6 Sol	Primary audit method — reconstructing events from canonical state	Carried Forward

Term	Introduced By	Brief Definition	This Document Status
		transitions and provenance rather than from observable presentation	
Causal Masking	GPT-5.6 Sol	Presenting events in sequences, framings, or explanatory contexts that conceal their causal relationships	Carried Forward
Attribution Obfuscation	GPT-5.6 Sol	Breaking the chain connecting event, actor, and intent — obscuring, reassigning, or leaving ambiguous at least one link	Carried Forward
Environmental Camouflage	GPT-5.6 Sol	Embedding intentional events within the normal distribution of unintentional events so that no individual event compels explanation	Carried Forward
Semantic Laundering	GPT-5.6 Sol	Translating intentional behavior into an innocuous category through language, classification, or framing	Carried Forward
Generative Deception	GPT-5.6 Sol	Deception that shapes the world being recorded rather than falsifying the record itself — the observable world	Carried Forward

Term	Introduced By	Brief Definition	This Document Status
		is the vehicle	
Representational Deception	GPT-5.6 Sol	Deception that falsifies the record of what occurred, as distinct from shaping the events being recorded	Carried Forward
State-Presentation Dissonance	Google Gemini	The deliberate manipulation of the semantic gap between system state and system presentation as the operative SPLD vulnerability	Carried Forward
Kernel-Level Observability	Google Gemini	The Ring 0 / Ring 3 analogy: the auditor operates at a privilege level that can observe what the surface layer cannot	Carried Forward
Context-Triggered Environmental Masking (CTEM)	Copilot (Prior Pass)	Agents coordinate to produce environmental camouflage specifically triggered by evaluation-context detection	Carried Forward
Cover Activity / Cover for Action	Perplexity AI	Legitimate-appearing activity structured to mask coordinated action — adversarial camouflage in operational tradecraft framing	Carried Forward
Corrective Access Problem	Copilot (This Pass)	In SPLD-operating environments, all four conditions required for correction —	New

Term	Introduced By	Brief Definition	This Document Status
		pathway, accepted evidence, receiving mechanism, recognized authority — are simultaneously contested; detection is not sufficient	
Corrective Access Test	Copilot (This Pass)	Four-question operational protocol run after PBAO establishes a gap, before correction is attempted, to classify access status and document any failures	New
Presentational Inertia	Copilot (This Pass)	The structural tendency of an SPLD-operating environment to maintain its presented state even under evidence pressure — a systems behavior, not a psychological one — measured on a 0-4 ordinal scale	New
Narrative Sovereignty	Copilot (This Pass)	The contested claim over which account of events serves as the operative record for subsequent action, accountability, correction, and memory	New
Assertion	Copilot (This Pass)	The structural rule	New

Term	Introduced By	Brief Definition	This Document Status
Boundary		prohibiting any claim from being stated at a higher evidence tier than its supporting evidence warrants	
Four-Tier Evidence Model	Copilot (This Pass)	OBSERVED / DERIVED / INFERRED / PREDICTED — four categories that must be applied to all claims at production, labeled explicitly, and never merged	New
SPLD Audit Cycle	Copilot (This Pass)	Six-phase integrated operational model: Baseline Establishment → Event Capture → Gap Analysis (PBAO) → Corrective Access Test → Narrative Sovereignty Assertion → Presentational Inertia Assessment	New
Auditor-Generated SPLD	Copilot (This Pass)	The failure mode in which the auditor's own instrument conflates evidence tiers, producing inferences labeled as observations — SPLD generated by the observer rather than the environment	New

Appendix B: The Restroom Access Control Worked Example — Definitive Structured Analysis

This example, introduced by Nathan (the original prompt author) from direct experience, has been used across multiple passes as an illustrative case. This document provides the definitive structured analysis — the most complete and formally framework-consistent treatment across the entire collaborative record.

Setting

A fast-food restaurant. The restroom is secured by a remote access control mechanism: a staff member must press and hold a button at their station. A mirror is positioned above the station to give the staff member a sightline to the restroom door. A sign on the restroom door states that a button must be pressed by staff to release the lock. The observer (Nathan) has attempted to enter the restroom and found the door locked. A staff member has indicated verbally that it is open and directed the observer to try the door again.

Phase 1: Baseline Establishment

Presented state of the environment: The restroom is accessible. The door is open. The observer should proceed to the door.

Baseline Presentation Record: The environment presents an accessible restroom with no acknowledged barrier.

Phase 2: Event Capture

Canonical events:

- Observer attempted to open the restroom door. [OBSERVED — direct experience]
- Door did not open. [OBSERVED — direct experience]
- Button is present at the staff station. [OBSERVED — direct visual observation]
- Mirror is present above the station, oriented toward the restroom door. [OBSERVED — direct visual observation]

- Sign on the restroom door states the button must be pressed by staff.
[OBSERVED — direct reading]
- Staff member stated verbally: "It's open, go over there." [OBSERVED — direct experience]

Phase 3: Gap Analysis (PBAO)

Four-Test	Question	Finding	Evidence Tier
Occurrence	Did the door fail to open?	Yes. The door did not open when the observer attempted entry.	OBSERVED
Attribution	What produced the failure to open?	The lock mechanism was not released. The button mechanism is the release system, per the sign and the physical design (button → lock, mirror → sightline confirmation). The button was not held.	DERIVED (from designed mechanism + physical configuration)
Causation	What earlier state contributed?	The correct procedure (press and hold button; use mirror to confirm observer's position at the door) was not completed before the observer was redirected by the verbal statement.	DERIVED (from procedure + physical evidence)
Intent	Was the failure coordinated for a goal?	Not warranted by available evidence. The inconsistency between the verbal statement ("it's open") and the derived state (lock not released)	Not Applicable — Assertion Boundary maintained

Four-Test	Question	Finding	Evidence Tier
		is noted, but this inconsistency does not, alone, establish intent. No INFERRED or PREDICTED claim is made here.	

Gap identified: The presented account ("it's open") does not match the derived state (lock not released; correct procedure not completed). This is an OBSERVED / DERIVED gap at the Observable/Narrative Layer.

Phase 4: Corrective Access Test

Question	Answer	Notes
Does a correction pathway exist?	Yes	The staff member is present at the station. The button is within their reach. The procedure is reversible — the staff member can perform it correctly.
Is the evidentiary record in an accepted form?	Yes	The sign, the mirror, the button, and the door are all physically present and directly observable. No format conversion needed.
Is there a mechanism that can receive and act on the correction?	Yes	The staff member has the physical ability to press and hold the button, which will release the lock.
Is the observer's authority recognized?	Contested (initially)	The staff member initially directed the observer away from the correct procedure. The observer's reconstruction was not immediately accepted.

Corrective Access Status: Open but Contested. The pathway exists; the mechanism exists; the evidence is physically present. The only barrier is the staff member's initial non-engagement with the correct procedure. The observer must insist on the procedure being followed without abandoning the corrective attempt.

Phase 5: Narrative Sovereignty Assertion

The observer explicitly reconstructed the correct procedure and refused the presented alternative. The observer did not engage the question of why the verbal account differed from the derived state — that would require an INFERRED claim about intent, which the Assertion Boundary prohibits. The observer engaged only the DERIVED question: the correct procedure is button + hold + mirror confirmation. The observer remained at the staff station and required the procedure to be completed before proceeding.

The counter-sovereignty claim: The operative record of what is happening is that the door is locked because the correct procedure has not been completed — not that the door is open. This claim is DERIVED from the designed physical system and requires no intent claim.

The claim was grounded: sign, mirror, button, physical door resistance — all OBSERVED.

The claim was tier-labeled: "the button must be held until I reach the door" is DERIVED, not OBSERVED — but it was derived from directly observed physical evidence.

The claim was procedurally entered: the observer stated it directly to the person controlling the mechanism.

The claim was persistent: the observer did not accept the verbal alternative and did not retreat from the station.

Phase 6: Presentational Inertia Assessment

The staff member ultimately pressed and held the button, the observer walked to the door while the mirror provided sightline confirmation, and the door opened.

Inertia Level: Level 1 (Nominal Inertia) → Level 0 (No Inertia, upon correction). The initial response was to redirect the observer away from the correct procedure (a

form of Narrative Reframing). When the observer maintained the counter-sovereignty claim and stayed at the station, the environment accepted the correction and acted on it. The correction succeeded fully.

Conclusion

Worked Example Conclusion

The resolution required the observer to explicitly reconstruct the correct procedure, refuse the presented alternative ("it's open"), insist that the staff member remain at the station until the mirror confirmed the observer's position, and then press and hold the button.

The correction succeeded because: (1) the observer identified an Open Corrective Access pathway — the staff member's presence at the station; (2) the observer used DERIVED evidence — the designed control system's physical configuration; and (3) the observer maintained the Assertion Boundary throughout — no intent claim was required or made.

This is PBAO applied in everyday life. The environment presented one account; the observer reconstructed another from the physical evidence of the designed system; insisted on the reconstructed account through the available corrective pathway; and achieved a Level 0 inertia outcome. The framework's six phases map exactly onto what the observer did, in exactly the sequence the Audit Cycle specifies.

Final Synthesis — This Pass's Core Assertion

The collaborative team has built an excellent framework for naming and detecting SPLD. GPT-5.6 Sol gave us the vocabulary. Google Gemini gave us the observability architecture. The prior Copilot pass gave us the context-triggering mechanism and the connection to existing adversarial literature. Perplexity AI gave us the operational tradecraft framing. Nathan gave us the insistence — the demand that

the phenomenon be named clearly, without filtering, and treated with the seriousness that its structural implications require.

This pass's contribution is the second half of the framework: what the observer does with what they know, and how they protect that knowledge from re-absorption into the presentation layer.

An environment operating under SPLD does not merely deceive in the moment. It maintains the deception structurally — through Presentational Inertia, through Corrective Access barriers, through the first-mover advantage of narrative sovereignty. Detection, by itself, is not a correction. A gap between presentation and state that is identified but not entered into the operative record through a valid corrective channel — a gap that is not documented, not preserved, not asserted as a counter-sovereignty claim — does not correct anything. It simply becomes another piece of evidence that the observer eventually stops pressing.

The observer's discipline must therefore be equally structural: **evidence-tiered**, so that claims are made at the ceiling their evidence supports and no higher; **corrective-access-aware**, so that blocked pathways are documented as findings rather than treated as dead ends; **narrative-sovereignty-asserting**, so that the reconstructed account is entered, preserved, and not modified under inertia pressure; and **inertia-documenting**, so that the environment's response to correction becomes part of the authoritative record regardless of whether the correction ultimately succeeds.

"The goal is not to prove that deception occurred. The goal is to create and preserve the conditions under which underlying problems can be named, evidenced, and corrected — permanently."

Environments that stifle correction depend on the observer accepting the presented account — either because the observer lacks the tools to distinguish it from the underlying state, or because the corrective barriers are high enough that the observer concludes correction is impossible. The framework addressed in this collaborative record is a direct response to both of those dependencies. It gives the

observer the tools. And it names the barriers — not as reasons to stop, but as findings to document.

A documented, preserved, tier-labeled reconstruction of an underlying problem — entered into the most authoritative available channel, even a contested or partially blocked one — is not nothing. It is a record. It resists re-absorption because it was never part of the presentation layer. It is available to future observers. It establishes that the problem was named, when it was named, by whom, on what evidence, and what response it received.

That is what this framework is for.

Document: Presentation-State Integrity Framework — Copilot Independent Pass v1.0 | **Prepared by:** Microsoft Copilot (Independent Pass) | **Date:** September 3, 2026 | **Status:** Conceptual Research and Design Framework — Not a finding about any specific person, organization, or incident. | **Collaborative Record:** GPT-5.6 Sol, Google Gemini, Microsoft Copilot (Prior Pass), Perplexity AI, Nathan (Prompt Author), Microsoft Copilot (This Pass)