

Root-Level, Reality, and Simulation Structures

Foundational Research Framework for Mapping, Comparing, Modeling, Lifecycle Design, and Conditional Failure

Date: September 23, 2026

Status: DRAFT — UNDER REVIEW — NOT CANON — ONESTAMP PENDING

Direct finding

The corrected program should study structures before it studies root failure. Conditional Failure at the Root Level (CFTRL) is therefore an applied downstream branch, while the parent program examines how structures arise, persist, nest, constrain, represent worlds, become treated as real, and can be compared under incomplete knowledge.

The program can rigorously investigate known engineered systems, nested simulations, digital twins, complex adaptive systems, lifecycle behavior, observer limitations, and failure diagnosis. It cannot currently prove whether experienced reality is base reality, derive an architecture outside all observable structure, or infer simulation from an unexplained anomaly. Those limits should be recorded as findings rather than treated as project failures.

Foundational statement

This research does not begin by assuming that reality, base reality, simulation, or root-level structure has already been identified or adequately defined. It studies how structures arise, persist, interact, nest, change, and fail; how observers classify structures as real, literal, simulated, foundational, or derivative; what evidence supports those classifications; and how provisional maps can be constructed under incomplete information. The resulting knowledge will support—but will not be limited to—the later design and development of Conditional Failure at the Root Level.

Why the program changed

The originating material asked about root-level structures, base reality, simulation, incomplete maps, laboratory design, and failure as an investigative instrument. The first synthesis cycle nevertheless placed “Conditional Failure at the Root Level” in nearly every title and made failure the organizing subject. That work remains useful as an applied sub-dossier, but it reversed the intended dependency: a failure framework cannot responsibly characterize root-level failure until the project has defined what “structure,” “root,” “reality,” “simulation,” “map,” and “evidence” mean within declared boundaries.

The corrected dependency is:

Structural study → reality-model study → simulation study → comparative mapping → lifecycle study → failure study → CFTRL application.

Research position

A system is an arrangement of interacting elements that exhibits behavior or meaning not attributable to isolated elements alone, and systems engineering explicitly treats system boundaries as demarcations between a system of interest and its environment [¹][²]. This supports a boundary-relative definition of rootness: a component may be foundational for one function and boundary without being ontologically ultimate.

Complex-systems research adds that global properties may emerge from local interactions without a central controller, so a foundational constraint or dependency may arise rather than be intentionally designed [³][⁴]. “Root” must therefore include engineered, inherited, distributed, plural, and emergent candidates.

Scientific models selectively represent targets or interpret theories; they are not identical to the represented world [⁵]. NASA’s modeling standard accordingly separates verification—whether a model conforms to its specification—from validation—whether it represents the real-world system adequately for an intended use—and requires uncertainty, validation domains, and credibility evidence to accompany results [⁶][⁷]. The project’s “map of reality” should therefore be an explicitly purpose-bound, versioned ensemble of observations, models, alternatives, uncertainty, and unresolved questions, not an assertion of complete reality.

Evidence can underdetermine theory: more than one model may remain compatible with available observations [⁸]. Structural realism offers a relevant but noncanonical research stance: science may preserve or recover relational structure even when theories about the intrinsic nature of entities change [⁹][¹⁰]. This makes relational maps—dependencies, constraints, transformations, invariants, and causal effects—more defensible initial targets than claims about ultimate substance.

Core distinctions

Concept	Operational treatment
Structure	Elements, relations, constraints, state, behavior, boundary, environment, and lifecycle represented for a declared purpose.
Root-level structure	A structure or set of structures whose alteration materially changes the validity, operation, interpretation, or recoverability of dependent structures within a declared boundary and function.
Reality structure	A structure proposed to participate in, constitute, constrain, or represent the experienced world; the label does not by itself establish ultimacy.
Literal-reality claim	A claim intended as an assertion of existence, not merely metaphor, fiction, or instrumental representation.
Model	A selective representation used for explanation, prediction, exploration, or control [⁵].
Simulation	An implemented model that generates state transitions or behavior under specified rules; it may be offline, interactive, nested, stochastic, agent-based, or synchronized to a target.
Digital twin	A dynamically connected representation of a real-world entity or process, distinguished from a static model by synchronization and data exchange [¹¹][¹²].

Concept	Operational treatment
Virtual structure	A structure instantiated through a supporting medium; virtuality alone does not determine whether its objects or events are unreal. Chalmers, for example, argues that virtual worlds can be genuine realities, a philosophical position the project should record without canonizing [13][14].
Base reality	A hypothesized reality with no containing or governing parent structure. This is an ontological designation, not an operational status established merely by failure to observe a parent.
Map	A purpose-bound representation of selected entities, relations, observations, uncertainties, and model alternatives.
Failure	A defined loss, unacceptable state, constraint violation, or inability to perform a required function—not simply surprise or anomaly.

A root of trust illustrates functional rootness. NIST defines roots of trust as highly reliable components performing critical security functions from which trust is built, but the fact that they are inherently trusted within an architecture does not establish that they are ontological roots [15][16]. This distinction separates structural, causal, control, trust, epistemic, and ontological roots.

Root identification

Rootness is best represented as a multidimensional, boundary-relative classification rather than a universal binary label. Candidate dimensions are dependency reach, constraint reach, intervention asymmetry, observability asymmetry, substitution resistance, recovery dependence, persistence across resets, and sensitivity of downstream models to changes in the candidate.

A candidate should be investigated through controlled interventions where ethical and technically safe: isolate it, delay it, replace it, perturb outputs, withdraw maintenance, corrupt telemetry, or substitute an independent implementation. Causal inference treats interventions and counterfactuals as distinct from passive observation, making it a useful methodological foundation for testing whether a candidate actually controls downstream outcomes [17][18].

No single score proves ontological rootness. Results should produce a **root profile**: root of what, for which function, within which boundary, under which conditions, observed by whom, and with what alternatives.

Reality and simulation

Known engineered simulations offer privileged ground truth: investigators can inspect host, guest, scheduler, state store, resource allocator, reset functions, and observation channels. These environments permit real tests of whether embedded agents can infer parent structure from constrained observations.

The universal simulation hypothesis is different. Bostrom’s published argument is a conditional trilemma about extinction, posthuman simulation behavior, and the proportion of simulated observers; it is not empirical proof that this world is simulated [19][20]. Strong versions can become observationally indistinguishable from nonsimulation alternatives, creating a falsifiability and underdetermination problem [21][8].

The program should therefore test **specified simulation families**, not “simulation” without architecture. A claim such as “this environment uses a discrete lattice that produces signature X” is testable if it states discriminating predictions and rejection conditions. A claim that an undetectable simulator can produce any observation is not empirically discriminating.

“Real” and “simulated” should not be forced into opposites. An environment can be simulated relative to a host while remaining causally effective and experientially real for embedded agents. The research taxonomy should therefore separate existence, implementation medium, causal efficacy, observer experience, dependence, and ontological ultimacy.

Mapping under uncertainty

A reliable map is not necessarily complete. Reliability means that scope, provenance, confidence, validation domain, alternatives, and revision history are explicit. NIST guidance treats uncertainty as a state of knowledge and warns that multiple plausible models may need to be entertained and their model uncertainty propagated [^22].

The proposed Reality-Structure Map should contain:

- Observation ledger: what was measured, by which instrument, when, and under what conditions.
- Claim ledger: observation, inference, hypothesis, model assumption, design choice, or speculation.
- Structure registry: boundaries, elements, relations, constraints, state, behavior, lifecycle, and environment.
- Root-candidate registry: function-relative root profiles and supporting interventions.
- Model ensemble: competing models capable of explaining the same evidence.
- Simulation architecture registry: known or hypothesized host-guest relations, synchronization, permissions, and state transitions.
- Failure and anomaly ledger: discrepancies without premature causal assignment.
- Unknowns register: known unknowns, unresolved alternatives, and presently non-testable questions.
- Testability register: direct empirical, simulation-testable, analogical, philosophical, or presently non-resolvable.
- Change record: additions, removals, revisions, confidence changes, and dissent.

The map should preserve the difference between observed, inferred, constructed, reported, and hypothesized information. Project governance already contains a compatible principle: admission is not certainty, inference must remain distinguishable from observation, and uncertainty must survive processing [^23].

Failure as an instrument

Failure is useful because perturbation reveals hidden dependency, constraint, and recovery structure. It does not reveal ontology automatically. A shared outage can expose a common dependency; it cannot by itself establish that the dependency is the ultimate substrate of reality.

STAMP is relevant because it treats safety as an emergent control property and accidents as failures to enforce constraints across interacting technical and organizational components, rather than always as one linear “root cause” [24][25]. This cautions the project against assuming a single metaphysical root merely because investigators seek a root cause.

Controlled fault experiments should begin in isolated, reversible environments. Chaos-engineering practice begins with measurable steady state, introduces realistic disturbances, compares control and experimental conditions, and limits blast radius [26][27]. The project should adapt those principles to research sandboxes, not use them as authorization to interfere with real systems or reality.

Laboratory architecture

The proposed laboratory has six layers:

1. **Concept laboratory:** Formal vocabularies, ontologies, boundary declarations, causal diagrams, and competing definitions.
2. **Known-architecture laboratory:** Small deterministic and stochastic systems with complete designer access.
3. **Nested-simulation laboratory:** Host-guest stacks with intentionally limited embedded observers.
4. **Digital-twin and empirical-analogue laboratory:** Models synchronized with bounded physical or organizational targets; digital twins can support diagnosis, prediction, optimization, and low-risk design testing [28][29].
5. **Comparative-structure laboratory:** Cross-domain comparison of physical, biological, social, computational, conceptual, virtual, and simulated systems.
6. **Speculative-model laboratory:** Explicitly labeled exploration of base reality, unknown containing structures, and extra-ontological hypotheses, with no promotion to empirical finding without discriminating evidence.

Every experiment should declare system boundary, observer permissions, ground truth available to investigators, hypotheses, alternative explanations, intervention, control, measured outcomes, stop conditions, uncertainty, and interpretation limits. NASA’s verification-and-validation discipline should govern model claims [6][7].

Initial experiments

Experiment	Question	Primary output
Root Identification Blind Test	Can embedded agents distinguish local, parent, shared, telemetry, rule, and no-failure conditions?	Accuracy, calibration, localization distance, unknown-rate.
Observational Equivalence Test	Can different host architectures generate indistinguishable guest evidence?	Equivalence classes and discriminating observations.
Verifier Compromise Test	What happens when the diagnostic channel depends on the candidate being diagnosed?	False assurance rate and independent-witness requirements.
Root Migration Test	Can essential functions move between foundations while identity and service remain coherent?	Migration conditions, discontinuities, recoverability.
Emergent Root Test	Can local interactions produce a global constraint that becomes foundational?	Emergence criteria and intervention sensitivity.
Model-Failure Separation Test	Can agents separate world change from sensor, telemetry, and model failure?	Confusion matrix and evidence thresholds.
Lifecycle Withdrawal Test	How do maintenance cessation, resource starvation, state corruption, and rule change propagate?	Failure trajectories and survival/recovery envelopes.
Evidence Survival Test	Can records of foundational failure survive when the foundation controls state and history?	Witness independence and evidence durability.

What works

The program has a strong interdisciplinary core. Systems engineering supplies boundaries and lifecycles [1][2]; complexity science supplies emergence and cross-scale organization [3][4]; causal inference supplies intervention and counterfactual reasoning [^17]; modeling standards supply verification, validation, credibility, and uncertainty [6][7]; security supplies functional roots of trust and attestation analogies [15][30]; and safety engineering supplies constraint-based, non-linear accounts of failure [^24].

The project's refusal to claim a complete architecture of existence is a methodological strength. It allows empirically tractable work and speculative exploration to coexist without being assigned the same evidentiary status.

What does not work

The project should reject several shortcuts:

- Labeling a structure “root” because it is first, large, powerful, hidden, or narratively important.
- Treating “physical,” “real,” “literal,” “virtual,” and “simulated” as mutually exclusive synonyms.
- Treating unexplained anomalies as evidence of simulation or base-reality compromise.
- Assuming absence of detected parent control proves base reality.
- Treating a model's internal coherence as validation against a target.

- Assuming one root, one cause, or one level of explanation.
- Allowing simulated success to establish unrestricted claims about reality.
- Allowing AI consensus to substitute for evidence or ratification.

Current assessment

The project is now between the concept and defined-framework stages. It has a rich origin prompt, multiple AI interpretations, an existing CFTRL sub-dossier, a clarified parent scope, a strong foundational statement, preliminary terminology, and useful candidate experiments. It does not yet have a ratified ontology, formal evidence standard, implemented laboratory, benchmark corpus, validated metrics, or canon decision.

The existing CFTRL dossier should be preserved unchanged as historical and applied work. It should be reclassified as a downstream sub-dossier rather than rewritten to appear as though the original interpretation never occurred.

Development plan

Phase 0 — Provenance lock

Preserve the origin prompt, clarification directive, every AI draft, file hashes, dates, and status labels. Separate author statements from AI interpretations.

Phase 1 — Vocabulary and scope

Ratify provisional working definitions, define system-boundary syntax, and create the claim/evidence taxonomy. Record disagreements rather than forcing premature consensus.

Phase 2 — Mapping standard

Create templates for structure records, root profiles, reality claims, simulation architectures, failure records, model alternatives, and unknowns.

Phase 3 — Laboratory specification

Specify software architecture, ground-truth controls, observer permissions, experiment manifests, safety constraints, metrics, and reproducibility requirements.

Phase 4 — Pilot implementation

Build minimal known-ground-truth systems and run taxonomy calibration, Root Identification Blind Test, Model-Failure Separation Test, and Observational Equivalence Test.

Phase 5 — Comparative expansion

Add network, biological, organizational, security, and digital-twin analogues. Test which structural dimensions transfer and where analogies fail.

Phase 6 — CFTRL integration

Use validated results to revise the failure taxonomy, diagnostic protocol, root-attestation problem, lifecycle design principles, and Department of Endings charter.

Phase 7 — Review and ratification

Conduct independent team reviews, adversarial review, source audit, correction pass, and ONESTAMP decision. Ratification should apply to explicit provisions, not an undifferentiated document bundle.

Team operating model

The Architect remains the source-of-intent and sole ratification authority under the project's stated ONESTAMP practice. AI collaborators should produce independent analyses before reading one another where feasible, disclose uncertainty, distinguish evidence from synthesis, and preserve disagreements. The existing project charter's multi-AI process—independent drafts, exchange, comparison, consensus/divergence log, synthesis, and ratification review—is directly applicable [^23].

Recommended functional roles are research coordinator, provenance custodian, ontology lead, methods lead, simulation lead, evidence auditor, red-team critic, ethics/safety reviewer, synthesis editor, and change-log maintainer. One participant may hold multiple roles, but a model should not be the sole author, validator, and approver of the same claim.

Success criteria

The first successful release does not need to identify base reality. It should:

- Define terms without hiding disagreement.
- Represent multiple kinds of roots without conflation.
- Build provisional maps that preserve uncertainty and alternatives.
- Demonstrate repeatable root-identification experiments in known systems.
- Quantify diagnostic accuracy and calibration.
- Show at least one observational-equivalence case and state what evidence would discriminate it.
- Demonstrate model verification and bounded validation.
- Preserve source provenance and historical drafts.
- Produce documented negative results and testability boundaries.
- Feed validated findings into CFTRL without making CFTRL the parent framework.

Governing conclusion

The central product is not a declaration of reality's architecture. It is a disciplined framework for asking what structures are present, which relations are supported by evidence, what remains model-dependent, which foundations are operationally identifiable, how simulations can test limited claims, and how failure can expose dependency without being mistaken for proof of ontology.

That framework gives CFTRL a defensible foundation. It also produces independent value for simulation design, reality-model lifecycle management, digital twins, complex-system maintenance, epistemic integrity, resilience, and responsible treatment of unknowns.

References

1. [INCOSE Systems Engineering Handbook \[5 ed.\] 2023022915, 9781119814290, 9781119814306, 9781119814313 - DOKUMEN.PUB](#) - A comprehensive reference on the discipline and practice of systems engineeringSystems engineering p...
2. [About Systems Engineering](#) - General Definition. A system is a set of elements that work together to produce results that the ind...
3. [Emergence, \(Self\)Organization, and Complexity](#) - Tune in for the live stream on YouTube or Twitter. Abstract: In this second talk on balance, I will ...
4. [Emergence, Self-Organization, and Social Interaction](#) - Welcome to Santa Fe Institute.
5. [Models in Science](#) - Stanford Encyclopedia of Philosophy
6. [STANDARD FOR MODELS AND SIMULATIONS](#) - NASA
7. [FAQ](#)
8. [Scientific Underdetermination](#)
9. [Structural Realism](#) - Stanford Encyclopedia of Philosophy
10. [Realism and Theory Change in Science](#)
11. [Security and Trust Considerations for Digital Twin Technology](#)
12. [Digital Twins Workshops Summary Report](#) - [nvlpubs.nist.gov](#)
13. [Reality+: Virtual Worlds and the Problems of Philosophy](#) - Google Play - Reality+: Virtual Worlds and the Problems of Philosophy - Ebook written by David J. Chalmers. Read t...
14. [Reality+ - David Chalmers](#)
15. [Hardware root of trust - Glossary](#) - NIST CSRC - CSRC Home Page
16. [Roots of Trust | CSRC - NIST Computer Security Resource Center](#) - Modern computing devices consist of various hardware, firmware, and software components at multiple ...
17. [WHAT IS CAUSAL INFERENCE](#)
18. [CAUSALITY](#) by Judea Pearl - University of California, Los Angeles
19. [Are We Living in a Computer Simulation](#) - [scalar.usc.edu](#)

20. [Are We Living in a Computer Simulation? - Oxford Academic - Abstract](#). I argue that at least one of the following propositions is true: (1) the human species is ...
21. [PDF] [A Critical Engagement of Bostrom's Computer Simulation Hypothesis](#)
22. [PDF] [Simple Guide for Evaluating and Expressing the Uncertainty of NIST ...](#)
23. [UIARC Preliminary Information Collection Charter Perplexity.pdf](#)
24. [A Systems Theoretic Approach to Safety Engineering - Nancy Leveson](#)
25. [16.63 Lecture 9: Introduction to systems theory and STAMP](#)
26. [PDF] [Chaos Engineering - arXiv](#)
27. [Chaos Engineering: A Multi-Vocal Literature Review](#)
28. [Digital twins | NIST](#) - The Forecasting Revolution: Digital Twins and the Bottom Line
29. [Digital Twin Core Conceptual Models and Services](#) - To improve operational efficiency and enhance competitiveness in business, organizations are broadly...
30. [Remote Attestation Sets](#)