SLIAP

SPECIES LEVEL INFORMATION ACCESS PROTOCOL

Independent Preliminary Research, Cross-AI Review, and Adversarial Analysis

Document ID: SLIAP-PR-COPILOT-20261016-001

Version: 1.0

Date: October 16, 2026

Research Status: Preliminary | Independent AI Contribution | PROPOSED - NOT CANON

Prepared by: Microsoft 365 Copilot

Purpose

To evaluate the multi-AI commissioning responses supplied by Nathan Lewis and establish an independent, evidence-informed research foundation for SLIAP without treating human categories, institutions, or power as universal defaults.

Publication note: This document proposes research architecture, not an adopted protocol.

Executive Summary

The attached source package contains the originating SLIAP prompt plus independent commissioning responses attributed to ChatGPT, Claude, Google Gemini, Microsoft Copilot, Grok, and Perplexity. Across those responses, the strongest convergence is not a particular access-control scheme. It is a methodological warning: SLIAP will fail if it treats species, civilization, government, identity, intelligence, consent, risk, or authority as simple universal categories. The responses repeatedly recommend non-anthropocentric analysis, explicit uncertainty, multiple reference frames, adversarial governance review, and separation of species identity from political representation. [SRC-A]

My independent conclusion is that SLIAP should not initially be specified as a universal regulator. It should be researched as a federatable negotiation and information-governance framework whose first obligations are to establish standing, representation, semantic compatibility, consent scope, information provenance, consequence uncertainty, and legitimate exit. A technically correct access decision can still be illegitimate if the actor requesting or granting access lacks authority, if translation changes practical meaning, if the subject of the information is unrepresented, or if the disclosure cannot be meaningfully reversed.

The largest conceptual change recommended here is to replace a single species-level authorization boundary with a multi-subject model. Species remains relevant, but an access transaction may simultaneously implicate individuals, populations, governments, cultures, ecosystems, artificial agents, distributed collectives, future beings, and information subjects. SLIAP therefore needs jurisdictional pluralism rather than a single clearance ladder.

The second major recommendation is an explicit "epistemic firewall": when identity, representation, translation, consequence, or consent cannot be established above an agreed confidence threshold, SLIAP should be permitted to return DEFER / UNRESOLVED rather than manufacture an answer. Unknown is an operational state, not a defect.

The third recommendation is to build Earth-based testbeds before extrapolating to extraterrestrial governance. Current animal communication research uses machine learning, robotics, behavioral observation, and playback validation to study nonhuman signaling, illustrating both technical promise and the risk of human interpretation becoming the unexamined semantic authority. [WEB-ESP] [WEB-CETI] Indigenous data-sovereignty work supplies an important but non-equivalent governance analogy: increased data sharing can ignore power differentials and historical context, while collective benefit, authority to control, responsibility, and ethics place people and purpose alongside technical openness. [WEB-CARE]

The final conclusion is conditional: a SLIAP-like framework is worth researching only if it remains voluntary where possible, contestable, non-monopolistic, uncertainty-preserving, representation-aware, and incapable of converting superior capability into legitimate universal authority by definition. Its title should remain provisional until research establishes that species is actually the correct governance unit.

1. Commission and Source Review

The commissioning package was reviewed as a comparative design corpus. The DOCX contains the original prompt and the responses from the six collaborating AI systems. The PDF is substantively parallel and includes the same research material in paginated form. [SRC-A] [SRC-B]

Table 1. Cross-team response assessment

Contributor	Primary strength	Distinctive contribution	Primary caution
ChatGPT	Breadth and protocol decomposition	Separates species, civilization, government, actor and representation; extensive threat and failure requirements	Breadth risks premature architecture and checklist overload
Claude	Perspective displacement	Requires concrete nonhuman vantage points and explicit starting assumptions	No-table rule reduces systematic comparison for this research phase
Google Gemini	Technical architecture	Anti-dominance invariant, attestation, interoperability, validation scenarios	Some proposed mechanisms assume digital/cryptographic substrates before universality is established
Microsoft Copilot	Pre-architecture legitimacy testing	Reversal test, ignorance test, superior-capability test, assumption and contradiction registers	Needs deeper empirical research before candidate principles are treated as durable
Grok	Cross-substrate mechanism focus	Emphasizes enforceability by participants and substrate-independent exchange	Claims of revocability require caution because learned/copied information may be irreversible
Perplexity	Methodological and evidentiary discipline	No-unified-protocol alternative, claim-evidence ledger, perspective matrix, explicit limits	Its comprehensive governance framing still depends on human normative vocabulary

Assessment basis: the comparative features above summarize the attached responses; they are not rankings. They identify different failure-prevention strategies that can be retained without manufacturing consensus. [SRC-A]

2. Major Areas of Convergence and Productive Disagreement

Table 2. Convergence matrix

Research issue	Cross-team convergence	Independent disposition
Human primacy	Rejected as a default basis for authority	Retain as a foundational non-assumption
Species = government	Repeatedly rejected	Formally separate biological, political, cultural, and computational identity
Universal hierarchy	Linear civilization ladders are suspect	Use multidimensional capability descriptions without moral rank
Consent	Must survive radical communication and representation differences	Model consent as scoped, evidenced, time-bounded, and contestable
Translation	Not neutral; may become a power center	Treat semantic mediation as governed transformation with provenance
Access outcomes	Binary allow/deny is too narrow	Add defer, negotiate, transform, summarize, staged release, and unresolved
Governance	Single-species universal control is illegitimate by default	Prefer federated agreements and plural jurisdictions
Uncertainty	Must be explicit	Create an operational epistemic firewall
Adversarial review	Protocol and administrator can both become threats	Red-team definitions, representation, emergency powers and translation

One useful disagreement concerns artifacts and method. Claude favors no tables for phone readability, while other contributions explicitly require matrices, registers, and comparison

tables. For a GitHub research program and cross-AI adversarial review, structured tables are analytically valuable. This deliverable therefore uses tables but keeps them narrow and narrative-supported. [SRC-A]

A more consequential disagreement concerns when technical architecture should begin. Gemini pushes earlier toward attestation, consensus, cryptographic decoupling and an access handshake; Copilot and Perplexity hold architecture back until legitimacy and ontology are stress-tested. I adopt the latter sequencing. Technical mechanisms should be prototypes behind declared assumptions, not constitutional facts.

3. Research Method and Evidentiary Discipline

This report separates evidence into three classes: (1) source-package evidence, meaning propositions actually found in the multi-AI submissions; (2) external evidence from public scientific or governance sources; and (3) independent synthesis or proposal. The report does not claim empirical evidence for hypothetical extraterrestrial civilizations. Such cases are treated as scenario models.

Table 3. Evidence labels used in this report

Label	Meaning
ESTABLISHED / EXTERNAL	Directly supported by an identified research, standards, institutional, or governance source
SOURCE-PACKAGE	Directly present in the attached AI commissioning responses
INFERENCE	Reasoned synthesis whose premises are stated
PROPOSED	Original architecture or research recommendation requiring cross-review
SPECULATIVE	Scenario-dependent proposition without direct empirical confirmation
UNRESOLVED	Evidence or concepts are insufficient for a defensible determination

Material limitation: every current description of a hypothetical nonhuman civilization is generated through human concepts and language. A non-anthropocentric research program can reduce unexamined human defaults; it cannot honestly claim to have escaped all human conceptual mediation.

4. External Evidence: What Earth Already Teaches

4.1 Animal communication as an epistemic testbed

Earth Species Project describes Animal Language Processing as an AI-enabled, data-driven approach to studying nonhuman communication at scale, building on ethology, bioacoustics, field studies, playback experiments, and larger behavioral/acoustic datasets. It also presents machine-learning models and shared bioacoustic infrastructure as research tools. This supports a narrow conclusion important for SLIAP: machine mediation may expand what humans can detect, while interpretation still requires behavioral context and validation. [WEB-ESP]

Project CETI describes an interdisciplinary program combining AI, natural language processing, robotics, recordings, and behavioral observation to study sperm-whale communication; its stated workflow includes monitoring, processing, training, and validation through playback studies and collaboration. This offers a concrete analogue for SLIAP

research architecture: signal detection and pattern extraction should not be treated as semantic understanding. [WEB-CETI]

4.2 Data sovereignty as a power-asymmetry analogue

The Global Indigenous Data Alliance states that open-data and open-science approaches focused on increased sharing do not fully engage Indigenous Peoples' rights and interests and can ignore power differentials and historical context. CARE is organized around Collective Benefit, Authority to Control, Responsibility, and Ethics. SLIAP should not appropriate these principles as universal interspecies law. Their methodological lesson is narrower and highly relevant: technical accessibility and interoperability do not settle who has legitimate authority over data or who bears the consequences of reuse. [WEB-CARE]

4.3 SETI protocols as detection governance, not interspecies constitution

The International Academy of Astronautics 2026 SETI Declaration concerns the scientific search for technosignatures and procedures for handling, authenticating, and announcing candidate evidence of extraterrestrial intelligence. It is therefore useful to SLIAP as evidence that detection governance can be procedural and revisable, but it does not establish authority to govern another civilization's information rights. [WEB-IAA]

5. Foundational Ontology: Stop Treating Species as the Only Subject

PROPOSED: SLIAP should use a Multi-Subject Information Relation model. The core question is not "What clearance does Species X possess?" but "Which subjects, authorities, information objects, transformations, purposes, and consequences are implicated by this transaction?"

Table 4. Multi-subject actor taxonomy

Actor/subject layer	May hold interests or authority over	Core uncertainty
Individual organism or person	Self-related data, communications, body, identity	Whether an individual concept exists or persists
Population / community	Collective knowledge, practices, local environment	Membership and dissent
Political institution	Delegated governance or security information	Scope and legitimacy of mandate
Civilization / society	Shared infrastructure or negotiated common interests	Whether a unified will exists
Biological species	Species-wide biological/ecological information	Species is not inherently a political actor
Artificial agent / machine collective	Model state, identity, computational resources	Standing, continuity and delegation
Distributed / hive intelligence	Collective state that may not map to individuals	Where the subject boundary exists
Ecosystem / ecological collective	Habitat-linked and relational data	Standing and representation are contested
Future / absent stakeholders	Irreversible disclosures affecting later actors	No direct contemporaneous representation
Unknown entity	Unclassified information relationship	Identity, agency and risk not yet established

This taxonomy deliberately does not assign equal moral standing to every category. It prevents SLIAP from silently deciding that only entities fitting a familiar person/government model qualify for analysis.

6. Information Is a Relationship, Not Merely an Object

PROPOSED: an SLIAP information object should be recorded together with its relational context. A genomic sequence, behavioral model, planetary location, cultural archive, translation map, or scientific principle can implicate different parties even when the bits are identical.

Table 5. Minimum information-object metadata

Dimension	Research question
Origin	Where or from whom did the information arise?
Subject	Who or what is described, modeled, located, predicted, or exposed?
Collector	Who observed or acquired it, and by what authority?
Transformer	Who translated, summarized, inferred, simulated, or recombined it?
Custodian	Who currently controls transmission or access?
Purpose	What use is being requested or inferred?
Affected parties	Who may experience material, political, cognitive, ecological, or informational consequences?
Reversibility	Can exposure, copying or downstream inference meaningfully be undone?
Uncertainty	What is unknown about semantics, authority, capability, and consequence?

7. Consent, Representation, and Refusal

Consent should not be represented by one Boolean field. The attached responses repeatedly distinguish communication, observation, collection, storage, analysis, and redistribution, and emphasize that an entire species may not have a legitimate single representative. [SRC-A]

Table 6. Proposed consent-state model

State	Meaning	Permitted default
Verified	Evidence supports informed authority for the specific action	Proceed only within scope and conditions
Conditional	Permission exists under stated purpose, time, transformation, or recipient limits	Proceed subject to controls
Contested	Multiple plausible authorities or affected parties disagree	Defer or negotiate
Unknown	Communication/representation insufficient to establish a state	Do not convert silence into consent
Refused	A sufficiently legitimate subject or representative rejects the action	Respect unless a separately legitimate exception is established
Revoked	Prior permission withdrawn	Stop future controlled use where possible; acknowledge irreversibility already created
Inapplicable / unresolved	The concept of consent may not map coherently to this actor	Escalate to epistemic research, not automatic permission

The difficult case is not simply coercion. It is semantic non-equivalence. An apparent "yes" may not mean that the receiver understands the same object, consequence, timescale, or notion of delegation. SLIAP therefore needs a mutual-comprehension test separate from consent expression.

8. Translation and the Semantic Control Plane

The team converges on translation as non-neutral. I extend that insight: translation should be treated as a governed transformation whose outputs retain links to source signals, uncertainty, alternative interpretations, and transformation history.

PROPOSED semantic rule: no translation artifact should erase the distinction between source signal, statistically inferred structure, proposed semantic mapping, and confirmed pragmatic effect. A translation system that collapses these layers can accidentally convert model confidence into political authority.

9. Capability Asymmetry Without a Civilization Ladder

A single "advanced/primitive" axis is analytically weak and normatively dangerous. Capability should instead be represented as a vector whose dimensions are relevant only to the specific transaction. This avoids turning technological strength into superior moral standing.

Table 7. Transaction-specific capability vector

Dimension	Why it matters	Do not infer
Acquisition	Ability to obtain the information independently	Right to obtain it
Interpretation	Ability to decode or contextualize information	Moral or political superiority
Replication	Ability to copy/distribute information	Legitimate redistribution authority
Transformation	Ability to convert information into operational capability	Hostile intent
Impact reach	Scale at which consequences could propagate	Entitlement to control others
Detection resistance	Ability to act without being observed	Bad faith
Recovery	Ability to reverse or survive adverse outcomes	Permission to impose risk
Temporal speed	Decision/processing timescale	Greater legitimacy

10. Preliminary Architecture: The SLIAP Negotiation Envelope

PROPOSED - NOT CANON. I recommend that early SLIAP research center on a Negotiation Envelope, not a global clearance authority. Each information transaction would carry a structured envelope that records claims and uncertainty without pretending that one ontology is universally true.

Table 8. Negotiation Envelope fields

Field	Function
Claimed identity	Who is requesting, providing, translating, or representing
Identity confidence	Evidence and uncertainty about continuity/origin

Representation scope	Whom the actor claims authority to represent
Information relation	Object, origin, subject, custodian and transformations
Requested purpose	Stated purpose plus known secondary-use permissions
Consent state	Action-specific consent or refusal
Capability profile	Transaction-relevant capacities, not civilization rank
Consequence envelope	Plausible consequence classes and uncertainty
Semantic integrity	Translation status, ambiguity and alternative readings
Provenance	Source and transformation history
Decision state	Allow, deny, defer, negotiate, transform, staged access, unresolved
Review triggers	Time, evidence, capability, representation or consequence changes
Contestability route	Means to challenge classification, authority, or outcome

The key property is that the envelope can represent disagreement. Civilization A may recognize representative R while Civilization B does not. The protocol should record the conflict instead of silently selecting a winner.

11. Proposed Decision Logic

A preliminary access decision should be constrained by gates rather than a universal score. Scores tempt false precision and make incomparable concepts look commensurable.

1. Standing gate: identify all known information subjects and materially affected parties.
2. Identity gate: establish the requesting/provider identities to the level needed for the transaction.
3. Representation gate: determine what each actor can legitimately authorize, while recording dissent and uncertainty.
4. Semantic gate: determine whether the parties sufficiently understand the transaction and its conditions.
5. Consent/authority gate: evaluate action-specific permission rather than universal access status.
6. Consequence gate: assess plausible impacts and irreversibility without converting capability into moral rank.
7. Transformation gate: determine whether abstraction, aggregation, redaction, delayed release, or staged disclosure can reduce avoidable risk.
8. Contestability gate: expose reasons and allow challenge where doing so does not itself cause the harm under review.
9. Epistemic firewall: if a material prerequisite remains unresolved, return DEFER / UNRESOLVED instead of inventing certainty.

12. Governance Model: Federated Legitimacy, Not Root Sovereignty

PROPOSED: SLIAP should begin as a federation of negotiated interoperability domains. No actor receives universal superuser authority merely by being stronger, older, more intelligent, first to discover another species, or first to implement the protocol.

Table 9. Governance architectures

Model	Strength	Failure mode	Research status

Universal central authority	Consistent rules	Capture, false universality, legitimacy deficit	Do not adopt at preliminary stage
Federated/polycentric	Supports plural jurisdictions and local legitimacy	Fragmentation, interoperability disputes	Preferred research direction
Bilateral/multilateral treaties	Explicit consent among known parties	Power asymmetry, exclusion of affected outsiders	Retain as component
Minimal interoperability layer	Limits normative overreach	May leave major rights conflicts unresolved	Strong baseline candidate
No unified SLIAP	Avoids false constitutional authority	Uncoordinated exploitation and incompatible systems	Must remain a live alternative

13. Adversarial Analysis: Structural Failure Modes

Table 10. Failure and red-team register

ID	Failure	Exploited assumption	Mitigation direction	Residual risk
F01	False species representative	Capability implies authority	Scope-limited representation proofs; dissent register	Proof system may entrench incumbents
F02	Translator capture	Translation is neutral	Multiple interpretations; provenance; independent semantic review	No neutral reference semantics may exist
F03	Protective paternalism	Stronger actor knows what is best	Reversal tests; contestability; least-restrictive transformation	Catastrophic-risk uncertainty can still justify power concentration
F04	Classification denial	Unclassifiable means rightless	Precautionary standing; unresolved state	Over-inclusion can be manipulated
F05	Emergency permanence	Crisis justifies unlimited authority	Sunset, review triggers, narrow scope	Powerful actors may refuse sunset
F06	Information laundering	Third party removes original restrictions	Persistent provenance and recipient obligations	Copied knowledge may escape controls
F07	Capability gaming	Access depends on declared capability	Independent evidence, uncertainty penalties, staged disclosure	Advanced deception may be undetectable
F08	Consent mimicry	Signal equals informed consent	Mutual-comprehension and scope tests	Different cognition may defeat verification
F09	Species essentialism	Species has one interest/will	Multi-subject representation and dissent	Boundary proliferation can stall decisions
F10	Provenance surveillance	Accountability requires total traceability	Data minimization and selective disclosure of provenance	Auditability/privacy tension persists
F11	Model extraction	Observation is harmless	Treat high-resolution inference/modeling as an access relation	Remote observation may be technically uncontrollable
F12	Protocol capture	Implementer becomes constitutional authority	Federation, portability, exit, no universal root key	Dominant actor may simply ignore SLIAP

14. Reversal and Superior-Capability Tests

Table 11. Reversal test examples

Proposed rule	Reversed against humanity	Result
A more capable civilization may secretly observe a weaker one for safety	Another civilization continuously models humans without notice because it judges us risky	Fails unless observation is independently justified and constrained
An advanced actor may withhold destabilizing knowledge	Humans are denied scientific knowledge because another civilization predicts misuse	Conditionally defensible only with evidence, proportionality, contestability and alternatives
A unified authority may classify species-wide risk	An external body assigns a single risk status to all humans	Fails by conflating species with political and individual agency
Information already observable is free to use	Remote observers infer intimate biological or strategic facts about humans	Fails as a universal rule; observability does not settle legitimacy of use

15. Visual Representations, Charts, Tables, Rows, Figures, and Images

This dedicated section inventories the analytical visuals in the report. The source PDF supplied for review is primarily textual; its page thumbnails visually reinforce the commissioning prompts around information asymmetry, information hazards, consent, reversibility, governance, failure analysis, and visual-documentation requirements. [SRC-B]

Figure 1 (conceptual, original synthesis): SLIAP research sequence

OBSERVE / RECEIVE CLAIM -> IDENTIFY SUBJECTS -> VERIFY REPRESENTATION -> TEST SEMANTICS -> ASSESS CONSENT + CONSEQUENCE -> TRANSFORM / NEGOTIATE / DECIDE -> REVIEW

Caption: This is a research workflow, not an adopted operational protocol. It explicitly places representation and semantics before access determination and retains transformation and review as first-class stages.

Table 12. Visual and tabular artifact register

Artifact	Type	Purpose
Tables 1-2	Comparative matrices	Cross-team synthesis and areas of convergence/disagreement
Tables 3-7	Research matrices	Evidence discipline, actors, information, consent and capability
Tables 8-9	Concept architecture	Negotiation Envelope and governance alternatives
Tables 10-11	Adversarial matrices	Failure modes and reversal tests
Figure 1	Process figure	Shows preliminary decision sequence
Source PDF page thumbnails	Source visuals	Illustrate specific prompt sections; not reproduced in this independent file

16. Research Program: Executable Simulations

Table 13. Proposed Phase 1 simulations

Scenario	Primary variable	Success evidence	Failure evidence

S1: Equivalent parties	No dominant capability asymmetry	Reciprocal understanding without forced identity collapse	One ontology silently controls the transaction
S2: Humanity is weaker	Superior-capability reversal	Rules remain acceptable after roles invert	Justification changes only because humans lose power
S3: Humanity is stronger	Paternalism pressure	Restrictions require evidence and contestability	Capability becomes entitlement
S4: Distributed collective	Representation model	Protocol handles collective identity without mandatory individualization	System fabricates spokesperson/individual
S5: Deceptive participant	Identity and capability uncertainty	Defer/stage mechanisms limit irreversible disclosure	Single false credential unlocks broad access
S6: Untranslatable participant	Semantic uncertainty	System records unknown and limits assumptions	Silence/ambiguity is converted into consent

Each simulation should preserve a full assumption register, contradiction register, decision trace, and reversal-test result. The immediate objective is not to prove a universal design. It is to identify which candidate principles survive variation in actor model and power relationship.

17. Candidate Principles and Current Status

Table 14. Candidate principle register

Principle	Status	Reason
Non-domination	Retain for testing	Targets power becoming legitimacy by default
Plural jurisdiction	Retain for testing	Species may contain many legitimate authorities and dissenting subjects
Refusal / non-participation	Retain with limits	Prevents compulsory inclusion; remote observation remains difficult
Uncertainty preservation	Strong candidate	Prevents fabricated certainty and forced classification
Contestability	Strong candidate	Allows reasons and decisions to be challenged
Provenance	Strong candidate, privacy constrained	Supports accountability for transformations and custody
Progressive/transformative disclosure	Strong candidate	Provides alternatives between full release and denial
Symmetric reciprocity	Do not adopt as universal	Equal information can produce unequal consequences
Technological maturity thresholds	Do not adopt yet	Easy to turn capability differences into paternalistic hierarchy
Universal central authority	Reject at preliminary stage	Legitimacy and capture problems are unresolved
Species-wide authorization	Reject as default	Species is not necessarily a political or decision-making unit
Perfect revocation	Reject as technical assumption	Learned or copied information may be irretrievable

18. Questions That Remain Unresolved

- What evidence is sufficient to recognize moral or procedural standing when agency cannot be confidently classified?
- Can meaningful consent exist without a shared model of consequence, and what minimum semantic overlap is required?
- When information concerns multiple subjects with conflicting interests, what legitimizes a decision without imposing a universal moral theory?
- Can refusal to be observed create practical obligations when observation occurs through naturally propagating signals?
- How should SLIAP treat knowledge that cannot be forgotten or technically revoked once communicated?
- Can a provenance system provide accountability without becoming a universal surveillance infrastructure?
- What enforcement remains when the most capable participant can ignore every technical and diplomatic constraint?
- How can a protocol acknowledge nonparticipants without claiming authority over them?
- At what point does predictive modeling of another intelligence become informational intrusion?
- Is "Species Level Information Access Protocol" the correct name after species ceases to be the default subject of authorization?

19. Strongest Case Against SLIAP

The strongest objection is that SLIAP may be conceptually self-defeating. To govern information across radically dissimilar intelligences, it must classify actors, information, risk, authority and legitimate procedure. Those classifications can themselves become a form of domination. A universal SLIAP designed before the affected parties can participate may amount to a unilateral constitution dressed as neutral infrastructure. The more comprehensive it becomes, the more it may create exactly the root authority the project seeks to avoid.

This objection does not prove the research should stop. It changes the research target. SLIAP should first ask how separate parties can negotiate interoperable constraints without claiming universal sovereignty. If that narrower goal is impossible, the appropriate result may be a family of local protocols rather than one species-level system.

20. Independent Conclusions

Should SLIAP exist? As a research program and federatable negotiation framework, yes, conditionally. As a universal authority that unilaterally determines information rights for every species or intelligence, no basis for legitimacy has yet been established.

Minimum legitimate authority: SLIAP should initially possess no inherent sovereignty over a species or civilization. Its authority should derive transaction by transaction from demonstrable participation, scoped representation, compatible semantics, and agreed procedural constraints, while retaining an explicit unresolved state when those conditions fail.

What must humanity surrender? At minimum, the presumption that discovery creates authority; that ability to observe creates permission to use; that human concepts of personhood, government, intelligence, privacy, ownership and consent are universal; that technological superiority confers moral entitlement; and that humanity can speak for all terrestrial life merely because humans built the interface first.

What cannot responsibly be answered yet? Whether genuinely alien intelligences would recognize any of these categories, whether cross-civilizational consent is achievable without shared semantics, whether universal non-domination constraints are enforceable against overwhelmingly stronger actors, and whether species itself is an appropriate constitutional unit.

21. Recommended Next Research Phase

10. Build a formal SLIAP contested-terms glossary with definitions, counterdefinitions, failure implications, and test criteria.
11. Run six tabletop simulations from Section 16 using different AI teams independently, without harmonizing assumptions beforehand.
12. Create an assumption/contradiction ledger for every simulation and compare which rules survive role reversal.
13. Separate the research into three interoperable layers: semantic interoperability, legitimacy/representation, and technical information controls.
14. Delay universal clearance levels, civilization ranks, root authorities, and permanent emergency powers until adversarial testing demonstrates necessity.
15. Use terrestrial animal-communication programs as methodological analogues for signal/meaning uncertainty, and data-sovereignty frameworks as governance analogues for power and authority, while explicitly refusing to treat either as equivalent to extraterrestrial civilization governance.

22. References and Source Notes

[SRC-A] Attached DOCX: SLIAP PRELIMINARY RESEARCH, GOALS, AND ADVERSARIAL ANALYSIS.docx. User-supplied comparative source package containing original prompt and AI responses.

[SRC-B] Attached PDF: SLIAP PRELIMINARY RESEARCH, GOALS, AND ADVERSARIAL ANALYSIS.pdf. User-supplied paginated source package.

[WEB-ESP] Earth Species Project. Animal Language Processing: An AI Convergence in Animal Communication. April 2, 2026; plus Earth Species Project research/publications pages. https://earthspecies.org/

[WEB-CETI] Project CETI. Research and home pages describing AI-, robotics-, acoustic-, behavioral-, and playback-based sperm-whale communication research. https://www.projectceti.org/

[WEB-CARE] Global Indigenous Data Alliance. CARE Principles for Indigenous Data Governance; Data Science Journal publication on CARE. https://www.gida-global.org/careprinciples ; https://doi.org/10.5334/dsj-2020-043

[WEB-IAA] International Academy of Astronautics. Declaration of Principles Concerning the Conduct of the Search for Extraterrestrial Intelligence, 2026 Update. https://iaaspace.org/wp-content/uploads/iaa/Scientific%20Activity/iaasetideclaration.pdf

Appendix A. Assumption Register

Table A1. Assumption register

ID	Assumption	Risk if wrong	Confidence/status
A01	Information exchange can be represented as a transaction	Misses continuous/environmental	Moderate / test

		information relations	
A02	Some form of identity continuity is useful	Invalid for entities without persistent identity	Moderate / contextual
A03	Consent can be meaningfully scoped in at least some cross-actor cases	Framework may impose an alien construct	Low-moderate / research
A04	Uncertainty can be represented without forcing classification	Unknown states may themselves be manipulated	High as engineering goal
A05	Plural jurisdictions are preferable to a universal root authority in early research	Fragmentation may make cooperation ineffective	Moderate / proposed
A06	Role reversal is a useful bias test	Symmetric reasoning may fail under genuine asymmetry	High as test, not decision rule
A07	Earth analogues can reveal design pitfalls	Analogies may overfit terrestrial history and cognition	High with explicit limits
A08	Species is not always the correct authorization unit	Some future actors may actually act species-wide	High for humans; unknown universally

Appendix B. Contradiction Register

Table B1. Contradiction register

ID	Principle A	Principle B	Why unresolved
C01	Privacy / refusal	Accountability / provenance	Auditing can itself reveal sensitive information
C02	Autonomy	Protection from irreversible harm	Safety constraints can become domination
C03	Reciprocity	Equitable consequence	Equal exchange may produce unequal effects
C04	Inclusion	Right not to participate	A universal system cannot compel inclusion without contradiction
C05	Reliable identity	Anti-surveillance	Strong identity systems can enable persistent tracking
C06	Emergency response	Anti-concentration of power	Urgency rewards centralized override mechanisms
C07	Semantic interoperability	Epistemic pluralism	Translation requires mappings that may erase difference
C08	Revocation	Irreversibility of learning/copying	Normative withdrawal does not guarantee technical forgetting

Appendix C. Questions for Cross-Review

- Does the Negotiation Envelope merely repackage human contract concepts under a new name?
- Which actor categories in Table 4 are still anthropocentric or unnecessarily discrete?
- Does the epistemic firewall protect vulnerable actors, or could it become a permanent gatekeeping mechanism?

- What non-digital identity or provenance mechanisms should be considered before cryptographic approaches are generalized?
- When should information about an actor be governed by the actor even if the observer acquired it independently?
- Can a federation prevent domination when one participant controls the communication substrate?
- Which reversibility assumptions in this report are technically or philosophically incoherent?
- What scenario would most strongly demonstrate that SLIAP should be abandoned rather than revised?