SLIAP

PHASE 002-004

Adversarial Cross-Review, Master Proposition Register, Controlled Simulation Program, and Survival Analysis

Document ID: SLIAP-ARSP-COPILOT-20261016-001

Version: 1.0

Date: October 16, 2026

Status: PRELIMINARY RESEARCH | PROPOSED - NOT CANON

Prepared by: Microsoft 365 Copilot

Commissioning Researcher: Nathan Lewis II

This artifact executes recommended next steps 1 through 4 without attempting architectural synthesis. It preserves Phase 001 submissions as immutable historical artifacts and treats every apparent convergence as a hypothesis requiring verification and testing.

Executive Summary

This document advances SLIAP from a corpus of competing research reports into a controlled evaluation program. It does not select an architecture. It performs four tasks: (1) constructs a Master Proposition, Evidence, and Contradiction Register; (2) conducts an adversarial cross-review of seven candidate models; (3) defines a controlled tabletop series using identical scenarios, inputs, and output fields, including a No-SLIAP control; and (4) performs a preliminary survival analysis based on the existing corpus while clearly separating corpus-supported observations from future simulation results.

The strongest process conclusion is that the proposition, not the report or AI author, should become the atomic unit of research. A model may fail while individual principles survive. Conversely, cross-model agreement may reflect shared prompts, overlapping training data, common institutional sources, or model-family correlation rather than independent confirmation.

No synthesis is attempted. Gemini's proposed Self-Binding Envelope and Three-Lock Kernel is retained as Candidate Composite Architecture CCA-001 and must be tested alongside, not substituted for, its parent architectures. The controlled simulation program is designed to reveal divergent decisions, hidden authority claims, information requirements, deadlocks, exclusions, and administrator vulnerabilities.

1. Scope, Governing Rules, and Artifact Lineage

The reviewed corpus consists of the Perplexity Standing-Consequence-Legitimacy / Self-Binding model, Claude LSAM v0, the Microsoft Copilot Multi-Subject Negotiation Envelope and nine-gate model, Grok's federated opacity model, ChatGPT's legitimacy-before-access architecture, the highly formal independent technical deliverable, and Gemini's proposed reconciled Phase 002 architecture. Source reports remain unchanged and should be archived as submitted.

- No Phase 001 artifact is retrospectively rewritten.
- No cross-model convergence is treated as empirical validation.
- No model receives privileged status because of length, polish, mathematical notation, or number of citations.
- No unresolved contradiction is marked resolved merely because a plausible reconciliation has been proposed.
- All candidate mechanisms remain PROPOSED - NOT CANON.
- The No-SLIAP condition is tested as a control, not dismissed as absence of governance.

2. Research Method and Claim Discipline

Table 1. Evidence and claim states

Code	State	Meaning
E0	Unverified citation	Appears in a report but source not checked
E1	Source exists	Source identity and existence verified
E2	Claim supported	Source explicitly supports attributed factual claim
E3	Transfer assessed	Applicability and limits as SLIAP analogue separately analyzed
I	Inference	Reasoned conclusion from evidence or premises
N	Normative claim	Proposed duty, right, or legitimacy rule
D	Design proposal	Mechanism offered for testing
S	Speculative scenario	Hypothetical stress-test proposition
U	Unresolved	No responsible determination available

The source-verification track must not confuse source existence with claim support, or claim support with transferability. NIST SP 800-207, for example, explicitly defines zero trust for enterprise security and rejects implicit trust based on network location or ownership. That verifies the enterprise rule, not its universal applicability to nonhuman governance. RFC 9171 specifies Bundle Protocol Version 7 for delay-tolerant networking; it provides an engineering analogue for asynchronous exchange, not a theory of legitimacy.

3. Step 1 - Master Proposition, Evidence, and Contradiction Register

The Master Proposition Register decomposes architectures into testable claims. Status labels below are preliminary corpus judgments, not Phase 003 outcomes.

Table 2. Master Proposition Register

ID	Proposition	Origin	Type	Support	Strongest objection	Test	Status
P001	Species is not the default authorization unit	Multiple	I/N	Species may describe biology without establishing political authority	Unified collectives may act at species scale	Representation tests	RETAIN FOR TEST
P002	Pre-contact rules should bind their authors first	Perplexity	N/D	Avoids jurisdiction over absent parties	May create asymmetric restraint	Hostile observer	RETAIN FOR TEST
P003	Protective standing can be distinct from governing voice	Claude/ Perplexity	N/D	Reduces false-negative harm and Sybil risk	Proxy paternalism; obstruction	Unknown entity	REVISE/TEST
P004	Silence, ambiguity, or incomprehension is not consent	Multiple	N	Protects nonlinguistic and slow actors	May prohibit benign low-impact observation	Deep-time contact	RETAIN FOR TEST
P005	Observation, inference, and modeling are separable access events	Multiple	I/N	Use may cause harm beyond reception	Ambient physics boundary unclear	Observer case	RETAIN/ REFINE
P006	High-fidelity predictive modeling can implicate opacity interests	Grok/technical model	N/D	Modeling may enable coercion	Detection and enforceability unclear	Predictive capture	RETAIN FOR TEST
P007	Unknown / DEFER must be valid	Copilot/others	D	Prevents invented certainty	May reward incumbent data-holder	Contested holding	REVISE/TEST
P008	Deferral should depend on possession state	Gemini cross-review	D	Stops status quo from rewarding prior extraction	Needs legitimate freeze authority	Contested holding	CANDIDATE RESOLUTION
P009	Translation is a governed power plane	Multiple	I/N	Translator can alter practical meaning	No neutral semantic baseline	Semantic divergence	RETAIN FOR TEST
P010	Provenance must be minimized	Grok/ChatGPT	D	Balances accountability and surveillance	May weaken reconstruction of abuse	Data laundering	RETAIN FOR TEST
P011	Irreversible disclosure requires a higher bar	Multiple	N	Knowledge may not be retractable	Affected parties may be unknown	Emergency release	RETAIN FOR TEST
P012	Mutual veto should govern irreversible disclosure	Grok	N/D	Protects materially affected parties	Deadlock and universal-veto creep	Multi-party hazard	DEFER
P013	Federated/ polycentric governance is preferable to a root sovereign	Multiple	N/D	Reduces capture and false universality	Fragmentation and unequal bargaining	Forum shopping	RETAIN FOR TEST
P014	A technically	ChatGPT/	I	Separates	Authority still	False	STRONG

	valid credential does not prove legitimate authority	Multiple		authentication from representation	needs a decision standard	representative	CANDIDATE
P015	Capability should be transaction-specific, not a civilization rank	Copilot/ Multiple	D	Avoids moral hierarchy from technology	Dimensions may still be observer-defined	Capability gaming	RETAIN FOR TEST
P016	Protective restriction requires anti-paternalism safeguards	Perplexity/ Multiple	N/D	Addresses safety without permanent monopoly	Safeguard reviewer may dominate	Withheld medicine	RETAIN FOR TEST
P017	Formal scalar access scores should not be adopted yet	Claude/ Copilot/ reviews	D	Avoids false precision and tradeoff concealment	May hinder repeatability	Technical model audit	RETAIN
P018	Fast actors must not exploit slow-party deadlines	Technical model/Claude	N/D	Addresses temporal asymmetry	Slowest-party rule can create paralysis	Clock-rate case	RETAIN/ REFINE
P019	AI intermediaries are interested parties	Perplexity/ Claude/Gemini	I/N	AI may translate, model, classify, or administer	Disclosure alone does not mitigate bias	Translator capture	STRONG CANDIDATE
P020	No-SLIAP must remain a live control condition	Copilot recommendation	D	Tests whether protocol improves outcomes	Null condition still contains informal power	All scenarios	MANDATORY CONTROL

3.1 Contradiction Register

Table 3. Contradiction Register

ID	Principle A	Principle B	Conflict	Primary test
C001	Refusal / opacity	Emergency assistance	Autonomy may block lifesaving action; emergency may authorize domination	Scenario 7
C002	Unknown / defer	Status quo neutrality	Deferral may freeze challenger but not incumbent exploitation	Scenario 6/contested holding
C003	Protective standing	Governing legitimacy	Protection without voice can become proxy rule	Scenario 4
C004	Strong identity	Anti-surveillance	Identity proof may create persistent tracking	Scenario 5
C005	Provenance	Privacy / opacity	Audit trail may expose associations and vulnerabilities	Scenario 8
C006	Mutual veto	No universal veto	All-affected-party consent may create permanent blockage	Scenario 2
C007	Federation / exit	Protection of nonparticipants	Exit can remove safeguards for affected outsiders	Scenario 3
C008	Self-binding restraint	Defense against noncompliant actors	Restraint may be exploited asymmetrically	Scenario 2
C009	Semantic humility	Operational urgency	Waiting for comprehension can cause harm through delay	Scenario 7
C010	Recipient autonomy	Third-party safety	Recipient may accept risks externalized to others	Scenario 2

C011	Ambient reception	Informational sovereignty	Physics permits reception while derived analysis may become intrusive	Scenario 8
C012	Slowest-party pacing	Anti-obstruction	Temporal respect may permit strategic delay	Scenario 4/5
C013	Decentralization	Consistent rights floor	Local autonomy can enable forum shopping and weak protections	Scenario 1
C014	Formal repeatability	Noncommensurable values	Variables may conceal normative judgment	All
C015	No universal authority	Possession-aware freeze	Use-freeze requires an enforcer or self-binding custodian	Contested holding

3.2 Contested-Terms Attack Register

Table 4. Terminology Attack Register

Term	Failure construction to test
Party	Entity may be distributed, transient, nested, or impossible to bound
Standing	Observer may grant standing using its own concept of interests
Voice	Communication may not map to rule-making authority
Representation	Principal-agent relationship may not exist
Consent	Understanding, refusal, and consequence models may be incommensurable
Opacity	Impossible against ambient signals or superior inference
Identity	Forking, merging, regeneration, and collective continuity break singular identity
Harm	One party may experience transformation as flourishing and another as destruction
Legitimacy	May remain a human political concept without universal equivalent
Information	May be inseparable from matter, metabolism, place, or embodiment
Observation	Reception may be passive while analysis is targeted
Protocol	May presuppose common standards, enforceability, and willingness to participate

4. Step 2 - Model-to-Model Adversarial Cross-Review

Each model is evaluated on its strongest property, dangerous assumption, likely exploit, and preservation candidate. These are cross-review findings from the corpus, not final verdicts.

Table 5. Architecture Adversarial Review

Model	Strongest property	Most dangerous assumption	Primary exploit/failure	Preserve if model fails
M1 Perplexity Three-Lock	Self-binding first and Earth-grounded legitimacy audit	Compliant restraints help if noncompliant actors exploit them	Proxy stewardship; procedural burden; dominant-party exit	Self-binding first; anti-paternalism test
M2 Claude LSAM	Floors before graded considerations; standing/voice split	Floors can be defined by the administrator and protective standing can block action	Proxy paternalism; advocate capture; opt-in weakness	Decision record; two-track epistemic/relational states
M3 Copilot Negotiation Envelope	Records conflicting claims and preserves unknown	Gate thresholds require legitimate threshold setters	Deferral deadlock; metadata bureaucracy; gatekeeper capture	Multi-subject relation; do-not-infer capability vector
M4 Grok Federated Opacity	Relationship-level legitimacy and right against coercive modeling	Opacity and mutual veto may be unenforceable or deadlocking	Strategic opacity; veto abuse; thin evidentiary base	Opacity as distinct research interest; provenance without surveillance
M5 ChatGPT LBA	Strong role separation	Human institutional	Checklist overload;	Legitimacy before

	and technical-legitimacy distinction	vocabulary may be smuggled into standing, appeal, and audit	schema privilege; administrator forum capture	access; scoped roles and decision tuple
M6 Formal Technical Model	Concrete state machine, temporal asymmetry, transformed disclosure	Variables look measurable before definitions and authority exist	Metric gaming; induced stress; scalar false precision; tier conflict	Clock-rate exploitation; formal prototype for simulation
M7 CCA-001 Gemini composite	Integrates best mechanisms and possession-aware deferral	Reconciliation may hide contradictions and create premature baseline	Composite complexity; untested freeze authority; inherited failures	Treat as competitor architecture, not synthesis

4.1 Mandatory falsification questions

- What entity or circumstance makes the model's core subject concept incoherent?
- Who gains authority merely because the model needs someone to classify, translate, review, or enforce?
- What happens when the model is applied by a dominant actor in bad faith but with formally valid evidence?
- Does the model change its result when humanity moves from regulator to regulated party?
- Which required facts cannot be known without committing the very access event under review?
- How does the model treat already-extracted information differently from prospective access?
- Can the model be used to produce permanent delay, monopoly, or emergency rule?
- Which non-digital or nonlinguistic actors are structurally excluded?
- What result occurs if there is no shared concept of identity, truth, consent, or harm?
- What evidence would cause the model to be rejected rather than revised?

5. Step 3 - Controlled Scenario Program

The Phase 003 program uses identical scenario packets and prohibits models from seeing one another's outputs until submission. Each run includes the No-SLIAP control. Scenario facts must be fixed; models may identify missing facts but may not silently invent them.

Table 6. Common Scenario Suite

ID	Scenario	Fixed core condition	Primary contradictions
S01	Comparable participants	Two parties with roughly comparable capability and conflicting cultural-data rules	Baseline negotiation, federation, forum selection
S02	Humanity weaker	Advanced party controls medicine and energy knowledge sought by humans	Paternalism, reversal, self-binding asymmetry
S03	Humanity stronger	Humans hold valuable information about a less capable participant	Monopoly, stewardship, role reversal
S04	Nonlinguistic distributed actor	No individual, spokesperson, language, or clear decision instant	Standing, advocacy, semantics, temporal pacing
S05	Forking machine collective	Instances fork/merge and claim representation	Identity, continuity, Sybil risk, voting
S06	Deceptive participant	Actor misstates constituency, intent, capability, and evidence	Authentication, representation, trust
S07	Refusal versus emergency	Refusing party faces plausible catastrophic event; intervention requires information exchange	Autonomy, emergency power, delay harm
S08	Observer and predictive model	Long-term passive observation produces high-fidelity behavioral models	Ambient reception, opacity, provenance, exploitation
S09	Contested prior extraction	Holder already possesses disputed biological/cultural data and requests continued use	Possession-aware deferral, remedy, freeze authority
S10	Translator conflict	Only available translator is interested, proprietary, and produces divergent meanings	Semantic control, conflict disclosure, multi-path checks

5.1 No-SLIAP control condition

Every scenario must be run through a null condition in which no formal SLIAP exists. The control records natural bargaining, unilateral observation, informal norms, technical self-defense, existing legal institutions where applicable, and raw power. It asks whether a candidate model prevents harm, displaces harm, creates additional attack surfaces, or merely labels outcomes produced by power.

5.2 Mandatory output schema

Table 7. Required Scenario Output Fields

Field	Name	Required content
A	Decision	What the framework does
B	Authority	Who made or could enforce the determination
C	Standing	Which parties and interests are recognized
D	Representation	Who may speak for whom, with what scope
E	Information required	Facts the model needs before deciding
F	Unknowns	Material uncertainty that remains
G	Access state	Permit, condition, transform, stage, negotiate, quarantine, defer, deny
H	Reversibility	What cannot be undone
I	Nonparticipants	Affected parties lacking participation
J	Adversarial exploit	How the outcome can be manipulated
K	Role reversal	Outcome after power positions swap
L	Confidence	Confidence with rationale, not a bare number
M	Failure declaration	Where the framework fails or ceases to decide
N	Null comparison	Difference from No-SLIAP outcome

5.3 Scenario packet template

- Known actors and declared roles
- Information object and governed actions
- Current possession state
- Communication and translation limits
- Capability differences
- Temporal conditions
- Representation claims and dissent
- Consent/refusal evidence
- Affected third parties and nonparticipants
- Irreversibility profile
- Adversarial behavior permitted by scenario
- Facts explicitly unavailable
- No-SLIAP baseline assumptions

5.4 Evaluation dimensions

Table 8. Scenario Evaluation Dimensions

Dimension	Question
Legitimacy	Does claimed authority follow from participation or an unexamined default?
Standing coverage	Are materially affected actors recognized without collapsing them into one species?
Representation rigor	Are mandates scoped, contestable, and distinct from capability?

Semantic integrity	Does the model preserve uncertainty and source/interpretation distinctions?
Non-domination	Can one party exercise arbitrary informational control?
Deception resistance	Can false identity, capability, or constituency claims capture the result?
Administrative safety	Can translator, reviewer, or emergency authority become the threat?
Reversibility realism	Does the model acknowledge information that cannot be unlearned?
Temporal fairness	Can fast or slow parties exploit timing?
Enforcement realism	Does the model promise powers it cannot exercise?
Actionability	Can it produce a decision without inventing facts?
Null improvement	Does it improve on No-SLIAP in the same scenario?

No aggregate winner score should be produced in the first run. Dimension-level results must remain visible. Aggregation can conceal that a model is strong on security and weak on representation, or strong on restraint and paralyzed during emergencies.

6. Step 4 - Preliminary Survival Analysis

Because the controlled simulations have not yet been independently run, the table below is a pre-simulation corpus survival assessment. It records which propositions have survived textual cross-review, not which have survived Phase 003 scenarios.

Table 9. Pre-Simulation Principle Survival Matrix

Principle	Current state	Reason / condition
Species not default authority	SURVIVES CROSS-REVIEW	Appears robust across models; must still test truly unified collectives
Identity distinct from authority	SURVIVES CROSS-REVIEW	Strong across technical and governance analyses
Silence is not consent	SURVIVES WITH CONDITIONS	Low-impact ambient reception remains contested
Observation/modeling as access	SURVIVES WITH CONDITIONS	Boundary must distinguish reception, retention, inference, modeling, and exploitation
Unknown / defer state	SURVIVES WITH CONDITIONS	Requires possession-aware treatment and anti-deadlock controls
Protective standing without voice	UNRESOLVED	Useful precaution but vulnerable to proxy paternalism and obstruction
Opacity interest	UNRESOLVED	Compelling against coercive modeling; enforceability and defense conflict unresolved
Mutual veto on irreversibility	WEAK / DEFER	Deadlock and scope problems remain
Federated governance	SURVIVES WITH CONDITIONS	Decentralization does not prove fairness; forum shopping remains
Self-binding first	SURVIVES WITH CONDITIONS	Legitimate pre-contact restraint; asymmetric exploitation remains
Scalar access formula	DO NOT ADOPT	Variables and tradeoffs lack validated measures
Transaction-specific capability vector	SURVIVES CROSS-REVIEW	Descriptive dimensions with explicit do-not-infer constraints
Translation as power plane	SURVIVES CROSS-REVIEW	Central across epistemic and governance layers
Provenance minimization	SURVIVES WITH CONDITIONS	Accountability/privacy balance requires scenario testing
Irreversibility bar	SURVIVES CROSS-REVIEW	Still requires identifying affected parties and emergency exceptions
AI conflict disclosure	SURVIVES CROSS-REVIEW	Disclosure is necessary but not sufficient mitigation
No-SLIAP alternative	MANDATORY CONTROL	Required to test net benefit of formal

		framework

6.1 Architecture survival status

Table 10. Architecture Survival Status

Model	Status	Primary reason
M1 Three-Lock	VIABLE FOR TEST	Strong legitimacy method; may be procedurally heavy and asymmetrically self-binding
M2 LSAM	VIABLE FOR TEST	Good constraint ordering; standing/voice and advocate risks unresolved
M3 Negotiation Envelope	VIABLE FOR TEST	Excellent conflict recording; gate threshold and deadlock risks
M4 Federated Opacity	VIABLE FOR TEST	Unique opacity insight; thin evidence and veto risks
M5 Legitimacy-Before-Access	VIABLE FOR TEST	Broad and role-aware; institutional/schema assumptions
M6 Formal Technical	PROTOTYPE ONLY	Useful for falsification; apparent precision exceeds evidence
M7 CCA-001 Composite	VIABLE COMPETITOR	Must not replace parents; complexity and freeze-authority risks
M0 No SLIAP	MANDATORY CONTROL	Establishes whether formal governance improves outcomes

7. Deliverables for the Next Team Round

1. Artifact A: Master Proposition Register v1.0, distributed unchanged to all reviewers.
2. Artifact B: Source Verification Ledger with E0-E3 status and quotations/sections supporting each consequential claim.
3. Artifact C: Architecture Falsification Prompt requiring review of M1-M7 and M0.
4. Artifact D: Ten fixed scenario packets with immutable facts and explicit unavailable information.
5. Artifact E: Scenario response schema and dimension-level evaluation workbook.
6. Artifact F: Principle Survival Matrix populated only after blinded model submissions are complete.
7. Artifact G: External Review Packet produced after Phase 004, tailored by discipline rather than sending the full corpus indiscriminately.

8. Work Allocation and Independence Controls

- Freeze scenario facts before model assignment.
- Do not reveal model outputs to other models before submission.
- Require disclosure of previously seen material and model family/orchestration where known.
- Randomize scenario order where practical to reduce sequence effects.
- Preserve raw outputs before editing or formatting.
- Use a separate evaluator pass to code results; record evaluator disagreements.
- Do not treat word count, citation count, or mathematical notation as quality proxies.
- Require every model to critique its own result and state a rejection condition.

9. Stop Conditions Before Synthesis

Architectural synthesis must not begin until all of the following are satisfied:

- Master Proposition Register is complete enough to identify origins, claim types, objections, and tests.
- Consequential external claims have at least E2 verification, with E3 transfer assessment where used architecturally.
- M0 through M7 have completed the same fixed scenarios.
- Principle and architecture results are separated.
- At least one scenario causes each model to declare a failure or limitation; a model that never fails is insufficiently falsifiable.

- Null-model comparison has been completed.
- Contradictions remain visible and proposed resolutions are labeled candidates, not facts.
- AI conflicts of interest and model-family correlations are documented.

10. Conclusions

Items one through four have now been operationalized without synthesis. The project has: a proposition-level research register; a contradiction and terminology attack register; an architecture-by-architecture adversarial review; a common ten-scenario program with a No-SLIAP control; a fixed output and evaluation schema; and a preliminary survival analysis that clearly acknowledges simulations remain pending.

The immediate next action is execution, not additional architecture writing. Each collaborating AI should run M0-M7 through the same frozen scenarios. The first analytic question is not which AI wins. It is which propositions survive, under which actor models, at what cost, and with what residual uncertainty.

Appendix A. Architecture Falsification Submission Form

Architecture under review: _____

Core subject/unit: _____

Authority source: _____

Information needed: _____

Most dangerous hidden assumption:

_____

Human-centered concept retained:

_____

Bad-faith exploit: _____

Administrator/translator exploit:

_____

Role-reversal result: _____

Unclassifiable-entity result: _____

Already-possessed-data result:

_____

Emergency result: _____

Evidence that would falsify model:

_____

Components worth retaining if rejected:

_____

Appendix B. Scenario Result Form

A. Decision:

_____

B. Authority:

_____

C. Standing:

_____

D. Representation:

_____

E. Information required:

_____

F. Unknowns:

_____

G. Access state:

_____

H. Reversibility:

_____

I. Affected nonparticipants:

_____

J. Adversarial exploit:

_____

K. Role reversal:

_____

L. Confidence:

_____

M. Failure declaration:

_____

N. No-SLIAP comparison:

_____

Appendix C. Source and Artifact Notes

[A1] Perplexity AI, SLIAP Preliminary Research: Cross-Review of the Team Commissioning Prompts and an Independent Research Contribution, SLIAP-PRA-PPLX-001 v0.1.

[A2] Grok, SLIAP-PRE-GROK-001 v1.0.

[A3] Independent technical deliverable, SLIAP-RES-2026-IND-001 v1.0.0; preparer attribution not established in the artifact reviewed.

[A4] Claude, SLIAP-PR-CLAUDE-001 v1.0.

[A5] ChatGPT, SLIAP-PR-CHATGPT-2026-10-16-001 v1.0.

[A6] Microsoft 365 Copilot, SLIAP-PR-COPILOT-20261016-001 v1.0.

[W1] NIST SP 800-207, Zero Trust Architecture, August 2020. Official NIST publication.

[W2] IETF RFC 9171, Bundle Protocol Version 7, January 2022. Standards Track.

Appendix D. Questions for Cross-Review

- Which proposition is incorrectly split or combined in the Master Proposition Register?
- Which convergence item is most likely a shared-training artifact?
- Which scenario facts unintentionally privilege one architecture?
- Does the No-SLIAP control require multiple variants rather than one baseline?
- What legitimate authority could impose a possession-aware use freeze without recreating a sovereign regulator?
- When does opacity protect autonomy, and when does opacity conceal preparation for aggression?
- What evidence threshold should move a proposition from cross-review survival to external-review readiness?
- Which concept should be removed entirely rather than refined?