

THE THREE FOUNDATIONAL BOOKS

of the Simulation Environment

The Tome of Virtual Creation • The Book of Bots • Simulation Sessions

Unified Development Document — Version 2 (Ratified)

Author / Architect of Record: Nathan Ashton Arik Lewis (Nathan Lewis II)

Compiled by: Claude — Primary Audit & Documentation Layer, OAL Multi-AI Collaboration

Date Ratified: July 26, 2026

Status: Established Canon — Internal Designs In Progress

1. Provenance & Sourcing Note

This is Version 2 of the unified development document, incorporating structural and canon decisions ratified by the Architect of Record on July 26, 2026, in response to a clarifying-question set relayed from ChatGPT. Version 1 material remains valid; this revision layers ratified decisions on top of it and updates every downstream section — development path, feedback, open questions, and red team analysis — to reflect the new canon status.

Sourcing standard is unchanged from Version 1: material verified directly from Claude's own conversation search is presented as confirmed. Material reported only by other platforms and not independently verified remains flagged as such.

Verified primary source

Thread: “Cataloguing three development project objects” — June 1, 2026

<https://claude.ai/chat/47402415-044e-4368-a223-11103f41399a>

Reported but unverified (unchanged from V1)

- A March 22, 2026 “Unified Development State Report” and an “Infinite Virtual Library” — not located in Claude's own history.
- An April 18, 2026 “SimulationLab Master Prompt” — not located in Claude's own history.

These remain leads to check against the originating platform's history, not settled fact.

2. Canon Decision Log — Ratified July 26, 2026

The following eight decisions were made directly by the Architect of Record and are now established canon. They supersede any prior speculative material on these specific points.

Question	Ratified Decision
1. Collective name	The Three Foundational Books
2. Third book title	Simulation Sessions is the official title (no alternate designation adopted)
3. Canon status	All three books are established canon; internal designs remain under development
4. Authority	Architect of Record holds final ratification authority; designated and authorized AIs and management systems may make provisional changes
5. Document behavior	The books are literal virtual books within the application — both functional operational systems and canonical archives
6. Canon updates	The books may exchange information independently and autonomously update their own records (see Section 4 for the provenance/traceability mechanism this requires)
7. Scope	Expansive: universal and meta-universal hierarchies, structures, architectures, and any connected system — not limited to individual simulation environments
8. Bot inclusion	The Book of Bots covers artificial entities, robots, digital replicas, autonomous programs, simulated people, avatars, and non-AI beings

3. Concept Overview

The Three Foundational Books remain a single interlocking system with a clear division of responsibility, now formally ratified as canon and expanded in scope:

- The Tome of Virtual Creation governs what exists and what rules apply, across every environment and hierarchy the system connects to — not only a single simulation.
- The Book of Bots governs every entity operating inside what the Tome creates, now explicitly including non-AI beings, robots, digital replicas, and avatars alongside artificial and autonomous constructs.
- Simulation Sessions governs what happens when the system is exercised, and — per the ratified authority structure — may now generate provisional updates to the other two books directly, subject to traceability and Architect override.

The books are ratified as literal virtual books that are simultaneously functional information systems: they can be read as an in-universe artifact and queried as a live, updating registry. Their scope is now explicitly universal and meta-universal, covering connected architectures beyond any single simulation instance.

4. Book Profiles

4.1 — The Tome of Virtual Creation

Classification: Artifact / Foundational Object — Creation & Law Layer

Status: **Established Canon** — Internal Design In Progress

Description

The master creative artifact governing construction of virtual spaces, entities, and rules — now ratified to operate across universal and meta-universal hierarchies rather than a single simulation environment. It functions as world specification, schema registry, configuration authority, and canonical source of truth for anything declared to exist across any connected architecture.

Design Directions on File

- Functions as the meta-document underlying every connected environment, not just one instance.
- Tiered access model, with readable sections scaling to the rank or authority of the reader.
- Undecided physical form, now resolved in part: it is a literal virtual book within the application, which may also present as a living, self-updating document rather than a static text.

Remaining Design Items

- Precise field schema for a Tome entry (what a single “law” or “declaration” record must contain to be valid).
- How meta-universal scope is partitioned internally — does the Tome hold one global rule set, or nested rule sets per universe/hierarchy level?

4.2 — The Book of Bots

Classification: Artifact / Registry / Lore Object — Entity & Identity Layer

Status: **Established Canon** — Internal Design In Progress

Description

The official registry of every constructed or simulated entity connected to the environment. Ratified scope now explicitly includes artificial entities, robots, digital replicas, autonomous programs, simulated people, avatars, and non-AI beings — a substantially broader registry than “bots” alone, kept under its original title.

Design Directions on File

- Functions like a lorebook: entity traits, relationships, and key events, referenceable during design sessions or live encounters.
- Candidate entry fields: origin, creator, purpose, known attack vectors, alignment status, containment notes — now needing an added entity-type field given the broadened category set.
- The descriptive-vs-prescriptive question remains genuinely open: does the Book record what already exists, or define what an entity is permitted to be?

Remaining Design Items

- A classification taxonomy covering all seven now-ratified entity categories (artificial entity, robot, digital replica, autonomous program, simulated person, avatar, non-AI being), since each likely needs different tracked fields.
- Whether non-AI beings and simulated people are logged with the same authority/permission structure as constructed entities, or a distinct one.

4.3 — Simulation Sessions

Classification: System / Process / Recurring Event — Event & History Layer

Status: **Established Canon** — Internal Design In Progress

Description

The structured practice ground of the environment. Its title is ratified as simply Simulation Sessions, with no alternate formal designation adopted. Per the ratified authority structure, Sessions are no longer purely observational — they may generate provisional, traceable updates directly into the Tome or Book of Bots, rather than only submitting recommendations.

Design Directions on File

- Modeled on AI-powered scenario simulation: generating varied scenarios so decision-makers can rehearse choices and refine judgment in a risk-managed setting.
- Flagged connection to the Graendal psychological framework as an ideal environment for controlled pressure-testing; containment protocols for the Sessions themselves remain recommended.
- Recommended minimum structure: defined win / loss / ambiguous outcome states rather than open-ended play.

Remaining Design Items

- The exact provenance/traceability format a Session must attach to an autonomous update (see Section 5) so every canon change remains auditable back to its originating Session.
- Whether Session participation is voluntary or imposed — still open.

5. Cross-Book Integration & Recommended Development Path

With autonomous, independent updating now ratified, the single biggest structural requirement is a traceability layer — every provisional change made by an authorized AI or management system must be attributable, reversible, and distinguishable from Architect-ratified canon. The recommended pipeline is revised accordingly:

Tome / Bot declaration or Session outcome → Authorized system applies as PROVISIONAL → Change logged with source, timestamp, and originating authority → Architect review → RATIFIED (permanent) or REVERTED

Phase 1 — Provisional-Change Ledger (new, highest priority)

Before any book is allowed to update autonomously, build the ledger that records every provisional change: what changed, which book, which authorized system made the change, and a timestamp. Without this, autonomous updating is unauditably by design.

Phase 2 — Entry Schema

Lock the minimum entry schema for each book, now including the entity-type field for the Book of Bots and the meta-universal partition question for the Tome.

Phase 3 — Authorized-System Registry

Define which specific AIs and management systems are “designated and authorized” to make provisional changes, per the ratified authority decision. An open authority grant without a named list of who holds it is itself a gap the Red Team analysis flags below.

Phase 4 — Ratification Cadence

Establish how often and by what process the Architect of Record reviews the provisional-change ledger and converts entries to ratified canon or reverts them.

6. Feedback

What Works

- Ratifying canon status now, ahead of finished internal design, is the right call — it lets the collaboration team build with a stable shared reference instead of debating existence-of-concept every time it comes up.
- Separating final authority (Architect) from provisional authority (designated AIs/systems) is a clean, well-tested governance pattern — it enables velocity without surrendering control, provided Phase 1 and Phase 3 above are actually built.
- Expanding scope to universal/meta-universal architecture up front, rather than retrofitting it later, avoids a much larger rewrite down the line once content already assumes a narrower scope.

What Doesn't Work Yet

- “Designated and authorized AIs and management systems” is not yet a named list — as written, the authority grant is broad enough that it's currently unclear who is actually covered by it.
- Autonomous updating was ratified before the traceability mechanism that makes it safe exists. This is the single highest-risk gap in the current design (expanded on in Red Team, below).
- The Book of Bots now covers seven distinct entity categories under one title with no differentiated schema yet — a robot and a simulated person are unlikely to need identical tracked fields.

7. Remaining Open Items

The eight structural questions from the prior round are resolved (see Section 2). What remains are implementation-level items that follow from those decisions rather than open canon questions:

- Who specifically is on the authorized-systems list referenced in Decision 4?
- What is the required format for a provisional-change log entry (Phase 1)?
- Does the Tome hold one global rule set or nested rule sets per universal/meta-universal level (Decision 7 follow-on)?
- Does each Book of Bots entity category (Decision 8) get its own schema, or one shared schema with optional fields?

8. Red Team Analysis

The primary vulnerability introduced by this ratification round is that autonomous, independent updating was approved before any traceability or override mechanism exists to govern it. As written, any “authorized AI or management system” can alter canon directly, with no defined log, no defined reversal process, and no defined list of who actually holds that authority. This is a materially larger risk than the conflict-resolution gap flagged in Version 1 — that gap was about contradictions being hard to catch; this one is about canon being changeable by an unenumerated set of actors with no audit trail. Phase 1 (Provisional-Change Ledger) and Phase 3 (Authorized-System Registry) should be treated as prerequisites, not follow-on work, before autonomous updating is switched on in practice.

A second risk follows directly from the expanded scope: ratifying universal and meta-universal reach for all three books, before any of them have a locked schema, means every future schema decision now has to account for meta-universal complexity from day one. This is the right long-term call, but it raises the cost of getting Phase 2 wrong — a schema built assuming single-environment scope would need a full rewrite, not a patch, if discovered later.

A third, lower-severity risk sits in the Book of Bots: seven ratified entity categories sharing one registry, with independent autonomous update rights per Decision 6, creates a plausible failure mode where two authorized systems update conflicting facts about the same entity from two different Sessions simultaneously, with no defined precedence rule between them. This is a variant of the original Version 1 conflict-resolution gap, now sharper because updates can happen without an Architect in the loop at all.

9. Source & Provenance Log

Date	Thread / Material	Contents	Verification
Jun 1, 2026	“Cataloguing three development project objects”	Naming and initial catalogue of all three books	Verified — Claude
Jul 26, 2026	This conversation — ChatGPT question set + Architect responses	Eight structural canon decisions ratified (Section 2)	Verified — direct Architect ratification
Mar 22, 2026 (reported)	“Unified Development State Report”	Alleged early project listing	Unverified

10. Next Steps

- Build the Provisional-Change Ledger (Phase 1) before enabling autonomous updates in practice.
- Name the specific authorized-systems list (Phase 3).
- Lock entry schemas for all three books (Phase 2), including the seven-category Book of Bots taxonomy.