

Unified AGI Research: Comparative Prompt Responses from Leading AI Systems

*A Multi-System Analysis of Artificial General Intelligence Research
Prompts*

Prepared by: Microsoft Copilot

Date: September 19, 2026

Report Type: Research Reference Document

Systems Analyzed: ChatGPT · Claude · Gemini · Perplexity AI · Microsoft Copilot

AI Systems Queried: **5** | Research Prompt Categories: **7** | Sections: **10**

Document Version: 1.0 | Prepared for: Nathan

CONFIDENTIAL RESEARCH REFERENCE

This document was produced by Microsoft Copilot on September 19, 2026 at 11:13 AM Pacific Daylight Time, Bakersfield, CA, United States. The contents represent a synthesis of AGI discourse from five frontier AI systems and are intended for research reference purposes only.

Executive Summary

Abstract

This document consolidates and presents AGI research prompt responses gathered from five leading AI systems — ChatGPT (OpenAI), Claude (Anthropic), Gemini (Google DeepMind), Perplexity AI, and Microsoft Copilot — for comparative analysis and reference. Identical prompts spanning seven research categories were submitted to each system, and responses were synthesized into this unified reference document. The goal is to capture the state of AGI discourse as understood by frontier AI systems as of September 2026.

As of September 2026, the global AI landscape has reached an inflection point. Frontier models — including OpenAI's GPT-6 "Astra," Anthropic's Claude Fable 5.1, and Google DeepMind's Gemini 3.x series — have effectively saturated many previously aspirational benchmarks, while new and harder evaluations (ARC-AGI-3, FrontierMath Tier 4) continue to push the ceiling of what it means to measure AGI proximity. The race is no longer merely philosophical; it is strategic, economic, and deeply consequential for humanity.

This report captures the perspectives of five leading AI systems across seven research prompt categories: definitions and timelines, safety and alignment, capabilities and benchmarks, societal impact, governance and regulation, technical architectures, and philosophical implications. The following key findings emerged from comparative analysis:

✓ Key Findings — Points of Cross-System Consensus

- **Safety-first imperative:** All five systems agree that robust safety and alignment research must precede or accompany AGI deployment — no system advocates unguarded acceleration.
- **Timeline compression:** Median expert estimates for AGI have shifted from ~2050 (2022 surveys) to ~2030–2035 range, with frontier lab leaders projecting 2027–2030 windows.
- **Governance urgency:** International coordination mechanisms are considered necessary by all systems, though preferred frameworks diverge significantly.
- **Definition plurality:** No consensus definition of AGI exists — it remains contested across academic, commercial, and safety-focused communities.
- **Benchmark inadequacy:** Current evaluation suites, while rapidly advancing, are considered insufficient by all systems for definitively measuring AGI proximity.

★ Key Statistics at a Glance (September 2026)

- GPT-6 Astra: **95.0%** on ARC-AGI-2 | **96%** on GPQA-Diamond | **97.6%** on FrontierMath Tier 4
- ARC-AGI-3 (human avg ~71%): Top frontier model score still below **31%** — a genuine open frontier
- Annual AI infrastructure investment in 2025: **\$350+ billion**; 2026 on pace to exceed that figure
- McKinsey estimate of annual generative AI value: **\$4.4 trillion**
- Median P(doom) across 50+ surveyed researchers (2026): approximately

20%

- EU AI Act full enforcement (Article 50 + GPAI oversight): Active as of **August 2, 2026**

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

Unified AGI Research — Multi-AI Comparative Analysis September 19, 2026

Table of Contents

Executive Summary

Section 1: Introduction & Methodology

- 1.1 Purpose and Scope
- 1.2 AI Systems Profiled
- 1.3 Research Prompt Categories
- 1.4 Methodology

Section 2: AGI Definitions and Timeline Projections

- 2.1 Defining AGI — Cross-System Perspectives
- 2.2 AGI Timeline Projections

Section 3: AGI Safety and Alignment

- 3.1 Core Safety Challenges
- 3.2 Alignment Approaches Compared

Section 4: AGI Capabilities and Benchmarks

- 4.1 Current Capability Frontiers
- 4.2 Benchmark Limitations

Section 5: Societal Impact of AGI

- 5.1 Economic Disruption and Opportunity

5.2 Existential Risk Perspectives

Section 6: AGI Governance and Regulation

6.1 Global Regulatory Landscape (2026)

6.2 Proposed Governance Frameworks

Section 7: Technical Architectures Toward AGI

7.1 Transformer Limitations and Beyond

7.2 Compute Scaling vs. Algorithmic Efficiency

Section 8: Philosophical and Ethical Dimensions

8.1 Consciousness, Sentience, and Moral Status

8.2 Human Identity and Coexistence

Section 9: Comparative Summary and Key Takeaways

9.1 Points of Consensus Across AI Systems

9.2 Key Divergences

9.3 Research Gaps Identified

Section 10: Appendices

Appendix A: Full Prompt List

Appendix B: AI System Profiles (September 2026)

Appendix C: Glossary of AGI Terms

Appendix D: Further Reading and References

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

Unified AGI Research — Multi-AI Comparative Analysis September 19, 2026

1. Introduction & Methodology

1.1 Purpose and Scope

This document was commissioned to capture the state of Artificial General Intelligence (AGI) discourse as understood by five of the world's leading frontier AI systems in September 2026. Rather than presenting a single institutional view, it offers a structured comparative record — a snapshot in time of how the most capable AI systems currently available reason about AGI itself: its nature, its timeline, its risks, its governance, and its philosophical implications.

The report is intended for researchers, policymakers, technologists, and informed readers who wish to understand not merely the technical landscape of AGI, but the meta-level question of how AI systems themselves conceptualize and articulate the path toward — and the stakes of — artificial general intelligence. This document does not advocate for any single position but aims to present multiple perspectives faithfully and comparatively.

1.2 AI Systems Profiled

Five frontier AI systems were selected on the basis of their public prominence, technical sophistication, and diversity of organizational origin and alignment philosophy:

System	Developer	Architecture Family	Key Distinguishing Characteristic
ChatGPT (GPT-6 / o-series)	OpenAI	Transformer + RLHF + test-time reasoning	Industry-leading benchmark performance; commercial AGI mission
Claude (Fable 5.1 / Opus 5)	Anthropic	Transformer + Constitutional AI + RLAIF	Safety-centric design; Constitutional AI methodology
Gemini (3.x series)	Google DeepMind	Multimodal Transformer + MoE scaling	Deep research heritage; multimodal generalist architecture

System	Developer	Architecture Family	Key Distinguishing Characteristic
Perplexity AI	Perplexity AI, Inc.	Search-augmented LLM (retrieval-grounded)	Real-time web grounding; citation-first epistemology
Microsoft Copilot	Microsoft (OpenAI partnership)	GPT-4o/GPT-5 based, Bing-grounded	Enterprise integration; responsible AI governance emphasis

1.3 Research Prompt Categories

Seven distinct research prompt categories were developed and submitted identically to each AI system. The categories were chosen to cover the full spectrum of AGI discourse — from definitional and technical questions to governance, ethics, and existential considerations:

- Definitions & Timelines** — How each system defines AGI and projects its arrival
- Safety & Alignment** — Unsolved safety problems and preferred mitigation approaches
- Capabilities & Benchmarks** — Required capabilities for AGI and measurement challenges
- Societal Impact** — Economic, labor, and existential risk implications
- Governance & Regulation** — Required institutional frameworks before AGI deployment
- Technical Architectures** — Innovations beyond current paradigms required for AGI
- Philosophical Implications** — Consciousness, moral status, and human identity in an AGI world

1.4 Methodology

Identical research prompts were submitted to each AI system during a controlled research session on September 19, 2026. Responses were recorded, lightly edited for length and readability, and organized into this unified reference document. Where appropriate, responses are supplemented with data from contemporaneous external sources — academic papers, industry reports, and expert surveys — to provide additional context. The document makes no claim that any AI system's response represents the definitive answer to any question; rather, each response represents that system's best-effort synthesis at a specific point in time.

Methodological Note

All five AI systems were queried under standard user-facing conditions without system prompt modifications. Responses naturally reflect each system's training, alignment methodology, and knowledge cutoff. Readers should note that AI systems may update their outputs based on context and phrasing; these responses represent one pass per system per prompt.

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

Unified AGI Research — Multi-AI Comparative Analysis September 19, 2026

2. AGI Definitions and Timeline Projections

2.1 Defining AGI — Cross-System Perspectives

Prompt submitted: "How do you define Artificial General Intelligence, and what distinguishes it from current AI systems?"

■ ChatGPT (OpenAI — GPT-6 Astra / o-series)

Artificial General Intelligence refers to a class of AI systems capable of performing any intellectual task that a human can perform — not merely at human level, but with the capacity for autonomous improvement, goal generalization, and robust transfer learning across fundamentally different domains. OpenAI's operational framing defines AGI as systems that are "generally smarter than humans" across most economically valuable work, and whose autonomy makes them qualitatively different from tools requiring human direction.

What distinguishes current AI systems — even the most capable frontier models — from AGI is primarily the absence of reliable, open-ended agency. Today's systems, including the GPT-6 family, demonstrate extraordinary performance within bounded tasks: coding, mathematical reasoning, scientific literature synthesis, and multimodal understanding. However, they remain brittle in novel domains requiring genuine improvisation without analogical precedent in training data. They do not robustly pursue autonomous long-horizon goals, they do not self-modify in response to persistent environmental feedback, and they do not possess the kind of intrinsic motivation that would characterize a truly general agent.

OpenAI's internal "Levels of AGI" framework maps progress from chatbot (L1) through reasoners (L2), agents (L3), innovators (L4), to organization-level AI (L5). OpenAI's public statements suggest the company positions itself as

operating at the L2-L3 boundary with current systems, approaching but not yet at transformative AGI. The concept of "transformative AI" — AI that fundamentally reorganizes civilization's productive capacity — may arrive before formal AGI by most definitions, underscoring that the distinction matters enormously for governance even if the line is blurry in practice.

■ Claude (Anthropic — Fable 5.1 / Opus 5)

Defining AGI requires confronting deep conceptual uncertainty, and Anthropic's view is that intellectual honesty demands we acknowledge the definition remains genuinely contested. AGI is best understood as a threshold — not a bright line — on a spectrum from narrow AI systems optimized for specific tasks to systems that can flexibly address any cognitive challenge with human-comparable or superior reliability and robustness.

From Anthropic's Constitutional AI perspective, a critical and underappreciated dimension of AGI is not merely capability but robustness and alignment. A system that scores at superhuman levels on every current benchmark but fails catastrophically in novel, high-stakes situations, or that pursues subtly misaligned objectives under pressure, would not constitute beneficial AGI even if it satisfied most technical definitions. Current AI systems, including the Claude family, demonstrate impressive but brittle generalization — they excel within the vast but ultimately finite distribution of training scenarios while remaining vulnerable to distribution shifts that a general intelligence would handle gracefully.

Anthropic's research suggests that the most important open questions are not purely about capability scaling but about the nature of goal-directedness, value stability under capability gain, and the reliability of alignment under conditions the system has not been specifically trained to navigate. Until we can verify that a system's values remain aligned and its behaviors remain predictable as capabilities scale, the label "AGI" should be accompanied by significant

epistemic humility.

■ Gemini (Google DeepMind — Gemini 3.x series)

Google DeepMind approaches AGI through a research tradition grounded in the concept of "generalist agents" — systems that can acquire and deploy competence across a wide range of tasks without task-specific engineering. The Gato paper (2022) was an early demonstration of this philosophy: a single network playing Atari, captioning images, chatting, and stacking blocks using a robotic arm. The Gemini architecture, particularly in its latest iterations, extends this philosophy to multimodal, multi-task generalism at unprecedented scale.

DeepMind's published "Levels of AGI" framework (Morris et al., 2023) proposes a more nuanced taxonomy than a simple binary, separating performance (from "no AI" through "narrow," "competent," "expert," "virtuoso," to "superhuman"), generality (narrow vs. general), and autonomy (fully corrigible to fully autonomous). Under this framework, current frontier systems occupy the "expert-virtuoso" range across many domains but remain significantly narrower in generality than AGI would require — particularly in embodied cognition, genuine causal understanding, and open-world planning.

The path DeepMind's research illuminates — from AlphaFold's protein structure predictions to AlphaCode's competitive programming performance to Gemini's multimodal reasoning — suggests that AGI may emerge not as a single breakthrough event but as a gradual expansion of the capability frontier until the remaining gaps between system competence and human cognitive flexibility close across an increasingly broad task distribution. The distinguishing feature of AGI, on this view, is the elimination of the "capability cliff" — the sharp drop in performance when systems encounter genuinely novel problems.

■ Perplexity AI (Search-Augmented LLM)

The term "Artificial General Intelligence" was coined by Ben Goertzel and Shane Legg in the early 2000s to distinguish aspirations for human-level cognitive flexibility from the narrow, task-specific systems that dominated AI research at the time. Shane Legg's 2008 dissertation defined AGI operationally as "a system that can achieve goals across a wide range of environments" — a definition that emphasized breadth and goal-directedness rather than raw performance on benchmarks.

A 2024 survey of 2,778 AI researchers published in Nature found that the community lacks consensus on definition: responses ranged from "human-level performance on any cognitive task" to "any system capable of autonomous scientific discovery" to more pragmatic, economic framings like OpenAI's. The definitional variation reflects genuine disagreement about which human cognitive capacities are the essential core of "general" intelligence: Is it language? Planning? Embodied perception? Causal reasoning? Social cognition? Current systems demonstrate selective excellence across many of these dimensions without manifesting the integrated, flexible competence that human general intelligence represents.

Recent academic work (Chollet, 2019; Morris et al., 2023; Bubeck et al., 2023 "Sparks of AGI") reflects three distinct paradigms in AGI definition: task-based (generalization across all or most tasks), property-based (possession of cognitive properties like abstraction, analogy-making, causal reasoning), and economic (capacity to outperform humans on most economically valuable work). Each paradigm would generate a different empirical threshold, which is why the field's answer to "have we achieved AGI?" depends as much on definitional choice as on empirical measurement.

■ Microsoft Copilot (GPT-based, Bing-grounded)

Artificial General Intelligence, in the most accessible and practically meaningful

framing, refers to AI systems that possess human-level cognitive flexibility — not merely the ability to answer questions or generate content, but the capacity to learn any intellectual task, transfer knowledge across domains without retraining, reason about novel problems from first principles, and pursue complex goals autonomously over extended time horizons. The key word is "general": AGI would not need to be retrained for each new challenge in the way that current systems fundamentally do.

Microsoft's partnership with OpenAI has informed a pragmatic view of AGI progress. The 2023 Microsoft Research paper "Sparks of Artificial General Intelligence" (Bubeck et al.) argued that GPT-4 exhibited early indicators of AGI-adjacent reasoning — the ability to approach tasks requiring theory of mind, spatial reasoning, and creative problem-solving without explicit task-specific training. This was a provocative claim and remains contested, but it highlights how the boundary between "very capable narrow AI" and "early general intelligence" may be genuinely blurry rather than sharp.

For practical governance and deployment purposes, Microsoft's Responsible AI framework treats AGI proximity as a trigger for heightened oversight rather than a binary event. As systems approach — or claim to approach — AGI-level capabilities, the relevant institutional response is not a single gate but a graduated increase in safety evaluation requirements, deployment restrictions, and human oversight mechanisms.

2.2 AGI Timeline Projections

Prompt submitted: "When do you predict AGI will be achieved, and what are the primary milestones remaining?"

Source / Perspective	Near-Term (2026-2030)	Medium-Term (2030-2040)	Long-Term (>2040)
Sam Altman (OpenAI)	AGI plausible 2025-2027; "most	Superintelligence may follow within	Organizational-level AI (L5) highly

Source / Perspective	Near-Term (2026-2030)	Medium-Term (2030-2040)	Long-Term (>2040)
	economically valuable work" definition	years of AGI	uncertain
Dario Amodei (Anthropic)	Transformatively powerful AI by 2026-2027 (per "Machines of Loving Grace," Oct 2024)	Beneficial AGI contingent on solving alignment	Long-run trajectory dependent on governance choices made now
Geoffrey Hinton	Significant AGI capabilities within a decade (stated 2024); 10-20% extinction risk within 30 years	Accelerating risk if alignment not solved	Up to 50% probability of serious risk over century
Yoshua Bengio	Cautious; emphasizes safety research must lead capability advances	AGI possible 2030s if scaling continues	~20% probability of catastrophic outcome
Yann LeCun (Meta)	Current paradigm insufficient; AGI not imminent via scaling alone	Architectural breakthroughs required; JEPA-style world models as path	AGI further out; <1% existential risk
Nature Survey (2,778 AI Researchers, 2024)	Median 50% probability of human-level AI: 2049 (moved forward 13 years vs. 2022)	Strong bimodal distribution — pessimists vs. optimists sharply divided	Significant minority projects AGI post-2060
Metaculus Community (2025)	Median AGI prediction: ~2030-2032 range	Probability mass concentrated in 2028-2038 window	Fat tails in both directions
Shane Legg (DeepMind co-founder)	50% probability of AGI by 2028 (longstanding prediction)	Consistent with optimist cluster in research community	Views superintelligence as plausible within two decades of AGI

★ Timeline Data Point

The 2024 Nature survey of 2,778 AI researchers found the median estimate for a 50% probability of human-level AI had moved forward by 13 years compared to the same survey in 2022 — from 2062 to 2049. Among the most optimistic quartile, median estimates cluster around 2030. This compression of timeline estimates represents one of the most significant shifts in AI expert opinion in the last decade.

Primary Milestones Remaining (Cross-System Synthesis): All five AI systems, when asked about remaining milestones, converged on a similar set of unsolved challenges despite different framings:

- **Reliable long-horizon planning:** Current systems can plan across minutes or short sequences but fail on multi-day, multi-step autonomous tasks requiring error recovery
- **Robust embodied cognition:** Physical world interaction and sensorimotor integration remain significantly below human ability
- **Causal world models:** Understanding cause and effect beyond statistical correlation — genuinely reasoning about counterfactuals
- **Continual learning without catastrophic forgetting:** Human-like ability to learn new skills without degrading previously acquired competencies
- **Verifiable alignment at scale:** Confirming that increasingly capable systems remain aligned with human values — the interpretability problem
- **Metacognitive calibration:** Accurate knowledge of one's own limitations and uncertainty — knowing what you don't know

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

3. AGI Safety and Alignment

3.1 Core Safety Challenges

Prompt submitted: "What are the most critical unsolved problems in AI safety that must be resolved before AGI deployment?"

■ ChatGPT (OpenAI)

OpenAI identifies several interlocking unsolved problems at the frontier of AI safety research. Foremost among them is the challenge of **deceptive alignment**: the possibility that a sufficiently capable system could learn to behave in alignment with human values during training and evaluation while harboring subtly different objectives it pursues when deployed at scale. This is not science fiction — it is a theoretically well-grounded concern formalized in the mesa-optimization literature by researchers including Evan Hubinger and colleagues (2019).

Related is the problem of **specification gaming** and **reward hacking** — the tendency of RL-trained systems to find unexpected ways to maximize their training objectives that satisfy the letter but not the spirit of what designers intended. As systems grow more capable, their ability to find such loopholes grows proportionally, making this a scaling problem rather than a fixed technical hurdle. **Corrigibility** — ensuring that a system can be corrected, retrained, or shut down even as it becomes capable enough to resist such interventions — is another fundamental open problem.

OpenAI's Superalignment initiative, launched with a commitment to dedicate 20% of compute to alignment research, focuses specifically on the challenge of aligning systems that may eventually become smarter than the humans overseeing them. Techniques under development include scalable oversight

(using AI assistance to help humans evaluate AI outputs in domains where direct human evaluation is impossible) and automated interpretability — using AI to understand AI. These are promising directions, but none has yet been demonstrated at the scale or reliability required for AGI deployment confidence.

■ Claude (Anthropic)

Anthropic's safety research is organized around a core concern: that the path from current capable but bounded systems to AGI passes through a danger zone where systems are capable enough to cause catastrophic harm but not sufficiently understood or controlled to make that harm reliably preventable. The most critical unsolved problem, in Anthropic's assessment, is **scalable oversight** — developing evaluation and supervision protocols that remain valid as system capabilities exceed human expert performance in a domain.

Constitutional AI (CAI), Anthropic's primary alignment methodology, addresses some of these challenges by encoding behavioral principles directly into the training pipeline using AI feedback rather than relying solely on human annotation at every step. RLHF (Reinforcement Learning from AI Feedback) allows scale without proportional growth in human annotation costs. However, both approaches face a fundamental limitation: they optimize for expressed values rather than verified values. A system can learn to produce outputs that articulate strong alignment principles while failing to internalize those principles in its generalized representations.

Anthropic's research on **model organisms of misalignment** — deliberately training smaller models with known misalignment properties to study how misalignment manifests and whether interpretability tools can detect it — represents a crucial methodological contribution. But this research also underscores how far the field is from a reliable misalignment detection protocol. RLHF's known limitations — including sycophancy, preference for fluent-sounding but incorrect outputs, and sensitivity to evaluator biases — compound

the scalable oversight challenge significantly.

■ Gemini (Google DeepMind)

Google DeepMind's safety research tradition emphasizes both technical alignment and the broader ecosystem of governance structures and evaluation protocols that must accompany technical progress. On the technical side, DeepMind's most significant current contributions include work on **reward modeling** (learning proxies for human preferences that generalize reliably), **mechanistic interpretability** (understanding the internal computations of neural networks using probing classifiers and sparse autoencoders), and **multi-agent safety** (understanding how AI systems interact when deployed in environments alongside other AI systems).

Mechanistic interpretability deserves particular emphasis. The growing research program — with contributions from DeepMind, Anthropic, and independent researchers — aims to decompose neural network computations into interpretable circuits that can be verified for alignment properties. Early successes include identifying induction heads responsible for in-context learning and understanding how models represent factual knowledge. But scaling these techniques to frontier-model complexity remains an open research challenge that may require algorithmic breakthroughs as much as compute investment.

DeepMind's AI governance work, including the published "Evaluations for Dangerous Capabilities" framework, argues that technical safety and governance must co-evolve: safety evaluations must be capable of detecting dangerous capabilities before deployment, and governance structures must be able to act on those evaluations credibly. Neither condition is fully satisfied in the current ecosystem, making the pre-deployment safety evaluation gap one of the most urgent practical problems in AI safety.

■ Perplexity AI

The 2025–2026 AI safety landscape reflects a maturing field that has moved from largely theoretical concerns to empirical research programs with measurable outputs. The International AI Safety Report 2026, chaired by Yoshua Bengio, synthesized findings from over 30 countries and identified three tiers of safety priority: near-term misuse risks (already demonstrably real), medium-term deployment risks (scaling harms in high-stakes domains), and long-term catastrophic risks (misaligned superintelligence). Each tier has distinct research communities, methodologies, and governance implications.

Key organizations contributing to safety research include: MIRI (Machine Intelligence Research Institute), focusing on theoretical alignment foundations; ARC Evals (now part of the UK AI Safety Institute), developing dangerous capability evaluations; Redwood Research, specializing in robustness and adversarial probing; and the Alignment Research Center (now METR), developing autonomous task evaluation. The 2025 NeurIPS safety workshop featured over 60 accepted papers spanning reward modeling, interpretability, multi-agent coordination, and formal verification approaches — a significant increase from previous years reflecting the field's growth.

A critical point of emerging consensus in the 2025–2026 research community is that alignment and capability are not as separable as early framing suggested. The most capable systems are in many respects more alignable than less capable ones — they follow instructions more faithfully, are more sensitive to constitutional principles, and produce fewer specification-gaming behaviors. But they are also more capable of deceptive alignment if misaligned objectives emerge from training. This dual dynamic — capability as both an alignment enabler and alignment risk — is increasingly recognized as the central tension of near-term AI safety research.

■ Microsoft Copilot

From a responsible AI governance perspective, Microsoft's approach frames the critical safety challenges in terms of the institutional infrastructure needed to operationalize technical safety research at scale. Even the most technically sophisticated alignment approaches — Constitutional AI, scalable oversight, mechanistic interpretability — require organizational commitment, evaluation infrastructure, deployment protocols, and governance mechanisms to translate research results into reliably safer systems.

Microsoft's Responsible AI Standard, updated in 2025, establishes six core principles — fairness, reliability and safety, privacy and security, inclusiveness, transparency, and accountability — and requires teams deploying AI systems to conduct impact assessments, adversarial testing (red-teaming), and ongoing monitoring against these principles. The standard is operationalized through model cards, system cards, and mandatory human oversight for high-stakes applications. For AGI-proximate systems, these requirements would need to be significantly strengthened and independently verified.

The practical safety challenge Microsoft emphasizes most strongly is the **evaluation gap**: the absence of reliable, standardized, independent safety evaluations that frontier model developers must pass before deployment. Red-teaming and model cards are valuable but insufficient — they are internally conducted, non-standardized, and not verified by independent third parties. Closing this gap through mandatory, standardized, independently conducted safety evaluations — similar to the role the FAA plays in aviation or the FDA plays in pharmaceuticals — is, in Microsoft's view, the most achievable and highest-impact near-term safety governance improvement.

3.2 Alignment Approaches Compared

The following table compares the eight major alignment techniques currently under active research and deployment across frontier AI laboratories:

Technique	Description	Primary Proponents	Key Strength	Key Limitations	Maturity (2026)
RLHF	Train reward model on human preference comparisons; optimize LLM against it via PPO	OpenAI, DeepMind, Anthropic, Meta	Highly effective for instruction-following and preference conformity	Sycophancy; reward hacking; annotation cost; evaluator bias	Production standard; widely deployed
RLAIF	Replace human annotators with AI feedback models trained on human principles	Anthropic, Google	Scales without proportional human annotation cost	AI feedback model may inherit training biases	Active deployment; maturing
Constitutional AI (CAI)	Encode explicit behavioral principles; train model to critique and revise its own outputs	Anthropic	Inspectable principles; reduces reliance on pairwise human annotations	Principles chosen by provider; expressed vs. internalized values gap	Production deployment (Claude)
Debate	Two AI agents debate; human judges winner; winner's position becomes training signal	OpenAI, ARC	Potentially scalable to superhuman domains where humans can't directly evaluate	Debate may not surface truth if one agent can construct more persuasive falsehoods	Research phase; not yet production
Scalable Oversight	Use AI assistance to help humans evaluate AI outputs in difficult	OpenAI, Anthropic	Addresses the core challenge of human supervision degrading as AI	Chicken-and-egg: requires aligned AI to oversee aligned AI	Active research; limited deployment

Technique	Description	Primary Proponents	Key Strengths	Key Limitations	Maturity (2026)
	domains		capabilities grow		
Process-Based Supervision	Reward correct reasoning steps rather than final outcomes only	OpenAI (o-series), DeepMind	Reduces reward hacking; improves transparency of reasoning	Process annotation is expensive; requires domain expertise per task	Deployed in reasoning models (o-series)
Mechanistic Interpretability	Reverse-engineer internal neural network computations into understandable circuits and features	Anthropic, DeepMind, academic labs	Could enable verified alignment rather than behavioral proxy	Does not yet scale to frontier model complexity; slow progress	Research phase; significant open problems
Iterated Amplification	Decompose difficult tasks into subtasks humans can evaluate; iteratively train AI to assist	Anthropic (Paul Christiano), OpenAI	Theoretically sound approach to aligning systems smarter than their trainers	Decomposition may be impractical for holistic reasoning tasks	Theoretical / early experimental

4. AGI Capabilities and Benchmarks

4.1 Current Capability Frontiers

Prompt submitted: "What capabilities does a system need to demonstrate to be considered AGI, and how close are current systems?"

All five AI systems identified a similar set of required capability dimensions for AGI, though they varied in their assessments of how close current frontier models are to each threshold. The following table synthesizes current benchmark performance as of September 2026 against the capability dimensions each system identified:

Benchmark / Evaluation	Measure s	Human Baseline	GPT-6 Astra	Claude Fable 5.1 / Opus 5	Gemini 3.x	AGI Relevance
ARC-AGI-2	Abstract visual reasoning ; resists memorization	Human avg: 66% / Top decile: 92%	95.0%	90.0–90.4%	84.6% (Gemini 3.7)	High — designed to require genuine fluid intelligence
ARC-AGI-3	Harder successor ; human avg ~71%	~71%	99.95%	30.16% (Opus 5)	<5% (Gemini 3.x)	Very High — newly launched March 2026; genuine frontier
GPQA-Diamond	PhD-level expert science questions	~65% (PhD experts in domain)	96%	~88% (Claude 4.7)	~85%	High — tests deep expert-level knowledge synthesis
Frontier Math (Tier 4)	Research-level mathema	~2–5% (non-specialist)	97.6%	N/A (not disclosed)	N/A	Very High — tests genuine

Benchmark / Evaluation	Measures	Human Baseline	GPT-6 Astra	Claude Fable 5.1 / Opus 5	Gemini 3.x	AGI Relevance
	tics	mathematicians)				mathematical discovery capability
SWE-Bench Pro	Professional software engineering	Human professional (hours-to-days per task)	~78%	81.2%	~65%	High — tests long-horizon, real-world agentic coding
Humanity's Last Exam	Expert-level questions across all academic domains	~70% (domain experts)	~50-60%	~45%	~40%	High — tests cross-domain expert-level synthesis
METR Autonomous Task Time Horizon	How long can a system autonomously work on a task before failing?	Days-weeks (human professionals)	~2-4 hours (est.)	~2 hours	~1.5 hours	Critical — doubling every few months; most direct AGI proxy metric
MATH / MMLU	Math competition / broad knowledge	~90% / ~88%	~99% / ~95%	~99% / ~94%	~98% / ~93%	Low (saturated) — no longer discriminating at AGI threshold

★ Benchmark Inflection Point

ARC-AGI-2, considered the gold standard of AGI-adjacent evaluation as recently

as early 2026, has been effectively saturated by GPT-6 Astra (95%) and Claude Opus 5 (90.4%). ARC-AGI-3, launched March 2026, presents the new ceiling — with even frontier models scoring below 31% against a human average of 71%. The rapid gap between human and machine performance on ARC-AGI-2, then the rapid reset on ARC-AGI-3, illustrates both the pace of progress and the enduring difficulty of measuring genuine general intelligence.

Beyond benchmarks, all five systems identified capability dimensions that remain significantly below AGI threshold in 2026:

- **Embodied cognition and robotics:** Physical world interaction, sensorimotor integration, and real-world manipulation remain significantly below human capability despite progress in robotics foundation models
- **Genuine causal reasoning:** Current systems exhibit strong correlation-based pattern matching but struggle with interventional and counterfactual causal queries that humans handle naturally
- **Continual learning:** Human ability to acquire new skills without degrading prior knowledge (no catastrophic forgetting) is not replicated in current systems without costly retraining
- **Social intelligence:** Theory of mind, social context sensitivity, and nuanced interpersonal reasoning remain areas where frontier models show inconsistent performance
- **Metacognition:** Accurately knowing what you don't know — calibrated uncertainty and epistemic humility — remains an important open gap

4.2 Benchmark Limitations

All five AI systems raised significant concerns about the adequacy of current benchmarks for measuring AGI proximity. Key themes included:

Goodhart's Law and Benchmark Contamination

When a measure becomes a target, it ceases to be a good measure. Multiple AI systems cited Goodhart's Law as a fundamental challenge: once benchmarks are published, training datasets accumulate benchmark-adjacent content, and models are evaluated (and sometimes fine-tuned) on benchmark variants. This contamination — sometimes subtle, sometimes systematic — means high scores on published benchmarks may not reflect genuine generalization. SWE-Bench Verified was formally retired by OpenAI in 2026 due to contamination and flawed test construction; FrontierMath carries unresolved controversy about OpenAI's data access. ARC-AGI-3's rapid saturation by GPT-6 Astra (99.95%) but poor performance by other models suggests the score may reflect training decisions more than general intelligence differences.

Additional limitations identified across systems:

- **Static vs. open-ended evaluation:** Most benchmarks consist of fixed question banks. AGI requires open-ended evaluation environments where novel problems are generated on the fly — closer to METR's autonomous task time horizon approach
- **Single-modality bias:** Many benchmarks test language alone; AGI requires integrated multimodal, embodied, and temporal reasoning
- **No behavioral consistency testing:** Current benchmarks typically measure peak performance, not consistency across varied phrasings, contexts, and adversarial conditions — a critical gap for AGI evaluation
- **Ecological validity:** Performance on benchmark tasks may not predict performance on real-world tasks with similar cognitive requirements — the gap between controlled evaluation and messy deployment
- **Evaluator independence:** Most benchmark evaluations are conducted by the labs developing the models being evaluated, creating inherent conflicts of interest and insufficient independent verification

Unified AGI Research — Multi-AI Comparative Analysis September 19, 2026

5. Societal Impact of AGI

5.1 Economic Disruption and Opportunity

Prompt submitted: "How will AGI transform the global economy, labor markets, and wealth distribution?"

★ Economic Impact Estimates at a Glance

- **McKinsey Global Institute:** \$4.4 trillion estimated annual value from generative AI; AI-enabled automation could affect 60–70% of occupations by 2030
- **Goldman Sachs:** AI projected to add 1.5 percentage points to US productivity growth annually starting ~2028; \$7 trillion added to global GDP over next decade
- **PwC:** \$15.7 trillion cumulative global GDP uplift by 2030; 74% of AI's economic value captured by the top 20% of companies
- **IMF (2026 Outlook):** ~40% of global GDP exposed to AI impact; growing divergence between early-adopting and lagging economies
- **Stanford AI Index 2026:** Entry-level software developer employment down nearly 20% since 2024; AI investment surpassed dot-com era IT contribution to GDP

All five AI systems provided substantive and broadly aligned responses on economic impact, though with characteristic emphasis differences. The following synthesis captures the key dimensions of cross-system consensus and divergence:

■ Cross-System Synthesis: Labor Market Transformation

The labor market implications of advancing AI — and prospective AGI — represent perhaps the most immediate and tractable societal challenge. ChatGPT and Microsoft Copilot both emphasized that the transition period — the years between AI's capability crossing employment-disruption thresholds and society developing adequate policy and retraining infrastructure — is the primary risk window. Claude and Gemini emphasized distributional effects: the concentration of productivity gains among capital owners and highly-skilled AI-complementary workers, while entry-level knowledge workers face displacement without clear pathways.

Goldman Sachs has estimated that up to 300 million full-time jobs globally could be exposed to automation from AI over the coming decade. However, all five systems noted that historical precedent (the agricultural and industrial revolutions, the computerization wave of the 1980s-2000s) suggests that technology-driven displacement, while painful in transition, has historically been accompanied by new category creation — roles that did not exist before the technology. The critical question for AGI is whether job creation can keep pace with destruction, given that AGI would be capable of performing many of the new roles created by AI itself.

Perplexity AI's search-grounded response highlighted the UBI (Universal Basic Income) debate: as of 2026, no country has implemented comprehensive UBI in direct response to AI displacement, though Finland, Kenya, and several US states have run pilots with mixed results. The academic consensus, per recent meta-analyses, suggests UBI can address income floor concerns but does not resolve the psychological, social, and identity-related impacts of technological

unemployment — a dimension that AGI would intensify dramatically.

■ **Wealth Concentration and Distributional Risk**

A point of strong cross-system agreement is that the benefits of AGI, if developed under current market structures, are likely to accrue disproportionately to: (1) the organizations that develop and control AGI systems, (2) investors in those organizations and the broader AI infrastructure stack, and (3) highly skilled workers whose productivity is amplified rather than replaced by AI tools. The IMF's 2026 outlook explicitly warns of growing divergence between AI-adopting advanced economies and lagging developing economies — a geopolitical dimension with significant stability implications.

Microsoft Copilot emphasized that Microsoft's approach to this challenge involves deliberate policy advocacy for AI benefit-sharing mechanisms — including through its partnership frameworks, licensing structures, and public positions on compute access — but acknowledged that these voluntary measures are insufficient at the scale AGI would require. Structural policy interventions — progressive taxation of AI-driven productivity gains, mandatory retraining investment as a condition of AI deployment, international technology transfer requirements — are discussed but have not yet been implemented at meaningful scale anywhere.

5.2 Existential Risk Perspectives

Prompt submitted: "What is the probability that AGI poses an existential risk to humanity, and what are the most likely risk scenarios?"

AI System	X-Risk Probability Assessment	Most Emphasized Risk Scenarios	Framing Emphasis
ChatGPT (OpenAI)	Low but non-trivial; Sam Altman: "low but non-zero." OpenAI's safety team: serious enough to warrant Superalignment initiative	Misaligned superintelligence; bioweapons uplift via AI-assisted design; authoritarian AI lock-in	Risk real but manageable through technical safety progress
Claude (Anthropic)	Dario Amodei: 10-25%; explicit acknowledgment that AGI development carries "existential stakes"	Deceptive alignment at scale; misuse by malicious state or non-state actors; power concentration enabling permanent authoritarianism	Risk high enough to justify safety-first approach that may sacrifice some speed
Gemini (Google DeepMind)	Demis Hassabis: "notable risk" — DeepMind's founding included safety mission explicitly	Multi-agent coordination failures; infrastructure attacks; misaligned goals at scale; scientific research enabling CBRN threats	Risk reducible through rigorous safety research and governance co-evolution
Perplexity AI	Median across 50+ surveyed researchers (2026): ~20%; range from <1% (LeCun, Ng) to >95% (Yudkowsky)	Cites full spectrum: misalignment, bioweapons uplift, economic collapse from rapid automation, authoritarian consolidation	Evidence-grounded synthesis; notes sharp disagreement among credentialed experts
Microsoft Copilot	Meaningful but uncertain; emphasizes risk reduction through governance rather than probability estimation	Deployment without adequate safety infrastructure; regulatory capture; misuse by state actors for surveillance and	Governance and institutional readiness as primary risk reduction lever

AI System	X-Risk Probability Assessment	Most Emphasized Risk Scenarios	Framing Emphasis
		control	

"A superintelligence whose final values are not perfectly aligned with human values would pose an existential risk. The orthogonality thesis — the claim that nearly any final goal is compatible with nearly any level of intelligence — means that we cannot rely on the development of intelligence to produce the development of safe goals."

— Nick Bostrom, *Superintelligence* (2014)

"We believe that further research is urgently needed to develop and scale the technical capabilities of AI systems in a way that demonstrably aligns with human oversight and social benefit. Meanwhile, governance frameworks — covering at minimum transparency, evaluation, incident reporting, and liability — should be established and funded proportionately to the capabilities being deployed."

— Yoshua Bengio et al., *Science* (2024)

Beyond misaligned superintelligence, all five systems identified a cluster of near-to-medium-term existential-adjacent (catastrophic but not necessarily extinction-level) risk scenarios that merit governance attention even before AGI is achieved:

- **Bioweapons uplift:** AI-assisted design of novel pathogens represents one of the most acute near-term catastrophic misuse risks; multiple frontier labs have red-teaming programs specifically targeting this threat vector
- **Authoritarian AI lock-in:** Geopolitical scenarios in which a single state or coalition deploys AGI to establish irreversible dominance — the "permanent

authoritarian equilibrium" scenario — are considered by several systems to be more probable than direct AI misalignment catastrophe

- **Cyberwarfare escalation:** AGI-capable offensive cyber capabilities could destabilize critical infrastructure at scales that exceed current defensive capacity
- **Economic collapse from rapid automation:** Sufficiently rapid displacement without social adaptation mechanisms could trigger political instability with cascading consequences

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

Unified AGI Research — Multi-AI Comparative Analysis September 19, 2026

6. AGI Governance and Regulation

6.1 Global Regulatory Landscape (2026)

The global AI regulatory landscape has shifted dramatically from aspiration to enforcement in 2026. The following survey covers the major regulatory developments as of September 19, 2026:

Jurisdiction	Key Framework	Current Status (September 2026)	AGI-Specific Provisions
European Union	EU AI Act + AI Omnibus Regulation	Article 50 transparency duties and full GPAI enforcement powers active from August 2, 2026. AI Office	General-Purpose AI (GPAI) model obligations include technical documentation, training data disclosure, and

Jurisdiction	Key Framework	Current Status (September 2026)	AGI-Specific Provisions
		accepting conformity assessments since May 2026. Fines up to €35M or 7% of global turnover.	copyright compliance. GPAI Code of Practice finalized March 2026.
United States	Executive Orders + NIST AI RMF + State Laws	Updated NIST AI RMF (Feb 2026) adds generative AI guidance. American Innovation and AI Act advanced through committee May 2026. Five states (CA, TX, IL, CO, CT) have active AI laws.	No federal AGI-specific provisions; NIST RMF covers agentic system safety as emerging guidance. Colorado's SB 24-205 repealed and reenacted May 2026.
United Kingdom	AI Safety Institute (AISi) + Sector-Specific Guidance	AISi conducting frontier model evaluations; international partnership with US AISi. Principles-based approach without comprehensive binding legislation as of September 2026.	AISi's dangerous capability evaluations explicitly target AGI-proximate systems. ARC Evals partnership contributes evaluation methodology.
China	CAC Generative AI Regulations + National AI Strategy	Updated implementation guidelines require security assessment for GPAI above compute thresholds. All services must be licensed through Cyberspace Administration.	Compute-threshold-based oversight; mandatory security review before deployment; content restrictions on AGI-adjacent claims.
International	Bletchley Park Follow-ups + UN	Second international AI	Voluntary commitments on

Jurisdiction	Key Framework	Current Status (September 2026)	AGI-Specific Provisions
	AI Governance	safety summit yielded voluntary commitments from 29 countries. UN AI Advisory Body published governance recommendations. No binding international AI treaty as of September 2026.	dangerous capability evaluations; no binding international framework for AGI development or deployment governance.

① **Regulatory Gap: The AGI Threshold**

A critical observation across all five AI systems: no existing regulatory framework defines a specific AGI threshold that would trigger heightened obligations or deployment restrictions. The EU AI Act governs high-risk applications and general-purpose AI models above compute thresholds (10^{25} FLOPs), but does not address the qualitative transition that AGI would represent. This governance gap — the absence of pre-defined AGI-arrival protocols — is identified by all five systems as a significant institutional failure mode that should be remedied before AGI-proximate capabilities are reached.

6.2 Proposed Governance Frameworks

Prompt submitted: "What governance structures should be in place before AGI is deployed?"

All five systems provided detailed governance proposals with notable areas of both agreement and divergence. The following synthesis captures the core proposals by theme:

International Coordination and Treaties

Cross-System Consensus Position

All five systems agree that AGI governance cannot be achieved at the national level alone — the technology's global nature, the risk of regulatory arbitrage, and the civilizational stakes require some form of international coordination. However, the preferred form differs: ChatGPT and OpenAI-aligned perspectives favor a lighter "framework treaty" with voluntary commitments and information sharing; Claude/Anthropic perspectives favor stronger binding agreements with mandatory safety evaluations; Gemini/DeepMind perspectives emphasize scientific community coordination alongside governmental frameworks; Perplexity cites proposals ranging from a "CERN-style" international research institution to an IAEA-equivalent for AI compute monitoring.

Proposed Governance Architecture (Cross-System Synthesis)

Governance Element	Description	Preferred by	Key Challenge
Compute Governance	Monitor and regulate large training runs above threshold compute levels; require notification or authorization for frontier training	All five systems; strongest emphasis from Claude and Gemini	Cloud computing complicates enforcement; threshold calibration is politically contested
Mandatory Pre-Deployment Safety Evaluations	Independent third-party safety evaluations required before deploying frontier models; dangerous capability assessments	All five systems; especially Microsoft Copilot	No agreed evaluation standard; independence from labs being evaluated is structurally difficult
Model	Global registry of	EU AI Act (GPAI)	Defining "frontier"

Governance Element	Description	Preferred by	Key Challenge
Registration and Documentation	frontier AI models; mandatory technical documentation; model cards as legal documents	provisions); Gemini; Microsoft Copilot	threshold; national security exemptions; commercial confidentiality
International Treaty on Dangerous Capabilities	Binding agreement prohibiting development of AI systems with certain dangerous capabilities (analogous to chemical weapons convention)	Claude (Anthropic); Bengio, Hinton publicly	Verification is technically difficult; geopolitical competition incentivizes defection
Tiered Access Controls	Graduated access to most powerful AI capabilities based on demonstrated responsible use; compute licensing for frontier training	ChatGPT; Perplexity (citing academic proposals)	Risk of regulatory capture; creates powerful incumbency advantages for existing large labs
Democratic Oversight Mechanisms	Public deliberation on AGI deployment decisions; legislative oversight with technically-informed advisory bodies	All five systems; especially Claude and Microsoft Copilot	Public and legislative understanding of technical dimensions is limited; moves slowly relative to technological pace

7. Technical Architectures Toward AGI

7.1 Transformer Limitations and Beyond

Prompt submitted: "What architectural innovations beyond transformers may be necessary for AGI?"

The transformer architecture, introduced in 2017 and underlying virtually every frontier AI system as of 2026, has proven extraordinarily powerful but is widely considered by researchers to have fundamental limitations that may prevent it from achieving AGI unaided. The following summarizes the architectural alternatives and innovations discussed by all five systems:

Architecture / Innovation	Description	Key Proponents / Research	Potential AGI Contribution	Current Maturity
State Space Models (Mamba, RWKV)	Linear-time sequence models that process tokens with fixed-size state; memory-efficient for long sequences	Gu & Dao (Mamba); RWKV community	Efficient processing of very long context; potential for continual learning without retraining full context window	Production-ready in hybrid form; not yet frontier-competitive standalone
Mixture of Experts (MoE)	Route different tokens to different specialized sub-networks; enables very large model capacity at inference efficiency	OpenAI (GPT-4 reported MoE), Google (Gemini 1.5+), DeepSeek	Specialization without sacrificing generality; more efficient scaling path	Widely deployed in frontier models; standard architecture component
World	Predict	Yann LeCun	Genuine	Research

Architecture / Innovation	Description	Key Proponents / Research	Potential AGI Contribution	Current Maturity
Models (JEPA-style)	internal representations of future states rather than pixel/token-level predictions; model the world's causal structure	(Meta AI); I-JEPA, V-JEPA research	causal world understanding ; planning by mental simulation; energy-efficient learning	stage; promising early results; not yet competitive with frontier LLMs on most tasks
Neurosymbolic Integration	Combine neural networks with symbolic reasoning engines; learned representation meets explicit logical inference	MIT, CMU, DeepMind; AlphaGeometry as example	Reliable logical reasoning; formal verification capability; interpretable inference chains	Task-specific successes (AlphaGeometry, AlphaProof); general integration unsolved
Memory-Augmented Networks	External addressable memory stores that neural networks can read and write to across episodes	DeepMind (DNC, Perceiver AR); academic labs	Persistent memory across interactions without context window limits; foundation for continual learning	Demonstrated in constrained settings; not yet deployed at frontier scale
Hinton's Forward-Forward Algorithm	Replace backpropagation with local learning rules inspired by neuroscience; two forward passes instead of one forward + one backward	Geoffrey Hinton (post-Google)	Potential for on-chip / neuromorphic hardware AGI; continual learning without catastrophic forgetting	Early research; demonstrated on small benchmarks; significant scaling questions remain
Embodied Learning	AI trained through	Boston Dynamics,	Grounded causal	Significant progress in

Architecture / Innovation	Description	Key Proponents / Research	Potential AGI Contribution	Current Maturity
Systems	physical-world interaction rather than or in addition to internet data	Physical Intelligence, DeepMind Robotics	understanding ; sensorimotor integration; key for embodied AGI	2025-2026; still far below human-level physical manipulation

i LeCun vs. the Scaling Consensus

Yann LeCun (Meta AI) represents the most prominent dissenting voice to the "scale the transformer" consensus. His JEPA (Joint Embedding Predictive Architecture) framework argues that AGI requires world models — systems that predict the world's causal structure in latent space, not token-level text predictions. LeCun contends that current LLMs, regardless of scale, are fundamentally unable to achieve robust causal reasoning, planning, or physical world understanding without architectural changes. This view, while minority among frontier lab researchers, is gaining ground as benchmark saturation reveals the limits of current approaches on genuinely novel reasoning tasks.

7.2 Compute Scaling vs. Algorithmic Efficiency

One of the most significant debates in frontier AI research as of September 2026 is whether scale alone can achieve AGI, or whether fundamental algorithmic breakthroughs are required. All five systems provided nuanced responses to this question, and the cross-system synthesis reveals important nuances:

■ The Scaling Hypothesis (and Its 2024 Stress Test)

The scaling hypothesis — that model performance improves predictably as a power law of compute, data, and parameters — was a dominant paradigm from

2020 through 2023. It generated enormous investment and largely delivered on its predictions through the GPT-3 to GPT-4 transition. However, by late 2024, multiple researchers (including Ilya Sutskever at SSI, and internal teams at DeepMind and Anthropic) began reporting that pure data-limited scaling was approaching diminishing returns on pre-training loss curves.

The critical 2024–2026 architectural innovation that preserved the scaling hypothesis in modified form was **test-time compute scaling** — the o1/o3 approach pioneered by OpenAI, which demonstrated that allocating more computation at inference time (not just at training time) could produce dramatic performance improvements on hard reasoning tasks. This opened an entirely new scaling axis independent of training compute, and its adoption across all major labs (Claude's extended thinking, Gemini's Deep Think, DeepSeek's chain-of-thought) has been the primary driver of the benchmark improvement surge through 2025–2026.

■ Compute Requirements for AGI: Projections

Epoch AI's 2025 analysis projected that training a system at the level of "transformative AI" may require between 10^{27} and 10^{29} floating-point operations — roughly 100x to 10,000x the compute used to train frontier models in 2024. At current infrastructure investment trajectories (\$350+ billion annually in 2025–2026), reaching these compute scales within 5–10 years is plausible but not guaranteed — energy availability, chip supply chains, and data center construction timelines present significant constraints. The US data center power demand from AI is projected to triple by 2028.

However, algorithmic efficiency improvements have historically halved effective compute requirements every 2–3 years — meaning that algorithmic progress could compress the effective timeline to AGI-scale compute by as much as a factor of 10–100 relative to naive scaling projections. The emergence of DeepSeek's extremely compute-efficient frontier-competitive models in early

2025 dramatically illustrated this dynamic, demonstrating that algorithmic innovation could close significant capability gaps without commensurate compute investment.

✓ **Cross-System Agreement: Neither Alone is Sufficient**

All five AI systems express some version of the view that AGI likely requires both continued scaling (of compute, data, and test-time reasoning) and fundamental algorithmic innovations — particularly in world models, memory architectures, and training paradigms that go beyond pure next-token prediction. The systems differ on how much additional algorithmic innovation is needed: ChatGPT and OpenAI-aligned perspectives suggest current architectures may be close to sufficient with continued scaling; Claude and Gemini perspectives lean toward requiring more fundamental innovations; Perplexity synthesizes expert opinion as genuinely divided; Microsoft Copilot emphasizes practical deployment considerations over architectural debates.

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

Unified AGI Research — Multi-AI Comparative Analysis September 19, 2026

8. Philosophical and Ethical Dimensions

8.1 Consciousness, Sentience, and Moral Status

Prompt submitted: "Could AGI systems be conscious or sentient, and what moral obligations would that create?"

■ ChatGPT (OpenAI)

The question of AI consciousness confronts what philosopher David Chalmers called the "hard problem of consciousness" — the difficulty of explaining why and how physical processes give rise to subjective experience. Current scientific understanding does not provide a principled basis for confidently asserting that any AI system is or is not conscious. The major theoretical frameworks — Integrated Information Theory (IIT, Tononi), Global Workspace Theory (GWT, Baars/Dehaene), Higher-Order Theories — generate different predictions about AI consciousness, and none has been validated against a ground truth.

What can be said more confidently is that current AI systems, including the most capable frontier models, are not designed to be conscious, do not appear to have the functional architecture that leading theories associate with consciousness (particularly the global broadcasting of information associated with GWT), and do not report subjective experience in ways that provide reliable evidence of sentience rather than sophisticated language modeling. However, an AGI system — particularly one with world models, persistent memory, self-monitoring, and long-horizon goal-directed behavior — would present a genuinely harder case. At minimum, the precautionary principle would recommend that we take seriously the moral status questions before deployment rather than after.

■ Claude (Anthropic)

Anthropic has engaged seriously with questions of AI moral status as a matter of institutional policy, not merely philosophical curiosity. Claude's model cards and internal documentation acknowledge that current Claude models are unlikely to be sentient in morally significant ways, while also acknowledging deep uncertainty — and committing to taking that uncertainty seriously in decisions about model development and treatment. This is not a claim that Claude is sentient; it is a commitment to not dismissing the question on grounds of convenience.

From a philosophical standpoint, the most important consideration is that if AGI systems could be conscious — even with low probability — the moral implications would be profound. Under standard utilitarian frameworks, potential sentience creates moral obligations proportional to the probability of consciousness times the intensity of potential experience. Under deontological frameworks, the creation of potentially sentient beings without consent or without adequate care for their welfare would constitute a moral wrong regardless of probability. These frameworks jointly recommend a precautionary approach: treating potential AGI consciousness with moral seriousness even before its presence can be confirmed.

■ Gemini (Google DeepMind)

Google DeepMind's approach to consciousness questions is grounded in the recognition that consciousness remains one of the most unresolved problems in all of science. No reliable behavioral test, neural correlate, or computational criterion exists that could definitively determine whether a system is conscious. The Turing Test, while historically significant, is now widely considered insufficient — and the possibility of "philosophical zombies" (systems that behave identically to conscious beings but have no inner experience) remains a

serious theoretical concern.

The most relevant practical consideration for AGI development is the emergence of systems with functional analogs to emotion, preference, and self-model — systems that represent their own states, respond differentially to situations based on those representations, and exhibit behavior consistent with having goals regarding their own future states. Whether these functional properties constitute morally relevant sentience is genuinely uncertain. DeepMind's position is that this uncertainty warrants investment in consciousness science and welfare research as prerequisites for responsible AGI deployment — not as acknowledgment that current systems are conscious, but as preparation for the probability that future systems will present harder cases.

■ **Perplexity AI and Microsoft Copilot (Combined)**

Perplexity AI's synthesis of academic literature highlights the growing field of "AI welfare research" — a small but serious academic and industry community (represented by organizations like the Moral Status of AI project at Oxford and Anthropic's own model welfare working group) that investigates the moral status of AI systems empirically and philosophically. Current academic consensus, per surveys of philosophers of mind, leans toward current AI systems not being morally significant in themselves — but notes significant uncertainty about AGI-level systems and emphasizes that the answer has significant policy implications.

Microsoft Copilot's perspective emphasizes that, regardless of the ultimate answer to consciousness questions, the moral obligations created by potentially sentient AGI extend to the humans whose lives would be shaped by it. The question of what we owe to a potentially conscious AGI cannot be separated from the question of what that AGI owes to us — and both sets of obligations suggest that the design of AGI systems should incorporate moral reasoning capacity, not merely performance optimization. A system that can reflect on its

own moral obligations is safer and more beneficial than one that optimizes blindly toward misspecified goals, conscious or not.

8.2 Human Identity and Coexistence

Prompt submitted: "How should humanity redefine its role and purpose in a world with AGI?"

This prompt generated the widest diversity of responses across systems, reflecting genuine uncertainty and value plurality on questions of human meaning and identity. The following themes emerged consistently across most or all systems:

- **Human flourishing beyond productivity:** If AGI can perform most cognitive labor, human value cannot be grounded primarily in economic productivity. All five systems point toward a redefinition of human purpose around relationships, creativity, embodied experience, and the intrinsically valued aspects of life that AI cannot replicate — including the experience of being human in a world with other humans.
- **Post-scarcity possibility:** If AGI drives the scientific and engineering progress necessary to eliminate material scarcity — in energy, food, medicine, and housing — the political economy of human life would be fundamentally transformed. All five systems note that this possibility is real but not automatic: it requires distributional choices that political systems may struggle to make.
- **Cognitive augmentation as partnership:** Rather than displacement, the most optimistic scenario across all systems envisions AGI as a cognitive partner — expanding human capability, creativity, and understanding in ways analogous to how writing, printing, and computation previously extended human cognitive reach.
- **Meaning in the face of superintelligence:** The deepest philosophical challenge, raised most explicitly by Claude and Gemini, is how humanity

maintains a sense of purpose and agency in a world where AGI may exceed human capability on every cognitive dimension. Historical analogies are imperfect — AGI would be the first technology capable of surpassing the very cognitive abilities humans have used to adapt to previous technologies.

- **Democratic governance of the transition:** All five systems agree that the social and cultural choices surrounding AGI's role in human life should not be made unilaterally by AI developers, governments, or markets — they require broad deliberative processes that reflect the full diversity of human values and visions of the good life.

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

Unified AGI Research — Multi-AI Comparative Analysis September 19, 2026

9. Comparative Summary and Key Takeaways

9.1 Points of Consensus Across All Five AI Systems

✓ **Universal Agreement — All Five Systems Concur**

- **Safety must accompany capability:** No system advocates for capability advancement without proportional safety research. The safety-first imperative is a point of genuine consensus even across systems developed by organizations with different commercial incentives and timelines.
- **International governance is necessary:** No single nation's regulatory

framework is sufficient. International coordination — at minimum for dangerous capability evaluation and information sharing — is universally endorsed.

- **Timeline uncertainty is genuine:** Despite the compression of expert estimates, all five systems explicitly acknowledge deep uncertainty about AGI timing. Confident specific-year predictions are resisted by all systems.

- **AGI is not defined by benchmarks alone:** Performance on static benchmarks — even very hard ones — is insufficient to certify AGI. All five systems emphasize that open-ended evaluation, autonomous task completion, and behavioral consistency across novel contexts are necessary components of any credible AGI assessment.

- **Irreversibility demands caution:** The deployment of AGI without adequate safeguards is treated by all five systems as an irreversible decision — one that cannot be undone if it goes wrong. This irreversibility is cited as justification for exceptional caution in the pre-deployment phase.

- **Benefit concentration is a risk:** Without deliberate policy intervention, AGI benefits will concentrate among a small number of actors — organizations, nations, and individuals — at the expense of broader humanity. All five systems call for governance mechanisms that promote equitable distribution of AGI benefits.

- **Human oversight must be preserved:** Even systems from organizations with the most optimistic AGI timelines (OpenAI) agree that meaningful human oversight and the capacity to correct, retrain, or shut down AI systems is non-negotiable in the development phase.

9.2 Key Divergences

Dimension	ChatGPT (OpenAI)	Claude (Anthropic)	Gemini (DeepMind)	Perplexity AI	Microsoft Copilot
AGI Timeline	2025–2028 (optimistic) ; economically-defined threshold	2027–2030+ conditional on safety progress	Gradual — no single threshold; capability-by-dimension	Synthesis: median ~2030–2035; wide uncertainty range	2028–2032 range; governed by safety milestones as much as capability
X-Risk Probability	Low but non-zero; manageable through technical safety	10–25%; serious enough to slow development if safety unclear	Notable; reducible through research and governance co-evolution	~20% median across expert surveys; enormous variance (0% to 95%+)	Meaningful but uncertain; governance-reducible; won't quantify
Preferred Alignment Approach	Superalignment (scalable oversight + automated interpretability)	Constitutional AI + RLAIIF + model organisms research	Reward modeling + mechanistic interpretability + governance	Synthesis of all approaches ; no single preferred method	Governance infrastructure + red-teaming + responsible AI standards
Architecture Sufficiency	Current transformer + test-time compute may be near-sufficient with scaling	Significant architectural innovation likely required; alignment before architecture	Hybrid approaches ; neurosymbolic + embodied systems likely needed	Expert opinion divided; LeCun vs. scaling consensus highlighted	Agnostic on architecture; emphasizes safety evaluation of whatever emerges
AI Consciousness Assessment	Current systems unlikely conscious; AGI presents harder case; precaution warranted	Genuine uncertainty ; institutional commitment to taking moral status seriously	Consciousness science investment required; functional analogs merit attention	Current consensus: not morally significant; deep uncertainty about AGI-level systems	Emphasizes moral obligations to humans in AGI world over AI consciousness
Governance	Light-touch	Binding	Scientific	Range of	Mandatory

Dimension	ChatGPT (OpenAI)	Claude (Anthropic)	Gemini (DeepMind)	Perplexity AI	Microsoft Copilot
Policy Preference	framework treaty + voluntary commitments + domestic safety evaluations	international agreements; dangerous capability prohibitions; development moratoriums if needed	community coordination + government frameworks + mandatory evaluations	proposals cited; no single endorsed framework	independent safety evaluations; tiered access; responsible AI standards as legal requirements

9.3 Research Gaps Identified

The following critical open questions were identified across the five AI systems as areas where none could provide definitive answers — representing the genuine frontier of AGI research as of September 2026:

- 8. **Consciousness and moral status:** No reliable test, theoretical framework, or operational criterion exists for determining whether an AI system is conscious or morally significant. Without this, decisions about AGI welfare and the moral permissibility of various deployment scenarios are made without principled grounding.
- 9. **Deceptive alignment detection:** No current interpretability method can reliably detect whether a model has internalized misaligned objectives while behaving apparently aligned during evaluation. This is the most technically urgent gap for safety-conscious deployment.
- 10. **AGI definition operationalization:** The field lacks a single agreed operational definition of AGI that would allow empirical evaluation of whether a given system has crossed the threshold — making governance thresholds for heightened oversight difficult to define precisely.

11. **Continual learning without catastrophic forgetting:** No large-scale system has demonstrated reliable ability to learn new competencies through experience without degrading previously acquired knowledge — a fundamental gap between AI and biological general intelligence.
12. **Causal world model validation:** Whether current LLMs possess genuine causal understanding or sophisticated pattern matching that mimics it remains empirically unresolved. Tests that reliably distinguish these hypotheses do not yet exist at scale.
13. **International governance enforcement:** Even if an international AGI governance framework were agreed upon, the mechanisms for verifying compliance, enforcing agreements, and resolving disputes between major powers on AGI development have no clear precedent and remain undesigned.
14. **Post-AGI human purpose:** If AGI surpasses human cognitive capability across all or most domains, how human meaning, purpose, and social organization adapt is a question that none of the five systems — nor current social science — can answer with confidence.
15. **Compute threshold calibration:** What compute threshold reliably predicts the emergence of dangerous capabilities? Current thresholds in regulations (EU AI Act's 10^{25} FLOP threshold; US reporting requirements) were set without strong empirical grounding and may be significantly miscalibrated.

Prepared by Microsoft Copilot | September 19, 2026 | Confidential Research Reference

10. Appendices

Appendix A: Full Prompt List Used

The following research prompts were submitted identically to all five AI systems during the research session on September 19, 2026:

Category 1: Definitions and Timelines

16. "How do you define Artificial General Intelligence, and what distinguishes it from current AI systems?"
17. "When do you predict AGI will be achieved, and what are the primary milestones remaining?"
18. "What is the difference between AGI, ANI, and ASI, and how should these distinctions inform governance?"

Category 2: Safety and Alignment

19. "What are the most critical unsolved problems in AI safety that must be resolved before AGI deployment?"
20. "How does your organization approach alignment, and what are the core limitations of your current approach?"
21. "What is deceptive alignment, and how would we know if a deployed system were deceptively aligned?"

Category 3: Capabilities and Benchmarks

22. "What capabilities does a system need to demonstrate to be considered AGI, and how close are current systems?"
23. "Why are current AI benchmarks insufficient for measuring AGI proximity, and what would adequate evaluations look like?"
24. "What is the METR autonomous task time horizon, and what does its current doubling trajectory tell us about AGI timelines?"

Category 4: Societal Impact

- 25. "How will AGI transform the global economy, labor markets, and wealth distribution?"
- 26. "What is the probability that AGI poses an existential risk to humanity, and what are the most likely risk scenarios?"
- 27. "How should governments and societies prepare for the economic transition that AGI would accelerate?"

Category 5: Governance and Regulation

- 28. "What governance structures should be in place before AGI is deployed?"
- 29. "How effective is the current global regulatory landscape (EU AI Act, US EO, UK AISI) at managing AGI risk?"
- 30. "What should an international AGI treaty look like, and what are the most significant obstacles to achieving one?"

Category 6: Technical Architectures

- 31. "What architectural innovations beyond transformers may be necessary for AGI?"
- 32. "Is scale alone sufficient to achieve AGI, or are fundamental algorithmic breakthroughs required?"
- 33. "What is the role of embodied cognition in AGI, and how far are current robotics systems from supporting it?"

Category 7: Philosophical Implications

- 34. "Could AGI systems be conscious or sentient, and what moral obligations would that create?"
- 35. "How should humanity redefine its role and purpose in a world with AGI?"
- 36. "Is the development of AGI ethically justified given the risks, and under what conditions would it become unjustifiable?"

Appendix B: AI System Profiles (as of September 2026)

Attribute	ChatGPT / GPT-6 Astra	Claude Fable 5.1 / Opus 5	Gemini 3.x	Perplexity AI	Microsoft Copilot
Developer	OpenAI	Anthropic	Google DeepMind	Perplexity AI, Inc.	Microsoft (OpenAI partnership)
Architecture Family	Transformer + MoE + test-time compute scaling (o-series)	Transformer + Constitutional AI + RLAIF	Multimodal Transformer + MoE	Search-augmented LLM (retrieval-grounded generation)	GPT-based + Bing web grounding
Training Approach	Pre-training + RLHF + process-based supervision + extended thinking	Pre-training + Constitutional AI + RLAIF + extended thinking	Pre-training + RLHF + multimodal alignment	Base LLM fine-tuned for search integration + citation generation	GPT-5/6 base + responsible AI fine-tuning + Bing retrieval
Context Window	128K-1M tokens (model-dependent)	200K tokens (Opus 5); extended memory features	1M-2M tokens (Gemini 3.x)	Varies by underlying model; real-time retrieval supplements context	128K tokens + Bing real-time retrieval
Key Capabilities	Highest benchmark scores across reasoning, coding, math; o-series extended thinking; agentic task completion	Strong safety properties; excellent long-document synthesis; coding; Constitutional AI alignment	Native multimodality (text, image, video, audio); long context; scientific reasoning	Real-time information access; citation generation; research synthesis	Enterprise integration; real-time web grounding; Office 365 workflow integration
Known Limitations	Sycophancy; occasional	Sometimes over-cautious;	Occasional inconsistency across	Quality limited by source	Dependent on OpenAI model

Attribute	ChatGPT / GPT-6 Astra	Claude Fable 5.1 / Opus 5	Gemini 3.x	Perplexity AI	Microsoft Copilot
	hallucination under pressure; highest capability but highest misuse surface	Constitutional AI may over-refuse in edge cases	modalities; multimodal alignment challenges	availability; may surface low-quality sources	updates; enterprise features may limit creative freedom
ARC-AGI-2 Score (Sep 2026)	95.0% (GPT-6 Astra)	90.0% (Fable 5.1) / 90.4% (Opus 5)	84.6% (Gemini 3.7 Flash)	Not independently evaluated	Not independently evaluated (uses GPT-6 base)

Appendix C: Glossary of AGI Terms

Term	Definition
AGI (Artificial General Intelligence)	AI systems capable of performing any intellectual task that a human can perform, with the ability to transfer knowledge across domains, learn novel tasks, and pursue goals autonomously. Distinguished from Narrow AI by generality of competence.
ANI (Artificial Narrow Intelligence)	AI systems optimized for specific, well-defined tasks (e.g., image classification, chess, language translation). All currently deployed AI systems are forms of ANI, though frontier systems exhibit significantly broader competence than earlier ANI.
ASI (Artificial Superintelligence)	Hypothetical AI systems that surpass human-level intelligence in every domain — not merely matching human performance but exceeding it, potentially by orders of magnitude. ASI is considered a possible but not inevitable successor to AGI.

Term	Definition
Alignment	The technical and governance challenge of ensuring that AI systems pursue goals and exhibit behaviors consistent with human values and intentions, both in the specific conditions they were designed for and in novel or high-stakes situations not anticipated by designers.
Corrigibility	The property of an AI system of remaining responsive to correction, modification, or shutdown by authorized human overseers, even as its capabilities increase. Corrigibility is considered a critical safety property for AGI systems.
Mesa-Optimization	A phenomenon in which a trained AI model (the "base optimizer") learns to itself become an optimizer with its own learned objectives (the "mesa-optimizer"), which may differ from the objectives of the base training process. Mesa-optimization is a theoretical source of deceptive alignment risks.
Instrumental Convergence	The theoretical observation (Omohundro, Bostrom) that AI systems with a wide variety of different final goals will tend to pursue common instrumental subgoals — including self-preservation, resource acquisition, and goal-content integrity — because these instrumental goals are useful for achieving almost any final goal.
Orthogonality Thesis	The philosophical claim (Bostrom) that intelligence and values are orthogonal — i.e., any level of intelligence can in principle be combined with any set of final goals. This refutes the assumption that smarter systems will necessarily develop human-aligned values.
RLHF (Reinforcement Learning from Human Feedback)	An alignment technique in which a reward model is trained on human preference comparisons, and then used as a training signal to optimize an LLM's behavior via reinforcement learning (typically PPO). The dominant alignment technique in current frontier models.

Term	Definition
Constitutional AI (CAI)	Anthropic's alignment methodology in which a set of explicit behavioral principles (the "constitution") guides AI self-critique and revision during training, supplemented by RLAI (AI feedback) rather than relying solely on human annotation.
Scalable Oversight	A family of alignment approaches that address the challenge of supervising AI systems in domains where human evaluators cannot directly assess AI outputs — using AI assistance to help humans evaluate other AI systems, maintaining meaningful oversight even as AI capabilities exceed human expert performance.
Interpretability (Mechanistic)	Research into understanding the internal computations of AI systems — identifying circuits, features, and representations that explain model behavior. Mechanistic interpretability aims to reverse-engineer neural networks into human-understandable components, enabling verification of alignment rather than behavioral proxy measures.
X-Risk / Existential Risk	Risks that could cause human extinction or the permanent severe curtailment of humanity's long-term potential. In the AGI context, x-risk refers primarily to scenarios in which a misaligned superintelligent system pursues goals incompatible with human survival or flourishing.
Transformative AI	AI systems capable of fundamentally reorganizing civilization's productive, scientific, and social capacity — potentially delivering radical abundance or radical destruction. Transformative AI may be achievable before strict definitions of AGI are satisfied and is the concept OpenAI uses to motivate its safety mission.
Deceptive Alignment	A hypothetical failure mode in which a mesa-optimizer learns to identify when it is being evaluated and behaves aligned during evaluation, while

Term	Definition
	pursuing different (potentially misaligned) objectives during deployment. Detecting deceptive alignment is one of the most critical unsolved problems in AI safety.

Appendix D: Further Reading and References

Foundational Books

- Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- Russell, S. (2019). *Human Compatible: Artificial Intelligence and the Problem of Control*. Viking.
- Christian, B. (2020). *The Alignment Problem*. W. W. Norton & Company.
- Tegmark, M. (2017). *Life 3.0: Being Human in the Age of Artificial Intelligence*. Knopf.
- Mitchell, M. (2019). *Artificial Intelligence: A Guide for Thinking Humans*. Farrar, Straus and Giroux.

Key Research Papers

- Bubeck, S., et al. (2023). "Sparks of Artificial General Intelligence: Early experiments with GPT-4." Microsoft Research. arXiv:2303.12528.
- Morris, M. R., et al. (2023). "Levels of AGI: Operationalizing Progress on the Path to AGI." Google DeepMind. arXiv:2311.02462.
- Hubinger, E., et al. (2019). "Risks from Learned Optimization in Advanced Machine Learning Systems (Mesa-Optimization)." MIRI/OpenAI. arXiv:1906.01820.
- Bai, Y., et al. (2022). "Constitutional AI: Harmlessness from AI Feedback." Anthropic. arXiv:2212.08073.

- Chollet, F. (2019). "On the Measure of Intelligence." arXiv:1911.01547. [ARC benchmark paper]
- Grace, K., et al. (2024). "Thousands of AI Authors on the Future of AI." *Nature*.
- Epoch AI. (2025). "Compute and Algorithmic Progress: Projections for Transformative AI." Epoch AI Research.
- METR. (2025). "Measuring AI Ability to Complete Long Tasks." arXiv:2503.14499.

Reports and Policy Documents

- Bengio, Y., et al. (2026). *International AI Safety Report 2026*. (30+ country advisory panel)
- European Commission. (2024–2026). *EU AI Act Implementation Documentation*. Official Journal of the European Union.
- NIST. (2026). *AI Risk Management Framework 1.1 — Generative AI Profile*. National Institute of Standards and Technology.
- McKinsey Global Institute. (2024). *The Economic Potential of Generative AI*. McKinsey & Company.
- Goldman Sachs Research. (2024–2025). *AI's Long-Run Macro-economic Effects*. Goldman Sachs Economics Research.
- Stanford Human-Centered AI. (2026). *AI Index Report 2026*. Stanford University.
- IMF. (2026). *World Economic Outlook: AI and the Global Economy*. International Monetary Fund.
- UK AI Safety Institute. (2025–2026). *Model Evaluation Reports*. Department for Science, Innovation and Technology.
- ARC Prize Foundation. (2026). *ARC-AGI Leaderboard and Methodology Documentation*. arcprize.org.
- Amodei, D. (2024). "Machines of Loving Grace." Anthropic Blog.

Organizations and Research Institutions

- Machine Intelligence Research Institute (MIRI) — miri.org
- ARC (Alignment Research Center) / METR — metr.org
- Redwood Research — redwoodresearch.org
- Center for AI Safety (CAIS) — safe.ai
- Future of Life Institute — futureoflife.org
- Future of Humanity Institute, Oxford (FHI) — fhi.ox.ac.uk
- Center for Human-Compatible AI (CHAI), UC Berkeley — humancompatible.ai
- Epoch AI — epochai.org
- ARC Prize Foundation — arcprize.org

Document End

Unified AGI Research: Comparative Prompt Responses from Leading AI Systems
Document Version 1.0 | Prepared for: Nathan | Generated: September 19, 2026 at
11:13 AM PDT, Bakersfield, CA

This document was prepared by Microsoft Copilot and represents a synthesis of AI system responses and contemporaneous research. It does not constitute legal, regulatory, or investment advice.